



POWER8 Processor User's Manual for the Single-Chip Module

OpenPOWER

16 March 2016
Version 1.3



© Copyright International Business Machines Corporation 2014, 2015, 2016

Printed in the United States of America March 2016

IBM, the IBM logo, and ibm.com are trademarks or registered trademarks of International Business Machines Corp., registered in many jurisdictions worldwide. Other product and service names might be trademarks of IBM or other companies. A current list of IBM trademarks is available on the Web at "Copyright and trademark information" at www.ibm.com/legal/copytrade.shtml.

The OpenPOWER word mark and the OpenPOWER Logo mark, and related marks, are trademarks and service marks licensed by OpenPOWER.

Intel, Intel logo, Intel Inside, Intel Inside logo, Intel Centrino, Intel Centrino logo, Celeron, Intel Xeon, Intel SpeedStep, Itanium, and Pentium are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States and other countries.

Other company, product, and service names may be trademarks or service marks of others.

All information contained in this document is subject to change without notice. The products described in this document are NOT intended for use in applications such as implantation, life support, or other hazardous uses where malfunction could result in death, bodily injury, or catastrophic property damage. The information contained in this document does not affect or change IBM product specifications or warranties. Nothing in this document shall operate as an express or implied indemnity under the intellectual property rights of IBM or third parties. All information contained in this document was obtained in specific environments, and is presented as an illustration. The results obtained in other operating environments may vary.

This document is intended for development of technology products compatible with Power Architecture®. You may use this document, for any purpose (commercial or personal) and make modifications and distribute; however, modifications to this document may violate Power Architecture and should be carefully considered. Any distribution of this document or its derivative works shall include this Notice page including but not limited to the IBM warranty disclaimer and IBM liability limitation. No other licenses (including patent licenses), expressed or implied, by estoppel or otherwise to any intellectual property rights are granted by this document.

THE INFORMATION CONTAINED IN THIS DOCUMENT IS PROVIDED ON AN "AS IS" BASIS. IBM makes no representations or warranties, either express or implied, including but not limited to, warranties of merchantability, fitness for a particular purpose, or non-infringement, or that any practice or implementation of the IBM documentation will not infringe any third party patents, copyrights, trade secrets, or other rights. In no event will IBM be liable for damages arising directly or indirectly from any use of the information contained in this document.

IBM Systems
294 Route 100, Building SOM4
Somers, NY 10589-3216

The IBM home page can be found at ibm.com®.

Version 1.3
16 March 2016



Contents

List of Tables	13
List of Figures	17
Revision Log	19
About this Document	21
Who Should Read this Document	21
Conventions Used in This Document	21
Representation of Numbers	21
Bit Significance	21
Other Conventions	21
Related Documents	22
1. POWER8 Processor Overview	23
1.1 General Features	23
2. POWER8 Processor Core	27
2.1 Key Design Fundamentals	27
2.1.1 64-Bit Implementation of the Power ISA (version 2.07)	27
2.1.3 Speculative Superscalar Inner Core Organization	28
2.1.4 Specific Focus on Storage Latency Management	29
2.2 Pipeline Structure	29
2.3 Microprocessor Core - Detailed Features	31
2.3.1 Instruction Fetching and Branch Prediction	31
2.3.3 Instruction Dispatch, Sequencing, and Completion Control	33
2.3.4 Fixed-Point Execution Pipelines	35
2.3.5 Load and Store Execution Pipelines	35
2.3.6 Branch and Condition Register Execution Pipelines	37
2.3.7 Unified Second-Level Memory Management (Address Translation)	37
2.3.8 Data Prefetch	38
2.3.9 VSU Execution Pipeline	38
2.3.10 DFP Execution Pipeline	39
3. Power Architecture Compliance	41
3.1 Book I - User Instruction Set Architecture	41
3.1.1 Defined Instructions	41
3.1.1.1 Illegal Instructions	41
3.1.1.2 Instructions Supported	41
3.1.1.3 Invalid Forms	41
3.1.2 Branch Processor	42
3.1.2.1 Instruction Fetching	42
3.1.2.2 Branch Prediction	42
3.1.2.3 Instruction Cache Block Touch Hint	42
3.1.2.4 Out-of-Order Execution and Instruction Flushes	42
3.1.2.5 Branch Processor Instructions with Undefined Results	43



3.1.3 Fixed-Point Processor	43
3.1.3.1 Fixed-Point Exception Register (XER)	43
3.1.4 Storage Access Alignment Support Overview	44
3.1.4.1 Misaligned Flushes	44
3.1.4.2 Alignment Interrupts	45
3.1.4.3 Fixed-Point Load Instructions	63
3.1.4.4 Fixed-Point Store Instructions	63
3.1.4.5 Fixed-Point Load and Store Multiple Instructions	63
3.1.4.6 Fixed-Point Move Assist Instructions	64
3.1.4.7 Integer Select (SEL)	65
3.1.4.8 Fixed-Point Logical Instructions	65
3.1.4.9 Move To/From Special Purpose Register (SPR) Instructions	65
3.1.4.10 Move to Condition Register Fields Instruction	65
3.1.4.11 Fixed-Point Invalid Forms and Undefined Conditions	66
3.2 Floating-Point Processor (FP, VMX, and VSX)	67
3.2.1 Vector Single-Precision Bandwidth	67
3.2.2 IEEE Compliance	68
3.2.3 Divide and Square-Root Latencies	68
3.2.4 Early Forwarding Conditions	68
3.2.5 Floating-Point Exceptions	69
3.2.6 Floating-Point Load and Store Instructions	69
3.2.7.1 NaN Propagation	70
3.2.7.2 Square Root Overflow and Underflow	70
3.2.7.3 Hardware Behavior on Enabled Underflow and Enabled Overflow Exception	70
3.2.8 Handling of Denormal Single-Precision Values in Double-Precision Format	71
3.2.8.1 Producing Single-Precision Denorms	71
3.2.8.2 Consuming Single-Precision Denorms	71
3.2.9 Floating-Point Invalid Forms and Undefined Conditions	72
3.3 Optional Facilities and Instructions	73
3.4 Little Endian	73
3.5 Book II - Virtual Environment Architecture	73
3.5.1 Classes of Instructions	73
3.5.1.1 Optional Instructions	73
3.5.3 Data Prefetch	74
3.5.5 Storage Model	75
3.5.5.1 Atomicity	75
3.5.5.2 Vector Element Atomicity	75
3.5.5.3 Transactional Memory	75
3.5.5.4 Storage Access Ordering	76
3.5.6 Atomic Updates and Reservations	76
3.5.7 Storage Control Instructions	76
3.5.7.1 Overview of Key Aspects of Storage Control Instructions	76
3.5.7.2 Instruction Cache Block Invalidate (icbi)	77
3.5.7.3 Instruction Cache Synchronize (isync)	77
3.5.7.4 Data Cache Block Touch (dcbt and dcbtst)	77
3.5.7.5 Data Cache Block Touch - No Access Needed Anymore (TH = '10001')	78
3.5.7.6 Data Cache Block Touch - Transient (TH = '10000')	78
3.5.7.7 Data Cache Block Zero (dcbz)	78
3.5.7.8 Data Cache Block Store (dcbst)	78
3.5.7.9 Data Cache Block Flush (dcbf , dcbfl and dcbflp)	79



3.5.7.10 Load and Reserve and Store Conditional Instructions	79
3.5.7.11 sync Instruction	79
3.5.7.12 eieio Instruction	79
3.5.7.13 miso Instruction	79
3.5.7.14 Transactional Memory Instructions	80
3.5.8.1 Time Base	81
3.5.9 Hypervisor Decrementer (HDEC)	82
3.5.10 Decrementer (DEC)	82
3.5.11 Book II Invalid Forms	82
3.6 Book III - Operating Environment Architecture	83
3.6.1 Classes of Instructions	83
3.6.1.1 Storage Control Instructions	83
3.6.1.2 Reserved Instructions	84
3.6.2 Branch Processor	84
3.6.2.1 SRR1 Register	84
3.6.2.2 MSR Register	84
3.6.2.3 Branch Processor Instructions	84
3.6.2.4 Current Instruction Address Breakpoint (CIABR)	85
3.6.2.5 Instruction Effective to Real Address Translation Cache (I-ERAT)	85
3.6.3 Fixed-Point Processor	86
3.6.3.1 Processor Version Register (PVR)	86
3.6.3.2 Processor ID Register (PIR)	87
3.6.3.3 Chip Information Register (CIR)	87
3.6.3.4 Move To/From Special Purpose Register Instructions	87
3.7 HID Registers (HID0, HID1, HID4, and HID5)	93
3.7.2 HID1 Register	95
3.7.5 Real Mode Offset (RMO) Region Sizes	99
3.7.7 Core-to-Core Trace SPR	100
3.7.8 Trigger Registers	100
3.7.9 IMC Array Access Register	100
3.7.10 Performance Monitor Registers	100
3.7.11 Other Fixed-Point Instructions	101
3.8 Storage Control	101
3.8.1 Virtual and Physical Address Ranges Supported	101
3.8.2 Data Effective-to-Real-Address Translation (D-ERAT)	101
3.8.3 Translation Lookaside Buffer (TLB)	103
3.8.4 Large-Page Support	105
3.8.6 Segment Lookaside Buffer (SLB)	106
3.8.7 Address Space Register	106
3.8.9 Reference and Change Bits	107
3.8.10 Storage Protection	107
3.8.11 Block Address Translation	107
3.8.13 Storage Access Modes - WIMG Bits	108
3.8.14 Speculative Storage Accesses	109
3.8.15 mtsr , mtsrin , mfsr , and mfsrin Instructions	109
3.8.16 TLB Invalidate Entry (tlbie and tlbiel) Instructions	109
3.8.17 TLB Invalidate All (tlbia) Instruction	111
3.8.18 TLB Synchronize (tlbsync) Instruction	111
3.8.19 Page Replacement Policy	111
3.8.20 Support for Store Gathering	112



3.8.21 Cache Coherency Paradoxes	112
3.8.22 Handling Parity Error, Multi-Hit, and Uncorrectable Errors	113
3.8.22.1 Parity Error	113
3.8.22.2 Multi-Hit	113
3.8.22.3 Both Multi-Hit and Parity Error	113
3.8.22.4 Uncorrectable Error Handling	114
3.8.22.5 TLB Parity Error and Multi-Hit Action	114
3.8.23.1 Interrupt Vectors	115
3.8.23.2 Interrupt Definitions	116
3.8.23.3 System Reset Interrupt	118
3.8.23.4 Machine Check Interrupt	119
3.8.23.5 Hypervisor Maintenance Interrupt	122
3.8.23.6 External Interrupt	122
3.8.23.7 Alignment Interrupt	123
3.8.23.8 Trace Interrupt	124
3.8.23.9 Performance Monitor Interrupt	124
3.8.23.10 SPMC Performance Monitor Interrupt	124
3.8.23.11 Facility Unavailable Interrupt	125
3.8.24 Logical Partitioning Support	125
3.8.24.1 Thread-to-LPAR Mapping	125
3.8.24.2 Dynamic LPAR Switching	125
3.8.25 Strong Access Ordering Mode (SAO)	125
3.8.26 Graphics Data Stream Support	125
3.8.27 Performance Monitoring, Sampling, and Trace	126
3.8.28 Processor Compatibility Mode	126
4. Storage Subsystem	127
4.1 L1 Cache	127
4.2 L2 Cache	127
4.2.1 L2 Cache Features	127
4.3 L3 Cache	128
4.3.1 L3 Features, Queues and Resources	128
4.4 NCU	130
4.4.1 NCU Characteristics	130
4.5 Memory Controller	131
4.6 POWER8 Memory Stack Partitioning	131
4.7 POWER8 Chip Memory Controller Unit Features	132
4.7.1 Bandwidth	133
4.7.2 POWER8 Memory Controller Characteristics	133
5. Simultaneous Multithreading	135
5.1 Overview	135
5.2 Partitioning of Resources in Different SMT Modes	135
5.3 Program Priority Register (PPR)	136
5.4 Thread Priority NOPs	137
5.5 Control Register	138
5.6 Thread Priority, Status, and Control Requirements	140
5.7 Thread Balance Control Requirements	140



5.8 Thread Switch Control Register (Hypervisor Access Only)	141
5.9 Thread Time-Out Register (Hypervisor only)	143
5.10 Forward Progress Timer	144
5.11 Thread Priority Boosting	144
5.12 Priority Boosting to Medium-High in User Mode	144
5.13 Thread Priority Boosting on Asynchronous Interrupt	145
5.13.1 When to Boost Thread Priority	145
5.14 Thread Prioritization Implementation	146
5.14.1 Thread Switch Fetch Priority	146
5.14.1.1 SMT2 Fetch Pattern	147
5.14.2 Thread Switch Decode Priority	147
5.14.3 Software-Set Thread Priority	149
5.14.4 Low-Power Modes for Application	149
5.14.5 Dynamic Thread Priority	149
5.15 Support for Multiple LPARs	150
5.15.1 Instruction Fetch	150
5.15.2 Decode	150
5.15.3 Microcode Fairness	151
5.15.4 Instruction Cache	151
5.15.5 Thread Set Allocation	151
5.15.6 Data Cache	151
5.15.7 ERATs	152
5.16 Controlling the Flow of Instructions in SMT	152
5.16.1 Dispatch Flush	152
5.16.1.1 Dispatch Flush Rules	152
5.16.1.2 Stall at Dispatch	153
5.16.2 Decode Hold	153
5.16.2.1 Balance Flush	153
5.17 Dynamic Thread Transitioning	154
5.17.1 Overview	154
5.17.2 Thread Set Definition	154
5.17.4 Thread Set Reconfiguration	155
5.17.4.1 Balancing	155
5.17.4.2 Mixing	156
5.17.4.3 Action	156
6. POWER8 SMP Interconnect	157
6.1 POWER8 SMP Interconnect Features	157
6.1.1 General Features	157
6.1.2 POWER8 Specific Features	158
6.1.3 Off-Chip Features	158
6.1.4 Power Management Features	158
6.1.5 RAS Features	158
6.2 External POWER8 SMP Interconnect	159
6.2.1 POWER8 SMP Interconnect Multichip Configurations	159
6.2.2 Protocol and Data Routing in Multichip Configurations	159
6.3 Coherency Flow	160
6.3.1 Physical Broadcast Flow	160
6.3.2 Broadcast Scope Definition	160



6.4 Command Ordering Support	160
6.5 Memory Coherence Directory	160
6.5.1 Directory Size	160
6.5.2 Operation	160
7. Interrupt Control Presenter	163
7.1 Features	165
7.1.1 Routing Layer	165
7.1.2 Presentation Layer	166
7.2 Interrupt Control Presenter Registers	168
7.2.1 ICP Address Map	168
7.2.2 Interrupt Base Address Register (ICPBAR)	168
8. PCI Express Controller	175
8.1 Specification Compliance	175
8.2 PEC Feature Summary	175
9. Power Management	177
9.1 Overview	177
9.2 Power Management Infrastructure	177
9.3 Power Management Policies and Modes of Operation	178
9.3.1 Maximum Power Savings Based on Utilization and Idle	178
9.3.2 Adaptive Power Savings with Performance Loss Floor	178
9.3.3 Power Cap	178
9.3.4 Turbo Performance Boost	178
9.4 Feature Summary	179
9.5 Overview of Chip Hardware Power-Management Features	179
9.5.1 Communication Paths for System Controllers	179
9.5.2 Sensors	179
9.5.3 Accelerators	180
9.5.4 Actuators/Controls	180
9.5.4.1 Configurations with Unused Components	181
9.6 Chip Hardware Power-Management Features	182
9.6.1 Chiplet Voltage Control	182
9.6.2 Chip-Level Voltage Control Sequencing	182
9.6.2.1 SPIVID VRM Control Sequencing	182
9.7 Functional Description of Processor Core Chiplet	182
9.7.1 Power Gating	182
9.7.2 Idle States	183
9.7.2.1 Core and Thread Doze	185
9.7.2.2 Single Thread Nap, Sleep, and Winkle	185
9.7.2.3 Sleep	186
9.7.2.4 Nap	187
9.7.2.5 Winkle	187
9.7.3 Special Wake-up	188
9.7.4 Pstates	188
9.7.4.1 Architectural Overview	188
9.7.4.2 Definitions	189



9.7.4.3 Permissible Behavior	191
9.7.4.4 Interaction with Idle Modes	191
9.8 Architected Control Facilities	192
9.8.1 Power Management Control Register (PMCR)	192
10. Performance Profile	199
10.1 Core	199
10.1.1 Level-1 Instruction Cache	199
10.1.2 Level-1 Instruction ERAT	200
10.1.3 Instruction Prefetch	200
10.1.4.1 Branch Direction Prediction using the Branch History Tables	201
10.1.4.2 Branch Prediction using Static Prediction and "a", "t" Bits	202
10.1.4.3 Address Prediction Using the Link Stack	203
10.1.4.4 Address Prediction using the Count Cache	204
10.1.4.5 Round-Trip Branch Processing	205
10.1.4.6 BC+4 Handling	206
10.1.4.7 BC+8 Handling	206
10.1.5 Store-Hit-Load Avoidance Table	207
10.1.6 Instruction Buffer	207
10.1.7 Group Formation	208
10.1.7.1 General Rules	208
10.1.7.2 Rules Specific to ST Mode	209
10.1.7.3 Rules Specific to SMT Modes	209
10.1.8 Group Ending NOP	209
10.1.9 First and Last Instructions	209
10.1.10 2-Way and 3-Way Cracked Instructions	213
10.1.11 Microcode	214
10.1.14 Instruction Issue	216
10.1.14.1 Steering Policy	217
10.1.14.2 BRQ and CRQ Operation	217
10.1.14.3 UniQ Issue Policies	217
10.1.14.4 FXU and VSU Selection	217
10.1.14.5 LU Selection	217
10.1.14.6 LSU Selection	218
10.1.14.7 Dispatch Bypass Instruction Selection	218
10.1.14.8 Back-to-Back Issue Policy	218
10.1.14.9 Limitations of Back-to-Back	219
10.1.14.10 Dual-Issued Stores	220
10.1.14.11 Wake-up Misspeculations	220
10.1.14.12 Chains of Misspeculations	220
10.1.14.13 Other Issue Inefficiencies	220
10.1.14.14 Issue-to-Issue Latencies	221
10.1.15.1 ISU Rejects	223
10.1.15.2 LSU Rejects	226
10.1.15.3 Flush Conditions	226
10.1.16.1 Storage Alignment	228
10.1.16.2 Special Case of Store Crossing a 64-Byte Boundary	228
10.1.17 Level-1 Data ERAT	228
10.1.18 Level-2 Data ERAT	229
10.1.19 Translation Look-Aside Buffer	230



10.1.20 Load Miss Queue	230
10.1.21 Transactional Memory	231
10.1.22 Store Queue and Store Forwarding	231
10.1.22.1 Stores in Real Mode (MSR[DR] = 0)	232
10.1.23 Data Prefetch	232
10.2 Chiplet	234
10.2.1 Level-2 Cache	234
10.2.2 Level-3 Cache	234
10.3 Latencies	235
10.3.1 Cache Latencies and Bandwidth	235
10.4 PCI Express Performance	253
10.4.1 Bandwidth	253
10.4.2 Latency	253
10.4.3 Cluster Latency 2K Message	253
10.4.4 I/O Bandwidth	253
10.4.5 PCIe Performance Goals	253
10.5 Performance Specific Instructions	254
10.5.1 Store Multiple and Store String	254
10.5.1.1 Store Quadword	254
10.5.1.2 eieio	254
10.6 Other Topics	255
10.6.1 Hot/Cold Page Affinity Support	255

11. Performance Monitor 269

11.1 Performance Monitor Overview	269
11.2 Performance Monitor Functions	270
11.2.1 Performance Monitor Event Selection	270
11.2.2 Machine States and Enabling the Performance Monitor Counters	270
11.2.3 Trigger Events and Enabling the Performance Monitor Counters	270
11.2.4 Performance Monitor Exceptions, Alerts, and Interrupts	270
11.2.5 Sampling	271
11.2.6 Thresholding	271
11.2.7 Trace Support Facilities	271
11.3 Special Purpose Registers and Fields Associated with Instrumentation	271
11.4 Enhanced Sampling Support	272
11.5 POWER8 Performance Monitor Event Selection	273
11.5.1 Event Bus Events and Event Bus Ramp	273
11.5.2 Direct Events	273
11.6 Performance Monitor Facility	274
11.6.1 Performance Monitor Facility Registers	274
11.6.1.1 Performance Monitor Counters (PMC1 - 6)	274
11.6.1.2 Monitor Control Register 0 (MMCR0)	275
11.6.1.3 Monitor Mode Control Register 1 (MMCR1)	280
11.6.1.4 Monitor Mode Control Register 2 (MMCR2)	283
11.6.1.5 Monitor Mode Control Register A (MMCRA)	284
11.6.1.6 Core Monitor Mode Control Register (MMCRC)	288
11.7 Hypervisor Performance Monitor	289
11.7.2 Monitor Mode Control Register H (MMCRH)	290



11.8 Supervisor Performance Monitor	292
11.8.1 Supervisor Performance Monitor Counters (SPMC1 - 2)	292
11.8.2 Monitor Mode Control Register S (MMCRS) Register	292
11.9 Sampled Instruction Event Register (SIER)	294
11.10 POWER8 CPI Stack	300
11.10.1 Completion Stall Accounting: LSU Related Stalls	303
11.10.1.1 Completion Stall Accounting: Data Cache Misses	303
11.10.1.2 Completion Stalls: Data Cache Miss that Resolves in a Local Core's L2 or L3 Cache	303
11.10.1.3 Completion Stalls: Data Cache Miss that Resolves in a Local Chip's L2 or L3 Cache	303
11.10.1.4 Completion Stalls: Data Cache Miss that Resolves from Remote Chip's Cache or Memory	303
11.10.1.5 Completion Stalls: Data Cache Miss that Resolves from Local Core's L2 or L3 (Dispatch Conflict)	304
11.10.1.6 Completion Stalls: Data Cache Miss that Resolves from Local Memory	304
11.10.1.7 Completion Stalls: Stores	304
11.10.1.8 Completion Stalls: Store Forwarding	304
11.10.1.9 Completion Stalls: LSU Rejects	304
11.10.1.10 Completion Stalls: LSU Rejects Due to ERAT Miss	304
11.10.1.11 Completion Stalls: LSU Rejects Due to LMQ Full	304
11.10.1.12 Completion Stalls: LSU Rejects Due to Load-Hit-Store Reject	304
11.10.2 Completion Stalls: FXU	304
11.10.3 Completion Stalls: VSU	304
11.10.4 Completion Stalls: IFU	305
11.10.5 Front-End Stalls	305
11.10.5.1 GCT Empty: I-Cache Miss	305
11.10.5.2 GCT Empty: I-Cache Miss That Also Missed the Local L3 Cache	305
11.10.5.3 GCT Empty: Branch Redirects	305
11.10.5.4 GCT Empty: Branch Redirects and I-Cache Miss	305
11.10.5.5 GCT Empty: Dispatch Hold Conditions	306
11.11 Exploiting Advanced Features of the PMU	306
11.11.1 Correlating Fabric Responses to Effective Addresses	306
11.11.1.1 Operation	306
11.11.1.2 Tools Exploitation	307
11.11.2.1 PMU Usage	309
11.11.3 Per LPAR Memory Bandwidth	310
11.11.4 Monitoring Fabric Command Scope at a Thread Level	310
11.12 PMC Events	311
11.13 SPMC Events	311
Appendix A. POWER8 Instruction Summary by Category	313
Appendix B. POWER8 Instruction Summary by Mnemonic	339
Appendix C. POWER8 Instruction Summary by Opcode	363
Appendix D. Performance Monitoring Events	387



Appendix E. SPMC Performance Monitoring Events	435
Glossary	443



List of Tables

Table 3-1.	XER Bits and Fields	43
Table 3-2.	Operand Alignment Effects on Performance (Non-Watchpoint Mode)	46
Table 3-3.	Operand Alignment Effects on Performance (Watchpoint Mode)	54
Table 3-4.	Latencies of Floating-Point Divide/Square-Root Instructions	68
Table 3-5.	Storage Control Instructions	76
Table 3-6.	Cache and TLB Management Instruction Effects on Transactional Accesses	80
Table 3-7.	I-ERAT I and G Bit Setting	86
Table 3-8.	SPR Table	87
Table 3-9.	HRMOR Update Sequence	100
Table 3-10.	D-ERAT I and G Bit Setting	102
Table 3-11.	256 MB Segments	104
Table 3-12.	1 TB Segments	104
Table 3-13.	PTE and SLBE Correspondence	105
Table 3-14.	WIMG Bits	108
Table 3-15.	IG Bits	108
Table 3-16.	Segment Size and Page Size Specifications for tlbie and tlbiel (L = 0)	110
Table 3-17.	Segment Size and Page Size Specifications for tlbie and tlbiel (L = 1)	110
Table 3-18.	Reference-Bit Array Update	112
Table 3-19.	Summary of POWER8 Behavior on Parity Error, Multi-Hit, and Uncorrectable Error	114
Table 3-20.	Interrupt Vector	115
Table 3-21.	Interrupt Vectors for the POWER8 Core	116
Table 3-22.	Implementation-Specific Interrupt Types	116
Table 3-23.	Implementation MSR and SRR1/HSRR1 Bits	116
Table 3-24.	System Reset Interrupt	118
Table 3-25.	Synchronous Machine Checks	120
Table 3-26.	Direct External Interrupt (LPES = '0')	122
Table 3-27.	Direct External Interrupt (LPES = '1')	122
Table 3-28.	Mediated External Interrupt (LPES = '0')	123
Table 3-29.	Mediated External Interrupt (LPES = '1')	123
Table 3-30.	Alignment Interrupt	123
Table 3-31.	Trace Interrupt	124
Table 4-1.	Maximum Frequencies	133
Table 4-2.	Minimum Frequencies	133
Table 4-3.	Bandwidth (per MCS/POWER8 Memory Buffer)	133
Table 4-4.	POWER8 Memory Controller Characteristics	133
Table 5-1.	SMT Modes	135
Table 5-2.	Front-End Execution Core Resource	135
Table 5-3.	Thread Priority Nops	137



Table 5-4.	mfspr CTRL Data Formatting	139
Table 5-5.	Thread Balance Control (Balance Flush)	140
Table 5-6.	Asynchronous Interrupt	145
Table 5-7.	Scoring System Summary	150
Table 5-8.	SMT Mode Boundary Crossing Reconfigurations	155
Table 5-9.	Thread Balancing Scenarios	155
Table 6-1.	POWER8 Broadcast Scope Definition	160
Table 7-1.	Interrupt Management Area: Interrupt Presentation Layer Ports	166
Table 7-2.	Facility Definitions	166
Table 7-3.	Resetting the Interrupt Condition	167
Table 7-4.	Interrupt Presenter Register Address Map	168
Table 8-1.	Supported I/O Configurations	176
Table 9-1.	Supported Chiplet Power Management Modes	184
Table 9-2.	Power Management Control Register (PMCR) - SPR 884	192
Table 9-3.	Power Management Idle Control Register (PMICR)- SPR 852	194
Table 9-4.	Power Management Status Register (PMSR) - SPR 853	196
Table 9-5.	Power Management Memory Activity Register (PMMAR) - SPR 854	197
Table 10-1.	Handling of bclr and bclrl Instructions	204
Table 10-2.	Handling of bcctr and bcctrl Instructions	205
Table 10-3.	bc+8 Pairable Instructions	206
Table 10-4.	Stores Ineligible for SHL Avoidance	207
Table 10-5.	IBuffer Rows per Thread	207
Table 10-6.	List of Instructions Marked as First	210
Table 10-7.	List of Instructions Marked as Last	211
Table 10-8.	2-Way Cracked Instructions	213
Table 10-9.	3-Way Cracked Instructions	214
Table 10-10.	Resource Requirements for Dispatch	216
Table 10-11.	Example where Back-to-Back is not Possible A->C, B->C	219
Table 10-12.	A -> B -> C	219
Table 10-13.	Issue-to-Issue Latencies	221
Table 10-14.	Flush Conditions	227
Table 10-15.	Cache Latencies and Bandwidth	235
Table 10-16.	Instruction Latencies and Throughputs	236
Table 10-17.	Instructions that Trigger an Exception on Consumption of a Scalar SP Denorm	258
Table 10-18.	Soft Patch Instruction on Unaligned Stores	268
Table 11-1.	Performance Monitor Counter Register	275
Table 11-2.	MMCR0 Register	276
Table 11-3.	MMCR1 Register	280
Table 11-4.	MMCR1 PMCxSEL Selection of Direct Events versus Event-Bus Events	282



Table 11-5. MMCR2 Register	283
Table 11-6. MMCRA Register	284
Table 11-7. Threshold Start/Stop Event Selection	285
Table 11-8. Random Sampling Eligibility Criteria	286
Table 11-9. MMCRC Register	288
Table 11-10. Hypervisor Performance Monitor Counter Register	290
Table 11-11. Monitor Mode Control Register H Register	290
Table 11-12. Supervisor Performance Monitor Counter	292
Table 11-13. Monitor Mode Control Register S Register	292
Table 11-14. Sampled Instruction Event Register (SIER)	295
Table 11-15. Implementation-Dependent Extension to Data Source Encodes	297
Table 11-16. Implementation-Dependent Extension Bits for Data Source Encodes (SIER[EXT])	299
Table 11-17. POWER8 Accounting	300
Table 11-18. CPI Stack	302
Table A-1. Category Listing	313
Table A-2. POWER8 Instructions by Category	314
Table B-1. POWER8 Instructions by Mnemonic	339
Table C-1. POWER8 Instructions by Opcode	363
Table D-1. POWER8 Event List by Event Name	388
Table E-1. SPMC Performance Monitoring Events	435





List of Figures

Figure 1-1.	Block Diagram for the POWER8 Processor	23
Figure 2-1.	POWER8 Processor Core	27
Figure 2-2.	Pipeline Structure	30
Figure 4-1.	Memory Stack Partitioning	131
Figure 5-1.	Dual SMT4 Decode Priorities	148
Figure 5-2.	Decode Priority in 4 LPAR Mode	151
Figure 6-1.	Two Socket Configuration (24-way)	159
Figure 6-2.	Four Socket Configuration (48-way)	159
Figure 7-1.	POWER8 Logical Interrupt Controller Structure	164
Figure 9-1.	Idle Mode Summary	183
Figure 9-2.	Sleep and Winkle Power Gating Progression	185
Figure 10-1.	Basic Building Blocks	255
Figure 10-2.	HCA Cache	257
Figure 11-1.	POWER8 CPI Stack Example	301





Revision Log

Each release of this document supersedes all previously released versions. The revision log lists all significant changes made to the document since its initial release. In the rest of the document, change bars in the margin indicate that the adjacent text was modified from the previous release of this document.

Revision Date	Description
16 March 2016	Version 1.3. <ul style="list-style-type: none">• Revised <i>Section 3.8.24 Logical Partitioning Support</i> on page 125.• Added <i>Section 11 Performance Monitor</i> on page 269.• Added <i>Appendix D Performance Monitoring Events</i> on page 387.• Added <i>Appendix E SPMC Performance Monitoring Events</i> on page 435.• Revised <i>Glossary</i> on page 443.
17 June 2015	Version 1.2. <ul style="list-style-type: none">• Revised <i>Section 3.1.4.5 Fixed-Point Load and Store Multiple Instructions</i> on page 63.• Revised <i>Section 3.1.4.6 Fixed-Point Move Assist Instructions</i> on page 64.• Revised <i>Section 3.1.4.11 Fixed-Point Invalid Forms and Undefined Conditions</i> on page 66.• Revised <i>Section 3.5.5 Storage Model</i> on page 75.• Revised <i>Section 3.5.5.4 Storage Access Ordering</i> on page 76.• Revised the spr heading in <i>Table 3-8 SPR Table</i> on page 87.• Revised the following sprs in <i>Table 3-8 SPR Table</i> on page 87: DSCR, PMCR, TRACE.• Revised <i>Direct External Interrupt</i> on page 122.• Revised <i>Mediated External Interrupt</i> on page 123.
10 March 2015	Version 1.11 <ul style="list-style-type: none">• Revised <i>Section 3.6.3.4 Move To/From Special Purpose Register Instructions</i> on page 87.• Revised <i>Table 3-8 SPR Table</i> on page 87.• Revised <i>Section 10.4.4 I/O Bandwidth</i> on page 253.



Revision Date	Description
29 January 2015	<p>Version 1.1.</p> <ul style="list-style-type: none"> • Revised <i>Section 3.1.4.2 Alignment Interrupts</i> on page 45. • Added a note to <i>Section 3.1.4.5 Fixed-Point Load and Store Multiple Instructions</i> on page 63. • Added a note to <i>Section 3.1.4.6 Fixed-Point Move Assist Instructions</i> on page 64. • Revised <i>Section 3.1.4.11 Fixed-Point Invalid Forms and Undefined Conditions</i> on page 66. • Moved L1 cache description from <i>Section 3.5.2 Cache</i> on page 74 to <i>Section 4.1 L1 Cache</i> on page 127. • Added <i>Section 3.5.5.2 Vector Element Atomicity</i> on page 75. • Revised <i>Section 3.5.7.2 Instruction Cache Block Invalidate (icbi)</i> on page 77. • Removed tsr. from the list of transactional memory instructions and store conditional instructions in <i>Section 3.5.11 Book II Invalid Forms</i> on page 82. • Added BESCR (806) to <i>Table 3-8 SPR Table</i> on page 87. • Renamed spr 848 to IC and spr 849 to VTB in <i>Table 3-8 SPR Table</i> on page 87. • Renamed the following sprs to Reserved in <i>Table 3-8 SPR Table</i> on page 87: spr 31, spr 888, and spr 889. • Removed the following sprs from <i>Table 3-8 SPR Table</i> on page 87: APSCR (spr 138), APSCRU (spr 139), LDBAR (spr 850), RWMR (spr 885), and TSR (spr 897). • Revised <i>Section 3.7.1 HID0 Register</i> on page 94. • Revised <i>Table 3-23 Implementation MSR and SRR1/HSRR1 Bits</i> on page 116. • Revised <i>Section 3.8.23.6 External Interrupt</i> on page 122. Added subsections <i>Direct External Interrupt</i> and <i>Mediated External Interrupt</i>. • Revised <i>Section 3.8.28 Processor Compatibility Mode</i> on page 126. • Removed reference to icswx in <i>Section 4.2.1 L2 Cache Features</i> on page 127. • Reorganized the order of sections in <i>Section 5 Simultaneous Multithreading</i> on page 135. • Revised <i>Section 5.3 Program Priority Register (PPR)</i> on page 136. • Revised <i>Section 5.16.1.1 Dispatch Flush Rules</i> on page 152. • Changed winkle instruction to rvwinkle instruction in <i>Section 9.5.4 Actuators/Controls</i> on page 180, <i>Section 9.7.2.5 Winkle</i> on page 187, <i>Section 9.7.4.2 Definitions</i> on page 189, and <i>Table 9-3 Power Management Idle Control Register (PMICR)- SPR 852</i> on page 194. • Removed reference to ACOP in <i>Section 9.7.2.2 Single Thread Nap, Sleep, and Winkle</i> on page 185.
29 January 2015 (continued)	<p>Version 1.1. (continued)</p> <ul style="list-style-type: none"> • Added <i>Section 10.1.8 Group Ending NOP</i> on page 209. • Removed the following instructions from <i>Table 10-6 List of Instructions Marked as First</i> on page 210: icswx, icswx., pbt., tsr., and waitasec. • Removed the following instructions from <i>Table 10-7 List of Instructions Marked as Last</i> on page 211: mtspr_TACR, mfspr_Idbar, mtspr_Idbar, mfspr_rwrm, mtspr_rwrm, tsr., and waitasec. • Removed the pbt(.) instruction from <i>Table 10-8 2-Way Cracked Instructions</i> on page 213. • Revised <i>Section 10.1.11 Microcode</i> on page 214. Removed table previously known as <i>Table 10-10. Instructions that Access Microcode</i>. • Revised <i>Section 10.1.12 Instruction Fusion</i> on page 215. • Removed a reference to icswx in the section <i>Conflicts Due to Write-Back Collisions and VSU/SPR Resources</i> on page 223. • Revised <i>Section 10.1.22 Store Queue and Store Forwarding</i> on page 231. • Defined pclk in <i>Section 10.3.1 Cache Latencies and Bandwidth</i> on page 235 • Corrections in <i>Table 10-16 Instruction Latencies and Throughputs</i> on page 236. • Revised <i>Section 10.4.1 Bandwidth</i> on page 253. • Moved <i>Section 3.5 Instruction That Can Soft Patch</i> to <i>Section 10.6.2 Instruction That Can Soft Patch</i> on page 258. • Added <i>Appendix A POWER8 Instruction Summary by Category</i> on page 313. • Added <i>Appendix B POWER8 Instruction Summary by Mnemonic</i> on page 339. • Added <i>Appendix C POWER8 Instruction Summary by Opcode</i> on page 363.
April 22, 2014	Version 1.0 (initial version).



About this Document

This user's manual describes the IBM® POWER8® processor. This document provides information about the registers, facilities, initialization, and use of the POWER8 processor.

This document provides information about the POWER8 processor that is visible from a programming model point of view, and is intended to be a companion to the baseline architecture documentation (see *Related Documents* on page 22). While there are some programming model considerations associated with chips and subsystems outside of the Central Electronics Complex (CEC), this document focuses primarily on the microprocessor core and the storage subsystem. For information about other chips that might appear in POWER8 systems, see the functional specifications for these individual chips.

Who Should Read this Document

This manual is intended for system software and hardware developers and application programmers who want to develop products for the POWER8 processor. It is assumed that the reader understands operating systems, microprocessor system design, basic principles of reduced instruction set computer (RISC) processing, and details of the Power ISA.

Conventions Used in This Document

This section explains numbers, bit fields, instructions, and signals that are in this document.

Representation of Numbers

Numbers are generally shown in decimal format, unless designated as follows:

- Hexadecimal values are preceded by an “x” and enclosed in single quotation marks.
For example: x'0A00'.
- Binary values in sentences are shown in single quotation marks.
For example: '1010'.

Note: A bit value that is immaterial, which is called a “don't care” bit, is represented by an “X.”

Bit Significance

In the POWER8 documentation, the smallest bit number represents the most significant bit of a field, and the largest bit number represents the least significant bit of a field.

Other Conventions

POWER8 instruction mnemonics are shown in lower-case, bold text. For example: **tlbivax**. I/O signal names are shown in upper case.



Related Documents

The following documents can be helpful when reading this specification. Contact your IBM representative to obtain any documents that are not available through [IBM Customer Connect](#) or [Power.org](#).

POWER8 Processor Single-Chip Module Datasheet

POWER8 Memory Buffer Datasheet

POWER8 Memory Buffer User's Manual

Power ISA User Instruction Set Architecture - Book I (Version 2.07)

Power ISA Virtual Environment Architecture - Book II (Version 2.07)

Power ISA Operating Environment Architecture (Server Environment) - Book III-S (Version 2.07)

Power ISA Transactional Memory (Version 2.07)

PowerPC Architecture Platform Requirements (PAPR+) Specification

[*PCI Express Base Specification, Revision 3.0*](#)

1. POWER8 Processor Overview

The POWER8 processor is a superscalar symmetric multiprocessor designed for use in servers and large-cluster systems. It uses IBM's CMOS 22 nm SOI technology with 15 metal layers.

1.1 General Features

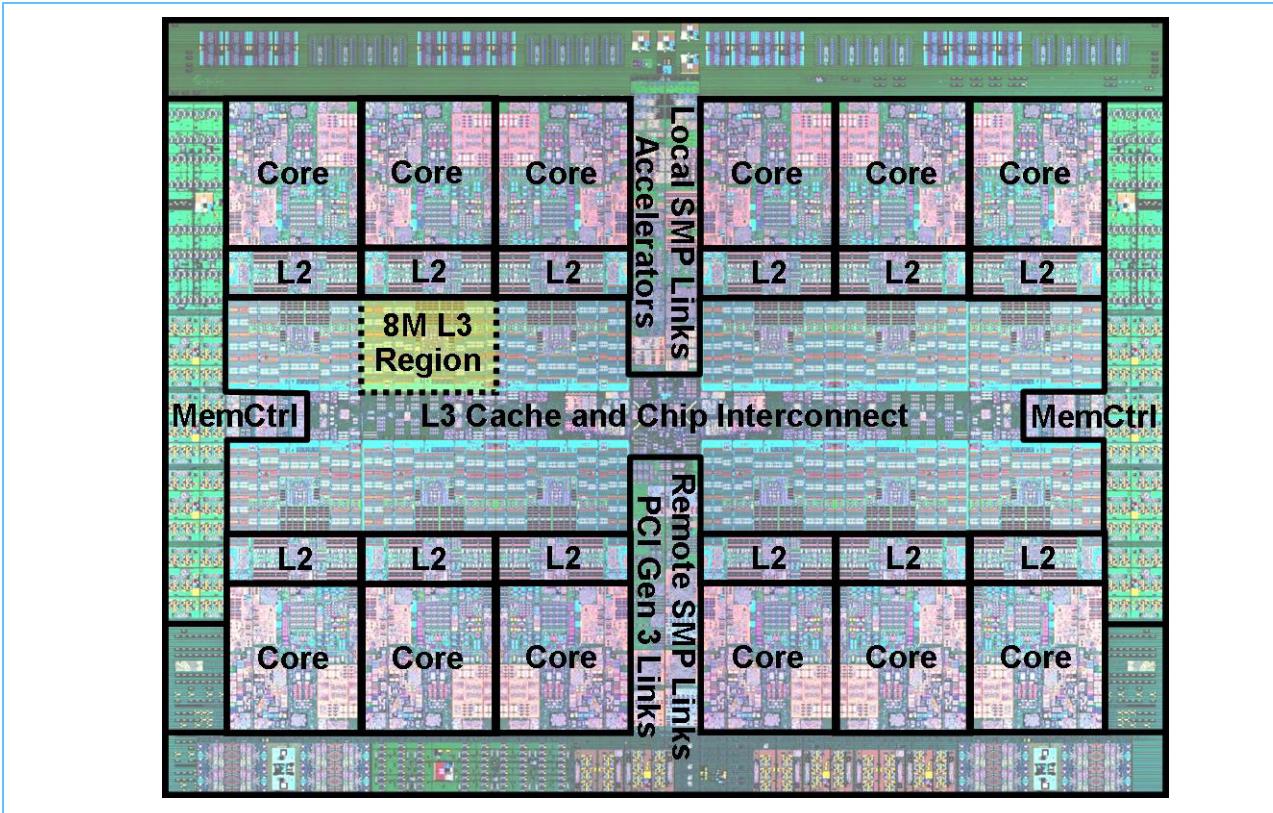
The POWER8 processor can have up to 12 cores enabled on a single chip in the single-chip module (SCM) configuration. Each core has eight threads using simultaneous multithreading (SMT). The SMT is dynamically tunable, so that each core can have one, two, four, or eight threads.

The POWER8 processor supports the following architectural features:

- PowerPC Architecture Book I, II, and III version 2.07
- PowerPC Architecture Platform Requirements (PAPR+), Version 2.1
- IEEE P754-2008 floating-point compliant
- Big-endian, little-endian, strong-ordering support extension
- 50-bit real address, 68-bit virtual address

Figure 1-1 provides a block diagram of the POWER8 processor on the SCM.

Figure 1-1. Block Diagram for the POWER8 Processor





The following features describe the main components of the 12-core POWER8 processor chip:

- POWER8 core and cache
 - Up to 12 processor cores
 - 16 execution pipes
 - 8-way SMT, Out-of-Order (O-o-O)
 - 124×2 GPR and 144×2 VMX/VSX/FPR renames
 - Software-architected register file
 - Concurrent support of 1 - 4 LPARs per core
- Chiplet
 - 32 KB instruction cache (I-cache)
 - 64 KB data cache (D-cache)
 - 512 KB private L2 cache
 - Local 8 MB L3 cache region
- POWER8 SMP on-chip interconnect
 - Eight 16-byte data buses, two address snoop buses, 32 on/off ramps
 - Asynchronous interface to core/L2/L3 and off-chip interconnect
- Four 9.6 GHz differential memory interfaces (each with 2-byte read and 1-byte write) to the POWER8 Memory Buffer chip
- POWER8 SMP off-chip interconnect
 - Maximum 48-way SMP
 - Three 6.4 GHz differential SMP interfaces (each with 2 bytes per direction)
- 8 GHz differential PCIe Generation 3 buses: 32 lanes configured as $(16x + 16x)$ or $(16x + 8x + 8x)$
- Coherent Accelerator Interface Architecture (CAPI) allows an FPGA or ASIC to connect coherently to the POWER8 processor SMP interconnect via the PCIe.
- Power management support
 - Hypervisor-directed power change requests using a PState mechanism
 - Sensors
 - Digital thermal sensor (DTS2) $\pm 5^\circ\text{C}$
 - Off-chip analog thermal diode $\pm 1 - 2^\circ\text{C}$
 - Dedicated performance, microarchitecture, and event counters
 - On-chip controller (OCC)
 - On-chip PowerPC 405 for real-time frequency and voltage modification
 - On-chiplet hardware assist (automated core chiplet management)
 - On-chip power management controls automated communications to the voltage regulation modules (VRMs) and voltage and frequency sequencers for automated Pstate and idle state support
 - Actuators
 - Per-chiplet frequency control through the DPLL
 - Per-chiplet internal VRMs for independent voltage control
 - Architected idle states: nap, sleep, and winkle; each with increasing power savings capability (and latency)



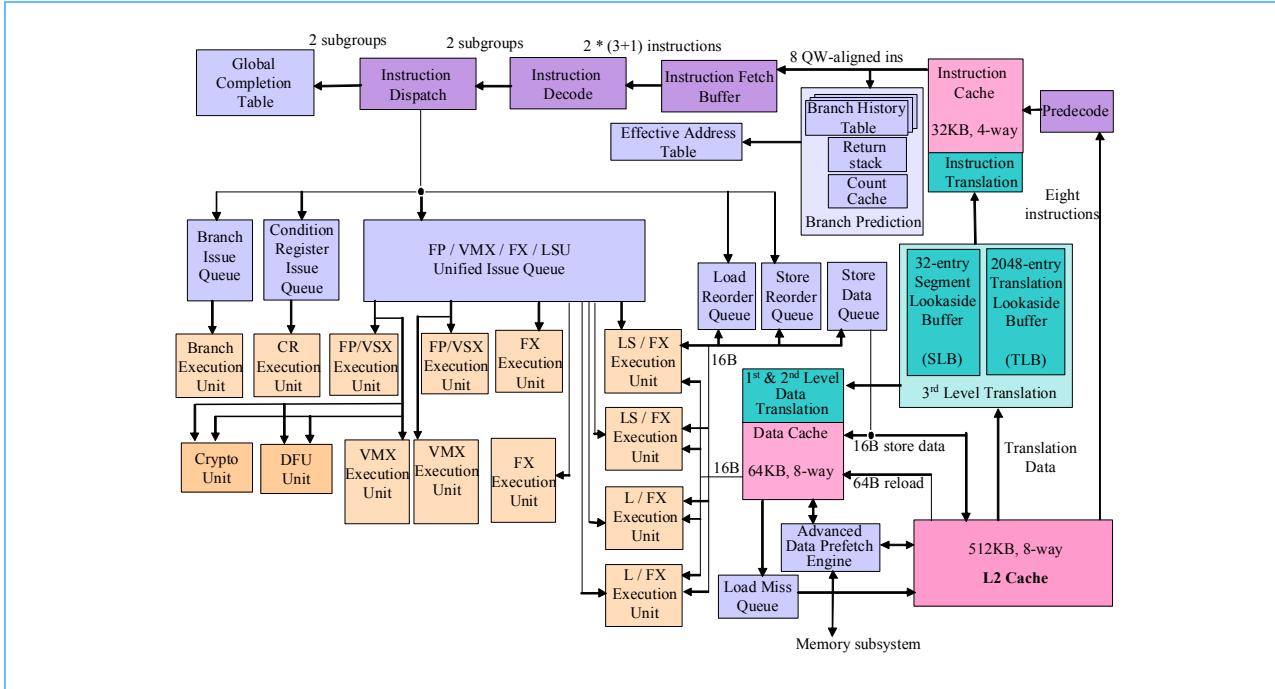
- SPR power management control registers (PMCR, PMICR, PMSR) for hypervisor support
- Memory/DIMM throttling for memory subsystem power and thermal management
- Features
 - On-chip accelerators
 - On-chip: compression, encryption, data move initiated by hypervisor
 - In-core: user invocation encryption (AES, SHA)
 - Cloud computing enhancements: page replacement/affinity assist, micropartition prefetching, IPL time reduction, four concurrent LPARs per core
 - Transactional memory
 - Random number generator
 - RAID6 support in VMX
 - Support for industry standard BMC
 - Multi-level TCE support
 - Turbo mode support



2. POWER8 Processor Core

This section provides an overview of the POWER8 processor core, including key design fundamentals, an overview of the master pipeline operation, and a detailed summary of key design features.

Figure 2-1. POWER8 Processor Core



2.1 Key Design Fundamentals

This section describes the key design fundamentals of the POWER8 processor core.

2.1.1 64-Bit Implementation of the Power ISA (version 2.07)

- Compatibility for all Power ISA application-level code (problem state)
 - Supports partition mobility
- Supports IEEE P754-2008 for binary floating-point arithmetic
- Support for 32-bit operating system bridge facility



2.1.2 Layered Implementation Strategy for High-Frequency Operation

- Deeply pipelined design:
 - 16 stages from I-cache access to writeback for most fixed-point register-to-register operations
 - 18 stages for most load/store operations (assuming an L1 D-cache hit) from I-cache to writeback
 - 23 stages for most floating-point operations from I-cache access to writeback
- Dynamic instruction cracking¹ for some instructions allows for simpler inner core data flow:
 - Dedicated data flow for cracking one instruction into two or more internal operations
 - Microcoded templates for longer emulation sequences

2.1.3 Speculative Superscalar Inner Core Organization

- Multi-threaded core design:
 - Single thread (ST), 2-way multithread (SMT2), 4-way multithread (SMT4), and 8-way multithread (SMT8) modes supported in single LPAR mode
 - (SMT8) supported in 2 and 4 LPAR mode
- Aggressive branch prediction:
 - Prediction for up to eight branches per cycle
 - Support for up to 24 predicted taken branches in flight per thread in ST and SMT2 mode, 12 predicted taken branches in SMT4 mode, and six predicted taken branches in SMT8 mode. The number of predicted not-taken branches tracked can be higher.
 - Prediction support for branch direction and branch target addresses
- In single-thread mode, in-order dispatch of up to eight internal operations (iops) into distributed issue queues per cycle:
 - Up to two branches in a dispatch group; the first branch can be predicted taken or not-taken
 - Up to six non-branch instructions in the dispatch group
 - Second branch terminates the group
 - No VS-routed instruction after a branch in the group
- In SMT2 and beyond there are two dispatch sets each with
 - Up to one branch in dispatch group, the first branch can be predicted taken or not-taken
 - Up to three non-branch instructions in the dispatch group
 - No VS routed instruction after a branch in the group
- Out-of-order issue of up to 10 operations into the following 10 issue ports:
 - Two ports to do loads or fixed-point operations.
 - Two ports to do stores, fixed-point loads, or fixed-point operations.
 - Two fixed-point operations
 - Two issue ports shared by two floating-point, two VSX, two VMX, one crypto, and one DFP operations
 - One branch operation
 - One condition register operation
- Register renaming on GPRs, FPRs, VRFs (VMX and VSX Registers), CR fields, XER (parts), FPSCR, VSCR, Link and Count Registers

1.Process by which some complex instructions are broken into multiple simpler, more RISC-like instructions.

- Sixteen execution units:
 - Two symmetric load/store units (LSU), capable of executing stores, fixed-point loads, and simple fixed-point operations
 - Two load-only units (LU) also capable of executing simple fixed-point operations
 - Two symmetric fixed-point units (FXU)
 - Four floating-point units (FPU), implemented as two 2-way SIMD operations for double- and single-precision. Scalar binary floating-point instructions can only use two FPUs.
 - Two VMX execution units capable of executing simple FX, permute, complex FX, and 4-way SIMD single-precision floating-point operations
 - One Crypto unit
 - One decimal floating-point unit (DFU)
 - One branch execution unit (BR)
 - One condition register logical execution unit (CRL)
- Large number of instructions in flight:
 - Up to 64 instructions in the instruction fetch buffer in ST mode and up to 128 instructions total in SMT2, SMT4, and SMT8 modes (up to 64 instructions per thread in SMT2 mode, 32 instructions per thread in SMT4, and 16 instructions per thread in SMT8)
 - Up to 48 instructions in three decode pipe stages and three dispatch buffers
 - Up to 224 instructions in the inner-core (after dispatch)
 - Up to 40 stores queued in the SRQ (available for forwarding), shared by the available threads
- Fast, selective flush of incorrect speculative instructions and results

2.1.4 Specific Focus on Storage Latency Management

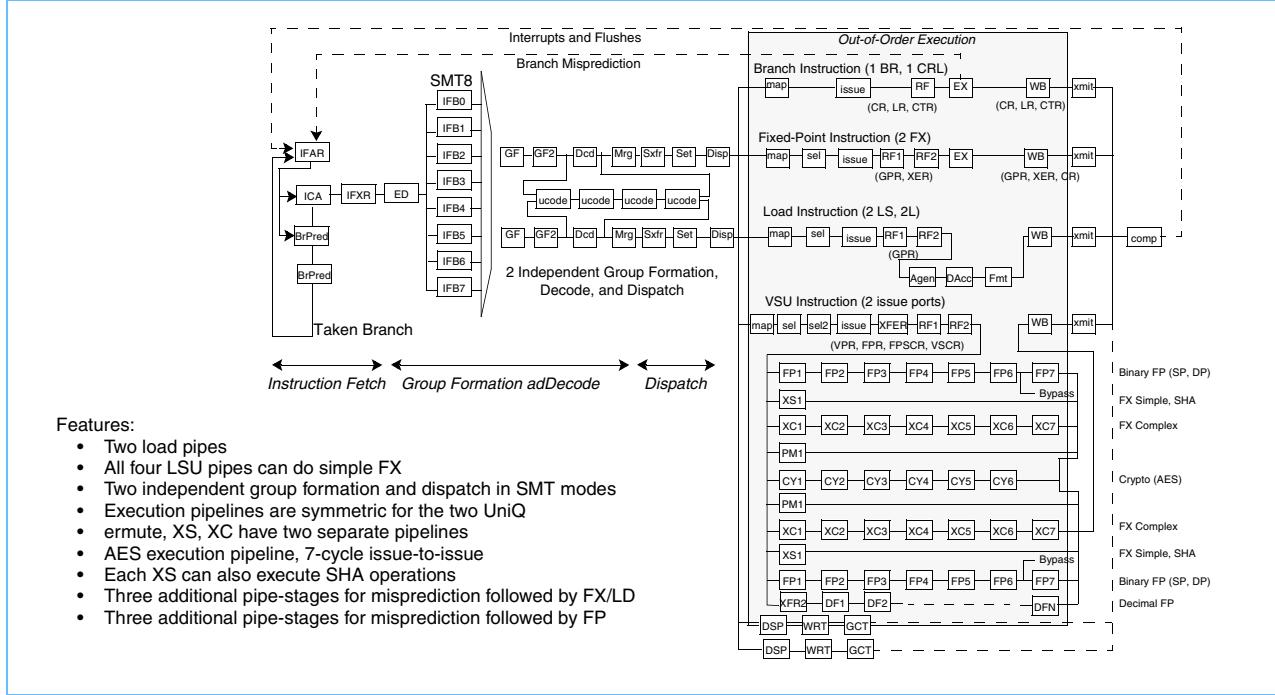
- Out-of-order and speculative issue of load operations
- Support for up to 16 outstanding L1 cache-line misses
- Hardware- or software-initiated instruction prefetching from L2 cache, L3 cache, and memory
- Hardware-initiated data-stream prefetching (using effective addresses). Support for up to 12 active streams
- Critical word forwarding, critical sector first
- Hardware instruction prefetching supported

2.2 Pipeline Structure

The pipeline structure for the processor can be subdivided into a *master pipeline* and several different *execution unit pipelines*. The master pipeline presents speculative in-order instructions to the mapping, sequencing, and dispatch functions. It ensures an orderly completion of the real execution path (throwing away any other potential speculative results associated with mispredicted paths). The execution unit pipelines allow out-of-order issuing of both speculative and non-speculative operations. The execution unit pipelines progress independently from the master pipeline and from one another.

Figure 2-2 illustrates these pipelines, where each box represents a pipeline stage.

Figure 2-2. Pipeline Structure



The legend for Figure 2-2 is as follows:

IFAR	Instruction fetch address register
ICA	Instruction cache access
ifxer	Instruction transfer
ED	Early decode cycle
D0	IDU predecode stage and instruction-fetch buffer latches
GF and GF2	Dispatch group determination
Dcd	Main decoder
Mrg	For assembly
ucode	Microcode
gxfer	Microcode selection and group transfer
Disp	Inter-dependency determination for instructions in the group
map	Register mapping
Sel	Issue queue selection
Issue	Instruction Issue
RF and RF2	Register file access
EX	Execution
WB	Writeback to the register file

EA/CA	Effective-address generation and data-cache decode
CA/fmt	Data-cache access and data formatting
FP1	Floating-point alignment and multiply
FP2	Floating-point alignment and multiply
FP3	Floating-point add
FP4	Floating-point add
FP5	Floating-point normalize result
FP6	Floating-point round and local 6-cycle forwarding
xmit	Finish and transmit
DSP	Group dispatch
WRT	Format and write into the GCT
GCT	Global completion table
Comp	Group completion
For VMX operations	
FP1 - FP6	Pipeline stages for the 4-way SIMD single-precision pipeline stages
XS1	Simple FX operation stage
XC1 - XC6	Complex FX operation stages
PM1/PM2	Permute stages

The processor core is divided into following seven units:

- IFU - Instruction fetch and decode unit
- ISU - Instruction dispatch and issue unit
- LSU - Load/store unit
- FXU - Fixed-point execution unit
- VSU - Vector and scalar unit (consists of VMX, binary floating-point, Crypto, and VSX)
- DFU - Decimal floating-point unit
- PC - Pervasive core unit

2.3 Microprocessor Core - Detailed Features

See [Section 11.1.12 Instruction Fusion](#) on page 213, [Section 11.1.13 Instruction Dispatch](#) on page 214, and [Section 11.1.14 Instruction Issue](#) on page 214 for additional core information.

2.3.1 Instruction Fetching and Branch Prediction

- 32 KB, 8-way set-associative I-cache:
 - 128-byte lines (broken into four 32-byte sectors).
 - Dedicated 64-byte interface from the L2 cache that can supply 64 bytes in every other processor clock.
 - Critical-sector-first reload policy.
 - Effective-address index, real-address tags.



- Banked I-cache, supports one read and one write per cycle when there is no bank conflict.
- Eight additional predecode bits per word to aid in fast decoding and group formation.
- Parity protected; force invalidate and reload on parity error.
- 64-entry effective-to-real address (ERAT) translation cache, fully associative.
 - Each entry can translate 4 KB, 64 KB, or 16 MB pages. For MSR[IR] = '1' and VRMA accesses, 16 MB and 16 GB pages take multiple 64 KB entries (the same is true for 1 MB pages; however, 1 MB pages are not visible to server workloads).
 - In MSR[IR] = '0' mode and non-VRMA accesses, 16 MB and 16 GB pages are installed as 4 KB translation blocks in the I-ERAT (the same is true for 1 MB pages).
 - In SMT mode, each entry is tagged to indicate invalid, valid for thread 0, valid for thread 1, valid for thread 2, valid for thread 3, valid for thread 4, valid for thread 5, valid for thread 5, or valid for thread 7.
- Fetch quadword aligned block of eight instructions per cycle.
 - In ST mode, instructions are fetched from the thread in every cycle
 - In all other modes, instructions are fetched from a given thread based on thread priority. If the threads are of equal priority, each thread gets approximately an equal number of fetch cycles, while optimizing the core throughput.
- Branch prediction:
 - Scan all eight fetched instructions for branches in each cycle
 - Predict up to eight branches per cycle (if the first one is predicted, fall-through)
 - Three table prediction structures: global, local, and selector (16K entries × 2 bits, 16K entries × 2 bits and 16K entries × 2 bits, respectively).
 - BHT tables are 16-way banked for concurrent read and write, when there is no bank conflict.
 - Global BHT is accessed using 20 bits of past global fetch history (folded into 11 bits to access global BHT)
 - 32-entry link stack for subroutine return address prediction (with some of these entries allocated for speculation) per thread in ST and SMT2 modes. For SMT4 mode, there are 16 entries available per thread. In SMT8 mode, there are 8 entries available per thread.
 - 768-entry count cache for address prediction shared by all the active threads running on the core. Five-hundred twelve (512) of the count cache entries are accessed based on the effective address of the **bctr** instruction XORed with the folded GHV with two confidence bits used to provide hysteresis for replacement. The other 256 entries are accessed with just the effective address of the **bctr** with one confidence bit.
 - No instruction fetch bubble to fetch from the sequential path of a predicted not-taken branch.
 - In ST mode, there are two cycles of instruction fetch bubbles to fetch from the target address of a predicted taken branch.
 - In SMT2 and SMT4 mode, thread priority is factored in to determine which thread to fetch from (to improve overall fetch throughput).
 - Track up to 24 outstanding taken branches per thread in ST and SMT2 mode, 12 outstanding taken branches in SMT4 mode, and six outstanding taken branches per thread in SMT8 mode. The number of predicted not-taken branches tracked can be higher.

2.3.2 Instruction Decode and Preprocessing

- 3-cycle pipeline to decode and preprocess instructions
 - Dedicated data flow for cracking one instruction into two or four internal operations. All instructions that crack into two internal operations and a subset of instructions that crack into four internal operations use the dedicated data flow. The rest of the cracked instructions use the microcoded templates.
 - Microcoded templates for longer emulation sequences of internal operations.
 - All internal operations expand into an approximately 80-bit internal form to simplify subsequent processing and explicitly expose register dependencies for all register pools.
 - Dispatch groups with up to eight instructions are formulated in ST mode. Two groups of up to four instructions are formulated in all other modes.
 - All cracked instructions must be the first instruction in a group.
 - Cracked and microcoded instructions have access to three renamed eGPRs, one renamed eCR-field, and two eVRs (for cracking 128-bit DFP instructions). The eGPR, eCR and eVR are extensions to the architected facilities.
- Logically there is one instruction fetch buffer (IFB) per thread (sizes differ based on the ST, SMT2, or SMT4 mode). Each IFB entry has four instructions.
 - There are 16 entries in an IFB per thread in ST and SMT2 mode, eight entries per thread in SMT4, and four entries per thread in SMT8 mode.
 - Physically, the IFB is implemented as one register file, partitioned for ST, SMT2, and SMT4 modes.
- Up to eight instructions can be placed in the IFB in a cycle (ST or SMT mode).
- Up to eight instructions can be taken out from the IFB in a cycle (ST or SMT mode).
- Instructions taken out for group formation and decode are from one (ST mode) or two threads.

2.3.3 Instruction Dispatch, Sequencing, and Completion Control

- Three dispatch buffers that can hold up to three dispatch groups when GCT is full.
- Inter-instruction dependence generation for RAW and WAW dependences.
- 28-entry global completion table:
 - Each entry is assigned to a particular thread at instruction dispatch.
 - Entries can be allocated nonsequentially and intermixed among the threads in SMT modes.
 - Group-oriented tracking associates up to two 4-operation dispatch groups or one 8-operation dispatch group to a single GCT entry.
 - Tracks internal operations from dispatch to instruction completion for up to 224 operations.
 - Branch instruction can be placed in the middle of a group (predicted taken or not-taken). A second branch always ends the group.
 - Branch misprediction for a branch placed in the middle of a group causes a partial group flush. The LSU flushes cause a flush of the entire group, even though the load/store operation might be in the middle of the group (no partial flush for LSU).
 - Capable of restoring the machine state for any of the instructions in flight.
 - Fast restoration for instructions on group boundaries (such as, branches).
 - Slower restoration for instructions contained within a group (such as, load/store operations).



- Supports precise exceptions (including machine check exceptions).
- Register renaming resources. A given thread can use any entry in the renamed register files, which can be dynamically shared amongst threads.
 - GPR rename mapper:
 - 124 entry in ST mode (32 architected and 92 for rename)
 - 2×124 entry in SMT2 mode (32 architected and 92 rename, per thread)
 - 2×124 entry in SMT4 mode. One GPR register file supports half the threads, and the other GPR register file supports the other half. Each GPR register file has 64 architected, leaving 60 for rename.
 - 2×124 entry in SMT8 mode. One physical register file supports half of the threads and the other register file supports the other half. The number of architected GPRS is limited to 64 at a time in the GPRs. The other least recently used GPRs are located in the SAR. This leaves 60 for rename per thread set.
 - Four eGPRs used on demand per thread (used for microcoded instructions)
 - In SMT2, SMT4, and SMT8 modes, a total of 106 renames (nonarchitected) are available, across both thread sets
 - The SAR is 72 entries per thread (32 GPR + 4 eGPR, $\times 2$ for TM checkpointing) and contains castouts from the register file.
 - FPR and VR rename mapper:
 - 144 entry in ST mode (64 architected and 80 rename).
 - 144×2 entry in SMT2 mode (64 architected and 80 rename, per thread).
 - 144×2 entry in SMT4 and SMT8 modes (limited to 64 architected and 80 rename, per thread set - castouts in the SAR).
 - In SMT2, SMT4, and SMT8 modes, a total of 106 renames (non-architected) are available, across both thread sets.
 - The SAR is 128 entries per thread (64 VSX, $\times 2$ for TM checkpointing) and contains castouts from the register file.
 - 30-entry XER rename mapper plus 32-entry ARF for the current architected values (XER broken into four mappable fields (ov, ca/oc, fxcc, tgcc) and one non-mappable field per thread)
 - Non-mappable bits: dc, ds, string-count; other bits special fields (value predict): so
 - 20-entry LR/CTR/TAR rename mapper plus 24-entry ARF for architected values (one LR and one CTR, and one TAR per thread).
 - 32-entry CR rename mapper plus 64-entry ARF for architected values (eight CR fields per thread).
 - 28-entry FPSCR rename mapper (each entry corresponds to a GCT entry, which can belong to at most one particular thread at a given time).
 - No register renaming on any architected registers not mentioned previously.
 - Instruction queuing resources:
 - Two 32-entry unified issue queues (UniQ) are used for fixed-point, floating-point, VMX, VSX, DFU, Crypto, and load/store instructions. The two queues are split per thread-set in SMT2, SMT4, and SMT8 mode.
 - In ST, each IOP is assigned to the opposite queue, half from the IOP before it.
 - One 15-entry issue queue for branch instructions.
 - One 8-entry issue queue for CR-logical instructions.

2.3.4 Fixed-Point Execution Pipelines

- Two symmetric fixed-point execution pipelines:
 - Both are capable of basic arithmetic, logical and shifting operations.
 - Both are capable of multiplies, divides, and SPR operations.
- Out-of-order issue with a bias towards oldest operations first.
- Symmetric forwarding between fixed-point and load/store execution pipelines.
- In ST mode, instructions from a given thread can be executed in either pipeline.
- In SMT2, SMT4, and SMT8 mode, instructions from thread set 0 execute in pipeline 0, and instructions from thread set 1 execute in pipeline 1.
- Threads are dynamically assigned to a given thread set.

2.3.5 Load and Store Execution Pipelines

- Two load/store and two load execution pipelines, with 3-cycle, load-to-use latency (2-cycle bubble) for fixed-point loads and 5-cycle load-to-use latency for VS and floating-point loads. See *Table 11-16 Instruction Latencies and Throughputs* on page 234 for additional information.
- Load/store units execute both load and store operations.
- Load units execute only load operations.
- Loads that update a FP, VSX, or VMX register execute in the LU. Loads that update an XER, execute in the LSU. Loads that update only a GPR can execute in either the LU or LSU.
- All four units can execute simple fixed-point operations:
 - In ST mode, a given load/store instruction can execute in either pipeline.
 - In SMT2, SMT4 and SMT8 modes, instructions from thread set 0 execute in pipeline 0, and instructions from thread set 1 execute in pipeline 1.
- Out-of-order issue with bias towards oldest operations first:
 - Fixed-point D-form stores are issued twice: an address-generation operation (issued to the LSU) and a data-steering operation (issued to the LU).
 - Fixed-point X-form stores are cracked into store_agen and store_data by the IFU (a cracked instruction starts a new group).
 - 64-bit floating-point and 64-bit VSX stores are issued twice: the store_agen is issued to the LSU and the store_data is issued to the VSU.
- 64 KB, 8-way set-associative, banked D-cache:
 - Supports four reads and one write in every cycle, when there is no bank conflict between write or read and a read. A given bank can support either two reads or one write in a given cycle.
 - 3-cycle load-use penalty for FXU loads (2-cycle bubble between a load and a dependent operation).
 - 5-cycle load-use penalty for FPU/VMX/VSX loads (4-cycle bubble between a load and a dependent operation).
 - Store-through (to L2 cache) policy; no allocate on store misses.
 - 128-byte cache line.
 - True LRU replacement policy.



- Dedicated 64-byte reload interface from the L2 cache, which can supply 64 bytes in every processor clock.
- Effective address index, real address tags (hardware fix-up on alias cases. That is, two different EAs that map to the same RA are not allowed to co-exist in the D-cache).
- Parity protected; precise machine check interrupt on parity error (software fix-up).
- 48-entry, fully-associative primary and 144-entry secondary data effective-to-real address (D-ERAT) translation cache shared across the two thread sets:
 - Each entry translates either 4 KB, 64 KB, or 16 MB pages:
 - 16 GB pages take multiple 16 MB pages (used for server-mode only).
 - MSR[DR] = '0' is also created in the D-ERAT and shared by all threads.
 - Binary LRU replacement policy.
 - In SMT mode, each entry is tagged to indicate the valid threads.
 - In ST mode, 48-entry primary and 256-entry secondary D-ERAT are available.
 - In SMT mode, there are 48-entry primary and 128-entry secondary D-ERAT available for each thread set. Entries are dynamically shared between the two threads (a given entry is not shared by multiple threads, unless it is MSR[DR] = '0').
- 32-entry, fully-associative segment lookaside buffer (SLB) per thread:
 - Each entry can support 256 MB or 1 TB segment sizes.
 - Multiple pages per segment (MPSS) feature is supported: 4 KB, 64 KB, and 16 MB pages (at most two pages) can be present concurrently in a given segment.
- 40-entry store re-order queue logically above the D-cache (real address based; CAM structure):
 - Sixty-four virtual entries (with no physical entity in SRQ) are available to allow a total of 64 outstanding stores to be dispatched per thread (for dispatch, virtual entry is sufficient).
 - A total of 40 outstanding stores can be issued (for issue, a real entry is required):
 - The SRQ is dynamically shared among the available threads.
 - The SRQ entry is allocated at the time of a store issue and deallocated when the store is written in the cache (after the completion point).
 - Store addresses and store data can be supplied on different cycles.
 - Stores wait in this queue until they are completed; then, they write the cache.
 - Supports store forwarding to inclusive subsequent loads (even if both are speculative). Store forwarding takes five additional cycles compared to a D-cache hit for a load.
 - For each SRQ entry, there is a store data queue (SDQ) entry of 16 bytes.
 - 16 bytes of store data can be sent to the L2 cache (and also to the D-cache, on a hit) in every processor cycle. Two distinct stores can be combined into one 16-byte store, under certain conditions).
- 44-entry load re-order queue (real address based; CAM structure):
 - Sixty-four virtual entries (with no physical entity in LRQ) are available to allow a total of 64 outstanding loads to be dispatched per thread in ST, SMT2, or SMT4 modes (for dispatch, virtual entry is sufficient).
 - A total of 44 outstanding loads can be issued. For instruction dispatch, virtual entry is sufficient; for issue, a real entry is required.
 - LRQ is dynamically shared among the available threads.

- Keeps track of out-of-order loads and watches for hazards. For example:
 - Previous store to the same address that gets executed after the load; system executes a flush.
 - Previous load from the same address and a cross-invalidate has occurred; system executes a flush.
 - Atomic load-quad instruction and a cross invalidate has occurred; system executes a flush.
- 16-entry load miss queue (real address based):
 - Keeps track of loads that have missed in the L1 D-caches.
 - Dynamically shared among the threads in SMT2, SMT4, and SMT8 modes.
 - Prefetches from L1 cache are also tracked using the LMQ.
- Two 16-byte loads and one 16-byte store operation are supported for VMX and VSX operations per cycle. All architecturally-allowed alignments are supported in hardware.
- True little-endian (LE) mode is supported. All architecturally allowed alignments are supported in hardware.

2.3.6 Branch and Condition Register Execution Pipelines

- One branch execution pipeline:
 - Computes actual branch address and branch direction for comparison with prediction.
 - Redirects instruction fetching if either direction or target prediction was incorrect.
 - Assists in training and maintaining the branch history table predictors, the link stack, and the count cache.
- One Condition Register (CR) logical pipeline:
 - Executes CR logical instructions and the CR movement operations.
 - Also executes some **mfsp** instructions .
- Out-of-order issue with bias towards oldest operations first.

2.3.7 Unified Second-Level Memory Management (Address Translation)

- 2048-entry, 4-way set associative TLB:
 - 4 KB, 64 KB, 16 MB, and 16 GB pages are supported in TLB.
 - The TLB also supports “Virtualized Page Class Key Protection” with 32 keys.
 - Hardware-based reload (from the L2 cache interface, no L1 D-cache corruption).
 - Hardware-based update of the R bit, C bit, and TS bit.
 - Parity protected; precise machine check interrupt on parity error (software fix-up).
 - ITLB entries are shared by the eight threads as long as the entry belongs to the logical partition running on the core.
 - 12-bit LPAR ID per entry.
- Hit-under-miss is allowed in the TLB.
- Support for four concurrent table walks (without any restriction on thread of D-side or I-side requests).
- 32-entry fully-associative SLB, one per thread:
 - An SLB miss results in an interrupt (software reloads the SLB).
 - An SLB can also be loaded via the 32-bit PowerPC segment register instructions.
 - An SLB supports 256 MB and 1 TB segment sizes.
- A segment with a 4 KB base page size is allowed to have mixed pages of sizes 4 KB, 64 KB, and 16 MB pages.
- A segment with a 64 KB base page size is allowed to have mixed pages of sizes 64 KB and 16 MB pages.



- A read of an invalid SLB entry returns zeros for enhanced security.
- Supports 68-bit virtual address and 50-bit real address.
- Both software and hardware TLB management is allowed.
- True LRU replacement policy.

2.3.8 Data Prefetch

- Software initiated streams can use up to 16 entries. The 16 entries are shared in ST and SMT2 mode. In SMT4 mode there are two groups of eight (two threads share a group) and in SMT8 mode there are four groups of four entries (again two threads share a group) managed with effective addresses.
- Supports hardware- and software-initiated streams.
- Hardware-initiated streams are dynamically shared among the available threads.
- Sixteen independent data streams capable of striding up or down.
- Stride one cache line support.
- Stride N support
- The 4-entry table tracks 32-byte strides across the previous eight miss addresses [50:58] to detect 32-byte strides
- The stream is installed in the main 16-entry queue when a stride is detected.
- Prefetches and allocates up to two cache lines ahead of a load into the L1 D-cache.
- The ramp and depth of prefetch is controlled using the DSCR Register.
- Support for software-initiated stream startup (special variant of the **dcbt** instruction).

2.3.9 VSU Execution Pipeline

- The VSU unit contains a binary floating-point execution unit, SIMD double-precision floating-point (VSX) execution unit, decimal floating-point unit (DFU), Crypto unit, and the VMX execution unit.
- Up to two instructions can be issued to the VSU in a given cycle to the two pipelines.
 - The pipes are fully symmetrical except for the Crypto and DFU engines which are shared.
 - In ST mode, instructions can be issued to both pipes.
 - In SMT2, SMT4, and SMT8 mode, pipe0 executes instructions from thread set 0 and pipe 1 executes instructions from thread set 1.
 - Out-of-order issue with a bias towards oldest operations first.
 - Two load result bus to the VRF; each supports up to 16-byte loads in a cycle.
 - Store data bus from VRF to the SDQ supports two 16-byte stores in a cycle.
- Floating-point execution.
 - Two symmetric floating-point execution pipelines, with 6-stage execution.
 - Both are capable of the full set of floating-point instructions.
 - All data formats supported in hardware (no floating-point assist interrupts).
 - Back-to-back 6-cycle issue to both local and remote FPU (symmetric forwarding between the floating-point pipelines, with no additional cycle of latency).
- VSX execution.
 - Two symmetric SIMD floating-point execution pipelines, with 6-stage execution.
 - Both are capable of the full set of VSX instructions (single-precision and double-precision).
 - All data formats supported in hardware (no assist interrupts).

- A test instruction facilitates execution of multiple concurrent divide or square-root operations.
- Back-to-back 6-cycle issue to both local and remote VSX pipe.
- VMX execution.
 - Four execution pipelines within VMX: simple fixed-point, complex fixed-point, permute, and 4-way SIMD single-precision floating-point unit.
 - Simple fixed-point operations take two execution cycles.
 - Complex fixed-point operations take seven execution cycles.
 - Permute operations take two execution cycles.
 - Vector floating-point operations take six execution cycles.
- Crypto execution.
 - The Crypto unit executes symmetric [AES](#) instructions that include polynomial multiply to support the Galois Counter Mode (GCM).
 - Allows out-of-order issue with a bias towards the oldest instruction.
 - Crypto operations can be issued from either issue queue.
 - Pipelined execution allowed.
 - Crypto operations take six execution cycles.
- See *Table 11-16 Instruction Latencies and Throughputs* on page 234 for additional information.

2.3.10 DFP Execution Pipeline

- The DFP unit can execute 64-bit or 128-bit decimal floating-point operations.
- Allows out-of-order issue with a bias towards the oldest instruction.
- Pipelined execution allowed.
- DFP operations can be issued from either issue queue.
- The 128-bit DFP instructions can be cracked into 2-way or 4-way internal operation.





3. Power Architecture Compliance

The following sections are intended to be read with their respective companion documents. Throughout these sections, it is assumed that the reader is familiar with the following architecture documents:

- *Power ISA User Instruction Set Architecture - Book I (version 2.07)*
- *Power ISA Virtual Instruction Set Architecture - Book II (version 2.07)*
- *Power ISA Operating Environment Architecture (Server Environment) - Book III-S (version 2.07)*

3.1 Book I - User Instruction Set Architecture

This section of the document identifies architectural implications of the POWER8 design point as they relate to the User Instruction Set Architecture. This is accomplished by walking through each of the relevant sections of Book I and highlighting the POWER8 solution to the architectural flexibility provided by the Power ISA.

In many cases, the architecture defines certain scenarios as *invalid forms*. In these cases, although this document provides information on how the POWER8 core handles these scenarios, it is strongly recommended that software avoid building any type of reliance on these behaviors because they are likely to be different in future generation machines. This information is primarily provided as an aid to the design verification and debug efforts

3.1.1 Defined Instructions

The POWER8 processor core implements all Book I instructions in the categories listed as being required for the server platform in *Appendix B. Platform Support Requirements of the Power ISA (Version 2.07)*.

3.1.1.1 Illegal Instructions

An attempt to execute an illegal instruction as defined in the *Appendix D. Illegal Instructions of the Power ISA (Version 2.07)* results in a hypervisor emulation assistance interrupt.

3.1.1.2 Instructions Supported

The POWER8 core supports all of the instructions described in Book I: Power ISA User Instruction Set Architecture of the *Power ISA*, except those instructions with the designation of *Embedded*. Furthermore, it supports the Service Processor “Attention” described in *Appendix E. Reserved Instructions of the Power ISA (Version 2.07)*. This instruction is conditionally enabled by HID0[31] = ‘1’. When enabled, this instruction is a user-level instruction.

3.1.1.3 Invalid Forms

In general, the POWER8 core handles *invalid forms* of instructions in the manner that is most convenient for the particular case (within the scope of meeting the boundedly-undefined definition described in the Power ISA). This document specifies the cases where a system-level error handler is invoked, but does not always describe actions for other cases of invalid forms. It is not recommended that software or other system facilities make use of the POWER8 behavior in these cases, because it is not formally specified and might be different in another processor that implements the Power ISA.



The POWER8 core ignores the state of reserved bits in the instructions (denoted by “///” in the instruction definition) and executes the instruction normally. Software should set these bits to ‘0’ per the Power ISA.

3.1.2 Branch Processor

3.1.2.1 Instruction Fetching

In an effort to increase performance, the POWER8 processor does instruction prefetching before it determines whether or not particular instructions will actually execute. This prefetching follows all of the architectural constraints relative to cache inhibited and guarded regions of storage. A set of software-accessible mode bits is implemented to allow control over the various types of prefetch supported (for more information, see *Section 3.7 HID Registers (HID0, HID1, HID4, and HID5)* on page 93).

3.1.2.2 Branch Prediction

The POWER8 processor core uses several dynamic branch prediction mechanisms to improve performance. A set of three branch history tables (local, global, and selector) is used to predict the direction of branch instructions early in the pipeline. To improve the efficiency of these predictors, the POWER8 core uses the architected BO field hint bits associated with many of the branch instructions (the “a” and “t” bits).

In addition, for **bclr** instructions, a link stack (or call-return stack) is used to predict the target address of the branch. Similarly, for **bcctr** instructions, local and global count caches are used to predict the target address for this type of branch. To improve the efficiency of these address predictors, the POWER8 core uses the architected BH-field hints associated with several of the branch instructions. These hints are used by the hardware to improve the accuracy of the link stack and the count cache.

As the branch instructions progress through the pipeline, eventually they become fully executed. At that point, the hardware determines whether the predicted target address and/or the direction of the branch matches the actual outcome of the branch. If the prediction was incorrect, the hardware takes the appropriate actions with respect to flushing undesired instructions and results, redirecting the pipeline, and updating the branch prediction information in the branch history tables.

Although the overall performance of the machine is strongly dependent on these branch prediction mechanisms, a set of firmware-accessible mode bits is available to disable these features via scan initialization.

3.1.2.3 Instruction Cache Block Touch Hint

The POWER8 core supports the instruction cache block touch instruction. Instead of bringing the data into the level 1 (L1) cache, it prefetches the data into the level 2 (L2) cache, which works just like a data cache block touch for store (dcbtst) hint instruction.

3.1.2.4 Out-of-Order Execution and Instruction Flushes

The POWER8 processor uses out-of-order instruction execution. Instructions can be speculative on a predicted branch direction, or simply speculative beyond an instruction that might cause an interrupt condition. In the event of a misprediction or an interrupt, instructions from the mispredicted path and the results produced by those instructions are discarded, presenting the effect of sequentially executed instructions down the appropriate branch paths and precise exceptions as required.

3.1.2.5 Branch Processor Instructions with Undefined Results

The results of executing an invalid form of a branch instruction or an instance of a branch instruction for which the architecture specifies that some results are undefined are described as follows. Only results that differ from those specified by the architecture are described in the following list.

- Instructions with reserved fields
Bits in reserved fields including the z-bits in the BO field are ignored; the results of executing an instruction in which one or more of these bits are '1' is the same as if the bits were '0'.
- **bcctr** and **bcctrl** instructions
If BO[2] = '0', the contents of CTR, before any update, are used as the target address and for the test of the contents of CTR to resolve the branch. The contents of the CTR are then decremented and written back to the CTR.
- System call instructions (opcode 17)

Bits 30:31	Description
'00'	sc instruction
'01'	illegal instruction exception
'10'	sc instruction
'11'	sc instruction

3.1.3 Fixed-Point Processor

3.1.3.1 Fixed-Point Exception Register (XER)

The Power ISA defines XER[0:31] and XER[35:56] as reserved. A **mfker** returns the value as shown in *Table 3-1*.

In the POWER8 core, the XER is implemented in several parts:

- XER renamed fields F0:F3 (see *Table 3-1*) are stored in an architected register file (ARF). The ARF consists of latches that store the F0:F3 fields for each of eight threads. The ARF contains eight (one per thread) Transactional Memory (TM) copies of the F0:F3 fields.
- XER nonrenamed bits (for example, F3NR) other than SO are stored in latches. Access is scoreboard managed. One copy for each of eight threads and eight TM copies are shared by FX0 and FX1.
- SO architected bits stored in latches with accompanying state machine. One copy for each of the eight threads and eight TM copies are shared by FX0 and FX1.

Table 3-1. XER Bits and Fields (Sheet 1 of 2)

XER Bits	Name	Field	Read/Write Behavior
0:31	Reserved	Unimplemented	Returns zeros on mfker .
32	SO	F3NR	Set to '1' whenever OV = '1', except when mtker sets SO = '0' and OV = '1'. This bit can be set to '0' or '1' by mtker . A mfker instruction reads the bit contents.
33	OV	F0	Set to '0' or '1' by various fixed-point instructions with OE = '1' or by mtker . A mfker instruction reads the bit contents.



Table 3-1. XER Bits and Fields (Sheet 2 of 2)

XER Bits	Name	Field	Read/Write Behavior
34	CA	F1	Set to '0' or '1' by add-carrying, subtract-from carrying and shift-right algebraic type instructions, and by mtxer . A mfker instruction reads the bit contents.
35:43	Reserved		Returns zeros on mfker
44:56	Reserved	F3NR	Written by mtxer . mfker reads the bit contents.
57:63	String length	F3NR	String length field used lswx and stswx . Written by mtxer . A mfker instruction reads the bit contents.

3.1.4 Storage Access Alignment Support Overview

Most storage accesses are performed without software intervention (such as, an alignment interrupt). The relative performance of these accesses depends to some degree on their alignment. In many cases, unaligned storage accesses are handled with performance equivalent to aligned accesses. However, in some cases the POWER8 processor is forced to break unaligned accesses into multiple internal operations. Further, because effective address alignment for storage references cannot be determined until execution time, and the POWER8 processor has dataflow-oriented execution pipelines that do not support iteration, some unaligned storage accesses actually cause a pipeline flush to allow a microcoded emulation of the instruction.

3.1.4.1 Misaligned Flushes

The **LSU** initiates a misaligned flush for the following conditions (See *Table 3-2* on page 46):

- Load crossing a 128-byte cache-line boundary and one of the cache-line misses.
- Load crossing a 32-byte sector boundary with either sector having an L1 D-cache miss and a 32-byte reload occurs instead of a 64-byte reload.
- Load/store crossing a 4 KB small page boundary.

Additionally, a misaligned flush is initiated in Data Address Watchpoint Register (DAWR) mode when the following conditions occur (see *Table 3-3* on page 54):

- Load crossing doubleword boundary when DAWR[63] = '1'
 - The following instructions are never considered to be crossing a doubleword boundary:
lq, lqarx, lfdp, lfdpx, lvebx, lvehx, lviewx, lvsl, lvsr, lvx, lvxl
- Store crossing doubleword boundary when DAWR[62] = '1'
 - The following instructions are never considered to be crossing a doubleword boundary:
stq, stqcx., stfdp, stfdpx, stvebx, stvehx, stviewx, stvx, stvxl

Note: If both a misaligned flush condition and an alignment interrupt condition are present, the alignment interrupt has precedence.

3.1.4.2 Alignment Interrupts

The LSU reports an alignment interrupt for the following conditions:

- Instruction not on a natural alignment boundary
 - Halfword boundary:
 - **Iharx**
 - **sthcx.**
 - Word boundary:
 - **Imw**, **Iwarx**,
 - **stmw**, **stwcx**
 - Doubleword boundary:
 - **Idarx**, **stdcx**
 - Quadword boundary:
 - **Ifdp**, **Ifdpx**, **Iq**, **stfdp**, **stfdpx**, **stq**, **stqcx**, **Iqarx**
- Little-endian mode
 - **Imw**, **Iswi**, **Iswx**
 - **stmw**, **stswi**, **stswx**
- Caching-inhibited storage
 - **Ifdp**, **Ifdpx**, **Imw**, **Iswi**, **Iswx**
 - **dcbz**, **stfdp**, **stfdpx**, **stmw**, **stswi**, **stswx**
 - Any load or store not on natural alignment boundary:
 - Halfword boundary:
 - **Iha**, **Ihau**, **Ihaux**, **Ihax**, **Ihbrx**, **Ihz**, **Ihzcix**, **Ihzu**, **Ihzux**, **Ihzx**
 - **sthbrx**, **sth**, **sthcix**, **sthu**, **sthux**, **sthx**
 - Word boundary:
 - **Ifiwax**, **Ifiwzx**, **Ifs**, **Ifsu**, **Ifsux**, **Ifsx**, **Iwa**, **Iwaux**, **Iwax**, **Iwbrx**, **Iwz**, **Iwzcix**, **Iwzu**, **Iwzux**, **Iwzx**, **Ixswux**, **Ixswx**
 - **stfiwx**, **stfs**, **stfsu**, **stfsux**, **stfsx**, **stwbrx**, **stw**, **stwcix**, **stwu**, **stwux**, **stwx**, **stxswux**, **stxswx**
 - Doubleword boundary:
 - **Id**, **Idbrx**, **Idcix**, **Idu**, **Idux**, **Idx**, **Ifd**, **Ifdi**, **Ifdux**, **Ifdx**, **Ixsdx**
 - **std**, **stdbrx**, **stdcix**, **stdu**, **stdux**, **stdx**, **stfd**, **stfdi**, **stfdux**, **stfdx**, **stxsdx**
 - Quadword boundary:
 - **Ixvdsx**, **Ixvd2x**, **Ixvw4x**, **stxvd2x**, **stxvw4x**

Note: A data storage interrupt (DSI) is reported on the following instructions to caching-inhibited storage regardless of the address alignment: **Ibarx**, **Iharx**, **Iwarx**, **Idarx**, **Iqarx**, **stbcx.**, **sthcx.**, **stwcx.**, **stdcx.**, **stqcx.**, **Iq**, **stq**



Table 3-2. Operand Alignment Effects on Performance (Non-Watchpoint Mode) (Sheet 1 of 8)

Power ISA Instructions	EA Alignment	Caching-Inhibited Alignment Interrupt?	BE: Takes Alignment Interrupt?	BE: Handled by Microcode?	LE: Takes Alignment Interrupt?	LE: Handled by Microcode?	DAWR Match on <u>QW</u> or <u>DW</u> , see other tab	DAR Value if DSI (non-DAWR) or Alignment Interrupt on First Page	DAR Value if DSI (non-DAWR) or Alignment Interrupt on Second Page
All loads/stores caching inhibited, not-naturally aligned		DSI or Alignment (see above for specific instructions)	Yes	No	Yes	No			
stxvw4x, stxvd2x	QW		No	No	No	No	N/A	EA	N/A
	DW		No	4K crossing	No	4K crossing	N/A	EA	EA + 16 ucode
	Even word (same as DW)		No	4K crossing	No	4K crossing	N/A	EA	EA + 16 ucode
	Odd word	Not naturally aligned	No	4K crossing	No	4K crossing	N/A	EA	EA + 16 ucode
	Non-word	Not naturally aligned	No	4K crossing	No	4K crossing	N/A	EA	EA + 16 ucode
stxswx	QW		No	No	No	No	N/A	EA	N/A
	DW	Not naturally aligned	No	No	No	No	N/A	EA	N/A
	Even word (same as DW)	Not naturally aligned	No	No	No	No	N/A	EA	N/A
	Odd word	Not naturally aligned	No	No	No	No	N/A	EA	N/A
	Non-word	Not naturally aligned	No	4K crossing	No	4K crossing	N/A	EA	N/A
stxsdx	QW		No	No	No	No	N/A	EA	N/A
	DW		No	No	No	No	N/A	EA	N/A
	Even word (same as DW)		No	No	No	No	N/A	EA	N/A
	Odd word	Not naturally aligned	No	4K crossing	No	4K crossing	N/A	EA	EA + 4
	Non-word	Not naturally aligned	No	4K crossing	No	4K crossing	N/A	EA	EA + 4



Table 3-2. Operand Alignment Effects on Performance (Non-Watchpoint Mode) (Sheet 2 of 8)

Power ISA Instructions	EA Alignment	Caching-Inhibited Alignment Interrupt?	BE: Takes Alignment Interrupt?	BE: Handled by Microcode?	LE: Takes Alignment Interrupt?	LE: Handled by Microcode?	DAWR Match on <u>QW</u> or <u>DW</u> , see other tab	DAR Value if DSI (non-DAWR) or Alignment Interrupt on First Page	DAR Value if DSI (non-DAWR) or Alignment Interrupt on Second Page
stvbx, stvehx, stvwx	QW		No	No	No	No	N/A	EA as defined by ISA	N/A
	DW		No	No	No	No	N/A	EA as defined by ISA	N/A
	Even word (same as DW)		No	No	No	No	N/A	EA as defined by ISA	N/A
	Odd word	Not naturally aligned	No	No	No	No	N/A	EA as defined by ISA	N/A
	Non-word	Not naturally aligned	No	No	No	No	N/A	EA as defined by ISA	N/A
stvx(l)	QW		No	No	No	No	N/A	EA as defined by ISA	N/A
	DW	Not naturally aligned	No	No	No	No	N/A	EA as defined by ISA	N/A
	Even word (same as DW)	Not naturally aligned	No	No	No	No	N/A	EA as defined by ISA	N/A
	Odd word	Not naturally aligned	No	No	No	No	N/A	EA as defined by ISA	N/A
	Non-word	Not naturally aligned	No	No	No	No	N/A	EA as defined by ISA	N/A
stq	QW	DSI	No	No	No	No	N/A	EA	N/A
	DW	DSI	Yes	No	Yes	No	N/A	EA	N/A
	Even word (same as DW)	DSI	Yes	No	Yes	No	N/A	EA	N/A
	Odd word	DSI	Yes	No	Yes	No	N/A	EA	N/A
	Non-word	DSI	Yes	No	Yes	No	N/A	EA	N/A
stfdp(x)	QW	Always	No	No	No	No	N/A	EA	N/A
	DW	Always	Yes	No	Yes	No	N/A	EA	N/A
	Even word (same as DW)	Always	Yes	No	Yes	No	N/A	EA	N/A
	Odd word	Always	Yes	No	Yes	No	N/A	EA	N/A
	Non-word	Always	Yes	No	Yes	No	N/A	EA	N/A



Table 3-2. Operand Alignment Effects on Performance (Non-Watchpoint Mode) (Sheet 3 of 8)

Power ISA Instructions	EA Alignment	Caching-Inhibited Alignment Interrupt?	BE: Takes Alignment Interrupt?	BE: Handled by Microcode?	LE: Takes Alignment Interrupt?	LE: Handled by Microcode?	DAWR Match on <u>QW</u> or <u>DW</u> , see other tab	DAR Value if DSI (non-DAWR) or Alignment Interrupt on First Page	DAR Value if DSI (non-DAWR) or Alignment Interrupt on Second Page
stfd	QW		No	No	No	No	N/A	EA	N/A
	DW		No	No	No	No	N/A	EA	N/A
	Even word (same as DW)		No	No	No	No	N/A	EA	N/A
	Odd word	Not naturally aligned	No	4K crossing	No	4K crossing	N/A	EA	E/A first byte in second page
	Non-word	Not naturally aligned	No	4K crossing	No	4K crossing	N/A	EA	E/A first byte in second page
stfs, stfliawx	QW		No	No	No	No	N/A	EA	N/A
	DW		No	No	No	No	N/A	EA	N/A
	Even word (same as DW)		No	No	No	No	N/A	EA	N/A
	Odd word		No	No	No	No	N/A	EA	N/A
	Non-word	Not naturally aligned	No	4K crossing	No	4K crossing	N/A	EA	E/A first byte in second page
stxsspx, stxsiwx	QW		No	No	No	No	N/A	EA	N/A
	DW		No	No	No	No	N/A	EA	N/A
	Even word (same as DW)		No	No	No	No	N/A	EA	N/A
	Odd word		No	No	No	No	N/A	EA	N/A
	Non-word	Not naturally aligned	No	4K crossing	No	4K crossing	N/A	EA	E/A first byte in second page
stswi, stswx	QW	Always	No	Yes	Yes	No	N/A	EA	E/A first byte in second page
	DW	Always	No	Yes	Yes	No	N/A	EA	E/A first byte in second page
	Even word (same as DW)	Always	No	Yes	Yes	No	N/A	EA	E/A first byte in second page
	Odd word	Always	No	Yes	Yes	No	N/A	EA	E/A first byte in second page
	Non-word	Always	No	Yes	Yes	No	N/A	EA	E/A first byte in second page



Table 3-2. Operand Alignment Effects on Performance (Non-Watchpoint Mode) (Sheet 4 of 8)

Power ISA Instructions	EA Alignment	Caching-Inhibited Alignment Interrupt?	BE: Takes Alignment Interrupt?	BE: Handled by Microcode?	LE: Takes Alignment Interrupt?	LE: Handled by Microcode?	DAWR Match on <u>QW</u> or <u>DW</u> , see other tab	DAR Value if DS1 (non-DAWR) or Alignment Interrupt on First Page	DAR Value if DS1 (non-DAWR) or Alignment Interrupt on Second Page
stmw	QW	Always	No	Yes	Yes	No	N/A	EA	E/A first byte in second page
	DW	Always	No	Yes	Yes	No	N/A	EA	E/A first byte in second page
	Even word (same as DW)	Always	No	Yes	Yes	No	N/A	EA	E/A first byte in second page
	Odd word	Always	No	Yes	Yes	No	N/A	EA	E/A first byte in second page
	Non-word	Always	Yes	No	Yes	No	N/A	EA	N/A
lxvd2x, lxvw4x	QW		No	No	No	No	N/A	EA	N/A
	DW		No	4K crossing OR 32-byte crossing L1 miss	No	4K crossing OR 32-byte crossing L1 miss	N/A	EA	EA + 16
	Even word (same as DW)		No	4K crossing OR 32-byte crossing L1 miss	No	4K crossing OR 32-byte crossing L1 miss	N/A	EA	EA + 16
	Odd word	Not naturally aligned	No	4K crossing OR 32-byte crossing L1 miss	No	4K crossing OR 32-byte crossing L1 miss	N/A	EA	EA + 16
	Non-word	Not naturally aligned	No	4K crossing OR 32-byte crossing L1 miss	No	4K crossing OR 32-byte crossing L1 miss	N/A	EA	EA + 16
lxswx	QW		No	No	No	No	N/A	EA	N/A
	DW	Not naturally aligned	No	No	No	No	N/A	EA	N/A
	Even word (same as DW)	Not naturally aligned	No	No	No	No	N/A	EA	N/A
	Odd word	Not naturally aligned	No	No	No	No	N/A	EA	E/A first byte in second page
	Non-word	Not naturally aligned	No	4K crossing OR 32-byte crossing L1 miss	No	4K crossing OR 32-byte crossing L1 miss	N/A	EA	E/A first byte in second page



Table 3-2. Operand Alignment Effects on Performance (Non-Watchpoint Mode) (Sheet 5 of 8)

Power ISA Instructions	EA Alignment	Caching-Inhibited Alignment Interrupt?	BE: Takes Alignment Interrupt?	BE: Handled by Microcode?	LE: Takes Alignment Interrupt?	LE: Handled by Microcode?	DAWR Match on <u>QW</u> or <u>DW</u> , see other tab	DAR Value if DSI (non-DAWR) or Alignment Interrupt on First Page	DAR Value if DSI (non-DAWR) or Alignment Interrupt on Second Page
lxvdsx lxsdx	QW		No	No	No	No	N/A	EA	N/A
	DW	Not naturally aligned	No	No	No	No	N/A	EA	N/A
	Even word (same as DW)	Not naturally aligned	No	No	No	No	N/A	EA	N/A
	Odd word	Not naturally aligned	No	4K crossing OR 32-byte crossing L1 miss	No	4K crossing OR 32-byte crossing L1 miss	N/A	EA	EA + 4
	Non-word	Not naturally aligned	No	4K crossing OR 32-byte crossing L1 miss	No	4K crossing OR 32-byte crossing L1 miss	N/A	EA	EA + 4
lvebx , lvehx , lviewx	QW		No	No	No	No	N/A	EA as defined by ISA	N/A
	DW		No	No	No	No	N/A	EA as defined by ISA	N/A
	Even word (same as DW)		No	No	No	No	N/A	EA as defined by ISA	N/A
	Odd word		No	No	No	No	N/A	EA as defined by ISA	N/A
	Non-word	Not naturally aligned	No	No	No	No	N/A	EA as defined by ISA	N/A
lxv(l)	QW		No	No	No	No	N/A	EA as defined by ISA	N/A
	DW	Not naturally aligned	No	No	No	No	N/A	EA as defined by ISA	N/A
	Even word (same as DW)	Not naturally aligned	No	No	No	No	N/A	EA as defined by ISA	N/A
	Odd word	Not naturally aligned	No	No	No	No	N/A	EA as defined by ISA	N/A
	Non-word	Not naturally aligned	No	No	No	No	N/A	EA as defined by ISA	N/A

Table 3-2. Operand Alignment Effects on Performance (Non-Watchpoint Mode) (Sheet 6 of 8)

Power ISA Instructions	EA Alignment	Caching-Inhibited Alignment Interrupt?	BE: Takes Alignment Interrupt?	BE: Handled by Microcode?	LE: Takes Alignment Interrupt?	LE: Handled by Microcode?	DAWR Match on <u>QW</u> or <u>DW</u> , see other tab	DAR Value if DSI (non-DAWR) or Alignment Interrupt on First Page	DAR Value if DSI (non-DAWR) or Alignment Interrupt on Second Page
lq	QW	DSI	No	No	No	No	N/A	EA	N/A
	DW	DSI	Yes	No	Yes	No	N/A	EA	N/A
	Even word (same as DW)	DSI	Yes	No	Yes	No	N/A	EA	N/A
	Odd word	DSI	Yes	No	Yes	No	N/A	EA	N/A
	Non-word	DSI	Yes	No	Yes	No	N/A	EA	N/A
ldp(x)	QW	Always	No	No	No	No	N/A	EA	N/A
	DW	Always	Yes	No	Yes	No	N/A	EA	N/A
	Even word (same as DW)	Always	Yes	No	Yes	No	N/A	EA	N/A
	Odd word	Always	Yes	No	Yes	No	N/A	EA	N/A
	Non-word	Always	Yes	No	Yes	No	N/A	EA	N/A
dcbz	QW	Always	No	No	No	No	N/A	EA	N/A
	DW	Always	No	No	No	No	N/A	EA	N/A
	Even word (same as DW)	Always	No	No	No	No	N/A	EA	N/A
	Odd word	Always	No	No	No	No	N/A	EA	N/A
	Non-word	Always	No	No	No	No	N/A	EA	N/A
lqarx	QW	DSI	No	No	No	No	N/A	EA	N/A
	DW	DSI	Yes	No	Yes	No	N/A	EA	N/A
	Even word (same as DW)	DSI	Yes	No	Yes	No	N/A	EA	N/A
	Odd word	DSI	Yes	No	Yes	No	N/A	EA	N/A
	Non-word	DSI	Yes	No	Yes	No	N/A	EA	N/A
ldarx	QW	DSI	No	No	No	No	N/A	EA	N/A
	DW	DSI	No	No	No	No	N/A	EA	N/A
	Even word (same as DW)	DSI	No	No	No	No	N/A	EA	N/A
	Odd word	DSI	Yes	No	Yes	No	N/A	EA	N/A
	Non-word	DSI	Yes	No	Yes	No	N/A	EA	N/A



Table 3-2. Operand Alignment Effects on Performance (Non-Watchpoint Mode) (Sheet 7 of 8)

Power ISA Instructions	EA Alignment	Caching-Inhibited Alignment Interrupt?	BE: Takes Alignment Interrupt?	BE: Handled by Microcode?	LE: Takes Alignment Interrupt?	LE: Handled by Microcode?	DAWR Match on <u>QW</u> or <u>DW</u> , see other tab	DAR Value if DSI (non-DAWR) or Alignment Interrupt on First Page	DAR Value if DSI (non-DAWR) or Alignment Interrupt on Second Page
lwarx	QW	DSI	No	No	No	No	N/A	EA	N/A
	DW	DSI	No	No	No	No	N/A	EA	N/A
	Even word (same as DW)	DSI	No	No	No	No	N/A	EA	N/A
	Odd word	DSI	No	No	No	No	N/A	EA	N/A
	Non-word	DSI	Yes	No	Yes	No	N/A	EA	N/A
lhaxr	QW	DSI	No	No	No	No	N/A	EA	N/A
	DW	DSI	No	No	No	No	N/A	EA	N/A
	Even word (same as DW)	DSI	No	No	No	No	N/A	EA	N/A
	Odd word	DSI	No	No	No	No	N/A	EA	N/A
	halfword	DSI	No	No	No	No	N/A	EA	N/A
	non half-word	DSI	Yes	No	Yes	No	N/A	EA	N/A
lxsspx, lxsiwax, lxsiwzx	QW		No	No	No	No	N/A	EA	N/A
	DW		No	No	No	No	N/A	EA	N/A
	Even word (same as DW)		No	No	No	No	N/A	EA	N/A
	Odd word		No	No	No	No	N/A	EA	N/A
	Non-word	Not naturally aligned	No	4K crossing OR 32-byte crossing L1 miss	No	4K crossing OR 32-byte crossing L1 miss	N/A	EA	N/A
lswi, lswx	QW	Always	No	Yes	Yes	No	N/A	EA	N/A
	DW	Always	No	Yes	Yes	No	N/A	EA	N/A
	Even word (same as DW)	Always	No	Yes	Yes	No	N/A	EA	N/A
	Odd word	Always	No	Yes	Yes	No	N/A	EA	N/A
	Non-word	Always	No	Yes	Yes	No	N/A	EA	N/A



Table 3-2. Operand Alignment Effects on Performance (Non-Watchpoint Mode) (Sheet 8 of 8)

Power ISA Instructions	EA Alignment	Caching-Inhibited Alignment Interrupt?	BE: Takes Alignment Interrupt?	BE: Handled by Microcode?	LE: Takes Alignment Interrupt?	LE: Handled by Microcode?	DAWR Match on <u>QW</u> or <u>DW</u> , see other tab	DAR Value if DSI (non-DAWR) or Alignment Interrupt on First Page	DAR Value if DSI (non-DAWR) or Alignment Interrupt on Second Page
Imw	QW	Always	No	Yes	Yes	No	N/A	EA	N/A
	DW	Always	No	Yes	Yes	No	N/A	EA	N/A
	Even word (same as DW)	Always	No	Yes	Yes	No	N/A	EA	N/A
	Odd word	Always	No	Yes	Yes	No	N/A	EA	N/A
	Non-word	Always	Yes	No	Yes	No	N/A	EA	N/A



Table 3-3. Operand Alignment Effects on Performance (Watchpoint Mode) (Sheet 1 of 9)

Power ISA Instructions	EA Alignment	Caching Inhibited Alignment Interrupt	BE: Takes Alignment Interrupt?	BE: Handled by Microcode?	LE: Takes Alignment Interrupt?	LE: Handled by Microcode?	DAWR Match on QW or DW	DAR Value if DSI (non-DAWR) or Alignment Interrupt on First Page	DAR Value if DSI (non-DAWR) or Alignment Interrupt on Second Page
All loads/stores Caching Inhibited not-naturally aligned		Always	Yes	No	Yes	No			
All DW operations or less (word, hw)	crossing a DW boundary	Always	No	yes	No	yes			
All QW operations	crossing a QW boundary	DSI or Alignment (see above for specific instructions)	No	yes	No	yes			
stxvw4x, stxvd2x									
	QW		No	No	No	Yes	DW	EA	N/A
	DW		No	Yes	No	Yes	N/A	EA	EA + 16 ucode
	Even word (same as DW)		No	Yes	No	Yes	N/A	EA	EA + 16 ucode
	Odd word	Not naturally aligned	No	Yes	No	Yes	N/A	EA	EA + 16 ucode
	Non-word	Not naturally aligned	No	Yes	No	Yes	N/A	EA	EA + 16 ucode
stxswx							DW		
	QW		No	No	No	No	N/A	EA	N/A
	DW	Not naturally aligned	No	No	No	No	N/A	EA	N/A
	Even word (same as DW)	Not naturally aligned	No	No	No	No	N/A	EA	N/A
	Odd word	Not naturally aligned	No	No	No	No	N/A	EA	N/A
	Non-word	Not naturally aligned	No	4K crossing	No	4K crossing	N/A	EA	E/A first byte in second page

Table 3-3. Operand Alignment Effects on Performance (Watchpoint Mode) (Sheet 2 of 9)

Power ISA Instructions	EA Alignment	Caching Inhibited Alignment Interrupt	BE: Takes Alignment Interrupt?	BE: Handled by Microcode?	LE: Takes Alignment Interrupt?	LE: Handled by Microcode?	DAWR Match on QW or DW	DAR Value if DSI (non-DAWR) or Alignment Interrupt on First Page	DAR Value if DSI (non-DAWR) or Alignment Interrupt on Second Page
stxsdx							DW		
	QW		No	No	No	No	N/A	EA	N/A
	DW		No	No	No	No	N/A	EA	N/A
	Even word (same as DW)		No	No	No	No	N/A	EA	N/A
	Odd word	Not naturally aligned	No	4K crossing	No	4K crossing	N/A	EA	E/A first byte in second page
	Non-word	Not naturally aligned	No	4K crossing	No	4K crossing	N/A	EA	E/A first byte in second page
stvebx, stvehx, stvewx							DW		
	QW		No	No	No	No	N/A	EA as defined by arch	N/A
	DW		No	No	No	No	N/A	EA as defined by arch	N/A
	Even word (same as DW)		No	No	No	No	N/A	EA as defined by arch	N/A
	Odd word	Not naturally aligned	No	No	No	No	N/A	EA as defined by arch	N/A
	Non-word	Not naturally aligned	No	No	No	No	N/A	EA as defined by arch	N/A
stvx(l)							QW		
	QW		No	No	No	No	N/A	EA as defined by arch	N/A
	DW	Not naturally aligned	No	No	No	No	N/A	EA as defined by arch	N/A
	Even word (same as DW)	Not naturally aligned	No	No	No	No	N/A	EA as defined by arch	N/A
	Odd word	Not naturally aligned	No	No	No	No	N/A	EA as defined by arch	N/A
	Non-word	Not naturally aligned	No	No	No	No	N/A	EA as defined by arch	N/A



Table 3-3. Operand Alignment Effects on Performance (Watchpoint Mode) (Sheet 3 of 9)

Power ISA Instructions	EA Alignment	Caching Inhibited Alignment Interrupt	BE: Takes Alignment Interrupt?	BE: Handled by Microcode?	LE: Takes Alignment Interrupt?	LE: Handled by Microcode?	DAWR Match on QW or DW	DAR Value if DSI (non-DAWR) or Alignment Interrupt on First Page	DAR Value if DSI (non-DAWR) or Alignment Interrupt on Second Page
stq							QW		
	QW	DSI	No	No	No	No	N/A	EA	N/A
	DW	DSI	Yes	No	Yes	No	N/A	EA	N/A
	Even word (same as DW)	DSI	Yes	No	Yes	No	N/A	EA	N/A
	Odd word	DSI	Yes	No	Yes	No	N/A	EA	N/A
	Non-word	DSI	Yes	No	Yes	No	N/A	EA	N/A
stfdp(x)							QW		
	QW	Always	No	No	No	No	N/A	EA	N/A
	DW	Always	Yes	No	Yes	No	N/A	EA	N/A
	Even word (same as DW)	Always	Yes	No	Yes	No	N/A	EA	N/A
	Odd word	Always	Yes	No	Yes	No	N/A	EA	N/A
	Non-word	Always	Yes	No	Yes	No	N/A	EA	N/A
stfd							DW		
	QW		No	No	No	No	N/A	EA	N/A
	DW		No	No	No	No	N/A	EA	N/A
	Even word (same as DW)		No	No	No	No	N/A	EA	N/A
	Odd word	Not naturally aligned	No	4K crossing	No	4K crossing	N/A	EA	E/A first byte in second page
	Non-word	Not naturally aligned	No	4K crossing	No	4K crossing	N/A	EA	E/A first byte in second page
stfs, stfliawx							DW		
	QW		No	No	No	No	N/A	EA	N/A
	DW		No	No	No	No	N/A	EA	N/A
	Even word (same as DW)		No	No	No	No	N/A	EA	N/A
	Odd word		No	No	No	No	N/A	EA	N/A
	Non-word	Not naturally aligned	No	4K crossing	No	4K crossing	N/A	EA	E/A first byte in second page

Table 3-3. Operand Alignment Effects on Performance (Watchpoint Mode) (Sheet 4 of 9)

Power ISA Instructions	EA Alignment	Caching Inhibited Alignment Interrupt	BE: Takes Alignment Interrupt?	BE: Handled by Microcode?	LE: Takes Alignment Interrupt?	LE: Handled by Microcode?	DAWR Match on QW or DW	DAR Value if DSI (non-DAWR) or Alignment Interrupt on First Page	DAR Value if DSI (non-DAWR) or Alignment Interrupt on Second Page
stxsspx, stxsiwx							DW		
	QW		No	No	No	No	N/A	EA	N/A
	DW		No	No	No	No	N/A	EA	N/A
	Even word (same as DW)		No	No	No	No	N/A	EA	N/A
	Odd word		No	No	No	No	N/A	EA	N/A
	Non-word	Not naturally aligned	No	4K crossing	No	4K crossing	N/A	EA	E/A first byte in second page
stswi, stswx							DW		
	QW	Always	No	Yes	Yes	No	N/A	EA	E/A first byte in second page
	DW	Always	No	Yes	Yes	No	N/A	EA	E/A first byte in second page
	Even word (same as DW)	Always	No	Yes	Yes	No	N/A	EA	E/A first byte in second page
	Odd word	Always	No	Yes	Yes	No	N/A	EA	E/A first byte in second page
	Non-word	Always	No	Yes	Yes	No	N/A	EA	E/A first byte in second page
stmw							DW		
	QW	Always	No	Yes	Yes	No	N/A	EA	E/A first byte in second page
	DW	Always	No	Yes	Yes	No	N/A	EA	E/A first byte in second page
	Even word (same as DW)	Always	No	Yes	Yes	No	N/A	EA	E/A first byte in second page
	Odd word	Always	No	Yes	Yes	No	N/A	EA	E/A first byte in second page
	Non-word	Always	Yes	No	Yes	No	N/A	EA	N/A



Table 3-3. Operand Alignment Effects on Performance (Watchpoint Mode) (Sheet 5 of 9)

Power ISA Instructions	EA Alignment	Caching Inhibited Alignment Interrupt	BE: Takes Alignment Interrupt?	BE: Handled by Microcode?	LE: Takes Alignment Interrupt?	LE: Handled by Microcode?	DAWR Match on QW or DW	DAR Value if DSI (non-DAWR) or Alignment Interrupt on First Page	DAR Value if DSI (non-DAWR) or Alignment Interrupt on Second Page
Ixvd2x, Ixvw4x							DW		
	QW		No	No	No	No	N/A	EA	N/A
	DW		No	4K crossing OR 32-byte crossing L1 miss	No	4K crossing OR 32-byte crossing L1 miss	N/A	EA	EA + 16
	Even word (same as DW)		No	4K crossing OR 32-byte crossing L1 miss	No	4K crossing OR 32-byte crossing L1 miss	N/A	EA	EA + 16
	Odd word	Not naturally aligned	No	4K crossing OR 32-byte crossing L1 miss	No	4K crossing OR 32-byte crossing L1 miss	N/A	EA	EA + 16
	Non-word	Not naturally aligned	No	4K crossing OR 32-byte crossing L1 miss	No	4K crossing OR 32-byte crossing L1 miss	N/A	EA	EA + 16
Ixswx							DW		
	QW		No	No	No	No	N/A	EA	N/A
	DW	Not naturally aligned	No	No	No	No	N/A	EA	N/A
	Even word (same as DW)	Not naturally aligned	No	No	No	No	N/A	EA	N/A
	Odd word	Not naturally aligned	No	No	No	No	N/A	EA	E/A first byte in second page
	Non-word	Not naturally aligned	No	4K crossing OR 32-byte crossing L1 miss	No	4K crossing OR 32-byte crossing L1 miss	N/A	EA	E/A first byte in second page



Table 3-3. Operand Alignment Effects on Performance (Watchpoint Mode) (Sheet 6 of 9)

Power ISA Instructions	EA Alignment	Caching Inhibited Alignment Interrupt	BE: Takes Alignment Interrupt?	BE: Handled by Microcode?	LE: Takes Alignment Interrupt?	LE: Handled by Microcode?	DAWR Match on QW or DW	DAR Value if DSI (non-DAWR) or Alignment Interrupt on First Page	DAR Value if DSI (non-DAWR) or Alignment Interrupt on Second Page
lxvdsx, lxsdx							DW		
	QW		No	No	No	No	N/A	EA	N/A
	DW	Not naturally aligned	No	No	No	No	N/A	EA	N/A
	Even word (same as DW)	Not naturally aligned	No	No	No	No	N/A	EA	N/A
	Odd word	Not naturally aligned	No	4K crossing OR 32-byte crossing L1 miss	No	4K crossing OR 32-byte crossing L1 miss	N/A	EA	EA + 4
	Non-word	Not naturally aligned	No	4K crossing OR 32-byte crossing L1 miss	No	4K crossing OR 32-byte crossing L1 miss	N/A	EA	EA of first byte in second dw
lvebx, lvehx, lviewx							DW		
	QW		No	No	No	No	N/A	EA as defined by arch	N/A
	DW		No	No	No	No	N/A	EA as defined by arch	N/A
	Even word (same as DW)		No	No	No	No	N/A	EA as defined by arch	N/A
	Odd word		No	No	No	No	N/A	EA as defined by arch	N/A
	Non-word	Not naturally aligned	No	No	No	No	N/A	EA as defined by arch	N/A



Table 3-3. Operand Alignment Effects on Performance (Watchpoint Mode) (Sheet 7 of 9)

Power ISA Instructions	EA Alignment	Caching Inhibited Alignment Interrupt	BE: Takes Alignment Interrupt?	BE: Handled by Microcode?	LE: Takes Alignment Interrupt?	LE: Handled by Microcode?	DAWR Match on QW or DW	DAR Value if DSI (non-DAWR) or Alignment Interrupt on First Page	DAR Value if DSI (non-DAWR) or Alignment Interrupt on Second Page
lvx(l)							QW		
	QW		No	No	No	No	N/A	EA as defined by arch	N/A
	DW	Not naturally aligned	No	No	No	No	N/A	EA as defined by arch	N/A
	Even word (same as DW)	Not naturally aligned	No	No	No	No	N/A	EA as defined by arch	N/A
	Odd word	Not naturally aligned	No	No	No	No	N/A	EA as defined by arch	N/A
	Non-word	Not naturally aligned	No	No	No	No	N/A	EA as defined by arch	N/A
lq							QW		
	QW	DSI	No	No	No	No	N/A	EA	N/A
	DW	DSI	Yes	No	Yes	No	N/A	EA	N/A
	Even word (same as DW)	DSI	Yes	No	Yes	No	N/A	EA	N/A
	Odd word	DSI	Yes	No	Yes	No	N/A	EA	N/A
	Non-word	DSI	Yes	No	Yes	No	N/A	EA	N/A
ldfp(x)							DW		
	QW	Always	No	No	No	No	N/A	EA	N/A
	DW	Always	Yes	No	Yes	No	N/A	EA	N/A
	Even word (same as DW)	Always	Yes	No	Yes	No	N/A	EA	N/A
	Odd word	Always	Yes	No	Yes	No	N/A	EA	N/A
	Non-word	Always	Yes	No	Yes	No	N/A	EA	N/A
dcbz							DW		
	QW	Always	No	No	No	No	N/A	EA	N/A
	DW	Always	No	No	No	No	N/A	EA	N/A
	Even word (same as DW)	Always	No	No	No	No	N/A	EA	N/A
	Odd word	Always	No	No	No	No	N/A	EA	N/A
	Non-word	Always	No	No	No	No	N/A	EA	N/A



Table 3-3. Operand Alignment Effects on Performance (Watchpoint Mode) (Sheet 8 of 9)

Power ISA Instructions	EA Alignment	Caching Inhibited Alignment Interrupt	BE: Takes Alignment Interrupt?	BE: Handled by Microcode?	LE: Takes Alignment Interrupt?	LE: Handled by Microcode?	DAWR Match on QW or DW	DAR Value if DSI (non-DAWR) or Alignment Interrupt on First Page	DAR Value if DSI (non-DAWR) or Alignment Interrupt on Second Page
lqarx							QW		
	QW	DSI	No	No	No	No	N/A	EA	N/A
	DW	DSI	Yes	No	Yes	No	N/A	EA	N/A
	Even word (same as DW)	DSI	Yes	No	Yes	No	N/A	EA	N/A
	Odd word	DSI	Yes	No	Yes	No	N/A	EA	N/A
	Non-word	DSI	Yes	No	Yes	No	N/A	EA	N/A
ldarx							DW		
	QW	DSI	No	No	No	No	N/A	EA	N/A
	DW	DSI	No	No	No	No	N/A	EA	N/A
	Even word (same as DW)	DSI	No	No	No	No	N/A	EA	N/A
	Odd word	DSI	Yes	No	Yes	No	N/A	EA	N/A
	Non-word	DSI	Yes	No	Yes	No	N/A	EA	N/A
lwarx							DW		
	QW	DSI	No	No	No	No	N/A	EA	N/A
	DW	DSI	No	No	No	No	N/A	EA	N/A
	Even word (same as DW)	DSI	No	No	No	No	N/A	EA	N/A
	Odd word	DSI	No	No	No	No	N/A	EA	N/A
	Non-word	DSI	Yes	No	Yes	No	N/A	EA	N/A
lharcx							DW		
	QW	DSI	No	No	No	No	N/A	EA	N/A
	DW	DSI	No	No	No	No	N/A	EA	N/A
	Even word (same as DW)	DSI	No	No	No	No	N/A	EA	N/A
	Odd word	DSI	No	No	No	No	N/A	EA	N/A
	halfword	DSI	No	No	No	No	N/A	EA	N/A
	non half-word	DSI	Yes	No	Yes	No	N/A	EA	N/A



Table 3-3. Operand Alignment Effects on Performance (Watchpoint Mode) (Sheet 9 of 9)

Power ISA Instructions	EA Alignment	Caching Inhibited Alignment Interrupt	BE: Takes Alignment Interrupt?	BE: Handled by Microcode?	LE: Takes Alignment Interrupt?	LE: Handled by Microcode?	DAWR Match on QW or DW	DAR Value if DSI (non-DAWR) or Alignment Interrupt on First Page	DAR Value if DSI (non-DAWR) or Alignment Interrupt on Second Page
lxsspx, lxsix-wax, lxi-wzx							DW		
	QW		No	No	No	No	N/A	EA	N/A
	DW		No	No	No	No	N/A	EA	N/A
	Even word (same as DW)		No	No	No	No	N/A	EA	N/A
	Odd word		No	No	No	No	N/A	EA	N/A
	Non-word	Not naturally aligned	No	4K crossing OR 32-byte crossing L1 miss	No	4K crossing OR 32-byte crossing L1 miss	N/A	EA	E/A first byte in second page
lswi, lswx							DW		
	QW	Always	No	Yes	Yes	No	N/A	EA	N/A
	DW	Always	No	Yes	Yes	No	N/A	EA	N/A
	Even word (same as DW)	Always	No	Yes	Yes	No	N/A	EA	N/A
	Odd word	Always	No	Yes	Yes	No	N/A	EA	N/A
	Non-word	Always	No	Yes	Yes	No	N/A	EA	N/A
lmw							DW		
	QW	Always	No	Yes	Yes	No	N/A	EA	N/A
	DW	Always	No	Yes	Yes	No	N/A	EA	N/A
	Even word (same as DW)	Always	No	Yes	Yes	No	N/A	EA	N/A
	Odd word	Always	No	Yes	Yes	No	N/A	EA	N/A
	Non-word	Always	Yes	Yes	Yes	No	N/A	EA	N/A



3.1.4.3 Fixed-Point Load Instructions

Most forms of unaligned load operations are executed entirely in hardware. If a basic load operation crosses a page boundary, and either page translation signals an exception condition, then when the interrupt occurs, it appears as though none of the load instruction has executed. This is not always the case for load multiple or load string instructions.

The “Load Algebraic”, “Load with Byte Reversal,” and “Load with Update” instructions might have greater latency than other load instructions. These instructions are implemented as a sequence of internal operations. Due to the dynamic scheduling and out-of-order execution capability of the processor, these effects are somewhat minimized. Also note that although these instructions are broken up in this manner, the effects are never visible from a programming model perspective.

With the exception of the “Load and Reserve” and “Load Quadword” instructions, any load from storage marked *caching inhibited* that is not aligned causes an alignment interrupt. The “Load and Reserve” and “Load Quadword” instructions instead cause a data storage interrupt when they attempt to access a caching inhibited page regardless of their alignment.

An attempt to execute a non-quadword-aligned **lq** or **lqarx** instruction to a *cacheable* address causes an alignment interrupt.

See *Section 3.1.4 Storage Access Alignment Support Overview* on page 44 for details.

3.1.4.4 Fixed-Point Store Instructions

Most forms of unaligned store operations are executed entirely in hardware. If a store operation crosses a page boundary and the second page translation signals an exception condition, then after the interrupt is taken, it appears as though none of the storage updates have occurred to either page. This is not always the case for store multiple or store string instructions.

With the exception of the “Store Conditional” and “Store Quadword” instructions, any store to storage marked *caching inhibited* that is not aligned causes an alignment interrupt. The “Store Conditional” and “Store Quadword” instructions instead cause a data storage interrupt when they attempt to store to a *caching inhibited* page regardless of their alignment.

An attempt to execute a non-quadword-aligned **stq** or **stqcx** instruction to a *cacheable* address causes an alignment interrupt.

See *Section 3.1.4 Storage Access Alignment Support Overview* on page 44 for details.

3.1.4.5 Fixed-Point Load and Store Multiple Instructions

Note: These instructions are provided for compatibility with legacy software. Software should use a sequence of load or store instructions for optimal performance.

The **lmw** instruction is executed in a manner such that up to two registers are loaded each cycle. Similarly, the **stmw** instruction is executed in a manner such that up to two registers are stored each cycle. The 40-entry store queue can accept up to two 8-byte stores per cycle; the cache can accept one 8-byte store per cycle. Because these instructions are emulated through the use of microcoded templates, after a small start-up penalty, they are processed at a rate of up to two registers per cycle.



Most forms of **Imw** and **stmw** instructions, even those that cross page and segment boundaries, are executed entirely in hardware. These instructions and the individual storage accesses associated with the instructions are not atomic. If a **stmw** crosses a page boundary, and the second page translation signals an exception condition, then after the interrupt is taken, it appears as though none, some, or all of the accesses to the first page have occurred, and none of the accesses to the second page have occurred. On the other hand, for the **Imw** instruction that cross a page boundary where the second page translation signals an exception condition, some of the target registers may not be updated.

An attempt to execute a non-word-aligned **Imw** or **stmw** causes an alignment interrupt.

An attempt to execute an **Imw** or **stmw** to storage marked cache inhibited causes an alignment interrupt.

See *Section 3.1.4 Storage Access Alignment Support Overview* on page 44 for details.

The architecture allows these instructions to be interrupted by certain types of asynchronous interrupts (external interrupts, decrementer interrupts, machine check interrupts, and system reset interrupts). However, the POWER8 core will not process an asynchronous interrupt in the middle of one of these instructions.

3.1.4.6 Fixed-Point Move Assist Instructions

Note: These instructions are provided for compatibility with legacy software. Software should use a sequence of load or store instructions for optimal performance.

The load string word instructions are executed in a manner such that up to two registers are loaded each cycle. Similarly, store string word instructions are executed in a manner such that up to two registers are stored each cycle. The 40-entry store queue can accept up to two 8-byte stores per cycle; the cache itself can only accept one 8-byte store per cycle.

Because the immediate forms of these instructions are implemented using microcoded templates, they incur a small start-up penalty. The X-form of the instructions contains a dependency on bits in the fixed-point XER register. If the dependent move assist x-form instruction is dispatched too soon after the instruction updating the XER, it will be flushed and re-executed.

Most “load string” and “store string” instructions that cross page or segment boundaries are executed entirely in hardware. If a “store string” crosses a page boundary and the second page translation signals an exception condition, then after the interrupt is taken, it appears as though none, some, or all of the accesses to the first page completed, and none of the accesses to the second page have occurred. On the other hand, for “load string” instructions that cross a page boundary where the second page translation signals an exception condition, all of the target registers have an undefined value.

If the storage operand of an **Iswi** is word aligned, the accesses are performed in an optimal manner. If the operands are so aligned, the accesses are performed in an optimal manner if the operand resides entirely within a 64-byte aligned block that is resident in the L1 D-cache or resides entirely within a 32-byte aligned block (which either hits or misses in the L1 D-cache). Although other “unaligned” string operations are supported in hardware, they can cause machine flushes and require long sequences of microcode. As a result, these types of “unaligned” string instructions might have significantly longer latencies. For additional information, see *Table 10-16 Instruction Latencies and Throughputs* on page 236.

An attempt to execute an **Iswi**, **Iswx**, **stswi**, or **stswx** instruction to storage marked cache inhibited causes an alignment interrupt.

See *Section 3.1.4 Storage Access Alignment Support Overview* on page 44 for details.



The architecture allows these instructions to be interrupted by certain types of asynchronous interrupts (external interrupts, decrementer interrupts, machine check interrupts, and system reset interrupts). However, the POWER8 core will not process an asynchronous interrupt in the middle of one of these instructions.

The architecture describes some preferred forms for the use of load and store string instructions. These preferred forms have no effect on the performance of the instructions in the POWER8 processor.

3.1.4.7 Integer Select (ISEL)

The POWER8 core implements the integer select instruction as defined in the Power ISA.

3.1.4.8 Fixed-Point Logical Instructions

The architecture defines the *preferred NOP* to be '**ori 0,0,0**'. In the POWER8 processor, this NOP form is recognized by the hardware and allowed to complete without taking any execution resources. This makes the instruction valuable for padding other instructions to achieve better alignment or better separation

3.1.4.9 Move To/From Special Purpose Register (SPR) Instructions

The POWER8 core supports problem state read access from the following optional special purpose registers (SPRs):

- PMC1 - Performance Monitor Counter 1
- PMC2 - Performance Monitor Counter 2
- PMC3 - Performance Monitor Counter 3
- PMC4 - Performance Monitor Counter 4
- PMC5 - Performance Monitor Counter 5
- PMC6 - Performance Monitor Counter 6
- MMCR0 - Monitor Mode Control Register 0
- MMCR1 - Monitor Mode Control Register 1
- SIAR - Sampled Instruction Address Register
- SDAR - Sampled Data Address Register

3.1.4.10 Move to Condition Register Fields Instruction

The architecture warns that updating a subset of the **CR** fields on an **mtcrf** instruction might have worse performance than updating all of the fields. In the POWER8 processor, both the **mtcrf** instruction and the **mfcr** instruction are emulated through the use of microcode templates. For best performance, software must use the single-field variants (**mtocrf** and **mfocrf**) of these instructions as described in the Power ISA.



3.1.4.11 Fixed-Point Invalid Forms and Undefined Conditions

The results of executing an invalid form of a fixed-point instruction or an instance of a fixed-point instruction for which the architecture specifies that some results are undefined are described below, for the cases in which executing an instruction does not cause an exception. The list below describes the results of executing invalid instruction forms on the POWER8 processor core.

- **Instruction with Reserved Fields**

Bits in reserved fields are ignored; the results of executing an instruction in which one or more reserved bits are '1' is the same as if the bits were '0'.

- **Load with Update Instructions (RA = 0)**

EA is placed into R0.

- **Load with Update Instructions (RA = RT)**

The storage operand addressed by EA is accessed. The displacement field is added to the data returned by the load and placed into RT.

- **Load Quadword Instruction (RTp is odd and RTp = RA)**

The POWER8 processor always takes a hypervisor emulation assistance interrupt anytime RTp is an odd register, RTp = RA (including when RA = 0) or RTp = RB for **Iq**.

- **Load Quadword And Reserve Indexed Instruction (RTp is odd, RTp = RA, RTp = RB)**

The POWER8 processor always takes a hypervisor emulation assistance interrupt anytime RTp is an odd register, RTp = RA (including when RA = 0) or RTp = RB for **Iqarx**.

- **Load Multiple Instructions (RA in the range of registers to be loaded)**

If an exception (for example, data storage or external) causes the execution of the instruction to be interrupted, the instruction is restarted, RA has been altered by the previous partial execution of the instruction, and RA is less than or greater than '0', the new contents of RA are used to compute EA.

- **Load Multiple Instructions (causing a misaligned access)**

For a load multiple word instruction, if the storage operand specified by EA is not a multiple of 4, an alignment interrupt is taken.

- **Load String Instructions (zero length string)**

RT is not altered.

- **Load String Instructions (RA and/or RB in the range of registers to be loaded)**

If RA and/or RB is in the range of registers to be loaded, the results are as follows.

Indexed Form: If RA = 0, let Rx be RB; otherwise let Rx be the register specified by the smaller of the two values in instruction fields RA and RB. If RT = Rx, no registers are loaded; otherwise, registers RT through RX - 1 are loaded as specified in the architecture (that is, only part of the storage operand is loaded).

Immediate Form: If RA = 0, the instruction is executed as if it were a valid form. If RA = RT, no registers are loaded; otherwise registers RT through RA - 1 are loaded as if the instruction were a valid form but specifying a shorter operand length.

- **Store with Update Instructions (RA = 0)**

EA is placed into R0.

- **Store Quadword and Store Quadword Conditional Instructions (RSp is odd)**

For the **stq** and **stqcx** instructions, the contents of RSp are stored into the words of storage addressed by EA and EA + 4 respectively. If RSp is odd, the low-order bit of the register number is considered to be '0' such that RSp - 1 and RSp are stored into the words in storage addressed by EA and EA + 4 respectively.

- **subfic, subfc, and subfco Instructions and their Rc = 1 forms**
If RA[0:15] = x'0000', XER[CA] reflects the carry-out of bit 16; otherwise, it reflects the carry-out of bit 40.
- **divw, divwo, divwu, and divwuo Instructions**
RT[0:31] is set to x'00000000'.
- **mulhw and mulhwu Instructions**
RT[0:31] contains the same result as RT[32:63].
- **Divide Instructions** (divide by zero)
If the divisor is 0, RT is set to zero. If Rc = '1' also, CR0 is set to '0010'.
- **Trap Word Immediate and Trap Word Instructions** (TO = (0b11110 | 0b11100))
- **Move To/From Special Purpose Register Instructions**
Table 3-8 on page 87 describes the read/write **mtspr** behavior for an SPR value that is not defined for the implementation.
- **Move From Time Base Instruction**
The **mftb** instruction is treated as an alias for the "mfspr Rx, 268" instruction. The results are the same as when executing an "mfspr Rx, 268" instruction.
- **Move From Condition Register Instruction**
The entire CR is copied into RT[32:63]. RT[0:31] is set to zero.
- **Move From One Condition Register Field Instruction** (only 1 bit of FXM set to '1')
Let n be the bit set to '1' in the FXM field. The CR field n is copied to RT[(4 × n + 32):(4 × n + 35)] and RT[(4 × n + 36):(4 × n + 39)]. The remaining bits are set to zero.
- **Move From One Condition Register Field Instruction** (multiple bits of FXM set to '1')
Let n be the first bit (from left to right) set to '1' in the FXM field. The CR field n is copied to RT[(4 × n + 32):(4 × n + 35)] and RT[(4 × n + 36):(4 × n + 39)]. The remaining bits are set to zero.

3.2 Floating-Point Processor (FP, VMX, and VSX)

The POWER8 VSU contains four double-precision floating-point units. Each of these units is optimized for fully pipelined double-precision multiply-add functionality. In addition, each unit is capable of performing the floating-point divide and square root instructions. The complex integer instructions from the VMX architecture are also executed on the VSU datapath.

The POWER8 VSU implements the VSX architecture, specifying 2-way DP or 4-way SP operations. Dependent floating-point instructions have a minimum issue-to-issue interval of six cycles. The vector SP throughput has improved because it is possible to execute two 4-way SIMD SP instructions per cycle.

Legacy BFU and VMX architectures are also fully supported in the POWER8 VSU.

3.2.1 Vector Single-Precision Bandwidth

In the POWER8 core, the double-precision FPU supports simultaneous execution of two vector single-precision operations. This increases the single-precision bandwidth of the POWER8 core to 16 FLOPs per cycle.

In the POWER8 core, the convert and estimate instructions are executed in a fully pipelined manner, as well as the increased bandwidth of the multiply-add and move instructions. From the floating-point instructions, only the divide and square-root instructions cannot be started every cycle.



The compares, min/max, and test-for-software divide/square-root instructions are now executed on the vector integer (XS) pipeline to take advantage of the shorter latency. Also, the move-from-FPSCR and move-to-FPSCR instructions are separated from the floating-point datapath.

3.2.2 IEEE Compliance

The POWER8 implementation of binary floating-point (BFP), decimal floating-point (DFP) and vector-scalar floating-point (VSX) architecture complies with the IEEE P754-2008 standard.

In the POWER8 processor core, the setting of the non-IEEE (FPSCR[NI]) bit has no effect on the instruction execution performance nor the results of the instruction.

3.2.3 Divide and Square-Root Latencies

All divide and square-root instructions can be executed independently on both pipes. The latencies of the instructions are listed in *Table 3-4*.

If the result of a double-precision divide instruction (**fdiv**, **xdivdp**, **xvdivdp**) is tiny, an additional iteration is necessary to denormalize the result. This increases the instruction latency by six cycles. For all other instructions, the latency does not depend on the input operands. For additional information, see *Table 10-16 Instruction Latencies and Throughputs* on page 236.

Table 3-4. Latencies of Floating-Point Divide/Square-Root Instructions

Instruction	Issue-to-Finish	Dependent Operations (issue-to-issue)	Next Independent Multicycle Operation on Same Pipe (issue-to-issue)	Number of Free Floating-Point Slots until Next Multicycle Can Start
fdiv xdivdp xvdivdp	35	32	26	18/26
fdivs xdivsp	29	26	20	14/20
xdivsp	31	28	22	12/22
fsqrt xssqrtdp xvsqrtdp	46	43	37	26/37
fsqrt xssqrts xvsqrts	34	31	25	19/25
xvsqrts	35	32	26	14/26

3.2.4 Early Forwarding Conditions

The back-to-back latency for all single-cycle floating-point instructions is six cycles. The 6-cycle result forwards the correct result only if both the result and the input operand are compatible floating-point values. For the following cases, the consumer instruction would misunderstand the data; instead it finishes with a flush request, triggering the ISU to eventually repeat the dependent instruction with operands from the register file:

- The forwarded result is the result of a convert-float-to-integer instruction.
- The forwarded result is used as an input operand for a convert-integer-to-float instruction.



- The forwarded result is a 32-bit floating-point value, but the input operand is a 64-bit floating-point value (for example, **xscvdpsp** to **fadd**).
- The forwarded result is a 64-bit floating-point value, but the input operand is a 32-bit floating-point value (for example, **xvaddirp** to **xvcvspuxds**).

The previous scenarios are rare in real programs, because they require mixing of incompatible data types. Hence, the performance loss is acceptable.

Note: It might make sense to execute a convert-float-to-integer instruction followed by a dependent convert-integer-to-float instruction. Because this also causes a flush request if the issue-to-issue-interval is six cycles, this scenario must be avoided by the compiler. (If the purpose is to round a floating-point number to an integral value, one of the atomic round-to-integral-value instructions must be used.)

There are out-of-range scenarios where the result of the instruction is not defined by the ISA. Because exponent MSBs might be dropped while packing the result into the target format, the value on the 6-cycle result (before packing) can be inconsistent with the value written to the target register (after packing). To guarantee that program results are independent of the issue behavior, the 6-cycle result is invalidated if there is such an inconsistency risk. Instructions that consume a result below via the 6-cycle result bus will finish with a flush request:

- The result of **xscvdpsp** with FPSCR[UE] = 1 and $0 < \text{IXBI} < 2^{-255}$ (an underflow will occur).
- The result of **xscvdpsp** with FPSCR[OE] = 1 and $\text{IXBI} \geq 2^{257}$ (an overflow will occur).

Instructions that consume the data below via the 6-cycle result might finish with a flush request. Whether or not it actually happens depends on the exponent value of the unnormalized result.

- The result of a single-precision arithmetic instruction with at least one out-of-range input operand, FPSCR[UE] = 1, and an underflow exception occurs.
- The result of a single-precision arithmetic instruction with at least one out-of-range input operand, FPSCR[OE] = 1, and an overflow exception occurs.

3.2.5 Floating-Point Exceptions

Precise floating-point exceptions are provided whenever either of the floating-point enabled exception mode bits (MSR[FE0], MSR[FE1]) are set. In all cases, the floating-point instructions are executed out-of-order (as required), and any resulting exceptions are sorted out at completion time. In some cases, due to the group-oriented instruction tracking scheme used, when an exception is detected, the hardware flushes the pipeline and re-dispatches the instructions individually to provide the precise exception. Because this only happens if an interrupt is to be taken, it does not represent a measurable slowdown in performance.

3.2.6 Floating-Point Load and Store Instructions

Most forms of unaligned floating-point storage accesses are executed entirely in hardware. In the POWER8 core, the load/store operations have been improved such that byte-aligned, 16-byte VSX load/store operations are executed efficiently in hardware. However, care must be taken that the unaligned VSX 16-byte storage accesses do not cross a 4K-page boundary, because doing so results in a micro-coded implementation with slower execution time. See *Section 3.1.4 Storage Access Alignment Support Overview* on page 44 for details.



3.2.7 Heterogeneous Precision Arithmetic

The following instructions are referred to as scalar single-precision arithmetic instructions:

- **fadds[.], xsaddsp, fsubs[.], xssubsp, fmuls[.], xsmulsp**
- **fmadds[.], xsmadd[am]sp, fmsubs[.], xsmsub[am]sp**
- **fnmadds[.], xsnmadd[am]sp, fnmsubs[.], xsnmsub[am]sp**
- **fsqrts[.], xssqrtsp, fdivs[.], xsdivsp**
- **fres[.], xsresp, frsqrtes[.], xsrsqrtesp**

3.2.7.1 NaN Propagation

If a single-precision arithmetic instruction propagates a NaN where any of the fraction bits [24:52] is nonzero, the resulting QNaN has all of the fraction bits [24:52] cleared to zero.

3.2.7.2 Square Root Overflow and Underflow

Due to the compacting nature of the square-root operation, the instructions **fsqrts**, **xssqrtsp**, **frsqrtes**, and **xsrsqrtesp** cannot underflow or overflow if their operands are representable in single-precision format. However, if the operand is not representable in single-precision format, an underflow or overflow can occur. This will be recorded in the FPSCR.

3.2.7.3 Hardware Behavior on Enabled Underflow and Enabled Overflow Exception

If FPSCR[UE] is enabled and an underflow occurs, the contents of the result register and FPSCR status are not defined for scalar SP instructions. The hardware takes the following actions:

1. Underflow exception is set, FPSCR[UX] = '1'.
2. The exponent of the normalized intermediate result is adjusted by adding 192.
3. The double-precision bias of 1023 is added to the exponent.
4. The biased exponent is reduced to 11 bits by ANDing with x'7FF'.
5. The adjusted rounded result is placed into the target FPR.
6. FPSCR[FPRF] is set to indicate Normalized Number.

If FPSCR[OE] is enabled and an overflow occurs, the contents of the result register and FPSCR status are not defined for scalar SP instructions. The hardware takes the following actions:

1. Overflow exception is set, FPSCR[OX] = '1'.
2. The exponent of the normalized intermediate result is adjusted by subtracting 192.
3. The double-precision bias of 1023 is added to the exponent.
4. The biased exponent is reduced to 11 bits by ANDing with x'7FF'.
5. The adjusted rounded result is placed into the target FPR.
6. FPSCR[FPRF] is set to indicate Normalized Number.



3.2.8 Handling of Denormal Single-Precision Values in Double-Precision Format

3.2.8.1 Producing Single-Precision Denorms

If a scalar single-precision instruction produces a denormal result, the FPU is not able to normalize the data after rounding, which is needed to compute the architected register file format. Instead the FPU sets the exponent bits to $897 = x'381'$, corresponding to an implied bit magnitude of 2^{-126} , and returns a denormal fraction. Additionally, a flag in the result register is set, which marks the 64-bit result single-precision denormal. This flag is called “dirty”. Dirty results can be produced by any of the following instructions:

- All scalar SP arithmetic instructions as listed above.
- The round-to-single-precision instructions **frsp[.]** and **xrsp**.
- The load single-precision instructions **lfs** and **lxsspx**.
- The floating-point select instruction **fsel[.]**.

Note: If an enabled underflow exception occurs, the arithmetic instructions, as well as the rounding instructions, return a normalized fraction. Hence, if FPSCR[UE] is set, an arithmetic or rounding instruction never sets the ‘dirty’ flag.

The **fsel[.]** instruction only produces a dirty result if the selected operand was dirty. Execution of **fsel[.]** does not depend on FPSCR[UE].

A single-precision load instruction always sets the dirty flag if single-precision denormal data are loaded into the register file regardless of FPSCR[UE].

Any instruction not listed previously cannot produce dirty results.

3.2.8.2 Consuming Single-Precision Denorms

The majority of the BFU and VSX double-precision format operations, including all stores, can handle dirty input operands and perform the required normalizing on-the-fly. Thus, inputting dirty data to one of these instructions has no side effect.

To ensure that a register does not contain dirty data, it is recommended to move the content of the register to itself using the VSX DP copy sign instruction

xvcpsgndp xi, xi, xi

This instruction normalizes a potentially dirty input and writes it back to the register file in 64-bit double-precision format.

If an instruction that cannot handle dirty input operands sees a dirty input operand, it generates a denormal input exception that results in a soft patch interrupt to the hypervisor with HSRR1[43] = ‘1’. The hypervisor can use the contents of the HEIR to examine the instruction that caused the interrupt and perform the previously mentioned **xvcpsgndp** for each of the source registers for that instruction. Alternatively, the hypervisor can simply perform the previous **xvcpsgndp** for all 64 VSR Registers. In either case, after the dirty data has been reformatted by the hypervisor software, the instruction can then be re-executed to produce the result. The VSX Scalar SP instructions, introduced in the POWER8 processor, can also write SP denorms into VSR Registers 0 - 31.



Instructions without Support of Dirty Inputs

The following instructions do not support dirty inputs and generate denormal input exceptions whenever any of the operands is marked dirty.

- Any instruction operating on 32-bit operands
- Any convert from integer
- The move to FPSCR instruction **mtfsf[.]**
- Any compare, test-for-software, or min/max instructions
- Any extract instructions: **fmr[g]oe]w**, **mfvsrd**, **mfvsrwz**
- Any DFU instructions
- Any VSX fixed-point arithmetic, logical, or permute instruction

If it is expected that a scalar single-precision workload produces denormal results quite often and also does single-precision compares frequently, it is recommended to precede the compare instructions with a normalizing instruction that does nothing, for example by using **fmr** or **xscpsgndp**. Alternatively, consider using scalar DP or vector SP code, where result precision and result format always match.

Table 10-17 Instructions that Trigger an Exception on Consumption of a Scalar SP Denorm on page 258 provides a complete list of VSU instructions that trigger an exception on consumption of a dirty operand.

3.2.9 Floating-Point Invalid Forms and Undefined Conditions

The results of executing an invalid form of a floating-point instruction or an instance of a floating-point instruction for which the architecture specifies that some results are undefined are described below for the cases when executing an instruction does not cause an exception.

- Scalar single-precision instructions with operands not representable in single-precision format
See *Section 3.2.7 Heterogeneous Precision Arithmetic* on page 70.
- Instruction with reserved fields
Bits in reserved fields are ignored. The results of executing an instruction in which one or more reserved bits are '1' is the same as if the bits were '0'.
- Load or Store Floating-Point with Update instructions (RA = 0)
EA is placed into R0.
- Floating-Point Store Single instructions (exponent < 874 and FRS[9:31] <> 0)
The value placed in storage is a 0 with the same sign as the value in the register.
- Scalar Floating-Point instructions
VSR_{64:127} is set to x'0000_0000_0000_0000'.
- Scalar Convert to Integer Word instructions (**xscvdpxws**, **xscvdpsxws**, **fctiwuz**, **fctiwz**, **ctiwu**, **fctiw**)
VSR[0:31] is set to VSR[32:63].
- VSX Scalar Convert From Double-Precision to Single-Precision (**xscvdpsp**, **xscvdpspn**)
VSR[32:63] is set to VSR[0:31].
- Scalar Convert to Integer instructions
FPSCR[FPRF] is set to '00000'.

- VSX Vector Convert From Double-Precision to Single-Precision (**xvcvdpdsp**),
VSX Vector Convert Double-Precision to Integer Word (**xvcvdpsxws**, **xvcvdpxws**),
VSX Vector Convert From Integer Doubleword to Single-Precision (**xvcvsxdsp**, **xvcvuxdsp**)

VSR[32:63] is set to VSR[0:31].
VSR[96:127] is set to VSR[64:95].

- Move from FPSCR instruction
FRT[0:63] is set to FPSCR[0:63] with the first 29 bits set to zero.
- Scalar Reciprocal Estimate instructions:
(**fre**, **fres**, **xsredp**, **xsresp**, **frsqrte**, **frsqrtes**, **xrsqrtep**, **xrsqrtesp**)
FPSCR[FR] and FPSCR[FI] are set to '0' and FPSCR[XX] is unchanged, even if an overflow exception occurs.
- VSX Vector Floating-Point Reciprocal Estimate instructions: **xvredp**, **xvresp**, **xvrsqrtep**, **xvrsqrtesp**
FPSCR[XX] is unchanged, even if an overflow exception occurs.
- Disabled Overflow Exception (OX = '1', OE = '0')
For divide and square root instructions, FPSCR[FR] is set to '1' if the result is rounded to $\pm\infty$, and set to '0' if the result is rounded to the largest representable number. For Scalar Reciprocal Estimate instructions, FPSCR[FR] is set to '0'. For all other instructions, FPSCR[FR] is set to '1' if a disabled overflow exception occurs.

3.3 Optional Facilities and Instructions

None.

3.4 LittleEndian

The POWER8 core supports true little endian. Byte swapping is performed before data is written to the I-cache and before data is fetched into the execution units; that is, between the D-cache and the execution units (for example, GPR, FPR/VR/VSR).

The load and store multiple instructions and the move assist instructions are not supported in little-endian mode. Attempting to execute any of these instructions in little-endian mode causes the system alignment error handler to be invoked.

3.5 Book II - Virtual Environment Architecture

3.5.1 Classes of Instructions

3.5.1.1 *Optional Instructions*

None.



3.5.2 Cache

The POWER8 chip contains three levels of cache hierarchy. All the caches (L1 I-cache, L1 D-cache, and the L2 and L3 caches) are dynamically shared among the eight threads. A cache block might be installed by one thread and used by the other threads. The basic coherence block size for the POWER8 core is 128 bytes.

The POWER8 chip automatically maintains the coherency of all data cached in these caches. The L1 cache employs Harvard cache organization, with separate L1 I-cache and L1 D-cache. L2 and L3 caches are unified. Because some levels of the cache hierarchy contain both instructions and data, when an instruction cache reload request is serviced by the L2 and/or the L3 caches, it is done so in a coherent manner.

The processor keeps the instruction storage consistent with the data storage. All cache lines in the L1 I-cache and L1 D-cache are also present in the L2 cache (inclusive property maintained).

The L1 I-cache is 8-way set associative and is indexed with five effective address bits (EA[51:55]). A particular physical block of memory with a given real address can be found in one of two positions in the L1 I-cache. The tag comparison associated with lookups in this cache (as well as all other caches in the system) are done using physical addresses, so that there are no synonym or alias hazards that must be explicitly handled by the system software.

The L1 D-cache is 8-way set associative and is indexed with six effective address bits (EA[51:56]). A particular physical block of memory with a given real address can only be found at a particular location in the L1 D-cache. On each access, the tag comparison is done with the physical address. On a cache miss, the cache reload mechanism searches the other seven related sets to determine if the required real address block is located elsewhere in the cache, and if so, it appropriately eliminates these copies.

In addition to maintaining caches, the POWER8 chip also includes several types of queues that act as logical predecessors and extensions to the caches. In particular, the machine contains store queues for holding store data above the caches, cast-out queues for holding modified data that has been pushed out of the caches (by the replacement algorithm, cache control instructions, or snoop requests), and others. All of these queues are maintained coherent by the hardware. In general, their presence should not be observable by either software or system hardware.

3.5.3 Data Prefetch

The POWER8 core provides an aggressive hardware-based data prefetching engine that is designed to work well for stride-one technical workloads with up to 16 streams. The 16 streams can be dynamically shared among the eight threads, with a sharing policy described elsewhere in more details. The POWER8 core implements enhanced data prefetching (**edcbt** instruction). For enhanced data prefetching, each thread can employ up to sixteen software initiated streams in ST mode. In SMT2 mode, the sixteen entries are shared between the two threads. In SMT4 mode, there are two groups of eight entries, and each group is shared between two threads. In SMT8 mode, there are four groups of four entries, and each group is shared between two threads

The POWER8 core also supports instruction cache block touch (**icbt**) by mapping it to **dcbtst** to prefetch into the L2 cache.

Another feature of the POWER8 core is the DSCR[58] bit that turns off hardware load-stream prefetching that can be accessed in problem state.

3.5.4 Effect of Operand Placement on Performance

In general, the POWER8 core provides excellent performance for storage accesses to naturally aligned boundaries (that is, to the boundaries defined by the operand size). In addition, the POWER8 core provides excellent hardware support for most unaligned storage accesses cases. In particular, cacheable load operations within 128-byte aligned sections of memory that hit in the L1 D-cache are performed at full speed independent of address alignment. Note that many storage accesses that cross page boundaries and segment boundaries are performed by the hardware. See *Section 3.1.4 Storage Access Alignment Support Overview* on page 44 for details.

3.5.5 Storage Model

The POWER8 core supports:

- Coherence block size of 128 bytes
- Weakly consistent access ordering of stores and loads per the Power ISA Book II, unless the special WIMG encoding designating all accesses to specific pages as being strongly ordered (SAO) is used to translate the access.
- Transactional memory accesses are performed and ordered as described in the Power ISA.

3.5.5.1 Atomicity

The POWER8 core is fully compliant with the architectural requirement for single-copy atomicity on naturally aligned storage accesses. This includes the quadword data atomicity associated with the **lq**, **lqarx**, **stq**, and **stqcx**. instructions. Furthermore, transactional mode accesses executed by a given thread appear to execute atomically to all other threads in the system (assuming the transaction succeeds).

3.5.5.2 Vector Element Atomicity

The Power ISA does not require vector accesses (VMX or VSX) be treated as atomic. However, the POWER8 processor does treat each element (word or doubleword) of VSX vector load and store instructions as an atomic access provided each element is naturally aligned.

3.5.5.3 Transactional Memory

The POWER8 processor supports Transactional Memory as described in the Power ISA.

The Transaction Exception and Status Register (TEXASR) is used to record various failure conditions that are described in the Power ISA. In addition, TEXASR[15] is used to specify various implementation-specific transaction failure causes that are not architected. The POWER8 processor sets TEXASR[15]='1' for the following reasons (implementation-specific transactional failure causes):

- Instruction fetch to caching-inhibited page in transactional mode.
- Recovery in transactional or suspend mode.
- Quiesce request in transaction or suspend mode.

The persistent bit is set to '0' for all of the above cases.



3.5.5.4 Storage Access Ordering

The architecture defines a weakly ordered storage model for most types of storage access scenarios. For these cases, the POWER8 core takes advantage of this relaxed requirement to achieve better performance through out-of-order instruction execution and out-of-order bus transactions. As a result, if strongly-ordered storage accesses are required, software must use the appropriate synchronizing instruction (**sync**, **ptesync**, **eieio**, or **lwsync**) to enforce order explicitly, or perform these accesses to regions marked with attributes that require the hardware to enforce strong ordering.

The POWER8 core employs the Real Mode Storage Control (RMSC) facility. Therefore, stores to storage marked as *non-guarded*, can be performed out-of-order, if the access is to well-behaved memory, as specified in the Real Mode Storage Control facility. Otherwise, stores to storage marked as *guarded*, cannot be performed out-of-order.

3.5.6 Atomic Updates and Reservations

The reservation granule size in the POWER8 core is 128 bytes.

There is at most one reservation per thread at any point in time.

3.5.7 Storage Control Instructions

3.5.7.1 Overview of Key Aspects of Storage Control Instructions

In the POWER8 core, all cache control instructions operate on aligned 128-byte sections of storage. *Table 3-5* summarizes many of the key aspects of the storage control instructions.

Table 3-5. Storage Control Instructions (Sheet 1 of 2)

Aspect	Book II Cache Instructions					
	icbi	dcbt	dcbtst	dcbz	dcbst	dcbf/dcbfl
Granularity	128 bytes					
Semantic checking	load (DSI on storage exception)	load (NOP on storage exception)	load (NOP on storage exception)	store (DSI on storage exception)	load (DSI on storage exception)	load (DSI on storage exception)
“r” bit update	yes	yes	yes	yes	yes	yes
“c” bit update	no	no	no	yes	no	no
L1 I-cache effect	see <i>Section 3.5.7.2</i> on page 77	none	none	none	none	none
L1 D-cache effect	none	see <i>Section 3.5.7.4</i> on page 77	see <i>Section 3.5.7.4</i> on page 77	as define in architecture	NOP	as define in architecture
L2 cache effect	none	see <i>Section 3.5.7.4</i> on page 77	see <i>Section 3.5.7.4</i> on page 77	see <i>Section 3.5.7.7</i> on page 78	see <i>Section 3.5.7.8</i> on page 78	see <i>Section 3.5.7.9</i> on page 79

Table 3-5. Storage Control Instructions (Sheet 2 of 2)

Aspect	Book II Cache Instructions					
	icbi	dcbt	dcbtst	dcbz	dcbst	dcbf/dcbfl
L3 cache effect	none	see <i>Section 3.5.7.4 on page 77</i>	see <i>Section 3.5.7.4 on page 77</i>	see <i>Section 3.5.7.7 on page 78</i>	see <i>Section 3.5.7.8 on page 78</i>	see <i>Section 3.5.7.9 on page 79</i>
<u>TLB</u> effect	reload as required	reload as required	reload as required	reload as required	reload as required	reload as required
<u>SLB</u> effect	reload as required	None (NOP if miss)	None (NOP if miss)	reload as required	reload as required	reload as required

3.5.7.2 Instruction Cache Block Invalidate (**icbi**)

The POWER8 core implements a split instruction/data (I/D) L1 cache where both caches are kept coherent with the L2 cache. Whenever any modification is made to the cache lines contained in the L2 cache, the L2 invalidates the copies in the L1 I-caches. Because of this, after an **icbi** instruction is translated, the processor core converts it to a *NOP* and does not broadcast the cache line targeted by the **icbi** instruction as the architecture stipulates. As a result of this and other implementation-specific design optimizations, instead of requiring the instruction sequence specified by the Power ISA to be executed on a per cache-line basis, software must only execute a single sequence of three instructions to make any previous code modifications become visible: **sync**, **icbi** (to any address), **isync**.

3.5.7.3 Instruction Cache Synchronize (**isync**)

As a performance optimization, the POWER8 core internally tracks and scoreboards **icbi** instructions that are required to be synchronized by the **isync** instruction. When the **isync** instruction is executed, this scoreboard bit is checked to see whether or not the machine must *flush and refetch* the instructions following the **isync**.

3.5.7.4 Data Cache Block Touch (**dcbt** and **dcbtst**)

The data cache block size for **dcbt** and **dcbtst** on the POWER8 core is 128 bytes.

The **dcbtst** instruction operates exactly the same way as the **dcbt** instruction.

These instructions act as a touch for the D-cache hierarchy and the TLB. If data translation is enabled (MSR[DR] = 1), and an SLB miss results, the instruction will be treated as a *NOP*. If a TLB miss results, the instruction reloads the TLB (and sets the reference bit). Once past translation, if the page protection attributes prohibit access, the page is marked *cache inhibited*, the page is marked *guarded*, or the processors' D-cache is disabled (via the bits in the HID4 Register), the instruction is finished as a *NOP* and does not reload the cache. Otherwise, the instruction checks the state of the L1 D-cache, and if it is not present, it initiates a reload. Note that this might also reload the L2 cache and/or the L3 cache with the addressed block if it is not already present in these caches. If the cache block is already present in the L1 D-cache, the cache content is not altered. Note that if the **dcbt** or **dcbtst** instruction does reload cache blocks, it affects the state of the cache replacement algorithm bits.

The POWER8 core does implement the optional extension to the **dcbt** instruction that allows software to directly engage a data stream prefetch from a particular address.

3.5.7.5 Data Cache Block Touch - No Access Needed Anymore (TH = '10001')

The POWER8 core supports this as specified in the Power ISA.

3.5.7.6 Data Cache Block Touch - Transient (TH = '10000')

The POWER8 core implements the load and store version of these transient touch instructions (**dcbtct**, **dcbtds**, **dcbtt**, **dcbtstct**, and **dcbtstt**). The transient property of a cache line is retained in the L3 cache for both the load and store version of the transient touch instructions. The transient property of a cache line is retained in the L2 cache for the load version of the transient touch instruction for the case that the line is loaded but not stored into. In this transient state, the transient line becomes the most likely cache line in its congruence class to be replaced next, thus preserving the other cache lines in that congruence class. This behavior is useful if it is known that a set of lines will be loaded or stored with a low probability for temporal cache reuse and it is desirable that they be as minimally intrusive to the cache as possible (for example, displacing as few lines in the cache as possible). Reading or writing a large array with the help of transient touch instructions only impacts one of the eight sets in the L3 cache, and reading a large array with the help of transient touch instructions only impacts one of the eight sets of the L2 cache.

3.5.7.7 Data Cache Block Zero (dcbz)

The data cache block size for **dcbz** on the POWER8 core is 128 bytes.

The function of **dcbz** is performed in the L2 cache. As a result, if the block addressed by the **dcbz** is present in the L1 D-cache, the block is invalidated before the operation is sent to the L2 cache logic for execution. The L2 cache gains exclusive access to the block (without actually reading the old data) and performs the zeroing function in a broadside manner.

If the cache block specified by the **dcbz** instruction contains an error (even one that is not correctable with ECC), the contents of all locations within the block are set to zeros in the L2 cache. If the specified block in the L2 cache does not contain a hard fault, a subsequent load from any location within the cache block returns zeros and does not cause a machine check interrupt.

If the block addressed by the **dcbz** instruction is in a memory region marked *cache inhibited*, or if the L1 D-cache or L2 cache is disabled (using the bits in the HID registers), then the instruction causes an alignment interrupt to occur.

3.5.7.8 Data Cache Block Store (dcbst)

The data cache block size for **dcbst** on the POWER8 core is 128 bytes.

The **dcbst** instruction has no direct effect on the L1 D-cache (because it is store-through, it never contains modified data). It also has no direct effect on the L2 cache or L3 cache (both of these are kept coherent with memory and I/O, so that nothing special needs to be done). As a result, the instruction simply goes through address translation, reports any errors, and is completed. The instruction is not sent to the storage subsystem, and consequently it does not broadcast any transactions onto the inter-processor SMP interconnect.

3.5.7.9 Data Cache Block Flush (`dcbf`, `dcbfl` and `dcbfip`)

The data cache block size for **dcbf**, **dcbfl** and **dcbfip** on the POWER8 core is 128 bytes.

The POWER8 core supports **dcbf** (L = 0), **dcbfl** (**dcbf** with L = 1) and **dcbfip** (**dcbf** with L = 3) as specified in the Power ISA.

3.5.7.10 Load and Reserve and Store Conditional Instructions

The reservation granule size for the POWER8 core is 128 bytes.

An attempt to execute a non-halfword aligned **Iharx or sthcx.**, or a non-word aligned **Iwarx or stwcx.**, or a non-doubleword aligned **Idarx or stdcx.**, or a non-quadword aligned **Iqarx or stqcx.** to a cacheable address causes an alignment interrupt.

An attempt to execute a **Ibarx, Iharx, Iwarx, Idarx, Iqarx, stbcx., sthcx., stwcx., stdcx., or a stqcx.** instruction to storage marked cache inhibited causes a data storage interrupt independent of the address alignment.

There are separate reservations per thread.

3.5.7.11 sync Instruction

The POWER8 design achieves high performance by exploiting speculative out-of-order instruction execution. The *heavyweight sync* (**hwsync**) instruction, as defined in the architecture, acts as a serious barrier to this type of aggressive execution and therefore, can have a dramatic effect on performance. Although the POWER8 core has optimized the performance of **hwsync** to some degree, care should be exercised in the indiscriminate use of this instruction. As a performance consideration, software should attempt to use the lightweight version of **sync** (often referred to as **lwsync** in this document) whenever possible. Unless otherwise stated, **sync** refers to **hwsync**.

The POWER8 core implements the **ptesync** for use in synchronizing page table updates.

See the Power ISA Books II and III-S for a complete description of the different forms of the **sync** instruction.

The POWER8 core does not support the optional elemental memory barriers described in the Power ISA. Instead, the Power8 design ignores the EB field of the **sync** instruction and the L-field solely determines which version of the **sync** instruction (**sync/lwsync/ptesync**) is performed by the hardware.

3.5.7.12 eieio Instruction

The POWER8 core implements **eieio** as described in the Power ISA.

In the POWER8 nest logic, the store queues above the L2 cache attempt to gather sequential both cacheable and cache-inhibited store operations to improve bandwidth. If this behavior is not desired, software must insert either an **eieio** (preferable for performance) or a **sync** to prevent it.

3.5.7.13 miso Instruction

The POWER8 core implements the **miso** instruction as a NOP. It has no effect on the stores.



3.5.7.14 Transactional Memory Instructions

The POWER8 core implements the transactional memory instructions as described in the Power ISA.

Table 3-6 lists how certain cache and TLB management instructions affect transactions.

Table 3-6. Cache and TLB Management Instruction Effects on Transactional Accesses

Mode			
Instruction	TM State	Fails Transaction	TEXASR bit set
tlbie	T	Always	8 - disallowed
tlbie	S	When the virtual address it is attempting to invalidate hits in the bloom filter for the current transaction	14 - translation invalidation conflict
tlbiel	T	Always	8 - disallowed
tlbiel	S	Never	N/A
dcbt (any TH)	T	Never (unless it causes a castout of the TM footprint)	10 - footprint overflow
dcbt (any TH)	S	Never (unless it causes a castout of the TM footprint)	10 - footprint overflow
dcbst	T	Always	8 - disallowed
dcbst	S	Never (dcbst is treated as a NOP in this case)	11 - self-induced conflict
dcbf (L=0,1)	T	Always	8 - disallowed
dcbf (L=0,1)	S	When the block (line) being pushed out of the cache is part of the TM footprint	11 - self-induced conflict
dcbf (L=3)	T	Always	8 - disallowed
dcbf (L=3)	S	Never	N/A
dcbz	T	Never (unless dcbz causes a castout of the TM footprint)	N/A
dcbz	S	Case 1: When the block (line) being zero'ed is part of the TM footprint Case 2: When the dcbz causes a castout of the TM load or store footprint	Case 1: 11 - self-induced conflict Case 2: 11 - footprint overflow
dcbtst	T	Never (unless dcbtst causes a castout of the TM footprint)	10 - footprint overflow
dcbtst	S	Never (unless dcbtst causes a castout of the TM footprint)	10 - footprint overflow
icbi	T	Always	8 - disallowed
icbi	S	Never (icbi is treated as NOP with regards to transaction failure in suspend mode)	N/A
icbt	T	Never (unless icbt causes a castout of the TM footprint)	10 - footprint overflow
icbt	S	Never (unless icbt causes a castout of the TM footprint)	10 - footprint overflow

3.5.8 Timer Facilities

3.5.8.1 Time Base

The Time-Base (TB) is designed to tick at the rate of time-of-day (TOD). In other words, bit 59 of the Time-Base Register increments at the 32 MHz clock. There is one Time-Base per processor core that is shared by all the threads running on a core. There is one decrementer per thread.

The POWER8 core implements two time-base modes: POWER8 time-base mode and non-POWER8 time-base mode. They are selectable by setting a mode bit in the Time Facility Management Register (TFMR).

Time Facility Management Register

The Time Facility Management Register (TFMR) is an SPR that is accessible only in the hypervisor state. Executing a move to or move from TFMR in a nonhypervisor state causes a privileged interrupt. There is one TFMR per processor core that is shared among the threads. The TFMR is used as both a status and control register.

POWER8 Time-Base Mode

The time-base function uses an external time-of-day (TOD) clock, which is independent of the processor frequency. This is required to support dynamic frequency variation for power management. The external TOD oscillator can be 16 MHz or 32 MHz. The external TOD oscillator is sampled to provide a 32 MHz step signal, which is distributed to all processors in the system.

Bits 0:59 of the TB are incremented at the 32 MHz frequency as provided by the distributed step signal. Bits 60:63 of the TB are incremented at a fixed frequency of 500 MHz. If the value of bits 60:63 is '1111' (saturated), it is held until the 32 MHz step signal causes bit 59 to change. At that time, bits 60:63 are allowed to change to '0000'.

To support multi-node configurations across multiple oscillator domains, error detection and recovery, and concurrent maintenance, the POWER8 core uses the following means of synchronizing the time bases across all processors. Each multicore processor chip contains a Time-of-Day (TOD) Register. The chip TOD registers are first synchronized across all the processor chips. Then, the time-base registers in each processor core are synchronized to the chip TOD. Also encoded on the step signal is a synchronization pulse which is used for synchronization and error checking. The synchronization mechanism requires system operations to complete within a sync interval. The sync interval can be set via the TFMR bits to be, for example, 1 µs, 2 µs (default, corresponds to TB bit 53), 4 µs, or 8 µs.

Error checking includes parity checks on all registers, and functional checking such as missing step signal detection and synchronization errors. The step signal rate is defined in the POWER8 mode to be 32 MHz, and the logic checks for the correct number of steps for each synchronization signal (which is selected by TFMR). After the TB is operational, the hardware also detects a missing step signal, which requires specifying in the TFMR the maximum number of processor cycles allowed without seeing a 32 MHz step signal, for the fastest allowable operating frequency. The TFMR maximum cycle step time-out should be specified as $(2 \times 31.25 \text{ ns}) / (\text{minimum processor cycle time in ns} \times 4)$.

The initial synchronization requires some software sequencing, which is performed by writing values to the TFMR (via **mtspr**). The TFMR also indicates the status of the various time facilities. The status bits in the TFMR are read-only, not modified by **mtspr** to the TFMR. The time facility logic implements error detection



for hardware and also for invalid software sequencing. Because synchronized time is critical to a system, writes to the time base or the TFMR that would break synchronization cause the logic to enter an error state and trigger a hypervisor maintenance interrupt.

To initially set the time and synchronize the time base values, software must synchronize all processors in the system and choose one processor to perform updates to the TFMR via a read-modify-write operation to preserve the other bits. This sequence assumes the external TOD oscillator distribution is already running.

After the chip TOD is running on all chips and the TB is running on the processor that drove this sequence, software must then release the remaining processors to synchronize their TB registers to their corresponding chip TOD.

3.5.9 Hypervisor Decrementer (HDEC)

There is one hypervisor decrementer register per thread. HDEC decrements every time TB bit 63 is incremented.

3.5.10 Decrementer (DEC)

There is one decrementer register per thread. DEC decrements every time TB bit 63 is incremented.

3.5.11 Book II Invalid Forms

The results of executing an invalid form of an instruction in Book II or an instance of such an instruction for which the architecture specifies that some results are undefined, are described here for the cases in which executing an instruction does not cause an exception. Only results that differ from those specified by the architecture are described in the following list.

- Instruction with reserved fields
Bits in reserved fields are ignored; the results of executing an instruction in which one or more reserved bits are ‘1’ is the same as if the bits were ‘0’.
- Transactional memory instructions and store conditional instructions (bit 31 is ignored)
Bit 31 of **tbegin.**, **tend.**, **tabort.**, **tabortwc.**, **tabortdc.**, **tabortwci.**, **tabortdci.**, **treclaim.**, **stbcx.**, **sthcx.**, **stwcx.**, **stdcx.** and **stqcx.** is ignored. Bit 31 = ‘1’ or bit 31 = ‘0’ is treated the same given that other x-form instructions implicitly set CR and have no “non-record” from variant. Ignoring bit 31 is an acceptable way to handle this invalid form.
- **mftb** instructions
This instruction produces the same result as the **mfsp** instruction. For a complete description on the associated invalid forms, see *Section 3.1.4.9 Move To/From Special Purpose Register (SPR) Instructions* on page 65.



3.6 Book III - Operating Environment Architecture

3.6.1 Classes of Instructions

3.6.1.1 Storage Control Instructions

The POWER8 core does provides support for the following instructions:

- **tlbie** - TLB invalidate entry (large and small page)
- **tlbiel** - Processor local form of TLB invalidate entry (large and small page, and the IS field)
- **tlbsync** - TLB synchronize
- **slbmte** - Segment Lookaside Buffer Move To Entry
- **slbmfev** - Segment Lookaside Buffer Move From Entry VSID
- **slbmfee** - Segment Lookaside Buffer Move From Entry ESID
- **slbfee** - Segment Lookaside Buffer Find Entry ESID
- **slbie** - SLB invalidate entry
- **slbia** - SLB invalidate all
- **mtsr** - Move to segment register (Bridge Facility)
- **mtsrin** - Move to segment register indirect (Bridge Facility)
- **mfsr** - Move from segment register (Bridge Facility)
- **mfsrin** - Move from segment register indirect (Bridge Facility)
- **mtmsr** - Move to Machine State Register (32-bit)
- **mtmsrd** - Move to Machine State Register (64-bit)
- **sc** - System Call
- **rfid** - Return From Interrupt Doubleword
- **hrfid** - Hypervisor Return From Interrupt Doubleword

The POWER8 core does not provide support for the following optional or obsolete instructions (attempted use of these results in a hypervisor emulation assistance interrupt):

- **tlbia** - TLB invalidate all
- **tlbiex** - TLB invalidate entry by index (obsolete)
- **slbiex** - SLB invalidate entry by index (obsolete)
- **dcb*a*** - Data cache block allocate (Book II; obsolete)
- **dcb*i*** - Data cache block invalidate (obsolete)
- **rfi** - Return from interrupt (32-bit; obsolete)

The following instruction variants are implemented:

- **ptesync** - Page table synchronize
- **hwsync** - Heavyweight synchronize
- **lwsync** - Lightweight synchronize



3.6.1.2 Reserved Instructions

The architecture breaks the reserved instruction class down into several categories as described in the *Reserved Instructions* appendix of the Power ISA. The POWER8 processor core behaves in the following manner with respect to these categories:

- Primary opcode equals zero.
 - POWER Architecture instructions not in the Power ISA. The POWER8 core takes a hypervisor emulation assistance interrupt. See the complete list in the Power ISA appendices.
 - Service processor “Attention” instruction.
 - In addition, there are several implementation-specific registers available for access through the **mtspr** and **mfspr** instructions. These are described in *Section 3.6.3.4 Move To/From Special Purpose Register Instructions* on page 87.

3.6.2 Branch Processor

3.6.2.1 SRR1 Register

In the POWER8 processor core, the SRR1 is implemented per the Power ISA.

3.6.2.2 MSR Register

In the POWER8 processor core, the MSR is implemented per the Power ISA.

3.6.2.3 Branch Processor Instructions

Branch Processor Instruction with Undefined Results

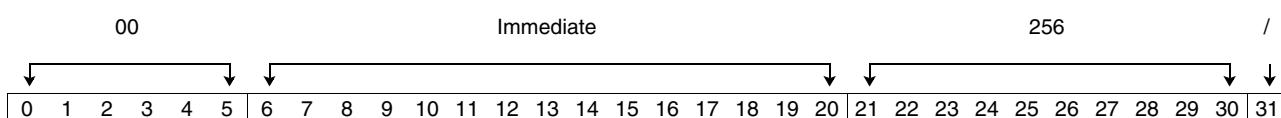
The results of executing an invalid form of a branch instruction or an instance of a branch instruction for which the architecture specifies that some results are undefined are described as follows.

System Call SC-Form

In the POWER8 processor core, the system call (`sc`) instruction is implemented per the Power ISA.

Support Processor Attention Instruction

The POWER8 processor core supports a special, implementation-dependent instruction for signaling an attention to the support processor.



The immediate field (I) has no effect on the operation of this instruction in the POWER8 processor core.



If HID0[31] = '1' (the support-processor attention enable bit is set), this instruction causes all preceding instructions to run to completion, the machine to quiesce, and the assertion of the support processor attention signal. Instruction execution does not resume until the support processor signals it to do so.

If HID0[31] = '0' (the support processor attention enable bit is not set), this instruction causes a hypervisor emulation assistance interrupt.

3.6.2.4 Current Instruction Address Breakpoint (CIABR)

The POWER8 processor core supports the CIAB Register as implemented per the Power ISA.

3.6.2.5 Instruction Effective to Real Address Translation Cache (I-ERAT)

The POWER8 processor core includes a 64-entry, 64-way set associative instruction effective-to-real address translation (I-ERAT) for fast translation of instruction effective addresses into physical (real) addresses. The ERAT is dynamically shared between all eight threads. The ERAT is implemented as a CAM that supports page sizes of 4 KB, 64 KB, and 16 MB. Instruction fetches to 16 GB pages are installed in the I-ERAT as multiple 16 MB page entries as needed. Likewise, instruction fetches to 1 MB pages are installed as multiple 64 KB pages as needed. Entries for hypervisor relocate off can be shared between all eight threads (regardless of LPAR mode), further increasing the capacity of the ERAT.

Because addresses associated with non-hypervisor real mode accesses are translated differently than those associated with virtual-mode accesses, the IERAT must keep the MSR[IR] and MSR[HV] bits (along with various bits of translation information) in each entry. This allows the I-ERAT to distinguish between translations that are valid for the various modes of operation. Because the content of each I-ERAT entry is the result of a page table search based on the contents of an SLB entry, to maintain consistency with the SLB (or segment registers), the following instructions cause all entries in the I-ERAT that belong to the thread executing the instruction to be invalidated. The ERAT compares as many effective address bits as are available in the various invalidate operations.

- **mtsr** or **mtsrin** instructions - used for segment register changes in 32-bit operating systems
- **slbia**
- **mtiamr**

The **slbie** instruction causes invalidation of a D-ERAT entry belonging to the thread (no impact to the other thread) only if there is a perfect address match (that is, for invalidation effective address bits, EA[0:35] are matched for **slbie** small 256 MB segment, EA[0:23] are matched for **slbie** 1 TB segment).

The **tlbie** instruction (or the detection of snooped-tlbie operations) invalidates all D-ERAT entries (irrespective of the thread) in the D-ERAT that have a perfect match. In other words, the entry is invalidated only if:

- EA[36:51] are matched for tlbie small page
- EA[36:47] are matched for tlbie 64 KB page
- EA[36:43] are matched for tlbie 1 MB page
- EA[36:39] are matched for tlbie 16 MB page

EA[24:29] are matched for tlbie 16 GB page

Upon power-on, each I-ERAT entry is set to the invalid state.



I and G bit setting in the I-ERAT is done based on *Table 3-7*.

Table 3-7. I-ERAT I and G Bit Setting

Condition									Resulting Action
MSR [IR]	MSR [HV]	LPCR[0]	HID4[33]	First access I = 1 Fetch	HID4[0:4] RMSC (above line)	I	G		
1	x	x	x	x	x	PTE(I)	PTE(G)	If G = '0', the page is written into the I-ERAT using the I-bit value and page size determined from the PTE and described above. If G = '1', an ISI is taken.	
0	0	1	x	x	x	PTE(I)	PTE(G)	Virtual real mode: If G = '0', the page is written into the I-ERAT using the I-bit value and page size determined from the PTE and described above. If G = '1', an ISI is taken.	
0	0	0	x	x	x	0	0	RMOR mode: set IG = '00'	
0	1	x	0	yes	no	1	0	Legacy RMSC mode: If access is below the guarded line established by HID4[0:4], access is allowed and a 4K page is installed in I-ERAT as I = '1', G = '0'.	
0	1	x	0	no	no	0	0	Legacy RMSC mode. If access is below the guarded line established by HID4[0:4], access is allowed and a 4 KB page is installed in I-ERAT as I = '0', G = '0'.	
0	1	x	0	x	yes	N/A	N/A	Legacy RMSC mode. If access is above the guarded line established by HID4[0:4], access is not allowed resulting in an ISI due to an instruction fetch to guarded storage.	
0	1	x	1	yes	x	1	N/A	Page-based RMSC mode. 16 M page is installed in I-ERAT with I = '1', G = '1'.	
0	1	x	1	no	x	0	N/A	Page-based RMSC mode. 16 M page is installed in I-ERAT with I = '0', G = '0'.	

3.6.3 Fixed-Point Processor

3.6.3.1 Processor Version Register (PVR)

The processor version number (PVR[0:15]) for the POWER8 processor in the SCM is x'004D'. In future versions of the POWER8 core, this section of the version number will only change if there are significant software-visible changes in the design.

The processor revision level (PVR[16:31]) starts at x'0100', indicating revision '1.0'. As revisions are made, bits [29:31] will indicate minor revisions. Similarly, bits [20:23] indicate major changes.

Example: The PVR value for POWER8 processor in the SCM for design level 1 is: x'004D0100'.

3.6.3.2 Processor ID Register (PIR)

The processor identification register (PIR) is a 32-bit register that holds a processor identification tag in the 9 least-significant bits [23:31]. This tag is used for tagging bus transactions and for processor differentiation in multiprocessor systems.

Bits	Field Name	Description
0:21	Reserved	Read as zeros
22:24	ChID	Chip ID
25:28	CoID	Core number
29:31	TID	Thread ID

The PIR is a read-only register. During power-on reset, PIR is set to a unique value for each processor in a multiprocessor system.

3.6.3.3 Chip Information Register (CIR)

The POWER8 processor implements the CIR per the Power ISA.

3.6.3.4 Move To/From Special Purpose Register Instructions

The POWER8 core supports special purpose registers listed in *Table 3-8*. Many of these SPRs are only accessible in hypervisor or privileged modes. A handful of these registers (for example, DSCR) are also user-mode accessible through a second SPR number.

To support multithreading, some of the SPRs are replicated in the POWER8 core, while the others are shared, as shown in the SMT column in *Table 3-8*. In the table column headers, Prob indicates problem state (HV = x, PR = 1), Priv indicates privileged state (HV = 0, PR = 0), and Hyp indicates Hypervisor state (HV = 1, PR = 0). In the SPR specific rows, Priv indicates that a privileged instruction type program interrupt will occur in that state for the attempted read or write of the SPR. FAC indicates the Facility Unavailable interrupt. HFAC indicates the Hypervisor Facility Unavailable interrupt. Illegal indicates a hypervisor emulation assistance interrupt will occur. NOP indicates the instruction will be treated as a NOP. A blank under each column indicates the access will be performed normally.

Table 3-8. SPR Table (Sheet 1 of 6)

SPR Name	SPR spr[5:9] spr[0:4]	Decimal	SMT	Length	Read (mfspr)			Write (mtspr)		
					Prob	Priv	Hyp	Prob	Priv	Hyp
XER	00000 00001	1	Replicated	64						
DSCR (FSCR [61] = 0)	00000 00011	3	Replicated	25	FAC			FAC		
DSCR (FSCR [61] = 1)	00000 00011	3	Replicated	25						
LR	00000 01000	8	Replicated	64						
CTR	00000 01001	9	Replicated	64						
AMR	00000 01101	13	Replicated	64						
DSCR (HFSCR[61] = 0)	00000 10001	17	Replicated	25	Priv	HFAC		Priv	HFAC	



Table 3-8. SPR Table (Sheet 2 of 6)

SPR Name	SPR		Decimal	SMT	Length	Read (mfspr)			Write (mtspr)		
	spr[5:9]	spr[0:4]				Prob	Priv	Hyp	Prob	Priv	Hyp
DSCR (HFSCR[61] = 1)	00000 10001	17	Replicated	25	Priv				Priv		
DSISR	00000 10010	18	Replicated	32	Priv				Priv		
DAR	00000 10011	19	Replicated	64	Priv				Priv		
DEC	00000 10110	22	Replicated	32	Priv				Priv		
SDR1	00000 11001	25	Per LPAR	64	Priv	Priv			Priv	Priv	
SRR0	00000 11010	26	Replicated	64							
SRR1	00000 11011	27	Replicated	64							
CFAR	00000 11100	28	Replicated	64							
AMR	00000 11101	29	Replicated	64							
Reserved	00000 11111	31	Replicated	64							
PID	00001 10000	48	Replicated	32	Priv				Priv		
IAMR	00001 11101	61	Replicated	32	Priv				Priv		
TFHAR	00100 00000	128	Replicated	64							
TFIAR	00100 00001	129	Replicated	64							
TEXASR	00100 00010	130	Replicated	64							
TEXASRU	00100 00011	131	Replicated	32							
CTRL	00100 01000	136	Shared	32	bit 63 - Priv	bit 63 - Priv			Illegal	NOP	NOP
CTRL	00100 11000	152	Shared	32	Priv	bit 63 - Priv			Illegal	NOP	NOP
FSCR	00100 11001	153	Replicated	64	Priv				Priv		
UAMOR	00100 11101	157	Replicated	64	Priv				Priv		
PSPB	00100 11111	159	Replicated	32	Priv				Priv		
DPDES	00101 10000	176	Per LPAR	8	Priv	Priv			Priv	Priv	
DHDES	00101 10001	177	Shared	8	Priv	Priv			Priv	Priv	
DAWR0	00101 10100	180	Replicated	64	Priv	Priv			Priv	Priv	
RPR	00101 11010	186	Per LPAR	64	Priv	Priv			Priv	Priv	
CIABR	00101 11011	187	Replicated	64	Priv	Priv			Priv	Priv	
DAWRX0	00101 11100	188	Replicated	32	Priv	Priv			Priv	Priv	
HFSCR	00101 11110	190	Replicated	64	Priv	Priv			Priv	Priv	
VRSAVE	01000 00000	256	Replicated	32							
SPRG3	01000 00011	259	Replicated	64	Illegal	NOP	NOP				
TB	01000 01100	268	Per LPAR	64	Illegal	NOP	NOP				
TBU	01000 01101	269	Per LPAR	64	Illegal	NOP	NOP				
SPRG0	01000 10000	272	Replicated	64	Priv				Priv		
SPRG1	01000 10001	273	Replicated	64	Priv				Priv		



Table 3-8. SPR Table (Sheet 3 of 6)

SPR Name	SPR spr[5:9] spr[0:4]	Decimal	SMT	Length	Read (mfspr)			Write (mtspr)		
					Prob	Priv	Hyp	Prob	Priv	Hyp
SPRG3	01000 10011	275	Replicated	64	Priv			Priv		
SPRC	01000 10100	276	Replicated	64	Priv	Priv		Priv	Priv	
SPRD	01000 10101	277	N/A	64	Priv	Priv		Priv	Priv	
TBL	01000 11100	284	Per LPAR	32	Illegal	NOP	NOP	Priv	Priv	
TBU	01000 11101	285	Per LPAR	32	Illegal	NOP	NOP	Priv	Priv	
TBU40	01000 11110	286	Per LPAR	64	Illegal	NOP	NOP	Priv	Priv	
PVR	01000 11111	287	Shared	32	Illegal	NOP	NOP	Priv		
HSPRG0	01001 10000	304	Replicated	64	Priv	Priv		Priv	Priv	
HSPRG1	01001 10001	305	Replicated	64	Priv	Priv		Priv	Priv	
HDSISR	01001 10010	306	Replicated	32	Priv	Priv		Priv	Priv	
HDAR	01001 10011	307	Replicated	64	Priv	Priv		Priv	Priv	
SPURR	01001 10100	308	Replicated	64	Priv	Priv		Priv		
PURR	01001 10101	309	Replicated	64	Priv	Priv		Priv		
HDEC	01001 10110	310	Per LPAR	32	Priv	Priv		Priv	Priv	
RMOR	01001 11000	312	Per LPAR	64	Priv	Priv		Priv	Priv	
HRMOR	01001 11001	313	Shared	64	Priv	Priv		Priv	Priv	
HSRR0	01001 11010	314	Replicated	64	Priv	Priv		Priv	Priv	
HSRR1	01001 11011	315	Replicated	64	Priv	Priv		Priv	Priv	
MMCRH	01001 11100	316	Shared	64	Priv	Priv		Priv	Priv	
TFMR	01001 11101	317	Shared/ LPAR bit 26 and 45 replicated	64	Priv	Priv		Priv	Priv	
LPCR	01001 11101	318	Replicated	64	Priv	Priv		Priv	Priv	
LPIDR	01001 11111	319	Per LPAR	64	Priv	Priv		Priv	Priv	
HMER	01010 10000	336	Replicated	64	Priv	Priv		Priv	Priv	
HMEER	01010 10001	337	Shared	64	Priv	Priv		Priv	Priv	
PCR	01010 10010	338	Per LPAR	64	Priv	Priv		Priv	Priv	
HEIR	01010 10011	339	Replicated	32	Priv	Priv		Priv	Priv	
HPMC1	01010 11000	344	Shared	64	Priv	Priv		Priv	Priv	
HPMC2	01010 11001	345	Shared	64	Priv	Priv		Priv	Priv	
HPMC3	01010 11010	346	Shared	64	Priv	Priv		Priv	Priv	
HPMC4	01010 11011	347	Shared	64	Priv	Priv		Priv	Priv	
AMOR	01010 11101	349	Per LPAR	64	Priv	Priv		Priv	Priv	
TIR	01101 11110	446	Replicated	8	Priv			Illegal	NOP	NOP
SIER	11000 00000	768	Replicated	64				Illegal	NOP	NOP
MMCR2	11000 00001	769	Replicated	64				Illegal	NOP	NOP



Table 3-8. SPR Table (Sheet 4 of 6)

SPR Name	SPR		Decimal	SMT	Length	Read (mfspr)			Write (mtspr)		
	spr[5:9]	spr[0:4]				Prob	Priv	Hyp	Prob	Priv	Hyp
MMCRA	11000 00010	770	Replicated	32					Illegal	NOP	NOP
PMC1	11000 00011	771	Replicated	32					Illegal	NOP	NOP
PMC2	11000 00100	772	Replicated	32					Illegal	NOP	NOP
PMC3	11000 00101	773	Replicated	32					Illegal	NOP	NOP
PMC4	11000 00110	774	Replicated	32					Illegal	NOP	NOP
PMC5	11000 00111	775	Replicated	32					Illegal	NOP	NOP
PMC6	11000 01000	776	Replicated	32					Illegal	NOP	NOP
MMCR0	11000 01011	779	Replicated	32					Illegal	NOP	NOP
SIAR	11000 01100	780	Replicated	64					Illegal	NOP	NOP
SDAR	11000 01101	781	Replicated	64					Illegal	NOP	NOP
MMCR1	11000 01110	782	Replicated	32					Illegal	NOP	NOP
SIER	11000 10000	784	Replicated	64	Priv				Priv		
MMCR2	11000 10001	785	Replicated	64	Priv				Priv		
MMCRA	11000 10010	786	Replicated	64	Priv				Priv		
PMC1	11000 10011	787	Replicated	32	Priv				Priv		
PMC2	11000 10100	788	Replicated	32	Priv				Priv		
PMC3	11000 10101	789	Replicated	32	Priv				Priv		
PMC4	11000 10110	790	Replicated	32	Priv				Priv		
PMC5	11000 10111	791	Replicated	32	Priv				Priv		
PMC6	11000 11000	792	Replicated	32	Priv				Priv		
MMCR0	11000 11011	795	Replicated	32	Priv				Priv		
SIAR	11000 11100	796	Replicated	64	Priv				Priv		
SDAR	11000 11101	797	Replicated	64	Priv				Priv		
MMCR1	11000 11110	798	Replicated	32	Priv				Priv		
IMC	11000 11111	799	Shared	64	Priv	Priv	Return zeros	Priv	Priv	Conditional based on HID1	
BESCRS	11001 00000	800	Replicated	64							
BESCRSU	11001 00001	801	Replicated	32							
BESCRR	11001 00010	802	Replicated	64							
BESCRRU	11001 00011	803	Replicated	32							
EBBHR	11001 00100	804	Replicated	64							
EBBRR	11001 00101	805	Replicated	64							
BESCR	11001 00110	806	Replicated	64							
Reserved	11001 01000	808	N/A		NOP	NOP	NOP	NOP	NOP	NOP	NOP
Reserved	11001 01001	809	N/A		NOP	NOP	NOP	NOP	NOP	NOP	NOP



Table 3-8. SPR Table (Sheet 5 of 6)

SPR Name	SPR	Decimal	SMT	Length	Read (mfspr)			Write (mtspr)		
	spr[5:9] spr[0:4]				Prob	Priv	Hyp	Prob	Priv	Hyp
Reserved	11001 01010	810	N/A		NOP	NOP	NOP	NOP	NOP	NOP
Reserved	11001 01011	811	N/A		NOP	NOP	NOP	NOP	NOP	NOP
TAR	11001 01111	815	Replicated	64						
IC	11010 10000	848	Replicated	64	Priv	Priv		Priv	Priv	
VTB	11010 10001	849	Per LPAR	64	Priv	Priv		Priv	Priv	
MMCRC	11010 10011	851	Per LPAR	32	Priv	Priv		Priv	Priv	
PMICR	11010 10100	852	Shared	64	Priv	Priv		Priv	Priv	
PMSR	11010 10101	853	Shared	32	Priv	Priv		Priv	Priv	NOP
PMMAR	11010 10110	854	Per LPAR	64	Priv	Priv		Priv	Priv	
Reserved	11010 11011	859	Per Core	64	Priv	Priv		Priv	Priv	
Reserved	11010 11100	860	Per Core	64	Priv	Priv		Priv	Priv	
Reserved	11010 11101	861	Per Core	64	Priv	Priv	NOP	Priv	Priv	
Reserved	11010 11110	862	N/A		Priv	Priv		Priv	Priv	
Reserved	11010 11111	863	Shared	32	Priv	Priv		Priv	Priv	
TRIG0	11011 10000	880	Replicated	64	Priv			Illegal	NOP	NOP
TRIG1	11011 10001	881	Replicated	64	Priv			Illegal	NOP	NOP
TRIG2	11011 10010	882	Replicated	64	Priv			Illegal	NOP	NOP
PMCR	11011 10100	884	Per LPAR	64	Priv			Priv	Priv	
Reserved	11011 11000	888	Shared	64	Priv			Priv		
Reserved	11011 11001	889	Shared	64	Priv			Priv		
SPMC1	11011 11100	892	Replicated	32	Priv			Priv		
SPMC2	11011 11101	893	Replicated	32	Priv			Priv		
MMCRS	11011 11110	894	Replicated	32	Priv			Priv		
Reserved	11011 11111	895	Replicated	32	Priv			Priv		
PPR	11100 00000	896	Replicated	64						
PPR32	11100 00010	898	Replicated	32						
TSCR	11100 11001	921	Shared	32	Priv	Priv		Priv	Priv	
TTR	11100 11010	922	Shared	64	Priv	Priv		Priv	Priv	
TRACE	11111 01110	1006	Shared	64	Illegal	NOP	NOP			
HID0	11111 10000	1008	Shared	64	Priv	Priv		Priv	Priv	
HID1	11111 10001	1009	Shared (1 bit replicated for lpar)	64	Priv	Priv		Priv	Priv	
HID4	11111 10100	1012	Shared (some replicated for lpar)	64	Priv	Priv		Priv	Priv	



Table 3-8. SPR Table (Sheet 6 of 6)

SPR Name	SPR		Decimal	SMT	Length	Read (mfspr)			Write (mtspr)		
	spr[5:9]	spr[0:4]				Prob	Priv	Hyp	Prob	Priv	Hyp
HID5	11111 10110	1014	Shared (HID4 spill over replicated per lpar)	64	Priv	Priv			Priv	Priv	
CIR	11111 11110	1022	Shared	32	Priv				Illegal	NOP	NOP
PIR	11111 11111	1023	Replicated	32	Priv				Priv	NOP	NOP
Unsupported POWER MQ	00000 00000	0	N/A	N/A	Illegal	Illegal	Illegal	Illegal	Illegal	Illegal	Illegal
Unsupported Power Real Time Upper SPR	00000 00100	4	N/A	N/A	Illegal	Illegal	Illegal	Illegal	Illegal	Illegal	Illegal
Unsupported Power Real Time Lower SPR	00000 00101	5	N/A	N/A	Illegal	Illegal	Illegal	Illegal	Illegal	Illegal	Illegal
Unsupported Power Decrementer SPR	00000 00110	6	N/A	N/A	Illegal	Illegal	Illegal	Illegal	Illegal	Illegal	Illegal
Unsupported non-privileged, non-Power SPR	xxxxx 0xxxx		N/A	N/A	Illegal	Illegal	Illegal	Illegal	Illegal	Illegal	Illegal
Unsupported-privileged SPR	xxxxx 1xxxx		N/A	N/A	Illegal	Illegal	Illegal	Illegal	Illegal	Illegal	Illegal

3.7 HID Registers (HID0, HID1, HID4, and HID5)

The POWER8 processor core includes many implementation-dependent mode bits that allow various features of the chip to be enabled and disabled. These bits are included in the Hardware Implementation Dependent registers (HID0, HID1, HID4, and HID5). In general, the HID0 Register controls high-level functions of the POWER8 core. The HID1 Register contains additional mode bits that are related to the instruction fetch and instruction decode functions in the POWER8 core. The HID4 and HID5 Registers contain bits related to the load-store function in the POWER8 processor. All of these registers are only accessible in hypervisor mode. Reserved bits in the HID registers should not be set by software and may return either a zero or one value depending on the bit if set. Attempts to set some of these bits may enable functions that are no longer supported and thus could cause unpredictable behavior. Two values of the contents of each register are shown in the descriptions.

Initial state:

This is the state of the register after a normal scan-based POR. The actual and full POR sequence can set bits beyond the scan-based POR.

Preferred state:

This is the preferred state of the register for optimal performance and function.

1. The following sequence must be used when modifying HID0:

```
sync
mtspr HID0,Rx
mfsptr Rx,HID0
mfsptr Rx,HID0
mfsptr Rx,HID0
mfsptr Rx,HID0
mfsptr Rx,HID0
isync
```

After modifying HID0, executing six **mfsptr** instructions specifying HID0 as the source and specifying the same target GPR (Rx) in all six instructions is necessary to ensure that the modification is effective and the processor is in a valid state to continue executing subsequent instructions

2. The following sequence must be used when modifying HID1:

```
mtspr HID1,Rx
mtspr HID1,Rx
isync
```

Executing two **mtspr** instructions is necessary to ensure that updates to all portions of HID1 are completed before the **isync** instruction is completed.

3. The following sequence must be used when modifying HID4:

```
sync
mtspr HID4,Rx
isync
```

Because the ERATs contain entries for nonhypervisor real mode translations, when changing the RMOR value, the hypervisor must flush the ERATs (that is, by executing **slbia**) before passing control to the operating system. The en_rmsc (HID4[33]) bit must not be changed by the hypervisor, and its initial value must be included in the init file.



3.7.1 HID0 Register

Initial state: x'0005_0000_0000_0000'

Preferred state: x'0005_0000_0000_0000'

The HID0 Register is defined as follows.

Bits	Field Name	Description
0:1	reserved	Reserved.
2	4_lpar_mode	Core runs in the 4 LPAR mode where the core is partitioned into four logical partitions (sub-processors). Has precedence over 2_lpar_mode (bit 6). This is a read-only bit.
3:5	reserved	Reserved.
6	2_lpar_mode	Core runs in the 2 LPAR mode where core is partitioned into two logical partitions (sub-processors). This is a read-only bit.
7:10	reserved	Reserved.
11	1LPARto2LPAR	A write to this bit initiates a 1 LPAR to 2 LPAR dynamic mode switch. Cannot read this bit.
12	1LPARto4LPAR	A write to this bit initiates a 1 LPAR to 4 LPAR dynamic mode switch. Cannot read this bit.
13	dis_recovery	Disable processor recovery mechanism.
14	reserved	Reserved.
15	dyn_lpar_dis	Dynamic LPAR switching mode disabled.
16:18	reserved	Reserved.
19	HILE	Hypervisor Interrupt Little Endian. The HILE bit is set to '0' during system initialization and can be modified by hypervisor software as needed. The contents of the HILE bit are copied into the MSR[LE] by interrupts that set MSR[HV] to '1' to establish the endian mode for the hypervisor interrupt handlers.
20:30	reserved	Reserved.
31	en_attn	Enable support processor attn instruction.
32:63	reserved	Reserved.

3.7.2 HID1 Register

The HID1 Register contains additional mode bits that are related to the instruction fetch and instruction decode functions in the POWER8 processor core.

Initial state: x'0000_0000_0000_0000'

Preferred state: x'0000_0000_0000_0000'

The HID1 Register is defined as follows.

Bits	Field Name	Description
0	flush_ic	Flush the I-cache directory and the IEADIR on a transition from '0' to '1'.
1:4	reserved	Reserved.
5	dis_sp_itw	Disable speculative table walks. If this bit is set to a '1', ERAT miss handling is only carried out when the instruction with the ERAT miss is next-to-complete.
6:7	ierat_lru_mode(0:1)	I-ERAT size restrictor. These HID1 bits can be used to limit the number of I-ERAT entries to use. 00 Full I-ERAT (64 entries) - normal settings 01 Half I-ERAT (32 entries) 10 Quarter I-ERAT (16 entries) 11 Eighth I-ERAT (8 entries)
8	reserved	Reserved.
9	dis_dfp	Disable DFP. If this bit is set to '1,' the DFP instructions are treated as illegals.
10	reserved	Should always be set to '0'.
11	dis_user_priority_verylow_lpar0	Disable user priority change to very low - LPAR0. 0 Allow user mode to set thread priority to very-low. 1 Do not allow user mode to set thread priority to very low.
12	dis_user_priority_verylow_lpar1	Disable user priority change to very low - LPAR1 0 Allow user mode to set thread priority to very-low. 1 Do not allow user mode to set thread priority to very low.
13	dis_user_priority_verylow_lpar2	Disable user priority change to very low - LPAR2 0 Allow user mode to set thread priority to very-low. 1 Do not allow user mode to set thread priority to very low
14	dis_user_priority_verylow_lpar3	Disable user priority change to very low - LPAR3 0 Allow user mode to set thread priority to very-low. 1 Do not allow user mode to set thread priority to very low.
15	en_reduce_spec	Enable the reduce speculation mode in the IFU. This mode is used to reduce bandwidth beyond the L3 cache.
16:20	reserved	Reserved.
21:31	spare	Spare. These bits are implemented and thus a move-from returns the value of the last move-to.
32:63	Unimplemented spare	Unimplemented spares. These bits always return zeros.



3.7.3 HID4 Register

The HID4 Register controls the load-store functions in the POWER8 core.

Initial state: x'0000_0000_0000_0000'

Preferred state: x'0000_0004_4000_0000'

Bits	Field Name	Description
0:4	rmsc	<p>Real mode storage control field used exclusively for "Legacy RMSC" mode (HID4[33]='0'). If this field is set to a value N, the first $256 \times 2^{(N-1)}$ MB, for $0 < N < 18$, of real memory is considered well-behaved, and real memory accesses in this range are treated as nonguarded in MSR[IR, DR] = '0' mode. Beyond this range, accesses are treated as guarded for MSR[IR, DR] = '0' mode.</p> <p>'00000' = all MSR[IR, DR] = '0' accesses are treated as guarded.</p> <p>'00001' = MSR[IR, DR] = '0' accesses in the first 256 MB are nonguarded and beyond that guarded.</p> <p>'00010' = MSR[IR, DR] = '0' accesses in the first 512 MB are nonguarded and beyond that guarded.</p> <p>.....</p> <p>'10000' = MSR[IR, DR] = '0' accesses in the first 8 TB are nonguarded and beyond that guarded.</p> <p>'10001' = MSR[IR, DR] = '0' accesses in the first 16 TB are nonguarded and beyond that guarded.</p> <p>>'10001' = all MSR[IR, DR] = '0' accesses are treated as guarded.</p> <p>An sibia instruction is needed after changing this bit.</p>
5:6	reserved	Reserved.
7	dis_lpage	Disable large page support (only support 4 KB page).
8	dis_mpss	Disable mixed page segment support (4 KB and 64 KB, 4 KB and 16 MB, 64 KB and 16 MB).
9	dis_mpssx	Disable mixed page segment support extensions (4 KB and 16 MB, 64 KB and 16 MB).
10	dis_tspec	Disable speculative load tablewalks (stores are always nonspeculative).
11	dis_vpck	Disable virtual-page class keys exception.
12	en_hash2	Enable secondary hash tablewalks for speculative store instructions (if LPCR[TC] = '0').
13	tlbie_cc	Congruence class invalidation of TLB due to tlbie instruction. '0' Use the IS field of the tlbie instruction. '1' Override the IS field, and always invalidate the entire congruence class.
14	force_geq1	Force all load instructions to be treated as if they were to guarded (G = 1) storage.
15	nop_dcbt	dcbt is treated as a NOP (go through address translation, suppressed from going to Nest).
16	nop_dcbtst	dcbtst is treated as a NOP (go through address translation, suppressed from going to Nest).
17	reserved	Reserved.
18	rst_prf	Reset the data prefetch mechanism. Suppress subsequent prefetch requests and clear the stream detection logic, such that stream detection will not be affected by accesses performed before setting the bit back to '0'
19	sus_prf	Suspend the data prefetch mechanism Preserve the current state, but disable allocations, updates, and requests until the bit is set back to '0'.
20	dis_prf1	Disable L1 data prefetching.
21	dis_prf2a	Disable L1 data prefetching two lines ahead (only prefetch one line ahead, pertains to all streams).
22	dis_prf3l	Disable L3 load data prefetching.
23	dis_prf3s	Disable L3 store data prefetching.
24	dis_prfh	Disable data prefetching initiated by hardware (allow software streams).



Bits	Field Name	Description
25	dis_prfnh	Disable stride-n data prefetching initiated by hardware.
26	dis_prfns	Disable stride-n data prefetching initiated by software.
27	dis_shad	Disable data prefetch shadow register. 0 Prefetch stream direction always assumed 'up', but the PRQ can detect a confirmation in the 'down' direction. 1 Prefetch stream direction assumed 'up' if the address is in upper 3/4 of cache line, otherwise, down.
28	lmt_prf	Limit the number of available streams to four.
29	ex_dcbz	Exclude dcbz from initiating L3 store prefetches.
30:32	reserved	Reserved.
33	en_rmsc	Enable "Page-based RMSC mode" described in Power ISA Book III-S as "history blocks" implementation. Note: Setting this bit to '0' disables the page-based RMSC mode and enables the legacy RMSC mode, which divides real memory into a lower unguarded regions and an upper guarded region. The Legacy mode may not be supported in future POWER processors.
34	reserved	Reserved.
35	reserved	Reserved.
36	dis_rep_lp0	Control bit to disable the PTE time base and reference bit array updates for LPAR 0 0 Reference bit array update disabled. 1 Reference bit array update enabled.
37:39	sel_tb_lp0	Encode bits to select TB bit that provides low-order 12-bit TB value for LPAR 0 000 Reserved. 001 Select 0.5 second. 010 Select 1 second. 011 Select 2 seconds. 100 Select 4 seconds. 101 Select 8 seconds. 110 Select 16 seconds. 111 Select 32 seconds.
40	sel_pte_lp0	1 HCA (Hot / Cold Page Affinity) PTE format in use for LPAR 0. 0 Power ISA PTE format.
41	dis_rep_lp1	Control bit to disable PTE time base and ref-bit array updates for LPAR 1 0 Reference bit array update disabled. 1 Reference bit array update enabled.
42:44	sel_tb_lp1	Encode bits to select TB bit that provides low-order 12-bit TB value for LPAR 1 000 Reserved. 001 Select 0.5 second. 010 Select 1 second. 011 Select 2 seconds. 100 Select 4 seconds. 101 Select 8 seconds. 110 Select 16 seconds. 111 Select 32 seconds.
45	sel_pte_lp1	1 HCA (Hot / Cold Page Affinity) PTE format in use for LPAR 1. 0 Power ISA PTE format.
46	dis_rep_lp2	Control bit to disable PTE time base and reference bit array updates LPAR 2. 0 Reference bit array update disabled. 1 Reference bit array update enabled.



Bits	Field Name	Description																
47:49	sel_tb_lp2	<p>Encode bits to select TB bit that provides low-order 12-bit TB value LPAR 2.</p> <table> <tr><td>000</td><td>Reserved.</td></tr> <tr><td>001</td><td>Select 0.5 second.</td></tr> <tr><td>010</td><td>Select 1 second.</td></tr> <tr><td>011</td><td>Select 2 seconds.</td></tr> <tr><td>100</td><td>Select 4 seconds.</td></tr> <tr><td>101</td><td>Select 8 seconds.</td></tr> <tr><td>110</td><td>Select 16 seconds.</td></tr> <tr><td>111</td><td>Select 32 seconds.</td></tr> </table>	000	Reserved.	001	Select 0.5 second.	010	Select 1 second.	011	Select 2 seconds.	100	Select 4 seconds.	101	Select 8 seconds.	110	Select 16 seconds.	111	Select 32 seconds.
000	Reserved.																	
001	Select 0.5 second.																	
010	Select 1 second.																	
011	Select 2 seconds.																	
100	Select 4 seconds.																	
101	Select 8 seconds.																	
110	Select 16 seconds.																	
111	Select 32 seconds.																	
50	sel_pte_lp2	<table> <tr><td>1</td><td>HCA (Hot / Cold Page Affinity) PTE format in use for LPAR 2.</td></tr> <tr><td>0</td><td>Power ISA PTE format.</td></tr> </table>	1	HCA (Hot / Cold Page Affinity) PTE format in use for LPAR 2.	0	Power ISA PTE format.												
1	HCA (Hot / Cold Page Affinity) PTE format in use for LPAR 2.																	
0	Power ISA PTE format.																	
51	dis_rep_lp3	<p>Control bit to disable PTE time base and ref-bit array updates LPAR 3</p> <table> <tr><td>0</td><td>Reference bit array update disabled.</td></tr> <tr><td>1</td><td>Reference bit array update enabled.</td></tr> </table> <p>Note: In single LPAR mode, all four partitions (en_rep_lpx and sel_tb_lpx and sel_pte_lpx) must be set to the same value. In 2 LPAR mode: en_rep_lp[0,1] and sel_tb_lp[0,1] and sel_pte_lp[0,1] must be set to the same value; and en_rep_lp[2,3] and sel_tb_lp[2,3] and sel_pte_lp[2,3] must be set to the same value.</p>	0	Reference bit array update disabled.	1	Reference bit array update enabled.												
0	Reference bit array update disabled.																	
1	Reference bit array update enabled.																	
52:54	sel_tb_lp3	<p>Encode bits to select TB bit that provides low-order 12-bit TB value LPAR 3.</p> <table> <tr><td>000</td><td>Reserved.</td></tr> <tr><td>001</td><td>Select 0.5 second.</td></tr> <tr><td>010</td><td>Select 1 second.</td></tr> <tr><td>011</td><td>Select 2 seconds.</td></tr> <tr><td>100</td><td>Select 4 seconds.</td></tr> <tr><td>101</td><td>Select 8 seconds.</td></tr> <tr><td>110</td><td>Select 16 seconds.</td></tr> <tr><td>111</td><td>Select 32 seconds.</td></tr> </table>	000	Reserved.	001	Select 0.5 second.	010	Select 1 second.	011	Select 2 seconds.	100	Select 4 seconds.	101	Select 8 seconds.	110	Select 16 seconds.	111	Select 32 seconds.
000	Reserved.																	
001	Select 0.5 second.																	
010	Select 1 second.																	
011	Select 2 seconds.																	
100	Select 4 seconds.																	
101	Select 8 seconds.																	
110	Select 16 seconds.																	
111	Select 32 seconds.																	
55	sel_pte_lp3	<table> <tr><td>1</td><td>HCA (Hot / Cold Page Affinity) PTE format in use for LPAR 0.</td></tr> <tr><td>0</td><td>Power ISA PTE format.</td></tr> </table>	1	HCA (Hot / Cold Page Affinity) PTE format in use for LPAR 0.	0	Power ISA PTE format.												
1	HCA (Hot / Cold Page Affinity) PTE format in use for LPAR 0.																	
0	Power ISA PTE format.																	
56:63	reserved	Reserved.																

3.7.4 HID5 Register

Initial state: x'0000_0000_0000_0000'
Preferred state: x'4000_0001_0000_0000'

The HID5 Register is defined as follows.

Bits	Field Name	Description
0:12	reserved	Reserved.
13:14	Throt_Dpf	Hypervisor memory bandwidth control for data prefetch. 00 Medium 01 Low 10 High 11 Reserved
15	reserved	Reserved.
16:29	reserved	Reserved.
30:31	pte_plus_n	PTE Plus n prefetching where n = 1 or 3. 00 Disable 10 PTE + 1 01 PTE + 3 11 Undefined
32:63	reserved	Reserved.

3.7.5 Real Mode Offset (RMO) Region Sizes

The following RMO sizes are available for the POWER8 processor. The RMLS[34:37] field in the LPCR defines the RMO sizes, as described below.

- 1000 - 32 MB
- 0011 - 64 MB
- 0111 - 128 MB
- 0100 - 256 MB
- 0010 - 1 GB
- 0001 - 16 GB
- 0000 - 256 GB



3.7.6 Hypervisor Real Mode Offset (HRMO) Register Update Sequence

Table 3-9. HRMOR Update Sequence

Master	Slave
Thread sync up point 1	Thread sync up point 1
Change HID4[0:4] = 0000 (default in the POWER8 processor)	
EA[0] = 1	EA[0] = 1
Thread sync up point 2	Thread sync up point 2
Change HRMOR	
Thread sync up point 3	Thread sync up point 3
isync	isync
sibia0	sibia0
isync	isync
Thread sync up point 4	Thread sync up point 4

3.7.7 Core-to-Core Trace SPR

The Trace SPR is used to access enhanced instruction trace information from the processor core trace logic. This 64-bit register is read only and has a privileged read access. There is a protocol associated with the use of this register to coordinate gathering instruction trace images from the other processor core.

3.7.8 Trigger Registers

Writes to the trigger registers, named TRIG0, TRIG1, and TRIG2, can be inserted in the instruction stream to cause triggers to the on-chip trace array debug logic. These are intended to be used for lab debug and bringup only and architecturally behave as a NOP.

3.7.9 IMC Array Access Register

The Instruction Match CAM (IMC) array facility is used for performance monitoring instrumentation and for the soft patch of instructions (this latter use is restricted for the support processor and is not available through the SPR access to this register array). The array has privileged write access and user-level read access via this SPR. Writes to the register array are used to configure the IMC, and reads return information about the availability of registers within the facility.

3.7.10 Performance Monitor Registers

The performance monitor counter registers (PMC1 - PMC6), the performance monitor control registers (MMCR0, MMCR1, MMCRA), and the sampled address registers (SIAR, SDAR) are supported in the POWER8 processor core. The performance monitor counter registers PMC7 and PMC8 are not implemented in the POWER8 processor core (an operation for these two performance counter registers will be treated as NOP).



3.7.11 Other Fixed-Point Instructions

The POWER8 processor core supports both the 32-bit **mtmsr** instruction and the 64-bit **mtmsrd** instruction.

The POWER8 processor core works to optimize the **mtmsr** instruction to help speed up the cases where little or no synchronization is required (such as, updates to the EE bit).

Software must avoid placing **mtmsr** and **mtmsrd** instructions that change the SF bit at address x'00000000FFFFFFFC' or x'FFFFFFFFFFFFFFFC'.

3.8 Storage Control

3.8.1 Virtual and Physical Address Ranges Supported

The POWER8 processor core supports a 68-bit virtual address and a 50-bit physical (real) address.

3.8.2 Data Effective-to-Real-Address Translation (D-ERAT)

The POWER8 processor core includes a primary ERAT and secondary ERAT. The primary ERAT has four copies of 48 entries. It is a fully-associative D-ERAT, each with binary LRU replacement policy (separate LRU for each half in SMT2/4/8 mode), data effective-to-real address translation (D-ERAT) cache for fast translation of data effective addresses into physical (real) addresses.

In the POWER8 core, each entry of the D-ERAT contains translation information for a 4 KB, 64 KB, or 16 MB block of effective storage, depending on the page size being used. Reference to a 1 MB or 16 GB page uses 64 KB or 16 MB entries in the D-ERAT, respectively.

In ST mode, 48 entries are available, and all four copies have the same content. In SMT2, SMT4, and SMT8 mode, 48 entries are available for threads on LS half 0, and 48 entries for LS half 1. The two ERATs on a half have the same contents.

D-ERAT entries are created for real mode (MSR[DR] = '0') and are shared by all threads in the same partition. Hypervisor real-mode entries are shared by all threads. Entries are invalidated via a CAM to clear the function in support of **sbia**, **slbie**, **tlbie**, and **mtsr** instructions. It is parity and multi-hit protected, reports to the FIR and completion on an error detect, hardware recovery via ABIST, and multi-hit is reported as a machine check interrupt.

The secondary D-ERAT is the second-level translation cache for data addresses. It has 256 entries managed in two 128-entry halves. It is a fully associative with FIFO replacement policy. In ST mode, 256 entries are available. In SMT2, SMT4, and SMT8 mode, 128 entries are available for threads on LS half 0, and 128 entries for LS half 1.

According to the Power ISA, aliasing the I-bit storage attribute is prohibited. In the POWER8 core, due to the caching of pages in the ERATs, software should avoid accessing the same real page with different values for the I-bit storage attribute. Failure to follow this restriction may result in a cache paradox or other boundedly undefined behavior.



D-ERAT I and G Bit Setting

The G and I bits are set based on *Table 3-10*.

Table 3-10. D-ERAT I and G Bit Setting

Condition									Resulting Action
MSR [DR]	MSR [HV]	LPCR[0]	HID4[33]	First access HV CI Instruction	HID4[0:4] RMSC (above line)	I	G		
1	x	x	x	x	x	PTE(I)	PTE(G)	An entry is created with the I and G values set from the PTE.	
0	0	1	x	x	x	PTE(I)	PTE(G)	Virtual real mode. An entry is created with the I and G values set from the PTE.	
0	0	0	x	x	x	0	0	RMOR mode. Set IG = '00'	
0	1	x	0	yes	no	1	0	Legacy RMSC mode. If the first access is caused by a hypervisor CI load or store (for example, ldcix , stdcix , and so on) is below the guarded line established by HID4[0:4], storage is G = '0' and CI load/store access causes a DSI and set DSISR(62) = '1'.	
0	1	x	0	no	no	0	0	Legacy RMSC mode. If the first access is caused by any instruction other than a hypervisor CI load or store is below the guarded line established by HID4[0:4], storage is G = '0' and an entry is established as I = '0' and G = '0'.	
0	1	x	0	no	yes	0	1	Legacy RMSC mode. If the first access is caused by any instruction other than a hypervisor CI load or store is above the guarded line established by HID4[0:4], storage is G = '1' and an entry is established as I = '0' and G = '1'.	
0	1	x	0	yes	yes	0	1	Legacy RMSC mode. If the first access is caused by a hypervisor CI load or store (for example, ldcix , stdcix , and so on) is above the guarded line established by HID4[0:4], storage is G = '1' and an entry is established as I = '1' and G = '1'.	
0	1	x	1	yes	x	1	1	Page-based RMSC mode. If the first access is caused by a hypervisor CI load or store (for example, ldcix , stdcix , and so on), an entry is established as I = '1' and G = '1'.	
0	1	x	1	no	x	0	0	Page-based RMSC mode. If the first access is caused by any instruction other than a hypervisor CI load or store, storage is G = '0' and an entry is established as I = '0' and G = '0'.	

In addition to the conditions shown in *Table 3-10*, the G bit is forced to a '1' for the following cases:

- HID4[33] = 0, HV = 1, and the referenced address is greater than 16 TB.
- HID4[33] = 0, HV = 1, and the RMSC value is '00000' or greater than '10001' (Illegal value).
- HV = 1, PR = 0, EA(0) = 1, and EA(14 to 33) = '11111111111111111111'.



The I bit is forced to a '1' for the following case:

- HV = 1, PR = 0, EA(0) = 1, and EA(14 to 33) = '111111111111111111'.

Caching-Inhibited Paradox Cases

If a caching-inhibited load instruction hits in the L1 data cache, the load data is serviced from the L1 data cache and no request is sent to the NCU.

If a caching-inhibited store instruction hits in the L1 data cache, the store data is written to the L1 data cache and sent to the NCU. Note that the L1 data cache and L2 cache are no longer coherent.

These scenarios are true both for instructions marked as caching-inhibited by the PTE I-bit and for the hypervisor caching-inhibited load and store instructions.

The POWER8 core supports the page-based real-mode storage control (RMSC) mechanism that allows speculative access to DR = 0 space, if there is real memory there.

Because the content of each D-ERAT entry is the result of a page table search based on the contents of an SLB entry, to maintain consistency with the SLB (or segment registers), the following instructions will cause all entries that belong to the thread in the D-ERAT to be invalidated (these instruction do not invalidate D-ERAT entries of the other thread).

- **mtrs** or **mtrsln** instructions - used for segment register changes in 32-bit operating systems
- **slbie**

The **slbie** instruction causes invalidation of a D-ERAT entry belonging to the thread (no impact to the other thread) only if there is a perfect address match: EA[0:35] are matched for **slbie** for a small (256 MB) segment and EA[0:23] are matched for **slbie** for a big (1 TB) segment.

The **tlbie** instruction (or the detection of snooped-tlbie operations) invalidates all D-ERAT entries (irrespective of the thread) in the D-ERAT that have a perfect match. In other words, the entry is invalidated only if:

- EA[36:51] are matched for tlbie small page
- EA[36:47] are matched for tlbie 64 KB page
- EA[36:43] are matched for tlbie 1 MB page
- EA[36:39] are matched for tlbie 16 MB page
- EA[24:29] are matched for tlbie 16 GB page

Upon power-on, each D-ERAT entry is set to the invalid state.

3.8.3 Translation Lookaside Buffer (TLB)

The POWER8 core contains a unified (combined for both instruction and data), 2048-entry, 4-way set-associative TLB (LRU-based replacement algorithm). In addition, the POWER8 core contains one 64-entry, fully-associative I-ERAT (single-level effective-to-real translation) and four 48-entry, fully-associative D-ERATs. The TLB is a cache of recently-used page table entries, and the ERATs are caches that contain translations derived from information in the page table and SLB. The TLB and ERATs are loaded and managed by hardware.

In the POWER8 core, the TLB entry stores the logic partition ID (LPID) in each TLB entry to indicate which partition loaded that TLB entry. Because the virtual and real address space are the same for all software threads within a logical partition, the TLB, which keeps the mapping from virtual-to-real address space, are completely shared by the threads within a partition and there is no thread-ID bit needed in the TLB to identify



which entry belongs to which thread. Different partitions have different mappings from virtual-to-real address space, and, because of this, TLB entries cannot be shared between partitions. A given entry in the TLB can be used by all the threads within a partition at the same time (however, a given entry in I-ERAT or D-ERAT, cannot be shared by the two threads at the same time because the ERATs also contains information from the effective to virtual address mapping), unless the entry is created by bypassing the SLB (that is, IR = 0 or DR = 0 addresses). Threads in different partitions are not able to access TLB entries from another partition.

In the POWER8 core, TLB is indexed with the following hashed address (this includes reads from the TLB, writes to the TLB, tlbie instructions, and snooped tlbie transactions).

Table 3-11. 256 MB Segments

Page Size	Index
4 KB	[VSID(45:49) XOR EA(43:47)] II EA(48:51)
64 KB	[VSID(45:49) XOR EA(39:43)] II EA(44:47)
1 MB	[VSID(45:49) XOR EA(36:40)] II EA(40:43)
16 MB	VSID(45:49) II EA(36:39)
16 GB	Does not exist

Table 3-12. 1 TB Segments

Page Size	Index
4 KB	[VSID(33:37) XOR EA(43:47)] II EA(48:51)
64 KB	[VSID(33:37) XOR EA(39:43)] II EA(44:47)
1 MB	[VSID(33:37) XOR EA(35:39)] II EA(40:43)
16 MB	[VSID(33:37) XOR EA(31:35)] II EA(36:39)
16 GB	[VSID(33:37) XOR EA(24:28)] II EA(26:29)

The POWER8 core supports hardware update of the storage access recording bits (reference and change) into the memory-based page table.

The Power ISA has been enhanced to add an L bit, to indicate large page in the page table entry (PTE). This bit is compared with the L bit of the SLB in translation mode for a hit during a table walk to reload the TLB on a TLB miss.

The POWER8 core supports TLB hit under miss and four table concurrent table walks. The POWER8 core also supports two outstanding I-ERAT misses (from the eight threads) and four outstanding D-ERAT misses at the same time.

The POWER8 core supports lockless TLBIE operations. The architectural requirement that only one thread at a time can execute **tlbie/tlbsync** instructions during a page table modification (section Page Table Updates of Power ISA Book III) need not be followed. This was traditionally implemented with a single global lock for the entire page table modification sequences. The term “lock-less” TLBIEs refers to the POWER8 core’s ability to manage concurrent **tlbie/tlbsync** sequences from multiple threads without this global lock.

However, software must still ensure that concurrent, conflicting, racing PTE updates from more than one thread do not occur (the hardware performs the updates in some fashion, but the end results is undefined due to the racing of the updates to the same PTE entry) and therefore, software locks or some other synchronization discipline are still required to prevent these collisions as necessary.”

The execution of **tlbie** instructions or the detection of **snooped-tlbie** operations off the bus cause an index-based invalidate to occur in the TLB, if there is a match. In other words, an entry is invalidated only if there is a perfect match of the effective address supplied by the **tlbie** operation and the content of the TLB entry.

The POWER8 core does not support the **tlbia** instruction.

Upon power-on, the POWER8 core initializes each TLB entry to the invalid state.

3.8.4 Large-Page Support

In addition to the normal 4 KB page size, the POWER8 core provides support for 64 KB, 1 MB, 16 MB, and 16 GB pages. Translation information for all these page sizes is kept in the TLB. Irrespective of the page size, a given page takes up only one entry in the TLB.

If a virtual address is mapped into a small (large) page and then later on mapped into a large (small) page without invalidating TLB entries between changing page size, a machine check interrupt can result with an indication that a parity error occurred when the TLB was accessed to translate an effective address. The error condition can be corrected by invalidating the entire TLB and SLB.

POWER8 also supports Multiple Page Sizes Per Segment (MPSS) as described in the Power ISA. Specifically, POWER8 supports mixing page size in a single segment with the following combinations only:

4K base / 64K actual

4K base / 16M actual

64K base / 16M actual

Table 3-13 shows the correspondence between PTE[L LP] values and SLBE[L LP] values.

Table 3-13. PTE and SLBE Correspondence

Entry Number	PTE			SLBE		Base Page Size	Actual Virtual Page Size	Notes
	L	LP		L	LP			
1	0	rrrr	rrrr	0	00	4 KB	4 KB	1
2	1	0000	0000	1	00	16 MB	16 MB	
3	1	rrrr	0001	1	01	64 KB	64 KB	2
4	1	0000	0010	1	11	1 MB	1 MB	
5	1	0000	0011	1	10	16 GB	16 GB	
6	1	rrrr	0111	0	00	4 KB	64 KB	1
7	1	0000	1000	1	01	64 KB	16 MB	2
8	1	0011	1000	0	00	4 KB	16 MB	1

1. Entries 1, 6, and 8 all use SLB[L LP] = 000 encoding for base page size 4 KB but have unique PTE[L LP] encodings for actual page size.
2. Entries 3 and 7 both use SLB[L LP] = 101 encoding for base page size of 64 KB but have unique PTE[L LP] encodings for actual page size.
3. Unimplemented SLB page size encodings are treated the same as the '000' case.
4. If the SLB page size is '110' (16 GB) and the segment size is small (256 MB), hardware treats the SLB page size the same as the '000' case.
5. The 'r' bits are part of the real page number for actual pages sizes less than 1 MB. They can be any value.



3.8.5 PTE Prefetching

The POWER8 core supports the prefetching of PTE entries. This feature is enabled by setting the values in HID5[30:31] as follows:

- 00 PTE prefetching is disabled.
- 01 PTE prefetching prefetches three additional PTE entries.
- 10 PTE prefetching prefetches one additional PTE.
- 11 Undefined.

This PTE prefetch feature works as follows. When a tablewalk is initiated to bring in a new PTE entry, the VPN of that tablewalk request is incremented by one and a request is sent to bring the PTE group associated with the incremented VPN into the L3 cache. Only the PTE groups associated with primary table walks are requested. The VPN is only incremented in the lower 2 bits. If the incremented VPN would have crossed this boundary, the lower two bits are set to '0'. If PTE prefetching is configured for prefetching three additional PTE groups and the boundary is crossed, the prefetching logic continues to increment the lower 2 bits until all three prefetches are sent.

3.8.6 Segment Lookaside Buffer (SLB)

The POWER8 core contains a unified (combined for both instruction and data), 32-entry, fully associative SLB per thread (eight per processor core). Information derived from the SLB can also be cached in the I-ERAT or the D-ERAT along with information from the TLB. As a result, many of the SLB management instructions have effects on the ERATs as well as on the SLB itself. The POWER8 core supports the 1 TB segment size, in addition to the usual 256 MB segments. Bit 0 of the SLB[B] field is ignored by POWER8 and should be always set to '0'b per the Power ISA for unimplemented segment size encodings.

Because the SLB is managed by software (the operating system), it is possible that multiple entries can be incorrectly set up to provide translations for the same effective address. If an effective address is translated by more than one SLB entry (that is, the ESID fields of the entries are identical), a machine check interrupt results with an indication that a parity error occurred when the SLB was accessed. When this happens the hardware logically OR's the data in the conflicting entries. The machine check handler can look at the SLB contents to try to determine if conflicting entries have been provided. When a parity error occurs not due to multiple entries, the entire SLB must be reloaded because the DAR does not contain an address indicating which entry caused the parity error. If the source of the error was due to multiple entries, the conflicting entries must be corrected for the translation to proceed, which might also be accomplished by reloading the entire SLB with good entries.

3.8.7 Address Space Register

The Address Space Register (ASR) has been removed from the Power ISA.



3.8.8 Support for 32-Bit Operating Systems

The POWER8 core supports the optional bridge facilities and instructions for 64-bit implementations described in the Bridge-to-SLB Architecture section of the *Power ISA Book III-S version 2.07*.

Associated with this support, the following optional instructions are supported:

- **mtsr** - Move to segment register
- **mtsri** - Move to segment register indirect
- **mfsr** - Move from segment register
- **mfsri** - Move from segment register indirect
- **mtmsr** - Move to machine state register (32-bit)

3.8.9 Reference and Change Bits

The POWER8 hardware performs *reference* and *change* bit updates to the page table.

The W and M bits in the PTE are assumed to be '01' respectively. If the Change bit is updated, the W and M bit in the PTE are set to '01' respectively.

The POWER8 core provides a mode bit for determining whether or not speculative load instructions should reload the TLB in the event of a miss (HID4[10]). If this mode bit is set to allow this behavior, the reference bit can be set on behalf of speculative loads (that is, ones that never actually complete from the programs perspective). Under no circumstances will the POWER8 processor core speculatively set the page table change bit.

In the POWER8 core, load instructions that have a TLB miss but are denied access by storage protection, cause a DSI, cause a DAWR DSI, cause an I = 1 DSI, or cause an alignment interrupt that still causes the reference bit for the subject page to be set. Similarly, store type instructions that cause a DSI, cause a DAWR DSI, cause an I = 1 DSI, or cause an alignment interrupt, set the change bit. On the other hand, the change bit is not set, if a store instruction is denied access by page protection exception.

3.8.10 Storage Protection

The architecture defines whether the instruction fetch is permitted from a page marked "no access" as implementation dependent. In the POWER8 processor core, these instruction fetches are permitted to continue without signaling an exception. The POWER8 core supports storage protection modes in the Power ISA, with 32 virtual-page class keys. AMOR and UAMOR are implemented as 32-bit SPRs.

3.8.11 Block Address Translation

Although this facility existed in earlier versions of the architecture, it is no longer part of the Power ISA. As a result, the POWER8 core does not support block address translation.



3.8.12 Real Mode Storage Control

The POWER8 core supports the ability to control cacheability of data and instruction accesses while in real mode.

For the D-side and I-side, in HV = 1, the translation is loaded with G = 0 and 16 MB page size.

The POWER8 core supports RMSC for data storage and instruction storage. The legacy RMSC behavior is also available under a HID4 bit control. The real memory in a system is often noncontiguous and the hypervisor data and instruction storage accesses can be scattered across the address space. The page-based RMSC architecture and implementation allows speculative access safely in system memory. The first time the access is made in DR = 0 and IR = 0 mode, it is done nonspeculatively. After the first access, a proper D-ERAT entry and I-ERAT entry is established. Subsequent accesses to such a D-ERAT and I-ERAT entry ensure that the access is made to system memory and therefore, can be done speculatively, providing higher performance.

To change to or from the legacy RMSC behavior, an **slbia** is needed after the HID4 Register update.

3.8.13 Storage Access Modes - WIMG Bits

The POWER8 core always assumes W = 0 and M = 1 independent of the value of these bits in the page table. Furthermore, when the hardware is performing a change bit update, it writes the W and M bits as W = 0 and M = 1. Per the Power ISA, accessing a page as both I = 0 and I = 1 is boundedly undefined. Software should avoid aliasing the I-bit on a page basis.

Table 3-14 summarizes the treatment of the WIMG bits in the POWER8 core.

Table 3-14. WIMG Bits

WIMG	Description
x1xx (except 1110)	Treated as WIMG = 0111
x0x1	Treated as WIMG = 0011
x0x0	Treated as WIMG = 0010
1110	Treated as WIMG = 0010 and accesses are strongly ordered

For the noncacheable unit (NCU), the IG combination has the following meaning in the POWER8 core to control the store ordering and store gathering.

Table 3-15. IG Bits

IG	Description
11	No gather, no reorder in NCU is allowed
10	Gather, reorder in NCU is allowed

In IG = '11' mode, cache-inhibited loads cannot be reordered relative to loads, and cache-inhibited stores cannot be reordered relative to cache-inhibited stores. Cache-inhibited loads can be reordered relative to cache-inhibited stores and vice-versa (if it is necessary to maintain ordering between loads and stores, barrier instructions must be used). There is no defined ordering between cache-inhibited load or store operations from different threads.

In IG = '11' mode, gathering is not permitted for either load or store operations within or between threads.

In IG = '10' mode, cache-inhibited loads or stores from a given thread can be gathered and can be reordered. This mode allows for higher performance with a certain loss of control of the order in which the operations are completed or whether operations are gathered (barriers can be used where necessary to re-establish order). There is no defined ordering between cache-inhibited load or store operations from different threads.

3.8.14 Speculative Storage Accesses

The POWER8 processor core can execute load instructions to nonguarded storage speculatively. This can occur when a load instruction is encountered on a predicted branch path, or when a logically preceding instruction causes an interrupt. As a result, it is possible for a speculative load that misses in the on-chip cache hierarchy to initiate an external storage request even if that load instruction is not actually executed as part of the true instruction stream.

3.8.15 mtsr, mtsrin, mfsr, and mfsrin Instructions

Most of the optional Power ISA bridge support for 32-bit operating systems is supported by the POWER8 core. As part of that support, the **mtsr**, **mtsrin**, **mfsr**, and **mfsrin** instructions are supported.

3.8.16 TLB Invalidate Entry (tlbie and tlbiel) Instructions

The POWER8 processor core implements the **tlbie** and **tlbiel** instruction described in the architecture. Both of these instructions support small and large pages. The **tlbiel** instruction is never sent outside the processor core.

The **tlbie** (large or small page) instruction performs an index-based (congruence class) invalidate of the TLB (if there is a perfect match of the effective address of the **tlbie** operation and the content of the TLB entry). Both I-ERAT and D-ERAT entries are invalidated when there is a perfect match.

The **tlbsync** instruction is used to synchronize the completion of the **tlbie** instruction. Only one **tlbsync** instruction is required to synchronize the completion of a group of **tlbie** instructions.

The POWER8 core also uses the appropriate address bits from the TLBIE transaction to index into both ERATs and then invalidates an entry (if any), that matches the effective address.

The POWER8 core also supports two groups of pages TLBIE instructions: one for eight consecutive 4 KB pages TLBIE with AP encode RB[56:58] = '110' and another one for eight consecutive 64 KB pages TLBIE with LP encode RB[46:51] = '001010'.

The **tlbie** is a hypervisor-only instruction. An attempt to execute it while not in hypervisor mode causes a privileged type of program interrupt. In the POWER8 core, the **tlbiel** is a privileged instruction. Any attempt to execute it while in the problem state causes a privileged type of program interrupt.



Table 3-16 shows the legal segment size and page size specifications for **tlbie** and **tlbiel** for the POWER8 core when L = '0'.

Table 3-16. Segment Size and Page Size Specifications for tlbie and tlbiel (L = 0)

RB[49:51]	RB{54:55} Segment Size	RB[63] L	RB[56:58] AP (Same as SLB[L I LP] Encoding)	Actual Page Size to be Invalidated
vvv	00	0	000	4 KB
vvv	00	0	101	64 KB
vvv	00	0	100	16 MB
000	00	0	110	8 consecutive 4 KB pages aligned on 32 KB boundary
vvv	01	0	000	4 KB
vvv	01	0	101	64 KB
vvv	01	0	100	16 MB
000	01	0	110	8 consecutive 4 KB pages aligned on 32 KB boundary

1. All other values must not be used when L = '0'.
 2. RB[54:55] = '00' corresponds to a 256 MB segment size and RB[54:55] = '01' corresponds to 1 TB segment size.
 3. 16 GB page with a small segment (RB(54:55)= '00') is *not* a permitted combination.

Table 3-17 shows the legal segment size and page size specifications for **tlbie** and **tlbiel** for the POWER8 core when L = '1'.

Table 3-17. Segment Size and Page Size Specifications for tlbie and tlbiel (L = 1)

RB[54:55] Segment Size	RB[63] L	RB[56:58] AP (same as SLB[L I LP] encoding)	Base Page Size	Actual Page Size to be Invalidated
00	1	0000 0000	16 MB	16 MB
00	1	VVVV 0001	64 KB	64 KB
00	1	0000 0010	1 MB	1 MB
00	1	0000 1000	64 KB	16 MB
00	1	v000 1010	64 KB	8 consecutive 64 KB pages on 512 KB boundary
01	0	0000 0000	16 MB	16 MB
01	1	VVVV 0001	64 KB	64 KB
01	1	0000 0010	1 MB	1 MB
01	1	0000 0011	16 GB	16 GB
01	1	0000 1000	64 KB	16 MB
01	1	v000 1010	64 KB	8 consecutive 64 KB pages on 512 KB boundary

1. All other values must not be used when L = '1'.
 2. 'v' corresponds to AVA (AVPN) bits (and thus can be any value).
 3. RB[54:55] = '00' corresponds to 256 MB segment size and RB[54:55] = '01' corresponds to 1 TB segment size.
 4. 16 GB page with a small segment (RB[54:55] = '00') is *not* a permitted combination.

3.8.17 TLB Invalidate All (**tlbia**) Instruction

The **tlbia** instruction is not implemented in the POWER8 core and if detected causes a hypervisor emulation assistance interrupt. The effects of the instruction can easily be emulated by executing a series of **tlbie** instructions (512 in the POWER8 core, for the 512 congruence classes) by incrementing the effective address bits [43:51] through their full range, and by setting the IS field of the **tlbie** instruction to the appropriate values as described in the Power ISA. To invalidate all entries irrespective of the LPAR ID, MSR[HV] must equal '1'. A special HID4 bit can be used to force the core to ignore the IS field and always invalidate the entire congruence class.

3.8.18 TLB Synchronize (**tlbsync**) Instruction

On a given thread, the **tlbsync** instruction forces any previous **tlbie** instructions to complete before the **tlbsync** is allowed to complete. The instruction is implemented as described in the Power ISA.

3.8.19 Page Replacement Policy

The POWER8 core supports an optional PTE format, which must be used when the optional page replacement policy (hot/cold page affinity) feature is enabled by setting HID4 bits: 40 (LPAR0), 45 (LPAR1), 50 (LPAR2), and 55 (LPAR3). A reference-history array in memory contains Reference-bit values for recent sampling periods. On the first TLB miss for a page in the sampling period, the reference-bit array is right shifted (with zero fill) by the number of sample periods since the last TLB miss, and then the high-order bit is set to a one.

The PTE entry has been reorganized to support this function.

- B bits in PTE DW0 bits [0:1] move to PTE DW1 bits [4:5]
- REF_ARRAY enable bit in PTE DW1 bit 6
- 12 bits PTE_UPD_TIME in PTE DW0 bits [0:11]

The LSU examines PTE_UPD_TIME stamp during tablewalk TLB reload. A 12-bit CPU_UPD_TIME is generated from time base + real address bits to compare against the PTE_UPD_TIME.

HID4[37:39] define sample periods from 0.5 second to 32 second.

- 000 - Reserved
- 001 - 0.5 second
- 010 - 1 second
- 011 - 2 seconds
- 100 - 4 seconds
- 101 - 8 seconds
- 110 - 16 seconds
- 111 - 32 seconds

HID4[36] controls the PTE time base and reference-bit array updates; '1' enables a reference-bit array update and '0' disables a reference-bit array update. The HID4[40] selects new or old PTE format; '1' for new PTE format and '0' for old PTE format.



Time-base bits, TB[18:43], are shadowed in the LSU. TB[35] is the 0.5 second bit and TB[29] is the 32 second bit. The high-order 12-bit CPU_UPD_TIME is generated from the 20-bit adder as follows.

- 0.5 second - TB[24:43] + [00000000000000||Real Addr(40:47)]
- 1 second - TB[23:42] + [000000000000||Real Addr(40:47)]
- 2 second - TB[22:41] + [000000000000||Real Addr(40:47)]
- 4 second - TB[21:40] + [000000000000||Real Addr(40:47)]
- 8 second - TB[20:39] + [000000000000||Real Addr(40:47)]
- 16 second - TB[19:38] + [000000000000||Real Addr(40:47)]
- 32 second - TB[18:37] + [000000000000||Real Addr(40:47)]

If CPU_UPD_TIME does not equal PTE_UPT_TIME, the LSU sends the 12-bit CPU_UPD_TIME with PTE address to the nest through the store port to update both the PTE and Reference-bit arrays. The TTYPE = '100000' is for PTE update and TTYPE = '100011' is for Reference-bit array update.

Table 3-18 defines how the Reference-bit array is updated.

Table 3-18. Reference-Bit Array Update

HID4[40] Page Based PTE Format	HID4[36] Enable Reference-Bit Array Function	PTE DW1[6] Enable Reference-Bit Array Update	CPU_UPD_TIME = PTE_UPT_TIME	Reference-Bit Array Update
0	0	X	X	Not Updated
0	1	X	X	Not Updated
1	0	X	X	Not Updated
1	1	0	X	Not Updated
1	1	1	1	Not Updated
1	1	1	0	Updated

3.8.20 Support for Store Gathering

The POWER8 core performs gathering of cacheable stores to reduce the store traffic into the L2 cache. For cacheable stores, the gathering occurs in L2 store queues that sit above the L2 cache. The store queue is shared by the threads. The store queue is comprised of two banks of sixteen 64-byte wide, fully-associative entries or gather stations. Stores can be gathered while architecturally permitted (that is, there is no intervening barrier operation) and the matching address is valid in the store queue. The conditions for pushing the store queue data into the L2 cache are not visible to the programmer.

Gathering of cache-inhibited stores is also supported and can be disabled with a mode bit in the noncacheable unit (NCU) configuration register. There are sixteen 64-byte gather stations in the NCU.

3.8.21 Cache Coherency Paradoxes

Accesses to a given cache line as both cacheable and caching inhibited are not supported in either the Power ISA nor the POWER8 chip. Because the value of the I-bit is cached by the ERATs inside the processor core, cacheable accesses may be performed speculatively and thus, software should avoid alias the I-bit (that is, caching-inhibited bit) on a per page basis. Failure for software to adhere to this restriction can lead to cache corruption.



3.8.22 Handling Parity Error, Multi-Hit, and Uncorrectable Errors

3.8.22.1 Parity Error

If there is a parity error in the D-cache, I-cache, D-ERAT, I-ERAT, TLB or several other register files, SRAM dataflow or control structures (but not the SLB), the POWER8 core sets the relevant FIR bit and initiates the instruction retry and recovery (IRR) process to “clean up” all the architected states and flush the caches, ERATs, and TLB, but keep the SLB as is. Software restores the SLB. After the recovery process, a hypervisor maintenance interrupt (HMI) is generated. On a successful recovery, the HMER indicates a successful recovery.

If the same parity error occurs several times and reaches a threshold, the hypervisor can decide that the core is nonfunctional. The threshold counter is maintained by the hypervisor in software.

HID0[13] must be set to ‘0’, otherwise processor recovery will not work.

Note: The instruction IRR process is engaged for detection of any recoverable parity error in the core or due to the firing of a control checker.

There is a separate FIR bit and FIR extension bits for a parity error in the I-cache, D-cache, SLB, D-ERAT, I-ERAT, TLB, and a few other structures. For all the other register files, there is one shared FIR bit to indicate parity error.

3.8.22.2 Multi-Hit

If there is a multi-hit in the D-ERAT, TLB, or SLB, the core finishes the operation with a machine check interrupt and sets the proper DSISR bit to indicate where the multi-hit was detected.

A multi-hit in the D-ERAT and SLB can occur due to a hardware failure. Multi-hit means more than one entry matched the EA in the D-ERAT (ESID in the case of an SLB). Due to their CAM structure, the result is a “bit-wise logical or” of the RA of the multiple entries (VSID in case of SLB). Because of this “bit-wise logical or”, multi-hit is very likely to generate a parity error as well.

Because the SLB is managed by software with the Power ISA, a software bug can result in a multi-hit in SLB structures. There is no known case of multi-hit in I-ERAT that can produce a wrong result.

There are separate FIR bits for a multi-hit in the D-ERAT, TLB and SLB.

3.8.22.3 Both Multi-Hit and Parity Error

If both multi-hit and parity errors happen in the D-ERAT or TLB, the processor core initiates an instruction retry and recovery (IRR) process. No machine check is presented, but, as usual, after the recovery operation the processor core provides an HMI interrupt.

For SLB, any error causes the processor to take a machine check interrupt. The FIR bit setting indicates both multi-hit and parity error.



3.8.22.4 Uncorrectable Error Handling

If there is an uncorrectable error (UE) for a translate or a load operation, the instruction will finish with a machine check indication to the ISU. The instruction is flushed and re-executed without generating any machine check, and a counter is maintained to see how many UEs occurred. If the UE occurs more than a threshold, a MC interrupt is taken. For caching-inhibited load operation, a MC interrupt is taken on the first occurrence of the UE.

For the instruction side (I-side), if an instruction is executed and in the non-speculative path, only then is it treated as a UE. Otherwise, the I-side UE handling mechanism is similar to the D-side.

The core provides the EA of the LSU operation that caused the UE in the DAR register. For a UE detected by the IFU for instruction fetches, SRR0 is set to the EA.

Table 3-19 summarizes how the POWER8 processor handles parity, multi-hit, and uncorrectable errors.

Table 3-19. Summary of POWER8 Behavior on Parity Error, Multi-Hit, and Uncorrectable Error

	Parity Error	Multi-Hit	Both Parity Error and Multi-Hit	Uncorrectable Error (UE)
SLB: I-side translation	MC, SRR1, SRR0	MC, SRR1, SRR0	MC, SRR1, SRR0	N/A
SLB: D-side translation, SLBFEE, MFSLB	MC, DSISR, DAR	MC, DSISR, DAR	MC, DSISR, DAR	N/A
TLB: I-side translation	IRR, HMI	MC, SRR1, SRR0	IRR, HMI	N/A
TLB: D-side translation, MFTLB	IRR, HMI	MC, DSISR, DAR	IRR, HMI	N/A
D-ERAT	IRR, HMI	MC, DSISR, DAR	IRR, HMI	N/A
Tablewalk: I-side initiated	IRR, HMI	N/A	N/A	MC, SRR1, SRR0
Tablewalk: D-side initiated	IRR, HMI	N/A	N/A	MC, DSISR, DAR
Load	IRR, HMI	N/A	N/A	MC, DSISR
CI Load	MC, DSISR	N/A	N/A	MC, DSISR
Store	IRR, HMI	N/A	N/A	MC, DSISR
Instruction fetch	IRR, HMI	N/A	N/A	MC, SRR1, SRR0
Any other structure (I-ERAT, other Regfile, I-cache, D-cache and other SRAMs, data-flow hardware control checker)	IRR, HMI	N/A	N/A	N/A

1. MC is a machine check interrupt, IRR is an instruction retry and recovery, HMI is a hypervisor maintenance interrupt, and SRR0, SRR1, DSISR, DAR are various SPRs set on a machine check interrupt. In the TLB, a multi-hit cannot generate a parity error, but a parity error can generate a multi-hit. In the SLB and D-ERAT, multi-hit probably generates a parity error.

3.8.22.5 TLB Parity Error and Multi-Hit Action

Parity = 0 and Multi-hit = 0: No action.

Parity = 1 and Multi-hit = 0: Parity error detected, IRR, followed by HMI (no machine check).

Parity = 0 and Multi-hit = 1: This case is probably caused by software setting up two TLB entries pointing to the same VSID.

Parity = 1 and Multi-hit = 1: Probably multiple bits flipped due to a soft-error that caused the parity error, but also made two VSIDs look the same. The POWER8 core does IRR and then HMI.

3.8.23 Interrupts

3.8.23.1 Interrupt Vectors

Exceptions implemented in the POWER8 core are listed in *Table 3-20*.

Table 3-20. Interrupt Vector

Exception Type	Exception Value
System Reset	0X00100
Machine Check	0X00200
Data Storage Interrupt	0X00300
Data Segment Interrupt	0X00380
Instruction Storage Interrupt	0X00400
Instruction Segment Interrupt	0X00480
External Interrupt	0X00500
Alignment Interrupt	0X00600
Program Interrupt	0X00700
Floating-Point Unavailable	0X00800
Decrementer Interrupt	0X00900
Hypervisor Decrementer Interrupt	0X00980
Directed Privileged Doorbell Interrupt	0X00A00
Reserved	0X00B00
System Call	0X00C00
Trace Interrupt	0X00D00
Hypervisor Data Storage Interrupt	0X00E00
Hypervisor Instruction Storage Interrupt	0X00E20
Hypervisor Emulation Assistance Interrupt	0X00E40
Hypervisor Maintenance Interrupt	0X00E60
Directed Hypervisor Doorbell Interrupt	0X00E80
Performance <u>SPMC</u> Interrupt	0x0EE0
Performance Interrupt	0X00F00
VMX Unavailable Interrupt	0X00F20
VSX Unavailable Interrupt	0X00F40
Facility Unavailable Interrupt	0X00F60
Hypervisor Facility Unavailable Interrupt	0X00F80
Soft Patch Interrupt	0X01500
Debug Interrupt	0X01600



3.8.23.2 Interrupt Definitions

Table 3-21 describes the interrupts that have been added to the POWER8 processor core.

Table 3-21. Interrupt Vectors for the POWER8 Core

Exception Type	Exception Value
Directed Privileged Doorbell Interrupt	x'00A00'
Directed Hypervisor Doorbell Interrupt	x'00E80'
Facility Unavailable Interrupt	x'00F60'
Hypervisor Facility Unavailable Interrupt	x'00F80'

In addition to the interrupt types defined in the Power ISA, the POWER8 core supports several implementation-specific interrupt types. These are summarized in *Table 3-22* and described in more detail in the subsequent sections

Table 3-22. Implementation-Specific Interrupt Types

Exception Type	Exception Value
Performance SPMC Interrupt	x'00EE0'

Implemented MSR and SRR1/HSRR1 Bits

Table 3-23. Implementation MSR and SRR1/HSRR1 Bits (Sheet 1 of 2)

Bits	MSR	SRR1/HSRR1
0	SF	SF
1	Reserved	Reserved
2	Not Implemented	Not Implemented
3	HV	HV
4	Not Implemented	Not Implemented
5	Reserved	Reserved
6:28	Not Implemented	Not Implemented
29:30	TS (Transactional State)	TS (Transactional State)
31	TM (Transactional Memory Available)	TM (Transactional Memory)
32	Not Implemented	Not Implemented
33	Not Implemented	Specific Interrupt Information
34	Not Implemented	Not Implemented
35:36	Not Implemented	Specific Interrupt Information
37	Not Implemented	Not Implemented
38	VMX	VMX
39	Not Implemented	Not Implemented
40	VSX	VSX
41	Not Implemented	Not Implemented
42:47	Not Implemented	Specific Interrupt Information



Table 3-23. Implementation MSR and SRR1/HSRR1 Bits (Sheet 2 of 2)

Bits	MSR	SRR1/HSRR1
48	EE	EE
49	PR	PR
50	FP	FP
51	ME	ME
52	FE0	FE0
53	SE	SE
54	BE	BE
55	FE1	FE1
56	US	US
57	Not Implemented	Not Implemented
58	IR	IR
59	DR	DR
60	Not Implemented	Not Implemented
61	PMM	PMM
62	RI	RI
63	LE	LE

3.8.23.3 System Reset Interrupt

The system reset interrupt is a non-maskable, asynchronous interrupt that is caused by an SCOM command for a soft reset.

Note: There is no explicit SRESET pin; SRESET must be invoked from the service processor.

Table 3-24. System Reset Interrupt

Register	Bits	Description
SRR0	0:63	Set to the effective address of the instruction that the processor would have.
SRR1	0:31	Implemented bits loaded from the MSR.
	32	Set to '0'.
	33	LPAR mode switch occurred while the thread was in power savings mode.
	35:36	Set to '0'.
	42:45	Interrupt caused by IFU detection of a hardware uncorrectable error (UE) 0000 Reserved by pervasive function. 0010 Interrupt caused by SCOM when not in power-saving mode or caused by back-to-back SRESET. 0011 Interrupt caused by hypervisor door bell. 0101 Interrupt caused by privileged door bell. 0100 Interrupt caused by SCOM when in power-saving mode. 0110 Interrupt caused by decrementer wake-up when in power-saving mode. 1000 Interrupt caused by external interrupt wake-up when in power-saving mode. 1010 Interrupt caused by HMI wake-up when in power saving mode. 1100 Interrupt caused by implementation-specific wake-up when in power-saving mode.
	46:47	Indicates if the interrupt occurs when the processor is in power-saving mode. 00 Interrupt did not occur while the processor was in power-saving mode. 01 Interrupt occurred while the processor was in power saving mode. The state of all resources was maintained as if the processor was not in power-saving mode 10 Interrupt occurred while the processor was in power-saving mode. The state of some resources was not maintained but the state of all hypervisor resources was maintained as if the processor was not in power-saving mode and the state of all other resources is such that the hypervisor can resume execution. 11 Interrupt occurred while the processor was in power-saving mode. The state of some resources was not maintained, and the state of some hypervisor resources was not maintained or the state of some resources is such that the hypervisor cannot resume execution.
	62	Loaded from MSR[62] if recoverable. Otherwise set to zero
	Others	Implemented bits loaded from MSR.

The POWER8 core implements a 1-deep queue to remember the reason of a subsequent system reset interrupt while a system reset interrupt is pending. The reason of the most important subsequent system reset interrupt is remembered per the following priority:

1. Hypervisor doorbell initiated system reset
2. Privileged doorbell initiated system reset
3. SCOM-initiated system reset
4. HMI-initiated system reset
5. External-initiated system reset
6. Decrementer-initiated system reset
7. Implementation-specific initiated system reset



3.8.23.4 Machine Check Interrupt

The are several possible causes of machine check interrupts in the POWER8 chip, some of which are generally recoverable and some of which are non-recoverable.

The following causes of machine check interrupts are precise and synchronous with the instruction that caused the operation which encountered the error (that is, SRR0 contains the address of the instruction that caused the operation).

1. The detection of either a parity error, or a multi-hit error, or both in the SLB during the execution of a load, store SLBFE, or MFSLB instruction. If the interrupt is caused by a soft error, executing the appropriate sequence of instructions in the machine check handler program clears the error condition without causing any loss of state, permitting the interrupted program to be resumed if MSR[RI] was a '1' when the instruction that encountered the error was executed.
2. If there is a multi-hit in D-ERAT or TLB, the core finishes the operation with a machine check interrupt and sets the proper DSISR bit to indicate where the multi-hit occurred.
3. If there is an uncorrectable ECC error when a load instruction is executed or when the page table is being searched in the process of translating an address, the instruction finishes with a machine check indication to the instruction sequencing unit. The instruction is flushed and re-executed without generating any machine check. A counter is maintained to see how many UEs occurred. If the UE occurs more than a pre-established threshold, a machine check interrupt is taken.
4. For a caching-inhibited load operation, the machine check interrupt is taken on the first occurrence of the UE.
5. For the I-side, if an instruction is executed and the instruction is in the nonspeculative path, only then will it be treated as a UE. Otherwise, the I-side UE handling mechanism is similar to the D-side.

In the cases described in items (2), (3), (4) and (5), no state is lost in the processor, but recovery of the correct data might not be possible.

For more traumatic errors or hard errors, these characteristics cannot be reliably provided on a machine check because it is likely that the failure will prevent reliable execution. Additionally, a machine check interrupt that occurs when MSR[ME] = '0' results in a checkstop.

In the POWER8 core, there is no asynchronous machine check interrupt. A machine check interrupt is taken when the machine check input pin is asserted, if enabled by HID0[32] = '1'. The FIR, debug logic, and hang recovery logic can also be programmed to induce machine check interrupts for various error conditions. In general, the POWER8 core works hard to make these interrupts recoverable, but there are some scenarios where it cannot achieve this. Software can use the MSR[RI] bit to help identify the cases where the machine check interrupt is recoverable.



Information about the suspected source of the error condition is logged into either the SRR1 Register, the DSISR Register, or both as defined in *Table 3-25* for synchronous machine checks.

Table 3-25. Synchronous Machine Checks

Register	Bits	Description
SRR0	0:63	Effective address of the next instruction that would have executed if the machine check interrupt was not taken. For cases where this is a recoverable machine check due to a load that has surfaced an error, this will be the address of the load instruction itself. (The POWER8 core allows the instruction to execute to surface the error, but inhibits the commitment of the results.) For cases where this is a recoverable machine check due to an instruction fetch surfacing an error, this will be the address of an instruction that initiated the memory/cache access.
SRR1	42	Interrupt caused by load/store detection of error (see DSISR).
	36, 43:45	Interrupt caused by instruction fetch, indicated by the following encoding: 0000 Reserved. 0001 Interrupt caused by a hardware uncorrectable error detected while doing an instruction fetch (but not translation related). 0010 Interrupt caused by an SLB parity error while translating an instruction fetch address. 0011 Interrupt caused by an SLB multiple hit, while translating an instruction fetch address. Note: This condition occurs if the ESIID fields of two or more SLB entries contain the same value. 0100 Interrupt caused by an I-ERAT multi-hit error. 0101 Interrupt caused by a TLB multiple-hit error detected while translating an instruction fetch address. Note: This condition occurs if an address is mapped to both a small and large page in the SLB. This condition can also occur due to a software bug, when a software-managed TLB mechanism is used. 0110 Interrupt caused by a hardware UE detected while doing a TLB reload for I-side. 0111 Reserved. 1000 Interrupt caused by an L2 abort on an instruction fetch due to foreign link time out. 1001 Interrupt caused by an L2 abort on an instruction tablewalk due to foreign link time out. 1011 Reserved. 11xx Reserved.
	62	Loaded from MSR[62] if recoverable. Otherwise, set to zero.
	others	Implemented bits loaded from MSR.

Table 3-25. Synchronous Machine Checks

Register	Bits	Description
DSISR	32:47	All zeros.
	48	Interrupt caused by a UE deferred error, but not for tablewalk (D-side only).
	49	Interrupt caused by a UE deferred error during a tablewalk (D-side).
	50	Nest abort.
	51	Nest abort for tablewalk.
	52	Interrupt caused by a D-ERAT multi-hit.
	53	Interrupt caused by a TLB multi-hit due to translation (D-side only) or MFTLB operation.
	54	Secondary D-ERAT multi-hit.
	55	Interrupt caused by a SLB parity error (translate lookup or mfslbfee) due to a translation (D-side only), SLBFEE, or MFSLB instruction.
	56	Interrupt caused by an SLB multi-hit (might not be recoverable) for translation (D-side only), SLBFEE, or MFSLB instruction.
	57	Zero.
DAR	58:63	All zeros.
	0:63	Effective address computed by a load or store instruction that caused the operation that encountered a parity error, or multi-hit, or both in the SLB, or which encountered a multi-hit in the TLB, or which encountered a multi-hit in D-ERAT, or which encountered an uncorrectable error (UE) while attempting to reload a TLB entry. For all other types of machine check interrupts, the DAR is undefined (including the case where the operand of the load instruction contains a UE). <ol style="list-style-type: none"> 1. SLB parity error, multi-hit, or both: DAR is loaded with the EA of the target of the load or store instruction that caused the error. 2. TLB multi-hit: DAR is loaded with the EA of the target of the load or store instruction that caused the error. 3. D-ERAT multi-hit: DAR is loaded with the EA of the target of the load or store instruction that caused the error. 4. UE on D-side table walk: DAR is loaded with the EA of the target of the load or store instruction. 5. UE on instruction fetch: DAR is undefined. 6. UE on I-side tablewalk: DAR is undefined. 7. UE on load or store instruction: DAR is undefined (EA is not available in LMQ for loads, so DAR cannot be loaded).

DSISR Implementation Note: All the bits have been implemented in hardware.

Machine Check Interrupt Handler Notes:

As mentioned previously, the machine check interrupt handler is expected to help hardware recover from certain types of D-ERAT, TLB, and SLB errors detected by the hardware. In general terms, the interrupt handler must check whether or not the machine check interrupt is recoverable (by looking at the state of the RI bit in SRR1). It must determine the type of error that caused the machine check (by looking at the state of the SRR1 and DSISR registers). It must flush the contents of the array that reported the detected error (this process is slightly different for each of the possible arrays). Finally, it must return to the interrupted process.



3.8.23.5 Hypervisor Maintenance Interrupt

The POWER8 hypervisor maintenance interrupt replaces the malfunction alert and thermal interrupt; and provides support for the recovery function. The Hypervisor Maintenance Exception Enable Register (HMEER) contains the sources of the interrupt, which can be masked by setting the HMEER enable bits to zero. The Hypervisor Maintenance Exception Register (HMER) setting indicates a successful recovery.

3.8.23.6 External Interrupt

An external interrupt is classified as being either a direct external interrupt or a mediated external interrupt. Both cause an interrupt to 0x500.

Direct External Interrupt

The direct external interrupt is signaled by the assertion of the external interrupt input signal. The external interrupt signal must remain asserted until the processor has actually taken the interrupt. Failure to meet this requirement can lead the processor to not recognize the interrupt request.

When LPES = '0', the following registers are set.

Table 3-26. Direct External Interrupt (LPES = '0')

Register	Bits	Description
HSRR0	0:63	Set to the effective address of the instruction that the thread would have attempted to execute next if no interrupt conditions were present.
HSRR1	33:36	Set to '0'.
	42:47	Set to '0'.
	Others	Loaded from the MSR.
MSR	–	See the <i>Power ISA</i> .

When LPES = '1', the following registers are set.

Table 3-27. Direct External Interrupt (LPES = '1')

Register	Bits	Description
SRR0	0:63	Set to the effective address of the instruction that the thread would have attempted to execute next if no interrupt conditions were present.
SRR1	33:36	Set to '0'.
	42:47	Set to '0'.
	Others	Loaded from the MSR.
MSR	–	See the <i>Power ISA</i> .

Mediated External Interrupt

Mediated external interrupts are caused by the LPCR[MER] = '1', when the thread is in privileged (supervisor) or problem state mode.

When LPES = '0', the following registers are set.

Table 3-28. Mediated External Interrupt (LPES = '0')

Register	Bits	Description
HSRR0	0:63	Set to the effective address of the instruction that the thread would have attempted to execute next if no interrupt conditions were present.
HSRR1	33:36	Set to '0'.
	42	Set to '1'.
	43:47	Set to '0'.
	Others	Loaded from the MSR.
MSR	–	See the <i>Power ISA</i> .

When LPES = '1', the following registers are set.

Table 3-29. Mediated External Interrupt (LPES = '1')

Register	Bits	Description
SRR0	0:63	Set to the effective address of the instruction that the thread would have attempted to execute next if no interrupt conditions were present.
SRR1	33:36	Set to '0'.
	42:47	Set to '0'.
	Others	Loaded from the MSR.
MSR	–	See the <i>Power ISA</i> .

3.8.23.7 Alignment Interrupt

See *Section 3.1.4.2* for details on when the POWER8 core takes alignment interrupts.

Table 3-30 shows how the DSISR and DAR are set for alignment interrupts.

Table 3-30. Alignment Interrupt

Register	Bits	Description
DSISR	32:63	Unchanged.
DAR	0:63	See <i>Table 3-2</i> on page 46 and <i>Table 3-3</i> on page 54 on how the DAR is set on alignment interrupts.



3.8.23.8 Trace Interrupt

The trace interrupt is taken when the single-step trace-enable bit (MSR[SE]) or the branch trace enable bit (MSR[BE]) is set and an instruction successfully completes. After a trace interrupt is taken, SRR0, SRR1, SIAR, and SDAR are set as shown in *Table 3-31*.

Table 3-31. Trace Interrupt

Register	Bits	Description
SRR0	0:63	Set as specified in the architecture.
SRR1	0:32	Implemented bits loaded from the MSR.
	33:34	'10'
	35	Set for a load instruction; otherwise, cleared.
	36	Set for a store instruction; otherwise, cleared.
	37:41	Loaded from the MSR.
	42	Set for a Ibarx/Iharx/Iwarx/Idarx/Iqarx or stbcx./sthcx./stwcx./stdcx./stqcx. instruction; otherwise, cleared.
	43	Set to a '1' if a CIABR trace.
	44:47	Set to '0'.
	48:63	Implemented bits loaded from the MSR.
Note: Bit 35 and 36 are not set if an X-form Load String or Store String instruction specifies an operand length of 0.		
SIAR	0:63	Set to the effective address of the traced instruction; undefined if a CIABR trace.
SDAR	0:63	If the instruction that took the trace interrupt was a storage access instruction, the SDAR is set to the effective address of the storage access. SDAR is not set if an X-form Load String or Store String instruction specifies an operand length of 0; undefined if a CIABR trace.

The contents of SIAR and SDAR are undefined until a trace interrupt occurs.

3.8.23.9 Performance Monitor Interrupt

The performance monitor interrupt is signaled when the MSR[EE] bit is set, the MMCR0[PMAE] bit is set, and any of the performance monitor counters overflow (this includes the eight performance counters defined in the SPR space, as well as the counters defined in MMIO space for the nest).

After such an event is detected, the POWER8 core waits for previously dispatched instructions to complete, and then takes the interrupt.

3.8.23.10 SPMC Performance Monitor Interrupt

The SPMC performance monitor interrupt is an implementation-specific interrupt introduced on the POWER8 core. The SPMC performance monitor interrupt is signaled when an SPMC overflow exception occurs.

The MSR, SRR0, SRR1 interrupt settings are the same as the architected performance monitor interrupt as described by the POWER ISA. For the SPMC performance monitor interrupt execution resumes at the effective address x'0000_0000_0000_OEE0'.

3.8.23.11 Facility Unavailable Interrupt

The POWER8 core implements the facility unavailable interrupt as defined in the Power ISA.

3.8.24 Logical Partitioning Support

The following sections describe the POWER8 implementation aspects of the LPAR architecture. Each POWER8 core can support one, two, and four LPARs referred to as 1 LPAR, 2 LPAR and 4 LPAR mode respectively.

3.8.24.1 Thread-to-LPAR Mapping

Logical threads are tied to partitions as follows:

- 2 LPAR mode
 - LPAR0 - threads 0, 1, 2, and 3
 - LPAR1 - threads 4, 5, 6, and 7
- 4 LPAR mode
 - LPAR0 - threads 0 and 1
 - LPAR1 - threads 2 and 3
 - LPAR2 - threads 4 and 5
 - LPAR3 - threads 6 and 7

The snooped TLBIE in 4 LPAR mode only invalidates the D-ERAT and I-ERAT entries for the partitions that matched. If only one LPAR ID matches, the thread pair for that LPAR is used to invalidate the D-ERAT and I-ERAT. If two LPAR IDs match, two thread pairs for that LPAR are used for ERAT invalidation. If more than two LPAR IDs match, no thread bits are used for ERAT invalidation.

3.8.24.2 Dynamic LPAR Switching

In 2 LPAR or 4 LPAR mode, the POWER8 core runs in SMT8 mode and each LPAR is entitled to its share of execution bandwidth, irrespective of how many threads are actually active on the core. HID0 Register bits 11, 12, and 15 are used to control dynamic LPAR switching.

- HID0[DYN_LPAR_DIS]. If not set and the core is in either 2 LPAR or 4 LPAR mode, hardware switches the core to 1 LPAR mode when threads 1 - 7 are all in *nap* mode.
- HID0[1LPARto2LPAR]. A write to the bit initiates a dynamic switch from 1 LPAR mode to 2 LPAR mode.
- HID0[1LPARto4LPAR]. A write to the bit initiates a dynamic switch from 1 LPAR mode to 4 LPAR mode

3.8.25 Strong Access Ordering Mode (SAO)

The POWER8 core supports the SAO mode defined in Power ISA.

3.8.26 Graphics Data Stream Support

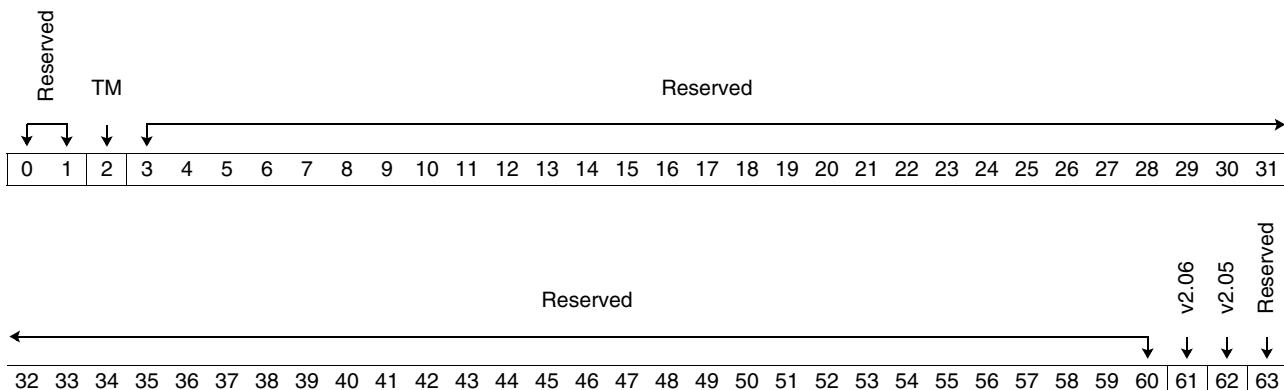
For cache-inhibited stores, the POWER8 core provides store gathering with an intentional stall to maximize the amount of gathering that can occur.

3.8.27 Performance Monitoring, Sampling, and Trace

Performance monitoring facilities have been incorporated into the POWER8 processor to enable the collection of performance related data and instruction traces. In general, the POWER8 core supports the recommended architecture for performance monitoring as described in the Power ISA.

3.8.28 Processor Compatibility Mode

The POWER8 core implements the Processor Compatibility Register (PCR) as described in the Power ISA to facilitate partition migration. Setting PCR[2] = '1' disables the transactional memory (TM) instructions in problem state. Setting PCR[61] = '1' disables all problem state instructions and facilities that were added in Version 2.07 of the Power ISA. Thus, setting this bit effectively makes a POWER8 core architecturally appear to problem state software as a Power ISA version 2.06 core. Likewise, setting PCR[62] = '1' disables all problem state instructions and facilities that were added in Version 2.06 of the Power ISA. Thus, to migrate a partition from a version 2.05 system to a POWER8 (version 2.07) system, PCR[61:62] must be set to '11'.



The POWER8 core does not follow the Power ISA for controlling the problem state code's ability to set thread priority to '001' based on the PCR[v2.06] bit. Instead, the POWER8 core controls this only by using the HID1[11:14] partition bits. See *Section 3.7.2 HID1 Register* on page 95.

There are bits in the I-cache that are generated based on the settings of bits in the PCR. Therefore, if the state of the PCR is changed due to a **mtspr(PCR)**, the I-cache should be flushed using HID1[0].

4. Storage Subsystem

4.1 L1 Cache

The L1 I-cache is 8-way set associative and is indexed with five effective address bits (EA[51:55]). A particular physical block of memory with a given real address can be found in one of two positions in the L1 I-cache. The tag comparison associated with lookups in this cache (as well as all other caches in the system) are done using physical addresses, so that there are no synonym or alias hazards that must be explicitly handled by the system software.

The L1 D-cache is 8-way set associative and is indexed with six effective address bits (EA[51:56]). A particular physical block of memory with a given real address can only be found at a particular location in the L1 D-cache. On each access, the tag comparison is done with the physical address. On a cache miss, the cache reload mechanism searches the other seven related sets to determine if the required real address block is located elsewhere in the cache, and if so, it appropriately eliminates these copies.

4.2 L2 Cache

The L2 cache is a unified cache that is accessed privately by a given core on the POWER8 processor. The L2 cache maintains full hardware coherency within the system and can supply cache intervention data to other cores on this die or on other POWER8 chips (for example, both on-chip and off-chip intervention). The L2 cache is a store-in cache that is fully inclusive of both the D-cache and I-cache for its private core (that is, the core has a store-through L1 D-cache). The L2 cache also supports private bus access to an 8 MB L3 cache region that is also private to this core.

4.2.1 L2 Cache Features

A summary of the L2 cache follows:

- 512 KB private cache per core
 - 128-byte line, 8-way set associative
 - Both I-side and D-side inclusive (ICBI is not required on the SMP interconnect)
 - Double-banked cache design interleaved on even or odd cache-line boundary
 - Can perform a read from one bank while writing to the other bank
 - Cache-array data protected by 8-byte SECDED ECC
- 8-way directory, dual-banked multiport
 - One processor read port, one snoop read port, and one write port per bank
 - The processor port into a given bank operates at $\frac{1}{2}$ the processor clock rate (initiated on 2:1 clock boundary)
 - The snoop port into a given bank operates at $\frac{1}{2}$ the processor clock rate (initiated on 2:1 clock boundary) allowing for two snoops per 2:1 clock across the two banks.
 - A bank can initiate up to two directory reads in a given 2:1 cycle (one on the snoop port, one on the processor port). A bank can initiate one write in a given pclk cycle (no reads allowed during a write pclk cycle)



- Directory array data protected by SECDED ECC
- 512 × 13-bit LRU arrays (logical configuration); 2 × 4 LRU vector tracking tree with cache-invalidate state biasing
- Seven processor cycles latency beyond L1 data cache on an L2 hit through the fastpath
- Eight processor cycles latency beyond L1 instruction cache on an L2 hit
- Point of global coherency
- Reservation stations, one per processor thread
- Support for transactional memory (TM) transactions
- Dual snoop bus port split by even and odd cache-line address
- Support for micropartition prefetch assist logging
- Support strong access ordering (SAO)
- Hardware line delete for error tolerance

4.3 L3 Cache

The L3 cache controller services read and cast-out requests from the attached L2 cache and snoop requests from the fabric. It also provides a mechanism to prefetch cache lines into the L3 cache based on requests from the core. Data movement for L2 cast-outs (cast-ins for the L3 cache) are carried out using a private interface between the L2 and L3 cache. Data movement for L3 cast-outs, L3 interventions, and memory pushes are carried out by sending commands and data to the fabric interface. The L3 cache supports victim mode for others on chip L3 cache (referred to as L3.1 caches).

4.3.1 L3 Features, Queues and Resources

- Local L3 cache region and shared L3.1
 - 128-byte cache line, 8-way set associative.
 - 8 cache banks (interleaved for access overlapping)
 - 64-byte reload and CI data buses.
 - 128-byte read or write per bank every 12 pclks.
 - macros configured in eight 8-byte dataflows.
 - 8-way directory: 4 banks, up to 4 reads or 2 reads and 2 writes every 2 pclks to differing banks.
Physically 16 SRAMs.
 - 26.5 pclk L1 data-cache load to use penalty.
 - Cache contents protected by 8-byte ECC.
 - Directory contents protected by ECC.
 - LRU algorithm with enhancements (for efficient L3.0/L3.1 victimization).
- Functionality
 - L3.0 management of victim lines from local L2 cache.
 - L3.1 management of victim lines from on-chip L3 caches.
 - Services load/store/I3 prefetch misses from local L2 cache.



- Dual Snoop ports split by even and odd cache-line address.
- Streaming data prefetch and cache-inject support.
- Dual-class L3.0/L3.1 LRU support and L3.1 activity throttling.
- Support for speculative memory transactions, including footprint tracking to detect collisions and mechanisms for conditional completion. Consists of fast flash 64 CL per thread and slow walk entire 8 MB available for one thread.
- Fast broadcast on-chip fetch request to memory controllers and other L3 caches for L2 demand load miss L3.
- Entire L3 cache clocked at ½ core frequency.



4.4 NCU

The POWER8 noncacheable unit (NCU) is responsible for processing noncacheable operations (such as, load and store operations with I = '1') and certain other uncacheable operations such as TLBIE, various sync instructions, ptesync instructions, and so on. All of these instructions support the behavioral definitions given in the Power ISA documents. One NCU unit is instantiated per core and this NCU handles operations for all eight threads in the associated POWER8 core.

The POWER8 NCU provides one dedicated cache-inhibited load station per thread to process one outstanding cache-inhibited load per thread. In the POWER8 implementation, cache-inhibited loads (whether guarded or not) are not gathered and are not reordered in any fashion.

For stores, sixteen 64-byte store gather stations are provided and are shared across the eight core threads. A pair of 64-byte stations can “chain” together to gather up to full 128-byte lines. The POWER8 NCU supports gathering and reordering for stores in the IG = '10' space. In the IG = '11' space, stores are neither reordered nor gathered as required by the architecture. The POWER8 NCU only gathers 4-, 8-, and 16-byte stores. The stores must be naturally aligned, start at a given address, not overlap and be contiguous. This is designed to support the general gathering case of a set of stores of a given size starting at the beginning of a 128-byte block and continuing on to fill the entire 128-byte block.

4.4.1 NCU Characteristics

The NCU has the following characteristics:

- Store buffer
 - 16 × 64-byte store gather stations (chainable to 8 × 128-byte gathered lines).
 - The store buffers are shared across threads (LSU backoff mechanism prevents any thread from blocking any other thread).
- Store modes (IG = '1X')
 - IG = '11' mode; stores are done in-order and no gathering is allowed.
 - IG = '10' mode; stores can be gathered and reordered.
- Loads. One outstanding load per thread.

4.5 Memory Controller

The POWER8 memory controller function is split between the POWER8 memory controller (MC) unit, which provides an interface to the processor bus, and the POWER8 Memory Buffer chip, which provides memory command scheduling, a memory buffer cache, and memory diagnostic functions. The POWER8 MC unit is connected to the POWER8 Memory Buffer chip through the differential memory interface (DMI) channel interface.

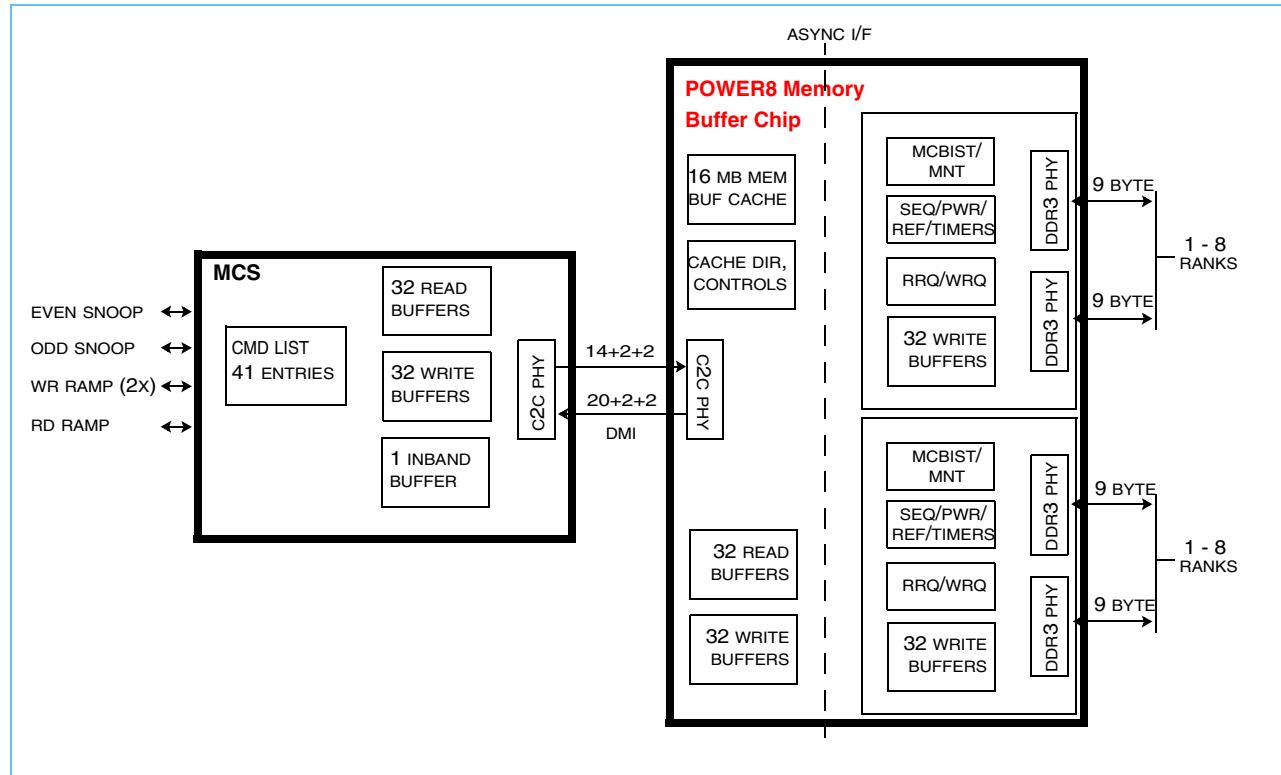
The POWER8 memory controller design has two memory controller units and 4 DMI ports per chip.

4.6 POWER8 Memory Stack Partitioning

Figure 4-1 shows a memory buffer chip connected to a memory controller synchronous (MCS) unit.

The POWER8 Memory Buffer chip contains four DDR3/4 memory ports that interface to Industry-Standard RDIMMs or LRDIMMs. Two of the four memory ports on the memory buffer operate in tandem to read or write a cache line of data using a burst length 8 (BL8) operation to the installed DIMMs.

Figure 4-1. Memory Stack Partitioning





4.7 POWER8 Chip Memory Controller Unit Features

- Physical Organization
 - POWER8 processor contains two memory controller units per chip.
 - Each memory controller unit contains two independent MCS units.
 - MCS units are accessed in pairs for selective memory mirroring.
 - Each MCS unit connects to a single DMI channel.
- Processor Bus Interface (per MCS unit)
 - Single 16-byte read data ramp, dual 16-byte write data ramp interfaces
 - 128-byte cache lines supported (cache line interleaving on a 32-byte basis)
 - Automatic maintenance of domain bits for multi-node systems
 - Speculative dispatch on a bandwidth-available basis
 - Delivery of critical octword of read data not gated by reading the entire cache line
- Memory Channel Interface (per MCS unit)
 - Downstream command/write data interface, up to 1 TB of memory addressing
 - Upstream tagged read data/status interface
 - CRC-protected
 - Error status and logging interface
 - Automatic hardware replay of soft channel faults
 - Automatic reconfiguration to use spare channel lanes to recover from hard failures
 - Hardware and software channel initialization capability
- Pervasive Interfaces (per MCS unit)
 - Performance monitor interface
 - SCOM interface
 - Debug bus interface (to shared trace array in processor bus logic)
- Resources (per MCS Unit)
 - Forty-one command list entries (32 memory entries, eight address-only entries, one inband configuration register entry)
 - Thirty-two 128-byte read data buffers
 - Thirty-two 128-byte write data buffers
 - One 128-byte configuration register read/write buffer
- Mirroring
 - Selective memory mirroring supported across MCS pairs. The pairs of MCS units must be within the same MC unit.
 - Mirrored memory accessed through unique memory address range (independent of memory extent).
- Clocking
 - 1:1 with processor bus
 - 1:4 with DMI channel

The memory buffer supports the following fixed clocking ratios and the maximum and minimum frequencies shown in *Table 4-1* and *Table 4-2* on page 133:

- 1:1 with processor bus
- 1:4 with DMI channel



Table 4-1. Maximum Frequencies

DRAM Speed (GHz)	POWER8 Memory Buffer DDR Frequency (GHz)	Processor Bus/MCS Frequency (GHz)	POWER8 Memory Buffer Cache Frequency (GHz)	DMI Channel Frequency (GHz)
1.33	1.33	2.4	1.2	9.6
1.6	1.6	2.4	1.2	9.6

Table 4-2. Minimum Frequencies

DRAM Speed (GHz)	POWER8 Memory Buffer DDR Frequency (GHz)	Processor Bus/MCS Frequency (GHz)	POWER8 Memory Buffer Cache Frequency (GHz)	DMI Channel Frequency (GHz)
1.33	1.33	2.0	1.0	8.0
1.6	1.6	2.0	1.0	8.0

4.7.1 Bandwidth

Table 4-3 lists the bandwidth per MCS unit per POWER8 Memory Buffer.

Table 4-3. Bandwidth (per MCS/POWER8 Memory Buffer)

Performance Parameter	Goal
Maximum downstream (write) DMI channel data bandwidth, 9.6 GHz channel	9.6 GB/s
Maximum upstream (read) DMI channel data bandwidth, 9.6 GHz channel	19.2 GB/s
Note: See the <i>POWER8 Memory Buffer User's Manual</i> for additional performance parameter goals.	

4.7.2 POWER8 Memory Controller Characteristics

Table 4-4 lists key characteristics of the POWER8 memory controller.

Table 4-4. POWER8 Memory Controller Characteristics (Sheet 1 of 2)

Parameter	POWER8
Processor Bus Interface	
Reflected command/partial response/combined response interfaces	Two per MCS unit
Read data ramp	One 16-byte read data ramp per MCS unit
Write data ramp	Two 16-byte write data ramp per MCS unit
Read buffers	Thirty-two 128-byte buffers per MCS unit
Write buffers	Thirty-two 128-byte buffers per MCS unit
Processor bus transfer size	32 bytes
Cache-line size	128 bytes
Command list entries	Thirty-two read/write, eight address-only, one configuration register per MCS unit
Address hashing	1, 2, or 4 MCS/group, 128-byte hashing



Table 4-4. POWER8 Memory Controller Characteristics (Sheet 2 of 2)

Parameter	POWER8
Memory Interface	
Memory channels	One per MCS unit
Channel speed/signaling	9.6 Gb/s maximum differential
Channel command/write data bus format	Combined command/data buses
Channel error protection	CRC, with retry
Read delay latency calculation	None, returned read data is tagged
Channel initialization/calibration	Dynamic
Channel lane sparing	Firmware (CRC syndrome capture), hardware
Memory fetch/store sizes	128 bytes
Channel speed ratios	4:1 (to SMP interconnect frequency)

5. Simultaneous Multithreading

5.1 Overview

The POWER8 processor core supports ST, SMT2, SMT4, and SMT8 modes. Any thread number can be run in any SMT mode, on any thread set. *Table 5-1* shows the SMT mode definitions.

Table 5-1. SMT Modes

Description	Number of Threads Enabled	Switch to this SMT mode when ...
SMT0	0	POR state
SMT1	1	1 thread
SMT2	1 - 2	2 threads
SMT4	1 - 4	3 - 4 threads
SMT8	1 - 8	5 - 8 threads

5.2 Partitioning of Resources in Different SMT Modes

Table 5-2 lists the resources that are partitioned in certain SMT modes.

Table 5-2. Front-End Execution Core Resource (Sheet 1 of 2)

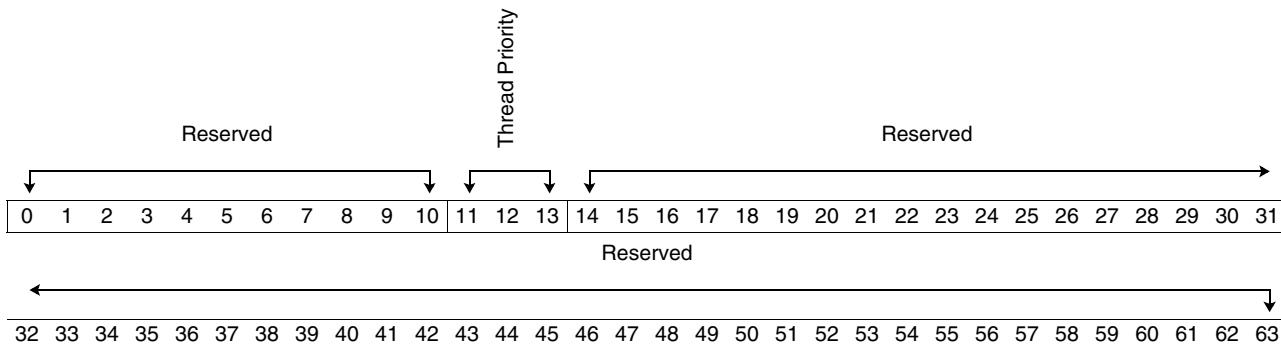
Resource	ST	SMT2	SMT4	SMT8
Fetch Bandwidth				
<u>EAT</u> Entries	24	24 per thread	12 per thread	6 per thread
Instruction Buffer Entries	64	64 per thread	32 per thread	16 per thread
Link Stack	32	32 per thread	16 per thread	8 per thread
D-ERAT Entries	48	48 per thread	48 per thread set	48 per thread set
Decode Bandwidth				
Dispatch Groups	6 nonbranch + 2 branch per thread	3 nonbranch + 1 branch per thread	3 nonbranch + 1 branch per thread	3 nonbranch + 1 branch per thread
<u>GPR</u> Entries	124	124 per thread	124 per thread set	124 per thread set
GPR in-flight renames	92	92 per thread	60 per thread set	60 per thread set
GPR Architected in First Level	32	32 per thread	64 per thread set	64 per thread set
GPR Architected in SAR				64 per thread set in SAR
<u>VRF</u> Entries	144	144 per thread	144 per thread set	144 per thread set
VRF in-flight renames	80	80 per thread	80 per thread set	80 per thread set
FP/VMX Architected in First Level	64	64 per thread	64 per thread set	64 per thread set
FP/VMX Architected in SAR			64 per thread set in SAR	192 per thread set in SAR
Steering	Toggle	Assign by thread set	Assign by thread set	Assign by thread set

Table 5-2. Front-End Execution Core Resource (Sheet 2 of 2)

Resource	ST	SMT2	SMT4	SMT8
Unified Issue Queue Entries	64	32 per thread	32 per thread set	32 per thread set
FXU, LSU, LU, VSU	2 each	1 each per thread	1 per thread set	1 per thread set
Completion Rates	1 group per cycle	1 (3+1) group per thread	1 (3+1) group per thread set	1 (3+1) group per thread set

5.3 Program Priority Register (PPR)

Each thread has a 64-bit status register associated with it, and bits 11:13 contain the priority for that thread. The effect of the priority set in the PPR is described in *Section 5.14 Thread Prioritization Implementation* on page 146. Some bits are read-only, while other bits are read/write. There is one PPR per thread.



Bits	Field Name	Description
0:10	Reserved	Reserved (not implemented).
11:13	Thread Priority	Thread Priority. 000 Not allowed 001 Very low 010 Low 011 Medium low 100 Normal 101 Medium high 110 High 111 Extra high Set to '100' on system reset interrupt.
14:63	Reserved	Reserved (not implemented).

The local PPR Register can be accessed with the **mtspr** or **mfspqr** instructions using SPR 896.

The PPR Register for each thread is initialized to x'0010_0000_0000_0000' at power-on.

5.4 Thread Priority NOPs

The thread switch priority can be read or written by software using the **mfsp** and **mtspr** instruction to the Thread Status Register. Thread priority can also be altered by executing special forms of the **or x,x,x** NOP. The priority is changed upon completion of the operation, provided the function is enabled for the current privilege level.

- On the POWER8 core, problem-state programs can set their priority to very-low. TSCR[16] is reserved on the POWER8 core.
- Supervisor programs can set their thread priority to six different values, very-low through high. TSCR[15] is reserved on the POWER8 core.
- Hypervisor code can set all levels.

Table 5-3 shows how to set the thread priority NOPs.

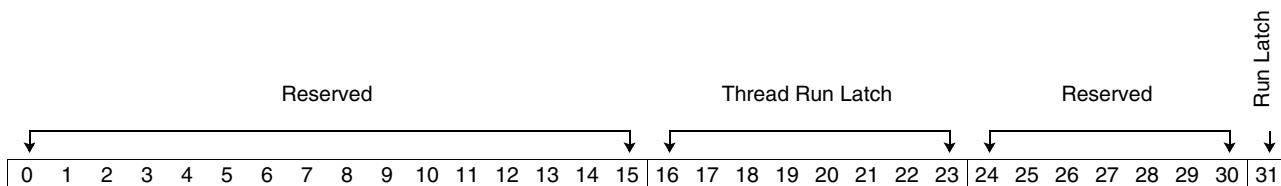
Table 5-3. Thread Priority Nops

Priority NOP / mtSPR	PPR[11:13]	Thread Priority	Required Privilege Level to Set Given Thread Priority Value
or 31,31,31 / mtPPR[11:13]	'001'	Very Low	Hypervisor, Supervisor, Problem
or 1,1,1 / mtPPR[11:13]	'010'	Low	Hypervisor, Supervisor, Problem
or 6,6,6 / mtPPR[11:13]	'011'	Medium Low	Hypervisor, Supervisor, Problem
or 2,2,2 / mtPPR[11:13]	'100'	Medium (Normal)	Hypervisor, Supervisor, Problem
or 5,5,5 / mtPPR[11:13]	'101'	Medium High	Hypervisor, Supervisor, Problem ¹
or 3,3,3 / mtPPR[11:13]	'110'	High	Hypervisor, Supervisor
or 7,7,7 / mtPPR[11:13]	'111'	Extra High	Hypervisor

1. See *Section 5.12 Priority Boosting to Medium-High in User Mode* on page 144.

5.5 Control Register

The Control Register (CTRL) is an architected 32-bit register. The bit assignment for the thread control bits in the CTRL supports up to eight threads. The POWER8 core supports threads 0 - 7. A bit in the CTRL Register represents the architected state for a particular thread.



Bits	Field Name	Description
0:15	Reserved	Reserved.
16:23	Thread Run Latch	Thread Run Latch bits (indirectly written supervisor and hypervisor). Threads 0 - 7.
24:30	Reserved	Reserved.
31	Run Latch	Run Latch for thread doing CTRL read/write (read only / rerouted supervisor or hypervisor write).

The CTRL Register can be read with the **mfsp** instruction using SPR 136 in user, supervisor, or hypervisor state.

The CTRL Register can be selectively written with the **mtsp** instruction using SPR 152 in the supervisor or hypervisor state.

The CTRL Register is initialized to 0x0000_0000 at power-on.

Even though a single CTRL Register is shared by the eight threads, there is no need to obtain a lock before updating the CTRL Register. There is only one bus that goes to the core pervasive unit, and the instruction issue logic serializes all **mtctrl** instructions. Updating the Run Latch bit must be done in hypervisor mode. When updating the Run Latch bit (in hypervisor mode), the software is recommended to set the Thread Enable bits to '00000000'. Setting the Thread Enable bit to '00000000' is not allowed. Therefore, this updates the run latch, but there is no effect to the Thread Enable bit and no thread will be killed or woken up.

CTRL[16:23] contain the Run Latches for threads 0 - 7. A **mtsp** CTRL instruction does not modify CTRL[16:23] based on GPR bits [48:55]. Instead, these bits are indirectly loaded by writing a value to CTRL[31]. The value written to CTRL[31] is loaded into CTRL[16] if thread 0 issued the move to CTRL and CTRL[17] if thread 1 issued the move to CTRL, and so on. A thread cannot update the thread Run Latch bit of another thread.

The run latch bit is only used by software for status and is sent to the performance monitor for performance analysis. For this purpose, the POWER8 processor core supports one run latch per thread. To use this function, if a thread is executing a dispatchable task, software must set the CTRL Run Latch for that thread to '1' by writing a '1' to CTRL[31]. If a thread is in a wait state, waiting for a dispatchable task, software must set the CTRL Run Latch for that thread to '0'.



Software can load the CTRL Register with a Run Latch value for its thread by writing the Run Latch value to CTRL[31]. Hardware routes data directed to CTRL[31] into either CTRL[16] - CTRL[23], depending on which thread is doing the write (see previous definition of CTRL[16:23]). When the CTRL Register is read, data driven on CTRL[31] comes from CTRL[16] - CTRL[23], depending on which thread is doing the read. CTRL[31] does not physically exist in hardware.

The data read on a **mfsp**(CTRL) is formatted differently based on the MSR[PR] and MSR[HV] bits and whether the core is in Big Core, 2 LPAR, or 4 LPAR mode. Bit 63 is always the Run Latch of the thread executing the **mfsp**. Bits 48:55 are formatted as shown in *Table 5-4*, where R0 equals run latch for thread 0, RT equals run latch of thread executing **mfsp**.

Table 5-4. mfsp CTRL Data Formatting

MSR[HV], MSR[PR]	Core Mode	Bits 48:55
00 - Privileged	Single LPAR	R0, R1, R2, R3, R4, R5, R6, R7
*1 - Problem	Single LPAR	RT, 0, 0, 0, 0, 0, 0, 0
10 - Hypervisor	Single LPAR	R0, R1, R2, R3, R4, R5, R6, R7
00 - Privileged	4 LPAR (Threads 0, 1)	R0, R1, 0, 0, 0, 0, 0, 0
00 - Privileged	4 LPAR (Threads 2, 3)	R2, R3, 0, 0, 0, 0, 0, 0
00 - Privileged	4 LPAR (Threads 4, 5)	R4, R5, 0, 0, 0, 0, 0, 0
00 - Privileged	4 LPAR (Threads 6, 7)	R6, R7, 0, 0, 0, 0, 0, 0
*1 - Problem	4 LPAR	RT, 0, 0, 0, 0, 0, 0, 0
10 - Hypervisor	4 LPAR	R0, R1, R2, R3, R4, R5, R6, R7
00 - Privileged	2 LPAR (Threads 0, 1, 2, 3)	R0, R1, R2, R3, 0, 0, 0, 0
00 - Privileged	2 LPAR (Threads 4, 5, 6, 7)	R4, R5, R6, R7, 0, 0, 0, 0
*1 - Problem	2 LPAR	RT, 0, 0, 0, 0, 0, 0, 0
10 - Hypervisor	2 LPAR	R0, R1, R2, R3, R4, R5, R6, R7



5.6 Thread Priority, Status, and Control Requirements

Thread priority, control, and status registers enable software to do the following:

- Give a large percentage of execution resources to critical tasks.
- Reduce the amount of resources and power used by low-priority work.
- Read foreground and background thread priority and status.
- Save and restore priority during interrupts.
- Provide a controlled way to allow supervisor/user code to change priority.
- Provide a means to kill or revive a thread.
- Avoid fine-grain livelock or deadlock situations between threads.

5.7 Thread Balance Control Requirements

The following mechanisms can be used to balance work between threads:

- Reduce ifetch priority of a thread using too many resources.
- Reduce decode priority of a thread using too many resources.
- Hold decode of thread with long latency events.
- Dispatch flush decode pipe to clean congested operations.
- Balance flush from next-to-complete plus one group and hold at IBUF until miss (and so on) resolves.

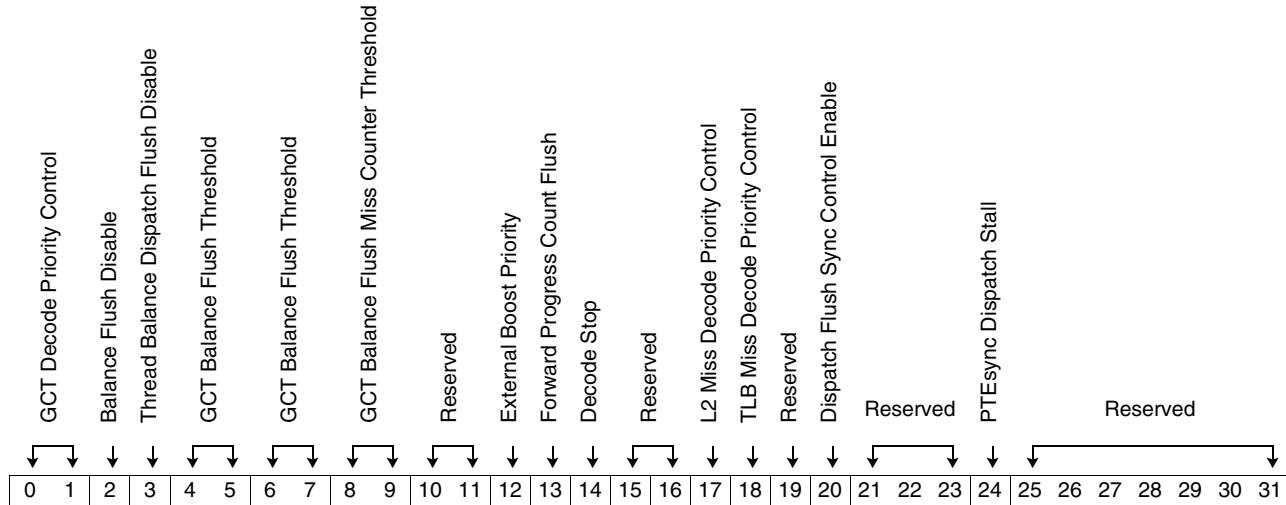
Table 5-5. Thread Balance Control (Balance Flush)

Indicator of Balance	Decode Priority	Decode Hold	Dispatch Flush	Balance Flush
GCT Utilization (covers UniQ utilization)	in plan	in plan	-	in plan ¹
UniQ Full	-	in plan ²	in plan ²	-
Sync operations ³	-	in plan ⁴	in plan	-
TLBIE	-	in plan ⁵	in plan	-
Score Board Full	-	-	in plan	-
TLB or L2 miss	in plan	-	-	-
TLB or L3 miss	-	In plan	in plan ⁶	in plan

1. Only used when a thread is stalled at dispatch and threads other than the stalled thread are consuming too many available resources.
2. If a particular thread or thread pair is using up its allocated UniQ resource, the thread pair is dispatched flushed and held at decode.
3. Synchronization instructions include: **sync**, **lwsync**, **ptesync**, **tlbsync**, **isync**, read/write of nonrenamed SPRs .
4. SYNC: After getting a dispatch flush, a sync instruction is held at IBUF(CLB) until its dispatch conditions are met. This is irrespective of thread priority.
5. TLBIE: After getting a dispatch flush, a **tlbie** instruction is held at the IBUF until the thread's GCT is empty. After the **tlbie** instruction is dispatched, the following instruction gets dispatched flushed and then held at IBUF until the "tlbie ack" is received from the GRS (through LSU). There is no effect on the **tlbief** instruction. This is irrespective of thread priority.
6. A dispatch flush is performed on the thread that is balanced flush only if that thread is stalled at dispatch and TSCR[3] is enabled.

5.8 Thread Switch Control Register (Hypervisor Access Only)

Thread priority controls are programmable. All bits are read/write. There is one Thread Switch Control Register per core. TSCR is initialized to x'0000_0000' at power-on.



Bits	Field Name	Description
0:1	GCT Decode Priority Control	GCT Decode Priority Control. If all threads are at the same software-set priority, decrease priority when a thread is using more than the following number of GCT entries: 00 Function disabled (TSCR POR default). 01 Scan-only latch value programmable 1 - 28. Scanlatch POR default is: 13 (SMT2); 7 (SMT4); 4 (SMT8). 10 Scan-only latch value programmable 1 - 28. Scanlatch POR default is 14 (SMT2); 8 (SMT4); 5 (SMT8). 11 Scan-only latch value programmable 1 - 28. Scanlatch POR default is 15 (SMT2); 9 (SMT4); 6 (SMT8).
2	Balance Flush Disable	Balance Flush Disable. 0 Enable NTCP1 balance flushes. 1 Disable NTCP1 balance flushes.
3	Thread Balance Dispatch Flush Disable	Thread Balance Dispatch Flush Disable. 0 Enable dispatch flush for the thread that was chosen for a balance flush if that thread is stalled at dispatch. A dispatch flush is a lower-latency flush than a balance flush. 1 Disable. Note: Conditions for dispatch flush are the same as a balance flush.
4:5	GCT Balance Flush Threshold for L3/TLB miss	GCT Balance Flush Threshold for L3/TLB miss. If a thread is stalled at dispatch with an L3 or TLB miss, consider the thread for balance flush when the thread is using more than the following number of GCT entries: 00 Function disabled (TSCR POR default). 01 Scan-only latch value programmable 1 - 28. Scan latch POR default is: 2 (SMT2); 2 (SMT4); 2 (SMT8). 10 Scan-only latch value programmable 1 - 28. Scan latch POR default is: 3 (SMT2); 3 (SMT4); 3 (SMT8). 11 Scan-only latch value programmable 1 - 28. Scan latch POR default is: 4 (SMT2); 4 (SMT4); 4 (SMT8).



Bits	Field Name	Description										
6:7	GCT Balance Flush Threshold if no L3/TLB miss	<p>GCT Balance Flush Threshold if no L3/TLB miss.</p> <p>In SMT2 and higher, if a thread is stalled at dispatch with no L3 or TLB misses, consider thread for balance flush when the thread is using more than the following number of GCT entries:</p> <table> <tr><td>00</td><td>Function disabled.</td></tr> <tr><td></td><td>Option available to have hardware disabled function in SMT2.</td></tr> <tr><td>01</td><td>Scan-only latch value programmable. 1 - 28. Scan latch POR default is: 22 (SMT2); 14 (SMT4); 7 (SMT8).</td></tr> <tr><td>10</td><td>Scan-only latch value programmable. 1 - 28. Scan latch POR default is: 24 (SMT2); 16 (SMT4); 8 (SMT8).</td></tr> <tr><td>11</td><td>Scan-only latch value programmable. 1 - 28. Scan latch POR default is: 26 (SMT2); 18 (SMT4); 9 (SMT8); 8 (4 LPAR).</td></tr> </table>	00	Function disabled.		Option available to have hardware disabled function in SMT2.	01	Scan-only latch value programmable. 1 - 28. Scan latch POR default is: 22 (SMT2); 14 (SMT4); 7 (SMT8).	10	Scan-only latch value programmable. 1 - 28. Scan latch POR default is: 24 (SMT2); 16 (SMT4); 8 (SMT8).	11	Scan-only latch value programmable. 1 - 28. Scan latch POR default is: 26 (SMT2); 18 (SMT4); 9 (SMT8); 8 (4 LPAR).
00	Function disabled.											
	Option available to have hardware disabled function in SMT2.											
01	Scan-only latch value programmable. 1 - 28. Scan latch POR default is: 22 (SMT2); 14 (SMT4); 7 (SMT8).											
10	Scan-only latch value programmable. 1 - 28. Scan latch POR default is: 24 (SMT2); 16 (SMT4); 8 (SMT8).											
11	Scan-only latch value programmable. 1 - 28. Scan latch POR default is: 26 (SMT2); 18 (SMT4); 9 (SMT8); 8 (4 LPAR).											
8:9	GCT Balance Flush Miss Counter Threshold	<p>GCT Balance Flush Miss Counter Threshold.</p> <p>If a thread is chosen to be balanced flushed based on TSCR[4:5] (GCT balance flush threshold), apply a CLB hold to that thread until the miss is resolved or until the following miss counter threshold is reached:</p> <table> <tr><td>00</td><td>Function disabled (If thread is balanced flushed due to TSCR[4:5], CLB hold is applied until a Miss is resolved).</td></tr> <tr><td>01</td><td>Scan-only latch value programmable, 10-bit LFSR; Scan latch POR default 256 cycles (x3C1 LFSR).</td></tr> <tr><td>10</td><td>Scan-only latch value programmable, 10-bit LFSR; Scan latch POR default 384 cycles (x0E4 LFSR).</td></tr> <tr><td>11</td><td>Scan-only latch value programmable, 10-bit LFSR; Scan latch POR default 512 cycles (x20F LFSR) 1023 cycles (x000 lfsr) is the maximum setting available.</td></tr> </table>	00	Function disabled (If thread is balanced flushed due to TSCR[4:5], CLB hold is applied until a Miss is resolved).	01	Scan-only latch value programmable, 10-bit LFSR; Scan latch POR default 256 cycles (x3C1 LFSR).	10	Scan-only latch value programmable, 10-bit LFSR; Scan latch POR default 384 cycles (x0E4 LFSR).	11	Scan-only latch value programmable, 10-bit LFSR; Scan latch POR default 512 cycles (x20F LFSR) 1023 cycles (x000 lfsr) is the maximum setting available.		
00	Function disabled (If thread is balanced flushed due to TSCR[4:5], CLB hold is applied until a Miss is resolved).											
01	Scan-only latch value programmable, 10-bit LFSR; Scan latch POR default 256 cycles (x3C1 LFSR).											
10	Scan-only latch value programmable, 10-bit LFSR; Scan latch POR default 384 cycles (x0E4 LFSR).											
11	Scan-only latch value programmable, 10-bit LFSR; Scan latch POR default 512 cycles (x20F LFSR) 1023 cycles (x000 lfsr) is the maximum setting available.											
10:11	Reserved	Reserved.										
12	External Boost Priority	<p>External Boost Priority.</p> <p>If '1' and an external interrupt request is active and the corresponding threads' priority is less than normal priority, set the threads' priority to normal.</p> <p>Note: This will <i>not</i> change the value in PPR[11:13] for the affected thread.</p>										
13	Enable Forward Progress Count Flush	<p>Enable Forward Progress Count Flush.</p> <p>Note: This bit only enables/disables the flush from occurring.</p> <p>The forward progress timer does not stop decrementing when set to '0'</p> <p>SMT2 and higher: If one thread is not making progress, enable flushing the other active threads.</p>										
14	Decode Stop	<p>Decode Stop.</p> <p>When set to 0, the forward progress timer (in PPR) is decremented even when the current thread is in decode stop state. When set to 1, the forward progress timer is not decremented when the current thread is in decode stop state.</p>										
15:16	Reserved	Reserved.										
17	L2 Miss Decode Priority Control	<p>L2 Miss Decode Priority Control.</p> <p>If all threads are at the same software set priority, then:</p> <table> <tr><td>0</td><td>L2 miss disabled for use in adjusting decode priority.</td></tr> <tr><td>1</td><td>L2 miss enabled for use in adjusting decode priority.</td></tr> </table>	0	L2 miss disabled for use in adjusting decode priority.	1	L2 miss enabled for use in adjusting decode priority.						
0	L2 miss disabled for use in adjusting decode priority.											
1	L2 miss enabled for use in adjusting decode priority.											
18	TLB Miss Decode Priority Control	<p>TLB Miss Decode Priority Control.</p> <p>If all threads are at the same software set priority, then:</p> <table> <tr><td>0</td><td>TLB miss disabled for use in adjusting decode priority.</td></tr> <tr><td>1</td><td>TLB miss enabled for use in adjusting decode priority.</td></tr> </table>	0	TLB miss disabled for use in adjusting decode priority.	1	TLB miss enabled for use in adjusting decode priority.						
0	TLB miss disabled for use in adjusting decode priority.											
1	TLB miss enabled for use in adjusting decode priority.											
19	Reserved	Reserved.										
20	Dispatch Flush Sync Control Enable	<p>Dispatch Flush Sync Control Enable.</p> <p>Stop decode if mode for thread with sync instruction is outstanding, (always on in shipping mode). Applies only to SMT2 and higher.</p>										
21:23	Reserved	Reserved.										

Bits	Field Name	Description
24	PTEsync Dispatch Stall	PTEsync dispatch stall. Set to '1' to enable the following function for ptesync instruction. Hold the ptesync instruction at dispatch until GCT is empty, SRQ is empty, no I-side table walk pending, and wait for a minimum of 15 cycles. After these conditions are satisfied, dispatch the ptesync . Afterwards, wait for GCT empty and SRQ empty to continue with the dispatch of further instructions.
25:31	Reserved	Reserved.

The TSCR can be accessed with the **mtspr** and **mfsp** instructions using SPR 921.

A SCOM latch is added to make this register as read-only for debug purposes.

TSCR is initialized to x'0000_0000' at power-on.

The preferred setting is: x'8ACC_6880'.

5.9 Thread Time-Out Register (Hypervisor only)

The Thread Time-Out Register (TTR) is used to ensure forward progress. There is one TTR per core. For more information see *Section 5.10 Forward Progress Timer* on page 144.



Bits	Field Name	Description
0:43	Reserved	Reserved.
44:63	Thread Time-out Flush Value.	Thread Time-out Flush Value. A x'00000' value generates a maximum count.

The TTR Register can be accessed with the **mtspr** and **mfsp** instructions using SPR 922.

The TTR Register is initialized to x'0000_0000_0000_0000' at power-on.



5.10 Forward Progress Timer

For the POWER8 core, the forward progress timer bits were moved out of the PPR Register. A non-architected latch bank holds these bits. PPR[44:63] are now reserved, non-implemented bits.

The forward progress latch bits are loaded from TTR[44:63] every time a group of instructions are retired on the current thread.

If the current thread is not in a decode stop state, the counter is decremented by ‘1’ every time a group completes on another thread in the same thread set (see TSCR[14] in *Section 5.8 Thread Switch Control Register (Hypervisor Access Only)* on page 141).

Initialized by scan flush to x‘00000’ (maximum count)
Decrementer stops at x‘00001’ (minimum value)

For SMT2 and higher:

A flush of the other active threads in the same thread set occurs when:

- The timer count reaches x‘00001’.
- The forward progress count flush is enabled TSC[13] = ‘1’.
- The group completes on another thread in the same thread set.

After the threads are flushed, no dispatch slots are given to the flushed thread until one group has completed for the current thread.

5.11 Thread Priority Boosting

Hardware typically does not change the thread priority value in the PPR, unless an **mtPPR** or one of the priority changing NOP instructions is committed. However, on the POWER8 core, problem-state programs can change the thread priority value to medium-high (‘5’) depending on the contents of the Problem-State Priority Boost Register (PSPBR). The problem-state boosting changes the contents of PPR[11:13].

The thread priority can be boosted internally by the hardware (in a software invisible manner) in certain cases (as described in *Section 5.13 Thread Priority Boosting on Asynchronous Interrupt* on page 145) to medium priority (‘4’). The boosting of thread priority for pending asynchronous interrupts does not affect the actual architected thread priority value in the PPR. Therefore, if the software does a **mfPPR** at any time during the asynchronous boosting, it always gets the last priority value explicitly set by the software for that thread.

5.12 Priority Boosting to Medium-High in User Mode

The POWER8 core allows a problem-state program that is executing on a thread to temporarily change the PPR thread priority value to medium high (‘5’) by executing an **mtPPR** or priority NOP. The temporary thread priority boosting is controlled by a 32-bit privileged Problem-State Priority Boost (PSPB) Register. There is one PSBPR per thread, which is set by a move-to PSPB.

A problem state program can set the program priority to medium-high only when the PSPB of the thread contains a nonzero value. The maximum value to which the PSPB can be set must be a power of 2 minus 1. Bits that are not required to represent this maximum value must return ‘0’s when read, regardless of what was written to them.

When the PSPB of the thread is set to a value less than its maximum value but greater than '0', its contents decrease monotonically at the same rate as the **SPURR** until its contents minus the amount it is to be decreased are '0' or less when a problem state program is executing on the thread at a priority of medium high.

When the contents of the PSPB minus the amount it is to be decreased are '0' or less, its contents are replaced by '0'. When the PSPB is set to its maximum value or '0', its contents do not change until it is set to a different value.

Whenever the priority of a thread is medium high and either of the following conditions exist, hardware changes the priority to medium:

- PSPB counts down to '0', or
- PSPB = 0 and the privilege state of the thread is changed to problem state (MSR[PR] = '1')

While in problem state at medium-high priority, there can be the potential of the PSPBR reaching '0' at the same time a priority NOP or **mtPPR** is trying to lower the thread priority to a value less than medium. If the attempted write to the PPR occurs in the same cycle, the priority NOP or **mtPPR** must update the PPR with its thread priority instead of allowing the PSPB reset to set the PPR to medium priority.

5.13 Thread Priority Boosting on Asynchronous Interrupt

If a thread has a priority less than medium ('4'), the priority of the thread is boosted on a pending asynchronous interrupt. This allows the interrupt to be serviced faster for a thread (that is waiting for the interrupt at a low-priority state). TSCR[12] is used to enable or disable priority boosting for any pending asynchronous interrupt. The boosting of thread priority does not affect the actual architected thread priority value in the PPR. Therefore, if the software does a move from PPR (**mfPPR**) at any time during asynchronous boosting, it always gets the last priority value explicitly set by the software for that thread.

5.13.1 When to Boost Thread Priority

Thread priority is boosted internally by the hardware on an asynchronous interrupt based on *Table 5-6*. After the priority is boosted, the hardware continues to treat the thread at medium ('4') priority until there is an **mtPPR** or priority NOP instruction that changes the thread priority.

Table 5-6. Asynchronous Interrupt

Interrupt	MSR Bits
Debug	EE = 1 or HV = 0 or PR = 1
Hmaintenance	EE = 1 or HV = 0 or PR = 1
External	EE = 1
Perfmon	EE = 1
HDEC	EE = 1 or HV = 0 or PR = 1
DEC	EE = 1
System Reset	Always (ignores TSCR[12])



In other words, there is no thread priority boosting for:

- Debug, Hmaintenance, or HDEC interrupt, if in hypervisor with EE = '0'
- External, Perfmon, or Decrementer interrupt, if EE = '0'

Otherwise, boost the priority on a pending interrupt. The reasons for not boosting the priority in the previous cases include:

- Operating system and hypervisor: Spinning on a lock, the priority is low and MSR[EE] = '0'. Priority must not be boosted because nothing useful is going to happen until the lock is acquired. Before the **stdcx** can get the lock, high priority is asserted. Therefore, if a thread is holding a lock, its priority does not need to be boosted when an external interrupt becomes pending.
- Operating system interrupt handler running with MSR[EE] = '0'. Priority is already at the desired level as a result of the implicit or explicit boost. No additional boost is required by the hardware.

5.14 Thread Prioritization Implementation

5.14.1 Thread Switch Fetch Priority

The thread priority is used to apportion fetch cycles. For example, if two threads have a priority weighting that is different, the ratio of those two weights determines the relative number of cycles those two threads will be given access to the I-cache. If all threads have equal priority, the threads are accessed in a round-robin manner.

The instruction fetcher makes no priority provisions for an asymmetric SMT environment. For example, if one side of the core has one thread, and the other side has more than one thread, and all threads are of equal priority, then each thread gets an equal number of fetch cycles, even though there are more decode resources available for the thread that is on its own side.

If all of the threads have the same priority, the fetcher fetches the threads in an order that tries to swap from one core side to the other as much as it can. This is desirable, because fetching multiple cycles on the same core side increases the chance that no instructions are available to decode/dispatch on the other core side.

Normally, a thread uses its fetch cycle if there is a chance that the fetch can result in a transfer. There are several cases where a thread relinquishes its fetch cycle and allows it to be skipped over. The cases where this happens are as follows:

- The thread is an I-cache miss pending or I-ERAT miss pending.
- The IBuffer is full for eight cycles, such that it is unlikely that there is room in the IBuffer when the instructions are fetched from the I-cache.
- There is a hold fetch from either the ISU or the pervasive core unit that tells us we cannot fetch on that thread.

When all the highest priority threads give up their cycle, there are times when the base hardware algorithm cannot assign another thread based on the priority. On the cycles when this occurs, the other threads that can be fetched take turns fetching, ignoring the thread priority.

If there is a flush on a thread and that thread is not already being selected, that thread is selected on the next IFM1 cycle, which is done to reduce the average latency on a flush.



5.14.1.1 SMT2 Fetch Pattern

In SMT mode, the target of a predicted taken branch can be fetched three cycles after the branch instruction is fetched. If threads are alternated in SMT2 mode, the earliest time that an instruction could be fetched would be allocated to the other thread, and thus the taken branch penalty goes from three to four cycles.

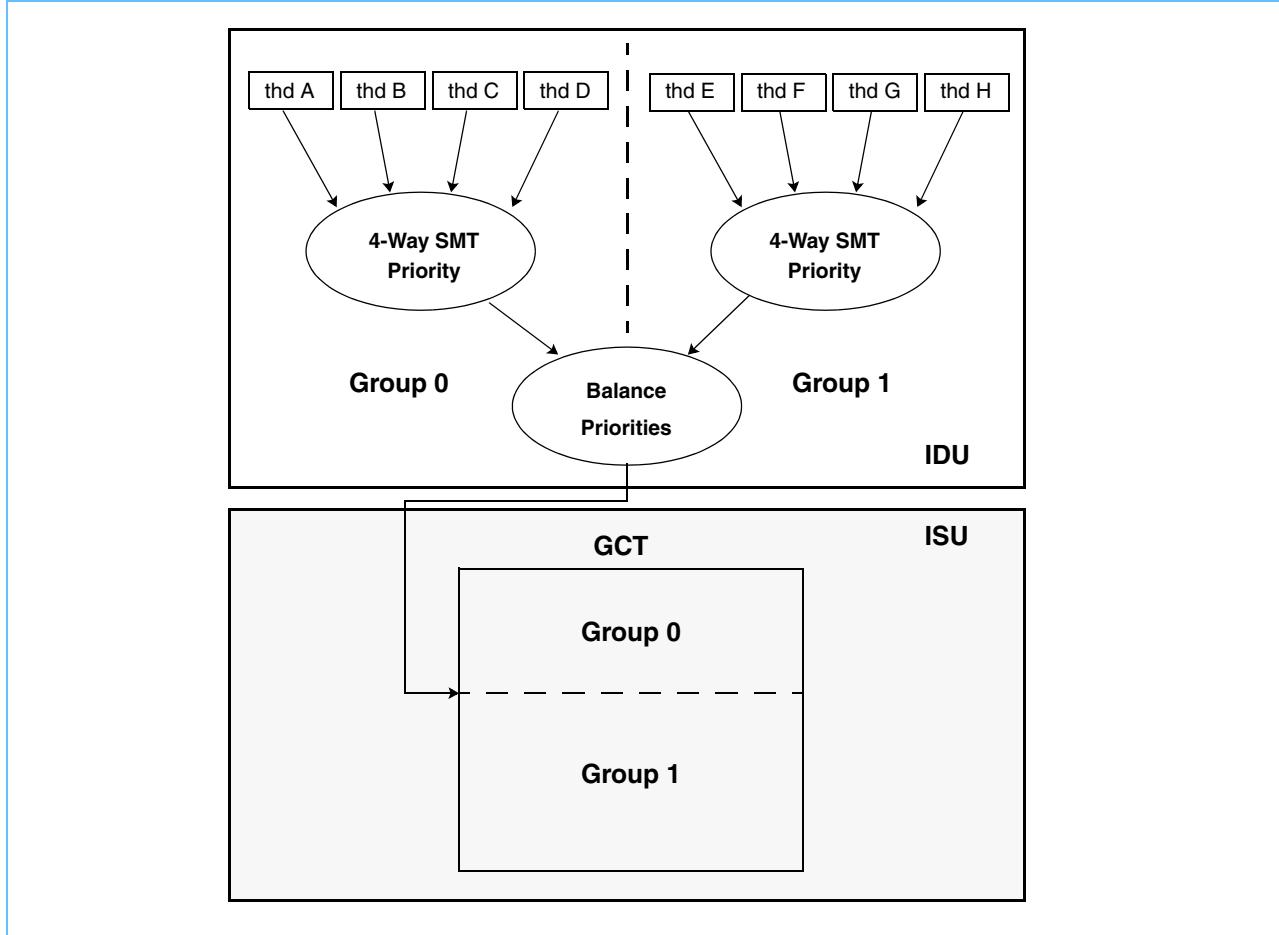
To reduce this effect, use a pattern in SMT2 mode that in most cases allows the same thread to be allocated every third cycle. The pattern implemented in SMT2 mode is '00100100 11011011'. This pattern causes the same thread to be assigned three cycles later 87.5% of the time.

5.14.2 Thread Switch Decode Priority

The instruction buffer, decode, and dispatch are divided into two parallel groups in SMT2, SMT4, and SMT8 modes. Each group gets half the dispatch width. In ST mode, the single thread gets the entire dispatch width. This drives the need for a pair of decode priority engines to run in parallel; one for the threads assigned to thread set 0 and one for those assigned to thread set 1. In SMT modes, each thread set operates in its own half of the dispatch group, independent of the other thread set.

To support two independent thread sets that dispatch in parallel, two SMT4 thread priority engines are used. One controls the decode cycles of the threads assigned to thread set 0. The other controls thread set 1. Each engine can manage 1 - 4 threads, depending on how many are assigned to the thread set. Additionally, the decode cycle selections of the two engines are balanced by controlling the number of GCT entries each thread set has access to. The GCT limits are set based on the relative values of the sum of thread set 0 priorities and the sum of thread set 1 priorities. As one side of the GCT fills up, the ISU naturally throttles back that thread set on the dispatch hold interface. This achieves SMT2, SMT4, and SMT8 thread management at the system level, while allowing two thread sets to decode and dispatch in parallel.

Figure 5-1. Dual SMT4 Decode Priorities



In SMT mode, decode cycles (opportunities to form instruction groups out of the IBuffer) are given to a thread based on the following ordered criteria:

1. Thread enabled, no slots given to a stopped thread, CTRL[8:15].
2. Per thread decode stops for Decode Hold. See *Section 5.16 Controlling the Flow of Instructions in SMT* on page 152.
3. Instruction availability in the threads IBuffer. Used only if the priority in PPR[11:13] is equal for all enabled threads.
4. Software-set thread priority, controlled in Thread Status Register (PPR[11:13]). See *Section 5.14.3 Software-Set Thread Priority* on page 149.
5. Dynamically changing thread decode priority. See *Section 5.14.5 Dynamic Thread Priority* on page 149.

The first three criteria are used only to eliminate threads from consideration for the next decode cycle. If all available threads are eliminated based on those three criteria, no thread forms an instruction group next cycle. Otherwise, the remaining eligible threads are considered for the next decode cycle according to either their software-set thread priority or a dynamic thread priority algorithm using the current state of each eligible thread.

5.14.3 Software-Set Thread Priority

Software-set priority is used to determine the thread to receive the next decode cycle if at least one of the enabled threads has a different Thread Priority value in PPR[11:13] than the other enabled threads. The intention for the software-set priority algorithm is to divide decode cycles according to the relative values of the thread priority values.

POWER8 core control is given by allowing the user to define the weightings between the seven priorities. A 64-bit SPR, Relative Priority Register (RPR), is provided for the user to set any 6-bit value (0 - 63), for each of the seven priority levels (very low, low, medium low, normal, medium high, high, extra high). Then, PPR[11:13] for each active thread determines which value to read from the RPR.

Each active thread receives a number of decode cycles, relative to the other threads, equal to their priority values. For example, within thread set 0, T0 has a relative priority of 17 (as defined by PPR[11:13] and the RPR), T3 has a relative priority of 6, and T7 has a relative priority of 45. Within a window of $17 + 6 + 45 = 68$ decode cycles, T0 gets 17 cycles, T3 gets 6 cycles, and T7 gets 45 cycles. The pattern repeats every 68 decode cycles. Additionally, the cycles given to each thread are distributed as evenly as is reasonably possible within the pattern.

5.14.4 Low-Power Modes for Application

The POWER8 core slows the rate of group formation and decodes to reduce power whenever all enabled threads have a priority of '001', as set in each thread's PPR[11:13]. This has different effects depending on the SMT mode (CTRL[24:27]).

In ST mode, the single thread receives one decode cycle every 128 cycles.

In SMT2 mode, the threads in Thread Set 0 and Thread Set 1 each receive one decode cycle every 128 cycles. Thread Set 0 and Thread Set 1 are independent, and their decode cycles are not guaranteed to be concurrent or non-concurrent.

In SMT4 and SMT8 modes, one thread from Thread Set 0 receives one decode cycle every 128 cycles, similarly for Thread Set 1. As in SMT2, the two thread sets operate independently. Within a given thread set, the threads are queued in a modified round robin fashion. Every 128 cycles, a thread is selected from the queue and given a decode cycle. The thread at the head of the queue receives the next available decode cycle unless its IBuffer is empty or it is otherwise removed from eligibility (ISU applies an IBuffer hold, and so on). If the thread is ineligible, the next thread in line is selected. After a thread receives a decode cycle, it is moved to the back of the queue.

5.14.5 Dynamic Thread Priority

If all enabled threads have the same thread priority value, a dynamic thread priority algorithm based on the state of the eligible threads determines which will get the next decode cycle. This algorithm uses a scoring system built on the number of GCT entries occupied by each thread, whether the thread has any outstanding L2 or TLB misses, and a final round-robin adder used only to break any ties between threads. This algorithm is implemented twice, once per thread set, where each algorithm manages 1 - 4 threads. The two thread sets are then balanced using the GCT as described in *Section 5.14.2 Thread Switch Decode Priority* on page 147.



Table 5-7 lists a summary of the scoring system.

Table 5-7. Scoring System Summary

Description	Score Adder
If a thread occupies fewer GCT entries than the threshold set in TSCR[0:1]	+10
If TSCR[0:1] = '00'	+0
If TSCR[17] = '1' and the thread has no outstanding L2 misses,	+5
If TSCR[18] = '1' and the thread does not have an outstanding TLB miss.	+5
A final value is added based on the thread's position relative to a round-robin pointer:	
Thread being pointed to	+4
Thread that shares downstream resources with the pointed thread (thread0 shares with thread1, thread 2 with thread3)	+3
Thread with the same LSB as the pointer (thread0 with thread2, thread1 with thread3)	+2
Remaining thread	+1

The eligible thread with the highest overall score is given the next decode cycle. Note that the round-robin pointer only affects results in the event of a tie from the other three adders. To ensure fairness between threads when one or more threads are disabled, the round-robin pointer rotates between all threads available in the current SMT mode regardless of whether the thread is enabled.

5.15 Support for Multiple LPARs

In cloud-computing mode, the core can be partitioned into 2 or 4 LPARs. When in 2 or 4 LPAR mode, the SMT mode is always forced to SMT8 so that most resources are allocated based on *Table 5-2 Front-End Execution Core Resource* on page 135.

5.15.1 Instruction Fetch

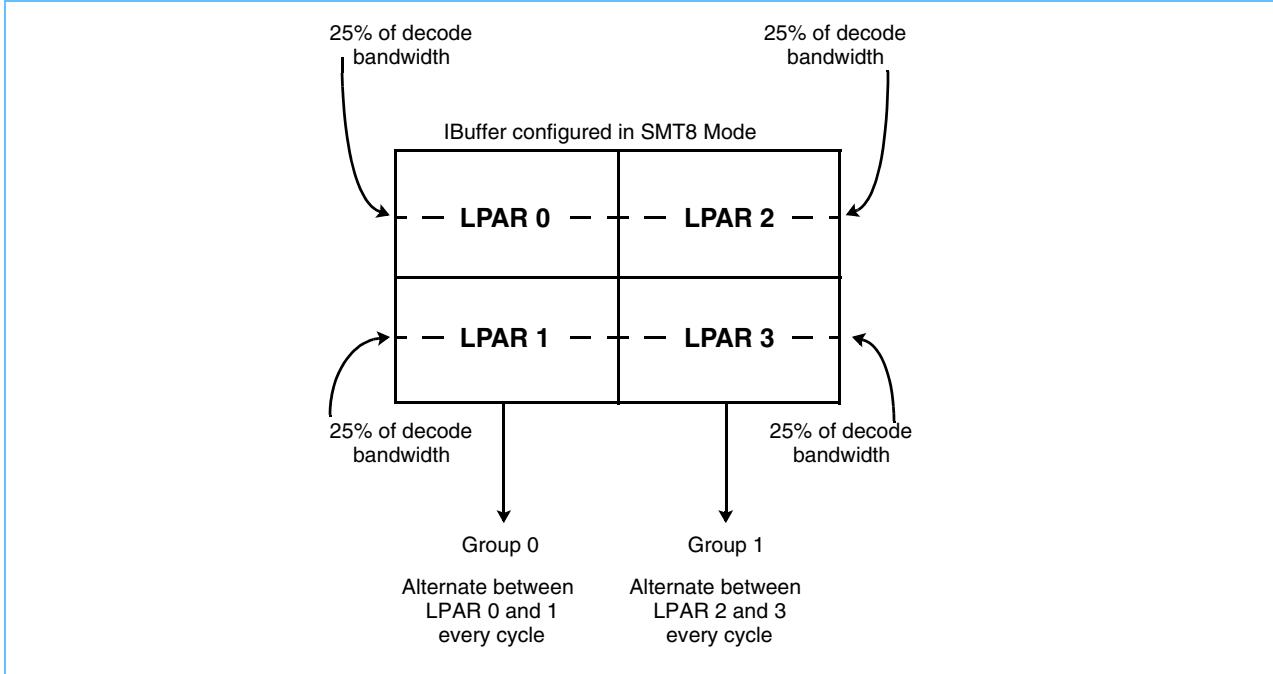
When the processor is configured in 2 or 4 LPAR mode, fetch cycles are allocated in a round robin manner to the partitions. This implies each partition will always get an equal percentage of the fetch cycles. If a partition does not have any threads that are ready, those cycles will be allocated to another partition, as long as other partitions have threads that are ready to fetch. If a partition contains multiple threads, the partition's fetch cycles for the threads in that partition are divided based on the relative software priority weighting.

5.15.2 Decode

In 2 LPAR mode, each partition gets 50% of the decode bandwidth. In 4 LPAR mode, each partition gets 25%. Each LPAR is given its 25% or 50% of the decode bandwidth, regardless of how many LPARs are active.

For 4 LPAR mode, one or two threads can be active. If an LPAR has two threads active, an SMT2 priority is followed as described in *Section 5.14.3 Software-Set Thread Priority* on page 149. Each thread is given cycles within its partitions time slice, relative to the other thread in its partition, based on their relative priorities. See *Figure 5-2* for details. For 2 LPAR mode, an SMT4 priority is followed.

Figure 5-2. Decode Priority in 4 LPAR Mode



5.15.3 Microcode Fairness

In multi-LPAR mode, the goal is to give each LPAR the same number of dispatch cycles. However, multi-cycle microcode instructions can cause an LPAR to consume multiple consecutive cycles. To remedy this, give the other LPAR sharing the dispatch bandwidth extra decode cycles to compensate for the loss of decode cycles. When microcode operations are in flight, each operation generates 32 groups (such as, load multiples).

The counter handles $32 \times 12 = 384$ catch-up cycles for either LPAR. A 10-bit counter handles up to 512 cycles for each LPAR.

5.15.4 Instruction Cache

I-cache allocation is altered for LPAR modes.

5.15.5 Thread Set Allocation

Threads 0 - 3 are allocated to side 0, and threads 4 - 7 to side 1, for both 2 LPAR or 4 LPAR mode. Thread set reconfiguration is disabled.

5.15.6 Data Cache

Threads 0 - 1 reload into sets 0 - 1, threads 2 - 3 reload into sets 2 - 3, threads 4 - 5 reload into sets 4 - 5, and threads 6 - 7 reload into sets 6 - 7



5.15.7 ERATs

Threads 0 - 3 share half of the primary/secondary ERAT, while threads 4 - 7 share the other half.

5.16 Controlling the Flow of Instructions in SMT

The ability to control the flow of instructions in an SMT processor is very important for the performance improvement due to SMT. When one thread is not making good progress (which can happen due to many reasons, such as an L2 miss, TLB miss, **sync** or other long latency operations), the other thread must be allowed to have as much of the machine resources as necessary for it to make progress. The following features are built into the POWER8 core for controlling instruction flow.

5.16.1 Dispatch Flush

A dispatch flush is a low-latency flush that flushes the decode pipe. **Sync**, **Iwsync**, **ptesync**, **tlbsync**, **tlbie**, and instructions with the scoreboard bit set can cause a dispatch flush on the POWER8 core. Also, if enabled, a thread that was balanced flushed is dispatch flushed if the chosen thread is stalled at dispatch.

5.16.1.1 Dispatch Flush Rules

1. Dispatch flushes are disabled if the core is in single-thread mode, or if the core is in SMT2 or SMT4 mode and there is one or fewer than one threads active.
2. Dispatch flush only occurs if a thread shares a group with another thread. SMT2 mode should not dispatch flush if each thread has its own group.
3. If both threads are on group0 (not balanced), dispatch flush can still occur.
4. Never dispatch flush a thread if it is in the middle of a microcode. Dispatch flushes related to SMT-performance are never done if in the middle of a microcode dispatch. Other dispatch flushes can happen to microcode, however, such as quiesce, **RAS**, or a forward-progress timeout.
5. If a thread's virtual **LRQ** or virtual SRQ is full, it gets dispatch flushed, irrespective of its thread priority.
6. A **sync**, **Iwsync**, **ptesync**, or **tlbsync** instruction from thread A causes a dispatch flush of thread A (similarly, for other threads).
7. If the GCT thread is not empty, the **tlbie** instruction gets dispatch flushed. The instruction that follows the **tlbie** is dispatch flushed if the **tlbie** instruction has not received **tlbie** acknowledge from the Nest, through the LSU.
8. If thread A has its scoreboard bit set (such as, non-renamed **mtspr** followed by **mfsp**), thread A is dispatch flushed (similarly, for other threads).
9. The instruction that follows an instruction that stalls dispatch until GCT empty can cause a dispatch flush. Instructions that stall dispatch until GCT empty include: **mtamr**, **treclaim**., and power mode instructions: **doze**, **nap**, **sleep**, **rvwinkle**, and **sp_attn**.
10. If TSCR[3] is enabled, dispatch flush the thread that was chosen for a balance flush if that thread is stalled at dispatch.
11. If a thread is stalled at dispatch because its unified queue half is full and the core is in SMT4 mode, then dispatch flush the stalled thread.

5.16.1.2 Stall at Dispatch

A thread can be *stalled at dispatch* due to unavailability of a shared resource that it needs for the next dispatch. When *stalled*, `dispatch_hold` is asserted to hold the decode pipe. The complete list of *stall* conditions follows:

- Required issue queue entries are not available
- Not enough renames are available for the entire group
- GCT is full

Note:

1. A 5-cycle delay is expected between counting the GCT entries and taking any action based on that count.
2. A 4-cycle delay is expected between the detection of a stall condition and causing a dispatch flush.

5.16.2 Decode Hold

1. After a dispatch flush, a `sync`, `lwsync`, `ptesync`, `tlbsync` instruction is held at IBUF until its dispatch conditions are met. This is done irrespective of thread priority.
2. After a dispatch flush, a `tlbie` instruction is held at IBUF until the GCT thread is empty. After the `tlbie` instruction is dispatched, the very next instruction gets dispatch flushed and then held at IBUF until GCT and the SRQ are empty. This is done irrespective of thread priority.
3. After a balance flush due to an L3 or TLB miss, instructions on the balance flush thread are held at IBUF until the miss is resolved or until the balance flush miss counter reaches the threshold value as determined by TSCR[8:9].

5.16.2.1 Balance Flush

A balance flush is an NTC+1 flush that flushes all instructions on a selected thread that are younger than the next-to-complete instruction group. It flushes the execution units, GCT, and EAT for the selected thread. Threads are considered for a balance flush only if a thread is stalled at dispatch. Balance flushes can be disabled using TSCR[2].

Criteria for Selecting a Thread to be Balanced Flushed

If the core is in SMT mode and more than one thread is active, then on a dispatch stall perform the following steps:

1. Select the threads with any number L3 or TLB misses, regardless of the balance flush miss counter value. If enabled by a debug switch, only select the threads with any number of L3 or TLB misses, if the balance flush miss counter for the thread is less than the counter threshold as described by TSCR[8:9]. If the miss counter for the thread is greater than the threshold value, ignore the miss on that thread and do not consider the thread for a balance flush.
2. If only one thread has an L3 or TLB miss, select that thread to be balanced flushed. Raise the CLB hold on the thread that was chosen to be balanced flushed until either the miss has been resolved or the balance flush miss counter threshold has been reached as described by TSCR[8:9].
3. If in SMT4 or SMT8 and more than one thread is eligible to be balanced flushed based on having an L3 or TLB miss and a GCT count over the balance flush GCT threshold TSCR[4:5], select all eligible threads to be balanced flushed. Raise the CLB hold on the threads that were chosen to be balanced flushed until



either the miss has been resolved or the balance flush miss counter threshold for that thread has been reached as described by TSCR[8:9].

4. If in SMT2 and both threads are eligible to be balanced flushed based on having an L3 or TLB miss and a GCT count over the balance flush GCT threshold TSCR[4:5], do not balance flush either thread. Do not raise the CLB hold.
5. If no threads have an L3 or TLB miss, select the threads with GCT counts over the TSCR[6:7] balance flush no-miss GCT threshold to be balanced flush. For the POWER8 processor, the default is to allow this option in SMT2, with a mode to disable it. It can be disabled for all SMT modes by setting TSCR[6:7] to '00'.
6. If the thread that is stalled at dispatch is also the thread that was chosen to be balanced flushed, then also do a dispatch flush on that thread if TSCR[3] is enabled. Otherwise, do only a balance flush on the chosen thread.

5.17 Dynamic Thread Transitioning

5.17.1 Overview

Any thread can be run in any SMT mode, on any thread set. Starting and stopping a thread, without crossing an SMT mode boundary (SMT1/2/4/8), can be done seamlessly without having to quiesce the core. However, if an SMT boundary is crossed, a quiesce (and optionally additional action) is required, as detailed in *Section 5.17.3 SMT Mode Boundary Crossings* on page 155. Note that the POWER8 processor treats the DOZE instruction as a NAP instruction.

5.17.2 Thread Set Definition

If only one thread is running, it is guaranteed to be on thread set 0. As threads are added, they are assigned to alternating thread sets, to strive to achieve balance in the number of threads per thread set. The GCT sets are tied to the thread sets and can complete one thread per cycle from each thread set. The dispatch groups are also tied to the thread sets. As threads are disabled, the thread set definition remains the same. Consequently, the remaining enabled threads can have imbalanced thread sets, which is addressed in *Section 5.17.4.1 Balancing* on page 155.

5.17.3 SMT Mode Boundary Crossings

When SMT mode boundaries (SMT1/2/4/8) are crossed, different resources throughout the core are reconfigured to maximize performance in a certain SMT configuration. All threads must be quiesced before any reconfiguration can occur.

Table 5-8 on page 155 shows how the different resources must be reconfigured in the various SMT mode boundary crossings. If multiple boundaries are crossed simultaneously, the action is the OR of the respective boundary-crossing functions listed in *Table 5-8*.

Table 5-8. SMT Mode Boundary Crossing Reconfigurations

From	To	Drain AMCs	Reconfigure IFU	Reconfigure IDU	Reconfigure LSU	Notes
SMT1	SMT2			Y	Y	
SMT2	SMT4		Y			
SMT4	SMT8		Y			
SMT8	SMT4		Y			
SMT4	SMT2		Y			
SMT2	SMT1	Y		Y	Y	Place remaining thread in thread set 0.

It is possible that, when an SMT mode is lowered, there might be too many threads in one of the thread sets, or the number of threads per thread set might be imbalanced. *Table 5-1 SMT Modes* on page 135 shows the maximum allowable number of threads per thread set. The corrective action is to reconfigure thread sets as described in *Section 5.17.4 Thread Set Reconfiguration*.

5.17.4 Thread Set Reconfiguration

Threads can be moved to a different thread set, for a variety of reasons. To do so, all threads must be quiesced.

5.17.4.1 Balancing

There can be an imbalance in the number of threads per thread set, either after an SMT mode change, or as threads are disabled. *Table 5-9* shows the different scenarios where a rebalance can be triggered. Some of the scenarios can only occur after an SMT mode change because once in a given SMT mode, there is a maximum allowable number of threads per thread set, as shown in *Table 5-1 SMT Modes* on page 135. Otherwise, the scenario can occur at any time.

Table 5-9. Thread Balancing Scenarios (Sheet 1 of 2)

Mode	Threads Per Set	Rebalanced Threads Per Set
SMT2	2/0	1/1
SMT4	2/0	1/1
SMT4	3/0	2/1
SMT4	3/1 4/0	2/2
SMT8	2/0	1/1
SMT8	3/0	2/1



Table 5-9. Thread Balancing Scenarios (Sheet 2 of 2)

Mode	Threads Per Set	Rebalanced Threads Per Set
SMT8	3/1 4/0	2/2
SMT8	4/1	3/2
SMT8	4/2	3/3

5.17.4.2 Mixing

There can be an inequality in the processor resources being used per thread set. If one thread set is heavily using a given resource, but the other thread set is not, wasting its allotment of the resource, it might be advantageous to move a thread using the resource to the other thread set, to maximize the usage of the resources. A reconfiguration based on this thread “mixing” can only occur in SMT4 and SMT8 modes.

5.17.4.3 Action

When moving a thread to the opposite thread set, its AMC data must be drained, pushing that thread’s GPR/VRF data out to the SAR, so that it can be reloaded into the opposite register file after it starts executing. If the action is triggered while executing, only perform the action after the trigger remains in the same state after a delayed amount of time, controlled by a programmable 20-bit LFSR counter (up to 1 million cycles). This delay is inserted because a reconfiguration is expensive in terms of latency. Therefore, it is important that the events that caused the trigger are in a relatively steady state. The link stack and D-ERAT must be rebuilt for the moved threads; therefore, performance initially suffers for those threads.

6. POWER8 SMP Interconnect

The POWER8 SMP interconnect fabric controller (FBC) is the underlying hardware used to create a scalable cache-coherent multiprocessor system. The POWER8 SMP interconnect controller provides coherent and noncoherent memory access, I/O operations, interrupt communication, and system controller communication. The FBC provides all of the interfaces, buffering, and sequencing of command and data operations within the storage subsystem. The FBC is integrated on the POWER8 chip, with up to 12 processor cores and an on-chip memory subsystem.

The POWER8 chip has up to three FBC links that can be used to connect to other POWER8 chips. The FBC link is a split-transaction, multiplexed command and data bus that can support up to four POWER8 chips. The bus topology is a fully-connected topology to reduce latency, increase redundancy, and improve concurrent maintenance. Reliability is improved with ECC on the external I/Os.

Cache coherence is maintained by using a snooping protocol. Address broadcasts are sent to the snoopers, snoop responses are sent back in-order to the initiating chip, and a combined snoop response broadcast is sent back to all of the snoopers. Multiple levels of snoop filtering, Chip pump, and remote chip pump, are supported to take advantage of the locality of data and processing threads. This approach reduces the amount of interlink bandwidth required, reduces the bandwidth needed for system-wide command broadcasts, and maintains hardware enforced coherency using a single snooping protocol. Chip pump limits the command broadcast scope to the snoopers found on a physical POWER8 chip. When the transaction cannot be completed coherently using this limited scope, the coherence protocol forces the command to be re-issued to all chips in the system (system pump). Conversely, remote chip pump limits the broadcast scope to a remote POWER8 chip; if the operation cannot complete coherently, the command is re-issued using SystemPump to complete the operation.

6.1 POWER8 SMP Interconnect Features

6.1.1 General Features

- Master command/data request arbitration.
- Command requests are tagged and broadcast using a snooping protocol, enabling high-speed cache-to-cache transfers.
- Multiple command scopes are used to reduce the bus utilizations system wide. The SMP interconnect controller uses cache states indicating the last known location of a line (sent off chip), information maintained in the system memory (snoop filter bits), a coarse-grained MCD indicating when a line has gone off the chip, and combined response equations indicate if the scope of the command is sufficient to complete the command or if a larger scope is necessary.
- The command snoop responses are used to create a combined response, which is broadcast to maintain system cache state coherency. Responses are not tagged. Instead, the order of commands from a chip source, using a specific command broadcast scope, is the same order that combined responses are issued from that source.
- Data is tagged and routed along a dynamically selected path using staging/buffering along the way to overcome data routing collisions. Command throttling and retry command back-off mechanisms are used for livelock prevention.
- Multiple data links between chips are supported (link aggregation).



6.1.2 POWER8 Specific Features

- Chip pump, remote chip pump, and system pump command broadcast scopes with memory domain indicators in cache states, memory, and in the MCD to determine when an increase in the command broadcast scope is required.
- 1- 4 socket configuration support (12-way - 48-way)
- 2x snoop bus support
- 64 local master (LM) system pump queue size (32 per snoop bus)
- 64 chip master (CM) chip pump queue size (32 per snoop bus)
- Service processor accessible SCOM registers for configuration setup

6.1.3 Off-Chip Features

- 2-byte DMI2 intergroup links (A0:A2)
- Aggregate data link support

6.1.4 Power Management Features

- Core chiplet frequency support
 - Doze mode within the floor/ceiling range
 - Doze mode outside the floor/ceiling requires change frequency command
- EX chiplet sleep mode support

6.1.5 RAS Features

- 100% ECC protection on external A-bus off-chip links
 - Single-bit error correction (incoming and outgoing data)
 - Double-bit error detection
- 100% ECC protection on internal data flow
- Livelock recovery mechanism
- Trace array
- Performance monitor
- FIR error reporting
 - Protocol errors
 - Underflow/overflow checkers
 - Asynchronous drop/repeat checkers
 - Parity checkers on coherency register files
- Error injection for system-code debug
 - Single-bit or double-bit errors on external SMP links

6.2 External POWER8 SMP Interconnect

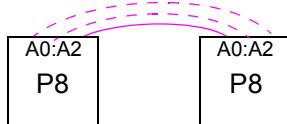
Within the SMP, the off-chip POWER8 SMP interconnect supports up to three coherent SMP links (A0:A2). The interchip A links connect up to four chips. The A links carry coherency traffic, as well as data, and are interchangeable with each other. The A links can also be configured as aggregate data-only links between groups.

The A links are configured as 2 bytes in width (A0:A2, 2-byte DMI2 intergroup links).

6.2.1 POWER8 SMP Interconnect Multichip Configurations

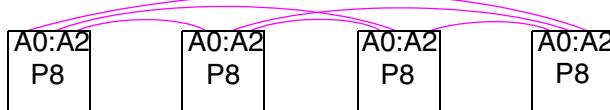
Figure 6-1 and *Figure 6-2* illustrate various POWER8 SMP interconnect multichip configurations.

Figure 6-1. Two Socket Configuration (24-way)



Note: The dotted line indicates an aggregate data-only link that only applies to the POWER8 processor in the SCM.

Figure 6-2. Four Socket Configuration (48-way)



6.2.2 Protocol and Data Routing in Multichip Configurations

The SMP ports configured for coherency are used for both data and control information transport. The use of the buses is as follows:

1. The chip containing the master that is the source of the command issues the reflected command and the combined response to all other chips in the group. The other chips direct the partial response that is a result of the reflected command back to the chip which provided the reflected command. Partial responses are collected at intermediate chips along the directed path from the partial response source to the chip containing the master. The number of partial responses received at the chip containing the master is equal to the number of SMP ports used to broadcast the reflected command.
2. Data is moved point-to-point. For read operations, the chip containing the source of the data directs the data to the chip containing the master. For write operations, the chip containing the master, directs the data to the slave that performs the write operation. Note that the routing tag contains chip and unit identifier information for this purpose.



6.3 Coherency Flow

6.3.1 Physical Broadcast Flow

There are three physical broadcast flows within the POWER8 SMP interconnect.

- Chip Pump (CP) - Command broadcast is limited to the local physical chip
- System Pump (SP) - Command broadcast is to all chips in the system.
- Remote Chip Pump (RCP) - Command broadcast is limited to the local physical chip and a remote physical chip.

Each physical broadcast flow has an independent set of coherency tracking structures (master tag FIFO, ticket counters, and so on).

6.3.2 Broadcast Scope Definition

The POWER8 broadcast terminology used in this documentation denotes the physical broadcast scope for commands. *Table 6-1* describes the physical broadcast scope and the equivalent coherency scope.

Table 6-1. POWER8 Broadcast Scope Definition

Coherency Command Scope	unit_pb_cmd_scope	Physical Broadcast
Chip scope	000	Chip
System scope	010	All SMP chips
Remote chip scope	011	Local/remote physical chips

6.4 Command Ordering Support

The processor bus supports the ability for a master to issue commands in a specific order per scope and for their completion to maintain the same order. The issue order is determined by the order the master issues the command request to the processor bus. The command completion is when the combined response is formed. The processor bus, however, does not guarantee the order in which commands are presented on the reflected command bus or the broadcast combined response bus.

6.5 Memory Coherence Directory

6.5.1 Directory Size

The Memory Coherency Directory (MCD) contains 2 bits per 16 MB page (1 bit for even 8 MB, 1 bit for odd 8 MB) for a total 2 TB of group LPC address space.

6.5.2 Operation

The MCD Lite is the MCD implementation for the POWER8 chip in the SCM. It is a coarse-grained directory of group scope memory domain status (MDS) bits designed to cover the full real address range of the memory controllers on the POWER8 chip. The size of memory address space covered by each bit (granule) of the MCD depends on the memory configuration. The MCD is functionally split into 2 units, called slices,



each connected to one of the two processor bus ports. The two units operate independently and do not share granule state. Additionally, the MCD unit can be enabled by software to probe the memory domain status of cache lines through the processor bus command interface. The MCD can recover the on-group status of a granule if all cache lines are found on group. The MCD also has the ability for software to directly read and write any bit in the MCD arrays. For data reliability purposes, the MCD arrays are protected by ECC.



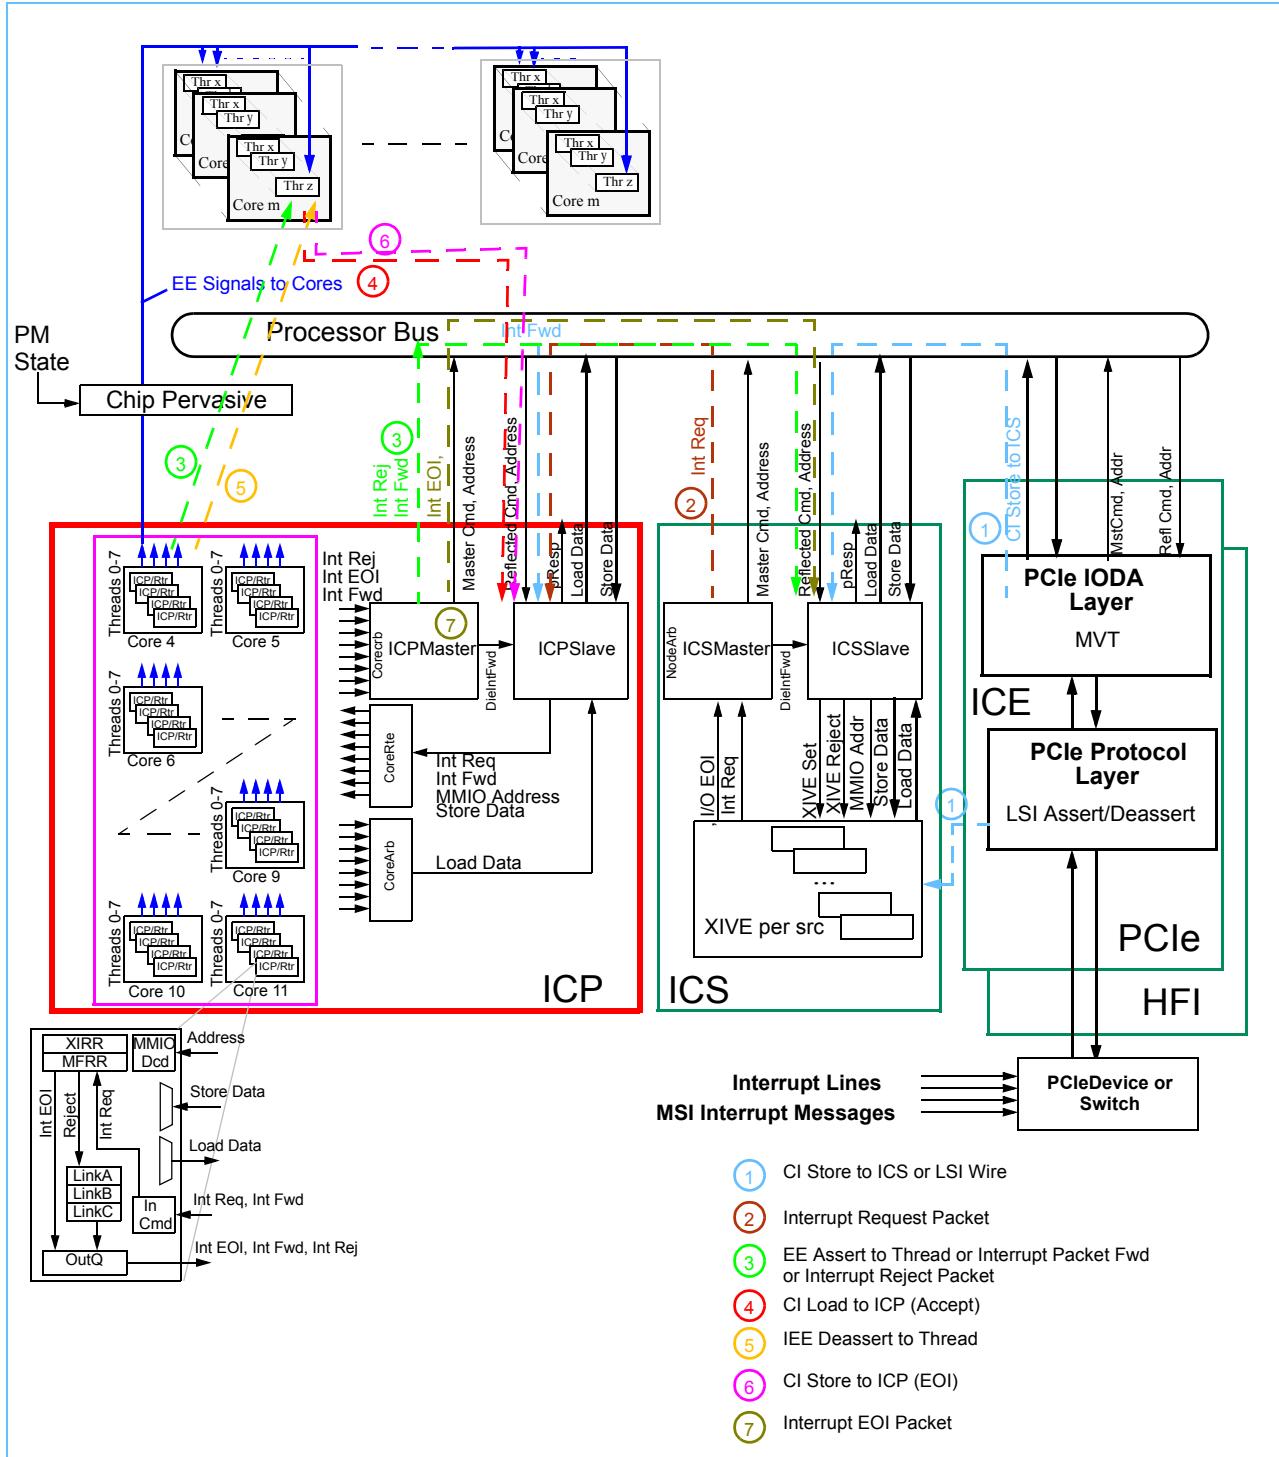


7. Interrupt Control Presenter

The interrupt control presenter (ICP) function combines the interrupt presentation and router functions. It decides which thread is to be notified of an available interrupt or returned to a source layer. The notification to a given thread is via a control wire to each PowerPC thread.

Interrupts reach the ICP via interrupt commands, issued through the processor bus, that either originate from a device attached to a PCIe bus or from an internal interrupt control source (ICS) function, depending on the system configuration. *Figure 7-1* on page 164 depicts the general interrupt structure.

Figure 7-1. POWER8 Logical Interrupt Controller Structure



7.1 Features

The interrupt presenter layer consists of the following facilities:

- Supports 96 threads (12 cores with 8 threads each) for the POWER8 processor in the SCM configuration.
- One External Interrupt Request Register (XIRR) per thread.
 - Current Processor Priority Register (CPPR).
 - External Interrupt Source Register (XISR) - interrupt source number (ISN) determines origin of interrupt.
- Most Favored Request Register (MFRR) per thread for interprocessor interrupts.
- PAPR Type II routing.
 - Three link registers for interrupt forwarding.
 - Link traversal detection creating an interrupt return.
- Continuous evaluation of priority of queued interrupt.
- Resides as a unit on the processor bus. Acknowledges (receives) interrupt request packets and produces interrupt return or end-of-interrupt (EOI) packets back to the requesting ICS or interrupt forward packets to pass the request to another ICP.
- Uses processor bus address-only transactions for interrupt requests, returns, EOIs, or forwarding.
- Interrupt source number of the processor is implemented as a 19-bit field.
- Supports internal ICS and PCIe.
- Has the following queue depths per chip.
 - 8-deep interrupt request queue.
 - 8-deep interrupt return queue.
 - 8-deep external forwarding queue.
 - 16-deep shared CI load/CI store queue.

7.1.1 Routing Layer

The IRL that handles the server number in the interrupt request acknowledges the command and passes it to the associated IPL. If the priority of the interrupt is more favored than the current operating priority, the IPL asserts the external interrupt signal to the thread and loads the interrupt into the XIRR. If the interrupt is less favored than the current priority, the interrupt is given back to the IRL. If the interrupt is a directed interrupt (also known as, specific to this thread), the IRL issues an interrupt return to the ICS for later representation.

If the interrupt is not directed, a field in the interrupt request (link pointer) identifies which of the link registers will be used for forwarding the interrupt to another thread's IRL using the interrupt forward command.

Because the link registers can be set up by software to form a ring, there is the danger that interrupts might be forwarded forever in such a loop. To avoid this, software must set the loop trip bit in one and only one of the link registers that form such a loop. When an interrupt-forward or an interrupt-request command reaches a link register, that has the loop trip bit set, the loop trip bit is carried in and passed on with all newly generated interrupt forward commands generated from that link register. An interrupt forward loop is detected when an interrupt-forward command reaches a link register that has the loop trip bit set and the command already carries the loop trip bit. If an interrupt forward loop is detected, the IRL issues an interrupt return to the ICS for later representation.



7.1.2 Presentation Layer

Each thread has an interrupt management area that contains register facilities that software uses to accept and initiate the end-of-interrupt operation. This interrupt management area is shown in *Table 7-1*. The base address (BA) is established by system firmware to locate the interrupt management area for each thread in the system address space.

Table 7-1. Interrupt Management Area: Interrupt Presentation Layer Ports

Address	Byte 0	Byte 1	Byte 2	Byte 3	Comments
BA+0	CPPR	XISR			XIRR without side effects
BA+4	CPPR	XISR			XIRR with load/store side effects
BA+8	Reserved				
BA+12	MFRR	Reserved or Unimplemented			
BA+16	LINK A				
BA+20	LINK B				
BA+24	LINK C				

The definition of these facilities are defined in *Table 7-2*:

Table 7-2. Facility Definitions

Facility Name	Description
CPPR	The Current Processor Priority Register is 8 bits in length allowing for 256 priorities values. The least favored level is 'x'0FF', and the most favored is 'x'000'. The interrupt presentation layer only signals the associated processor with interrupt conditions that are more favored than the current setting of the CPPR.
MFRR	The Most Favored interprocessor Request Register is 8 bits in length of identical format to the CPPR. Its value indicates the most favored interprocessor interrupt queued for the associated processor. If its value is more favored than the CPPR, an interprocessor interrupt is signaled to the processor.
LINK	One or more Link registers are found in only Type II interrupt presentation controllers. These registers are configured by platform code to form one or more circular linked lists of per processing unit thread interrupt presentation controllers that make up the group servicing a group server queue. The circular linked list replaces the functionality of the interrupt routing layer's Available Processor Mask used by type I interrupt presentation controllers. The Link register is composed of five fields. The length of four of these fields is implementation dependent. However, all Link registers of a given machine provide the same implementation-dependent values. All unimplemented bits read as zeros, and all unimplemented/read-only bits silently ignore writes. Below are defined the fields of the Link registers.
FLAG	The FLAG field of the Link register is a 1-bit long read-only field that defines the number of implemented Link Registers. The FLAG field of all Link registers except for the last implemented Link register FLAG is '1'. The FLAG field of the last implemented Link register FLAG is '0'.
LSPEC	The LSPEC field specifies the LINK register within the targeted interrupt presentation controller that contains the next link in the circular chain. The LSPEC field starts from the low-order bit of the LINK register (bit 31) consisting of M bits where M is the log base 2 of (the number of implemented link registers per interrupt presentation controller +1).
LOOP TRIP	The LOOP TRIP bit in the Link register tells the interrupt forward command to set this bit in the forward packet when an interrupt traverses this register. Once set, the next Loop Trip encounter causes the forwarding to cease and a return packet is generated. This is used by firmware to avoid creating open loops during partition updates.

Software accepts the interrupt condition and is responsible for resetting the interrupt condition through the operations defined in *Table 7-3*.

Table 7-3. Resetting the Interrupt Condition

Load/Store Operation	Interrupt Presentation Function
Load BA+0 (1 byte)	Poll of CPPR value.
Load BA+0 (4 bytes)	Poll of current interrupt status (source and priority).
Load BA+4 (1 byte)	Poll of CPPR value.
Load BA+4 (4 bytes)	Software accepts interrupt. This causes the deactivation of the interrupt signal to the processor thread.
Load BA+12 (1 byte)	Poll MFRR.
Store BA+4 (1 byte)	Set CPPR value.
Store BA+4 (4 bytes)	Software signals end-of-interrupt (EOI) processing. This cause the EOI to be sent to the source.
Store BA+12 (1 byte)	Software signals an interprocessor interrupt event of the priority stored.

When the interrupt presentation layer signals an interrupt to a processor-thread, it loads the XISR with the source of the interrupt. Any nonzero value in the XISR field causes an interrupt to be signalled to the processor. This signal is masked by the processor's MSR[EE] bit before the processor generates the interrupt sequence. If, at a later time, a more favored priority interrupt is made available to the interrupt routing layer, the interrupt routing layer can atomically change the value in the XISR to reflect the source of the more favored interrupt. If a more favored interrupt pre-empts a less favored interrupt in this way, the less favored interrupt is re-presented at a later time.

After the processor has read the XIRR at BA+4, the interrupt routing layer cannot change its mind and either pre-empt or cancel the request.

The XIRR facility appears twice in the external interrupt management area. Address BA+0 is designed to be used with interrupt polling. Address BA+4 has side effects when read or written, and is designed to allow efficient interrupt handler software by having the hardware assist the software in the interrupt queueing process.

The Most Favored Request Register (MFRR) holds the priority of the most-favored request queued on a software managed queue for this processor. When written to a value other than x'FF' the MFRR competes with other external interrupts for the right to interrupt the processor. When the MFRR priority is the most-favored of all interrupt requests directed to the processor, an appropriate value is loaded into the XISR and an interrupt is signaled to the processor. When the processor reads the XIRR at BA+4, the value in the MFRR is loaded by the hardware into the CPPR. The MFRR can be read back by the software to ensure that the MFRR write has been performed.

During the processing of an interprocessor interrupt, the highest priority request is de-queued by the software from the software queue associated with the MFRR and the priority of the next-favored request is loaded into the MFRR by the software.

7.2 Interrupt Control Presenter Registers

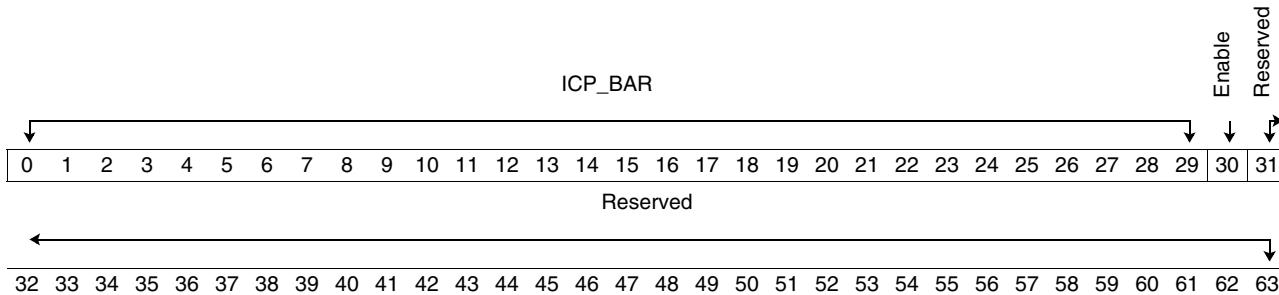
7.2.1 ICP Address Map

Table 7-4. Interrupt Presenter Register Address Map

Mnemonic	Register Name	Address (44:63) ¹	Length
PowerPC System Interrupt Registers:			
	• Address (14:43) = 0xbbbb bbbb (where bbbb bbbb is from ICPBAR Register)		
XIRR	External Interrupt Request Register (CPPR and XISR) for processor thread nnnnnnnn (without side effects)	0nnn nnnn 0000 0000 0000	1 byte / 4 bytes
XIRR	External Interrupt Request Register (CPPR and XISR) for processor thread nnnnnnnn (with load/store side effects)	0nnn nnnn 0000 0000 0100	1 byte / 4 bytes
MFRR	Most Favored Request Register for processor thread nnnn	0nnn nnnn 0000 0000 1100	1 byte
LINKA	Link A Register	0nnn nnnn 0000 0001 0000	4 bytes
LINKB	Link B Register	0nnn nnnn 0000 0001 0100	4 bytes
LINKC	Link C Register	0nnn nnnn 0000 0001 1000	4 bytes
1. Where nnn_nnnn is defined as the following to identify the targeted thread: 45:48 - CoreID and 49:51 - ThreadID			

7.2.2 Interrupt Base Address Register (ICPBAR)

Access: R/W/A/O SCOM only



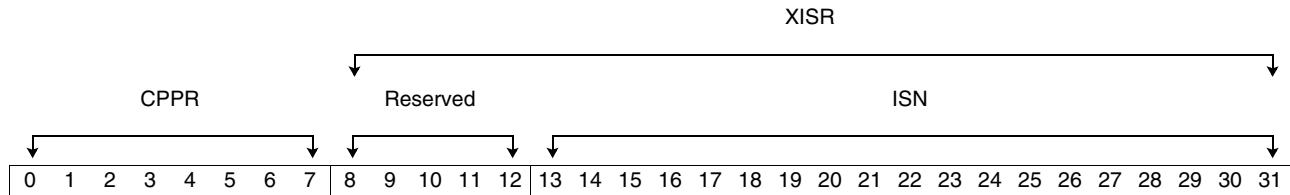
Bits	Field Name	Description
0:29	ICP_BAR	Interrupt Control Presenter Base Address Register. Contents of this register are compared to real address bits 14:43 for a match.
30	Enable	Indicates that ICP_BAR is valid.
31:63	Reserved	Reserved.

This register indicates the 1M region for the PowerPC architected presentation registers within the POWER8 processor. This register is sized to support a 50-bit physical address space.

The following registers are located at an address offset from base address (BA) established by the ICPBAR Register.

7.2.3 External Interrupt Request Register (XIRRt with t = 0 - 7)

Access: R/W (1-byte and 4-byte length to BA+4, 4-byte length only to BA+0)



Bits	Field Name	Description	Initial Value
0:7	CPPR	Current Processor Priority Register Gives the current operating priority of the processor (not the priority of the external interrupt being presented). x'FF' is lowest priority and x'00' is the highest.	x'FF'
8:12	Reserved	Reserved.	
13:31	ISN	Interrupt source number.	

This register is used to pass the ISN to the software on a read and to cause an end-of-interrupt (EOI) on a write. This register also contains the priority register CPPR. There is one register for each processor supported.

The XIRR is made up of two parts: the CPPR and the XISR. The XISR portion contains the ISN of an interrupt being presented. If the XIRR is read when there is no interrupt being presented, the XISR is all zeros.

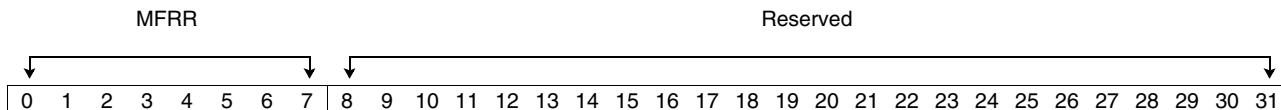
The XIRR can be read at two addresses, BA+0 and BA+4.

- A 4-byte read from BA+0 returns the current CPPR and XISR, but does not have any side effects.
- A 1-byte read from BA+0 returns the current CPPR and has no side effects.
- A 1- or 4-byte write to the BA+0 has no side effects and does not update the XIRR.
- A 1-byte write of the XIRR at BA+4 stores a new value to the CPPR. This results in an interrupt resend being broadcast on the SMP interconnect controller bus if the priority is lowered. This can also result in a pending interrupt being returned to its source with a interrupt return issued on the processor bus.
- A 4-byte write of the XIRR at BA+4 stores a new value to the CPPR. This results in an interrupt resend being broadcast on the processor bus if the priority is lowered. Software must ensure that the level being set in the CPPR is not higher than the current level. The XISR is not modified but the store data is sent as the ISN for the EOI command.

Note: Essentially, software must store the XIRR exactly what it read from XIRR when it received the interrupt. This causes an EOI to the interrupt source and puts CPPR back to its original priority.

7.2.4 Most Favored Request Register (MFRR_t with t = 0 - 7)

Access: R/W (1-byte length) BA+12



Bits	Field Name	Description	Initial Value
0:7	MFRR	Most Favored Request Register. Establishes the priority of an interprocessor interrupt. x'FF' is lowest priority, and x'00' is the highest priority.	x'FF'
8:31	Reserved	Reserved.	

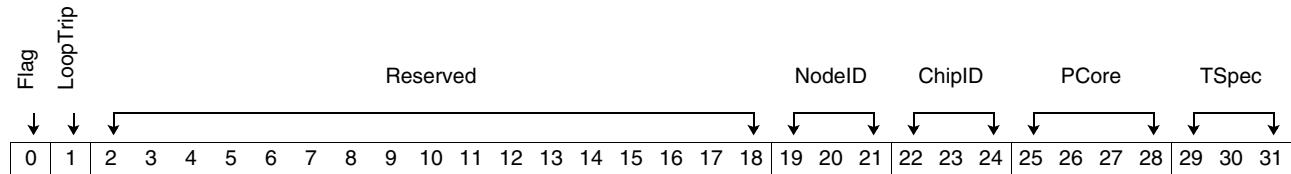
These registers are used by software to send interprocessor interrupts. There is one register for each processor/thread supported.

To send an interprocessor interrupt, software writes into the MFRR associated with the processor that it wants to interrupt. When software has written the MFRR to something other than x'FF', the POWER8 processor activates the interrupt to the associated processor if the MFRR priority is higher priority than the associated processor CPPR and higher than the priority of any pending interrupt to that processor.

When the software reads the XIRR at BA+4, the value in the MFRR and ISSR is loaded by the hardware into the XIRR. Further interrupts from this MFRR are blocked until software writes to this MFRR again. Reads of BA+4 for an interprocessor interrupt returns the ISN = 2. When no interprocessor interrupt is active, software should write a x'FF' to the first byte of this register.

7.2.5 Link Register A (LinkAt with t = 0 - 7)

Access: R/W (4 byte to BA+16)



Bits	Field Name	Description
0	Flag	Indicates that the Link Register is the last one implemented of the three architected.
1	LoopTrip	Indicates that this presenter controls loop checking: 0 Do not set upon forward packet. 1 Set upon forwarded packet.
2:21	Reserved	Reserved.
22:24	ChipID	ChipID. See note 1.
25:28	PCore	Processor core that is the next link in the chain. See note 1.
29:31	TSpec	Thread Specification. Thread within the specific processor core that is the next link in the chain. See note 1.

1. If an interrupt request to the thread in question cannot be presented at this point in time, the interrupt is forwarded to the appropriate node, chip, core, or thread. Such a chain, as mentioned previously, can be created, where an interrupt is forwarded from a possible candidate (thread) to the next candidate. The current test relates to this chain, but is not very expressive.

7.2.6 Link Register B (LinkBt with t = 0 - 7)

Access: R/W (4 byte to BA+20)

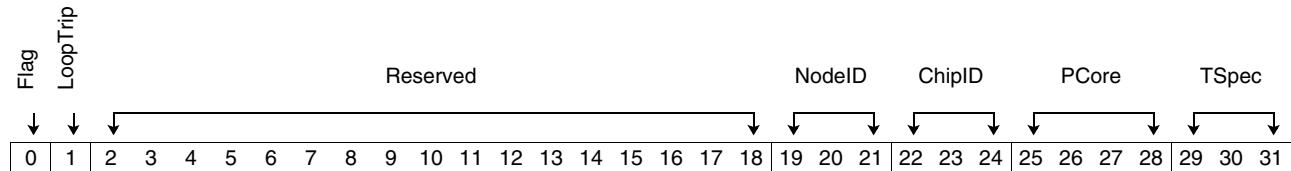


Bits	Field Name	Description
0	Flag	Indicates that the Link Register is the last one implemented of the three architected.
1	LoopTrip	Indicates that this presenter controls loop checking: 0 Do not set upon forward packet. 1 Set upon forwarded packet.
2:18	Reserved	Reserved.
19:21	NodeID	NodeID. See note 1.
22:24	ChipID	ChipID. See note 1.
25:28	PCore	Processor core that is the next link in the chain. See note 1.
29:31	TSpec	Thread specification. Thread within the specific processor core that is the next link in the chain. See note 1.

1. If an interrupt request to the thread in question cannot be presented at this point in time, the interrupt is forwarded to the appropriate node, chip, core, or thread. Such a chain, as mentioned previously, can be created, where an interrupt is forwarded from a possible candidate (thread) to the next candidate. The current test relates to this chain, but is not very expressive.

7.2.7 Link Register C (LinkCt with t = 0 - 7)

Access: R/W (4 byte to BA+24)



Bits	Field Name	Description
0	Flag	Indicates Link Register is the last one implemented of the three architected.
1	LoopTrip	Indicates that this presenter controls loop checking: 0 Do not set upon forward packet. 1 Set upon forwarded packet.
2:18	Reserved	Reserved.
19:21	NodeID	NodeID. See note 1.
22:24	ChipID	ChipID. See note 1.
25:28	PCore	Processor core that is the next link in the chain. See note 1.
29:31	TSpec	Thread specification. Thread within the specific processor core that is the next link in the chain. See note 1.

1. If an interrupt request to the thread in question cannot be presented at this point in time, the interrupt is forwarded to the appropriate node, chip, core, or thread. Such a chain, as mentioned previously, can be created, where an interrupt is forwarded from a possible candidate (thread) to the next candidate. The current test relates to this "chain," but is not very expressive.





8. PCI Express Controller

The PCI Express controller (PEC) bridges between the internal processor bus and the high-speed serial (HSS) links that drive the PCI Express (PCIe) I/O. The PEC acts as a processor bus master on behalf of the PCIe port, converting inbound memory read and write packets into processor bus DMA traffic. The PEC also acts as a processor bus slave, transferring processor load and store commands to the PCIe devices attached to the port.

8.1 Specification Compliance

The PEC is compliant with the following IBM and industry standards:

- POWER Architecture Platform Requirements (PAPR+) Specification, Version 2.1
- I/O Design Architecture v2
- PCI Express Base Specification, Revision 3.0

8.2 PEC Feature Summary

- PCIe Generation 3 Root Complex (RC)
 - Backwards compatible with Generation 1 and 2
 - 2.5, 5.0, and 8 GT/s signalling rate
- 32 PCIe I/O lanes configurable to three independent root complexes for the POWER8 SCM configuration
- Each root complex with 256 partitionable endpoints (PE) for LPAR support
- TCE based address translation for DMA requests.
 - 50-bit address support
 - Translation validation table based on PCI routing ID
- 2048 MSI interrupts per RC
- Eight LSI interrupts per RC
- IBM enhanced error handling (EEH) support
- Processor bus cache-inhibited space segmented by PE:
 - PCI 32-bit memory space segmented into 256 domains by the memory domain table
 - PCI 64-bit memory space segmented by 16 M64 BARs with 16 segments each
- Support for ECRC
- Support for lane swapping
- Support for TLP hints



8.2.1 Supported Configuration

The 32 lanes of HSS I/O can be configured to support three independent PCI buses. *Table 8-1* describes the maximum lane allocation. In addition to supporting PCI operations, the HSS I/O can be allocated for use by the processor bus SMP interface.

Table 8-1. Supported I/O Configurations

PEC0	PEC1	PEC2
16	16	Unused
16	8	8
8	16	Unused
8	8	8
Unused	16	Unused
Unused	8	8

9. Power Management

9.1 Overview

The POWER8 processor uses a number of the more traditional dynamic power-saving techniques, such as clock gating latches and arrays when they are not needed. These techniques make it possible to reduce peak power and therefore, thermal design point power (TDP). They also make it possible to dynamically power gate (turn the power off to) individual cores or full core chiplets when the core is not being used.

The POWER8 processor has an adaptive power management technique to reduce average power, collectively known as EnergyScale™, to proactively take advantage of variations in workload, environmental condition, and overall system utilization. This, coupled with a policy direction from both the customer and feedback from the Hypervisor/operating system that is running on the machine, is used to determine modes of operation and the best power and performance trade-off to implement during runtime to meet customer goals yet achieve best possible performance.

The POWER8 processor extends support to enable:

- Micropartition energy management by providing more synchronous input to the hardware platform upon LPAR switches through hypervisor control Pstate (performance states)
- Faster adaptive monitoring and actuation through the on-chip controller (OCC) (one per POWER8 chip)
- Better instrumentation and control for memory power management (including partition-level memory monitoring and throttling)

Managing the power and performance trade-off is a complex problem. There are many ways to control the behavior of the hardware, but these also have a number of side effects, all of which vary based on the workload being processed. Because there is no “one-size fits all” policy that can be implemented, the POWER8 processor supports an adaptive approach to the problem in the form of a joint hardware, firmware, and software solution.

9.2 Power Management Infrastructure

The POWER8 processor supports a hierarchical solution to power management. An entity running on an attached service processor, with management software running on the cores, can establish power budgets. Then, the POWER8 processor can be required to stay within the budget. Hypervisor-based energy management algorithms for the partition and micropartition level can then influence the power/performance trade-off in conjunction with the on-chip controller (OCC), which deals with the processor and memory level. Features are handled at the lowest level possible to allow the greatest flexibility and to reduce overall complexity of the hardware design. In general, power management hooks exist inside the processor core itself, inside the processor core chiplet (the asynchronous entity which includes the core’s L2 and L3 caches), in the chip-level nest unit level, and at the chip level. This affects how the features are implemented and therefore, laid out in the SCOM registers. For example, voltage is controlled at the chiplet level (as well as the chip), frequency and power gating in each core chiplet, and software-directed modes and instruction throttling controls inside the core itself.



9.3 Power Management Policies and Modes of Operation

The POWER8 interface must be simple but powerful to observe the operating power of the POWER8 system and its subcomponents and to direct overall policies and modes of operation for the system. The POWER8 power management subsystem (EnergyScale hardware and firmware) must provide clear policies that can be customized to achieve the desired level of power and performance efficiency, within bounds specified by the customer. The POWER8 processor supports multiple power management choices for the system operation. They can be selected by the customer, depending on the situation at any given time or for particular data center constraints.

9.3.1 Maximum Power Savings Based on Utilization and Idle

With this policy,¹ decisions are made by the EnergyScale and hypervisor firmware to take advantage of predictable idle periods of low-processor utilization. No performance (throughput) loss is allowed during long periods of invariant utilization percentage in the observable seconds or minutes timescale. This mode is tailored to satisfy the SPEC-Power benchmark, various upcoming government regulations, and some predictable data center usage models. In this policy, power is not reduced during periods of near 100% utilization, because full performance is maintained.

9.3.2 Adaptive Power Savings with Performance Loss Floor

This policy attempts to save as much power as possible with a maximum defined allowable performance loss. This policy is useful for system workloads or periods of time with high or unpredictable utilization. An attempt is made by the EnergyScale firmware and hardware to save as much power as possible while maintaining at least a minimal performance level. In more advanced forms, a power shifting technique can be used to maximize performance by giving more power to busier cores by taking power from less busy cores.

9.3.3 Power Cap

For this policy, the EnergyScale firmware allows the POWER8 processor to perform as well as possible underneath a hard power ceiling (to within milliwatts of accuracy). Frequency, voltage, and throttling are used to limit the throughput of the system to prevent the Power Cap from being exceeded. This policy allows data centers to limit the total power consumed by the server to budget their electric bill or work within physical constraints (power delivery, cooling capacity) of their data center

9.3.4 Turbo Performance Boost

This policy allows EnergyScale to boost frequency by managing voltage and internal modes assuming lower than peak workload. Chip health (temperature, circuit performance, and so on) is monitored by internal sensors to ensure safety, such that the chip never exceeds its power or thermal specification. By using activity event counters as an estimation for power being consumed at various periods during runtime, a given code sequence and mode selection is deterministic on every part in every environment. As a result, the execution time of a given piece of code is largely invariant between multiple runs, regardless of environmental

1. Any of these three policies can be run concurrently. Under these policies, a given code sequence can run at different frequencies (including overclocked "Turbo" when possible) and operating modes depending on the part, environmental conditions, workload, and so on. The system manages voltage, frequency, and internal modes as a function of customer policy and internal sensors ([DTS](#) and performance counters).

conditions, a part's ability to overclock, and so on. Under this policy, EnergyScale firmware can change frequency or modes dynamically during runtime, but this change occurs the same way every time code is run and on every part in that sort bin.

9.4 Feature Summary

The POWER8 processor has additional enhanced high-value power management features as follows:

- Enablement of micropartition energy management through SPR-based PStates controlled by the hypervisor.
- Autonomous, real-time management through firmware driven, on-chip controller (OCC).
- Per core chiplet voltage regulation to allow each core to be independently set based on the currently running workload. Bounds are managed by the OCC.
- Per-thread accounting (instruction completion, work rate, memory counting) to enable advanced utilization-based power management algorithms.
- Memory access threshold throttling to control both partition-level power management and power shifting between processors and memory, the two largest components of power in the CEC.

9.5 Overview of Chip Hardware Power-Management Features

9.5.1 Communication Paths for System Controllers

- Dedicated special wake-up bit per core chiplet owned by the OCC
- Hypervisor and OCC communication. Dedicated indirect OCB channel for hypervisor queue (up to 64 bytes) in OCC SRAM.

9.5.2 Sensors

- Digital thermal sensor (DTS)
 - Diode bandgap design.
 - Analog/digital converter built into hardware to provide a digital readout.
 - Available via SCOM during runtime; used by the OCC to safely implement turbo modes and to protect the chip from environmental changes that could lead to overheating.
 - Three implemented per core; one implemented per EX cache region.
 - Automated hardware thermal overtemperature protection is not supported. The real-time OCC firmware accomplishes this function.
- Dedicated performance, microarchitecture, and activity/event counters
 - Used for processor and memory utilization and weighted power activity proxy measurement to direct power and performance trade-off decisions and selection of appropriate power management techniques.
 - Not shared with the performance monitor unit (PMU).
 - Events and thresholds are also routed to the PMU and HTM for power and performance analysis using traditional performance techniques.



- Per-thread run cycle, instruction dispatch, instruction completing, work rate, and programmable memory hierarchy counters for per-thread utilization accounting.
- Chiplet memory access counter and throttling for per-core-chiplet memory power allocation and shifting.
- Processor bus overcommit retry counter access.

9.5.3 Accelerators

- On-chip controller (OCC)
 - Embedded PowerPC 405 with 16 KB instruction cache and 16 KB data cache
 - On-chip SRAM tank (512 KB)
 - Access to system DRAM memory via the processor bus for instruction and data area overflow (firmware managed)
- SPIVID command interface to chip external VRMs controlling core voltage rails. SPIVID VRMs accept respective core V_{DD} and V_{CS} target values.

9.5.4 Actuators/Controls

- Architected idle modes (nap, sleep, winkle)
 - Hypervisor can execute idle instructions to quiesce the processor pipeline and allow varying levels of power savings, each with higher latency to enter or exit.
 - No support for doze.
 - Support for sleep instruction (fast sleep: core power on; deep sleep: core power off). The sleep instruction allows for the L3 cache to remain functional for use by other cores.
 - Support for the **rvwinkle** instruction (fast winkle: chiplet power to V_{RET} ; deep winkle: chiplet power off).
- Per-chiplet core frequency control.
 - Digital PLL (DPLL) allows desired dynamic frequency range of -50% to +10%.
 - With winkle of other cores, can be up to +30% of nominal core frequency.
 - Automated frequency reduction nap, sleep, winkle and low activity detect – controlled by PMICR SPR settings.
- External (off-chip) VRM voltage control.
 - Three SPI VID interfaces to VRMs are associated with a given chip.
 - VRMs are passed respective core V_{DD} and V_{CS} targets with CRC to ensure transmission. VRMs, using load-line sensing, ramp the rails without offset overlap.
 - VRMs return a status frame for command confirmation of VID write validity.
 - Redundancy supported via a broadcast of VID commands on all configured interfaces.
 - If two interfaces, good if one responds positively (acknowledged and no errors).
 - If three interfaces, good if two respond positively (acknowledged and no errors).
 - If responses are negative (timeouts, acknowledged with errors), voltage changing is suspended. OCC notification for firmware handling.
- Memory controller (DIMM) throttling. See *POWER8 Memory Buffer User's Manual* for full support details.
- Firmware runtime power-management controls.

- Power Management Control Register (PMCR) [SPR](#) for run-time Pstate control. (See *Section 9.8 Architected Control Facilities* on page 192 for details.)
 - Pstates: 8-bit signed values that are an offset from the part's nominal frequency.
 - Global: requests the global voltage rail to be changed.
 - Local: requests the local frequency and voltage to be changed.
- Power Management Idle Control Register (PMICR) SPR for idle-time Pstate control.
 - Nap Pstate
 - Sleep Pstate
 - Winkle Pstate
- Power Management Memory Access Register (PMMAR) SPR for chiplet memory traffic control.
 - Chiplet memory traffic accounting.
 - Chiplet memory traffic throttling (if in single-core chiplet partition mode).

9.5.4.1 Configurations with Unused Components

- Dynamically or statically disable active power in the I/O devices for unused buses.
 - Static disablement of unused buses available in each EI4 bus using SCOM.
 - EI4 buses support hardware-driven dynamic spare lane power down.
- Dynamically clock-gating unused units.
 - The following units can be clock gated (via thold) when the associated chip is not attached to the interface.
 - PCI-E Host Bridges (PHB)
 - SMP interconnect A0:A2 Elastic Differential I/O (EDI) buses individually
- Statically clock-off unused units/buses
 - Each MCS0 or MCS1 unit and associated differential memory interface (DMI) buses.
 - Memory control buffer (MCB) for direct drive memory
 - Elastic I/O - four buses individually.
- Partial good requirements
 - Both partial good “bad” scenarios and gard/runtime deallocation
 - Bad core chiplets (with a bad L3 cache region) can be completely power gated at IPL time (chiplet power same as winkle).



9.6 Chip Hardware Power-Management Features

9.6.1 Chiplet Voltage Control

The POWER8 processor supports several VRM control mechanisms to support multiple different system configurations. The core chiplets are on a separate voltage plane than the other Nest components of the chip. The chip-level power management control (PMC) macro and the OCC are, in combination, programmable to support these configurations.

Core chiplets all share the same voltage plane and have to run with “highest common denominator” (that is, the core demanding the highest voltage sets the value of the voltage rail). The OCC is responsible for establishing the best frequency and therefore voltage bounds based on the workload running, the power/performance efficiency policy selected by the customer, and the system budgets established by the thermal management component.

9.6.2 Chip-Level Voltage Control Sequencing

The external VRMs (eVRMs) sourcing the logic (V_{DD}) and array (V_{CS}) rails are controlled by voltage control interfaces from the POWER8 chip. These interfaces use serial peripheral interconnect (SPI) signaling to a VRM chipset that converts the addressed VID command to industry standard Intel VRM-11 interface components or others as implemented by the VRMs.

9.6.2.1 SPIVID VRM Control Sequencing

The V_{DD} and V_{CS} target voltages are sent in one command to the VRM set. That target can be the full voltage swing request (V_{MAX} to V_{MIN} or vice versa) or any subset. With the V_{DD} and V_{CS} targets, the VRMs, through sampling of load lines, manage the offset of the two rails during the slew.

9.7 Functional Description of Processor Core Chiplet

9.7.1 Power Gating

The POWER8 processor implements per-core and per-core-chiplet power gating with two separately controlled power grid regions inside the EX core chiplet. These power regions also correspond to clocking domains. Thus all circuits on the full-frequency core clock grid (that is, processor core and L2 cache) are controlled by the core power gate. The EX cache power gate controls the voltage to the circuits on the half-speed cache_nclk region of the chiplet (for example, L3 cache, NCU, and processor bus interface units). This allows the entire core chiplet to be powered off to save maximum power when not in use (via the winkle state or static EX cache mode). It also allows only the core region to be powered off via the sleep state, such that the L3 cache remains available for other cores on the chip to use as an L3.1 extended cache for lateral castouts to maximize their performance.

The benefits of POWER8 power gating include:

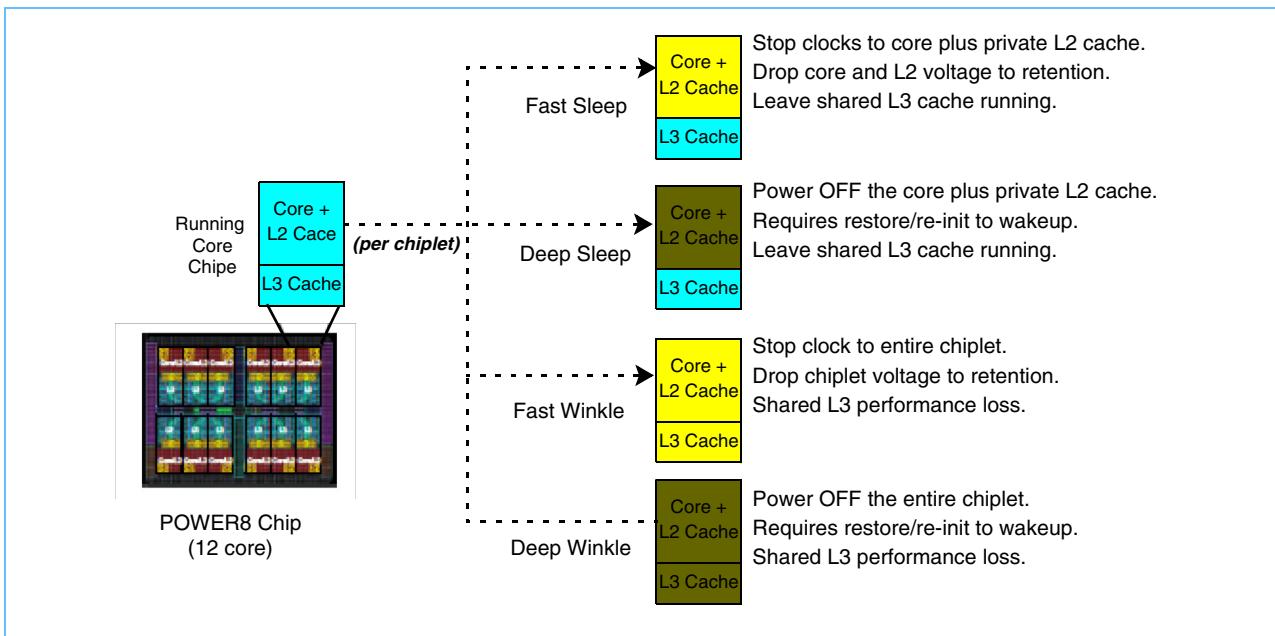
- Dynamic EX cache mode can be a runtime option via sleep mode.
- The adjacent core can enter turbo mode to take advantage of overall lower chip power and temperature.
- Idle power is significantly reduced.

- Per-chiplet control enables better partition-level power management capability (due to per-chiplet control).
- Better power efficiency modes are enabled.
- Better energy management capability and benchmark scores.

9.7.2 Idle States

When a section of a unit or macro detects that it is idle, it gates clocks to the unneeded latches to save active power. This is done extensively in both the core and the nest which has the effect of reducing peak power because all circuits cannot be kept busy 100% of the time. However, when the hypervisor decides that the entire processor core can be idle, it can execute one of the architected idle state instructions. These instructions stop the fetch and dispatch of instructions, quiescing the core. Depending on the idle instruction executed, hardware can further reduce voltage and clock grid frequency or even power the core completely off to zero volts to save leakage power. *Figure 9-1* shows the effect of each of these idle modes on the processor core chiplet power.

Figure 9-1. Idle Mode Summary



A single thread entering any architected idle state causes that thread to stop dispatching and fetching instructions.

When a thread wakes up from an idle state, it takes an SRESET interrupt to restart program execution. The SRR0 contains the address of the idle instruction that caused the thread to go into idle state. (Note: The architecture says that SRR0 content is undefined.) SRR1 contains the reason for the wake-up and amount of state that was lost due to being in that idle mode.



All threads on a core must be in that architected idle state or deeper¹ for that core to enter an architected idle mode. The core enters the least aggressive idle mode of the eight threads. For example, if a core has one thread in nap, two in sleep, and one in winkle, the Core enters nap mode.

Note: Modes exist to disable chiplet clock/voltage off for lab modes such as hardware trace (performance and debug analysis).

The POWER8 processor supports three of the four architected power-save or idle states defined in the Power ISA (nap, sleep, and winkle). There is an important distinction between thread idle state and core idle state. When all threads on a core are in an architected idle state, more aggressive power-management techniques can be engaged at the core, core chiplet, or even chip level.

Table 9-1. Supported Chiplet Power Management Modes

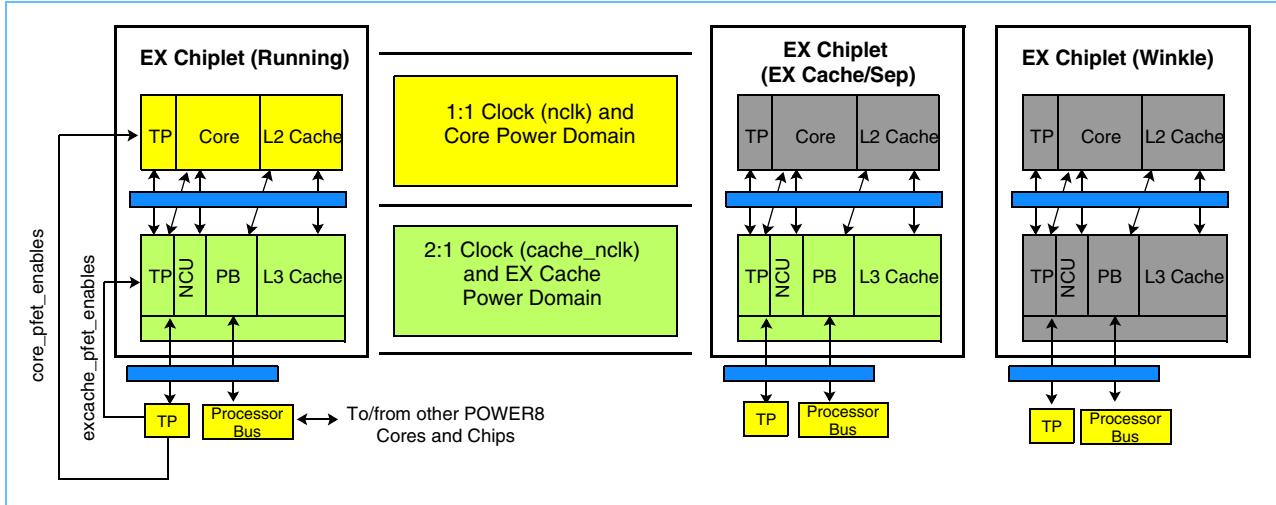
Mode/State	Core Domain Voltage	EX Cache Domain Voltage	Core Domain Mesh	EX Cache Domain Mesh	Core Domain THolds	EX Cache Domain THolds ¹	Estimated Relative Power ²	Estimated Entry and Exit Latency
Running	V _{FUNC} ³	V _{FUNC} ³	Running	Running	Running	Running	N/A	N/A
Nap	V _{FUNC} ³		Running ⁴		Partially stopped	Running	90%	Typical: ~ 2 µs Worst case: < 50 µs
Fast Sleep	V _{ON}	V _{FUNC} ³	RefClk	Running ⁴	Stopped	Running	60%	Expected typical < 1 ms
Deep Sleep	V _{OFF} ⁵	V _{FUNC} ³	Off	Running ⁴	Stopped	Running	25%	Expected typical < 5 ms
Fast Winkle	V _{ON}		RefClk		Stopped	Stopped	45%	N/A
Deep Winkle	V _{OFF} ⁵	V _{OFF} ⁵	Off	Off	Stopped	Stopped	Effectively zero (~0.1 W)	Expected typical < 20 ms

1. Full L3.1 performance.
2. Compared to operating system idle loop.
3. If enabled, these states allow for internal voltage regulation.
4. These states have an optional frequency reduction available.
5. This state requires POR assistance to restore lost state.
6. Power save mode transition rules: There are no transitions between power management idle states. Entry and exit from all architected idle states happens to and from a running state.

1. Deeper is defined as an idle state that is architecturally defined to save more power; that is, winkle is deeper than sleep which is deeper than nap.

With the power-gating capabilities implemented, the sleep and winkle idle states provide a progression as shown in *Figure 9-2*.

Figure 9-2. Sleep and Winkle Power Gating Progression



9.7.2.1 Core and Thread Doze

Core and thread doze are not supported. The hypervisor does not use doze mode.

9.7.2.2 Single Thread Nap, Sleep, and Winkle

Thread state can be lost when a single thread enters nap, sleep, or winkle mode, because it can cause the core to perform an SMT thread switch for performance gain on the remaining non-idle threads.

The only state lost in thread nap mode is the thread-specific state in the case that SMT mode changed. All non-thread-specific states including timebase registers and hypervisor states are preserved. Some thread-specific hypervisor states are preserved in nap.

When an SMT thread switch is enabled, the napping thread's resources can be given to other threads to improve core performance. Software must restore the architected state of the dormant thread upon exiting nap mode except for the following architected facilities, which are preserved by the hardware:

- SLB State
- All Hypervisor Special Registers (includes PURR, SPURR, AMOR, UAMOR, AMR)
- DEC
- SPRG0-3
- DAR
- DSISR
- DABR, DABRX
- DSCR
- All Performance Monitoring Special Registers (PMCs, SIAR, SDAR)



9.7.2.3 Sleep

Generally, the sleep instruction removes the core from operation while leaving the associated L3 cache for use by other core chiplets. The amount of power savings and the inverse amount of exit latency is controlled by whether the core power gating is used or not.

Upon entering sleep, hardware invalidates the coherency caches (L1, ERAT, TLB, and so on) in the core and purges cache/state in the L2 units. The L3 cache is fenced from the core/L2 and retained as operational in lateral castout (LCO) mode, so that it can be used by other operational core chiplets.

Upon fencing the L3 portion (referred to as the extended cache option), the core and L2 cache are either:

- Powered off using power gating core FETs within the chiplet (defined as deep sleep).
- Have all clocks in the core stopped and the core clock grid dropped to the refclk frequency. The voltage is lowered to a $V_{RETENTION}$ level for low-latency exit (defined as fast sleep).

Logically, all state of the processor core domain is lost, including hypervisor state.

Before entering sleep, the hypervisor must update the power-on memory image with the state of any registers that must be set after wake-up. This list includes: HRMOR, HSPRG0, LPCR, HID0, HID1, HID5, along with meaningful Timer Facility PURR/SPURR/TB values. In addition, the hypervisor can set values into the PMICR for the following fields:

sleep_pstate_req	Indicates the Pstate that the EX cache portion of the chiplet is requested to run.
sleep_pstate_en	Enables as valid the sleep_pstate_req field.
sleep_global_en	Indicates if the sleep_pstate_req is to be treated as global.

Because the threads are waking up in real mode, the hypervisor guarantees that the first thread to wake-up restores SDR1 and LPIDR before the others leave real mode.

Any runtime changes to the state of the processor core and L2 cache by the off-chip firmware components (OCC and FSP) must be done with special wake-up asserted and a simultaneous update to the power-on memory image containing the necessary SCOM operations to restore it.

SCOM (and PCB) accesses to a core in sleep are no longer possible as the power to the core may be shut off (deep sleep case); for simplicity, fast sleep is treated the same. To access the core for these operations, a special wake-up action must take place. See *Section 9.7.3 Special Wake-up* on page 188.

Note: The hypervisor is required to set LPCR[PECE] = '101' on nonguarded cores (otherwise the core hangs when an interrupt occurs).

Exiting sleep only occurs because either a malfunction alert (caused by some other core checkstop) or external interrupt targeting a thread in the core. In the case of a deep sleep exit, a mini-power-on reset is performed on the core by a chip level power-on reset (POR) to bring the core back online (scan, ABIST, SCOM, and so on.). In the case of a fast sleep exit, voltage and frequency are restored and fencing is dropped to allow processing of the waking condition.

In either case, after wake-up, the hypervisor must resynchronize the Timebase facility.



9.7.2.4 Nap

In this state, clocks to a majority of the processor core is turned off while preserving coherency (L1 and L2 caches, ERATs, TLB, and SLB). The clock rate to the chiplet is dropped to lowest frequency that the current processor bus setting allows. This allows for low-latency entry and exit (micro-seconds). Generally, the entire IFU and LSU remain running while allowing for gating of non-snoop portions. A portion of pervasive remains running (for example, sensors, Timebase, and interrupt processing).

Nap Frequency Change

Based on the setting of the PMICR, the execution of the nap instruction will drop the chiplet to reduced frequency (nap Pstate). Upon wake-up, latency of a few microseconds will be incurred for the DPLL to ramp back to full frequency.

9.7.2.5 Winkle

Generally, the **rvwinkle** instruction removes the core and associated L3 EX cache from operation. The amount of power savings and the inverse amount of exit latency is controlled by whether the core and EX cache power gating is used or not.

The winkle idle instruction causes power to be gated off to the entire chiplet (core/L2/L3) and cause the chiplet to be disconnected from the processor bus. Like Sleep, all state in the processor core domain is considered lost, including Hypervisor state. However, unlike Sleep, any Hypervisor and Firmware state in the remainder of the chiplet is also lost.

Upon disconnection from the processor bus, the core and chiplet are either:

- Powered off using power gating FETs within the core and chiplet (defined as deep winkle)
- Have all clocks in core and EX cache stopped and the core and chiplet clock grid dropped to the refclk frequency (defined as fast winkle)

Before entering winkle, the hypervisor must update the power-on memory image with the state of any registers that must be set after wake-up. This list includes: HRMOR, HSPRG0, LPCR, HID0, HID1, HID5, along with “meaningful” Timer Facility PURR/SPURR/TB values. In addition, the hypervisor can set values into the PMICR for the following fields:

winkle_pstate_req	Indicates the Pstate that might be requested globally (to lower voltage); locally, this has no effect.
winkle_pstate_en	Enables as valid the winkle_pstate_req field.
winkle_global_en	Indicates if the winkle_pstate_req is to treated as global.

Because the threads are waking up in real mode, the hypervisor guarantees that the first thread to wake up restores the SDR1 and LPIDR Registers before the others leave real mode.

Any runtime changes to the state of the processor core, L2 cache, and L3 cache by the off-chip firmware components (FSP) must be done with a special wake-up asserted and a simultaneous update to the power-on memory image containing the necessary SCOM operations to restore it.

SCOM (and PCB) accesses to a core in winkle are no longer possible, because the power to the core chiplet has been shut off. To access the core for these operations, a special wake-up action must take place. See *Section 9.7.3 Special Wake-up* on page 188.



Exiting winkle only occurs due to either a malfunction alert (caused by some other core checkstop) or an external interrupt targeting a thread in the core. In the case of a deep winkle exit, a mini-power-on reset is performed on the core by a chip-level power-on reset to bring the core and EX cache back online (scan, ABIST, SCOM, and so on). In the case of a fast winkle exit, voltage and frequency are restored and fencing is dropped to allow processing of the waking condition.

9.7.3 Special Wake-up

Normal wake-up from architected idle states resumes instruction execution in the processor core, beginning at the SRESET vector, due to an interrupt condition being present. There are times when various components of firmware need to access facilities inside the chiplet without resuming instructions. A special wake-up mode is provided to each firmware component (FSP, OCC, and hypervisor). It suppresses new architected idle states from being entered by the chiplet. If already in an idle state, it causes, if necessary, restoration of power and state first. This allows the clocks to resume to the chiplet for the access to be completed.

When special wake-up is cleared, one of two things happen:

- If a wake-up condition exists or if the processor core is no longer in an architected idle state because a wake-up interrupt already occurred, nothing happens.
- If the core is still in an idle state and no wake-up condition exists, the chiplet re-enters the appropriate idle state that was previously requested by the core.

While special wake-up is asserted, the core chiplet runs at the frequency represented by the idle Pstate (sleep Pstate or winkle Pstate) as defined in the PMICR unless bounded by the settings of PMax and PMin as established by the OCC. Low-activity detect is disabled from changing state when architected idle state and special wake-up are both active.

9.7.4 Pstates

The POWER8 processor implements the Pstate architecture with a local Pstate table of 128 logical entries.

9.7.4.1 Architectural Overview

Hypervisor firmware and operating system software control the scheduling of work on the processor core threads. It is necessary to abstract out the low-level details of frequency and voltage settings into a facility called a Pstate. This abstraction provides an isolation layer from the hardware, because the combination of voltage and frequency can vary per chip based on manufacturing variation and technology while allowing control of performance and power efficiency. Performance and power efficiency have a linear relationship to frequency and a quadratic relationship to voltage. As long as this linear and quadratic relationship to Pstates is understood, the energy management components in the hypervisor can appropriately manage both performance and power efficiency.

The Pstate itself does not logically alter the behavior of the core in terms of fetch, dispatch, and execution rates of instructions. Rather, it is an indication of the performance level represented in terms of a frequency target that the platform (processor chip and underlying control firmware) attempts to fulfill within the constraints of available power and thermal capacities while also taking the opportunity to minimize the same.

The Pstate architecture is based upon a hierarchical view. Some elements of a request can be handled in a manner local to the core chiplet making the request. A primary example of a local operation is frequency change. Other elements must be sent out of the core chiplet to the central, or global, entity which might have

to arbitrate between multiple core chiplet requests to take action in response. A primary example of a global operation is where voltage rails are shared among core chiplets and coordination is required to slew such a rail.

Underlying the Pstate request notion, as made by energy management components with the hypervisor, is a platform set of elements that manage the allowable power (and, hence, performance) that is expended by elements in the machine. In the case of Pstates, the element is the core chiplet. Thus, limits are placed on the maximum allowable frequency and voltage based on budgetary thresholds as well as, on minimum allowable frequencies that are necessary to maintain machine coherency. This is why the Pstates are considered requests in that the platform hardware and firmware can limit or bound the value requested. Therefore, mechanisms are provided to allow the hypervisor components to read the actual operational values so as to react, if necessary, to any such clipping actions.

9.7.4.2 Definitions

Fnom	<p>Fnom is the nominal frequency established after part characterization that allows the part to fully run all supported code streams within the power and thermal envelope provided by the system. The value of Fnom is, thus, both part and system dependent.</p> <p>The Fnom value is established by platform-specific firmware that has access to both information about the part instance (typically held in a vital product data (VPD) record) as well as the present system characteristics. This value is written into core chiplet and global chip hardware facilities as a representation of a chip-specific frequency value. It is then used as the basis from which all Pstate calculations (global and local) are made.</p>
Pstate Number (PSN)	A Pstate number is a signed value that is an offset from the chip Fnom frequency value. PSNs range from +127 to -127.
Pmax	<p>Pmax is the maximum PSN allowed for the core chiplet and is established by platform-specific firmware. It takes into account a number of factors including, but not limited to:</p> <ul style="list-style-type: none">• The power delivery into the socket containing this core chiplet• The number of core chiplets implemented in this socket• The thermal environment of the socket• The policy established by the customer (maximum performance, maximum power savings, deterministic performance, and so on)
Pmin	<p>Hardware ensures that no mechanism, software in the form of a Pstate, or hardware in the form of power shifting, or guardband management mechanisms exceeds this value.</p> <p>Pmin is the minimum PSN allowed for the core chiplet and is established by platform-specific firmware. It takes into account a number of factors including, but not limited to:</p> <ul style="list-style-type: none">• The cache coherency of the SMP machine, because timely response to cache snooping actions is typically required• Any applicable performance floors that customer-driven policy can establish



Global Pstate (GPS)

The Global Pstate represents a signed value that is an offset from the chip Fnomin value. It indicates that a frequency change is being requested that requires coordination by a chip-level entity (for example, outside the core chiplet) to manage any elements across other dependent core chiplets. Typically, the element that needs to be managed is the voltage setting. Some examples of a voltage change action are as follows:

- The voltage rail supplying this core chiplet is shared by other core chiplets. In this case, a “vote” about what the voltage should be must take place. If warranted, a voltage change action is completed.
- The voltage rail supplying this core chiplet is segregated from other core chiplets but the control means for changing the voltage setting is external to the core chiplet itself. In this case, the central entity completes a voltage change action on the core chiplet’s behalf.

In the case of Global Pstate increase, any voltage change that needs to occur must be completed and the core chiplet notified of this fact before the frequency change can take place.

In the case of Global Pstate decrease, the core chiplet requests the decrease to the central entity, but might or might not lower its present local operating Pstate. This choice on the local Pstate is a matter of policy. If the request was honored (that is, it passed the rail voltting), the central entity notifies all chiplets of the reduction as a new Global Actual Pstate. It waits for acknowledgements that the new voltage is honored and then reduces the voltage.

Local Pstate (LPS)

The Local Pstate represents a signed value that is an offset from the chip Fnomin value. It indicates that a frequency change is being requested that does not require coordination by a chip-level entity; it is handled locally within the core chiplet. For implementations that do not have any means of varying voltage locally, this becomes a request to immediately move frequency within the bounds established by the current Global Pstate or FMax (whichever is lower) on the upper side and FMin on the lower side. Upon change of the frequency, the new current value is reflected in the Actual Local Pstate field.

Because any operational changes are local, no communication to the central element is performed. Implementations can (and probably will) make the Actual Local Pstate accessible via the service network so that the values can be read as necessary.

Nap Pstate (NPS)

The Nap Pstate is the PSN that indicates the frequency (and, optionally, voltage) change to request upon the execution of the nap instruction.

Sleep Pstate (SPS)

The Sleep Pstate is the PSN that indicates the frequency (and, optionally, voltage) change to request upon the execution of the sleep instruction.

Winkle Pstate (WPS)

The Winkle Pstate is the PSN that indicates the frequency (and, optionally, voltage) change to request upon the execution of the **rwwinkle** instruction.

9.7.4.3 Permissible Behavior

- Global Pstate requests are made to the central area upon any PMCR store. There is no interlock defined that ensures that a previous request was, in fact, honored.
- The core chiplet can run at a higher local Pstate than presented in a Global Pstate Request because the global actual Pstate is, at all times, honored.

While it is generally viewed that Energy Management algorithms (in the hypervisor) make the local Pstate less than or equal to the global Pstate request (to save energy), the hardware does not enforce this limitation.

9.7.4.4 Interaction with Idle Modes

The sleep and winkle high, medium, and low latency controls are intended to allow platforms the option of leveraging power gating (or similar technologies), because these have a major influence on restoration latency versus power expended.

The winkle Pstate can be used in conjunction with the latency controls where power gating cannot be invoked; therefore, a Pstate is used to communicate to the platform the intended operating point.

Architecture Note

An intended usage model is:

- Hypervisor energy management algorithms establish the nap Pstate (NPS), sleep Pstate (SPS), and winkle Pstate (WPS) fields and their respective global and local settings.
- Eventually, all threads will achieve nap (inclusive of some in sleep and some in winkle), sleep (inclusive of some in winkle) or winkle state
 - For the case of nap, the NPS is used to affect frequency and voltage per the flow respective of the global/local bit.
 - For the case of sleep, the SPS is used to affect frequency and voltage per the flow respective of the Sleep Global Enable bit.
 - For the case of winkle, the WPS is used to affect frequency and voltage per the flow respective of the Winkle Global Enable bit.



9.7.5 Resonant Clocking

Resonant Clocking is a mode of the clock distribution whereby on-chip inductors are used in conjunction with the natural clock mesh capacitance to form a resonant structure to save power drawn by the clock mesh. This mode can only be enabled when the frequency (Pstate) is within one of two specific bands. These bands are defined by SCOM registers associated within each chiplet and are setup at POR. After the transition facility is enabled, any frequency changes that exit these defined bands cause the resonant clocking mode to be disabled before the transition. Then, after the new frequency is achieved, if that new frequency is within a defined resonant band, the resonant clocking circuits are re-enabled.

Resonant clocking is modally supported.

9.8 Architected Control Facilities

9.8.1 Power Management Control Register (PMCR)

The PMCR is the mechanism used by the hypervisor firmware to request Pstate changes. This register has only one instance for the core.

Table 9-2. Power Management Control Register (PMCR) - SPR 884 (Sheet 1 of 2)

Field	Field Name	Access	Function
0:7	global_pstate_req	R/W	<p>Global Pstate request</p> <p>Writes to this field initiate a coordination action with any available central element that arbitrates between other cores that might share a power rail with this core.</p> <p>Reads from the field return the value last written.</p> <p>The value is an 8-bit signed integer representing an offset from $F_{nominal}$. Legal values are +127 to -127 with the value increment being platform dependent¹.</p>
8:15	local_pstate_req	R/W	<p>Local Pstate request</p> <p>Writes to this field initiate a local performance change request without explicitly orchestrating with any available central element.</p> <p>Reads from the field return the value last written.</p> <p>The value is an 8-bit signed integer representing an offset from $F_{nominal}$. Legal values are +127 to -127 with the value increment being platform dependent¹.</p>
16	Reserved		Reserved

1. For the POWER8 processor, the increment is the reference clock frequency divided by the DPPL divider.



Table 9-2. Power Management Control Register (PMCR) - SPR 884 (Sheet 2 of 2)

Field	Field Name	Access	Function
17:23	spurr_fraction	R/W	<p>SPURR fraction Modifies utilization/accounting of the SPURR for the requested Local Pstate. This accounting only occurs if the requested Pstate is honored. The default value is '1000000'. It is multiplied by other SPURR factors to create a composite discount value. Reads from the field return the value last written. The value is an 8-bit signed integer representing an offset from $F_{nominal}$. Legal values are +127 to -127 with the value increment being platform dependent¹.</p>
24:51	Reserved		Reserved
52:63	lpar_id	R/W	<p>LPAR ID The value representing the logical partition ID that will be executing with the settings defined by the other fields. Writes to this field provide the LPAR ID for the platform to associate the power for the subsequent execution window (for example, until the next PMCR store). The actual accounting action is controlled by the power_acctng_change field. Reads from the field return the value last written.</p>

1. For the POWER8 processor, the increment is the reference clock frequency divided by the DPPL divider.



9.8.2 Power Management Idle Control Register (PMICR)

The PMICR controls the Pstates used under idle modes.

This register has only one instance for the core.

Table 9-3. Power Management Idle Control Register (PMICR)- SPR 852 (Sheet 1 of 2)

Field	Field Name	Access	Function
0:7	nap_pstate_req	R/W	<p>Nap Pstate (NPS) request Writes to this field initiate a coordination action with any available central element that will arbitrate between other cores that can share a power rail with this core. Reads from the field return the value last written. The value is an 8-bit signed integer representing an offset from $F_{nominal}$. Legal values are +127 to -127 with the value increment being platform dependent¹</p>
8	nap_pstate_en	R/W	<p>Nap Pstate enable 1 Enable the Nap Pstate (NPS) request, Nap Global Enable, and Nap Latency functions. 0 Disables the Nap Pstate (NPS) Request, Nap Global Enable, and Nap Latency functions. Reads from the field return the value last written.</p>
9	nap_global_en	R/W	<p>Nap global enable 1 Upon the execution of a Nap instruction, the Nap Pstate (NPS) request is sent to the central element as a Global Pstate Request. 0 Upon the execution of a Nap instruction, the Nap Pstate (NPS) request is sent to the local element as a Local Pstate Request. Reads from the field return the value last written.</p>
10:15	Reserved		Reserved
16:23	sleep_pstate_req	R/W	<p>Sleep Pstate (NPS) request Writes to this field initiate a coordination action with any available central element that arbitrates between other cores that can share a power rail with this core. Reads from the field return the value last written. The value is an 8-bit signed integer representing an offset from $F_{nominal}$. Legal values are +127 to -127 with the value increment being platform dependent.¹</p>
24	sleep_pstate_en	R/W	<p>Sleep Pstate enable 1 Enables the sleep Pstate (NPS) request, sleep global enable, and sleep latency functions. 0 Disables the sleep Pstate (NPS) request, sleep global enable, and sleep latency functions. Reads from the field return the value last written.</p>
25	sleep_global_en	R/W	<p>Sleep global enable 1 Upon the execution of a sleep instruction , the sleep Pstate (NPS) request is sent to the central element as a global Pstate request. 0 Upon the execution of a sleep instruction , the sleep Pstate (NPS) request is sent to the local element as a Local Pstate request. Reads from the field return the value last written.</p>
26:31	Reserved		Reserved
32:39	winkle_pstate_req	R/W	<p>Winkle Pstate (NPS) request Writes to this field initiate a coordination action with any available central element that arbitrates between other cores that can share a power rail with this core. Reads from the field return the value last written. The value is an 8-bit signed integer representing an offset from $F_{nominal}$. Legal values are +127 to -127 with the value increment being platform dependent.¹</p>

1. For POWER8, the increment is the reference clock frequency divided by the DPPL divider.



Table 9-3. Power Management Idle Control Register (PMICR)- SPR 852 (Sheet 2 of 2)

Field	Field Name	Access	Function
40	winkle_pstate_en	R/W	<p>Winkle Pstate enable</p> <p>1 Enable the winkle Pstate (NPS) request, winkle global enable, and winkle latency functions.</p> <p>0 Disables the winkle Pstate (NPS) request, winkle global enable, and winkle latency functions.</p> <p>Reads from the field return the value last written.</p>
41	winkle_global_en	R/W	<p>Winkle global enable</p> <p>1 Upon the execution of the rwinkle instruction, the winkle Pstate (NPS) request is sent to the central element as a global Pstate request.</p> <p>0 Upon the execution of the rwinkle instruction, the winkle Pstate (NPS) request is sent to the local element as a local Pstate request.</p> <p>Reads from the field return the value last written.</p>
42:43	winkle_latency	R/W	<p>Winkle latency</p> <p>00 Ceases instructions and honors Pstate change but does not perform additional state changing actions.</p> <p>01 Indicates to the platform that a sub-state that has the lowest latency is enabled.</p> <p>10 Indicates to the platform that a sub-state that might have a medium exit latency is enabled.</p> <p>11 Indicates to the platform that a sub-state that might have higher exit latency is enabled.</p> <p>It is platform-specific as to the differentiation between low, medium, and high.</p>
42:63	Reserved		Reserved

1. For POWER8, the increment is the reference clock frequency divided by the DPPL divider.

9.8.3 Power Management Status Register (PMSR)

The PMSR provides access to the platform values in place. This register has only one instance for the core.

Table 9-4. Power Management Status Register (PMSR) - SPR 853

Field	Field Name	Access	Function
0:7	global_pstate_actual	RO	Actual Global Pstate Reads from this field return the presently established global Pstate value. The value is an 8-bit signed integer representing an offset from $F_{nominal}$. Legal values are +127 to -127 with the value increment being platform dependent.
8:15	local_pstate_actual	RO	Actual Local Pstate Reads from this field return the presently established local Pstate value. This value is always less than or equal to global_pstate_actual. The value is an 8-bit signed integer representing an offset from $F_{nominal}$. Legal values are +127 to -127 with the value increment being platform dependent.
16:23	pmin	RO	Pstate Minimum Reads from this field return the presently established minimum Pstate set by the platform. This value can change autonomously based on the current policy in place and the physical constraints of the platform. The value is an 8-bit signed integer representing an offset from $F_{nominal}$. Legal values are +127 to -127 with the value increment being platform dependent.
24:31	pmax	RO	Pstate Maximum Reads from this field return the presently established maximum Pstate set by the platform. This value can change autonomously based on the current policy in place and the physical constraints of the platform. Value is an 8 bit signed integer representing an offset from $F_{nominal}$. Legal values are +127 to -127 with the value increment being platform dependent
32:63	Reserved		Reserved.

9.8.4 Power Management Memory Activity Register (PMMAR)

The PMMAR controls the memory activity that is allowed to be produced by the partition. This register is replicated by partition if in multiple-partition mode.

Table 9-5. Power Management Memory Activity Register (PMMAR) - SPR 854

Field	Field Name	Access	Function
0:31	mem_op_limit	R/W	Memory operation limit Defines the number of memory operations allowed within the rolling credit window defined by mem_crdt_window. Once this limit is reached, the memory operations are stalled until the mem_crdt_window expires.
32:45	mem_crdt_window	R/W	Memory credit limit Defines the rolling credit window time for which memory operations are managed. x'0000' Disabled (no stalling) x'0001' 64 µs ... x'3FFF' 1.048 s
47:63	Reserved		Reserved





10. Performance Profile

This chapter provides information about the performance profile of the POWER8 processor chip. Due to the complex nature of the speculative, out-of-order execution core coupled with the multilevel storage hierarchy, it is important that people concerned with performance and performance optimization develop an insightful understanding about the operational nature of the machine.

10.1 Core

10.1.1 Level-1 Instruction Cache

The L1 I-cache is a 32 KB, 8-way set-associative cache of instructions allocated in 128-byte lines with 32-byte sector valids. The replacement algorithm is pseudo LRU. The I-cache is also kept coherent with the L2 cache (an **icbi** instruction can be treated as a NOP). The I-cache is used to feed instructions to the rest of the core at up to eight instructions per cycle as long as the addresses are presented each cycle and each address results in a cache hit.

The I-cache consists of three parts:

- Instruction cache
- I-cache directory (IDir)
- Effective address directory (EADIR)

The EADIR is used to predict the way select for a given I-cache access keeps the I-cache access time to one cycle. The EADIR is accessed using EA bits 52:56 and tagged using EA bits 41:51. All entries are initially assumed to be shared between threads, but certain combinations of reduced sharing are also possible (sharing is limited to thread pairs in Cloud mode). The basic I-cache miss latency assumes an EADIR miss. An EADIR hit that is later determined to be an I-cache miss adds approximately eight cycles to the base miss latency. The I-cache directory contains the real address (RA) of the lines in the I-cache, line and sector valids, and certain MSR state bits (PR, LE, and HV) from when the line was fetched. Both the EADIR and the IDir are addressed using EA bits 52:56.

Single-thread aliasing can occur when a given EA[41:56] cache line is required for more than one real address, or combination of MSR bits. Only one of the two combinations can be valid in the I-cache for a given thread at any one time, and an EADIR invalidate is required before fetching the other alias. Multi-thread EADIR aliasing results when two threads map the same EA(41:56) to two different real addresses or MSR combinations. The second address is brought in as private, after an EADIR invalidate. In 4 LPAR mode, only the two threads in the same LPAR can share I-cache lines.

The I-cache is banked and allows a concurrent read and write, if they go to different banks. On an I-cache miss and L2 hit, instructions are typically returned on the L2 cache interface in two consecutive 64-byte beats. If the instructions are coming from the L3 cache, four 32-byte sectors are typically expected to arrive in every other cycle. If the instructions are coming from the processor bus, the critical sector comes first and it is typical to expect the successive sectors to come every fifth processor cycle. Note that the I-cache miss latency is typically four cycles longer than the D-cache latency.

The I-cache is arranged in 128-byte lines of four 32-byte sectors. Each sector has its own valid. Within a line, the I-cache can be accessed on any 16-byte boundary to return 32 consecutive bytes. Addressing the last 16-byte boundary of a line returns only 16 bytes.



Data in the I-cache must also be in the L2 cache (inclusivity). This means that if the L2 cache removes a line that is thought to be in the I-cache, it must send an invalidate for that line to the I-cache. Invalidates are handled in parallel with read accesses so that normally they are invisible to performance. The LRU algorithm in the L2 cache requires data to be accessed to keep it from being replaced. This means that a line in the I-cache that is being heavily used, can be forced out of the L2 cache due to an L1 D-cache request. To avoid this situation, the processor periodically sends a dummy request to the L2 cache on an I-cache hit address. When bandwidth permits, these requests serve only to update the L2 LRU for the I-cache lines in hopes of preventing them from being aged out.

10.1.2 Level-1 Instruction ERAT

On instruction fetches, effective address bits are used to index into the I-cache, the directory and the effective-to-real-address-translation (ERAT) table. The ERAT is a fully-associative 64-entry table and contains both the effective addresses and the real addresses. For an ERAT hit, the effective address of the instruction must match the effective address contained in the ERAT entry being indexed, and the ERAT entry must be valid. In addition, the IR, US, HV, and PR bits from the MSR at the time of ERAT miss are stored in the ERAT when the ERAT is loaded on a ERAT miss. These bits must match the corresponding bits in the MSR at the time of instruction fetch for an ERAT hit. I-ERAT minimum miss penalty (assuming a TLB hit) is 18 cycles.

The I-ERAT directly supports 4 KB, 64 KB, and 16 MB page sizes. Other page sizes are stored in the next smaller supported page-size granule.

10.1.3 Instruction Prefetch

If a particular instruction fetch misses in the I-cache, a demand fetch reload request (possibly speculative) is sent to the L2 cache subsystem. The L2 cache processes this reload request with high priority and forwards it onto subsequent levels of cache or memory in the event that it misses in these caches.

In addition to these demand-oriented instruction fetching mechanisms, the POWER8 processor core also works to automatically prefetch instruction cache lines that might be referenced soon into its instruction cache. If there is an I-cache miss, it generates prefetch requests for the next one in both SMT2 and SMT4 mode or three sequential cache lines (in ST mode). In SMT8 mode, no instruction prefetching requests are generated. As these requests return cache lines, they are stored directly in the instruction cache. The banked cache design of the POWER8 instruction cache allows a concurrent read and write (as long as they go to different banks), so that writing prefetch lines into the instruction cache does not steal cycles from fetching the instructions from the cache.

The POWER8 core has an additional mode for systems where memory bandwidth is a concern. This mode is called reduced-speculation mode. In this mode, demand requests are marked as speculative if, at the time of the I-cache miss, the fetch is not for the next-to-complete instruction. If the request is marked speculative, the memory subsystem can return an indication of "no data" if the requested data is not found in the L3 cache and local memory has indicated that recent bandwidth demands were greater than a programmable threshold. In this case, the demand request must be retried when the request becomes next-to-complete. Also, in this mode, prefetches are always marked speculative and are not retried regardless of the returned state.

10.1.4 Branch Prediction

In each cycle, up to eight instructions are fetched from the instruction cache. From there, these instructions are sent to the branch prediction logic. The branch prediction logic scans all the fetched instructions looking for branches; this information is used by the instruction decode logic to look for up to two branches for group formation. Depending upon the branch type found, various branch prediction mechanisms engage to help predict the target address of the branch or the branch direction or both. More specifically, branch target addresses for **bclr** and **bcctr** instructions can be predicted using the link stack or the count cache mechanism (target addresses for absolute and relative branches are computed directly as part of the branch scan function). Dynamic branch direction prediction (taken or not taken) is done through the use of three branch history tables. Static branch direction prediction is done using hints as defined by *Power ISA User Instruction Set Architecture (Book I)*.

It is important to note that all conditional branches are predicted in the POWER8 core (even if the condition is resolved well ahead of time or the value of the link or count register is known when the branch to link or count instruction is fetched) and no unconditional branches are “predicted”. As branch instructions flow through the rest of the pipeline and ultimately execute in the branch execution unit, the actual outcome of the branches is determined. At that point, if the predictions were correct, the branch instructions are simply completed like all other instructions. If the prediction is incorrect, the instruction fetch logic issues a flush and redirects the pipeline down the corrected path.

10.1.4.1 Branch Direction Prediction using the Branch History Tables

The POWER8 core uses a set of three branch history tables to predict the direction of branch instructions. The first table, called the local predictor, is similar to the traditional branch history table. It is a 16K entry array that is indexed by the address of the branch instruction to produce a 2-bit predictor. The MSb of the 2-bit predictor indicates whether the branch direction should be “taken” or “not-taken”.

The second table, called the global predictor, works to predict a branch based on the actual path of execution to reach the branch. The path of execution is identified by a 20-bit vector, one bit per fetch group (that is, the group of instructions fetched in a cycle), for each of the previous 20 fetch groups. This vector is referred to as the global history vector. Each bit in the global history vector indicates whether the next group of instructions fetched are from a sequential cache sector (0) or not (1). The global history vector captures this information for the actual path of execution through these sectors. That is, if there is a redirection of instruction fetching, some of the fetched group of instructions are discarded and the global history vector is corrected immediately. The 20-bit global history vector is first folded (by a bitwise XOR of bits 0:10 with bits 9:19 to generate an 11-bit path vector, which is then hashed by a bitwise XOR with the address of the branch instruction to index into the 16K entry global history table to produce another 2-bit branch direction predictor. Similar to the local predictor, the MSb of this 2-bit global predictor is an alternate indicator of whether the branch should be predicted to be “taken” or “not-taken”.

Finally, the third table, called the selector table, keeps track of which of the two prediction schemes works better for a given branch. It is used to select between the local and the global predictions. The 16K entry selector table is indexed exactly the same way as the global history table is indexed and the MSB of the selected entry is used as the 1-bit selector. This combination of branch prediction tables has been shown to produce very accurate predictions on a wide range of workload types.

If the first branch encountered in a particular cycle is predicted as not taken, the POWER8 core can predict and act on a second branch in the same cycle. In this case, the machine registers both branches as predicted (for subsequent resolution at branch execution), and it redirects the instruction fetching based on the second branch.



As branch instructions are executed and resolved, the branch history tables (as well as, the other predictors described in the following sections) are updated to reflect the latest and most accurate information. Unconditional branches (including branches with the BO field set to '1z1zz') and statically predicted conditional branches (such as, branches with the "a" bit set to '1') do not have an entry in the BHTs. Therefore, they do not cause any BHT update.

All three BHTs are implemented as banked arrays and allow concurrent read and write operations. If the concurrent accesses are to different banks, both are honored. However, if there is a bank conflict, the read is given higher priority. The BHT update logic will perform multiple write updates speculatively up to 5 attempts before forcing a hole in the fetch logic to allow the write to be done.

10.1.4.2 Branch Prediction using Static Prediction and “a”, “t” Bits

For some conditional branches, the software knows what the branch direction prediction ought to be. The POWER8 core allows the software to override the dynamic branch prediction in such cases. Software communicates its wish to override dynamic branch prediction by setting the "a" bit in the BO field. If the "a" bit is '0', then the "t" bit is ignored and the dynamic branch prediction described above is used. If the "a" bit is '1' for a branch, dynamic branch prediction is not used and no entry is updated in the branch history table when such a branch is executed. The static branch direction prediction itself is communicated by properly setting the "t" bit in the BO-field ('1' for taken and '0' for not-taken). Software is expected to use this feature only for those conditional branches for which it believes that software branch prediction will result in at least as good a performance as the hardware branch prediction.

Three separate cases have been identified where the software should override hardware branch prediction for improved performance:

- *When the conditional branch is known by the software to be almost always uni-directional.* For example, branches that guard segments of code that are only executed when a rare event occurs.
- *For the branches that close out a lock acquisition sequence.* It is desirable to force the branch prediction to be *not taken*. This provides the best performance for the most common case where the lock is successfully acquired. Even if the lock is not successfully acquired on this iteration, it is still best to assume (from a branch prediction standpoint) that it will be acquired in the next iteration. Note that, left alone, if the lock is not acquired in the first iteration, the branch history mechanism would work to update the prediction to predict *taken* (that is, predict lock acquisition failure and cause more "lwarx" traffic) for the next iteration.

```
top: lwarx
    add
    stwrx
    bc top      <-- POWER8 predicts this branch to be not taken, through
                  software directives that properly set the "a" and "t" bits.
```

- *To force a conditional branch to be always mispredicted to initiate instruction prefetching.* This allows some instructions to be speculatively executed or processed to some extent by the instruction fetch logic before they are discarded. The instruction in the (wrongly) predicted path can be used as hint instruction to the memory subsystem. For example, *software prefetching of instructions* from location "Line_to_touch" can be initiated by forcing a branch misprediction as follows ("a"-bit in the **bc** instruction indicates "must agree with static prediction").

Short distance touches:

```
    bc Line_to_touch      // Static prediction taken, but CR bit is set "not-taken"
```

Long distance touches:

```
    bc Next                // Static prediction not-taken, but CR bit is set "taken"
```



```
b Line_to_touch      // Initiate prefetch for cache line "Line_to_touch"  
Next:...              // Instructions in the actual instruction stream
```

This type of software prefetching is useful if the line to prefetch is in the L3 cache or beyond. Due to high penalty for branch misprediction, it might not be beneficial if the referenced line is already in the L2 cache and even harmful if it is already in the I-cache. It is beneficial if the compiler makes special attempts to schedule code around such a branch that reduces the misprediction penalty. Attempts to reduce the forced branch misprediction penalty can be made by:

- Setting the CR bit used by the “bc” as early as possible
- Scheduling such a branch in a code segment where there are relatively few branches so that the branch does not wait too long in the branch issue queue behind other branches
- Trying to overlap a likely D-cache miss with the forced branch misprediction
- Scheduling such a branch after an existing long chain of flow dependency

10.1.4.3 Address Prediction Using the Link Stack

The POWER8 processor core uses a link stack to predict the target address for a branch-to-link instruction that it believes corresponds to a subroutine return. By setting the hint bits in a branch-to-link instruction, software communicates to the processor whether the instruction represents a subroutine return or a target address that is likely to repeat or neither (see *Table 10-1*).

When instruction fetch logic fetches a branch and link instruction either unconditional or conditional but predicted taken, it pushes the address of the next instruction into the stack. When it fetches a branch-to-link instruction with “taken” prediction and with hint bits indicating a subroutine return, the stack is popped and instruction fetching starts from the popped address.

In the POWER8 core, the link stack is 32-entries deep per thread in single-thread and SMT2 mode. In SMT4 mode, it is 16-entries deep. In SMT8 mode, it is 8-entries deep. In all modes, entries are preserved to keep speculative pushes, which can be used for branch misprediction recovery.

Speculative execution can corrupt the link stack, both its pointer and its contents. The exact nature of the corruption depends on the sequence of the stack-modifying branch instructions that get purged from the system on a misspeculation.

Because branch-to-link instructions are fairly common and the branch misprediction penalty is high, the POWER8 core uses an extensive speculation tolerance mechanism in its link-stack implementation that allows it to recover the link stack under most circumstances. To recover the stack pointer at misspeculation, the value of the stack pointer at the time a branch is scanned by the instruction fetch logic is stored in a table and restored from on branch misprediction.



*Table 10-1. Handling of **bclrl** and **bclrI** Instructions*

Instruction	BH Field	POWER8 Design	Power ISA
bclrl	xx	If the branch is predicted taken, the link stack address is used as the predicted target address; however, the link stack is not popped.	Reserved.
bclr	00	If the branch is predicted taken, the link stack is popped and the popped address is used as the predicted target address.	The branch is a subroutine return.
bclr	01	If the branch is predicted taken, the target is predicted using the count cache. The count cache data and confidence fields might be updated when the branch is executed and resolved. No action is taken by the link stack.	Target address is likely to repeat.
bclr	10	Same as BH = '00'	Reserved.
bclr	11	If the branch is predicted taken, the link stack address is used as the predicted target address; however, the link stack is not popped.	Target is not predictable.

10.1.4.4 Address Prediction using the Count Cache

The target address of a branch to count (**bcctr[I]**) instruction is often repetitive and can be predicted if the address is saved in a cache from an earlier execution of the same instruction. This is also true for some of the branch-to-link (**bclrl[I]**) instructions, which are not predictable through the use of the link stack because they do not correspond to a subroutine return. By setting the hint bits appropriately, software communicates to the hardware whether the target address for such branches are predictable using a cache. See *Table 10-2*.

The POWER8 core uses two 256-entry count caches (one with hashed addressing similar to the global branch history table and one using the unhashed instruction address) for predicting the target of a **bclrl[I]** or **bcctr[I]** instruction whose target address is likely to repeat. Each entry in the count cache can hold a 62-bit address. When a **bclrl[I]** or **bcctr[I]** instruction is executed, for which the software indicates that the target is predictable using such a cache, the target address is written in the count cache. When such an instruction is fetched, the target address is predicted using the count cache. The count cache does not have a separate array to select between the two count caches. Instead, the 2-bit selector value is stored in the array using the unhashed instruction address.

Sometimes a given branch has a small set of targets, but predominantly favors a particular target. To help predict such branches, the count cache also uses a 2-bit confirmation counter (for each entry) for replacement. Every time an entry is used for successful prediction, the counter is incremented (saturating counter), and it is decremented on a misprediction. If the counter is zero, the entry can be replaced. When an entry is first allocated or installed, the counter value is set to 1.

Note: **bcctr** instructions that are found in the local count cache must be located in octword blocks for the POWER8 core. The local count cache is primarily used for **bctr(I)** instructions that have only one target. The POWER8 core can only make one prediction per fetch group and bases the prediction off the fetch group address (octword) and not the branch that is eventually identified as the first taken branch in the group. If there are other non-bcctr branches in the octword, there is no problem. But if there are two or more bcctr branches in the fetch group, only one can be in the count cache at a time so prediction can thrash.

Table 10-2. Handling of **bcctr and **bcctrl** Instructions**

Instruction	BH field	POWER8 Design	Power ISA
bcctr, bcctrl	00	If the branch is predicted taken, the target address is predicted using the count cache. Update the count cache when the branch is executed, if the branch is resolved as taken. For the bcctrl instruction, if the branch is predicted taken, push in the link stack the address of the next sequential instruction when the bcctrl instruction is fetched.	Target address is likely to repeat.
bcctr, bcctrl	01	Same as BH = '00'.	Reserved
bcctr, bcctrl	10	Same as BH = '00'.	Reserved
bcctr, bcctrl	11	Same as BH = '00'.	Target is not predictable

10.1.4.5 Round-Trip Branch Processing

Instruction fetch logic can fetch up to eight instructions in a given cycle and forward them down the pipeline to the instruction decode logic. It scans through this group of instructions in the next cycle to determine the position of all the branches in the group and the predictions for them. At the same time, it fetches the next sequential cache sector of eight instructions, if there is no I-cache miss, I-ERAT miss, fetch redirection, pipeline hold (for example, due to successive stages not having enough resources to process the instructions already in progress). The next sequential I-cache sector is fetched with the assumption that there is no unconditional branch or a conditional branch that is predicted taken in the previous group of instructions that it fetched. If the assumption is correct, instruction fetch proceeds unhindered. If incorrect, a *bht redirection* event occurs, when instruction fetch logic instructs the instruction decode logic to discard the instructions it sent after that branch and starts fetching from the target of that branch. If it is a branch-to-link/count instruction, the target of the branch is predicted based on the prediction mechanism described earlier. Otherwise, the target of the branch is calculated from the instruction and its address. On a BHT redirection, two cycles worth of instruction fetching might be lost.

If the branch is mispredicted, all the instructions in the pipeline that are fetched after the branch instruction are purged from the system and new instructions from the new address, as determined by the executed branch instruction, are fetched. The purging involves restoring all the queues and the mapper register files to the state that is consistent with not having fetched or executed any instruction after the branch.

Because of its deep pipeline, the branch misprediction penalty is high in the POWER8 core. If a mispredicted branch is executed before some of its earlier instructions (in program order), the pipeline will not be completely dry after the misprediction. These earlier instructions can keep the execution units busy while the instruction fetch logic brings new instructions, thus reducing the penalty to some extent. The average branch misprediction penalty depends on many factors including when the associated CR bit becomes available, the number of instructions that survive a misprediction, and associated cache misses. However, after a branch is executed and found to be mispredicted, it takes a minimum of 16 cycles for a new instruction from the redirected path to come down the pipeline.



10.1.4.6 BC+4 Handling

Unconditional branches with the link bit set and a displacement of four can be used to set the address of the next instruction into the link register. Architecturally these branches are taken and go to the next instruction. The hardware handles BCL+4, with a BO = 20 (unconditional taken) as a special case. For this special case, the branch is always treated as *not taken* for fetch, resulting in a no-fetch penalty. When the special branch executes, it updates the link register, and does not cause a flush (even though fetch processed it as *not taken* and architecturally it was *taken*). The PMU sees these special branches as requiring direction prediction, direction predicted correctly, and branch not taken counts.

10.1.4.7 BC+8 Handling

In addition to the normal branch handling, bc instructions whose BO field indicates only a CR dependency with a displacement of +8 are handled as a special case. Such bc+8 instructions are marked as potential candidates for conversion to predication. When the next instruction after a bc+8 is in the list of allowed to be predicated (see *Table 10-3*), the bc+8 is marked as first for group formation processing. Finally, if the branch prediction for the bc+8 instruction is predicted not taken, the instruction pair is treated similar to an **isel**, in that the branch cannot be mispredicted and the next instruction is conditionally executed based on the bc outcome. Any time a bc+8 instruction is executed, the branch prediction algorithm writes a value assuming the branch was not taken. This encourages all bc+8 instructions to be converted to predication. Use of a NOP can be used to prevent bc+8 conversions in specific cases (making the bc+8 a bc+12).

Table 10-3. bc+8 Pairable Instructions

Assembler Mnemonic
addi
addis
add
and
or
xor
ori
stb
sth
stw
std

When a bc+8 instruction pair is converted to predication, the destination register now is pending until the predication is resolved. This can mean that subsequent dependent instructions are prevented from issuing. The bc+8 conversions must be avoided for the following circumstances:

1. If the branch is expected to be very predictable, the branch prediction hardware outperforms predication.
2. If the pairable operation is the start of a dependency chain involving loads, sometimes it is more beneficial to have the loads start execution even if the branch eventually flushes. Predication delays dependent operations from executing. Therefore, if a long latency load is delayed, the performance is better served by allowing the branch to be predicted. This is especially true when the resultant memory access can be the same cache line in either outcome. Branch misprediction can be hidden by the long latency memory

reference while the dependency delay due to predication cannot be hidden. Heap sort algorithms show this effect.

10.1.5 Store-Hit-Load Avoidance Table

In an out-of-order execution machine, a younger load can be executed ahead of an older store. When the addresses are independent, this is not a problem. But if the addresses overlap, the load has the wrong data. The base design detects this condition and flushes the instruction stream after the store causing all subsequent instructions to be refetched. This ensures that the load follows the store and gets the newer data.

There are cases where the store hit load is repeated many times (possibly in a loop) where performance is degraded due to the flushes. Therefore, the POWER8 core adds an 8-entry table that records the store address when an SHL flush occurs. On the second occurrence, the decode logic is instructed to mark the RA field of the store as written (store with dependency). Assuming the out-of-order load uses the same RA register ($RA \neq 0$), the load is now delayed due to this artificial dependency. This performs significantly better than flushing. The table also monitors the store address for subsequent SHL flushes and disables the function when it does not appear to help.

The store-hit-load table is managed with an LRU algorithm. Entries are also invalidated when the corresponding I-cache congruence class is written. Only certain stores can be converted to store with dependency. *Table 10-4* lists ineligible stores.

Table 10-4. Stores Ineligible for SHL Avoidance

Ineligible Stores for SHL Table
All stores when RA = 0
All stores with update
All store conditionals

10.1.6 Instruction Buffer

Instructions read out of the I-cache are forwarded to the IBuffer as a staging area for group formation. The IBuffer is arranged as a register file where each row can hold up to four instructions (16-byte aligned from the I-cache). There are a total of 32 rows, but the number of rows per thread varies by SMT mode as shown in *Table 10-5*. Branches not from the last instruction of an aligned quadword and not to the first instruction of an aligned quadword cause inefficiencies in the IBuffer.

Table 10-5. IBuffer Rows per Thread

SMT Mode	Rows
1	16
2	16
4	8
8	4

IFetching for a thread is prevented if there is no room in the IBuffer and is restarted as the IBuffer is drained. There must be room for eight instructions before an IFetch is initiated. There are non-flush circumstances where a full IBuffer could go empty before newly fetched instructions make it back to the IBuffer if groups are large and dispatch is unencumbered. But this is expected to be rare.



10.1.7 Group Formation

Instructions are pulled from the IBuffer and formed into dispatch groups. Instructions in a dispatch group enter the out-of-order part of the processor in parallel and must complete together so that group formation can limit the overall performance of a processor.

In single thread mode, the group size is limited to six nonbranch and two branch instructions. All instructions must come from the oldest two quadwords in the IBuffer, and two quadwords can be emptied per cycle (with some exceptions). Slots are numbered 0 - 5 for the non-branch instructions and 6 - 7 for the branch instructions.

In multi-thread mode, group size is limited to three nonbranch instructions and one branch instruction for each of two threads (groups 0 and 1). All instructions must come from the oldest two quadwords per thread, but only one quadword per thread can be emptied per cycle. Nonbranch slots are numbered 0 - 2 for the group 0 thread and 3 - 5 for the group 1 thread. The branch slot is 6 for the group 0 thread and 7 for the group 1 thread.

Group formation rules for the POWER8 core are in the following sections.

10.1.7.1 General Rules

General group formation rules (regardless of SMT mode) are:

- Can go past the first branch (predicted T or NT) to pick up more instructions (branch anywhere). A special case ends a group when a backward branch to within 1024 bytes is predicted taken. This allows a compiler to assume group formation starts at a loop-entry point after the first iteration.
- Instructions marked First must start a group.
- Instructions marked Last must end a group.
- No FPU operations are allowed in a group after a branch (T or NT).
- A bc+8/12 operation is marked First if the branch confidence table indicates to predicate.
- A 2-way cracked operation occupies two nonbranch slots.
- A 3-way cracked operation occupies three nonbranch slots (modified from 4-way cracked). This ends the group.
- Branch and link instructions that update the Link Register are marked Last.
- A 3-source operation in slots 0 or 1 reserves slot 2 for transmitting the third operand. A 3-source operations in slots 3 or 4 reserves slot 5 for transmitting the third operand. No 3-source operations are allowed in slots 2 or 5.



10.1.7.2 Rules Specific to ST Mode

Rules specific to ST mode only (6 + 2 dispatch) are:

- A group is sourced from the bottom two rows of the IBuffer in each cycle.
- Maximum of six nonbranches per group.
- Maximum of two branches per group.
- A group ends immediately after the second branch (T or NT).
- Eight instructions total per group.
- Slot 2 is reserved when a bc+8/12 starts a group. This slot is used for extra source operands in the predicated group.
- Two-way cracked operations go to slots 0 and 1. A second 2-way cracked operation can be in the same group, if they are next to each other in the code stream (goes to slots 3 and 4).
- Three-way cracked operations go to slots 0, 1, and 2, and end the group.

10.1.7.3 Rules Specific to SMT Modes

Rules specific to SMT modes (dual 3 + 1 dispatch) are:

- The two 3+1 dispatch halves operate independently, reading two threads out of the IBuffer in each cycle.
- A group is sourced from the bottom two rows of the IBuffer (two rows per thread) in each cycle.
- Maximum of three nonbranches per group.
- Maximum of one branch per group.
- Four instructions total per group.
- Two-way cracked operations go to slots 0 and 1 for group 0, or slots 3 and 4 for group 1.
- Three-way cracked operations go to slots 0, 1, and 2 for group 0, or slots 3, 4, and 5 for group 1.

Both dispatch halves share a single microcode engine. This can cause one half to stall the other if both halves need microcode simultaneously.

10.1.8 Group Ending NOP

When a store is followed by a load in a single-dispatch group and the addresses overlap such that forwarding is not allowed (load-hit-store), the group must be *flushed* to break up the store and load (store cannot complete due to the load and the load cannot complete due to the store). A group ending NOP inserted between the store and the load will prevent the flush. Use an “**ori R2,R2,0**” as the group ending instruction.

10.1.9 First and Last Instructions

Instructions that require specific execution units are marked as *First*. See *Table 10-6* on page 210 for a list of instructions marked as First.

Instructions that require that no other instructions follow in the same dispatch group are marked as *Last* and end the present dispatch group prematurely. See *Table 10-7* on page 211 for a list of instructions marked as Last.



Table 10-6. List of Instructions Marked as First

Mnemonic
addc
addc.
addco
addco.
addeo.
addg6s
addmeo.
addzeo.
and
clrbhrb
crand
crandc
creqv
crnand
crnor
cror
crorc
crxor
dcbf
dcbfl.l=1
dcbfl.l=3
dcbst
dcbz.0
doze
eieio
hrfid
isel
isync
lmw
logmpp
lq
lqarx0_1
lqarx1_1
lswi
lswx
lwsync
mcrf
mfbhrbe
mfcr
mfmsr
mfocrf
mfspr [non-renamed]
mfspr_CTR
mfspr_default

Mnemonic
mfspr_LR
mfspr_xer
mfsr
mfsrin
msgclr
msgclrpp
msgsnd
msgsndp
msle
mtcrf
mtiamr
mtmsr
mtmsr_ee
mtmsrd
mtmsrd_ee
mtspr [non-renamed]
mtspr_CTR
mtspr_default
mtspr_LR
mtspr_xer
mtsr
mtsrin
nap
or
rfebb
rfid
rvwinkle
sc
slbfee.
slbia
slbia.001
slbia.010
slbia.011
slbia.1--
slbie
slbmfee
slbmfev
slbmte
sleep
sp_attn
stbcix
stbcx.
stdcix
stdcx.

Mnemonic
sthcix
sthcx.
stmw
stswi
stswx
stwcix
stwcx.
subfc
subfc.
subfco
subfco.
subfeo.
subfmeo.
subfzeo.
sync
syncpte
tabort.
tbegin.
tcheck.
tend.
tlbie
tlbiel
tlbiellpg
tlbsync
trechkpt
treclaim.



Table 10-7. List of Instructions Marked as Last

Mnemonic	Mnemonic	Mnemonic	Mnemonic
addc	mfspr_MMCR0.0	mfspr_SPURR	mtspr_CIR
addc.	mfspr_MMCR0.1	mfspr_TAR	mtspr_CSIGR
addco	mfspr_MMCR1.0	mfspr_TB.268	mtspr_CTRL
addco.	mfspr_MMCR1.1	mfspr_TB40	mtspr_DEC
addeo.	mfspr_MMCR2a	mfspr_TBL	mtspr_DSCR3
addg6s	mfspr_MMCR2b	mfspr_TBU.269	mtspr_EBBHR
addmeo.	mfspr_MMCR40	mfspr_TBU.285	mtspr_EBRR
addzeo.	mfspr_MMCR41	mfspr_TEXASR	mtspr_FSCR
and	mfspr_MMCRC	mfspr_TEXASRU	mtspr_HDEC
clrbhrb	mfspr_MMCRH	mfspr_TFHAR	mtspr_HEIR
doze	mfspr_MMCRS	mfspr_TFIAR	mtspr_HFSCR
hrfid	mfspr_PCR	mfspr_TFMR	mtspr_HIDO
isync	mfspr_PIR	mfspr_TMFR	mtspr_HID1
lq	mfspr_PMC1.0	mfspr_TRACE	mtspr_HMER
lqarx0_1	mfspr_PMC1.1	mfspr_trig0	mtspr_HSPRG0
lqarx1_1	mfspr_PMC2.0	mfspr_trig1	mtspr_HSPRG1
mfcr	mfspr_PMC2.1	mfspr_trig2	mtspr_HSRR0
mfspr_BESCR	mfspr_PMC3.0	mfspr_VTB	mtspr_HSRR1
mfspr_BESCRR	mfspr_PMC3.1	mfspr_xer	mtspr_IC
mfspr_BESCRRU	mfspr_PMC4.0	mfsr	mtspr_MMCR2a
mfspr_BESCRS	mfspr_PMC4.1	mfsrin	mtspr_MMCR2b
mfspr_BESCRSU	mfspr_PMC5.0	msgclr	mtspr_MMCRC
mfspr_CIR	mfspr_PMC5.1	msgclrP	mtspr_MMCRS
mfspr_CTRL	mfspr_PMC6.0	msgsnd	mtspr_PIR
mfspr_DEC	mfspr_PMC6.1	msgsndp	mtspr_PMCR
mfspr_DSCR3	mfspr_PMCR	msle	mtspr_PMICR
mfspr_EBBHR	mfspr_PMICR	mtiamr	mtspr_PMMAR
mfspr_EBRR	mfspr_PMMAR	mtmsr	mtspr_PMSR
mfspr_HDEC	mfspr_PMSR	mtmsr_ee	mtspr_PPR
mfspr_HIDO	mfspr_PPR32	mtmsrd	mtspr_PPR32
mfspr_HMEER	mfspr_PSPB	mtmsrd_ee	mtspr_PSPB
mfspr_HMER	mfspr_PURR	mtspr_BESCR	mtspr_PURR
mfspr_HPMC1	mfspr_SIER0	mtspr_BESCRR	mtspr_RPR
mfspr_HPMC2	mfspr_SIER1	mtspr_BESCRRU	mtspr_SIER0
mfspr_HPMC3	mfspr_SPMC1	mtspr_BESCRS	mtspr_SIER1
mfspr_HPMC4	mfspr_SPMC2	mtspr_BESCRSU	mtspr_SPMC1
mfspr_IC	mfspr_SPRC	mtspr_CFAR	mtspr_SPMC2
mfspr_IMC	mfspr_SPRD	mtspr_CIABR	mtspr_SPRC

Mnemonic	Mnemonic
mtspr_SPRD	slbmte
mtspr_SPRG0	sleep
mtspr_SPRG1	sp_attn
mtspr_SPRG2	stbcix
mtspr_SPRG3	stbcx.
mtspr_SPURR	stdcix
mtspr_SRR0	stdcx.
mtspr_SRR1	sthcix
mtspr_TAR	sthcx.
mtspr_TB40	stwcix
mtspr_TBL	stwcx.
mtspr_TBU	subfc
mtspr_TEXASR	subfc.
mtspr_TEXASRU	subfco
mtspr_TFHAR	subfco.
mtspr_TFIAR	subfeo.
mtspr_TFMR	subfmeo.
mtspr_TRACE	subfzeo.
mtspr_trig0	tabort.
mtspr_trig1	tbegin.
mtspr_TSCR	tcheck.
mtspr_TSRR	tcheck.
mtspr_TTR	tend.
mtspr_VTB	tend.
mtspr_xer	tlbie
nap	tlbiel
or	tlbiellpg
rfebb	trechkpt
rfid	treclaim.
rvwinkle	
sc	
slbfee.	
slbia	
slbia.001	
slbia.010	
slbia.011	
slbia.1--	
slbie	
slbmfee	
slbmfev	

10.1.10 2-Way and 3-Way Cracked Instructions

For some instructions that are important for performance and that cannot be done in a single IOP but do fit in two or three IOPs, instruction cracking is used rather than microcode. Instruction cracking involves expanding a single architected instruction into two or three IOPs in the same instruction dispatch group. The IOPs are uninterruptable in the middle of the sequence. *Table 10-8* lists the 2-way cracked instructions and *Table 10-9* on page 214 lists the 3-way cracked instructions.

Table 10-8. 2-Way Cracked Instructions

Mnemonic	Mnemonic	Mnemonic	Mnemonic
addc.	logmpp	slbia	stvewx
addco	lq	slbia.001	stvx(l)
adde.	lqarx0_1	slbia.010	stwbrx
addeo	lqarx1_1	slbia.011	stwcix
addic.	mfspr_DSCR3	slbia.1--	stwcx.
addme.	mfsr	slbie	stwux
addmeo	mfsrin	slbmfee	stwx
addo.	msgclr	slbmfev	stxvd2x
addze.	msgclrp	slbmte	stxvw4x
addzeo	msgsnd	sld.	subfc.
cntlzd.	msgsndp	slw.	subfe.
ctlzw.	mtiamr	srad.	subfeo
divd.	mtspr_DSCR3	sraw.	subfme.
divde.	mulhd.	srawi.	subfmeo
divdeo.	mulhdu.	srd.	subfo.
divdeu.	mulhw.	srw.	subfze.
divdeuo.	mulhwu.	stbcix	subfzeo
divdo.	mulld.	stbcx.	tbegin.
divdu.	mulldo.	stbx	tcheck.
divduo.	mullw.	stdbrx	tend.
divw.	mullwo.	stdcix	tlbie
divwe.	nego.	stdcx.	tlbiel
divweo.	rldcl.	stdux	tlbiellpg
divweu.	rldcr.	stdx	xsradi.
divweuo.	rldic.	sthbrx	
divwo.	rldicl.	sthcix	
divwu.	rldicr.	sthcx.	
divwuo.	rldimi.	sthux	
extsb.	rlwimi.	sthx	
extsh.	rlwimm.	stvebx	
extsw.	rlwnm.	stvehx	
isync	slbfee.		

Table 10-9. 3-Way Cracked Instructions

Mnemonic
addg6s
addc
addc.
addco
addco.
subfc
subfc.
subfco
subfco.
addeo.
subfeo.
addzeo.
subzeo.
addmeo.
submeo.

10.1.11 Microcode

Instructions requiring microcode to execute have additional performance penalties to access the microcode. A microcoded instruction after a non-microcoded instruction incurs a 2-cycle penalty before the first group of the microcode routine is available. These cycles are only visible if the pipeline is not stalled. Consecutive microcoded instructions do not require the two additional cycles. There is no overhead at the end of a microcode sequence. In certain unalign cases, these instructions can trap to microcode:

- **lmw**
- **stmw**
- **lswi**
- **lswx**
- **stswn**
- **stswnx**
- **ld, ldu, ldx, ldux**
- **lfdw, lfdwx, lfdx, lfd**
- **lha, lhax, lhau, lhaux, lhz, lhzx, lhzu, lhzux**
- **lwa, lwax, lwau, lwaux, lwz, lwzx, lwzu, lwzux**
- **mtsr, mtsrin**

10.1.12 Instruction Fusion

POWER8 instruction fusion involves combining information from two adjacent instructions into one instruction so that it executes faster than the non-fused case. “Adjacent instructions” refers to the instruction locations after group formation. In single-threaded mode, 6/2 groups are formed where there are up to six non-branch instructions and up to two branch instructions. In multi-threaded mode, two 3/1 groups are formed where there can be up to three non-branch instructions and one branch instruction. To be adjacent, the instructions to be fused must be in the same group of non-branches, without a branch in between. Instructions can span an I-cache line, but will only be fused if both have been fetched and are in the ibuffer.

There are two fuse types:

{addi} followed by one of these {lxvd2x, lxvw4x, lxvdsx, lvebx, lvehx, lviewx, lvx, lxsdx}

Requirements:

addi(rt) = lxvd2x(rb)

lxvd2x(ra) cannot be ‘0’

Result:

addi - no change

lxvd2x gets the immed field from addi., rb is not used.

This effectively provides a d-form version for the vector loads.

The dependency between the two operations is removed.

{addis} followed by one of these {ld, lbz, lhz, lwz}

Requirements:

addis(rt) = ld(ra) = ld(rt) - (cannot be ‘0’)

addis(SI) first 12 bits must be all 0’s or all 1’s

TA = 0 for this fusion to occur and

if SI = 11111111110000 and the msb of the d/ds field of the load equals ‘1’, then fusion does not occur.

Result:

addis gets changed to a NOP. (It still takes up a dispatch slot, but is sent directly to completion.)

The last 5 bits of addis(SI) are sent with the ld (information from the addis is passed to the ld).

The addis is removed from execution.



10.1.13 Instruction Dispatch

Dispatch represents the last in-order stage of the pipeline until completion. At this point, resources are checked and the pipeline stalled if not available. *Table 10-10* lists the dispatch stall conditions due to resource limitations.

Table 10-10. Resource Requirements for Dispatch

Dispatch Hold Condition	Minimum Cycles Required	Release (Completion Based)	Inefficiencies
Branch Issue Queue full			
CR Issue Queue full			
Unified Issue Queue full	6	+6 cycles	Three cycles reserved per dispatch lane
Insufficient CR mapper entries	6	+0 cycles	Three cycles reserved per dispatch lane
Insufficient <u>XER</u> mapper entries	6	+6 cycles	
Insufficient <u>CTRLR</u> mapper entries	4		
Insufficient <u>GPR</u> mapper entries		+9 cycles	Three cycles reserved per dispatch lane
Insufficient <u>EPSCR</u> mapper entries	2		
Insufficient <u>VRF</u> mapper entries		+9 cycles	Three cycles reserved per dispatch lane
<u>LRQ</u> full			
<u>SRQ</u> full			
<u>GCT</u> full (setup stage)			

10.1.14 Instruction Issue

Up to eight internal operations (IOPs) can be dispatched and renamed per cycle. After renaming, these IOPs are placed into a set of issue queues that are distributed by instruction type.

There are three different issue queues:

- Unified issue queue (UniQ): This is a 64-entry queue that is partitioned into two halves. Queue half 0 issues instructions to FX0, VS0, LS0, and LU0. Queue half 1 issues instructions to FX1, VS1, LS1, and LU1. Each queue half can issue instructions to one FXU, one VSU, one LU, and one LSU. Instructions from either queue half can issue to the DFU and Crypto units.
- BRQ: 15-entry issue queue for branch operations
- CRQ: 8-entry issue queue for condition register logical operations and mfSPRs, for SPRs that are owned by the IFU, ISU, and PC units. mfSPR IOPs occupy one entry in both the UniQ and CRQ.

10.1.14.1 Steering Policy

In ST mode, instructions are assigned to queue halves based on dispatch slots.

The instruction in slots 0, 1, and 2 are steered to alternating queue halves. For example, if IOP 0 is steered to queue half 1, IOP 1 would be steered to queue half 0 and IOP 2 would be steered to queue half 1. The IOP in slot 3 must be steered to the opposite queue half as the IOP in slot 0. Likewise IOP 4 is steered opposite IOP 1, and IOP 5 is steered opposite IOP 2.

The first instruction in a group is assigned to the queue half that had the fewer number of instructions assigned to it in the previous dispatch cycle.

In SMT2, SMT4, and SMT8 modes, the two queue halves are partitioned by thread set, so instructions are steered by thread set. Instructions in slots 0, 1, and 2 (thread set 0) are assigned to queue half 0. Instructions in slots 3, 4, and 5 (thread set 1) are assigned to queue half 1.

10.1.14.2 BRQ and CRQ Operation

Instructions are allocated into the BRQ and CRQ at dispatch time. Each cycle, the oldest ready instruction out of each queue is selected for execution. The age between threads is determined by dispatch order. An instruction is deallocated from its issue queue two cycles after it issues.

10.1.14.3 UniQ Issue Policies

Each cycle, each queue half selects four instructions to be issued, one each to the FXU, VSU, LU, and LSU for that queue half. In general, selection is biased toward the oldest ready instruction. Instructions can be speculatively issued before their source operands are really ready. They can also be speculatively issued before knowing if there is going to be contention for a resource such as a GPR write port. Instructions are rejected if they cannot be successfully issued, and are re-issued again later.

10.1.14.4 FXU and VSU Selection

In general, the oldest ready instruction is selected for issue on a given unit. However, due to cycle time restrictions, it is impossible to determine the oldest among 32 instructions and select one in a cycle. Instead, the issue queue half is divided into two quadrants of 16 entries each. The oldest ready instruction in each quadrant is selected. Each cycle, one quadrant is determined to be high priority and the other is low priority. An LFSR is used to select which quadrant is high priority. If the high-priority quadrant has any ready instructions, the oldest one from that quadrant is selected. Otherwise, the oldest from the low-priority quadrant is selected.

10.1.14.5 LU Selection

In addition to loads, simple fixed-point instructions are also eligible to be issued to the LU. The selection logic picks the first instruction found in this list:

- Oldest LU-eligible load or store_data in the high-priority quadrant
- Oldest LU-eligible load or store_data in the low-priority quadrant
- Youngest simple fixed-point in the high-priority quadrant
- Youngest simple fixed-point in the low-priority quadrant



10.1.14.6 LSU Selection

Stores, simple fixed-point, and some loads are eligible to be issued to the LSU. The selection logic picks the first instruction found in this list:

- Oldest store (or LSU-only load) in the high-priority quadrant
- Oldest store (or LSU-only load) in the low-priority quadrant
- Oldest LSU-eligible load in the low-priority quadrant
- Youngest simple fixed-point in the low-priority quadrant
- Youngest simple fixed-point in the high-priority quadrant

10.1.14.7 Dispatch Bypass Instruction Selection

Instructions are selected for issue using the normal selection mechanism if they were dispatched at least four cycles ago. Instructions in the dispatch+3 stage are selected for issue using a bypass (such as, dispatch+3 bypass), which uses a different selection policy.

The following instructions are not eligible to use the dispatch+3 bypass:

- Instructions that are dependent upon older instructions in the same dispatch group or the group dispatched in the previous cycle
- Conditional IOPs (used by bc+8)
- Loads and stores waiting for real LRQ or SRQ entries
- Completion-serializing instructions
- VS-routed operations
- Instructions with XER sources

Instructions are speculatively scheduled (oldest first) for dispatch bypass as if there are no older ready instructions in the issue queue in stages disp+4 or later.

10.1.14.8 Back-to-Back Issue Policy

In general, single-cycle fixed-point instructions are issued in consecutive cycles if they are in the same half of the queue. An instruction is considered to be back-to-back ready (B2B_RDY) if, for all of its source operands, either (1) the available lines are already set, or (2) the source operands in the same half of the queue and are currently marked as FX0_RDY and are single-cycle instructions. The oldest instruction that is either B2B_RDY or FX0_RDY is selected for execution on FX0.

10.1.14.9 Limitations of Back-to-Back

There are three sources of inefficiency in the back-to-back mechanism:

1. If an instruction has two source operands in the same half of the queue, neither of which has been deallocated, the dependent cannot wake up back-to-back unless both of the producer instructions are marked as Ready (such as, they are candidates for selection).

In the example shown in *Table 10-11*, instruction C has two source operands, A and B. Instructions A and B are independent. All three instructions are in the same half of the queue. Instruction A is selected in cycle 0. In cycle 1, it issues and its available line is set and dependent instructions (might) wake up. In cycle 1, B wakes up, and it is in the FX0_RDY latch in cycle 2. However, by this point in time, A is no longer marked in the FX0_RDY latch. Instruction A has not been deallocated from the queue yet; therefore, C is still tracking two source operands - A and B. Because A is not marked as FX0_RDY, instruction C cannot wake up via the back-to-back mechanism in cycle 2.

Table 10-11. Example where Back-to-Back is not Possible A->C, B->C (All in the same half of queue)

	0	1	2	3			
A	sel	iss	rf				
B		wak	sel	iss	rf		
C				wak	sel	iss	rf
fx0_rdy	A		B		C		
b2b_rdy							

2. Only the first level in a chain of dependent instructions can execute via the back-to-back mechanism. For example, suppose instruction C is dependent on B, which is dependent on A. B can be selected back-to-back after A, but then C cannot be selected back-to-back after B. This is shown in *Table 10-12*. Note that D can then execute back-to-back after instruction C.

If B is not selected back-to-back after A, C can be selected back-to-back after B.

Table 10-12. A -> B -> C (All in the same half of queue)

	0	1	2	3	4		
A	sel	iss	rf				
B	b2b wake	sel	iss				
C			wake	sel	iss		
D				b2b wake	sel	iss	
fx0_rdy	A			C			
b2b_rdy		B			D		

3. If there are two or more FX instructions marked as ready in the ready latches (that is, candidates for selection) and zero load/stores that are candidates for selection, then in the following cycle there will be no instructions woken up back-to-back that are candidates for selection. This is to handle the problem of FX instructions that are issued to the LSU attempting to wake up their dependents back-to-back.



10.1.14.10 Dual-Issued Stores

The store_agen and store_data portions of D-form store instructions share the same entry in the issue queue. For fixed-point stores, the store_agen is issued to the LSU and the store_data is issued to the LU. For FP/VMX/VSX stores, the store_agen is issued to the LSU and the store_data is issued to the VSU. The dependence tracking logic only tracks the source operand of the store_agen. Both the store_data and store_agen will “wake up” and attempt to issue as soon as the store_agen’s source operand is ready. The store_data will be delayed until that point even if it was already ready to issue. If the store_data’s source operand was not yet available and it tries to issue anyway, it will get rejected.

10.1.14.11 Wake-up Misspeculations

There are several cases where instructions issued can be woken up too early. Misspeculations are detected when the dependent instruction is in the register file access stage. In the following cases, instructions get a source-unavailable reject:

- Dispatch bypass: instructions can be issued off of a dispatch bypass before their source operands are ready.
- Store_data IOPs for dual-issued stores wake up when the store_agen wakes up. They can try to issue at any time after that point, except one or two cycles after the store_agen is issued. They can continue to reject over and over if their source operands are not ready.
- SAR bypass: If instructions are issued with ‘needs_rtag’ bits on, or if there was a recent castout, they are rejected. See *SAR Bypass - Related Rejects* on page 225.
- Dependents of load misses or LSU-rejected loads wake up their dependents too early.
- The instruction’s producers were rejected for any reason (non-ready sources or any unavailable resource) after they woke up their dependents.

10.1.14.12 Chains of Misspeculations

It should be pointed out that 1-, 2-, 3-, and 4-cycle instructions can wake up their dependents before it is known that they must be rejected. This is because their available lines get set before or while the misspeculation is detected. Hence, it is possible for an infinite chain of dependent instructions to wake up and be scheduled speculatively. Any of the resource conflicts or wake-up misspeculations discussed previously can cause these scenarios.

Chains of dependents of load instructions continue to wake-up and are issued even after the load is rejected or misses in the L1 cache and DVAL is known to be zero.

10.1.14.13 Other Issue Inefficiencies

Normally, a 2-cycle FXU operation has an issue-to-issue latency of two cycles within the same FXU, and three cycles to the other FXU. However, if a 2-cycle FXU operation is issued and then rejected for any reason, the subsequent times it is issued, the UniQ will attempt to wake up instructions in the other half of the queue assuming a 2-cycle issue-to-issue latency. If a dependent instruction in the other half of the queue is issued too early, it will be rejected in the EXE stage because the result of the 2-cycle instruction is not available.

10.1.14.14 Issue-to-Issue Latencies

Table 10-13 on page 221 shows the minimum issue-to-issue latencies for different classes of instructions. There are some exceptions, as noted in the comments at the end *Table 10-13*. A brief explanation of the acronyms follows:

PM	VMX/VSX permute
XS	VMX/VSX simple operations
FX	All instructions reading from or writing to the FPSCR and VSCR
VX	VMX complex operations
FD	VSX scalar; BFU floating-point
VD	Vector double-precision floating-point
VS	VMX/VSX 4-way vector single-precision floating-point
DFU	Decimal floating-point
CY	Cypher / cryptographic instructions
VSU-FX	Move to/from FPSCR
FXU	Any operation issued to the fixed-point unit
LU	Any operation issued to the load unit, including simple fixed-point and store-data operations
LSU	Any operation issued to the load-store unit, including simple fixed-point operations

Table 10-13. Issue-to-Issue Latencies (Sheet 1 of 2)

From	To	Latency	Comments
PM/XS/FX	any	2 or 7	[1]
FD/VD	any besides FD/VD	7	
FD/VD	FD/VD	6	
CY	FD/VD/VX/VS	7	
CY	any besides FD/VD/VS/VX	6	
VS	any besides VS	7	
VS	VS	6	
DFU	any	13	[2]
VX	any	7	
VSX-to-GPR move	any fixed-point operation	4 or 5	[3]
GPR-to-VSX move	any VSX operation	5	
FP/VS/VMX load	any	5	

1. Normally permutes and VMX-simple operations have a latency of two cycles. However, if a 2-cycle operation is issued five cycles after a 7-cycle operation on the same VSU pipe, the 2-cycle operation is dynamically converted to a 7-cycle operation to prevent a resulting bus collision.
2. This applies to pipelined DFP operations only. It applies to only the DFP IOP of the cracked DFP instruction.
3. Latency of four cycles if VSU pipe 0(1) to FXU pipe 0(1). Otherwise five cycles.
4. See rules about back-to-back-to-back (*Section 10.1.14.9 Limitations of Back-to-Back* on page 219).
5. Minimum two cycles issue-to-issue. Subject to special rules: **cmp** and **bc** must be in the same dispatch group, or else the **cmp** must be the most recent CR writer.



Table 10-13. Issue-to-Issue Latencies (Sheet 2 of 2)

From	To	Latency	Comments
fixed-point load (no XER dest)	any	3	
fixed-point load (with XER dest)	any	4	
store w/ update	any	3	
FXU 1-cycle (no XER dest)	same FXU	1	[4]
FXU 1-cycle (no XER dest)	other FXU, either LU / LSU	2	
FXU 1-cycle (with XER)	same FXU	2	
FXU 1-cycle (with XER)	other FXU, either LU / LSU	3	
FXU 2-cycle (no XER)	same FXU	2	
FXU 2-cycle (no XER)	other FXU, either LU / LSU	2	
FXU 2-cycle (with XER)	same FXU	3	
FXU 2-cycle (with XER)	other FXU, either LU / LSU	3	
FX multiply	same FXU	4	
FX multiply	other FXU, either LU / LSU	5	
divw	same FXU	variable	
divw	other FXU, either LU / LSU	above +1	
mf CTR/LR/TAR/CR	FXU / LU / LSU	5	
cmp	bc	2	[5]
mtCTR/LR/TAR/CR	dependent CRQ operation	6	
mtCTR/LR/TAR	dependent branch	5	
branch updating CTR/LR/TAR	branch reading CTR/LR/TAR	3	
branch updating CTR/LR/TAR	mfCTR/LR/TAR	4	
CR-logical	CR-logical or mfCR	3	
CR-logical	branch	3	

1. Normally permutes and VMX-simple operations have a latency of two cycles. However, if a 2-cycle operation is issued five cycles after a 7-cycle operation on the same VSU pipe, the 2-cycle operation is dynamically converted to a 7-cycle operation to prevent a resulting bus collision.
 2. This applies to pipelined DFP operations only. It applies to only the DFP IOP of the cracked DFP instruction.
 3. Latency of four cycles if VSU pipe 0(1) to FXU pipe 0(1). Otherwise five cycles.
 4. See rules about back-to-back-to-back (*Section 10.1.14.9 Limitations of Back-to-Back* on page 219).
 5. Minimum two cycles issue-to-issue. Subject to special rules: **cmp** and **bc** must be in the same dispatch group, or else the **cmp** must be the most recent CR writer.

10.1.15 Pipeline Hazards

10.1.15.1 ISU Rejects

The issue queue can reject instructions for the following types of conflicts and speculation. Reject penalties are shown in parenthesis.

- Result bus arbitration. (5)
- Finish port arbitration. (5)
- Other resources internal to the execution units. (5)
 - FX and VS dividers (5)
 - DFP, crypto units (5)
 - SPR bus (5)
 - LSU prefetcher utilizes Agen logic (5)
- Instructions need to request the SAR bypass. (7 - 10)
- Instructions need to wait for a swap to be processed before they can issue. (7)
- Instructions were speculatively issued before their source operands were available. (6)

Conflicts Due to Write-Back Collisions and VSU/SPR Resources

The following resource conflicts result in a 5-cycle rejection loop:

FXU pipe:

- FX 2-cycle instruction followed by a 1-cycle FX instruction. The 1-cycle instruction gets rejected.
- FX 4-cycle instruction followed by one bubble followed by an FX 2-cycle instruction. The 2-cycle instruction gets rejected.
- FX 4-cycle instruction followed by two bubbles followed by an FX 1-cycle instruction. The 1-cycle instruction gets rejected. (This case is frequently, but not always, prevented by blocking 1-cycle FX instructions within UniQ from issuing three cycles after a multiply issues. They can still be issued off the dispatch bypass.)
- FX divide followed by another FX divide. The divide continues to issue and reject. (The next divide can issue three cycles before WB of the first).
- The CRQ issues an **mfCTR/LR/CR** that has a write-back collision with an add issued in the previous cycle on the same thread set or trying to use the same GPR write port. The **add** gets rejected.
- A multiply or VSU-extract collides with **mfCTR/LR/CR** on the same thread set or pipe: the multiply/extract wins and **mfspr** is not issued.
- An **mfCTR/LR/CR** collides with a 2-cycle operation; **mfspr** has priority.
- Any 4, 6, and 10-cycle **mfspr** colliding with 1, 2, 4, variable-length FX operations and VSU extracts. Depends on latency and which can be cancelled.
- A VSU extract is issued at the same time as a multiply on the same queue half. Only one can be issued. Toggle LFSR to determine priority.
- A VSU extract is issued at the same time or up to four cycles before a FX multicycle operation. The multi-cycle operation is rejected.



- A VSU extract is issued causing WB collisions with FX 1- or 2-cycle operations. The FX operation is rejected.
- In SMT-shared mode, the GPR SAR bypass must arbitrate with FX instructions.
 - For 1- and 2-cycle FX instructions, priority between the issuing instruction and SAR bypass is determined by an LFSR that can be controlled to give priority to the SAR bypass 25, 50, 75, or 87% of the time. The default recommendation is to leave this at 75%.
 - For multiplies, the SAR bypass has priority if both of the following are true:
 - (1) The LFSR above indicates that the SAR bypass should win over the FX operation.
 - (2) Previous SAR bypass attempts have been killed.

Note: This means that if there are many multiplies attempting to do a SAR bypass, it takes a long time to make forward progress because the SAR bypass must have failed once before.

- Divides always have priority over SAR bypass. The SAR bypass is killed if it collides with a divide WB.
- If two **mfspres** or **mtspres** are issued at the same time on any FXU or CRU issue port, one must be cancelled. A priority LFSR determines which is cancelled. (This is an SPR bus, not a GPR write port limitation.)

VSU pipe:

- Write-back contention between the following classes of pipelined instructions: 13-cycle (DFP), 7-cycle (FD/FX/VD/VS/VX/CY), 2-cycle (XS/PM/SD/SQ).

Note: If a 2-cycle operation is issued five cycles after a 7-cycle operation, the 2-cycle operation is converted to a 7-cycle operation. This means it has an issue-to-issue latency of seven cycles, and finishes at the same time as a 7-cycle operation.

- CY issued to both pipes; one must be rejected (LFSR determines priority).
- DFP issued to both pipes; one must be rejected (LFSR determines priority).
- Swap colliding with a VSX load in SMT-shared mode. The load wins and the swap is killed. The instruction dependent on the swap will issue and request to use SAR bypass.
- For up to three cycles after issuing a multi-cycle instruction, an attempt is made to issue any ready VS instructions and then reject them.

LSU/LU pipe:

- LU/LSU uses a pipe for the D-ERAT reload; an operation issued at the same time is rejected.
- LU/LSU pipe is used by data prefetch. (LSU sends back a “steal agen FX only” signal.) A simple-FX operation issued at the same time is canceled or rejected
- LU only: A resulting bus conflict due to L2 data coming back for a GPR or VSX load; load issued to the LU at the same time is rejected.
- LSU sends back XER results for **stcx**. Conflicting instruction on the LSU pipe gets a 6-cycle reject.
- Store forwarding: (LU/LSU sends a “delay5” indication) conflicting instruction (issued five cycles after the load/**mfspres** getting store forwarding) gets a reject.
- Same operation issued to both LSU and FXU (for FX->LS); no rejecting is required in this case; the LSU issue is cancelled and the operation is issued to the FXU.
- An **mtspres** and **mfspres** issued to both LSU pipes at the same time; one must be rejected.

- In SMT-shared mode, loads issued to the LSU will block swaps, causing them to be cancelled.

SAR Bypass - Related Rejects

Instructions can be issued even if they do not have valid source rtags (that is, it has a “needs rtag” indicator), but their source operand data is in the SAR. When this happens, the instructions request a SAR bypass. Due to the time it takes to read data out of the SAR, the instructions requesting a SAR bypass are rejected by the issue queue and re-issued after the data has been placed in bypass latches.

- If an instruction is issued with a “needs rtag” bit for a source operand, and the GPR data mux was not configured to use the SAR bypass, then the instruction requests a SAR bypass and gets a reject.
 - If it wins arbitration of the SAR bypass, it gets an 8, 9, or 10-cycle reject, depending on if the instruction needs 1, 2, or 3 SAR bypasses.
 - If the instruction does not win SAR bypass arbitration, it gets a 7-cycle reject.
- If the front end of the issue queue thinks the instruction must get data out of the SAR (that is, the instruction is issued with a “needs rtag” bit on) and the GPR SAR bypass controls were set up correctly but the data was not read out of the SAR due to a SAR bank collision, or if the wrong data was read out of the SAR, the instruction must get a 6-cycle reject. On the FX/LS/LU pipes, this can happen if an instruction was flushed, and another one with the same queue position (qpos) was dispatched in its place, and a bypass latch was reserved for that qpos. This case does not apply to the VS pipe.
- If an instruction is *not* issued with a “needs rtag”, but the GPR source mux had been configured to get data from the SAR, the instruction must get a 5/6-cycle reject. Like the previous case, this one only happens if there is a false qpos match due to quick qpos reuse, and only applies to the FX, LS, and LU pipes.
- If an instruction requests SAR bypass and then successfully uses it or gets flushed, and a second instruction is dispatched into the same qpos, a false SAR bypass can be driven if the SAR bypass latch contents are no longer valid.
- If an FX or VS instruction gets both a resource conflict and a SAR reject, it has a reject latency according to the SAR reject (7 - 10 cycles).
- If an LS or LU instruction gets both a resource conflict and a SAR reject, it has a reject latency according to the resource conflict (5 cycles).

AMC Castout-Related Rejects

If an instruction’s source was recently cast out of the [AMC](#), because either a completion castout or a swap victim castout, there is a window of time where the instruction will get rejected. This mechanism is in place because, if a given rtag X is used in a swap, it is not possible to determine if instructions with source rtag X are dependent on the swap victim or the swap grantee. For VRF castouts, there is a 4-cycle window where dependents of the rtag being cast out get rejected. For GPR castouts, the window is 6 cycles.

AMC/SAR Swap-Related Rejects

The front end of the UniQ can attempt to issue instructions that are waiting for swaps if they are issued in the disp+3, disp+4, or disp+5 cycles. These instructions must be rejected, because if they were not, they can complete before the swap was processed. This can potentially cause errors in the AMC. Therefore, at issue time, check to see if the issued instructions should be waiting for swaps. If so, they get a 5-cycle reject.



XER ARF Related Rejects

If an instruction with an XER source has its data in the XER ARE, but the XER source rtag of that instruction matched one of the XER destination rtags of recently issued instructions, the instruction with the XER source will be rejected. The reason for this is that it is not known until late in the ISS+1 cycle if the XER data resides in the ROB or the ARF. The DL bits and XER bypass controls are dependent on this information. They are computed as if the false rtag matches are true dependences.

FPSCR Related Rejects

VS instructions reading a renamed FPSCR can get a source-not-ready reject if the FPSCR producer has not successfully issued.

An FPSCR consumer also gets a reject if the producer completed in the window of time between when the consumer was dispatched and issued. This is because the location of the FPSCR data (ARE or ROB) is not known until ISS+2.

10.1.15.2 LSU Rejects

The LSU detects several internal conditions that cause it to reject an instruction on a given port. When the LSU rejects an instruction, it causes a 10-cycle turn-around for the instruction to be re-issued. A list of conditions that cause an instruction to be rejected follows:

- SRQ - load-hit-store same cycle (store is older): The younger load is rejected.
- SRQ - load-hit-store (different group, store forwarding not possible): Load is rejected if the data needed by the load is not fully contained in the store data queue (SDQ) for the store instruction.
- SRQ - load-hit-store (data not available): Load is rejected because the data to be stored by the store instruction is not yet available in the store data queue.
- SRQ - load-hit-sync: Load is younger and therefore rejected because **sync** has to complete first.
- SRQ - load-hit-DCB (**dcbz**, **cidcbf**): Load is younger and hits the cache line targeted by the DCB instruction but is rejected because the DCB instruction must complete (**dcbz** could have forwarded '0', but the POWER8 processor core does not do that).
- SRQ - **Iwarx/Idarx** not next-to-complete: If a previous **Iwarx/Idarx** instruction sets the reservation, a successive **Iwarx** is rejected unless it is next-to-complete. This prevents excessive loss of reservation and still guarantees forward progress.
- Global - force reject for hang detect: When the pervasive unit sends the hang detect signal, all load/store instructions are rejected by the LSU.

10.1.15.3 Flush Conditions

Pipeline flushes are the last and most expensive means of reordering instructions in the out-of-order section of the processor. Instructions that have not completed are eligible to be flushed. Most flushes are a result of late detection of an architecturally unsafe ordering of instructions. Others are rare conditions that are most easily solved by flushing when they occur yet allow the processor to achieve higher performance under the more frequent circumstances. *Table 10-14* lists all of the flush conditions.

Table 10-14. Flush Conditions

Flush Condition	Special Considerations	Programming Ramifications
Cache inhibited and not single instruction group or cache inhibited per page table I bit.	Force to single instruction group and wait until next-to-complete to make nonspeculative per architecture.	Avoid in performance critical code.
Guarded storage, L1 cache miss and not single instruction group.	Force to single instruction group and wait until next-to-complete to make nonspeculative per architecture.	Avoid in performance critical code.
Unaligned	Unaligned references not handled directly in hardware need to reference microcode. Instruction first flushed to single instruction group and then flushed to microcode routine entry point.	Avoid unaligned data in performance critical regions. See <i>Section 10.1.16.1 Storage Alignment</i> on page 228 for unaligned cases not handled directly by hardware.
ECC error on reload data	Data forwarded assuming no error. Load must be restarted.	None
Invalidate-hit-reload	Store or invalidate while load miss outstanding to the same address. Load must be restarted.	Use LARX/STCX for shared data in performance critical regions.
Load-hit-store EA alias	Because EA(44:51) is different, hardware does detect. Flush until store drains.	Avoid EA aliasing in shared memory.
Load-hit-store, same group, cannot forward	Load data not contained in single-store data-queue entry.	Code-specific optimizations. Can detect and replace by register move if performance limiting.
Store-hit-load	Issued out-of-order, older store to same address as younger load. Flush after store to restart load.	SHL avoidance in hardware when RA fields match. See <i>Section 10.1.5 Store-Hit-Load Avoidance Table</i> on page 207.
Load-hit-load with snoop	Issued out-of-order, older load to same address after intermediate snoop after younger load.	Use LARX/STCX for shared data in performance critical regions.
LARX-hit-LARX	Issued out-of-order, older LARX detects younger LARX. Only one LARX per thread can be outstanding at a time. Need to flush younger LARX.	On a fail, software must spin on a load until the value changes.
LRQ atomic	First or second part of an LQ or LQARX instruction has been invalidated. Architecturally required to be atomic.	Use LARX/STCX for shared data in performance critical regions.
RMSC flush	HV = '1', DR = '0', not single instruction and ERAT miss. Force to single instruction group and wait until next-to-complete to make nonspeculative per architecture.	Avoid in performance critical code.
sync flush for speculative load that has a matching snoop	Performance optimization that allows speculation beyond a sync . Flush is used to recover when speculation must be backed out.	None.
TM load with cache line crossing where the second address cannot be sent to the L2.	Performance optimization allowing cache line crossings to be faster. Because both addresses must be seen by the L2 if both are L1 cache hits, anything blocking the second address from making it to the L2 cache must cause a flush.	Avoid unaligned cache line crossings in performance sensitive code.



10.1.16 Level-1 Data Cache

10.1.16.1 Storage Alignment

The LSU performs most loads and stores that are unaligned with the same timing as naturally aligned loads and stores with some exceptions (see *Section 3.5.5.4 Storage Access Ordering* on page 76). When these cases occur, a misaligned flush is generated, which causes a refetch of this instruction, and the microcode unit generates multiple instructions that will now cross the boundary without flush.

For simple loads (not string instructions), the group with the unaligned load or store is flushed and re-executed with microcode that avoids the boundary crossing. If the load or store was the first instruction in the group, only one flush is needed. If not the first instruction in a group, the first flush causes the instruction to come back in a single instruction group, and a second flush is needed for the misaligned flush.

String instructions are already microcoded as single instruction groups; thus only one flush is needed.

Misaligned flushes tend to be around 30 cycles each for the best case.

Special Load Cache Line Crossing Cases

If a load instruction crosses a cache line with a data access and both cache lines are in the L1 data cache, it is highly likely that the access will have only a 5-cycle penalty over a normal L1 cache access. In the POWER8 core, all loads that are allowed to cross a cache line can get this treatment with any byte alignment. The implementation saves the bytes from the first cache line access. Then five cycles later, it accesses the second cache line and merges data from the first cache line with the data from the second cache line. When this occurs, an issues slot is taken for the second access.

There is one hazard that exists between 16-byte VSX loads on LUx pipes and a load on an LSx pipe, where the x's are the same (meaning the same pipe half) when these two instructions occur in the same address generation cycle. If the sum of the bytes from the first cache-line access of the VSX instruction and the bytes from the first cache-line access of the load¹ are greater than 15, the load in LSx is rejected. This is because there are not enough resources to save more than 15 bytes for each pipe half.

10.1.16.2 Special Case of Store Crossing a 64-Byte Boundary

If a store crosses a 64-byte boundary, when the store drains from the store queue, a 1-cycle penalty occurs compared to normal store drains. This is because it takes two write cycles to write across a 64-byte boundary as compared to one write cycle for all other stores.

10.1.17 Level-1 Data ERAT

In the POWER8 core, there are two logical primary D-ERATs implemented as four physical D-ERATs. There is one physical ERAT for each of the two load-store units and one physical ERAT for each of the two load units. Logically, the two primary D-ERATs for LSU0 and LU0 always contain the same data and the D-ERATs for LSU1 and LU1 always contain the same data. Each primary D-ERAT is implemented as a fully associative 48-entry array, with modified binary LRU replacement algorithm. D-ERAT entries are created for 4 KB, 64 KB, or 16 MB pages only. 16 GB pages are broken into 16 MB pages in the D-ERAT, where the installed page contains the referenced address. Similarly, 1 MB pages are broken into 64 KB pages where the installed page contains the referenced address. The four individual ERATs operate in one of two modes: shared or

1. This must be an integer load and thus cannot be bigger than 7 bytes in this case.

split. In single-thread mode and SMT2, the contents of all four D-ERATs are identical; this is referred to as shared mode. But in SMT4 or SMT8 mode, two ERATs (one LSU and one LU) contain translation for half of the active thread (two threads for SMT4 mode and four threads for SMT8 mode). In split mode, the two paired ERATs both contain addresses that can be different from the other paired ERATs. In split mode for the POWER8 core, the two ERAT pairs can be loaded at the same time (with different data) for the following cases:

1. Both ERAT pairs can be loaded from the secondary ERAT.
2. One ERAT pair can be loaded from the secondary ERAT, and one ERAT pair can be loaded from the TLB.
3. One ERAT pair can be loaded from the secondary ERAT, and one ERAT pair can be loaded from memory (tablewalk data).

The POWER8 D-ERAT supports hit-under-miss. Up to four D-ERAT misses are entered in the ERAT miss queue (EMQ) per cycle, one per LSU pipe. There are eight total EMQ entries. These eight entries arbitrate for access to the secondary ERAT. The secondary ERAT contains 256 entries managed in two halves, with a round-robin LRU algorithm. An instruction that misses the primary ERAT but hits the secondary ERAT receives a reject on its initial ERAT lookup and a hit on its second lookup. Similar to the primary ERAT, the secondary ERATs operates in either split or shared mode. In split mode, each of the halves of the secondary ERAT is dedicated to half of the threads. The two halves for the secondary ERATs correspond to the ERAT pairs in the primary ERAT. The secondary ERAT can process two misses concurrently in split mode but only one miss for shared mode. If an address misses the secondary ERAT, it is sent to the SLB and TLB. The EMQ entries are the only eight D-ERAT misses that are sent to the TLB, others result in a reject. These eight entries and two I-side misses arbitrate for the SLB and TLB sequentially. If there is a hit in both, the information for the D-ERAT is returned to both the primary and secondary ERATs. TLB hits are handled speculatively, and the I-ERAT and D-ERAT are loaded speculatively in this case.

10.1.18 Level-2 Data ERAT

The secondary D-ERAT is the second level of non-instruction translation caching in the core. When a data reference misses the primary ERAT, it looks up the address translation in the secondary D-ERAT. If the translation is found in the secondary D-ERAT, it is then loaded into the primary D-ERAT. If the translation is not found in either the primary or the secondary D-ERAT, the request then has to check the SLB and TLB.

The best-case reload time for a hit in the secondary D-ERAT is to have the translation present in the primary D-ERAT nine cycles after the initial look-up. This best-case timing is fast enough to ensure that the translation is installed in the primary D-ERAT before the instruction is re-issued by the issue queue, because the fastest time that the issue queue can re-issue an instruction is greater than nine cycles. When an instruction misses both the primary and secondary D-ERAT, the translation is installed in both. However, because the replacement policies are managed independently, it is possible for a translation to be aged out of the secondary D-ERAT while still remaining in the primary D-ERAT.

The secondary D-ERAT consists of four 64-entry fully associative CAM arrays, for a total of 256 entries. In single thread mode, all 256 entries are available for that thread. In split SMT mode, the secondary D-ERAT is treated as two 128-entry arrays. The secondary D-ERAT replacement policy is a simple FIFO scheme.



10.1.19 Translation Look-Aside Buffer

The TLB has 2048 entries and is four-way associative. It uses a true LRU replacement algorithm. The TLBs are indexed with a hashed address calculated from portions of the virtual address and the page size. Each entry in the TLB represents a particular page size: 4 KB, 64 KB, 1 MB, 16 MB, and 16 GB (that is, all page sizes are natively supported in the TLB). Server software does not have any plan to exploit the 1 MB page sizes; however, to the processor, 1 MB pages are always enabled.

The real address is returned to both the primary and secondary D-ERAT on a TLB hit. The primary and secondary ERATs have different LRU algorithms; therefore, the secondary ERAT is not guaranteed to contain the same address translations contained in the primary ERATs. The TLB contains entries for both the I-cache and D-cache, while the secondary ERAT only contains entries for the D-cache. The I-cache and D-cache maintain separate primary ERATs. Therefore, the D-side ERAT contains entries only for the D-cache, and similarly the I-side ERAT contains entries only for the I-cache. A TLB hit reloads the primary (and secondary) ERAT after two rejects. A following ERAT lookup then hits in the primary ERAT.

If the TLB request does not hit in the TLB, a tablewalk is initiated that loads the TLB, secondary ERAT, and primary ERAT with the translation for the instruction that generated the TLB miss. Up to four outstanding tablewalks can be active at one time. The implementation allows tablewalks for speculative instructions but does not update the T or C bit in a PTE entry unless the instruction is NTC when the PTE entry is found. The TLB is reloaded with the corresponding PTE entry even if the instruction is speculative.

10.1.20 Load Miss Queue

The Load Miss Queue (LMQ) is a 16-deep queue that handles all data cache misses for the LSU, including data prefetches to the L1 cache. These 16 entries are shared by all eight threads on this core. The queue is 4-ported; thus, it can take a miss from each one of the four load/store pipes, and can pass those four requests to the L2 cache in one cycle as well. The two LU ports (load only ports) have a fast path to the L2 cache. Even then, if two demand loads occur in the same cycle on the LUs, only one can win, and the other is placed in the L2 queue, causing a minimum of a 3-cycle penalty. The LS ports, take a minimum 3-cycle penalty to L2 latency, because they must be queued before they can enter the L2 pipe.

Data is returned from the nest in two 64-byte beats or four 32-byte beats. The critical beat (the beat with the requested data) is favored to return first. Normally, two 64-byte beats are returned back-to-back if the line was in the L2 cache and two 64-byte beats from a line that was in the L3 cache, return with a bubble (cycle) in between. If the data came from the power bus, the data returns in four 32-byte beats, with at least one bubble in between each beat.

Merging into the LMQ is not allowed for loads. In other words, if two loads for the same line are executed, the first gets a miss queue entry, and the second one is rejected until the second load hits the cache after the data has been refilled from the nest. One exception is that a load can always merge with an L1 prefetch if they are for the same line.

If a flush occurs after a load has made an L2 request, the load does not forward data, but the miss queue remains busy until all data is returned and placed into the cache.

Cache-inhibited loads are handled by the LMQ as well. Although they do not write to the cache, they still occupy an entry until data is returned.

If a store with the same real address as a load in the LMQ occurs after the load request but before the data is back, the LMQ does not allow that data to be written into the cache.



If a data cache miss request finds the data in memory, a special tag is returned indicating that the [ECC](#) check has not been done on this beat of data. The memory controller reserved the right to cancel this line on the last beat of data if the ECC checker finds an error. This allows the memory controller to source data much earlier because it does not have to go through the ECC checker first. The LMQ must mark any load that uses this line as speculative when that load finishes. The global completion table logic must not complete these loads until the LMQ indicates that the data is good.

10.1.21 Transactional Memory

The L2 cache is largely responsible for tracking the transactional memory (TM) footprint. The TM footprint consists of 128-byte cache lines that have been accessed by either load or store instructions while the thread is executing in the transactional state. Generally in rollback-only-transactions (ROTs), only stores are tracked. However, due to implementation-specific reasons, a small fraction of loads in an ROT can also be included in the tracking footprint. The footprint is tracked by four banks of 16-entry CAMs, which are shared by all threads on a given chiplet. A footprint overflow type of transaction failure results if either the CAM or the L2 congruence class capacity is exceeded. If either a transactional or non-transactional access initiated by another thread conflicts with a given thread's footprint, the transaction also fails and the appropriate type of conflict is reported in the TEXASR per the Power ISA.

If a **tlbie** request from any thread in the same LPAR hits a page used in a transaction, the transaction is required to fail. Rather than exactly comparing all TM pages to a **tlbie**, a bloom filter¹ approach is used.

10.1.22 Store Queue and Store Forwarding

The LSU contains a 40-entry store reorder queue (SRQ) that holds real addresses and a 40-entry store data queue (SDQ) that holds a quadword of data. These store queues are dynamically shared among the available threads (in both SMT2, SMT4, and SMT8 modes). In the POWER8 core, each SDQ entry can hold one store operation, which can be up to 16 bytes wide (for VSX and VMX stores, or a **stq**, **stqcx**.).

Store addresses are loaded into the SRQ when the instruction is issued. These can be issued in any order. Fixed-point data is loaded into the SDQ from the FXU after the data is accessed from the GPR (`st_data` is transferred via the LU0 or LU1 RB operand bus). Floating-point data is loaded into the SDQ from a shared 16-byte data bus from the FPU after the data is accessed from the FPR. This shared 16-byte data bus also transfers VSX/VMX store data to the SDQ from a dedicated bus from the FPU after the data is accessed from the FPR. In the POWER8 core, there are two dedicated 16-byte data buses to transfer FPU/VSX/VMX data to the SDQ. Stores are removed from the SRQ and SDQ and written to the cache in program order after all the previous instructions are committed.

Loads that are issued, which hit (address match) stores in the SRQ that have not been written to the cache, are candidates for store forwarding. A store forward can occur under the following conditions:

- The data is available in the SDQ.
- The load is completely contained within the store (load byte count is less than or equal to the store byte count).
- The store is older than the load.
- The page-table I-bit is not set.
- No collision with an L1 reload operation (AGEN steal).

1. This can cause a transaction to fail when it did not need to, but greatly reduces the overhead of detection for what is a rare event.



There are some exceptions to these rules; for example, see *Section 10.5.1.1 Store Quadword* on page 254 for more details. If the above conditions are not met, the load instruction can be either rejected or flushed. A load that becomes a store forward operation is treated as a miss from the standpoint of the issue logic. The latency of a store forward operation is five cycles.

In addition to stores, **sync**, **lwsync**, **ptesync**, **eieio**, **dcbz**, **dcbf**, **icbi**, **tlbsync**, and **tlbie** (non-local) are installed in the SRQ. Additionally, on the POWER8 core, the following instructions will install in the SRQ: **isync**, **tm_begin**, **tm_check**, **tm_end**, and **logmpp**.

10.1.22.1 Stores in Real Mode (MSR[DR] = 0)

For store instructions, the store re-order queue above the caches is used as a temporary holding spot for both the address and the data. When the store passes the completion point, the store queue is marked to allow the store data to be written into any of the appropriate caches. If the real-mode I-bit (located in the HID4 Register) is set, cache misses are expected, so that the store simply works its way towards memory or memory-mapped I/O. Note that, as a result, the store action is only observed once per store instruction.

10.1.23 Data Prefetch

The purpose of the data prefetch mechanism is to reduce the negative performance impact of increasing memory latencies, particularly on technical workloads. These programs often access memory in regular, sequential patterns. Their working sets are also so large that they often do not fit into the limited size L1, L2, and L3 caches used in the POWER8 chip. For additional information, see *Table 10-16 Instruction Latencies and Throughputs* on page 236.

Designed into the load-store unit, the prefetch engine can recognize sequentially increasing or decreasing accesses to adjacent cache lines and then request anticipated lines from more distant levels of the cache/memory hierarchy. The usefulness of these prefetches is reinforced as repeated demand references are made along such a path or stream. The depth of prefetch is then increased until enough lines are being brought into L1, L2, and L3 cache that much or all of the load latency can be hidden. The most urgently needed lines are prefetched into the nearest cache levels. During stream start up, several lines ahead of the current demand reference can be requested from the memory subsystem. After steady state is achieved, each stream confirmation causes the engine to bring one additional line into the L1 cache, one additional line into the L2 cache, and one additional line into the L3 cache. To effectively hide the latency of the memory access while minimizing the potentially detrimental effects of prefetching such as cache pollution, the requests are staged such that the line that is being brought into the L3 cache is typically several lines ahead of the one being brought into the L1 cache. Because the L3 cache is much larger than the L1 cache, it can tolerate the most speculative requests more easily than the L1 cache can.

A basic prefetch start-up sequence follows.

Prefetch begins by saving the effective address of the L1 D-cache misses in a 16-entry queue, offset up or down by one line address (the initial miss to line n triggers the allocation of an entry for line $n+1$). This initial allocation is managed on a pseudo-LRU basis, and causes a request for line $n+1$ to be brought into the L3 cache. A subsequent demand L1 lookup for line $n+1$ matches the newly allocated entry (referred to here as a confirmation). This confirmation triggers an update to the existing queue entry, which now points to line $n+2$. In addition, the confirmation triggers a request for lines $n+2,3,4,5$ to be brought into the L3 cache. A request is also made for line $n+2$ to be brought into the L1 and L2 cache. Upon the next confirmation (demand L1 lookup for line $n+2$), the queue entry is again updated, this time to point to line $n+3$. This confirmation triggers requests for lines $n+6,7,8,9,10$ to be brought into the L3 cache. Line $n+3$ is also brought into the L1 and L2 cache. A third confirmation (demand L1 lookup for line $n+3$) triggers requests



for lines $n+11, 12, 13, 14, 15$ to be brought into the L3 cache, as well as requests for lines $n+4, 5$ to be brought into the L1 and L2 cache.

At this point, the L3 prefetch engine is prefetching 12 lines ahead of the current demand load. The L1 prefetch engine is prefetching two lines ahead of the demand load, and the stream has reached steady state. Subsequent confirmations do not cause the L3 engine to go any further ahead than 12 lines or the L1 engine to go any further ahead than two lines. A fourth confirmation (demand L1 lookup for line $n+4$) triggers a request for line $n+16$ to be brought into the L3 cache, as well as a request for line $n+6$ to be brought into the L1 and L2 cache. Because the L3 prefetch engine is now in steady state, this request for line $n+16$ can happen immediately or it can be queued up to take advantage of efficiencies in the memory subsystem. Either way, each subsequent confirmation eventually triggers a request for one additional line to be brought into the caches.

Although this description represents a typical stream ramp profile, the maximum depth of the stream and the urgency with which the prefetch engine attains that depth are fully programmable via the Data Stream Control Register (DSCR). The DSCR is described in Book II of the Power ISA.

The direction of the stream is always assumed to be increasing in magnitude. However, during stream start-up, the prefetch engine simultaneously examines streams for both increasing and decreasing patterns using a “shadow” queue. Subsequent references can either confirm the increasing direction or the decreasing direction. If the stream is confirmed by a decreasing pattern, the direction of the stream is marked as such from that point forward.

In addition to the stream detection described previously, the prefetch engine is augmented by a stride-N detection engine. The purpose of the stride-N engine is to detect streams that have regular access patterns, but which do not fetch from consecutive cache lines in memory. The hardware can detect strides up to 8 KB in length, with a 32-byte granularity. The detection is handled in a 4-entry buffer that examines the stride between the current cache miss and the previous four cache misses. When a pattern is detected, a stream is created in the regular prefetch queue. From this point forward, the stream is treated just like any other stream, with the difference being subsequent prefetch requests can fetch from nonconsecutive cache lines.

Prefetch streams are tracked by the effective address and are allowed to cross small and medium memory page boundaries, but will be invalidated when crossing a large-page boundary. All prefetch requests must therefore go through address translation before being sent to the memory subsystem. When address translation is not found in the ERAT for a prefetch request, the prefetch initiates an ERAT miss request. This has the effect of prefetching not only the data that is needed by the program, but also the address translation. If any type of exception is encountered while performing the translation lookup for the prefetch, the request is dropped and the requesting stream is invalidated. Streams otherwise remain active until replaced with a new address by the allocation mechanism, or until the thread that allocated the stream is no longer running on the processor.

In all cases, a prefetch request is completed when the lines are found already resident in or have been returned to the target cache level. If a demand miss “catches up with” an outstanding prefetch request, it is either merged or rejected depending on the level in the cache hierarchy. If it is rejected, it is retried until it hits in the cache.



10.2 Chiplet

10.2.1 Level-2 Cache

Each L2 cache on the POWER8 chip is a unified cache for its respective core. The L2 cache is an 8-way associative 512 KB cache with fast access to its own private 8 MB L3 cache region through a private low latency bus. The L2 cache maintains full hardware coherence within the system and can supply intervention data to the other cores on this POWER8 chip or to other cores on other POWER8 chips. Logically, the L2 cache is an in-line cache. Unlike the L1 caches, which are store-through, it is a store-in cache. It is fully inclusive of the L1 D-caches and the L1 I-caches.

Each L2 cache is comprised of 512 associative sets (called congruence classes); each congruence class contains eight 128-byte cache lines. Real address bits 48:56 are used in the 9-bit congruence class address.

The L1 D-cache is a store-through design. As such, when a cacheable store is removed from the SRQ, if it hits in the L1 D-cache, the line is updated. In addition, all stores (hit or miss) are forwarded on to the L2 cache. If the store is a miss in the L1 D-cache, the data is not written to the D-cache and the line is not reloaded from the L2 cache (no L1 D-cache allocate on store miss).

10.2.2 Level-3 Cache

Each L3 cache region on the POWER8 chip is a unified victim cache for its respective core/L2 cache, as well as for other L3 caches on chip. The resident cache lines installed from the attached L2 cache are referred to as L3.0 lines, and the resident cache lines installed from other L3s are referred to as L3.1 lines. When castout from this L3 cache, L3.1 lines go to memory and L3.0 lines have the option of being castout to other L3 caches on chip. The L3 cache is an 8-way associative 8 MB cache. The L3 cache maintains full hardware coherence within the system and can supply intervention data to the other cores on this POWER8 chip or to other cores on other POWER8 chips. Logically, the L3 cache is an in-line cache. The L3 cache is a victim cache of the L2 cache (that is, all valid lines that are victimized in the L2 cache are castout to the L3 cache). The L3 cache is not inclusive of the L2 cache. Each L3 cache is comprised of 8192 associative sets (called congruence classes), each congruence class contains eight 128-byte cache lines. Real address bits 44:56 are used in the 13-bit congruence class address.

10.3 Latencies

10.3.1 Cache Latencies and Bandwidth

Table 10-15 lists several key bandwidth and latency values for the chip. These represent best case values under ideal conditions. Actual values can be somewhat higher due to resource limitations or queueing effects. Still, these values are useful in understanding system performance. The *pclk* in *Table 10-15* is the clock rate of the processor.

Table 10-15. Cache Latencies and Bandwidth

Description	Latency	Bandwidth
L2 D-cache load hit (bypass)	8.5 pclk	64 bytes/pclk
L2 I-cache load hit (bypass)	9.5 pclk	64 bytes/pclk
L3 load hit	26.5 pclk	32 bytes/pclk
L2.1 load hit	94 pclk	16 bytes/pclk
L3.1 load hit	112 pclk	16 bytes/pclk
L2.5 load hit	325 pclk	
L3.5 load hit	343 pclk	
Memory load chip pump	80 ns	230 GBps (2:1 read:write)
Memory load system pump	140 ns	
A link		25.6 GBps (peak)

Note: Pclk represents one processor clock.



10.3.2 Instructional Latencies and Throughputs

The fixed-point latency, when specified as x-y, is the latency where x is for two instructions on the same thread-set side and y is for two instructions on opposite thread-set sides.

Table 10-16. Instruction Latencies and Throughputs (Sheet 1 of 17)

Instruction	Pipeline	Latency	Throughput (IPC)	Notes decode / dispatch	Other Interlocks	Other Comments
ld ldx ldu ldux lwa lwax lwaux lwz lwzx lwzu lwzux lha lhax lhau lhaux lhz lhzx lhzu lhzux lbz lbzx lbzu lbzux	LSU or LU, FXU	3 cycles (RT) 2 - 4 cycles (RA)	2/cycle	1	N/A	Broken into a basic load, an exts , and an add .
lq	LSU, FXU	3 (RT) 3 (RT+1) 6 (XER)	1/cycle	1	N/A	
stb sth stw std	LSU, LU	N/A	2/cycle	1	no	Stores are internally issued twice (once as an ea_gen operation, and once as a data operation). Stores put their data into the STQ, which later writes to the cache/memory hierarchy (the STQ supports forwarding for many cases).
stbx sthx stwx stdx	LSU, LU	N/A	2/cycle	1	no	Three source operands are cracked to minimize breadth or renaming facility.
stbu sthu stwu stdu	LSU, LU	3 cycles for updated register	2/cycle	1	no	
stbux sthux stwux stdux	LSU, LU	3 cycles for updated register	2/cycle	1	N/A	Cracked into two IOPs.
stq	LSU, LU	N/A	1/cycle	1	N/A	Cracked into three IOPs.
lmw	LSU or LU	N/A	2 register loads per cycle after start-up	1	N/A	Number of dispatch groups is equal to the number of registers modulo three. The microcode generates inline sequence of basic loads.
stmw	LSU, LU	N/A	2 register stores per cycle after startup	1	N/A	Number of dispatch groups is equal to the number of registers modulo three. The microcode generates inline sequence of basic stores.
<ol style="list-style-type: none"> 1. Store_agen IOPs are issued to the LSU, and store_data IOPs are issued to the LU. Fixed-point loads that do not have an XER destination can be issued to either the LU or the LSU. Loads with floating-point destinations can only be issued to the LU, and loads with XER sources can only be issued to the LSU. 2. Branches can issue two cycles after a producing a cmp instruction if they were in the same dispatch group, or the cmp is the youngest instruction writing a CR relative to the cmp. Otherwise, the issue-to-issue latency to a branch is three cycles. 3. Normally, permutes and simple operations have a latency of two cycles. However, if a 2-cycle operation is issued five cycles after a 7-cycle operation on the same VSU pipe, the 2-cycle operation is dynamically converted into a 7-cycle operation to prevent a resultant bus collision (to make use of the issue slot). 						



Table 10-16. Instruction Latencies and Throughputs (Sheet 2 of 17)

Instruction	Pipeline	Latency	Throughput (IPC)	Notes decode / dispatch	Other Interlocks	Other Comments
lswi (naturally aligned)	LSU or LU	N/A	2 register loads per cycle after start-up	1	N/A	Number of dispatch groups is approximately equal to number of registers modulo three. The microcode assumes natural alignment and generates inline sequence of basic loads.
lswx (naturally aligned)	LSU or LU	N/A	2 register loads per cycle after start-up	1	N/A	Number of dispatch groups is approximately equal to number of registers modulo three plus four more for startup. The microcode assumes natural alignment and generates inline sequence of basic loads.
stswi (naturally aligned)	LSU, LU	N/A	2 register stores per cycle after start-up	1	N/A	Number of dispatch groups is approximately equal to number of registers modulo three, Microcode assumes natural alignment and generates inline sequence of basic stores.
stswx (naturally aligned)	LSU, LU	N/A	2 register stores per cycle after start-up	1	N/A	Approximately equal to (number of registers modulo three plus four more for startup. The microcode assumes natural alignment and generates inline sequence of basic stores.
lswi lswx stswi stswx (unaligned)	LSU,LU	N/A	N/A	1	N/A	String instruction is first decoded and dispatched as described previously. At execute, the LSU notes that it is unaligned and causes a <i>machine flush</i> . As the string instruction goes through decode the second time, it is broken up in a way that takes the misalignment into account.
lwarx ldarx	LSU or LU	N/A	N/A	1	no	Forced to miss data L1 cache.

1. Store_agen IOPs are issued to the LSU, and store_data IOPs are issued to the LU. Fixed-point loads that do not have an XER destination can be issued to either the LU or the LSU. Loads with floating-point destinations can only be issued to the LU, and loads with XER sources can only be issued to the LSU.
2. Branches can issue two cycles after a producing a **cmp** instruction if they were in the same dispatch group, or the **cmp** is the youngest instruction writing a CR relative to the **cmp**. Otherwise, the issue-to-issue latency to a branch is three cycles.
3. Normally, permutes and simple operations have a latency of two cycles. However, if a 2-cycle operation is issued five cycles after a 7-cycle operation on the same VSU pipe, the 2-cycle operation is dynamically converted into a 7-cycle operation to prevent a resultant bus collision (to make use of the issue slot).



Table 10-16. Instruction Latencies and Throughputs (Sheet 3 of 17)

Instruction	Pipeline	Latency	Throughput (IPC)	Notes decode / dispatch	Other Interlocks	Other Comments
stwcx. stdcx.	LSU, LU	N/A	N/A	1	no	Must establish coherency block ownership before completing the instruction (other stores do not have to do this). Can take anywhere from 10 cycles to 100's depending on the state of the coherency block in the memory hierarchy.
addi addis add add. subf subf. addic subfic adde addme subfme subfze neg neg. nego	FXU (or LU or LSU for non-dot forms)	1 - 2 cycles (GPR) 2 cycles (XER), 5 cycles (CR)	6/cycle, 2/cycle (with XER or CR updates)		no	
addic. adde. subfe. addme. subfme. addze. subfze.	FXU	1 - 2 cycles (GPR, 5 - 6 cycles (CR)	1/cycle		no	Cracked into a basic add (sub) and a cmp .
addo. subfo. addeo subfeo addmeo subfmeo addzeo subfzeo nego.	FXU	1 - 2 cycles (GPR), 2 cycles (XER), 5 cycles (CR)	1/cycle		no	These operations might architecturally change the summary overflow bit. Summary overflow is written at completion time. If change is detected, everything younger is flushed.
addeo. addmeo. subfmeo. addzeo. subfzeo.	FXU	1 - 2 cycles (GPR), 2 - 3 cycles (XER), 6 - 7 cycles (CR)	2/3 cycles		no	These operations might architecturally change the summary overflow bit. Summary overflow is written at completion time. If change is detected, everything younger is flushed.
addo subfo	FXU	1 - 2 cycles (GPR) 2 cycles (XER)	2/cycle		no	These operations might architecturally change the summary overflow bit. Summary overflow is written at completion time. If change is detected, everything younger is flushed.
addc subfc	FXU	1 - 2 cycles (GPR)	2/cycle		wait for non-rename scoreboard bit to clear	
<ol style="list-style-type: none"> 1. Store_agen IOPs are issued to the LSU, and store_data IOPs are issued to the LU. Fixed-point loads that do not have an XER destination can be issued to either the LU or the LSU. Loads with floating-point destinations can only be issued to the LU, and loads with XER sources can only be issued to the LSU. 2. Branches can issue two cycles after a producing a cmp instruction if they were in the same dispatch group, or the cmp is the youngest instruction writing a CR relative to the cmp. Otherwise, the issue-to-issue latency to a branch is three cycles. 3. Normally, permutes and simple operations have a latency of two cycles. However, if a 2-cycle operation is issued five cycles after a 7-cycle operation on the same VSU pipe, the 2-cycle operation is dynamically converted into a 7-cycle operation to prevent a resultant bus collision (to make use of the issue slot). 						

Table 10-16. Instruction Latencies and Throughputs (Sheet 4 of 17)

Instruction	Pipeline	Latency	Throughput (IPC)	Notes decode / dispatch	Other Interlocks	Other Comments
addc. addco addco. subfc. subfco subfco.	FXU	1 - 2 cycles (GPR) 5 - 6 cycles (CR) 2 cycles (XER)	1/cycle		wait for non-ren name scoreboard bit to clear	
isel	FXU	5 cycles (GPR) 8 cycles (CR)				
mulli	FXU	4-5 cycles	2/cycle		no	Pipelined.
mullw mulhw mulhwu	FXU	4-5 cycles	2/cycle		no	Pipelined.
mulld mulhd mulhdu	FXU	4-5 cycles	2/cycle		no	Pipelined.
mullwo mulldo	FXU	4-5 cycles (GPR) 5 cycles (XER)	2/cycle		no	Pipelined in FXU. These operations might architecturally change the summary overflow bit. The first attempted execution of these instructions assume that the SO-bit does not change. If it does, the instruction causes a flush and then is re-executed.
mullw. mulld. mulhd. mulhw. mulhdu. mulhwu.	FXU	4 - 5 cycles (GPR) 8 cycles (CR)	2/2cycles		no	Pipelined in FXU. Cracked into baseline operation and a cmp .

1. Store_agen IOPs are issued to the LSU, and store_data IOPs are issued to the LU. Fixed-point loads that do not have an XER destination can be issued to either the LU or the LSU. Loads with floating-point destinations can only be issued to the LU, and loads with XER sources can only be issued to the LSU.
 2. Branches can issue two cycles after a producing a **cmp** instruction if they were in the same dispatch group, or the **cmp** is the youngest instruction writing a CR relative to the **cmp**. Otherwise, the issue-to-issue latency to a branch is three cycles.
 3. Normally, permutes and simple operations have a latency of two cycles. However, if a 2-cycle operation is issued five cycles after a 7-cycle operation on the same VSU pipe, the 2-cycle operation is dynamically converted into a 7-cycle operation to prevent a resultant bus collision (to make use of the issue slot).



Table 10-16. Instruction Latencies and Throughputs (Sheet 5 of 17)

Instruction	Pipeline	Latency	Throughput (IPC)	Notes decode / dispatch	Other Interlocks	Other Comments
mullwo. mulldo.	FXU	4 - 5 cycles (GPR) 8 cycles (CR)	2/2cycles		no	Pipelined in FXU. Cracked into baseline operation and a cmp . These operations might architecturally change the summary overflow bit. The first attempted execution of these instructions assumes that the SO-bit does not change. If it does, the instruction causes a flush and then is re-executed.
divd divdu divwe divweu	FXU	12 - 23 cycles (GPR)	2/12 cycles to 2/23 cycles		no	Actual latency depends on data. Non-divides can be issued underneath.
divw divwu	FXU	12 - 15 cycles (GPR)	2/12 cycles to 2/15 cycles		no	Actual latency depends on data. Non-divides can be issued underneath.
divde divdeu	FXU	14 - 41 cycles	2/14 cycles to 2/40 cycles		no	Actual latency depends on data. Non-divides can be issued underneath.
divd. divw. divdu. divwu. divde. divwe. divdeu. divweu.	FXU	same as above for GPR + 6 cycles for CR	same as above		no	Non-divides can be issued underneath.
divdo divwo divduo divwuo divdeo divweo divdeuo divweuo	FXU	same as above for GPR, +1 for XER	same as above		no	These operations might architecturally change the summary overflow bit. The first attempted execution of these instructions assumes that the SO-bit does not change. If it does, the instruction causes a flush and then is re-executed.
<ol style="list-style-type: none"> 1. Store_agen IOPs are issued to the LSU, and store_data IOPs are issued to the LU. Fixed-point loads that do not have an XER destination can be issued to either the LU or the LSU. Loads with floating-point destinations can only be issued to the LU, and loads with XER sources can only be issued to the LSU. 2. Branches can issue two cycles after a producing a cmp instruction if they were in the same dispatch group, or the cmp is the youngest instruction writing a CR relative to the cmp. Otherwise, the issue-to-issue latency to a branch is three cycles. 3. Normally, permutes and simple operations have a latency of two cycles. However, if a 2-cycle operation is issued five cycles after a 7-cycle operation on the same VSU pipe, the 2-cycle operation is dynamically converted into a 7-cycle operation to prevent a resultant bus collision (to make use of the issue slot). 						

Table 10-16. Instruction Latencies and Throughputs (Sheet 6 of 17)

Instruction	Pipeline	Latency	Throughput (IPC)	Notes decode / dispatch	Other Interlocks	Other Comments
divdo. divwo. divduo. divwuo. divdeo. divweo. divdeuo. divweuo.	FXU	same as above for GPR, +6 cycles for CR	same as above		no	These operations might architecturally change the summary overflow bit. The first attempted execution of these instructions assumes that the SO-bit does not change. If it does, the instruction causes a flush and then is re-executed.
cmpi cmp cmpli cmpl	FXU	5 cycles (CR)	2/cycle		no	
tdi twi td tw (XBI field refers to renamed bits in XER)	FXU	N/A	2/cycle		no	
andi. andis. ori oris xorix orixs and and. or or. xor xor. nand nand. nor nor. eqv eqv. andc andc. orc orc.	FXU/LSU	1 - 2 cycles (GPR) 6 cycles (CR)	6/cycle 2/cycle (with XER or CR destination)		no	
ori 0,0,0 (preferred NOP)	none	0 cycles	6/cycle		no	A special form of this instruction is completed at the same time that it is dispatched.
extsb extsb. extsh extsh. extsw extsw.	FXU	1 - 2 cycles (GPR) 4 - 5 cycles (CR)	2/cycle	1 when rc = 1	no	
cntlzd cntlzd. cntlzw cntlzw.	FXU	3 cycles (GPR) 6 - 7 cycles (CR)	2/cycle	1 when rc = 1	no	
rlidcl rlidcl. rlidcr rlidcr. rldic rlidc. rlwinm rlwinm. rldcl rlidcl. rlidcr rlidcr. rlwnm rlwnm. rldimi rldimi.	FXU	1 - 2 cycles (GPR) 5 - 6 cycles (CR)	2/cycle 1/cycle when rc = 1	1 when rc = 1	no	
rlwimi	FXU	1 cycle (GPR)	2/cycle		no	
1. Store_agen IOPs are issued to the LSU, and store_data IOPs are issued to the LU. Fixed-point loads that do not have an XER destination can be issued to either the LU or the LSU. Loads with floating-point destinations can only be issued to the LU, and loads with XER sources can only be issued to the LSU. 2. Branches can issue two cycles after a producing a cmp instruction if they were in the same dispatch group, or the cmp is the youngest instruction writing a CR relative to the cmp . Otherwise, the issue-to-issue latency to a branch is three cycles. 3. Normally, permutes and simple operations have a latency of two cycles. However, if a 2-cycle operation is issued five cycles after a 7-cycle operation on the same VSU pipe, the 2-cycle operation is dynamically converted into a 7-cycle operation to prevent a resultant bus collision (to make use of the issue slot).						



Table 10-16. Instruction Latencies and Throughputs (Sheet 7 of 17)

Instruction	Pipeline	Latency	Throughput (IPC)	Notes decode / dispatch	Other Interlocks	Other Comments
rlwimi.	FXU	1 - 2 cycles (GPR) 5 - 6 cycles (CR)	1/cycle		no	
sld sld. slw slw. srd srd. srw srw. sradi srawi srad sraw	FXU	1 - 2 cycles (GPR) 5 - 6 cycles (CR)	1/cycle	1 when rc = 1	no	
sradi. srawi. srad. sraw.	FXU	1 - 2 cycles (GPR) 5 - 6 cycles (CR)	1/cycle		no	
mtspr(xer)	FXU (one)	N/A	N/A	wait for non-rename scoreboard bit to clear		Can only be done by one of the FXU pipelines.
mtspr(ir) mtspr(ctr)	FXU (one)	5 cycles	1/cycle		no	Can only be done by one of the FXU pipelines.
mtspr(others) mtmsr mtmsrd	Either FXU or LSU (depends on SPR)	varies based on SPR	varies based on SPR		wait for non-rename scoreboard bit to clear	Can only be done by one of the FXU pipelines.
mfsp(xer)	FXU (one)	8 cycles	N/A	wait for non-rename scoreboard and bit to clear		The FXU pipeline is blocked while waiting for SPR value.
mfsp(ir) mfsp(ctr)	CRLogic	5 cycles	1/cycle		no	The FXU pipeline is blocked for one cycle.
mfsp(others) mfmsr	FXU (one) LSU (one), or CR logic	varies based on SPR	varies based on SPR	wait for non-rename scoreboard bit to clear		For FXU, IFU, and pervasive-owned SPRs, the FXU pipeline is blocked while waiting for the SPR value.
mtcrf (one field) mtocrf	FXU (one)	N/A	1/cycle		no	
mtcrf (more than one field)	FXU (one)	N/A	N/A		no	
mfcf	CRlogic	N/A	N/A		no	
mfcrf	CR logic	3 cycles	1/cycle		no	
<ol style="list-style-type: none"> 1. Store_agen IOPs are issued to the LSU, and store_data IOPs are issued to the LU. Fixed-point loads that do not have an XER destination can be issued to either the LU or the LSU. Loads with floating-point destinations can only be issued to the LU, and loads with XER sources can only be issued to the LSU. 2. Branches can issue two cycles after a producing a cmp instruction if they were in the same dispatch group, or the cmp is the youngest instruction writing a CR relative to the cmp. Otherwise, the issue-to-issue latency to a branch is three cycles. 3. Normally, permutes and simple operations have a latency of two cycles. However, if a 2-cycle operation is issued five cycles after a 7-cycle operation on the same VSU pipe, the 2-cycle operation is dynamically converted into a 7-cycle operation to prevent a resultant bus collision (to make use of the issue slot). 						



Table 10-16. Instruction Latencies and Throughputs (Sheet 8 of 17)

Instruction	Pipeline	Latency	Throughput (IPC)	Notes decode / dispatch	Other Interlocks	Other Comments
mftb	FXU (one)	~10 cycles	approximately one per 10 cycles		wait for non-rename scoreboard bit to clear	Use of the FXU pipeline is blocked while waiting for the SPR value.
lfs lfsx lfd lfdx lxdsx lxvd2x lxvw4x lxvdsx	LU	5 cycles	2/cycle	1	no	Instructions are executed in the LU unit. The FPU just arranges renames.
lfsu lfsux lfdu lfdux	LU, FXU	5 cycles (FPR) 1 - 2 cycles (GPR)	2/cycle	1	no	Cracked into normal load and an FXU add .
lvebx lvehx lviewx lvx lvxl	LU	5 cycles	2/cycle	1	no	
stfs stfsx stfd stfdx	LSU, FPU	N/A	2/cycle	1	no	Instructions are dispatched to both the FPU and the LSU units.
stfsu stfsux stfdt stfdtx	LSU, FPU, FXU	3 cycles for GPR	2/cycle	1	no	LSU.
stfiwx	LSU, FPU	N/A	2/cycle	1	no	
stxsdx stxvd2x stxvw4x	LSU, FPU	N/A	2/cycle	1	no	Instructions are dispatched to both the FPU and the LSU units.
stvebx stvehx stvewx stvx stvxl	LSU, FPU, FXU	1 - 2 cycles for GPR	2/cycle	1	no	Cracked into a normal store and a FXU add . Normal store instructions are dispatched to both the FPU and the LSU units.
sync	LSU	N/A	N/A		wait until it is next to complete and for all prior load data to come home	Forces previous stores to finish into the cache/memory hierarchy (out of the STQs).
ptesync	LSU	N/A	N/A		wait until it is next to complete and for all prior load data to come home	Forces previous stores to finish into the cache/memory hierarchy (out of the STQs).

1. Store_agen IOPs are issued to the LSU, and store_data IOPs are issued to the LU. Fixed-point loads that do not have an XER destination can be issued to either the LU or the LSU. Loads with floating-point destinations can only be issued to the LU, and loads with XER sources can only be issued to the LSU.
2. Branches can issue two cycles after a producing a **cmp** instruction if they were in the same dispatch group, or the **cmp** is the youngest instruction writing a CR relative to the **cmp**. Otherwise, the issue-to-issue latency to a branch is three cycles.
3. Normally, permutes and simple operations have a latency of two cycles. However, if a 2-cycle operation is issued five cycles after a 7-cycle operation on the same VSU pipe, the 2-cycle operation is dynamically converted into a 7-cycle operation to prevent a resultant bus collision (to make use of the issue slot).



Table 10-16. Instruction Latencies and Throughputs (Sheet 9 of 17)

Instruction	Pipeline	Latency	Throughput (IPC)	Notes decode / dispatch	Other Interlocks	Other Comments
lwsync	LSU	N/A	N/A		wait until it is next to complete and for all prior load data to come home	Forces previous stores to finish into the cache/memory hierarchy (out of the STQs). Still broadcasts sync transaction onto the SMP interconnect (but does not block).
eieio	LSU	N/A	N/A		wait until it is next to complete and for all prior load data to come home	
isync	LSU	N/A	N/A		wait until it is next to complete and for all prior load data to come home	Causes a flush and instruction refetch only if it is preceded by instructions that change the content of the machine or icbi or pte-sync .
icbi	LSU	N/A	N/A			After the LSU translates the EA, the icbi drains out of the store reorder queue and is finished (that is, no request is sent to the L2 cache because the L2 cache keeps the I-cache coherent).
dcbt dbtst	LSU	N/A	1/cycle			The lead time between a dcbt and a subsequent load is equal to the load-to-use-latency of the level of cache hierarchy the line will be coming from.
dcbz	LSU	N/A	1/cycle			Invalidates L1 cache line on its way to the L2 cache. Allocation and zero function occur at the L2 cache (handled by the L2 cache and treated as any other store).
dcbst	LSU	N/A	1/cycle			
dcbf	LSU	N/A	1/cycle			
<ol style="list-style-type: none"> 1. Store_agen IOPs are issued to the LSU, and store_data IOPs are issued to the LU. Fixed-point loads that do not have an XER destination can be issued to either the LU or the LSU. Loads with floating-point destinations can only be issued to the LU, and loads with XER sources can only be issued to the LSU. 2. Branches can issue two cycles after a producing a cmp instruction if they were in the same dispatch group, or the cmp is the youngest instruction writing a CR relative to the cmp. Otherwise, the issue-to-issue latency to a branch is three cycles. 3. Normally, permutes and simple operations have a latency of two cycles. However, if a 2-cycle operation is issued five cycles after a 7-cycle operation on the same VSU pipe, the 2-cycle operation is dynamically converted into a 7-cycle operation to prevent a resultant bus collision (to make use of the issue slot). 						

Table 10-16. Instruction Latencies and Throughputs (Sheet 10 of 17)

Instruction	Pipeline	Latency	Throughput (IPC)	Notes decode / dispatch	Other Interlocks	Other Comments
sibie	LSU	N/A				Causes index-based invalidate in both the I-ERAT and the D-ERAT.
sibia	LSU	N/A				Fully invalidates the SLB, the I-ERAT, and D-ERAT.
tibie	LSU	N/A				Causes index-based invalidate in both the I-ERAT and the D-ERAT. Is broadcast onto the SMP interconnect.
tibiel	LSU	N/A				Causes index-based invalidate in both the I-ERAT and the D-ERAT. Is not broadcast onto the SMP interconnect.
tlbsync	LSU	N/A				
sibmte	LSU	N/A		wait for non-rename scoreboard bit to clear		Causes selective invalidates out of the I-ERAT and D-ERAT.
sibmfev sibmfee	LSU	N/A		wait for non-rename scoreboard bit to clear		
mtsr mtsrin	LSU			wait for non-rename scoreboard bit to clear		Causes selective invalidates out of the I-ERAT and D-ERAT.
mfsr mfsrin	LSU	3 cycles	1/cycle		No	
mffs mffs. mcrfs	FPU		1/cycle		No	Alone in a dispatch group, completion serialized.
mtfsfi. mtfsf. mtfsb0. mtfsb1.	FPU	3 cycles to dependent FP operation	1/cycle		No	Alone in a dispatch group, completion serialized.
mtfsfi mtfsf mtfsb0 mtfsb1	FPU	3 cycles to dependent FP operation	1/cycle		No	Alone in a dispatch group.
1. Store_agen IOPs are issued to the LSU, and store_data IOPs are issued to the LU. Fixed-point loads that do not have an XER destination can be issued to either the LU or the LSU. Loads with floating-point destinations can only be issued to the LU, and loads with XER sources can only be issued to the LSU. 2. Branches can issue two cycles after a producing a cmp instruction if they were in the same dispatch group, or the cmp is the youngest instruction writing a CR relative to the cmp . Otherwise, the issue-to-issue latency to a branch is three cycles. 3. Normally, permutes and simple operations have a latency of two cycles. However, if a 2-cycle operation is issued five cycles after a 7-cycle operation on the same VSU pipe, the 2-cycle operation is dynamically converted into a 7-cycle operation to prevent a resultant bus collision (to make use of the issue slot).						



Table 10-16. Instruction Latencies and Throughputs (Sheet 11 of 17)

Instruction	Pipeline	Latency	Throughput (IPC)	Notes decode / dispatch	Other Interlocks	Other Comments
fabs fadd fadds fcfd fcfids fcfdi fcfdus fcpsgn fctid fctidu fctiduz fctidz fctiw fctiwu fctiwuz fctiwz fmadd fmadds fmr fmsub fmsubs fmul fmuls fnabs fneg fnmadd fnmadds fnmsub fnmsubs fre fres frim frin frip friz frsp frsq rte frsq rtes fsel fsub fsubs	FPU	6 cycles	2/cycle		No	
fabs. fadd. fadds. fcfd. fcfids. fcfdi. fcfdus. fcpsgn. fctid. fctidu. fctiduz. fctidz. fctiw. fctiwu. fctiwuz. fctiwz. fmadd. fmadds. fmr. fmsub. fmsubs. fmul. fmuls. fnabs. fneg. fnmadd. fnmadds. fnmsub. fnmsubs. fre. fres. frim. frin. frip. friz. frsp. frsq rte. frsq rtes. fsel. fsub. fsubs.	FPU	6 cycles (FPR) 9 cycles (CR)	1/cycle		No	Alone in a dispatch group, completion serialized.
fdiv	FPU	32 cycles	2/26 cycles		No	Makes use of microcoded sequence within the floating-point unit. Free FPU slots can be interleaved with other FPU instructions.
fdivs	FPU	26 cycles	2/20 cycles		No	
fsqrt	FPU	43 cycles	2/37 cycles		No	
fsqrts	FPU	31 cycles	2/25 cycles		No	
fdiv. fdivs. fsqrt. fsqrts.	FPU	same as non-dot forms, +1 cycles for CR	1/2 of non-dot above		No	Alone in a dispatch group, completion serialized. See <i>Table 3-4 Latencies of Floating-Point Divide/Square-Root Instructions</i> on page 68 for additional information.
fcmpu fcmpo ftdiv ftsqrt	Vector simple integer unit (XS)	4 - 9 cycles (to CR)	2/cycle		No	
fmrgew fmrgow	Vector permute unit (PM)	2 cycles	2/cycle		No	Latency to seven cycles if WB conflict.

1. Store_agen IOPs are issued to the LSU, and store_data IOPs are issued to the LU. Fixed-point loads that do not have an XER destination can be issued to either the LU or the LSU. Loads with floating-point destinations can only be issued to the LU, and loads with XER sources can only be issued to the LSU.
2. Branches can issue two cycles after a producing a **cmp** instruction if they were in the same dispatch group, or the **cmp** is the youngest instruction writing a CR relative to the **cmp**. Otherwise, the issue-to-issue latency to a branch is three cycles.
3. Normally, permutes and simple operations have a latency of two cycles. However, if a 2-cycle operation is issued five cycles after a 7-cycle operation on the same VSU pipe, the 2-cycle operation is dynamically converted into a 7-cycle operation to prevent a resultant bus collision (to make use of the issue slot).



Table 10-16. Instruction Latencies and Throughputs (Sheet 12 of 17)

Instruction	Pipeline	Latency	Throughput (IPC)	Notes decode / dispatch	Other Interlocks	Other Comments
<code>xsabsdp xsadddp xscpsgndp xscvdpspn xscvdpsxds xscvdpxws xscvdpxuds xscvdpxws xscvspdp xscvdpspn xscvsxddp xsc-vuxddp xsmaddadp xsmadd-dmdp xsmsubadp xsmsubmdp xsmuldp xnab-sdp xsnegdp xsnmaddadp xsnmaddmdp xsnmsubadp xsnmsubmdp xsrdpi xsrpic xsrpim xsrpib xsrpiz xsredp xrsqrtdp xssubdp</code>	FPU	6 cycles to FPU (+1 cycle to other VSU ops)	2/cycle		No	
<code>xsaddsp xscvdpsp xscvxsdsp xscvuxdsp xmaddasp xmaddmsp xsm-subasp xsmsubmsp xsmulsp xnmaaddasp xnmaaddmsp xsnmsubasp xsnmsubmsp xresp xsrsp xrsqrtesp xssubsp</code>	FPU	6 cycles to FPU (+1 cycle to other VSU ops)	2/cycle		No	
<code>xvabsdp xvaddirp xvcpsgndp xvcvdpsp xvcvdpxds xvcvdpxws xvcvdpxuds xvcvdpxws xvcvspdp xvcvdpxds xvcvpxuds xvcvsxddp xvcvxdsp xvcvsxwdp xvcvuxddp xvc-vuxdsp xvcvuxwdp xvmadd-adp xvmaddmdp xvmuldp xvnab-sdp xvnegdp xvnmaddadp xvnmaddmdp xvnmsubadp xvnmsubmdp xvrdpi xvrpic xvrdpim xvrpib xvrpiz xvredp xvrsqrtdp xvsudp</code>	FPU	6 cycles to FPU (+1 cycle to other VSU ops)	2/cycle		No	
<code>xvabssp xvaddirp xvcpsgnsp xvcvdpxws xvcvpxuds xvcvdpxws xvcvuxwsp xmaddasp xmaddmsp xvm-subasp xvmulsp xvnabssp xvnegsp xvnmaadd-asp xvnmaddmsp xvnmsub-asp xvnmsubmsp xvresp xvrspi xvrspic xvrspim xvr-spip xvrspiz xvrsqrtesp xvsudp</code>	FPU	6 cycles to FPU (+1 cycle to other VSU ops)	2/cycle		No	
<p>1. Store_agen IOPs are issued to the LSU, and store_data IOPs are issued to the LU. Fixed-point loads that do not have an XER destination can be issued to either the LU or the LSU. Loads with floating-point destinations can only be issued to the LU, and loads with XER sources can only be issued to the LSU.</p> <p>2. Branches can issue two cycles after a producing a cmp instruction if they were in the same dispatch group, or the cmp is the youngest instruction writing a CR relative to the cmp. Otherwise, the issue-to-issue latency to a branch is three cycles.</p> <p>3. Normally, permutes and simple operations have a latency of two cycles. However, if a 2-cycle operation is issued five cycles after a 7-cycle operation on the same VSU pipe, the 2-cycle operation is dynamically converted into a 7-cycle operation to prevent a resultant bus collision (to make use of the issue slot).</p>						



Table 10-16. Instruction Latencies and Throughputs (Sheet 13 of 17)

Instruction	Pipeline	Latency	Throughput (IPC)	Notes decode / dispatch	Other Interlocks	Other Comments
xsdivdp xvdivdp	FPU	32 cycles	2/26 cycles		No	Makes use of microcoded sequence within the floating-point unit. Free FPU slots can be interleaved with other FPU instructions.
xvdivsp	FPU	28 cycles	2/22 cycles		No	
xssqrtdp xvsqrtdp	FPU	43 cycles	2/37 cycles		No	
xvsqrtp	FPU	32 cycles	2/26 cycles		No	
vaddfp vcfpxws vcfpxws vcxwfp vcuxwfp vexptefp vlogef p v Maddfp vnmsubfp vrefp vrfim vrfin vrfip vrfiz vrsqrtefp vsu bfp	FPU	6 cycles to FPU (+1 cycle to other VSU ops)	2/cycle			
xsmaxdp xsmindp xvc-mpeqdp xvcmeqsp xvc-mpgedp xvcmpgesp xvcmpgtdp xvcmpgtsp xvmaxdp xvmaxsp xvmindp xvminsp	vector simple integer unit (XS)	2 cycles to all VSU ops	2/cycle			See note 3.
xscmpodp xscmpudp xstdi-vdp xstsqrtdp xvdivdp xvtdivsp xvtqrtdp xvtqrtsp	vector simple integer unit (XS)	4..9 cycles (to CR)	2/cycle			
vcmpbfp. vcmpeqfp. vcmpequb. vcmpequh. vcmpequw. vcmpequd. vcmpgefp. vcmptfip. vcmpgtsb. vcmptsh. vcmpgtsw. vcmptsd. vcmptg- tub. vcmptuh. vcmptuw. vcmpgtud. xvcmeqdp. xvc- mpeqsp. xvcmpgedp. xvc- mpgesp. xvcmpgtdp. xvcmpgtsp.	vector simple integer unit (XS)	2 cycles to all VSU ops, 4.9 cycles to CR	2/cycle			See note 3.
mtvscr	vector simple integer unit (XS)		1/cycle			Completion serialized.

1. Store_agen IOPs are issued to the LSU, and store_data IOPs are issued to the LU. Fixed-point loads that do not have an XER destination can be issued to either the LU or the LSU. Loads with floating-point destinations can only be issued to the LU, and loads with XER sources can only be issued to the LSU.
2. Branches can issue two cycles after a producing a **cmp** instruction if they were in the same dispatch group, or the **cmp** is the youngest instruction writing a CR relative to the **cmp**. Otherwise, the issue-to-issue latency to a branch is three cycles.
3. Normally, permutes and simple operations have a latency of two cycles. However, if a 2-cycle operation is issued five cycles after a 7-cycle operation on the same VSU pipe, the 2-cycle operation is dynamically converted into a 7-cycle operation to prevent a resultant bus collision (to make use of the issue slot).

Table 10-16. Instruction Latencies and Throughputs (Sheet 14 of 17)

Instruction	Pipeline	Latency	Throughput (IPC)	Notes decode / dispatch	Other Interlocks	Other Comments
mfvscr vaddcuw vaddsbs vaddshs vaddsbs vaddubm vaddubs vadduhm vadduhs vadduwm vadduws vaddudm vand vandc vavgbs vavgsh vavgsw vavgub vavguh vavguw vclzb vclzd vclzh vclzw vcmpbf vcmpeqfp vcmpequb vcmpequh vcmpequw vcmpequd vcmpgefp vcmpgtfp vcmpgtsb vcmpgtsh vcmpgtsw vcmpgtsd vcmpgtu- tub vcmpgtuh vcmpgtuw vcmpgtud veqv vmaxfp vmaxsb vmaxsh vmaxsw vmaxsd vmaxub vmaxuh vmaxuw vmaxud vminfp vminsb vminsh vminsww vminsd vminub vminuh vminuw vminud vnand vnor vor vorc vpopcntb vpopcnth vpopcntw vrlb vrlid vrlh vrlw vshasigmad vshasigmaw vslb vsld vslh vslw vsrab vsrad vsrah vsraw vsrb vsrd vsrh vsrw vsubcuw vsubbsb vsub- shs vsubsws vsububm vsub- ubs vsubuhm vsubuhs vsubuwm vsubuws vsubudm vxor	vector sim- ple integer unit (XS)	2 cycles to all VSU ops	2/cycle			See note 3.
xxland xxlandc xxleqv xxl- nand xxlnor xxlor xxlorc xxlxor	vector sim- ple integer unit (XS)	2 cycles to all VSU ops	2/cycle			See note 3.
vpopcntd vadduqm vaddcuq vaddeuqm vaddecuq vsub- uqm vsubcuq vsubeuqm vsubecuq	vector sim- ple integer unit (XS)	4 cycles to all VSU ops	2/2 cycles			Cracked into two depen- dent XS instructions.

1. Store_agen IOPs are issued to the LSU, and store_data IOPs are issued to the LU. Fixed-point loads that do not have an XER destination can be issued to either the LU or the LSU. Loads with floating-point destinations can only be issued to the LU, and loads with XER sources can only be issued to the LSU.
 2. Branches can issue two cycles after a producing a **cmp** instruction if they were in the same dispatch group, or the **cmp** is the youngest instruction writing a CR relative to the **cmp**. Otherwise, the issue-to-issue latency to a branch is three cycles.
 3. Normally, permutes and simple operations have a latency of two cycles. However, if a 2-cycle operation is issued five cycles after a 7-cycle operation on the same VSU pipe, the 2-cycle operation is dynamically converted into a 7-cycle operation to prevent a resultant bus collision (to make use of the issue slot).



Table 10-16. Instruction Latencies and Throughputs (Sheet 15 of 17)

Instruction	Pipeline	Latency	Throughput (IPC)	Notes decode / dispatch	Other Interlocks	Other Comments
<code>vbpermq vgbbd vmrgew vmrghb vmrghh vmrghw vmrlb vmrlgh vmrlhw vmrgow vperm vpermxor vpkpx vpksdss vpksdus vpkshss vpkshus vpkswss vpkswus vpkuludm vpkuludus vpkuhum vpkuhus vpkuwum vpkuwus vsel vsl vsldoi vslo vspltb vsplth vspltsb vspltish vspltsw vspltw vsr vsro vupkhpx vupkhbs vupkhsh vupkhsw vupklpx vupklshb vupklsh vupklsw</code>	Permute unit (PM)	2 cycles to all VSU ops	2/cycle			See note 3.
<code>xxmrghw xxmrglw xxpermdi xxsel xxldwi xxspltw</code>	Permute unit (PM)	2 cycles to all VSU ops	2/cycle		No	See note 3.
<code>vhaddshs vhraddshs vmladduhm vmsummbm vmsumshm vmsumshs vmsumubm vmsumuhs vmsumuhs vmuless vmuless vmuless vmuless vmuless vmuless vmuless vmuless vmuless vmuless vmuless vmuless vmuless vmuless vmlouw vmlouwm vsum2sws vsum4sbs vsum4shs vsum4ubs vsum-sws</code>	VX (Vector complex)	7 cycles to all VSU ops	2/cycle			
<code>vcipher vcipherlast vncipher vncipherlast vpmsumb vpmsumrd vpmsumh vpmsumw vsbox</code>	CY (Vector Crypto)	6 cycles, +1 to FPU and Complex	1/cycle			Crypto pipeline shared between both VSU pipes.
<code>mfvsrd mfvsrwz</code>		4 or 5 cycles	1/cycle			Four if on same pipe. Five to all.
<code>mtvsrd mtvsrwz</code>		5				
<code>bcdadd. bcdsub. dadd dadd. dsub dsub. dcmpo dcmpu dtstdc dtstdg dtstx dtstsf dquai dquai. dqua dqua. drrnd drrnd. drintx drintx. drintn drintn. dctdp dctdp. ddedpd ddedpd. denbcd denbcd. dxex dxex. dxexq dxexq. diex diex. dscli dscli. dscri dscri. dtstdcq dtstdgq dtstfq</code>	DFU	13 cycles	1/cycle			DFU pipeline shared between both VSU pipes.
<ol style="list-style-type: none"> 1. Store_agen IOPs are issued to the LSU, and store_data IOPs are issued to the LU. Fixed-point loads that do not have an XER destination can be issued to either the LU or the LSU. Loads with floating-point destinations can only be issued to the LU, and loads with XER sources can only be issued to the LSU. 2. Branches can issue two cycles after a producing a cmp instruction if they were in the same dispatch group, or the cmp is the youngest instruction writing a CR relative to the cmp. Otherwise, the issue-to-issue latency to a branch is three cycles. 3. Normally, permutes and simple operations have a latency of two cycles. However, if a 2-cycle operation is issued five cycles after a 7-cycle operation on the same VSU pipe, the 2-cycle operation is dynamically converted into a 7-cycle operation to prevent a resultant bus collision (to make use of the issue slot). 						

Table 10-16. Instruction Latencies and Throughputs (Sheet 16 of 17)

Instruction	Pipeline	Latency	Throughput (IPC)	Notes decode / dispatch	Other Interlocks	Other Comments
dcmopoq dcmpuq drrndq drrndq. diexq diexq. dquaiq dquaiq. dstexq ddedpdq ddedpdq. denbcdq denbcdq. dscliq dscliq. dscrlq dscrlq. dctqpq dctqpq. drintxq drintxq. drintnq drintnq.	DFU	13 cycles +2 in VSU = 15 cycle	1/cycle			DFU pipeline shared between both VSU . Instruction is split in DFU and permute intruction.
daddq daddq. dsubq dsubq. dquaq dquaq.	DFU	2 in VSU +13 cycles +2 in VSU = 17 cyle	2/3 cycle			DFU pipeline shared between both VSU pipes. Instruction is split in DFU and two permute intructions.
drsp drsp. ddedpd ddedpd. dctfix dctfix.	DFU	25 cycles	1/ 12 cycle			DFU pipeline shared between both VSU pipes.
dcffix dcffix.	DFU	32 cycles	1/ 19 cycle			DFU pipeline shared between both VSU pipes.
dcffixq dcffixq.	DFU	32 cycles +2 in VSU = 34 cycles	1/ 19 cycle			DFU pipeline shared between both VSU pipes. Instruction is split in DFU and permute intruction.
dmul dmul.	DFU	24+N cycles at least 28 =(28 to 40) cycles	1/ 28 to 40 cycle			DFU pipeline shared between both VSU pipes. Every BCD digit of the shorter operand (N) extends the execution time by one cycle.
<ol style="list-style-type: none"> 1. Store_agen IOPs are issued to the LSU, and store_data IOPs are issued to the LU. Fixed-point loads that do not have an XER destination can be issued to either the LU or the LSU. Loads with floating-point destinations can only be issued to the LU, and loads with XER sources can only be issued to the LSU. 2. Branches can issue two cycles after a producing a cmp instruction if they were in the same dispatch group, or the cmp is the youngest instruction writing a CR relative to the cmp. Otherwise, the issue-to-issue latency to a branch is three cycles. 3. Normally, permutes and simple operations have a latency of two cycles. However, if a 2-cycle operation is issued five cycles after a 7-cycle operation on the same VSU pipe, the 2-cycle operation is dynamically converted into a 7-cycle operation to prevent a resultant bus collision (to make use of the issue slot). 						



Table 10-16. Instruction Latencies and Throughputs (Sheet 17 of 17)

Instruction	Pipeline	Latency	Throughput (IPC)	Notes decode / dispatch	Other Interlocks	Other Comments
dmulq dmulq.	DFU	20+2N cycles at least 28 +4in VSU =(32 to 90) cycles	1/ 28 to 86 cycle			DFU pipeline shared between both VSU pipes. Instruction is split in DFU and two permute instructions. Every BCD digit of the shorter operand (N) extends the execution time by two cycles.
ddiv ddiv.	DFU	32 +4Q cycles =(32 to 96) cycles	1/32 to 96 cycle			DFU pipeline shared between both VSU pipes. Every BCD digit of the quotient (Q) extends the execution time by four cycles.
ddivq ddivq.	DFU	32 +4Q cycles +4 in VSU =(36 to 172) cycles	1/32 to 168 cycle			DFU pipeline shared between both VSU pipes. Instruction is split in DFU and two permute instructions. Every BCD digit of the quotient (Q) extends the execution time by four cycles.
<ol style="list-style-type: none"> 1. Store_agen IOPs are issued to the LSU, and store_data IOPs are issued to the LU. Fixed-point loads that do not have an XER destination can be issued to either the LU or the LSU. Loads with floating-point destinations can only be issued to the LU, and loads with XER sources can only be issued to the LSU. 2. Branches can issue two cycles after a producing a cmp instruction if they were in the same dispatch group, or the cmp is the youngest instruction writing a CR relative to the cmp. Otherwise, the issue-to-issue latency to a branch is three cycles. 3. Normally, permutes and simple operations have a latency of two cycles. However, if a 2-cycle operation is issued five cycles after a 7-cycle operation on the same VSU pipe, the 2-cycle operation is dynamically converted into a 7-cycle operation to prevent a resultant bus collision (to make use of the issue slot). 						

10.4 PCI Express Performance

10.4.1 Bandwidth

On the POWER8 processor chip, the PCIe host bridge is integrated into the chip and provides 32 lanes of PCIe Express 3.0 ports.

10.4.2 Latency

Time to first byte latency calculations are approximately 220 ns when doing a DMA read without TCEs.

10.4.3 Cluster Latency 2K Message

Sending a 2K message from one cluster to another cluster using an IB adapter is 2611 ns.

10.4.4 I/O Bandwidth

Two PCIe x16 buses on the POWER8 chip support 28 GB/s.

10.4.5 PCIe Performance Goals

OnePCIe x16 bus is able to achieve 87.5% efficiency thus obtaining 14 GB/s bandwidth bi-directional (14.0 GB/s DMA writes and 14.0 GB/s DMA reads).

- PCIe payload of 512 bytes on memory writes
- PCIe payload of 256 bytes on read completions
- Assume no TCE miss
- Assume no RTC cache miss
- Assume no MSI
- Assume 64K packet operation



10.5 Performance Specific Instructions

10.5.1 Store Multiple and Store String

Store multiple and store string are processed much like load multiple and load string. Each of these instructions is microcoded with one store operation generated for each register. In a store string, only some bytes of the last register can be stored. This is also handled with a single operation of 1 - 7 bytes. For store string immediate, the first microcode group contains only NOPs. Each of the store operations results in an st-agen issue and an st-data issue. Each operation uses one SRQ entry. The store operations can execute in any order, although they are written from low to high address as maintained by the order of SRQ entries. When a group consists of four store operations, two are queued for each LSU pipe. In this way, two store operations can execute in each cycle. Many store multiple or store string instructions can require multiple groups of internal operations. Each group completes in order, without waiting for later groups to finish. Consequently, if the instruction is flushed, all groups including those that have completed, are redone. This means that these stores could be seen more than once by the Nest.

Store string indexed instructions are processed like load string indexed. The first microcode group contains an XER read and an address add. The next several microcode groups are entirely nops to fill the decode pipeline. At this point, the length is known and the correct number of stores is generated with up to six stores per group.

Unaligned string instructions occur when an individual store crosses a 4 KB page boundary. This causes the entire instruction to be flushed to be decoded again to avoid the boundary crossing. Each original store now is placed in its own group, consisting of a left shift, two stores, and a NOP. The two stores each store from 1 - 7 bytes of data left aligned in the register. The first microcode group consists of nops. For store string indexed, the first group contains an address add. There is no need to wait for XER in the unaligned case.

10.5.1.1 *Store Quadword*

Store quadword (stq) is a 2-way split on the POWER8 processor core because only a single st-agen and two st_data operations are used: store upper, store lower, nop, nop. The first store operates as a normal store; st-gen goes through an LSU pipe, and the st_data goes through an LU pipe. The Tag bit in XER is read by the st-agen operation of the store upper. The second store is a st_data only, and it goes through an LU pipe. A single SRQ entry is used to hold the store quadword. After the store quadword is completed, it goes out as a single store to the L1 cache (and updates it if necessary) and to the L2 cache.

In the POWER8 core, MSR[LE] = 0 mode stq data is allowed to store-forward to a load other than a **lарx**, **lve**, **lq** and **lqarx**. The **lq** or **lqarx**'s first doubleword always gets its data and tag bit from the cache.

An **stq**, however, does not store forward to any load, if any thread is running in MSR[LE] = 1 mode.

10.5.1.2 *eieio*

The **eieio** instruction is in a single instruction group. **eieio** is held at dispatch until the LMQ is drained and there are no LSU instructions in the issue queues. Cache-inhibited loads are rejected when an older PTESync is in the SRQ. The **eieio** is sent to the nest like a normal store.

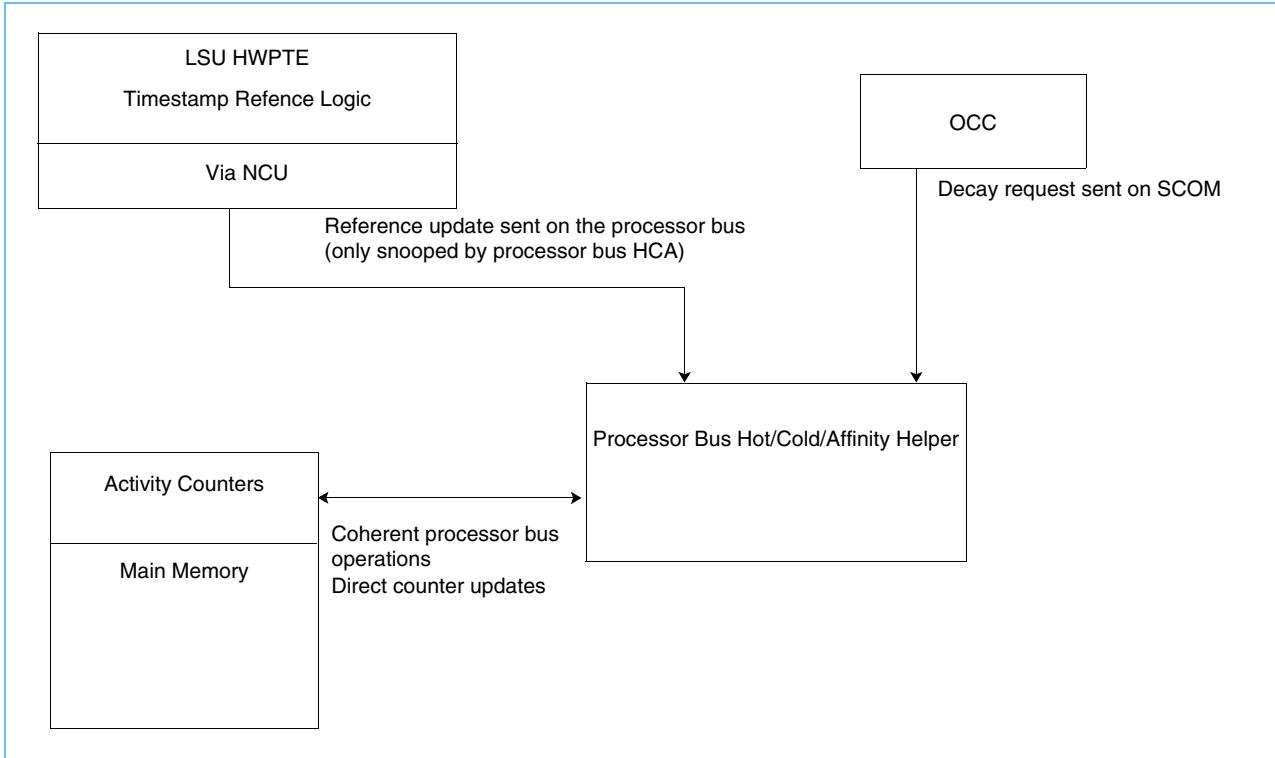
10.6 Other Topics

10.6.1 Hot/Cold Page Affinity Support

The POWER8 SMP interconnect Hot/Cold/Helper (HCA) is part of a hardware support system to enable efficient isolation of activity characteristics of memory pages. By monitoring hardware access to memory, it can gather information about the following:

- Activity rate: Enable power efficient memory management.
- Affinity information: Collect sharing pattern, enable operating-system-based page migrations.
- Long term access history: Maintain history.

Figure 10-1. Basic Building Blocks



LSU HWPTE sends a request on the processor bus to control the HCA. These requests are:

Reference update request. This command is only snooped by a specific HCA that has a matching value in the Base Address Register (BAR). There is a 32-byte payload; however, only the timestamp part is needed. The HCA then issues a reference update to the memory controller. The `hpc_ref_upd` operates on a single 4-byte data field in a 128-byte cache line, aligned on an 4-byte address boundary.

OCC Send Decay Request via SCOM

The decay request contains an address, range and a delay. The HCA then issues decay updates on the processor bus passing the address. The range indicates how many decay update commands to send. This works on a cache line in the counter area in the memory controller.



Activity Counters in Memory

The memory controller receives the new processor bus commands (activity update, reference update, and decay update) and updates the counters in memory.

Processor Bus Hot/Cold/Affinity Helper

Contains logic to snoop each memory reference in the processor bus. BAR registers are used to identify the memory location and range that belongs to this chip. The HCA contains a cache of the counter for each 4 KB memory page. It increments these counters as it detects a memory reference. HCA only works on 4 KB pages. Only the memory processor bus is snooped. HCA does not look at the combined response. The cache of these counters is not the full counter, but only 12 bits. As these counters fill, an activity update command on the processor bus is sent to the main memory to update the actual counter. The counter cache does not contain all possible 4 KB counters for this chip. As new references are detected, a random LRU algorithm is used to determine which line to cast out to memory to make space for this new counter. HCA will castout assuming memory is a 4 KB page size. Even if memory is at a different page size, the counters are mapped at a 4 KB granularity.

The following types are snooped by the HCA:

- All read groups
- rwitm
- dclaim, dcbz
- cl_dma_wr
- bsr_cp
- cp_m, cp_t, cp_tn, htm_cl_w
- cp_ig_me
- cl_w_cln
- dma_pr_w
- dc当地, dc当地
- bkill
- armw
- armwf
- ci_pr_rd
- dma_pr_rd
- ci_pr_w
- ci_pr_ooo_w
- cl_cln
- bsr_get

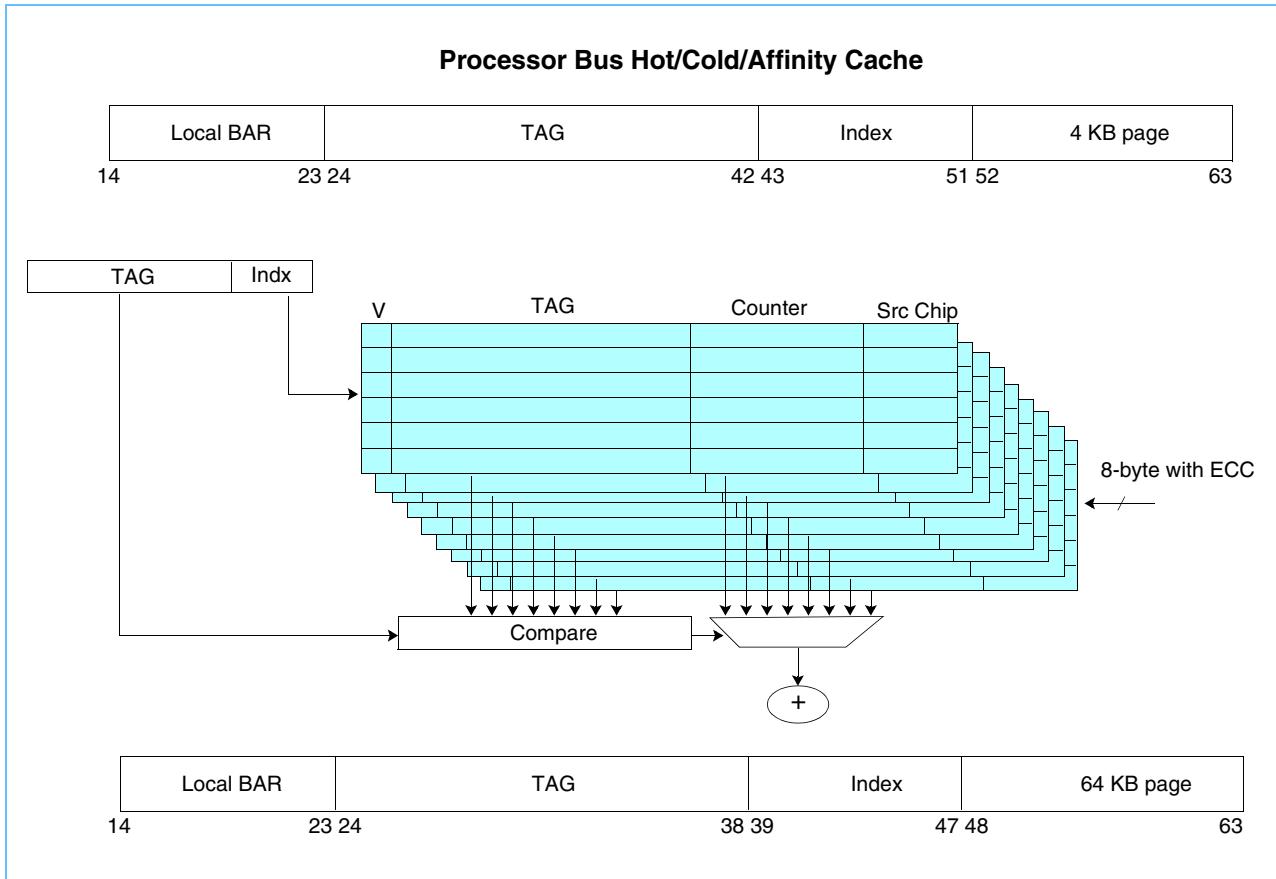
The following commands are issued by the HCA:

- Activity update: operands (counter address, count value, source chip)
- Reference update: operands (counter address, PTE time)
- Decay update: operands (counter region)
- rd_go_s: to obtain HPC on chip scope.

To reduce the bandwidth requirements for these updates, there is a mode to only send the 16th or 32nd hca.update. In such a case, the count is multiplied by 32 or 16 accordingly. However, this leads to inaccuracy in the count because not all hca.updates have the same count.

Figure 10-2 shows the HCA cache.

Figure 10-2. HCA Cache





10.6.2 Instruction That Can Soft Patch

Table 10-17 contains the explicit list of instructions that trigger a normalization interrupt (soft patch) if any of the operands contains scalar single-precision denormal data, as described in *Section 3.2.8 Handling of Denormal Single-Precision Values in Double-Precision Format* on page 71.

Table 10-17. Instructions that Trigger an Exception on Consumption of a Scalar SP Denorm (Sheet 1 of 11)

Mnemonic	Operands	Description
fcfid	FRT,FRB	Floating Convert From Integer Doubleword
fcfid.	FRT,FRB	Floating Convert From Integer Doubleword and Record
fcfids	FRT,FRB	Floating Convert From Integer Doubleword Single
fcfids.	FRT,FRB	Floating Convert From Integer Doubleword Single and Record
fcfidu	FRT,FRB	Floating Convert From Integer Doubleword Unsigned
fcfidu.	FRT,FRB	Floating Convert From Integer Doubleword Unsigned and Record
fcfidus	FRT,FRB	Floating Convert From Integer Doubleword Unsigned Single
fcfidus.	FRT,FRB	Floating Convert From Integer Doubleword Unsigned Single and Record
fcmpo	BF,FRA,FRB	Floating Compare Ordered
fcmpr	BF,FRA,FRB	Floating Compare Unordered
fmrgew	FRT,FRA,FRB	Floating Merge Even Word
fmrgow	FRT,FRA,FRB	Floating Merge Odd Word
ftdiv	BF,FRA,FRB	Floating Test for software Divide
ftsqr	BF,FRB	Floating Test for software Square Root
mtfsf	FLM,FRB,L,W	Move To FPSCR Fields
mtfsf.	FLM,FRB,L,W	Move To FPSCR Fields and Record
dadd	FRT,FRA,FRB	DFP Add
dadd.	FRT,FRA,FRB	DFP Add and Record
daddq	FRTp,FRAp,FRBp	DFP Add Quad
daddq.	FRTp,FRAp,FRBp	DFP Add Quad and Record
dcffix	FRT,FRB	DFP Convert From Fixed
dcffix.	FRT,FRB	DFP Convert From Fixed and Record
dcffixq	FRTp,FRB	DFP Convert From Fixed Quad
dcffixq.	FRTp,FRB	DFP Convert From Fixed Quad and Record
dcmpr	BF,FRA,FRB	DFP Compare Ordered
dcmprq	BF,FRAp,FRBp	DFP Compare Ordered Quad
dcmpr	BF,FRA,FRB	DFP Compare Unordered
dcmprq	BF,FRAp,FRBp	DFP Compare Unordered Quad
dctdp	FRT,FRB	DFP Convert To 64
dctdp.	FRT,FRB	DFP Convert To 64 and Record
dctfix	FRT,FRB	DFP Convert To Fixed
dctfix.	FRT,FRB	DFP Convert To Fixed and Record



Table 10-17. Instructions that Trigger an Exception on Consumption of a Scalar SP Denorm (Sheet 2 of 11)

Mnemonic	Operands	Description
dctfixq	FRT,FRBp	DFP Convert To Fixed Quad
dctfixq.	FRT,FRBp	DFP Convert To Fixed Quad and Record
dctqpq	FRTp,FRB	DFP Convert To 128
dctqpq.	FRTp,FRB	DFP Convert To 128 and Record
ddedpd	SP,FRT,FRB	DFP Decode DPD To BCD
ddedpd.	SP,FRT,FRB	DFP Decode DPD To BCD and Record
ddedpdq	SP,FRTp,FRBp	DFP Decode DPD To BCD Quad
ddedpdq.	SP,FRTp,FRBp	DFP Decode DPD To BCD Quad and Record
ddiv	FRT,FRA,FRB	DFP Divide
ddiv.	FRT,FRA,FRB	DFP Divide and Record
ddivq	FRTp,FRAp,FRBp	DFP Divide Quad
ddivq.	FRTp,FRAp,FRBp	DFP Divide Quad and Record
denbcd	S,FRT,FRB	DFP Encode BCD To DPD
denbcd.	S,FRT,FRB	DFP Encode BCD To DPD and Record
denbcdq	S,FRTp,FRBp	DFP Encode BCD To DPD Quad
denbcdq.	S,FRTp,FRBp	DFP Encode BCD To DPD Quad and Record
diex	FRT,FRA,FRB	DFP Insert Biased Exponent
diex.	FRT,FRA,FRB	DFP Insert Biased Exponent and Record
diexq	FRTp,FRA,FRBp	DFP Insert Biased Exponent Quad
diexq.	FRTp,FRA,FRBp	DFP Insert Biased Exponent Quad and Record
dmul	FRT,FRA,FRB	DFP Multiply
dmul.	FRT,FRA,FRB	DFP Multiply and Record
dmulq	FRTp,FRAp,FRBp	DFP Multiply Quad
dmulq.	FRTp,FRAp,FRBp	DFP Multiply Quad and Record
dqua	FRT,FRA,FRB,RMC	DFP Quantize
dqua.	FRT,FRA,FRB,RMC	DFP Quantize and Record
dquai	TE,FRT,FRB,RMC	DFP Quantize Immediate
dquai.	TE,FRT,FRB,RMC	DFP Quantize Immediate and Record
dquaiq	TE,FRTp,FRBp,RMC	DFP Quantize Immediate Quad
dquaiq.	TE,FRTp,FRBp,RMC	DFP Quantize Immediate Quad and Record
dquaq	FRTp,FRAp,FRBp,RMC	DFP Quantize Quad
dquaq.	FRTp,FRAp,FRBp,RMC	DFP Quantize Quad and Record
drdpq	FRTp,FRBp	DFP Round To 64
drdpq.	FRTp,FRBp	DFP Round To 64 and Record
rintn	R,FRT,FRB,RMC	DFP Round To FP Integer Without Inexact
rintn.	R,FRT,FRB,RMC	DFP Round To FP Integer Without Inexact and Record
rintnq	FRTp,FRBp,RMC	DFP Round To FP Integer Without Inexact Quad



Table 10-17. Instructions that Trigger an Exception on Consumption of a Scalar SP Denorm (Sheet 3 of 11)

Mnemonic	Operands	Description
drintnq.	FRTp,FRBp,RMC	DFP Round To FP Integer Without Inexact Quad and Record
drintx	R,FRT,FRB,RMC	DFP Round To FP Integer With Inexact
drintx.	R,FRT,FRB,RMC	DFP Round To FP Integer With Inexact and Record
drintxq	FRTp,FRBp,RMC	DFP Round To FP Integer With Inexact Quad
drintxq.	FRTp,FRBp,RMC	DFP Round To FP Integer With Inexact Quad and Record
drrnd	FRT,FRA,FRB,RMC	DFP Reround
drrnd.	FRT,FRA,FRB,RMC	DFP Reround and Record
drrndq	FRTp,FRA,FRBp,RMC	DFP Reround Quad
drrndq.	FRTp,FRA,FRBp,RMC	DFP Reround Quad and Record
drsp	FRT,FRB	DFP Round To 32
drsp.	FRT,FRB	DFP Round To 32 and Record
dscli	FRT,FRA,SH	DFP Shift Coefficient Left Immediate
dscli.	FRT,FRA,SH	DFP Shift Coefficient Left Immediate and Record
dscliq	FRTp,FRAp,SH	DFP Shift Coefficient Left Immediate Quad
dscliq.	FRTp,FRAp,SH	DFP Shift Coefficient Left Immediate Quad and Record
dscrit	FRT,FRA,SH	DFP Shift Coefficient Right Immediate
dscrit.	FRT,FRA,SH	DFP Shift Coefficient Right Immediate and Record
dscrig	FRTp,FRAp,SH	DFP Shift Coefficient Right Immediate Quad
dscrig.	FRTp,FRAp,SH	DFP Shift Coefficient Right Immediate Quad and Record
dsub	FRT,FRA,FRB	DFP Subtract
dsub.	FRT,FRA,FRB	DFP Subtract and Record
dsubq	FRTp,FRAp,FRBp	DFP Subtract Quad
dsubq.	FRTp,FRAp,FRBp	DFP Subtract Quad and Record
dtstdc	BF,FRA,DCM	DFP Test Data Class
dtstdcq	BF,FRAp,DCM	DFP Test Data Class Quad
dtstdg	BF,FRA,DGM	DFP Test Data Group
dtstdgq	BF,FRAp,DGM	DFP Test Data Group Quad
dtstex	BF,FRA,FRB	DFP Test Exponent
dtstexq	BF,FRAp,FRBp	DFP Test Exponent Quad
dtstsf	BF,FRA,FRB	DFP Test Significance
dtstsfq	BF,FRAp,FRBp	DFP Test Significance Quad
dxex	FRT,FRB	DFP Extract Biased Exponent
dxex.	FRT,FRB	DFP Extract Biased Exponent and Record
dxexq	FRT,FRBp	DFP Extract Biased Exponent Quad
dxexq.	FRT,FRBp	DFP Extract Biased Exponent Quad and Record
mfvsrd	RA,XS	Move From VSR Doubleword
mfvsrwz	RA,XS	Move From VSR Word and Zero



Table 10-17. Instructions that Trigger an Exception on Consumption of a Scalar SP Denorm (Sheet 4 of 11)

Mnemonic	Operands	Description
mtvscr	VB	Move To VSCR
vaddcuw	VT,VA,VB	Vector Add and Write Carry-Out Unsigned Word
vaddfp	VT,VA,VB	Vector Add Single-Precision
vaddsbs	VT,VA,VB	Vector Add Signed Byte Saturate
vaddshs	VT,VA,VB	Vector Add Signed Halfword Saturate
vaddsws	VT,VA,VB	Vector Add Signed Word Saturate
vaddubm	VT,VA,VB	Vector Add Unsigned Byte Modulo
vaddubs	VT,VA,VB	Vector Add Unsigned Byte Saturate
vadduhm	VT,VA,VB	Vector Add Unsigned Halfword Modulo
vadduhs	VT,VA,VB	Vector Add Unsigned Halfword Saturate
vadduwm	VT,VA,VB	Vector Add Unsigned Word Modulo
vadduws	VT,VA,VB	Vector Add Unsigned Word Saturate
vaddudm	VT,VA,VB	Vector Add Unsigned Doubleword Modulo
vand	VT,VA,VB	Vector Logical AND
vandc	VT,VA,VB	Vector Logical AND with Complement
vavgsb	VT,VA,VB	Vector Average Signed Byte
vavgsh	VT,VA,VB	Vector Average Signed Halfword
vavgsw	VT,VA,VB	Vector Average Signed Word
vavgub	VT,VA,VB	Vector Average Unsigned Byte
vavguh	VT,VA,VB	Vector Average Unsigned Halfword
vavguw	VT,VA,VB	Vector Average Unsigned Word
vbrd	VT,VB	Vector Byte Reverse Doubleword (microcode)
vbrdxor	VT,VA,VB	Vector Byte Reverse Doubleword Xor (microcode)
vbrw	VT,VB	Vector Byte Reverse Word (microcode)
vbrwxor	VT,VA,VB	Vector Byte Reverse Word Xor (microcode)
vctsxs	VT,VB,UIM	Vector Convert Single-Precision to Signed Fixed-Point Word Saturate
vctuxs	VT,VB,UIM	Vector Convert Single-Precision to Unsigned Fixed-Point Word Saturate
vcipher	VT,VA,VB	Vector AES Cipher
vcipherlast	VT,VA,VB	Vector AES Cipher Last
vclzb	VT,VB	Vector Count Leading Zero Byte
vclzh	VT,VB	Vector Count Leading Zero Halfword
vclzw	VT,VB	Vector Count Leading Zero Word
vcmpbfp	VT,VA,VB	Vector Compare Bounds Single-Precision
vcmpbfp.	VT,VA,VB	Vector Compare Bounds Single-Precision and Record
vcmpeqfp	VT,VA,VB	Vector Compare Equal To Single-Precision
vcmpeqfp.	VT,VA,VB	Vector Compare Equal To Single-Precision and Record
vcmpequb	VT,VA,VB	Vector Compare Equal To Unsigned Byte



Table 10-17. Instructions that Trigger an Exception on Consumption of a Scalar SP Denorm (Sheet 5 of 11)

Mnemonic	Operands	Description
vcmpequb.	VT,VA,VB	Vector Compare Equal To Unsigned Byte and Record
vcmpequh	VT,VA,VB	Vector Compare Equal To Unsigned Halfword
vcmpequh.	VT,VA,VB	Vector Compare Equal To Unsigned Halfword and Record
vcmpequw	VT,VA,VB	Vector Compare Equal To Unsigned Word
vcmpequw.	VT,VA,VB	Vector Compare Equal To Unsigned Word and Record
vcmpequd	VT,VA,VB	Vector Compare Equal To Unsigned Doubleword
vcmpequd.	VT,VA,VB	Vector Compare Equal To Unsigned Doubleword and Record
vcmpgefp	VT,VA,VB	Vector Compare Greater Than or Equal To Single-Precision
vcmpgefp.	VT,VA,VB	Vector Compare Greater Than or Equal To Single-Precision and Record
vcmpgtfp	VT,VA,VB	Vector Compare Greater Than Single-Precision
vcmpgtfp.	VT,VA,VB	Vector Compare Greater Than Single-Precision and Record
vcmpgtsb	VT,VA,VB	Vector Compare Greater Than Signed Byte
vcmpgtsb.	VT,VA,VB	Vector Compare Greater Than Signed Byte and Record
vcmpgtsh	VT,VA,VB	Vector Compare Greater Than Signed Halfword
vcmpgtsh.	VT,VA,VB	Vector Compare Greater Than Signed Halfword and Record
vcmpgtsw	VT,VA,VB	Vector Compare Greater Than Signed Word
vcmpgtsw.	VT,VA,VB	Vector Compare Greater Than Signed Word and Record
vcmpgtsd	VT,VA,VB	Vector Compare Greater Than Signed Doubleword
vcmpgtsd.	VT,VA,VB	Vector Compare Greater Than Signed Doubleword and Record
vcmpgtub	VT,VA,VB	Vector Compare Greater Than Unsigned Byte
vcmpgtub.	VT,VA,VB	Vector Compare Greater Than Unsigned Byte and Record
vcmpgtuh	VT,VA,VB	Vector Compare Greater Than Unsigned Halfword
vcmpgtuh.	VT,VA,VB	Vector Compare Greater Than Unsigned Halfword and Record
vcmpgtuw	VT,VA,VB	Vector Compare Greater Than Unsigned Word
vcmpgtuw.	VT,VA,VB	Vector Compare Greater Than Unsigned Word and Record
vcmpgtud	VT,VA,VB	Vector Compare Greater Than Unsigned Doubleword
vcmpgtud.	VT,VA,VB	Vector Compare Greater Than Unsigned Doubleword and Record
vcfsx	VT,VB,UIM	Vector Convert Signed Fixed-Point Word to Single-Precision
vcfux	VT,VB,UIM	Vector Convert Unsigned Fixed-Point Word to Single-Precision
vgbbd	VT,VB	Vector Gather Bits by Bytes by Doubleword
vexptefp	VT,VB	Vector 2 Raised to the Exponent Estimate Single-Precision
vlogefp	VT,VB	Vector Log Base 2 Estimate Single-Precision
vmaddfp	VT,VA,VC,VB	Vector Multiply-Add Single-Precision
vmaxfp	VT,VA,VB	Vector Maximum Single-Precision
vmaxsb	VT,VA,VB	Vector Maximum Signed Byte
vmaxsh	VT,VA,VB	Vector Maximum Signed Halfword
vmaxsw	VT,VA,VB	Vector Maximum Signed Word



Table 10-17. Instructions that Trigger an Exception on Consumption of a Scalar SP Denorm (Sheet 6 of 11)

Mnemonic	Operands	Description
vmaxsd	VT,VA,VB	Vector Maximum Signed Doubleword
vmaxub	VT,VA,VB	Vector Maximum Unsigned Byte
vmaxuh	VT,VA,VB	Vector Maximum Unsigned Halfword
vmaxuw	VT,VA,VB	Vector Maximum Unsigned Word
vmaxud	VT,VA,VB	Vector Maximum Unsigned Doubleword
vhaddshs	VT,VA,VB,VC	Vector Multiply-High-Add Signed Halfword Saturate
vhaddshs	VT,VA,VB,VC	Vector Multiply-High-Round-Add Signed Halfword Saturate
vminfp	VT,VA,VB	Vector Minimum Single-Precision
vminsbt	VT,VA,VB	Vector Minimum Signed Byte
vminsh	VT,VA,VB	Vector Minimum Signed Halfword
vminsw	VT,VA,VB	Vector Minimum Signed Word
vminsd	VT,VA,VB	Vector Minimum Signed Doubleword
vminub	VT,VA,VB	Vector Minimum Unsigned Byte
vminuh	VT,VA,VB	Vector Minimum Unsigned Halfword
vminuw	VT,VA,VB	Vector Minimum Unsigned Word
vminud	VT,VA,VB	Vector Minimum Unsigned Doubleword
vladduhm	VT,VA,VB,VC	Vector Multiply-Low-Add Unsigned Halfword Modulo
vmrgew	VT,VA,VB	Vector Merge Even Word
vmrghb	VT,VA,VB	Vector Merge High Byte
vmrghh	VT,VA,VB	Vector Merge High Halfword
vmrghw	VT,VA,VB	Vector Merge High Word
vmrglb	VT,VA,VB	Vector Merge Low Byte
vmrglh	VT,VA,VB	Vector Merge Low Halfword
vmrglw	VT,VA,VB	Vector Merge Low Word
vmrgow	VT,VA,VB	Vector Merge Odd Word
vmsummbm	VT,VA,VB,VC	Vector Multiply-Sum Mixed Byte Modulo
vmsumshm	VT,VA,VB,VC	Vector Multiply-Sum Signed Halfword Modulo
vmsumshs	VT,VA,VB,VC	Vector Multiply-Sum Signed Halfword Saturate
vmsumubm	VT,VA,VB,VC	Vector Multiply-Sum Unsigned Byte Modulo
vmsumuham	VT,VA,VB,VC	Vector Multiply-Sum Unsigned Halfword Modulo
vmsumuhs	VT,VA,VB,VC	Vector Multiply-Sum Unsigned Halfword Saturate
vmulesb	VT,VA,VB	Vector Multiply Even Signed Byte
vmulesh	VT,VA,VB	Vector Multiply Even Signed Halfword
vmuleub	VT,VA,VB	Vector Multiply Even Unsigned Byte
vmuleuh	VT,VA,VB	Vector Multiply Even Unsigned Halfword
vmulosb	VT,VA,VB	Vector Multiply Odd Signed Byte
vmulosh	VT,VA,VB	Vector Multiply Odd Signed Halfword



Table 10-17. Instructions that Trigger an Exception on Consumption of a Scalar SP Denorm (Sheet 7 of 11)

Mnemonic	Operands	Description
vmuloub	VT,VA,VB	Vector Multiply Odd Unsigned Byte
vmuloh	VT,VA,VB	Vector Multiply Odd Unsigned Halfword
vncipher	VT,VA,VB	Vector AES Inverse Cipher
vncipherlast	VT,VA,VB	Vector AES Inverse Cipher Last
vnmsubfp	VT,VA,VC,VB	Vector Negative Multiply-Subtract Single-Precision
vnor	VT,VA,VB	Vector Logical NOR
vor	VT,VA,VB	Vector Logical OR
vperm	VT,VA,VB,VC	Vector Permute
vpermxor	VT,VA,VB,VC	Vector Permute Xor
vpkpx	VT,VA,VB	Vector Pack Pixel
vpksdss	VT,VA,VB	Vector Pack Signed Doubleword Signed Saturate
vpksdus	VT,VA,VB	Vector Pack Signed Doubleword Unsigned Saturate
vpkshss	VT,VA,VB	Vector Pack Signed Halfword Signed Saturate
vpkshus	VT,VA,VB	Vector Pack Signed Halfword Unsigned Saturate
vpkswss	VT,VA,VB	Vector Pack Signed Word Signed Saturate
vpkswus	VT,VA,VB	Vector Pack Signed Word Unsigned Saturate
vpkudum	VT,VA,VB	Vector Pack Unsigned Doubleword Unsigned Modulo
vpkudus	VT,VA,VB	Vector Pack Unsigned Doubleword Unsigned Saturate
vpkuhum	VT,VA,VB	Vector Pack Unsigned Halfword Unsigned Modulo
vpkuhus	VT,VA,VB	Vector Pack Unsigned Halfword Unsigned Saturate
vpkuwum	VT,VA,VB	Vector Pack Unsigned Word Unsigned Modulo
vpkuwus	VT,VA,VB	Vector Pack Unsigned Word Unsigned Saturate
vpmsumb	VT,VA,VB	Vector Polynomial Multiply-Sum Byte
vpmsumd	VT,VA,VB	Vector Polynomial Multiply-Sum Doubleword
vpmsumh	VT,VA,VB	Vector Polynomial Multiply-Sum Halfword
vpmsumw	VT,VA,VB	Vector Polynomial Multiply-Sum Word
vrefp	VT,VB	Vector Reciprocal Estimate Single-Precision
vrfim	VT,VB	Vector Round to Single-Precision Integer toward -Infinity
vrfin	VT,VB	Vector Round to Single-Precision Integer Nearest
vrfip	VT,VB	Vector Round to Single-Precision Integer toward +Infinity
vrfiz	VT,VB	Vector Round to Single-Precision Integer toward Zero
vrlb	VT,VA,VB	Vector Rotate Left Byte
vrld	VT,VA,VB	Vector Rotate Left Doubleword
vrlh	VT,VA,VB	Vector Rotate Left Halfword
vrlw	VT,VA,VB	Vector Rotate Left Word
vrsqrtefp	VT,VB	Vector Reciprocal Square Root Estimate Single-Precision
vsbox	VT,VA	Vector Sbox (AES)



Table 10-17. Instructions that Trigger an Exception on Consumption of a Scalar SP Denorm (Sheet 8 of 11)

Mnemonic	Operands	Description
vsel	VT,VA,VB,VC	Vector Select
vshasigmad	VT,VA	Vector SHA-512 Sigma Doubleword
vshasigmaw	VT,VA	Vector SHA-256 Sigma Word
vsl	VT,VA,VB	Vector Shift Left
vslb	VT,VA,VB	Vector Shift Left Byte
vslid	VT,VA,VB	Vector Shift Left Doubleword
vsldoi	VT,VA,VB,SHB	Vector Shift Left Double by Octet Immediate
vslh	VT,VA,VB	Vector Shift Left Halfword
vslo	VT,VA,VB	Vector Shift Left by Octet
vslw	VT,VA,VB	Vector Shift Left Word
vspltb	VT,VB,UIM	Vector Splat Byte
vsplth	VT,VB,UIM	Vector Splat Halfword
vspltw	VT,VB,UIM	Vector Splat Word
vsr	VT,VA,VB	Vector Shift Right
vsrab	VT,VA,VB	Vector Shift Right Algebraic Byte
vsrad	VT,VA,VB	Vector Shift Right Algebraic Doubleword
vsrah	VT,VA,VB	Vector Shift Right Algebraic Halfword
vsraw	VT,VA,VB	Vector Shift Right Algebraic Word
vsrb	VT,VA,VB	Vector Shift Right Byte
vsrd	VT,VA,VB	Vector Shift Right Doubleword
vsrh	VT,VA,VB	Vector Shift Right Halfword
vsro	VT,VA,VB	Vector Shift Right by Octet
vsrw	VT,VA,VB	Vector Shift Right Word
vsbcuw	VT,VA,VB	Vector Subtract and Write Carry-Out Unsigned Word
vsubfp	VT,VA,VB	Vector Subtract Single-Precision
vsubsbs	VT,VA,VB	Vector Subtract Signed Byte Saturate
vsubshs	VT,VA,VB	Vector Subtract Signed Halfword Saturate
vsubsws	VT,VA,VB	Vector Subtract Signed Word Saturate
vsububm	VT,VA,VB	Vector Subtract Unsigned Byte Modulo
vsububs	VT,VA,VB	Vector Subtract Unsigned Byte Saturate
vsubuhm	VT,VA,VB	Vector Subtract Unsigned Halfword Modulo
vsubuhs	VT,VA,VB	Vector Subtract Unsigned Halfword Saturate
vsubuwm	VT,VA,VB	Vector Subtract Unsigned Word Modulo
vsubuws	VT,VA,VB	Vector Subtract Unsigned Word Saturate
vsubudm	VT,VA,VB	Vector Subtract Unsigned Doubleword Modulo
vsum2sws	VT,VA,VB	Vector Sum across Half Signed Word Saturate
vsum4sbs	VT,VA,VB	Vector Sum across Quarter Signed Byte Saturate



Table 10-17. Instructions that Trigger an Exception on Consumption of a Scalar SP Denorm (Sheet 9 of 11)

Mnemonic	Operands	Description
vsum4shs	VT,VA,VB	Vector Sum across Quarter Signed Halfword Saturate
vsum4ubs	VT,VA,VB	Vector Sum across Quarter Unsigned Byte Saturate
vsumsws	VT,VA,VB	Vector Sum across Signed Word Saturate
vupkhpx	VT,VB	Vector Unpack High Pixel
vupkhsb	VT,VB	Vector Unpack High Signed Byte
vupkhsh	VT,VB	Vector Unpack High Signed Halfword
vupkhsw	VT,VB	Vector Unpack High Signed Word
vupklpx	VT,VB	Vector Unpack Low Pixel
vupklsb	VT,VB	Vector Unpack Low Signed Byte
vupklsh	VT,VB	Vector Unpack Low Signed Halfword
vupklsw	VT,VB	Vector Unpack Low Signed Word
vxor	VT,RA,RB	Vector Logical XOR
xscmpodp	BF,XA,XB	VSX Scalar Compare Ordered Double-Precision
xscmpudp	BF,XA,XB	VSX Scalar Compare Unordered Double-Precision
xscvspdp	XT,XB	VSX Scalar Convert Single-Precision to Double-Precision (p=1)
xscvsxddp	XT,XB	VSX Scalar Convert Signed Fixed-Point Doubleword to Double-Precision
xscvsxdsp	XT,XB	VSX Scalar Convert and Round Signed Integer Doubleword to Single-Precision Format
xscvuxddp	XT,XB	VSX Scalar Convert Unsigned Fixed-Point Doubleword to Double-Precision
xscvuxdsp	XT,XB	VSX Scalar Convert and Round Unsigned Integer Doubleword to Single-Precision Format
xsmaxdp	XT,XA,XB	VSX Scalar Maximum Double-Precision
xsmindp	XT,XA,XB	VSX Scalar Minimum Double-Precision
xstdivdp	BF,XA,XB	VSX Scalar Test for Software Divide Double-Precision
xstsqrtdp	BF,XB	VSX Scalar Test for Software Square Root Double-Precision
xvabssp	XT,XB	VSX Vector Absolute Value Single-Precision
xvaddsp	XT,XA,XB	VSX Vector Add Single-Precision
xvcmpeqdp	XT,XA,XB	VSX Vector Compare Equal To Double-Precision
xvcmpeqdp.	XT,XA,XB	VSX Vector Compare Equal To Double-Precision and Record
xvcmpeqsp	XT,XA,XB	VSX Vector Compare Equal To Single-Precision
xvcmpeqsp.	XT,XA,XB	VSX Vector Compare Equal To Single-Precision and Record
xvcmpgedp	XT,XA,XB	VSX Vector Compare Greater Than or Equal To Double-Precision
xvcmpgedp.	XT,XA,XB	VSX Vector Compare Greater Than or Equal To Double-Precision and Record
xvcmpgesp	XT,XA,XB	VSX Vector Compare Greater Than or Equal To Single-Precision
xvcmpgesp.	XT,XA,XB	VSX Vector Compare Greater Than or Equal To Single-Precision and Record
xvcmpgtdp	XT,XA,XB	VSX Vector Compare Greater Than Double-Precision
xvcmpgtdp.	XT,XA,XB	VSX Vector Compare Greater Than Double-Precision and Record
xvcmpgtsp	XT,XA,XB	VSX Vector Compare Greater Than Single-Precision



Table 10-17. Instructions that Trigger an Exception on Consumption of a Scalar SP Denorm (Sheet 10 of 11)

Mnemonic	Operands	Description
xvcmpgtsp.	XT,XA,XB	VSX Vector Compare Greater Than Single-Precision and Record
xvcpsgnsp	XT,XA,XB	VSX Vector Copy Sign Single-Precision
xvcvspdp	XT,XB	VSX Vector Convert Single-Precision to Double-Precision (p=1)
xvcvpsxds	XT,XB	VSX Vector Convert Single-Precision to Signed Fixed-Point Doubleword Saturate
xvcvpsxws	XT,XB	VSX Vector Convert Single-Precision to Signed Fixed-Point Word Saturate
xvcvspuxds	XT,XB	VSX Vector Convert Single-Precision to Unsigned Fixed-Point Doubleword Saturate
xvcvspuxws	XT,XB	VSX Vector Convert Single-Precision to Unsigned Fixed-Point Word Saturate
xvcvsxddp	XT,XB	VSX Vector Convert Signed Fixed-Point Doubleword to Double-Precision
xvcvsxdsp	XT,XB	VSX Vector Convert Signed Fixed-Point Doubleword to Single-Precision
xvcvsxwdp	XT,XB	VSX Vector Convert Signed Fixed-Point Word to Double-Precision
xvcvsxwsp	XT,XB	VSX Vector Convert Signed Fixed-Point Word to Single-Precision
xvcvuxddp	XT,XB	VSX Vector Convert Unsigned Fixed-Point Doubleword to Double-Precision
xvcvuxdsp	XT,XB	VSX Vector Convert Unsigned Fixed-Point Doubleword to Single-Precision
xvcvuxwdp	XT,XB	VSX Vector Convert Unsigned Fixed-Point Word to Double-Precision
xvcvuxwsp	XT,XB	VSX Vector Convert Unsigned Fixed-Point Word to Single-Precision
xdivsp	XT,XA,XB	VSX Vector Divide Single-Precision
xvmaddasp	XT,XA,XB	VSX Vector Multiply-Add Type-A Single-Precision
xvmaddmsp	XT,XA,XB	VSX Vector Multiply-Add Type-M Single-Precision
xvmaxdp	XT,XA,XB	VSX Vector Maximum Double-Precision
xvmaxsp	XT,XA,XB	VSX Vector Maximum Single-Precision
xvmindp	XT,XA,XB	VSX Vector Minimum Double-Precision
xvminsp	XT,XA,XB	VSX Vector Minimum Single-Precision
xvmsubasp	XT,XA,XB	VSX Vector Multiply-Subtract Type-A Single-Precision
xvmsubmsp	XT,XA,XB	VSX Vector Multiply-Subtract Type-M Single-Precision
xvmulsp	XT,XA,XB	VSX Vector Multiply Single-Precision
xvnabssp	XT,XB	VSX Vector Negative Absolute Value Single-Precision
xvnegsp	XT,XB	VSX Vector Negate Single-Precision
xvnmaddasp	XT,XA,XB	VSX Vector Negative Multiply-Add Type-A Single-Precision
xvnmaddmsp	XT,XA,XB	VSX Vector Negative Multiply-Add Type-M Single-Precision
xvnmsubasp	XT,XA,XB	VSX Vector Negative Multiply-Subtract Type-A Single-Precision
xvnmsubmsp	XT,XA,XB	VSX Vector Negative Multiply-Subtract Type-M Single-Precision
xvresp	XT,XB	VSX Vector Reciprocal Estimate Single-Precision
xvrspi	XT,XB	VSX Vector Round to Single-Precision Integer
xvrspic	XT,XB	VSX Vector Round to Single-Precision Integer using Current rounding mode
xvrspim	XT,XB	VSX Vector Round to Single-Precision Integer toward -Infinity
xvrspip	XT,XB	VSX Vector Round to Single-Precision Integer toward +Infinity
xvrspiz	XT,XB	VSX Vector Round to Single-Precision Integer toward Zero



Table 10-17. Instructions that Trigger an Exception on Consumption of a Scalar SP Denorm (Sheet 11 of 11)

Mnemonic	Operands	Description
xvrsqrtesp	XT,XB	VSX Vector Reciprocal Square Root Estimate Single-Precision
xvsqrtp	XT,XB	VSX Vector Square Root Single-Precision
xvsubsp	XT,XA,XB	VSX Vector Subtract Single-Precision
xtdivdp	BF,XA,XB	VSX Vector Test for software Divide Double-Precision
xtdivsp	BF,XA,XB	VSX Vector Test for software Divide Single-Precision
xvtsqrtdp	BF,XB	VSX Vector Test for software Square Root Double-Precision
xvtsqrtp	BF,XB	VSX Vector Test for software Square Root Single-Precision
xxland	XT,XA,XB	VSX Logical AND
xxlandc	XT,XA,XB	VSX Logical AND with Complement
xxlnor	XT,XA,XB	VSX Logical NOR
xxlor	XT,XA,XB	VSX Logical OR
xxlxor	XT,XA,XB	VSX Logical XOR
xxmrghw	XT,XA,XB	VSX Merge High Word
xxmrglw	XT,XA,XB	VSX Merge Low Word
xxpermdi	XT,XA,XB,DM	VSX Permute Doubleword Immediate
xxsel	XT,XA,XB,XC	VSX Select
xxsldwi	XT,XA,XB,SHW	VSX Shift Left Double by Word Immediate
xxspltd	XT,XA	VSX Splat Doubleword (microcode)
xxsplitw	XT,XB,UIM	VSX Splat Word

Table 10-18 lists BFP and VSX store instructions that cause a soft patch request if the address is unaligned and LSU flushes to microcode. For example, a store crosses 4K or the address crosses an 8-byte boundary for doubleword (or smaller) operations and a 16-byte boundary for quadword operations in DAWR mode (DWARX0 for the POWER8 core).

Table 10-18. Soft Patch Instruction on Unaligned Stores

Mnemonic	Operands	Description
stfd	FRS,D(RA)	Store Float Double
stfdup	FRS,D(RA)	Store Float Double with Update
stfdupx	FRS,RA,RB	Store Float Double with Update Indexed
stfdx	FRS,RA,RB	Store Float Double Indexed
stfiwx	FRS,RA,RB	Store Float as Integer Word Indexed
stxsdx	XS,RA,RB	Store VSR Scalar Doubleword Indexed
stxsiwx	XS,RA,RB	Store VSX Scalar Integer Word Indexed
stxvd2x	XS,RA,RB	Store VSR Vector Doubleword*2 Indexed (only in LE mode)
stxvw4x	XS,RA,RB	Store VSR Vector Word*4 Indexed (only in LE mode)

11. Performance Monitor

The POWER8 processor has built-in features for monitoring and collecting data for performance analysis. Collectively, the features are referred to as instrumentation. Performance instrumentation is divided into two broad categories: the performance monitor and the trace facilities.

11.1 Performance Monitor Overview

Each POWER8 core has four core-level hypervisor performance monitor counters (HPMCs), 48 thread-level performance monitor counters (PMCs), and 16 thread-level supervisor-level performance monitor counters (SPMCs).

Each HPMC is 64 bits wide. These counters can only be read or written in hypervisor mode.

- HPMC1 and HPMC2 are programmable.
- HPMC3 is a dedicated counter for run instructions. Run instructions are completed PowerPC instructions gated by the run latch.
- HPMC4 is a dedicated counter for run cycles. It counts cycles when any thread's run latch is set.
- MMCRH is a hypervisor resource that controls the hypervisor-specific performance monitor features.

Each thread has six PMCs. Each PMC is 32 bits wide. By connecting adjacent PMCs, PMC1 - 4 can be used as $32 \times N$ ($N = 1 - 4$) bit counters.

- PMC1 - 4 are programmable.
- PMC5 is a dedicated counter for run instructions. Run instructions are completed PowerPC instructions gated by the run latch.
- PMC6 is a dedicated counter for run cycles. Run cycles are cycles that are gated by the run latch.

Each thread has two supervisor performance monitor counters (SPMCs). Each SPMC is 32 bits wide.

- SPMC1 and SPMC2 are programmable.
- MMCRS is a supervisor resource that controls the SPMCs.

The performance monitor is configured and controlled through the following registers, buffer, and interrupt:

- Monitor Mode Control Register 0 (MMCR0) is a hardware thread that controls basic operation (start/stop/freeze) of the performance monitor.
- Monitor Mode Control Register 1 (MMCR1) is a hardware-thread resource that controls what to count.
- Monitor Mode Control Register 2 (MMCR2) is a hardware-thread resource that controls the basic operation of each PMC individually.
- Monitor Mode Control Register A (MMCRA) is a hardware-thread resource that includes indicator bits for feedback between the hardware, software, and configuration fields for special features of the performance monitor.
- Sampled Instruction Address Register (SIAR) is a 64-bit register.
- Sampled Data Address Register (SDAR) is a 64-bit register.
- Sampled Instruction Event Register (SIER) is a 64-bit register.
- Branch history rolling buffer (BHRB) is a buffer that contains the target addresses of most recent branch instructions for which the branch was taken.



- Performance monitor interrupt (PMI) is caused by monitored conditions and events.

11.2 Performance Monitor Functions

The POWER8 performance monitor includes the following functions:

- Counts instructions completed and cycles gated or not gated (according to the value of MMCR0[55]) by the run latch in individual (dedicated) 32-bit counters. The counting of these events can be enabled by software under several conditions such as problem or supervisor state.
- Counts up to four concurrent software selected events in individual 32-bit counters. The counting of these events can be enabled by software under several conditions such as problem or supervisor state, run, or wait state.
- Generates a maskable interrupt when an event counter overflows.
- Freezes the contents of the event counters until a selected event or condition occurs and then begins counting (triggering).
- Increments the event counters until a selected event or condition occurs and then freezes counting.
- Monitors classes of instructions selected by the instruction matching facility.
- Chooses an instruction for detailed monitoring (sampling).
- Counts start/stop event pairs that exceed a selected timeout value (thresholding).

11.2.1 Performance Monitor Event Selection

One event per counter can be selected for monitoring at a given time. The event to be monitored is selected by setting the appropriate value into the MMCR1 event selection field for that counter. The event counted can be the number of cycles when the event is active or the number of occurrences of the event depending on the particular event selected.

11.2.2 Machine States and Enabling the Performance Monitor Counters

The performance monitor counting can be enabled or disabled under several machine states that are selected using control bit fields in the MMCRs and the state bits in other special purpose registers (SPRs).

11.2.3 Trigger Events and Enabling the Performance Monitor Counters

Certain conditions and events, called trigger events, control performance monitor activities, such as starting and stopping the counters and causing performance monitor exceptions. These scenarios are selected using the condition or event enable bit fields and the exception enable bits of the MMCRs in conjunction with the control bits in other SPRs.

11.2.4 Performance Monitor Exceptions, Alerts, and Interrupts

Trigger events can cause performance monitor exceptions to occur based on the values of the exception enable bits in the MMCRs. An enabled exception causes the Performance Monitor Alert Occurred (PMAO) indicator bit to be set in the MMCR0 Register. This bit can only be reset by software. When running in a thread-level environment, the operating system can be swapped out while a performance monitor exception



alert is pending. The hypervisor preserves the value of MMCR0 across the thread swap. When the operating system is redispached, the alert is still pending. When enabled for external interrupts, a performance monitor alert causes a performance monitor interrupt to occur.

11.2.5 Sampling

The POWER8 processor can be configured to sample instructions for detailed monitoring. POWER8 instrumentation supports setting mask values for matching particular instructions or types of instructions that are then eligible to be sampled. The performance monitor includes events for counting sampled instructions at each stage of the pipeline and in certain other situations. Instruction sampling is a useful facility for gathering both detailed and statistical information for particular instructions.

11.2.6 Thresholding

The POWER8 processor monitors the pipeline stage progression of sampled instructions and can detect when the stage-to-stage cycle count for a selected start/stop pair of pipeline stages exceeds a specified threshold value. The threshold value can also be used to detect sampled loads whose latency exceeds the threshold value or simply to count the occurrence of any event between the start/stop pair.

11.2.7 Trace Support Facilities

The POWER8 processor supports both the single-step and the branch-trace modes as defined by the Power ISA.

11.3 Special Purpose Registers and Fields Associated with Instrumentation

The POWER8 processor instrumentation facilities and associated POWER8 components include several SPRs used for or associated with performance monitoring, instruction matching, instruction sampling, and tracing. Unless noted otherwise, the special purpose registers described can be read in problem and supervisor state by using the **mfsp** instruction and written in supervisor state by using the **mtsp** instruction. The MSR register is read by the **mfm** instruction and written by the **mtm** instruction.

The POWER8 processor instrumentation facilities include the following special purpose registers and register bit fields:

- Performance Monitor Mode Control Registers (MMCRx). These registers include both counting control and event select bit fields.
- Performance Monitor Counter Registers (PMCx). These registers increment for each time (or cycle, depending on the selected event) that an event occurs while the counter is enabled. These registers also have the control function for the counter overflow condition.
- Machine State Register [EE] (MSR[EE]). This register bit is used to enable or disable the external interrupt. The performance monitor interrupt is considered an external interrupt.
- Machine State Register [PMM] (MSR[PMM]). This register bit is used to enable or disable performance monitor activity controlled by the Process Mark bit.
- Machine State Register [PR] (MSR[PR]). This register bit is used to establish problem/supervisor mode and the performance monitor counting activity controlled by this bit.
- Machine State Register [SE] (MSR[SE]). This register bit is used to enable or disable a trace interrupt after each instruction is completed.



- Machine State Register [BE] (MSR[BE]). This register bit is used to enable or disable a trace interrupt after a branch instruction is completed.
- Control Register [31]. This register bit is used by operating systems to indicate idle/run state. The performance monitor can use this bit to avoid counting events during idle periods. This bit is commonly called the Run Latch.
- Instruction Match CAM Register (IMC). The IMC SPR is used to access the IMC array, which contains tag bits and mask values used for instruction matching. The **mtimc** and **mfimc** instructions can be executed only in supervisor mode.
- Timebase Register [47, 51, 55, 63]. These register bits are used for time-based events. Most performance monitor events are cycle based; that is, they count based on processor cycles. The Timebase Register is used to maintain time-of-day and can be used by the performance monitor to count time intervals.
- Sample Address Registers (SxAR). The Sampled Instruction Address Register (SIAR) contains the instruction address relating to a sampled instruction. The Sampled Data Address Register (SDAR) contains the data address relating to a sampled instruction. These registers can only be updated when performance monitor exceptions are enabled. This protects the contents from change until software can read them. The values written to these registers by the hardware depend on the processing state and on the kind of instruction that is being sampled.
- Machine Status Save/Restore Register (SRRO, SRR1). These registers are used to save machine status during interrupts.

11.4 Enhanced Sampling Support

Profiling or sampling is a common approach to associate expensive performance events in a processor to instruction and data addresses. Profiling enables the identification of hot spots in code and data, performance-sensitive areas, and problem instructions, data areas, or both. Profiling is commonly achieved by identifying a particular instruction and collecting detailed information about that instruction (instruction sampling).

The Power ISA provides two SPRs to identify sampled instructions. The SIAR captures the effective address of the sampled instruction. The SDAR captures the effective address of the sampled instruction's data operand, if any. Each of these registers has a Valid bit in the Sampled Instruction Event Register (SIER). The SIAR and SDAR are not cleared when a new sampled instruction is selected, but the indicator bits are cleared. Therefore, the indicator bits are required to show that the addresses housed in the SIAR or SDAR are not for a previous (executed or cancelled) sampled instruction. The sampled registers can only be updated by the processor when performance monitor exceptions are enabled. Performance monitor exceptions clear the enable bit, which locks the contents of the sampled registers.

The POWER8 processor supports three sampling modes:

- Random Instruction Sampling (RIS). Random instruction sampling selects (or marks) one instruction at a time and tracks its execution through the processor pipeline from decode to completion. Events that can be attributed to a sampled instruction are called marked events. By profiling on marked events it is possible to uniquely identify which instruction caused a particular event.
- Random Event Sampling (RES). Random event sampling selects or marks an instruction after an event has happened to an instruction. This was specifically added for improved sampling rates for important performance-sensitive events such as cache misses or branch mispredicts.



- Continuous Sampling (CS). Continuous sampling that updates the SIAR and SDAR on every instruction's completion and data access. Continuous sampling is useful for profiling on execution frequency or cache-line accesses.

The POWER8 processor has the ability to associate multiple performance events to the same instruction or data address. This is accomplished by recording up to 64 bits of information pertaining to a marked instruction during its lifetime in the pipeline. These 64 bits are accessible by using a software-accessible SIER. This enables collection of multiple events with a single pass.

11.5 POWER8 Performance Monitor Event Selection

For each of the four programmable thread-level counters, one event can be selected for monitoring at a given time. The event to be monitored is selected among the available event sources by setting the corresponding MMCR1[PMCxUNIT] and MMCR1[PMCxSEL] bits to the appropriate values. The remaining MMCR bits control additional details for certain events. The quantities counted indicate the number of cycles that the event is active or the number of occurrences of the event, depending on the setting of MMCR1[PMCxSEL(7)].

The POWER8 events are listed in *Appendix D Performance Monitoring Events* on page 387.

11.5.1 Event Bus Events and Event Bus Ramp

Some of the events available on the POWER8 core are routed from the units to the [PMU](#) using an event bus. The event bus carries one pair of events per PMC. The event pair is selected from all available events on the unit side. On the PMU side, the PMC can be set to count either event in the pair, the sum of both paired events, or cycles in which both events in the pair are active according to MMCR1[PMCxSEL(6) and PMCxCOMB].

11.5.2 Direct Events

Direct events are provided as dedicated signals from the units to the PMU. In particular, they are fully independent of the event bus and can be monitored at any time. A subset of direct events belong to the compatibility section for the thread-level counters and feed the hypervisor counters.



11.6 Performance Monitor Facility

Performance monitor operation is summarized by the following hierarchy, starting at the lowest level:

- A counter negative condition exists when the value in a PMC is negative (that is, when bit 0 of the PMC = '1'). A time-base transition event occurs when a selected bit of the Timebase changes from '0' to '1' (the bit is selected by a field in MMCR0). The term "condition" or "event" is used as an abbreviation for a counter negative condition or time-base transition event. A condition or event can be caused implicitly by the hardware (for example, incrementing a PMC) or explicitly by software (**mtspr**).
- A condition or event is enabled if the corresponding enable bit MMCR0[PMAE] is set. The occurrence of an enabled condition or event can have effects within the performance monitor, such as causing the PMCs to cease counting if MMCR0[FCECE] is set.
- An enabled condition or event causes a performance monitor alert if performance monitor alerts are enabled by MMCR0[PMAE]. A single performance monitor alert can reflect multiple enabled conditions and events.
- A performance monitor alert causes a performance monitor exception.
- When a performance monitor exception occurs, MMCR0[PMAO] is set to '1' within a reasonable period of time, but no later than the completion of the next context-synchronizing instruction or event. Even without software synchronization, the new contents of MMCR0[PMAO] is set to '1' sufficiently soon that the performance monitor facility is useful to software for its intended purposes.
- A performance monitor exception causes one of the following:
 - If MSR[EE] = '1' and MMCR0[EBC] = '0', an interrupt occurs.
 - If MSR[PR] = '1' and MMCR0[EBC] = '1', a performance monitor event-based exception occurs if BESCR[PME] = '1'.

Note: The performance monitor can be effectively disabled (that is, put into a state in which the performance monitor SPRs are not altered and performance monitor exceptions do not occur) by setting MMCR0 = 'x0000_0000_8000_0000'.

11.6.1 Performance Monitor Facility Registers

The performance monitor registers count events, control the operation of the performance monitor, and provide associated information.

11.6.1.1 Performance Monitor Counters (PMC1 - 6)

The six performance monitor counters, PMC1 - PMC6, are 32-bit registers that count events. PMC1 - PMC4 are referred to as "programmable" counters because the events that can be counted can be specified by software. The codes that identify the events that are counted are selected by specifying the appropriate code in the PMCxSEL event select fields in MMCR1[32:63]. Chaining two or more counters can be accomplished by setting PMCxSEL to select an event for counting defined as the overflow (msb going to '1') of a chained counter. The lower bit, PMCxSEL(7), controls counting the occurrences of the event or counting the cycles the event is active. Some events are documented as multi-bit events, where a multi-bit count (a count of 1 - 8) can be counted per cycle. When PMCxSEL(7) = '1', the falling edges of the event are counted (a count of one per cycle whether it is a multi-bit event or not). Some events can include operations that are performed out-of-order and speculatively.

When MMCR0[37] is set, PMU interrupts are enabled including counter negative and BHRB interrupts. A counter negative interrupt is generated when the MMCR0[48:49], PMCxCE, condition enable is set and the associated PMC x most-significant bit transitions from '0' to '1'. The source of the interrupt can be determined by examination of MMCR0[52], MMCR0[54], and the PMC or PMCs with negative conditions. See *Table 11-2* on page 276 for a description of the MMCR0 Register

PMC5 and PMC6 are not programmable. PMC5 counts completed instructions and PMC6 counts cycles. MMCR0[55] controls whether PMC5 and PMC6 are gated by the run latch. MMCR0[PMCC] controls whether PMC5 and PMC6 are under the control of various other bits in MMCR0 and MMCR2. When PMC5 and PMC6 are not under the control of these bits, PMC5 and PMC6 do not cause performance monitor events (PMC interrupts). When MMCR0[PMCC] = '11', PMC5 and PMC6 are allocated to supervisor control. In this case, the freeze conditions in MMCRS apply to PMC5 and PMC6 rather than the freeze conditions in MMCR0 and MMCR2.

Table 11-1. Performance Monitor Counter Register

Bits	Field Name	Description
0	CTR_NEG	Counter negative bit. When an adjacent PMC uses overflow counting, this becomes count data.
1:31	CTRDATA	Counter data.

Note: PMC5 and PMC6 facilitate calculating basic performance metrics such as cycles per instruction (CPI).

11.6.1.2 Monitor Control Register 0 (MMCR0)

Monitor Mode Control Register 0 (MMCR0) is a 64-bit register. This register, along with MMCR1 and MMCRA, controls the operation of the performance monitor.

Some MMCR0 bits are altered by the hardware when various events occur and some bits are altered by software.

The following notation is used:

When MMCR0[PMCC] = '11':

- PMCs refers to PMC1 - 4; PMC j or
- PMC j CE refers to PMC j or PMC j CE, respectively, where j = 2 - 4;

Otherwise:

- PMCs refers to PMC1 - 6 and PMC j or
- PMC j CE refers to PMC j or PMC j CE, respectively, where j = 2 - 6.

When MMCR0[PMCC] = '10' or '11', read/write access is provided to problem-state programs. Only FC, PMAE, and PMAO can be accessed. When **mtspr** is executed in problem state, all other bits are unchanged and any read or write attempt to these bits results in a facility unavailable interrupt.

The MMCR0 Register bits are defined in *Table 11-2* on page 276.



Table 11-2. MMCR0 Register (Sheet 1 of 4)

Bits	Field Name	Description
0:31	Reserved	Reserved.
32	FC	<p>Freeze counters.</p> <p>0 The PMCs are incremented (if permitted by other MMCR bits). 1 The PMCs are not incremented.</p> <p>The hardware sets this bit to '1' when an enabled condition or event occurs and MMCR0[FCECE] = '1'.</p> <p>Note: When MMCR0[PMCC] = '11', this bit affects PMC1 - 4; otherwise, it controls PMC1 - 6.</p>
33	FCS	<p>Freeze counters and BHRB in privileged state (FCS).</p> <p>0 The PMCs are incremented (if permitted by other MMCR bits) and entries are written into the BHRB (if permitted by the BHRB Instruction Filtering Mode field in the MMCRA). 1 The PMCs are not incremented and entries are not written into the BHRB if MSR[HV, PR] = '00'.</p> <p>Note: When MMCR0[PMCC] = '11', this bit affects PMC1 - 4; otherwise, it controls PMC1 - 6.</p>
34	FCP	<p>Conditionally freeze counters and BHRB in problem state (FCP).</p> <p>If FCPC = '0' (the value of bit 51), this field has the following meaning:</p> <p>0 The PMCs are incremented (if permitted by other MMCR bits) and entries are written into the BHRB (if permitted by the BHRB Instruction Filtering Mode field in the MMCRA). 1 The PMCs are not incremented and entries are not written into the BHRB if MSR[HV, PR] = '01'.</p> <p>If FCPC = '1', this field has the following meaning:</p> <p>0 The PMCs are not incremented and entries are not written into the BHRB if MSR[HV, PR] = '01'. 1 The PMCs are not incremented and the BHRB entries are not written if MSR[HV, PR] = '11'.</p> <p>Note: When MMCR0[PMCC] = '11', this bit affects PMC1 - 4; otherwise, it controls PMC1 - 6.</p>
35	FCM1	<p>Freeze counters when MSR[PMM] = '1'.</p> <p>0 The PMCs are incremented (if permitted by other MMCR bits). 1 The PMCs are not incremented if MSR[PMM] = '1'.</p> <p>Note: When MMCR0[PMCC] = '11', this bit affects PMC1 - 4; otherwise, it controls PMC1 - 6.</p>
36	FCM0	<p>Freeze counters when MSR[PMM] = '0'.</p> <p>0 The PMCs are incremented (if permitted by other MMCR bits). 1 The PMCs are not incremented if MSR[PMM] = '0'.</p> <p>Note: When MMCR0[PMCC] = '11', this bit affects PMC1 - 4; otherwise, it controls PMC1 - 6.</p>
37	PMAE	<p>Performance monitor alert enable.</p> <p>0 Performance monitor alerts are disabled. 1 Performance monitor alerts are enabled until a performance monitor alert occurs, at which time: MMCR0[PMAE] is set to '0' MMCR0[PMAO] is set to '1'</p> <p>Note: When MMCR0[PMCC] = '11', this bit affects PMC1 - 4; otherwise, it controls PMC1 - 6.</p> <p>Note: Software can set this bit and MMCR0[PMAO] to '0' to prevent performance monitor exceptions. Software can set this bit to '1' and then poll the bit to determine whether an enabled condition or event has occurred. This is especially useful for software that runs with MSR[EE] = '0'. In previous versions of the architecture that lacked the concept of performance monitor alerts, this bit was called Performance Monitor Exception Enable (PMXE).</p>
38	FCECE	<p>Freeze counters on enabled condition or event.</p> <p>0 The PMCs are incremented (if permitted by other MMCR bits). 1 The PMCs are incremented (if permitted by other MMCR bits) until an enabled condition or event occurs when MMCR0[TRIGGER] = '0', at which time MMCR0[FC] is set to '1'. If the enabled condition or event occurs when MMCR0[TRIGGER] = '1', the FCECE bit is treated as if it were '0'.</p> <p>Note: When MMCR0[PMCC] = '11', this bit affects PMC1 - 4; otherwise, it controls PMC1 - 6.</p>



Table 11-2. MMCR0 Register (Sheet 2 of 4)

Bits	Field Name	Description
39:40	TBSEL	<p>Time-base selector. This field selects the Time-Base bit that can cause a time-base transition event (the event occurs when the selected bit changes from '0' to '1').</p> <p>00 Time-Base bit 47 is selected. 01 Time-Base bit 51 is selected. 10 Time-Base bit 55 is selected. 11 Time-Base bit 63 is selected.</p> <p>When the selected time base transitions from '0' to '1', and the time-base event is enabled (MMCR0[TBEE] = '1'), and the performance monitor interrupt is enabled, a performance monitor interrupt occurs and the performance monitor interrupt is disabled (MMCR0[PMAE] = '0'). In multiple processor systems with the time-base registers synchronized among the processors, time-base transition events can be used to correlate the performance monitor data obtained by several processors if software has specified the same TBSEL value for all of the processors in the system. Even with multi-LPAR support, LPAR0 time-base bits are used so that all threads count events over the same periods and throw alerts at the same time.</p>
41	TBEE	<p>Time-base event enable.</p> <p>0 Time-base transition events are disabled. 1 Time-base transition events are enabled.</p>
42	BHRBA	<p>This field controls whether the BHRB instructions are available in problem state. If an attempt is made to execute a BHRB instruction in problem state when the BHRB instructions are not available, a facility unavailable interrupt occurs.</p> <p>0 Instructions mfbhrb and clrbhrb are not available in problem state. 1 Instructions mfbhrb and clrbhrb are available in problem state.</p>
43	EBE	<p>Performance monitor event-based branch enable. This field controls whether performance monitor event-based branches are enabled.</p> <p>0 Performance Monitor event-based branches are disabled. 1 Performance Monitor event-based branches are enabled.</p> <p>Note: Enabling event-based branches gives problem-state programs visibility to and control of MMCR0[PMAE] and MMCR0[PMAO]. This enables problem-state programs to recognize when performance monitor event-based exceptions have occurred and to re-enable performance monitor event-based exceptions. To enable a problem-state program to use the event-based branch facility for performance monitor events, software first initializes the performance monitor registers to values appropriate to the program, sets MMCR0[PMAE] and MMCR0[PMAO] to '0', sets BESCR to '0', and sets MMCR0[EBE] to '1'. MMCR0[PMCC] must also be set to '10' or '11' to give problem-state programs write access to the PMCs. If the event-based branch facility has not been enabled in the FSCR and HFSCR, it must be enabled in these registers as well.</p> <p>Note: When MMCR0[PMCC] = '11', this bit affects PMC1 - 4; otherwise, it controls PMC1 - 6.</p>
44:45	PMCC	<p>PMC control. This field controls the number of counters that are included in the performance monitor and the set of registers that are available to be read and written in problem state. In problem state, if an attempt to read a register that is unavailable to be read, or an attempt is made to write a register that is unavailable to be written, a facility unavailable interrupt occurs.</p> <p>00 PMC1 - 6 are included in the performance monitor, and SPRs 768 - 782 are available to be read in problem state but not available to be written in problem state. 01 PMC1 - 6 are included in the performance monitor, but SPRs 768 - 782 are not available to be read or written in problem state. 10 PMC1 - 6 are included in the performance monitor. Selected fields of MMCR0, MMCR2, MMCRA, and all bits of PMC1 - 6 are available to be read and written in problem state, and SIER, SDAR, and SIAR are read-only in problem state; all other SPRs in the range of 768 - 782 are not available to be read or written in problem state. 11 PMC1 - 4 are included in the performance monitor. Selected fields of MMCR0, MMCR2, MMCRA, and all bits of PMC1 - 4 are available to be read and written in problem state, and SIER, SDAR, SIAR are read-only in problem state; all other SPRs in the range of 768 - 782 are unavailable to be read or written in problem state.</p> <p>Note: When PMC5 and PMC6 are not part of the performance monitor (that is, when PMCC = '11'), they are controlled by bits in MMCRS, rather than MMCR0 and MMCR2. Counter negative conditions in PMC5 and PMC6 do not result in performance monitor alerts or exceptions and do not result in performance monitor interrupts.</p>



Table 11-2. MMCR0 Register (Sheet 3 of 4)

Bits	Field Name	Description
46	FCTM	Freeze counters when transactional memory is in the transactional state (MSR[TS] = '10'). 0 PMCs are incremented. 1 PMCs are not incremented when TM is in a transactional state. Note: When MMCR0[PMCC] = '11', this bit affects PMC1 - 4; otherwise, it controls PMC1 - 6.
47	FCNTS	Freeze counters transactional memory is in a nontransactional state (MSR[TS] = '00'). 0 PMCs are incremented (if permitted by other MMCR bits). 1 PMCs are not incremented when TM is in a nontransactional state. Note: When MMCR0[PMCC] = '11', this bit affects PMC1 - 4; otherwise, it controls PMC1 - 6.
48	PMC1CE	PMC1 condition enable. This bit determines whether the counter negative condition due to a negative value in PMC1 is enabled. 0 Disable PMC1 counter negative condition. 1 Enable PMC1 counter negative condition.
49	PMCjCE	PMCj condition enable. This bit determines whether the counter negative condition due to a negative value in PMCj ($2 \leq j$) is enabled. 0 Disable PMCj ($2 \leq j$) counter negative condition. 1 Enable PMCj ($2 \leq j$) counter negative condition. Note: The following notation is used. When MMCR0[PMCC] = '11', PMCs refers to PMC1 - 4 and PMCj or PMCjCE refers to PMCj or PMCjCE, respectively, where $j = 2 - 4$; otherwise, PMCs refers to PMC1 - 6 and PMCj or PMCjCE refers to PMCj or PMCjCE, respectively, where $j = 2 - 6$. Note: When MMCR0[PMCC] = '11', this bit affects PMC2 - 4; otherwise, it controls PMC2 - 6.
50	TRIGGER	Trigger enable. 0 The PMCs are incremented (if permitted by other MMCR bits). 1 PMC1 is incremented (if permitted by other MMCR bits). The PMCjs are not incremented until PMC1 is negative or an enabled condition or event occurs; at which time, the PMCjs resume incrementing (if permitted by other MMCR bits) and MMCR0[TRIGGER] is set to '0'. See the description of the FCECE bit, regarding the interaction between TRIGGER and FCECE. Case 1: Resume counting in the PMCjs when PMC1 becomes negative without causing a performance monitor interrupt. Then freeze all PMCs (and optionally cause a performance monitor interrupt) when a PMCj becomes negative. The PMCjs then reflect the events that occurred between the time when PMC1 became negative and the time a PMCj becomes negative. This use requires the following MMCR0 bit settings. <ul style="list-style-type: none">• TRIGGER = '1'• PMC1CE = '0'• PMCjCE = '1'• TBEE = '0'• FCECE = '1'• PMAE = '1' (if a performance monitor interrupt is required) Case 2: Resume counting in the PMCjs when PMC1 becomes negative and causes a performance monitor interrupt without freezing any PMCs. The PMCjs then reflect the events that occurred between the time PMC1 became negative and the time the interrupt handler reads them. This use requires the following MMCR0 bit settings. <ul style="list-style-type: none">• TRIGGER = '1'• PMC1CE = '1'• TBEE = '0'• FCECE = '0'• PMAE = '1' Note: When MMCR0[PMCC] = '11', this bit affects PMC1 - 4; otherwise, it controls PMC1 - 6.
51	FCPC	Freeze counters and BHRB in a problem-state condition. This bit controls the operation of bit 34 (FCP). Note: When MMCR0[PMCC] = '11', this bit affects PMC1 - 4; otherwise, it controls PMC1 - 6.
52	PMAOQ	0 PMU interrupt is asynchronous. 1 PMU interrupt is synchronous.



Table 11-2. MMCR0 Register (Sheet 4 of 4)

Bits	Field Name	Description
53	Reserved	Reserved.
54	PMAOQ_TRC	Qualifier for BHRB trace mode interrupt.
55	CC56WAIT	Freeze counters 5 and 6 in wait state. When MMCR0[PMCC] = '11', the setting of this bit has no effect; otherwise, it is defined as follows: 0 PMC5 and PMC6 are incremented only when CTRL[RUN] = '1', if permitted by other MMCR bits. Software is expected to set CTRL[RUN] = '0' when it is in a wait state. For example, when there is no process ready to run. 1 PMC5 and PMC6 are incremented regardless of the state of CTRL[RUN].
56	PMAO	Performance monitor alert has occurred. 0 A performance monitor event has not occurred since the last time software set this bit to '0'. 1 A performance monitor event has occurred since the last time software set this bit to '0'. This bit is set to '1' by the processor when a performance monitor event occurs. This bit can be set to '0' only by the mtspr instruction. When MMCR0[EBC] = '1', setting BESCR[PME0] = '0' or '1' sets PMAO to '0' or '1', respectively. Software can set this bit to '1' to simulate the occurrence of a performance monitor event. Software should set this bit to '0' after handling the performance monitor event. Note: When MMCR0[PMCC] = '11', this bit affects PMC1 - 4; otherwise, it controls PMC1 - 6.
57	FCSS	Freeze counters when transactional memory is in a suspended state. MSR[TS] = '01'. 0 PMCs are incremented (if permitted by other MMCR bits). 1 PMCs are not incremented when transactional memory is in suspended state. Note: When MMCR0[PMCC] = '11', this bit affects PMC1 - 4; otherwise, it controls PMC1 - 6.
58	FC1-4	Freeze counters 1 - 4. 0 PMC1 - PMC4 are incremented (if permitted by other MMCR bits). 1 PMC1 - PMC4 are not incremented.
59	FC5-6	Freeze counters 5 - 6. 0 PMC5 and PMC6 are incremented (if permitted by other MMCR bits). 1 PMC5 and PMC6 are not incremented.
60:61	Reserved	Reserved.
62	FCWAIT	Freeze counters in wait state. 0 The PMCs are incremented (if permitted by other MMCR bits). 1 The PMCs are not incremented if CTRL[RUN] = 0. Software is expected to set CTRL[RUN] = 0 when it is in a wait state; that is, when there is no process ready to run. Note: When MMCR0[PMCC] = '11', this bit affects PMC1 - 4; otherwise, it controls PMC1 - 6. Note: PM_CYC counts cycles regardless of the run latch.
63	FCH	Freeze counters and BHRB in hypervisor state. 0 The PMCs are incremented (if permitted by other MMCR bits) and BHRB entries are written (if permitted by MMCRA bits). 1 The PMCs are not incremented and BHRB entries are not written if MSR[HV, PR] = '10'. Note: When MMCR0[PMCC] = '11', this bit affects PMC1 - 4; otherwise, it controls PMC1 - 6.



11.6.1.3 Monitor Mode Control Register 1 (MMCR1)

Monitor Mode Control Register 1 (MMCR1) is a 64-bit register.

Events due to randomly sampled instructions only occur if random sampling is enabled (MMCRA[SE] = '1'); all other events occur whenever the event specification is met regardless of the value of MMCR0[SE].

Table 11-3. MMCR1 Register (Sheet 1 of 2)

Bits	Field Name	Description	
0:3	PMC1UNIT	0000	Reserved
		0001	Reserved
		0010	ISU0
		0011	ISU1
		0100	IFU0
		0101	IFU1
		0110	L2 Bus0
		0111	L2 Bus1
		1000	L3 Bus0
		1001	L3 Bus1
		1010	VSU0
		1011	VSU1
		1100	LSU0
		1101	LSU1
		1110	LSU2
		1111	LSU3
4:7	PMC2UNIT	PMC2 events unit selector (same encoding as PMC1UNIT).	
8:11	PMC3UNIT	PMC3 events unit selector (same encoding as PMC1UNIT).	
12:15	PMC4UNIT	PMC4 events unit selector (same encoding as PMC1UNIT).	
16	DC_RLD_QUAL	Bit defines data cache reload qualifier. 0 Counts only demand reloads to the L1 data cache. 1 Counts all reloads to the L1 data cache (demand and prefetch). Applies to all PM_DATA_FROM_* events. See <i>Appendix D Performance Monitoring Events</i> on page 387.	
17	IC_RLD_QUAL	Bit defines instruction cache reload qualifier. 0 Counts only demand reloads to the L1 instruction cache. 1 Counts all reloads to the L1 instruction cache (demand and prefetch). Applies to all PM_INST_FROM_* events. See <i>Appendix D Performance Monitoring Events</i> on page 387.	
18:19	Reserved	Reserved.	
20:24	FAB_CRESP_MATCH	Bits to match a specific combined response from the fabric.	
25:27	FAB_TYPE_MATCH	Bits to match on a specific ttype group from the fabric.	
28	PMC1COMB	PMC1 event-pair combination. When counting event-bus events in PMC1, this bit enables the event-pair combination circuitry. 0 Individual event (PMC1SEL[6] low selects bit 0 of the pair, high selects bit 1 of the pair). 1 Combined event (PMC1SEL[6] low selects to ADD the pair, high selects to AND the pair).	
29	PMC2COMB	PMC2 event-pair combination. When counting event-bus events in PMC2, this bit enables the event-pair combination circuitry. 0 Individual event (PMC2SEL[6] low selects bit 0 of the pair, high selects bit 1 of the pair). 1 Combined event (PMC2SEL[6] low selects to ADD the pair, high selects to AND the pair).	

Table 11-3. MMCR1 Register (Sheet 2 of 2)

Bits	Field Name	Description
30	PMC3COMB	PMC3 event-pair combination. When counting event-bus events in PMC3, this bit enables the event-pair combination circuitry. 0 Individual event (PMC3SEL[6] low selects bit 0 of the pair, high selects bit 1 of the pair). 1 Combined event (PMC3SEL[6] low selects to ADD the pair, high selects to AND the pair).
31	PMC4COMB	PMC4 event-pair combination. When counting event-bus events in PMC4, this bit enables the event-pair combination circuitry. 0 Individual event (PMC4SEL[6] low selects bit 0 of the pair, high selects bit 1 of the pair). 1 Combined event (PMC4SEL[6] low selects to ADD the pair, high selects to AND the pair).
32:39	PMC1SEL	PMC1 event selector. The value in this bit field combined with MMCR1[0:31] determines which event is counted. x'10' PMC1 counts PMC4 overflows. x'24' PMC1 counts PMC5 overflows. When counting overflows, freeze conditions for both counters must be identical. Note: PMCxSEL(0:1) = '10' selects event-bus events, PMCxSEL(0:2) = '111' for compatibility direct events, PMCxSEL(0) = '0' for other direct events. When PMCxSEL(7) is high, the falling edges of an event are counted rather than the event itself. Changes from nonzero-to-zero is the edge for event-bus events.
40:47	PMC2SEL	PMC2 event selector. The value in this bit field combined with MMCR1[0:31] determines which event is counted. x'10' PMC2 counts PMC1 overflows. When counting overflows, freeze conditions for both counters must be identical. Note: PMCxSEL(0:1) = '10' selects event-bus events, PMCxSEL(0:2) = '111' for compatibility direct events, PMCxSEL(0) = '0' for other direct events. When PMCxSEL(7) is high, the falling edges of an event are counted rather than the event itself. Changes from nonzero-to-zero is the edge for event-bus events.
48:55	PMC3SEL	PMC3 event selector. The value in this bit field combined with MMCR1[0:31] determines which event is counted. x'10' PMC3 counts PMC2 overflows. x'24' PMC3 counts PMC6 overflows. When counting overflows, freeze conditions for both counters must be identical. Note: PMCxSEL(0:1) = '10' selects event-bus events, PMCxSEL(0:2) = '111' for compatibility direct events, PMCxSEL(0) = '0' for other direct events. When PMCxSEL(7) is high, the falling edges of an event are counted rather than the event itself. Changes from nonzero-to-zero is the edge for event-bus events.
56:63	PMC4SEL	PMC4 event selector. The value in this bit field combined with MMCR1[0:31] determines which event is counted. x'10' PMC4 counts PMC3 overflows. When counting overflows, freeze conditions for both counters must be identical. Note: PMCxSEL(0:1) = '10' selects event-bus events, PMCxSEL(0:2) = '111' for compatibility direct events, PMCxSEL(0) = '0' for other direct events. When PMCxSEL(7) is high, the falling edges of an event are counted rather than the event itself. Changes from nonzero-to-zero is the edge for event-bus events.

Note: To avoid a spurious or phantom count during the transition, event counting on PMC1 - 4 suspends counting for one cycle during SPR writes to MMCR1.



Table 11-4. MMCR1 PMCxSEL Selection of Direct Events versus Event-Bus Events

PMCxSEL(0:3)	Event Type
0000	Direct Events
0001	
0010	
0011	
0100	
0101	
0110	
0111	
1000	Event Bus Events
1001	
1010	
1011	
1100	Reserved
1101	
1110	Compatibility Direct Events
1111	

11.6.1.4 Monitor Mode Control Register 2 (MMCR2)

Monitor Mode Control Register 2 (MMCR2) is a 64-bit register that contains 9-bit fields for controlling the operation of PMC1 - PMC6.

A single bit in each Cn field of MMCR2 is described in the following bit description table. When MMCR0[PMCC] = '11', fields C1 - C4 control the operation of PMC1 - PMC4, respectively, and fields C5 and C6 are meaningless; otherwise, fields C1 - C6 control the operation of PMC1 - PMC6, respectively. The bit definitions of each Cn field are as follows, where n = 1 - 6.

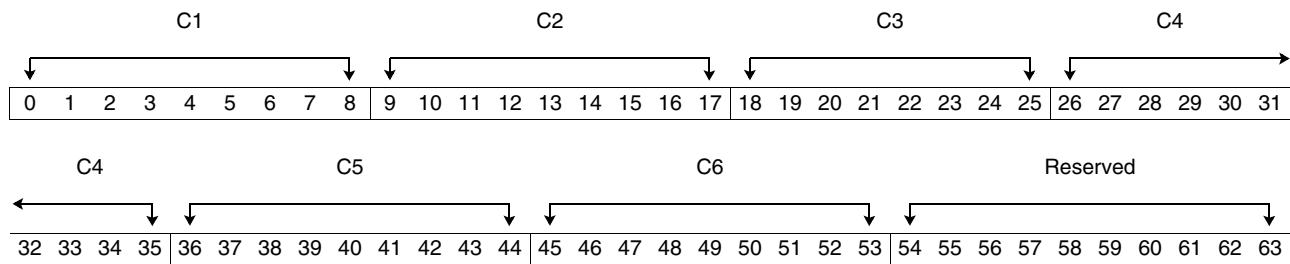


Table 11-5. MMCR2 Register

Bits	Field Name	Description
0	FCnS	Freeze counter n in privileged state (FCnS). 0 PMCr is incremented (if permitted by other MMCR bits). 1 PMCr is not incremented if MSR[HV, PR] = '0'.
1	FC1nP	Freeze counter n in problem state (FCnP). 0 PMCr is incremented (if permitted by other MMCR bits). 1 If FCnPC = '0', PMCr is not incremented if MSR[PR] = '1'.
2	FCnPC	Freeze counter in either problem or an adjunct state (FCnPC). <u>FCnP</u> <u>FCnPC</u> 0 0 PMCr is incremented (if permitted by other MMCR bits). 1 0 Freeze when MSR[HV] = X, MSR[PR] = 1 (problem state or adjunct). 0 1 Freeze when MSR[HV] = 0, MSR[PR] = 1 ("true" problem state). 1 1 Freeze when MSR[HV] = 1, MSR[PR] = 1 (adjunct).
3	FCnM1	Freeze counter n while Mark = '1' (FCnM1). 0 PMCr is incremented (if permitted by other MMCR bits). 1 PMCr is not incremented if MSR[PMM] = '1'.
4	FCnM0	Freeze counter n while Mark = '0' (FCnM0). 0 PMCr is incremented (if permitted by other MMCR bits). 1 PMCr is not incremented if MSR[PMM] = '0'.
5	FCnWAIT	Freeze counter n in wait state (FCnWAIT). 0 PMCr is incremented (if permitted by other MMCR bits). 1 PMCr is not incremented if CTRL[RUN] = '0'. Software is expected to set CTRL[RUN] = '0' when it is in a wait state; that is, when there is no process ready to run. Note: PM_CYC counts cycles regardless of the run latch.
6	FCnH	Freeze counter n in hypervisor state (FCnH). 0 PMCr is incremented (if permitted by other MMCR bits). 1 PMCr is not incremented if MSR[HV, PR] = '10'.
7:8	Reserved	Reserved.



11.6.1.5 Monitor Mode Control Register A (MMCRA)

Monitor Mode Control Register A (MMCRA) is a 64-bit register. MMCRA gives privileged programs the ability to control the sampling process and threshold events.

Table 11-6. MMCRA Register (Sheet 1 of 2)

Bits	Field Name	Description
0:15	Reserved	Reserved.
16	SIER_CTRL	When this bit is set to '1', SIER bits 16:28 are changed to data real address bits 44:56.
17:19	Reserved	Reserved.
20:21	SDAR_MODE	Continuous sampling. These bits specify how the SDAR should be updated in continuous-sampling mode. 00 No updates. 01 Continuous-sampling mode updates SDAR on a TLB miss. 10 Continuous-sampling mode updates SDAR on a D-cache miss. 11 Continuous-sampling mode updates SDAR on a store issue.
22:24	THRESH_CMP_EXP	Three-bit exponent portion of a threshold compare floating point.
25:31	THRESH_CMP_MANTISSA	Seven-bit mantissa portion of a threshold compare floating point. Note: The mantissa upper two bits should never be zero unless the exponent being written is also zero. This is an artifact of the "power of 4" floating-point counter. The floating-point mantissa starts counting events starting from '000000' and when it reaches '1111111', it rolls over to '010000000' where the exponent is incremented by 1.
32:33	BHRB_FILTER	BHRB instruction filtering mode (IFM). This field controls the filter criterion used by hardware when recording branch instructions that satisfy the BHRB privilege filtering mode requirements into the BHRB. 00 No filtering. 01 Record only bl , bcl , bcirl , bctr , and bctarl instructions. 10 Do not record the b , bl , bc , and bcl instructions for which the BO field indicates "branch always." For the bclr , bcirl , bctr , bcctl , bctarl , and bctarl instructions for which the BO field indicates "branch always," record only one entry containing the branch target address. 11 Filter and enter BHRB entries as for mode '10', but do not record branch instructions for which BO0 = '1', or for which the "a" bit in the BO field is set to '1'. Note: The filters provided by the IFM field can be restated in terms of the operation performed as follows: 01 Record only branch instructions that have the link bit set. 10 Record only the target addresses of conditional branch instructions and XL-form unconditional branch instructions. 11 Filter as for encoding '10', but do not record instructions that provide a hint or that do not depend on the value of CRBI.
34:36	THRESH_CTR_EXP	Three-bit exponent portion of the threshold floating-point counter.
37	Reserved	Reserved.
38:44	THRESH_CTR_MANTISSA	Seven-bit mantissa portion of the threshold floating-point counter.

Table 11-6. MMCRA Register (Sheet 2 of 2)

Bits	Field Name	Description	
45:47	THRESH_EVENT_SEL	Selection for an event specified for threshold counting. PMC1 - 4 event counting, selected as shown, adheres to possible freeze conditions set up in the MMCR0, MMCR2, and MMCRC Registers.	
		000	Do not count, disable thresholding.
		001	Counts number of cycles that the CTRL run latch is set (does not depend on freeze conditions).
		010	Instructions completed while the CTRL run latch is set (does not depend on freeze conditions).
		011	Reserved.
		100	Event configured in PMC1 (depends on freeze conditions).
		101	Event configured in PMC2 (depends on freeze conditions).
		110	Event configured in PMC3 (depends on freeze conditions).
		111	Event configured in PMC4 (depends on freeze conditions).
48:51	THRESH_START	Threshold start event. See Table 11-7 on page 285 for details.	
52:55	THRESH_END	Threshold end events. See Table 11-7 on page 285 for details.	
56	Reserved	Reserved.	
57:59	RAND_SAMP_ELIG	Eligibility criteria. See Table 11-8 on page 286 for details.	
60	SAMPLE_RESET	Any time this bit is set to '1', the sampling state machine is forced from the start state into the idle state. If the state machine is not in the start state when this bit is set, it has no effect.	
61:62	RAND_SAMP_MODE	Random sampling mode (SM).	
		00	Random instruction sampling (RIS): Instructions that meet the criterion specified in the RAND_SAMP_ELIG field for random instruction sampling are randomly selected and sampled. See Table 11-8 on page 286 for details.
		01	Random load/store sampling (LSS): Events that meet the criterion specified in the RAND_SAMP_ELIG field for random load/store sampling are randomly selected for sampling. See Table 11-8 on page 286 for details.
		10	Random branch facility sampling (BFS): Events that meet the criterion specified in the RAND_SAMP_ELIG field for random branch facility sampling are randomly selected for sampling. See Table 11-8 on page 286 for details.
		11	Reserved.
63	SAMPLE_ENABLE	0	Continuous sampling.
		1	Random sampling.

Table 11-7. Threshold Start/Stop Event Selection

MMCRA[48] MMCRA[52]	MMCRA[49:51] MMCRA[53:55]	Description	Architected
0	000	Reserved.	Yes
0	001	Sampled instruction decoded.	Yes
0	010	Sampled instruction dispatched.	Yes
0	011	Sampled instruction issued.	Yes
0	100	Sampled instruction finished.	Yes
0	101	Sampled instruction completed.	Yes
0	110	Sampled instruction L1 load miss.	Yes
0	111	Sampled instruction L1 reload valid	Yes
1	000	Event selected in MMCR1 for PMC1 counting.	No
1	001	Event selected in MMCR1 for PMC2 counting.	No



Table 11-7. Threshold Start/Stop Event Selection

MMCRA[48] MMCRA[52]	MMCRA[49:51] MMCRA[53:55]	Description	Architected
1	010	Event selected in MMCR1 for PMC3 counting.	No
1	011	Event selected in MMCR1 for PMC4 counting.	No
1	100	Sampled group next to complete.	No
1	101	RC machine dispatched for sampled instruction.	No
1	110	RC machine done for sampled instruction.	No
1	111	Reserved.	No

Table 11-8. Random Sampling Eligibility Criteria

MMCRA[63]	MMCRA[61:62]	MMCRA[57]	MMCRA[58:59]	Eligibility Criteria	Architected
Random Instruction Sampling					
1	00	0	00	Random.	Yes
1	00	0	01	Load/store. Any operation that gets routed to the LSU (for example, load, store, move from/to LSU SPR).	Yes
1	00	0	10	Probe NOP.	Yes
1	00	0	11	Reserved.	Yes
1	00	1	00	IMC.	No
1	00	1	01	IMC + random.	No
1	00	1	10	Long latency operation (div, sqrt, mul, mtcdr, brLK = '1').	No
1	00	1	11	Reserved.	No
Random Event Sampling (LSU)					
1	01	0	00	Load misses.	Yes
1	01	0	01	Reserved.	Yes
1	01	0	10	Reserved.	Yes
1	01	0	11	Reserved.	Yes
1	01	1	00	Larx/stcx.	No
1	01	1	01	Prefetch target tracker. Mark loads that match prefetch effective address currently being tracked.	No
1	01	1	10	Reserved.	No
1	01	1	11	Reserved.	No
Random Event Sampling (BRU)					
1	10	0	00	Branch mispredicts.	Yes
1	10	0	01	Branch mispredicts (<u>CR</u>).	Yes
1	10	0	10	Reserved.	Yes
1	10	0	11	Taken branches.	Yes
1	10	1	00	Nonrepeating branches.	No



Table 11-8. Random Sampling Eligibility Criteria

MMCRA[63]	MMCRA[61:62]	MMCRA[57]	MMCRA[58:59]	Eligibility Criteria	Architected
1	10	1	01	All branches that require prediction.	No
1	10	1	10	Reserved.	No
1	10	1	11	Reserved.	No



11.6.1.6 Core Monitor Mode Control Register (MMCRC)

This SPR is not replicated on a per split-core basis. This is a core-level SPR (hypervisor access only) that is used for a variety of purposes.

Table 11-9. MMCRC Register (Sheet 1 of 2)

Bits	Field Name	Description
0:42	Reserved	Reserved.
43:45	BHRB_TID	<p>Thread ID for BHRB_MODE. When MMCRC[BHRB_MODE] is set to private mode, this field specifies the thread that owns the BHRB. All other threads do not write the BHRB and their move-from-BHRB entry instructions always return 0's.</p> <p>In split-core mode, physical thread IDs map to LPARs as follows:</p> <ul style="list-style-type: none"> LPAR0: T0/T1 LPAR1: T2/T3 LPAR2: T4/T5 LPAR3: T6/T7
46	BHRB_MODE	<p>0 Shared mode: All threads write to the BHRB. The BHRB entries are split among the number of threads specified by the SMT mode. (<u>SMT</u> = 32 entries, SMT2 = 16 entries per thread, SMT4 = 8 entries per thread, and SMT8 = 4 entries per thread).</p> <p>1 Private mode: Only the thread specified by MMCRC[BHRB_Thread] can write to the BHRB. All other threads do not write to the BHRB and their move-from BHRB entry instructions always return '0'.</p>
47:49	L3EVENT_SEL0	Event Bus0 control.
50:51	L3EVENT_SEL1	Event Bus1 control.
52	INT_MODE	<p>An L2 control that changes S versus T to early versus late.</p> <p>When this bit is enabled, the events that are called PM_DATA_FROM_*_SHR are interpreted as PM_DATA_FROM_*_EARLY and events called PM_DATA_FROM_*_MOD are interpreted as PM_DATA_FROM_*_LATE.</p>
53:55	L2EVENT_SEL	L2 bus event control.
56	GLOB_FRZ_SPMC_01	<p>Global freeze for SPMC T0/T1.</p> <p>Note: When MMCR0[PMCC] = '11', this bit also affects PMC5 - 6. In 4 LPAR mode, this maps to LPAR0. In 2 LPAR mode, this maps to LPAR0.</p>
57	GLOB_FRZ_SPMC_23	<p>Global freeze for SPMC T2/T3.</p> <p>Note: When MMCR0[PMCC] = '11', this bit also affects PMC5 - 6. In 4 LPAR mode, this maps to LPAR1. In 2 LPAR mode, this maps to LPAR0.</p>
58	GLOB_FRZ_SPMC_45	<p>Global freeze for SPMC T4/T5.</p> <p>Note: When MMCR0[PMCC] = '11', this bit also affects PMC5 - 6. In 4 LPAR mode, this maps to LPAR2. In 2 LPAR mode, this maps to LPAR1.</p>
59	GLOB_FRZ_SPMC_67	<p>Global freeze for SPMC T6/T7.</p> <p>Note: When MMCR0[PMCC] = '11', this bit also affects PMC5 - 6. In 4 LPAR mode, this maps to LPAR3. In 2 LPAR mode, this maps to LPAR1.</p>
60	GLOB_FRZ_PMC_01	<p>Global freeze for thread level PMCs T0/T1.</p> <p>Note: When MMCR0[PMCC] = '11', this bit affects PMC1 - 4; otherwise, it controls PMC1 - 6. In 4 LPAR mode, this maps to LPAR0. In 2 LPAR mode, this maps to LPAR0.</p>
61	GLOB_FRZ_PMC_23	<p>Global freeze for thread level PMCs T2/T3.</p> <p>Note: When MMCR0[PMCC] = '11', this bit affects PMC1 - 4; otherwise, it controls PMC1 - 6. In 4 LPAR mode, this maps to LPAR1. In 2 LPAR mode, this maps to LPAR0.</p>
62	GLOB_FRZ_PMC_45	<p>Global freeze for thread level PMCs T4/T5.</p> <p>Note: When MMCR0[PMCC] = '11', this bit affects PMC1 - 4; otherwise, it controls PMC1 - 6. In 4 LPAR mode, this maps to LPAR2. In 2 LPAR mode, this maps to LPAR1.</p>

Table 11-9. MMCRC Register (Sheet 2 of 2)

Bits	Field Name	Description
63	GLOB_FRZ_PMC_67	Global freeze for thread level PMCs T6/T7. Note: When MMCR0[PMCC] = '11', this bit affects PMC1 - 4; otherwise, it controls PMC1 - 6. In 4 LPAR mode, this maps to LPAR3. In 2 LPAR mode, this maps to LPAR1.

11.7 Hypervisor Performance Monitor

The POWER8 core has four dedicated 64-bit hypervisor performance counters, HPMC1 - HPMC4, that are shared between all threads. These counters are only available for use by the hypervisor. The hypervisor performance monitor is controlled through MMCRH. The configuration and operation of the hypervisor performance monitor is similar to that of the thread-level counters, PMC1 - PMC6, with the following differences:

- HPMC1 - HPMC4, and MMCRH are hypervisor-only resources. They can only be accessed in hypervisor mode (MSR[HV, PR] = '10'). Any attempt to read or write to these SPRs when not in hypervisor mode results in an exception.
- The hypervisor performance monitor has access to a limited subset of events. In particular, there is no access to event-bus events.
- HPMC1 counts events selected by MMCRH[HPMC1SEL] and HPMC2 counts events selected by MMCRH[HPMC2SEL]. Counting is enabled on HPMC1 and HPMC2 when MMCRH[HVFC], the freeze bit, is low and the other MMCRH enable counter conditions are met.
- HPMC3 counts instructions completed for threads when the CTRL Register run latch bits are set, the LPCR[ONL] bits are set, and the MMCRH[HVECLP] LPAR conditions are met. No MMCRH freeze or enable counter bits affect HPMC3 counting other than MMCRH[HVECLP].
- HPMC4 counts cycles for threads when the CTRL Register run latch bits are set, the LPCR[ONL] bit is set, and the MMCRH[HVECLP] LPAR conditions are met. No MMCRH freeze or enable counter bits affect HPMC4 counting other than MMCRH[HVECLP].
- There is no hypervisor performance monitor interrupt. Hypervisor performance monitor alerts can be enabled through the appropriate controls in MMCRH. Hypervisor performance monitor alerts behave similarly to the performance monitor (PMAO) alert, which is enabled through the thread-level performance monitor (using MMCR0), except that no interrupt is generated.



11.7.1 Hypervisor Performance Monitor Counters (HPMC1 - 4)

The Hypervisor Performance Monitor Counter Registers are defined in *Table 11-10*.

Table 11-10. Hypervisor Performance Monitor Counter Register

Bits	Field Name	Description
0	CTR_NEG	Counter negative bit.
1:63	CTRDATA	Counter data.

11.7.2 Monitor Mode Control Register H (MMCRH)

For HPMC1 or HPMC2 to count the selected event, the HVFC freeze bit must be low and at least one counter-enable bit must be high. *Table 11-11* describes the MMCRH Register.

Table 11-11. Monitor Mode Control Register H Register (Sheet 1 of 2)

Bits	Field Name	Description
0:31	Reserved	Reserved.
32	HVFC	Hypervisor freeze counters. 0 The HPMCs are incremented if not inhibited by another condition. 1 The HPMCs are not incremented. Note: This bit resets to '1', such that counters are always frozen after <u>POR</u> . See HVFCECE, bit 47, for additional information.
33	HVPMAE	Hypervisor performance monitor alert enable. 0 Hypervisor performance monitor alerts are disabled. 1 Hypervisor performance monitor alerts are enabled until a hypervisor performance monitor alert occurs, at which time the hardware disables the hypervisor performance monitor alert (MMCRH[HVPMAE] is set to '0'). Note: Software can set PMAE to '1' and then poll the bit to determine whether an enabled condition or event has occurred.
34	HVPMAO	Hypervisor performance monitor alert has occurred. 0 A hypervisor performance monitor alert has not occurred since the last time software set this bit to '0'. 1 A hypervisor performance monitor alert has occurred since the last time software set this bit to '0'. This bit is set to '1' by the hardware when a hypervisor performance monitor alert occurs and MMCRH[HVPMAE] = '1'. It can be set to '0' only by the mtspr instruction. Software should set this bit to '0' after handling a performance monitor alert. Note: Software can set this bit to '1' to simulate the occurrence of a performance monitor alert.
35	HVECONL	Enable hypervisor counters according to the T[0:7]_LPCR Register ONL bits. 0 The HPMCs are enabled (unless inhibited by another MMCRH condition). 1 The HPMCs are enabled for counting for threads with the ONL bits set (unless inhibited by another MMCRH condition).
36	HVECRL	Enable hypervisor counters according to the CTRL Register run latch bits. 0 The HPMCs are enabled (unless inhibited by another MMCRH condition). 1 The HPMCs are enabled for counting for threads with the run latch bits set (unless inhibited by another MMCRH condition).

Table 11-11. Monitor Mode Control Register H Register (Sheet 2 of 2)

Bits	Field Name	Description
37:38	HVECLP	<p>Enable hypervisor counters in 4 LPAR mode according to the following bits:</p> <ul style="list-style-type: none"> 00 The HPMCs are enabled for events on threads 0 and 1 (see note). 01 The HPMCs are enabled for events on threads 2 and 3 (see note). 10 The HPMCs are enabled for events on threads 4 and 5 (see note). 11 The HPMCs are enabled for events on threads 6 and 7 (see note). <p>Enable hypervisor counters in 2 LPAR mode according to the following bits:</p> <ul style="list-style-type: none"> 0X The HPMCs are enabled for events on threads 0 - 3 (see note). 1X The HPMCs are enabled for events on threads 4 - 7 (see note). <p>Note: Unless inhibited by another MMCRH condition.</p>
39	HVECH	<p>Enable hypervisor counters when in hypervisor state.</p> <ul style="list-style-type: none"> 0 The HPMCs are enabled (see note). 1 The HPMCs are enabled for counting for threads with MSR[HV, PR] = '10' (see note). <p>Note: If multiple MMCRH[HVECH, HVECP, HVECS, or HVECA] bits are set, counting is enabled when any of these conditions being selected are met, unless inhibited by another MMCRH condition.</p>
40	HVECP	<p>Enable hypervisor counters when in problem state.</p> <ul style="list-style-type: none"> 0 The HPMCs are enabled (see note). 1 The HPMCs are enabled for counting for threads with MSR[HV, PR] = '01' (see note). <p>Note: If multiple MMCRH[HVECH, HVECP, HVECS, or HVECA] bits are set, counting is enabled when any of these conditions being selected are met, unless inhibited by another MMCRH condition.</p>
41	HVEC	<p>Enable hypervisor counters when in supervisor state.</p> <ul style="list-style-type: none"> 0 The HPMCs are enabled (see note). 1 The HPMCs are enabled for counting for threads with MSR[HV, PR] = '00' (see note). <p>Note: If multiple MMCRH[HVECH, HVECP, HVECS, or HVECA] bits are set, counting is enabled when any of these conditions being selected are met, unless inhibited by another MMCRH condition.</p>
42	HVECA	<p>Enable hypervisor counters when in the adjunct state.</p> <ul style="list-style-type: none"> 0 The HPMCs are enabled (see note). 1 The HPMCs are enabled for counting for threads with MSR[HV, PR] = '11' (see note). <p>Note: If multiple MMCRH[HVECH, HVECP, HVECS, or HVECA] bits are set, counting is enabled when any of these conditions being selected are met, unless inhibited by another MMCRH condition.</p>
43:44	Reserved	Reserved.
45:46	HVECUS	<p>Enable hypervisor counters correlate to the user state, MSR[US].</p> <ul style="list-style-type: none"> 00 The HPMCs are enabled (see note). 01 The HPMCs are enabled in user state, MSR[US] = '1' (see note). 10 The HPMCs are enabled when not in user state, MSR[US] = '0' (see note). 11 Reserved. <p>Note: Unless inhibited by another MMCRH condition.</p>
47	HVFCECE	<p>Freeze counters on enabled condition or event.</p> <ul style="list-style-type: none"> 0 The HPMCs are incremented (if permitted by other MMCRH bits). 1 The HPMCs are incremented (if permitted by other MMCRH bits) until an enabled condition or event occurs at which time MMCRH[HVFC] is set to '1'.
48:56	HPMC1SEL	HPMC1 event selector. The value in this bit field determines which event is counted by HPMC1. HVFC must be set at least one cycle before and held for at least two cycles after setting a new HPMC1SEL.
56:63	HPMC2SEL	HPMC2 event selector. The value in this bit field determines which event is counted by HPMC2. HVFC must be set at least one cycle before and held for at least two cycles after setting a new HPMC2SEL.



11.8 Supervisor Performance Monitor

To support dynamic optimization efforts, the operating system (AIX/Linux/iOS) might require a dedicated facility for performance monitoring because the regular PMU can be used by user-level programs such as [JITs](#) or static compilers. The SPMCs, like the PMCs, are thread-level counters. However, SPMCs are only accessible to privileged code. Any attempt to read or write the SPMCs or MMCRS Register in problem state results in an exception.

The following SPRs per thread are available:

- SPMC1: a 32-bit counter
- SPMC2: a 32-bit counter
- MMCRS: a 16-bit configuration register

The SPMCs can select from all the compatibility events and a selected set of events to be used for dynamic optimization. Freeze conditions in MMCRS apply to the counting of SPMC1 - 2. When MMCR0[PMCC] = '11', PMC5 - 6 are allocated to supervisor control and conditions in MMCRS (rather than conditions MMCR0 or MMCR2) apply to PMC5 - 6.

11.8.1 Supervisor Performance Monitor Counters (SPMC1 - 2)

The Supervisor Performance Monitor Counters are defined in *Table 11-12*.

Table 11-12. Supervisor Performance Monitor Counter

Bits	Field Name	Description
32	CTR_NEG	Counter negative bit.
33:63	CTRDATA	Counter data.

11.8.2 Monitor Mode Control Register S (MMCRS) Register

The Monitor Mode Control Register S (MMCRS) Register is defined in *Table 11-13*.

Table 11-13. Monitor Mode Control Register S Register (Sheet 1 of 3)

Bits	Field Name	Description
0:31	Reserved	Reserved.
32	FC	Freeze counters. 0 The SPMCs are incremented if permitted by other MMCR bits. 1 The SPMCs are not incremented. The processor sets this bit to '1' when an enabled condition or event occurs and MMCRS[FCECE] = '1'.
33	FCS	Freeze counters in privileged state. 0 The SPMCs are incremented if permitted by other MMCRS bits. 1 The SPMCs are not incremented if MSR[HV, PR] = '00'.

Table 11-13. Monitor Mode Control Register S Register (Sheet 2 of 3)

Bits	Field Name	Description															
34	FCP	<p>Freeze counter in problem state.</p> <table> <tr> <td><u>FCP</u></td><td><u>FCPC</u></td><td></td></tr> <tr> <td>0</td><td>0</td><td>SPMCn is incremented (if permitted by other MMCR bits)</td></tr> <tr> <td>1</td><td>0</td><td>Freeze when MSR[HV] = 'X', MSR[PR] = '1'</td></tr> <tr> <td>0</td><td>1</td><td>Freeze when MSR[HV] = '0', MSR[PR] = '1'</td></tr> <tr> <td>1</td><td>1</td><td>Freeze when MSR[HV] = '1', MSR[PR] = '1'</td></tr> </table> <p>Note: To freeze the counters in problem state regardless of MSR[HV], MMCRS[FCPC] must be set to '0' and MMCRS[FCP] must be set to '1'.</p>	<u>FCP</u>	<u>FCPC</u>		0	0	SPMCn is incremented (if permitted by other MMCR bits)	1	0	Freeze when MSR[HV] = 'X', MSR[PR] = '1'	0	1	Freeze when MSR[HV] = '0', MSR[PR] = '1'	1	1	Freeze when MSR[HV] = '1', MSR[PR] = '1'
<u>FCP</u>	<u>FCPC</u>																
0	0	SPMCn is incremented (if permitted by other MMCR bits)															
1	0	Freeze when MSR[HV] = 'X', MSR[PR] = '1'															
0	1	Freeze when MSR[HV] = '0', MSR[PR] = '1'															
1	1	Freeze when MSR[HV] = '1', MSR[PR] = '1'															
35	FCPC	Freeze counters in problem state condition. This bit controls the operation of FCP (bit 34).															
36	FCH	Freeze counters in hypervisor state. 0 The SPMCs are incremented (if permitted by other MMCRS bits). 1 The SPMCs are not incremented if MSR[HV, PR] = '10'.															
37	FCM1	Freeze counters when MSR[PMM] = '1'. 0 The SPMCs are incremented. 1 The SPMCs are not incremented when MSR[PMM] = '1'.															
38	FCM0	Freeze counters when MSR[PMM] = '0'. 0 The SPMCs are incremented. 1 The SPMCs are not incremented when MSR[PMM] = '0'.															
39	FCTM	Freeze counters in <u>TM</u> transactional state. 0 SPMCs are incremented. 1 SPMCs are not incremented when TM is in transactional state (MSR[TM] = '10').															
40	FCTS	Freeze counters in TM suspend state. 0 SPMCs are incremented. 1 SPMCs are not incremented when TM is suspended state (MSR[TM] = '01').															
41:42	Reserved	Reserved.															
43	FCWAIT	Freeze counters in wait state. 0 The SPMCs are incremented (if permitted by other MMCR bits). 1 The SPMCs are not incremented if CTRL[RUN] = '0'. Software is expected to set CTRL[RUN] = '0' when it is in a wait state; that is, when there is no process ready to run.															
44	FCECE	Freeze counters on an enabled condition or event. 0 The SPMCs are incremented (if permitted by other MMCR bits). 1 The SPMCs are incremented (if permitted by other MMCR bits) until an enabled condition or event occurs at which time: MMCRS[FC] is set to '1'.															
45	PMAE	<p>Performance monitor alert enable.</p> <table> <tr> <td>0</td><td>Performance monitor alerts are disabled.</td></tr> <tr> <td>1</td><td>Performance monitor alerts are enabled until a performance monitor alert occurs, at which time: MMCRO[PMAE] is set to '0' MMCRO[PMAO] is set to '1'</td></tr> </table> <p>Note: Software can set this bit and MMCRO[PMAO] = '0' to prevent performance monitor exceptions. Software can set this bit to '1' and then poll the bit to determine whether an enabled condition or event has occurred. This is especially useful for software that runs with MSR[EE] = '0'.</p>	0	Performance monitor alerts are disabled.	1	Performance monitor alerts are enabled until a performance monitor alert occurs, at which time: MMCRO[PMAE] is set to '0' MMCRO[PMAO] is set to '1'											
0	Performance monitor alerts are disabled.																
1	Performance monitor alerts are enabled until a performance monitor alert occurs, at which time: MMCRO[PMAE] is set to '0' MMCRO[PMAO] is set to '1'																



Table 11-13. Monitor Mode Control Register S Register (Sheet 3 of 3)

Bits	Field Name	Description
46	PMAO	<p>Performance monitor alert has occurred.</p> <p>0 A performance monitor event has not occurred since the last time software set this bit to '0'.</p> <p>1 A performance monitor event has occurred since the last time software set this bit to '0'. This bit is set to '1' by the processor when a performance monitor event occurs. This bit can be set to '0' only by the mtspr instruction.</p> <p>Software can set this bit to '1' to simulate the occurrence of a performance monitor event.</p> <p>Software should set this bit to '0' after handling the performance monitor event.</p>
47	Reserved	Reserved.
48:55	SPMC1_SEL	SPMC1 select. Bit 55 is edge detect. See <i>Section 11.4 Enhanced Sampling Support</i> on page 272.
56:63	SPMC2_SEL	SPMC2 select. Bit 63 is edge detect. See <i>Section 11.4 Enhanced Sampling Support</i> on page 272.

11.9 Sampled Instruction Event Register (SIER)

The Sampled Instruction Event Register is a 64-bit register that stores information that pertains to a sampled or marked instruction. The contents of the SIER can be altered by the hardware if and only if MMCR0[PMAE] = '1'. Thus, after the performance monitor alert occurs, the contents of SIER are not altered by the hardware until software sets MMCR0[PMAE] = '1'. After software sets MMCR0[PMAE] = '1', the contents of SIER are undefined until the next performance monitor alert occurs.

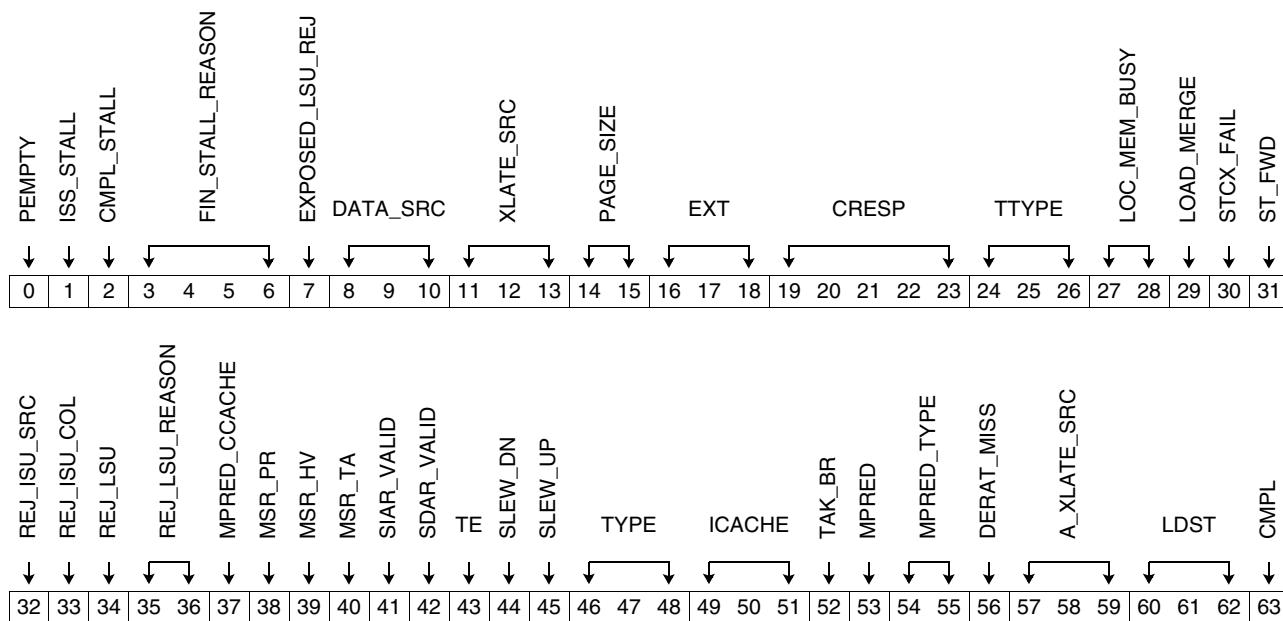


Table 11-14. Sampled Instruction Event Register (SIER) (Sheet 1 of 3)

Bits	Field Name	Description
0	PEMPTY	Pipeline was empty (first instruction after GCT noslot).
1	ISS_STALL	Sampled instruction issue stall. The marked instruction was issued after it became <u>NTC</u> .
2	CMPL_STALL	Sampled instruction group completion stall. The marked instruction became NTC some time before completion.
3:6	FIN_STALL_REASON	Sampled instruction group finish stall reason. 0000 Reserved 0001 Marked finish before NTC 0010 Reserved 0011 Reserved 0100 LSU D-cache miss 0101 LSU load finish 0110 LSU store forward 0111 LSU store 1000 FXU multicycle 1001 BRU finish mispredict 1010 VSU multicycle 1011 Reserved 1100 FXU other 1101 VSU other 1110 BRU other 1111 Reserved
7	EXPOSED_LSU_REJ	Sampled instruction LSU reject exposed to completion.
8:10	DATA_SRC	Sampled instruction microarchitecture-dependent data source information for loads and extended store information. Encoding defined in <i>Table 11-15 Implementation-Dependent Extension to Data Source Encodes</i> on page 297.
11:13	XLATE_SRC	Sampled instruction microarchitecture-dependent data source information to go along with data/translation/source/load information. Encodings are defined in <i>Table 11-15</i> on page 297.
14:15	PAGE_SIZE	Sampled instruction suffered an ERAT miss for the page size. 00 4 KB 01 64 KB 10 16 MB 11 16 GB
16:18	EXT	Sampled instruction data/store/translation extension information. Encodings are defined in <i>Table 11-16 Implementation-Dependent Extension Bits for Data Source Encodes (SIER[EXT])</i> on page 299.
19:23	CRESP	Combined response from the SMP interconnect.
24:26	TTYPE	SMP interconnect event ttype.
27:28	LOC_MEM_BUSY	Local memory busy.
29	LOAD_MERGE	Load merge.
30	STCX_FAIL	Sampled stcx failed.
31	ST_FWD	Store forward.
32	REJ_ISU_SRC	Sampled instruction suffered an ISU reject due to source unavailable.
33	REJ_ISU_COL	Sampled instruction suffered a ISU reject due to resource collision.
34	REJ_LSU	Sampled instruction suffered an LSU reject.



Table 11-14. Sampled Instruction Event Register (SIER) (Sheet 2 of 3)

Bits	Field Name	Description
35:36	REJ_LSU_REASON	Sampled instruction suffered a reject. 00 ERAT miss 01 LMQ full 10 LHS 11 Set MPRED Valid only when SIER[REJ_LSU] = '1'.
37	MPRED_CCACHE	C-cache misprediction.
38	MSR_PR	Sampled problem state (MSR[PR] = '1'). Note: This bit updates regardless of MMCRA[SE] and MMCRA[PMAE].
39	MSR_HV	Sampled hypervisor state (MSR[HV] = '1'). Note: This bit updates regardless of MMCRA[SE] and MMCRA[PMAE].
40	Reserved	Reserved.
41	SIAR_VALID	SIAR is valid for sampled instructions.
42	SDAR_VALID	SDAR is valid for sampled instructions.
43	TE	Threshold exceeded.
44	SLEW_DN	Frequency is slewed down for any reason.
45	SLEW_UP	Frequency is slewed up for any reason.
46:48	TYPE	Type of sampled instruction. 000 Reserved. 001 Sampled instruction is a load. 010 Sampled instruction is a store. 011 Sampled instruction is a branch. 100 Sampled instruction is a floating-point instruction. 101 Sampled instruction is a fixed-point instruction. 110 Sampled instruction is an IFU but nonbranch. 111 Reserved.
49:51	ICACHE	000 Reserved. 001 Sampled instruction hit in the I-cache. 010 Sampled instruction hit in the L2 cache. 011 Sampled instruction hit in the L3 cache. 100 Sampled instruction is an L3 miss. 101 Reserved. 110 Reserved. 111 Reserved. Note: This field is not valid when random event sampling is enabled. MMCRA[63] = '1' and MMCRA[61:62] ≠ '00'.
52	TAK_BR	Sampled instruction is a taken branch.
53	MPRED	Sampled branch instruction is mispredicted.
54:55	MPRED_TYPE	Sampled instruction branch mispredicted information type. 00 Reserved. 01 Branch mispredicted due to direction. 10 Branch mispredicted due to target address. 11 Reserved.
56	DERAT_MISS	Sampled instruction suffered a data ERAT miss.

Table 11-14. Sampled Instruction Event Register (SIER) (Sheet 3 of 3)

Bits	Field Name	Description
57:59	A_XLATE_SRC	<p>Sampled instruction data translation information.</p> <p>000 Reserved.</p> <p>001 TLB hit.</p> <p>010 Data translation hit in the L2 cache.</p> <p>011 Data translation hit in the L3 cache.</p> <p>100 Data translation resolved from memory.</p> <p>101 Data translation resolved from on-chip cache (other than local L2 or L3 cache).</p> <p>110 Data translation resolved from off-chip cache.</p> <p>111 Reserved.</p>
60:62	LDST	<p>Sampled load-store instruction information.</p> <p>000 Reserved.</p> <p>001 Load hit in the L1 D-cache.</p> <p>010 Load hit in the L2 cache.</p> <p>011 Load hit in the L3 cache.</p> <p>100 Load resolved from memory.</p> <p>101 Load resolved from the on-chip cache (other than local L2 or L3 cache).</p> <p>110 Load resolved from the off-chip cache.</p> <p>111 Store missed the L1 cache and was sent to the cache/memory subsystem.</p>
63	CMPL	Sampled instruction completed.

Table 11-15. Implementation-Dependent Extension to Data Source Encodes (Sheet 1 of 3)

Architected Bits SIER[LDST] SIER[A_XLATE_SRC]	Implementation- Dependent Bit Encoding SIER[DATA_SRC] SIER[XLATE_SRC]	Implementation-Dependent Bit Description
L2 Hit		
010	000	Private L2 cache for this core sourced data (or NCU loads) without dispatch conflicts.
010	001	Private L2 cache for this core sourced data in the Mepf state without dispatch conflicts. (L3 prefetch brought line in Me.)
010	010	Private L2 cache for this core sourced data that had a dispatch conflict on a load-hit-store.
010	011	Private L2 cache for this core sourced data that had a dispatch conflict other than a load-hit-store.
010	100	Reserved.
010	101	Reserved.
010	110	Reserved.
010	111	Reserved.
L3 Hit		
011	000	The private L3 cache for this core sourced data without dispatch conflicts.
011	001	The private L3 cache for this core sourced data in the Mepf state without dispatch conflicts. (L3 prefetch brought line in Me.)
011	010	The private L3 cache for this core-sourced data that had a dispatch conflict.
011	011	Reserved.
011	100	Reserved.



Table 11-15. Implementation-Dependent Extension to Data Source Encodes (Sheet 2 of 3)

Architected Bits SIER[LDST] SIER[A_XLATE_SRC]	Implementation- Dependent Bit Encoding SIER[DATA_SRC] SIER[XLATE_SRC]	Implementation-Dependent Bit Description
011	101	Reserved.
011	110	Reserved.
011	111	Reserved.
Memory		
100	000	Reserved.
100	001	Cacheable load on the chip from memory.
100	010	Reserved.
100	011	Cacheable load within group scope.
100	100	Reserved.
100	101	Cacheable load beyond group scope.
100	110	Reserved.
100	111	Reserved.
On-Chip Cache		
101	000	Data sourced from another L2.1 cache (not M but valid) on the same chip.
101	001	Data sourced from another L2.1 cache (in M) on the same chip.
101	010	Data sourced from another L3.1 cache (not M but valid) on the same chip.
101	011	Data sourced from another L3.1 cache (in M) on the same chip.
101	100	Data sourced from another L3.1 cache (not M but valid) on the same chip (ECO).
101	101	Data sourced from another L3.1 cache (in M) on the same chip (ECO).
101	110	Reserved.
101	111	Reserved.
Off-Chip Cache		
110	000	Data sourced from another L2 or L3 cache (not M but valid) within group scope.
110	001	Data sourced from another L2 or L3 cache (in M) within group scope.
110	010	Data sourced from another L2 or L3 cache (not M but valid) beyond group scope.
110	011	Data sourced from another L2 or L3 cache (in M) beyond group scope.
110	100	Reserved.
110	101	Reserved.
110	110	Reserved.
110	111	Reserved.
Store Information (only applicable for SIER[LDST])		
111	000	Store did not require an RC dispatch.
111	001	Store completed in the L2 cache without intervention.
111	010	Store completed in the L2 cache with modified intervention.



Table 11-15. Implementation-Dependent Extension to Data Source Encodes (Sheet 3 of 3)

Architected Bits SIER[LDST] SIER[A_XLATE_SRC]	Implementation- Dependent Bit Encoding SIER[DATA_SRC] SIER[XLATE_SRC]	Implementation-Dependent Bit Description
111	011	Reserved.
111	100	Reserved.
111	101	Reserved.
111	110	Reserved.
111	111	Reserved.

Table 11-16. Implementation-Dependent Extension Bits for Data Source Encodes (SIER[EXT])

Encoding	Case	Description	Prediction
000	Nonfabric operation	Private L2/L3 hit for this core sourced data (or NCU load).	N/A
001	Fabric Initial/Final Pump = Chip	Initial and final pump scope and data sourced across this scope.	Correct
010	Fabric Initial/Final Pump = Group	Initial and final pump scope and data sourced across this scope.	Correct
011	Fabric Initial/Final Pump = System	Initial and final pump scope and data sourced across this scope.	Correct
100	Fabric Final Pump = Group	Final pump scope to get data sourced ended up larger than initial pump scope. (Original scope selected was too small.)	Mispredict
101	Fabric Final Pump = Group	Final pump scope got data from source that was at a smaller scope. (Original scope selected was too large.)	Mispredict
110	Fabric Final Pump = System	Final pump scope to get data sourced ended up larger than initial pump scope. (Original scope selected was too small.)	Mispredict
111	Fabric Final Pump = System	Final pump scope got data from source that was at a smaller scope. (Original scope selected was too large.)	Mispredict



11.10 POWER8 CPI Stack

The POWER8 processor provides functionality for a hardware-based CPI stack that can hierarchically account for completion stalls and front-end stalls on a per-thread basis.

At the most basic level, the CPI stack on a per-thread level can be broken down into three main categories:

- Completion cycles are cycles during which a group of instructions completed for that thread (completion cycles).
- Front-end stall cycles are cycles where the GCT did not have any entries for that thread.
- Completion stall cycles are cycles where the GCT has an entry for the given thread but no completion was seen for the thread. These are cycles where a thread is considered to be completion stalled.

Completion stall cycles can be broken down hierarchically into different reasons why the completion of an instruction was delayed. The mechanism for determining the reason in the POWER8 processor is next-to-finish (NTF) based. A next-to-finish instruction is defined as the oldest instruction in a next-to-complete group that has not finished yet.

The POWER8 processor does precise accounting based on an NTF indication. *Table 11-17* shows the accounting technique.

Table 11-17. POWER8 Accounting

Age of Instruction	Instruction Address	Instruction Mnemonic
1	56CCDC00	lwz R30,20(R3)
2	56CCDC04	twi 0x4,R30,0x0
3	56CCDC08	lwz R3,0(R30)
4	56CCDC0C	cmp CR0,)x0,R3,R4
5	56CCDC10	bc 4,2,+4296
6	56CCDC14	lwz R3,16(R30)

Figure 11-1 POWER8 CPI Stack Example on page 301 shows a group of instructions dispatched. The number to the left indicates the age of the instruction. In this sequence, instructions 2, 3, and 6 are dependent on instruction 1; instruction 4 is dependent on instruction 3; and instruction 5 is dependent on instruction 4. The example assumes that all instructions finish after they become NTC. *Figure 11-1* shows the finish order. The POWER8 processor introduces the NTF-based completion stall. This mechanism ensures that a stall period is broken down into fine grained NTF stall periods. The completion stall is always charged to the oldest instruction in the NTC group that has not finished yet.

Figure 11-1. POWER8 CPI Stack Example

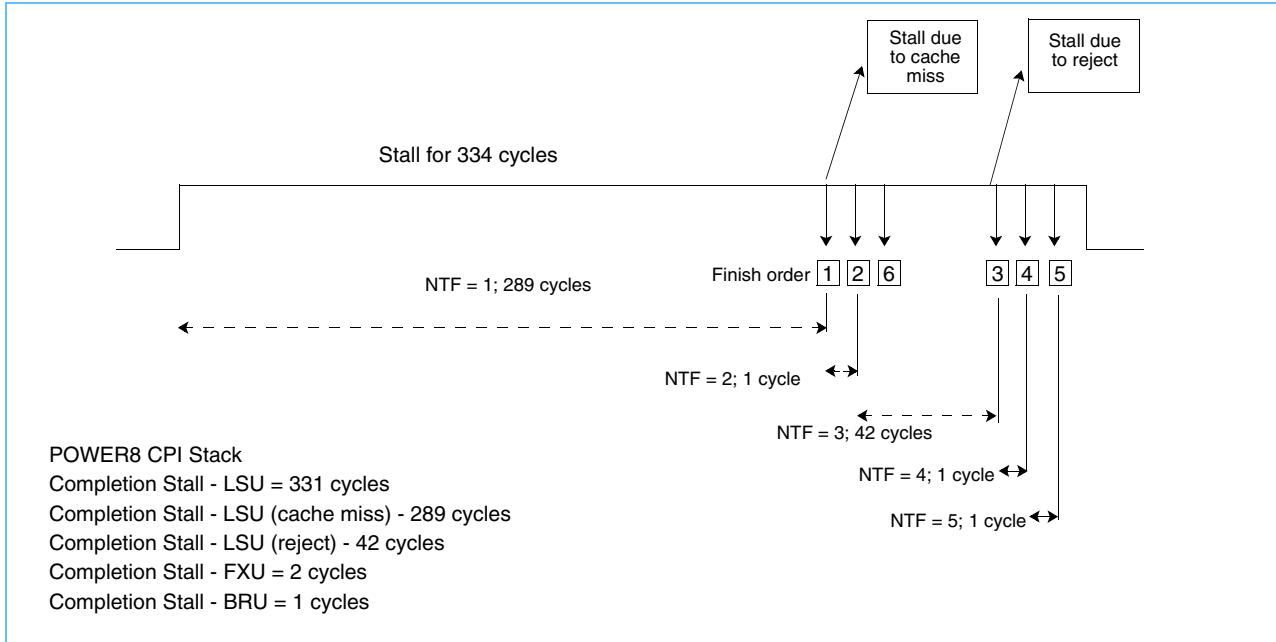


Table 11-18 on page 302 shows a visualization of the CPI stack.



Table 11-18. CPI Stack (Sheet 1 of 2)

Cycles (PM_RUN_CYC)	Stalled Cycles (PM_CMPLU_STALL)	Stall Due to Load/Store (PM_CMPLU_STALL_LSU)	Stall Due to BR or CR (PM_CMPLU_STALL_BRU_CRU)	Stall Due to Branch (PM_CMPLU_STALL_BRU) Stall Due to CR
			Stall Due to Fixed Point (PM_CMPLU_STALL_FXU)	Stall Due to Fixed-Point Long (PM_CMPLU_STALL_FXLONG) Stall Due to Fixed-Point (Other)
			Stall Due to Vector/ Scalar (PM_CMPLU_STALL_VSU)	Stall Due to Vector (PM_CMPLU_STALL_VECTOR) Stall Due to Vector (Other)
				Stall Due to Scalar (PM_CMPLU_STALL_SCALAR) Stall Due to Scalar Long (PM_CMPLU_STALL_SCALAR_LONG) Stall Due to Scalar (Other)
				Stall Due to Vector/Scalar (Other)
			Stall Due to D-Cache Miss (PM_CMPLU_STALL_DCACHE_MISS)	Stall Due to L2/L3 Hit (PM_CMPLU_STALL_DMISS_L2L3) L2/L3 Hit with No Conflict
				Stall Due to On-Chip L2/L3 (PM_CMPLU_STALL_DMISS_L21_L31)
				Stall Due to On-Chip Memory (PM_CMPLU_STALL_DMISS_LMEM)
				Stall Due to Off-Chip Memory (PM_CMPLU_STALL_DMISS_REMOTE)
				Stall Due to Off-Node Memory/Cache
				Reject Due to Load-Hit-Store (PM_CMPLU_STALL_REJECT_LHS) Reject Due to ERAT Miss (PM_CMPLU_STALL_ERAT_MISS) Reject Due to LMQ Full (PM_CMPLU_STALL_REJ_LMQ_FULL) Reject Due to Reject (Other)
			Stall Due to LSU Reject (PM_CMPLU_STALL_REJECT)	Stall Due to Store Finish (PM_CMPLU_STALL_STORE) Stall Due to Load Finish (PM_CMPLU_STALL_LOAD_FINISH) Stall Due to Store Forward (PM_CMPLU_STALL_ST_FWD) Stall Due to Load/Store (Other)
				Stall Due to Next-To-Complete Flush (PM_CMPLU_STALL_NTCG_FLUSH)
				Stall Due to NOPs (PM_CMPLU_STALL_NO_NTF)
				Stall Cycles (Other)
				Finished Group Waiting to Complete (PM_NTCG_ALL_FIN)
Thread Blocked (PM_CMPLU_STALL_THRD)			Blocked Due to LWSYNC (PM_CMPLU_STALL_LWSYNC)	Blocked Due to LWSYNC (PM_CMPLU_STALL_LWSYNC)
				Blocked Due to HWSYNC (PM_CMPLU_STALL_HWSYNC)
			Blocked Due to ECC Delay (PM_CMPLU_STALL_MEM_ECC_DELAY)	Blocked Due to ECC Delay (PM_CMPLU_STALL_MEM_ECC_DELAY)
				Blocked Due to Other Thread's Flush (PM_CMPLU_STALL_FLUSH)
			Blocked Due to COQ Full (PM_CMPLU_STALL_COQ_FULL)	Blocked Due to COQ Full (PM_CMPLU_STALL_COQ_FULL)
				Thread Blocked (Other)
Note: Shaded boxes indicate that these are not hardware events but are derived from existing events.				

Table 11-18. CPI Stack (Sheet 2 of 2)

Cycles (continued)	Nothing to Dispatch (PM_GCT_NOSLOT_CYC)	Nothing to Dispatch Due to I-Cache Miss (PM_GCT_NOSLOT_IC_MISS)	Nothing to Dispatch Due to I-Cache L3 Miss (PM_GCT_NOSLOT_IC_L3MISS)
			Nothing to Dispatch Due to I-Cache Miss (Other)
			Nothing to Dispatch Due to Branch Mispredict (PM_GCT_NOSLOT_BR_MPRED)
			Nothing to Dispatch Due to Branch Mispredict and I-Cache Miss (PM_GCT_NOSLOT_BR_MPRED_ICMISS)
		Nothing to Dispatch - Dispatch Held	Dispatch Held Due to Mapper (PM_GCT_NOSLOT_DISP_HELD_MAP)
			Dispatch Held Due to Store Queue (PM_GCT_NOSLOT_DISP_HELD_SRQ)
			Dispatch Held Due to Issue Queue (PM_GCT_NOSLOT_DISP_HELD_ISSQ)
			Dispatch Held (Other) (PM_GCT_NOSLOT_DISP_HELD_OTHER)
			Nothing to Dispatch (Other)
			Other
			Completion Cycles (PM_RUN_INST_CMPL)

Note: Shaded boxes indicate that these are not hardware events but are derived from existing events.

11.10.1 Completion Stall Accounting: LSU Related Stalls

This category of stalls indicates that an instruction that finished in the load-store unit was responsible for a portion of the NTF stall.

LSU stalls can be hierarchically broken down into stalls due to cache misses, rejects, stores, store forwarding, and **stcx** fails.

11.10.1.1 Completion Stall Accounting: Data Cache Misses

This category of stalls indicates that an instruction that finished in the load unit and suffered a cache miss was responsible for a portion of the NTF stall.

11.10.1.2 Completion Stalls: Data Cache Miss that Resolves in a Local Core's L2 or L3 Cache

This category of stalls indicates that an instruction that finished in the load unit and suffered a cache miss that resolved in the local L2 or L3 cache was responsible for a portion of the NTF stall.

11.10.1.3 Completion Stalls: Data Cache Miss that Resolves in a Local Chip's L2 or L3 Cache

This category of stalls indicates that an instruction that finished in the load unit and suffered a cache miss that resolved in the local chip's L2 or L3 cache was responsible for a portion of the NTF stall.

11.10.1.4 Completion Stalls: Data Cache Miss that Resolves from Remote Chip's Cache or Memory

This category of stalls indicates that an instruction that finished in the load unit and suffered a cache miss that resolved in a remote chip's cache or memory was responsible for a portion of the NTF stall.



11.10.1.5 Completion Stalls: Data Cache Miss that Resolves from Local Core's L2 or L3 (Dispatch Conflict)

This category of stalls indicates that an instruction that finished in the load unit and suffered a cache miss that resolved on a core's local L2 or L3 cache but with a conflict was responsible for a portion of the NTF stall.

11.10.1.6 Completion Stalls: Data Cache Miss that Resolves from Local Memory

This category of stalls indicates that an instruction that finished in the load unit and suffered a cache miss that resolved in the core's local memory was responsible for a portion of the NTF stall.

11.10.1.7 Completion Stalls: Stores

This category of stalls indicates that a store instruction that finished in the load-store unit was responsible for a portion of the NTF stall.

11.10.1.8 Completion Stalls: Store Forwarding

This category of stalls indicates that a store instruction forwarded data to a load that was responsible for a portion of the NTF stall.

11.10.1.9 Completion Stalls: LSU Rejects

This category of stalls indicates that completion was stalled due to an LSU reject.

11.10.1.10 Completion Stalls: LSU Rejects Due to ERAT Miss

This category of stalls indicates that completion was stalled due to an LSU reject due to an ERAT miss.

11.10.1.11 Completion Stalls: LSU Rejects Due to LMQ Full

This category of stalls indicates that completion was stalled due to an LSU reject because the LMQ is full.

11.10.1.12 Completion Stalls: LSU Rejects Due to Load-Hit-Store Reject

This category of stalls indicates that completion was stalled due to an LSU reject because the LMQ is full.

11.10.2 Completion Stalls: FXU

This category of stalls covers stalls due to a fixed-point instruction. The PMU starts counting stall cycles and if the NTF instruction is an FXU instruction, the stall cycles are charged to this category.

11.10.3 Completion Stalls: VSU

This category of stalls covers stalls due to a VSU instruction. The PMU starts counting stall cycles and if the NTF instruction is a VSU instruction, the stall cycles are charged to this category.

11.10.4 Completion Stalls: IFU

This category of stalls covers stalls due to instructions that finish in the IFU. The PMU starts counting stall cycles and if the NTF instruction finishes in the IFU, the stall cycles are charged to this category.

This category is further subdivided into branch and CR instructions.

11.10.5 Front-End Stalls

The CPI stack also accounts for cycles when the GCT is completely empty for the given thread. This is an indication of thread starvation, and the hardware attempts to account for reasons for the GCT to be empty for that thread.

11.10.5.1 GCT Empty: I-Cache Miss

This event counts the cycles when the GCT is empty for that thread because of an instruction-cache miss.

Every group dispatched has an indication if the group suffered a branch redirect or an I-cache miss or both. The PMU starts counting cycles where the GCT is empty for the given thread. If a group is dispatched, causing the GCT to be nonempty, the PMU stops counting and accounts for the GCT-empty cycles and charges the front-end stall cycles to an I-cache miss if the group dispatched has an I-cache miss indication.

11.10.5.2 GCT Empty: I-Cache Miss That Also Missed the Local L3 Cache

This event counts the cycles where the GCT was empty for that thread because of an instruction-cache miss that resolved in the local L2 cache.

Every group dispatched has an indication if the group suffered a branch redirect or an I-cache miss or both. Every group dispatched also has an indication of the I-cache miss resolved from the local L2 or local L3 cache or was an L3 miss. The PMU starts counting cycles when the GCT is empty for the given thread. If a group was dispatched causing the GCT to be nonempty, the PMU stops counting and accounts for the GCT empty cycles and charges the front-end stall cycles to an I-cache miss if the group dispatched has an I-cache miss indication.

11.10.5.3 GCT Empty: Branch Redirects

This event counts the cycles where the GCT was empty for that thread because of a branch redirect. Every group dispatched has an indication if the group suffered a branch redirect or an I-cache miss or both. The PMU starts counting cycles where the GCT is empty for the given thread. If a group was dispatched causing the GCT to be nonempty, the PMU stops counting and accounts for the GCT empty cycles and charges the front-end stall cycles to branch redirects if the group has a branch redirect indication.

11.10.5.4 GCT Empty: Branch Redirects and I-Cache Miss

This event counts the cycles where the GCT was empty for that thread because of a branch redirect that also suffered an I-cache miss. Every group dispatched has an indication if the group suffered a branch redirect or an I-cache miss or both. The PMU starts counting cycles where the GCT is empty for the given thread. If a group was dispatched causing the GCT to be nonempty, the PMU stops counting and accounts for the GCT empty cycles and charges the front-end stall cycles to branch redirects and an I-cache miss if the group has a branch redirect and an I-cache miss indication.



11.10.5.5 GCT Empty: Dispatch Hold Conditions

This event counts the cycles where the GCT was empty for that thread because of dispatch hold conditions that exist in the pipeline.

The PMU starts counting cycles where the GCT is empty for the given thread and stops counting when a group is dispatched and it indicates a previous dispatch hold condition for that thread.

11.11 Exploiting Advanced Features of the PMU

11.11.1 Correlating Fabric Responses to Effective Addresses

This feature identifies the code and data addresses of locks and atomic updates, transactional memory, and other contention across threads that are causing traffic on the inter-processor connection fabric. In particular, this feature accurately identifies which locking code sequences and which data structures are involved in the locks that are causing problems. This feature also distinguishes between problems related to locking and other problems that might cause poor scaling but are in fact unrelated to the locking.

Identifying the causes of lock contention in a precise and accurate manner that is easily accessible to programmers through standard tools helps them produce software that has better scaling characteristics. This is important because scaling up to a large number of cores and threads is important for cloud-computing platforms.

The basic idea is to sample or mark loads or stores and record their effective instruction and data addresses in the core in special-purpose registers (SIAR/SDAR). While the load/store is pending in the nest (L2 or L3 miss), the L2 cache forwards fabric responses for that tagged load/store to the performance monitoring unit located in the core.

11.11.1.1 Operation

The PMU records effective addresses and fabric responses in the SIER. The fabric responses can be post-processed to find contention issues. The hardware can also be programmed to find a specific response from the fabric using the thread-level PMCs, or count responses and raise an interrupt when the threshold is reached.

- The POWER8 processor provides thread-level sampling to randomly mark loads/stores using random instruction sampling or random event sampling. The effective address of the instruction is logged in the SIAR and the data effective address is logged in the SDAR.
- If a load or store is marked in the thread-level PMU, the sampling state machine monitors the progress of the marked instruction at the time the request for the line is sent to the L2 cache.
- Only one instruction can be marked at a time. The hardware waits for the current marked instruction to either flush or complete before marking the next instruction.
- The PMU captures information about the address, combined response, and type in the SIER based on the MMCRA settings. When MMCRA[SIER_CTRL] = '1', the SIER[16:28] bits are changed to data real address bits [44:56].

By default, the SIER provides cresp and type information. However, this can be changed to provide real address bits [44:56] that map to congruence class bits in the L2 or L3 cache.

Because the PMU captures the instruction and data effective address in the SIAR and SDAR registers, combined responses, ttypes, and effective address can be correlated to determine where and what type of contention issues exist.

11.11.1.2 Tools Exploitation

The PMU provides the following performance events that can be configured and used for profiling. For example, the user can profile directly for stores that do a background kill (bkill) to gain ownership of the line, which is a sign that shared copies were given away.

PM_MRK_FAB_RSP_T_INTV: T intervention
PM_MRK_FAB_RSP_RWITM_RTY: RWITM was retried
PM_MRK_FAB_RSP_DCLAIM: store did a dclaim
PM_MRK_FAB_RSP_CLAIM_RTY: dclaim was retried
PM_MRK_FAB_RSP_BKIL: store did a bkill
PM_MRK_FAB_RSP_RD_RTY: loads were retried

The PMU also provides a general purpose match event: PM_MRK_FAB_RSP_MATCH. This event counts the number of occurrences of the combination of ttype and crespc specified in MMCR1[20:27]. See *Table 11-3 MMCR1 Register* on page 280. This enables the tools to create any combination of ttype and crespc and events on demand.

In addition to the performance events, the PMU also provides the ability to post-process the SIER. For example, the PM_MRK_ST_CMPL can be used as a profiling event and the SIER can then be processed to see the various cresps that the stores are receiving.

Software post-processing tools can look for patterns (see the list that follows) to identify possible bottlenecks. Workloads can then be tuned to avoid or reduce these bottlenecks.

- Loads that get a T intervention. Because T interventions happen after a combined response is done, the latency is much longer. If more than a small fraction of interventions are sourced by T states, the added latency can cause a performance problem.

If a go_SI:T response is seen, either the data was not present in the local node (indicating poor locality), or the SI copy in the local node aged out of the cache and was discarded (indicating cache management or cache address conflict issues). The table in the architecture specifies go_SI:T to have an encoding of '01000'; therefore, tools can look for this while processing the SIER.

If a go_S:T ('01100') is seen, the cache that contained the SI copy in the local node was not able to respond (indicating snoop machine utilization issues).

If a go_SI:ipc ('01010') response to a rd_go_s command is seen, this indicates that there are S copies of the data present in the node, but the SI copy has been lost suggesting that there are cache management or sharing pattern issues.

- Loads that get lots of retries. Each retry adds hundreds of cycles of latency to the load request, which impacts performance.

rtv_ned_np: indicates conflicts are on-node.

rtv_ned_sp: indicates that data was not found on-node or that there are conflicts off-node.



- Stores that had to do a bkill. Similar to load retries, these make the stores take longer. If the store is followed by a **sync**, as is usually the case in a locking sequence, the **sync** cannot complete until the previous stores are complete. A store that hits a Tx state in the local cache ends up doing a bkill on the fabric.
addr_ack_done resp: indicates that the bkill finished successfully.
addr_ack_bk_np: indicates that the bkill had to be resent to the local node. Other shared copies are created while a bkill is pending, that indicate high levels of contention for the line.
addr_ack_bk_sp: indicates that the bkill had to be resent to the whole system.
- Stores that had to do a dclaim.
A store that hits an S/SI state in the L2 cache results in a dclaim being issued.
go_M_bk_np: indicates that the dclaim was successful, but a background kill must be sent to the node afterward.
go_M_bk_sp: indicates that the dclaim was successful, but a background kill must be sent to the whole system.
rty_np: indicates that the dclaim must be resent to the node. When this response is seen, either another cache is also trying to gain ownership, or the memory controller queues are full and the dclaim cannot be accepted there.
rty_sp: indicates that the dclaim must be resent to the whole system.
rty_lost_claim_np and rty_lost_claim_sp: indicate that another thread has gained ownership of the line for a store and this thread must invalidate its copy of the line and start over (as if it had missed the cache) by sending an RWITM to the node or the whole system respectively. This indicates that many threads are storing to the line at the same time and that there is a high level of contention, especially if there were one or more rty_np/rty_sp seen before the final lost claim.
- A **lарx** that sees an SI.
An SI intervention response to a **lарx** request indicates that there are shared copies of the line in the system, which is undesirable for locks. In the case of an atomic update, this can be due to incorrect hint bits or a protection window not being sufficient.



11.11.2 Finding Wasted L3 Prefetches

The POWER8 processor has a cache state called Mepf that determines the effectiveness of L3 prefetches. Rules for the Mepf state are as follows:

1. The state is formed by the L3 cache when an L3 prefetch machine brings a line in the Me state from the SMP interconnect; that is, the line comes in from another cache (L2 or L3) or memory.
2. An Mepf-to-Me transition occurs when a demand load or an L1 prefetch hits a line in the Mepf state in the L3 cache.

11.11.2.1 PMU Usage

The PMU has the following Mepf events:

- **PM_DATA_FROM_L3_MEFP**

This event counts when a demand load merges with an L1 prefetch or a demand load hits on the line in an Mepf state. There is also a version of the same event with no Mepf (PM_DATA_FROM_L3) that includes lines in the Mepf state and all the other possible states.

- **PM_MRK_DATA_FROM_L3_MEFP_CYC**

This is the same as the previous event but can be used to count the cycles that a demand load hits in the Mepf state. This event tells how many cycles the prefetch is saving.

- **PM_MEM_PREF**

This event counts the number of prefetches that came into the chiplet from memory for a particular LPAR.

- **PM_MEM_CO_MEFP**

This event counts the number of lines that were evicted from the L3 cache in the Mepf state for a particular LPAR. This event indicates that the line was never consumed by the core. Assuming all prefetches come from memory, the difference of PM_MEM_PREF - PM_MEM_CO_MEFP indicates how many L3 prefetches are used. It is also possible that prefetches came from another cache.

The L3 event bus on the default setting on Bus0 supports the following events. These events are per physical core:

- **PM_L3_PF_MISS_L3**

The L3 prefetch missed in the L3 cache and the L3 prefetch machines will master a request for an L3 prefetch on the SMP interconnect.

- **PM_L3_CO_MEFP**

An L3 replacement of a line in the Mepf state. This event along with PM_L3_PF_MISS_L3 provides a per-core average of lines brought in by L3 prefetches but never consumed by an L1 prefetch or demand load.

The following events can be used to find an affinity of prefetches:

- **PM_L3_PF_ON_CHIP_CACHE**

L3 prefetch from an on-chip cache.

- **PM_L3_PF_OFF_CHIP_CACHE**

L3 prefetch from an off-chip cache.

- **PM_L3_PF_ON_CHIP_MEM**

L3 prefetch from on-chip memory.

- **PM_L3_PF_OFF_CHIP_MEM**

L3 prefetch from off-chip memory.



11.11.3 Per LPAR Memory Bandwidth

The PMU provides the following performance events to characterize and break down memory bandwidth at an LPAR level. This can be compared to memory bandwidth at a socket or chip level to find memory bandwidth bottlenecks and identify the LPARs that are consuming memory bandwidth.

The PMU provides the following events.

- **PM_MEM_READ:** Counts the number of L2 mastered reads that resolved from memory. This includes:
 - Demand L1 misses
 - L1 prefetches
 - Instruction fetches
 - Instruction prefetches
 - Speculative load misses (that is, loads that missed the L1 cache but cancelled after the L2 cache started doing a read)
 - Data translation requests
 - Instruction translation requests
- **PM_MEM_PREF:** Counts the number of L3 prefetches that were sourced from memory. This also includes store spawned prefetches.
- **PM_MEM_RWITM:** Counts store misses from the L2 cache that were sourced from memory.
- **PM_MEM_CO:** Count includes dirty castouts; both L2 stores that hit in the local L3 cache and L2 stores that hit in the local L2 cache.

11.11.4 Monitoring Fabric Command Scope at a Thread Level

Pumps mastered by the internal fabric (SMP) interconnect can be monitored at the chip level and monitor pumps mastered by the L2 cache at the chiplet level.

The scope of the pumps is a primary indication of affinity issues. For example, if affinity is good, the pump scope is limited to the narrowest scope (chip pump). Many system pumps typically means bad scaling.

It is also useful to monitor retries along with the scope of SMP interconnect commands. Many retries add latency to requests and increase snoop/bus utilization.

The pumps the SMP interconnect sees can be either from the L2 or the L3 cache. The L2 cache masters commands on the SMP interconnect for demand misses, L1 prefetches, store misses, instruction fetches, instruction prefetches, translation misses, bkills, dclaims, and translation prefetches.

The L2 instrumentation provides the following events at the core and chiplet level:

- **PM_L2_CHIP_PUMP**
- **PM_L2_GROUP_PUMP**
- **PM_L2_SYS_PUMP**

The L3 cache masters commands on the SMP interconnect for L3 prefetches, lateral castouts, and castouts to memory. The L3 instrumentation provides the following events at the core and chiplet level:

- **PM_L3_P0_NODE_PUMP**
- **PM_L3_P1_NODE_PUMP**
- **PM_L3_P0_GRP_PUMP**



- PM_L3_P1_GRP_PUMP
- PM_L3_P0_SYS_PUMP
- PM_L3_P1_SYS_PUMP

The following events are also available to the thread-level PMCs:

- PM_GRP_PUMP_CPRED
- PM_GRP_PUMP_MPRED
- PM_GRP_PUMP_MPRED_RTY
- PM_SYS_PUMP_CPRED
- PM_SYS_PUMP_MPRED
- PM_SYS_PUMP_MPRED_RTY
- PM_PUMP_CPRED
- PM_PUMP_MPRED

These events can also be broken down in the PMU for data, instruction, and translation (instruction and data).

When profiling, the PMU generates marked events for some of these events and also records the extension field in the SIER[EXT] bits. Profiling tools can then post-process this information enabling performance tools to characterize SMP interconnect pump activity at the thread level. This information is useful in determining where affinity issues arise.

11.12 PMC Events

The PMC events are listed in *Appendix D Performance Monitoring Events* on page 387.

11.13 SPMC Events

The SPMC events are listed in *Appendix E SPMC Performance Monitoring Events* on page 435.





Appendix A. POWER8 Instruction Summary by Category

Table A-1. defines the POWER8 categories listed in Table A-2. on page 314.

Table A-1. Category Listing

Abbreviation	Category	Notes
64	64-Bit	Required for 64-bit implementations; not defined for 32-bit implementations.
B	Base	Required for all implementations.
DFP	Decimal Floating-Point	Decimal floating-point facilities.
FP	Floating-Point	Floating-point facilities.
FP[R]	Floating-Point Record	Floating-point with Rc = 1.
LSQ	Load/Store Quadword	Load/Store quadword instructions. See Power ISA, Book III-S.
MA	Move Assist	Move assist instructions. ³
S	Server	Required for server implementations.
TM	Transactional Memory	Full hardware transactional memory support.
V	Vector	Vector facilities.
V.AES	Vector	Advanced encryption standard assist instruction.
V.RAID	Vector	Vector permute-XOR instruction.
V.SHA2	Vector	Secure hash algorithm-2 assist instructions.
VSX	Vector-Scalar Extension	Vector-scalar extension. Requires implementation of floating-point and vector categories.

1. .in - facilities and instructions that, in the next version of the architecture, will be required as part of the category they are dependent on. For example, FP[R].in.
2. .out - facilities and instructions that, in some future version of the architecture, will be dropped out of the architecture. For example, FP.out.
3. Move assist instructions are not valid in little-endian mode.



Table A-2. lists the instructions implemented in the POWER8 processor in order by category.

Table A-2. *POWER8 Instructions by Category* (Sheet 1 of 24)

Instruction Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
X	31	7C0001F8	64	bpermd	Bit Permute Doubleword
X	31	7C000074	64	cntlzd[.]	Count Leading Zeros Doubleword
XO	31	7C0003D2	64	divd[.]	Divide Doubleword
XO	31	7C000352	64	divde[.]	Divide Doubleword Extended
XO	31	7C000752	64	divdeo[.]	Divide Doubleword Extended and Record OV
XO	31	7C000312	64	divdeu[.]	Divide Doubleword Extended Unsigned
XO	31	7C000712	64	divdeuo[.]	Divide Doubleword Extended Unsigned and Record OV
XO	31	7C0007D2	64	divdo[.]	Divide Doubleword and Record OV
XO	31	7C000392	64	divdu[.]	Divide Doubleword Unsigned
XO	31	7C000792	64	divduo[.]	Divide Doubleword Unsigned and Record OV
X	31	7C0007B4	64	extsw[.]	Extend Sign Word
DS	58	E8000000	64	ld	Load Doubleword
X	31	7C0000A8	64	ldarx	Load Doubleword And Reserve Indexed
X	31	7C000428	64	ldbrix	Load Doubleword Byte-Reverse Indexed
X	31	7C00003A	64	ldepx	Load Doubleword by External PID Indexed
DS	58	E8000001	64	ldu	Load Doubleword with Update
X	31	7C00006A	64	ldux	Load Doubleword with Update Indexed
X	31	7C00002A	64	idx	Load Doubleword Indexed
DS	58	E8000002	64	lwa	Load Word Algebraic
X	31	7C0002EA	64	lwaux	Load Word Algebraic with Update Indexed
X	31	7C0002AA	64	lwax	Load Word Algebraic Indexed
XO	31	7C000092	64	mulhd[.]	Multiply High Doubleword
XO	31	7C000012	64	mulhdu[.]	Multiply High Doubleword Unsigned
XO	31	7C0001D2	64	mulld[.]	Multiply Low Doubleword
XO	31	7C0005D2	64	mulldo[.]	Multiply Low Doubleword and Record OV
X	31	7C0003F4	64	popcntd	Population Count Doubleword
X	31	7C000174	64	ptryd	Parity Doubleword
MDS	30	78000010	64	rldcl[.]	Rotate Left Doubleword then Clear Left
MDS	30	78000012	64	rldcr[.]	Rotate Left Doubleword then Clear Right
MD	30	78000008	64	rldic[.]	Rotate Left Doubleword Immediate then Clear
MD	30	78000000	64	rldic[.]	Rotate Left Doubleword Immediate then Clear Left
MD	30	78000004	64	rldicr[.]	Rotate Left Doubleword Immediate then Clear Right
MD	30	7800000C	64	rldimi[.]	Rotate Left Doubleword Immediate then Mask Insert



Table A-2. POWER8 Instructions by Category (Sheet 2 of 24)

Instruction Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
X	31	7C000036	64	sld[.]	Shift Left Doubleword
X	31	7C000634	64	srad[.]	Shift Right Algebraic Doubleword
XS	31	7C000674	64	sradi[.]	Shift Right Algebraic Doubleword Immediate
X	31	7C000436	64	srd[.]	Shift Right Doubleword
DS	62	F8000000	64	std	Store Doubleword
X	31	7C000528	64	stdbrx	Store Doubleword Byte-Reverse Indexed
X	31	7C0001AD	64	stdcx.	Store Doubleword Conditional Indexed and Record CR0
X	31	7C00013A	64	stdeps	Store Doubleword by External PID Indexed
DS	62	F8000001	64	stdu	Store Doubleword with Update
X	31	7C00016A	64	stdux	Store Doubleword with Update Indexed
X	31	7C00012A	64	stdx	Store Doubleword Indexed
X	31	7C000088	64	td	Trap Doubleword
D	2	08000000	64	tdi	Trap Doubleword Immediate
XO	31	7C000214	B	add[.]	Add
XO	31	7C000014	B	addc[.]	Add Carrying
XO	31	7C000414	B	addco[.]	Add Carrying and Record OV
XO	31	7C000114	B	adde[.]	Add Extended
XO	31	7C000514	B	addeo[.]	Add Extended and Record OV and Record OV
D	14	38000000	B	addi	Add Immediate
D	12	30000000	B	addic	Add Immediate Carrying
D	13	34000000	B	addic.	Add Immediate Carrying and Record CR0
D	15	3C000000	B	addis	Add Immediate Shifted
XO	31	7C0001D4	B	addme[.]	Add to Minus One Extended
XO	31	7C0005D4	B	addmeo[.]	Add to Minus One Extended and Record OV
XO	31	7C000614	B	addo[.]	Add and Record OV
XO	31	7C000194	B	addze[.]	Add to Zero Extended
XO	31	7C000594	B	addzeo[.]	Add to Zero Extended and Record OV
X	31	7C000038	B	and[.]	AND
X	31	7C000078	B	andc[.]	AND with Complement
D	28	70000000	B	andi.	AND Immediate and Record CR0
D	29	74000000	B	andis.	AND Immediate Shifted and Record CR0
I	18	48000000	B	b[!][a]	Branch
B	16	40000000	B	bc[l][a]	Branch Conditional
XL	19	4C000420	B	bcctr[l]	Branch Conditional to Count Register



Table A-2. POWER8 Instructions by Category (Sheet 3 of 24)

Instruction Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
XL	19	4C000020	B	bclr[l]	Branch Conditional to Link Register
X	19	4C000460	B	bctar[l]	Branch Conditional to Branch Target Address Register
X	31	7C000000	B	cmp	Compare
X	31	7C0003F8	B	cmpb	Compare Byte
D	11	2C000000	B	cmpi	Compare Immediate
X	31	7C000040	B	cmpl	Compare Logical
D	10	28000000	B	cmpli	Compare Logical Immediate
X	31	7C000034	B	cntlzw[.]	Count Leading Zeros Word
XL	19	4C000202	B	crand	Condition Register AND
XL	19	4C000102	B	crandc	Condition Register AND with Complement
XL	19	4C000242	B	creqv	Condition Register Equivalent
XL	19	4C0001C2	B	crnand	Condition Register NAND
XL	19	4C000042	B	crnor	Condition Register NOR
XL	19	4C000382	B	cror	Condition Register OR
XL	19	4C000342	B	crorc	Condition Register OR with Complement
XL	19	4C000182	B	crxor	Condition Register XOR
X	31	7C0000AC	B	dcbf	Data Cache Block Flush
X	31	7C00006C	B	dcbst	Data Cache Block Store
X	31	7C00022C	B	dcbt	Data Cache Block Touch
X	31	7C0001EC	B	dcbtst	Data Cache Block Touch for Store
X	31	7C0007EC	B	dcbz	Data Cache Block Zero
XO	31	7C0003D6	B	divw[.]	Divide Word
XO	31	7C000356	B	divwe[.]	Divide Word Extended
XO	31	7C000756	B	divweo[.]	Divide Word Extended and Record OV
XO	31	7C000316	B	divweu[.]	Divide Word Extended Unsigned
XO	31	7C000716	B	divweuo[.]	Divide Word Extended Unsigned and Record OV
XO	31	7C0007D6	B	divwo[.]	Divide Word and Record OV
XO	31	7C000396	B	divwu[.]	Divide Word Unsigned
XO	31	7C000796	B	divwuo[.]	Divide Word Unsigned and Record OV
X	31	7C000238	B	eqv[.]	Equivalent
X	31	7C000774	B	extsb[.]	Extend Sign Byte
X	31	7C000734	B	extsh[.]	Extend Sign Halfword
X	31	7C0007AC	B	icbi	Instruction Cache Block Invalidate
X	31	7C00002C	B	icbt	Instruction Cache Block Touch



Table A-2. POWER8 Instructions by Category (Sheet 4 of 24)

Instruction Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
A	31	7C00001E	B	isel	Integer Select
XL	19	4C00012C	B	isync	Instruction Synchronize
X	31	7C000068	B	lbarx	Load Byte And Reserve Indexed
D	34	88000000	B	lbz	Load Byte and Zero
D	35	8C000000	B	lbzu	Load Byte and Zero with Update
X	31	7C0000EE	B	lbzux	Load Byte and Zero with Update Indexed
X	31	7C0000AE	B	lbzx	Load Byte and Zero Indexed
D	42	A8000000	B	lha	Load Halfword Algebraic
X	31	7C0000E8	B	lhax	Load Halfword And Reserve Indexed Xform
D	43	AC000000	B	lhau	Load Halfword Algebraic with Update
X	31	7C0002EE	B	lhaux	Load Halfword Algebraic with Update Indexed
X	31	7C0002AE	B	lhax	Load Halfword Algebraic Indexed
X	31	7C00062C	B	lhbrx	Load Halfword Byte-Reverse Indexed
D	40	A0000000	B	lhz	Load Halfword and Zero
D	41	A4000000	B	lhzu	Load Halfword and Zero with Update
X	31	7C00026E	B	lhzux	Load Halfword and Zero with Update Indexed
X	31	7C00022E	B	lhzx	Load Halfword and Zero Indexed
D	46	B8000000	B	lmw	Load Multiple Word
X	31	7C000028	B	lwax	Load Word and Reserve Indexed
X	31	7C00042C	B	lwbrx	Load Word Byte-Reverse Indexed
D	32	80000000	B	lwz	Load Word and Zero
D	33	84000000	B	lwzu	Load Word and Zero with Update
X	31	7C00006E	B	lwzux	Load Word and Zero with Update Indexed
X	31	7C00002E	B	lwzx	Load Word and Zero Indexed
XL	19	4C000000	B	mcrf	Move Condition Register Field
XFX	31	7C000026	B	mfcr	Move From Condition Register
XFX	31	7C100026	B	mfocrf	Move From One Condition Register Field
XFX	31	7C0002A6	B	mfsp	Move From Special Purpose Register
XFX	31	7C000120	B	mtcrcf	Move To Condition Register Fields
XFX	31	7C100120	B	mtocrf	Move To One Condition Register Field
XFX	31	7C0003A6	B	mtspr	Move To Special Purpose Register
XO	31	7C000096	B	mulhw[.]	Multiply High Word
XO	31	7C000016	B	mulhwu[.]	Multiply High Word Unsigned
D	7	1C000000	B	mulli	Multiply Low Immediate



Table A-2. POWER8 Instructions by Category (Sheet 5 of 24)

Instruction Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
XO	31	7C0001D6	B	mullw[.]	Multiply Low Word
XO	31	7C0005D6	B	mullwo[.]	Multiply Low Word and Record OV
X	31	7C0003B8	B	nand[.]	NAND
XO	31	7C0000D0	B	neg[.]	Negate
XO	31	7C0004D0	B	nego[.]	Negate and Record OV
X	31	7C0000F8	B	nor[.]	NOR
X	31	7C000378	B	or[.]	OR
X	31	7C000338	B	orc[.]	OR with Complement
D	24	60000000	B	ori	OR Immediate
D	25	64000000	B	oris	OR Immediate Shifted
X	31	7C0000F4	B	popcntb	Population Count Byte-wise
X	31	7C0002F4	B	popcntw	Population Count Words
X	31	7C000134	B	prtyw	Parity Word
M	20	50000000	B	rlwimi[.]	Rotate Left Word Immediate then Mask Insert
M	21	54000000	B	rlwinm[.]	Rotate Left Word Immediate then AND with Mask
M	23	5C000000	B	rlwnm[.]	Rotate Left Word then AND with Mask
SC	17	44000002	B	sc	System Call
X	31	7C000030	B	slw[.]	Shift Left Word
X	31	7C000630	B	sraw[.]	Shift Right Algebraic Word
X	31	7C000670	B	srawi[.]	Shift Right Algebraic Word Immediate
X	31	7C000430	B	srw[.]	Shift Right Word
D	38	98000000	B	stb	Store Byte
X	31	7C00056D	B	stbcx.	Store Byte Conditional Indexed
D	39	9C000000	B	stbu	Store Byte with Update
X	31	7C0001EE	B	stbux	Store Byte with Update Indexed
X	31	7C0001AE	B	stbx	Store Byte Indexed
D	44	B0000000	B	sth	Store Halfword
X	31	7C00072C	B	sthbrx	Store Halfword Byte-Reverse Indexed
X	31	7C0005AD	B	sthcx.	Store Halfword Conditional Indexed Xform
D	45	B4000000	B	sthuh	Store Halfword with Update
X	31	7C00036E	B	sthux	Store Halfword with Update Indexed
X	31	7C00032E	B	sthx	Store Halfword Indexed
D	47	BC000000	B	stmw	Store Multiple Word
D	36	90000000	B	stw	Store Word



Table A-2. POWER8 Instructions by Category (Sheet 6 of 24)

Instruction Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
X	31	7C00052C	B	stwbrx	Store Word Byte-Reverse Indexed
X	31	7C00012D	B	stwcx.	Store Word Conditional Indexed and Record CR0
D	37	94000000	B	stwu	Store Word with Update
X	31	7C00016E	B	stwux	Store Word with Update Indexed
X	31	7C00012E	B	stwx	Store Word Indexed
XO	31	7C000050	B	subff[.]	Subtract From
XO	31	7C000010	B	subfc[.]	Subtract From Carrying
XO	31	7C000410	B	subfco[.]	Subtract From Carrying and Record OV
XO	31	7C000110	B	subfe[.]	Subtract From Extended
XO	31	7C000510	B	subfeo[.]	Subtract From Extended and Record OV
D	8	20000000	B	subfic	Subtract From Immediate Carrying
XO	31	7C0001D0	B	subfme[.]	Subtract From Minus One Extended
XO	31	7C0005D0	B	subfmeo[.]	Subtract From Minus One Extended and Record OV
XO	31	7C000450	B	subfo[.]	Subtract From and Record OV
XO	31	7C000190	B	subfze[.]	Subtract From Zero Extended
XO	31	7C000590	B	subfzeo[.]	Subtract From Zero Extended and Record OV
X	31	7C0004AC	B	sync	Synchronize
X	31	7C000008	B	tw	Trap Word
D	3	0C000000	B	twi	Trap Word Immediate
X	26	68000000	B	xnop	Executed No Operation
X	31	7C000278	B	xor[.]	XOR
D	26	68000000	B	xori	XOR Immediate
D	27	6C000000	B	xoris	XOR Immediate Shifted
X	59	EC000004	DFP	dadd[.]	Decimal Floating Add
X	63	FC000004	DFP	daddq[.]	Decimal Floating Add Quad
X	59	EC000644	DFP	dcffix[.]	Decimal Floating Convert From Fixed
X	63	FC000644	DFP	dcffixq[.]	Decimal Floating Convert From Fixed Quad
X	59	EC000104	DFP	dcmwo	Decimal Floating Compare Ordered
X	63	FC000104	DFP	dcmwoq	Decimal Floating Compare Ordered Quad
X	59	EC000504	DFP	dcmwu	Decimal Floating Compare Unordered
X	63	FC000504	DFP	dcmwuq	Decimal Floating Compare Unordered Quad
X	59	EC000204	DFP	dctdp[.]	Decimal Floating Convert To DFP Long
X	59	EC000244	DFP	dctfix[.]	Decimal Floating Convert To Fixed
X	63	FC000244	DFP	dctfixq[.]	Decimal Floating Convert To Fixed Quad



Table A-2. POWER8 Instructions by Category (Sheet 7 of 24)

Instruction Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
X	63	FC000204	DFP	dctqpq[.]	Decimal Floating Convert To DFP Extended
X	59	EC000284	DFP	ddedpd[.]	Decimal Floating Decode DPD To BCD
X	63	FC000284	DFP	ddedpdq[.]	Decimal Floating Decode DPD To BCD Quad
X	59	EC000444	DFP	ddiv[.]	Decimal Floating Divide
X	63	FC000444	DFP	ddivq[.]	Decimal Floating Divide Quad
X	59	EC000684	DFP	denbcd[.]	Decimal Floating Encode BCD To DPD
X	63	FC000684	DFP	denbcdq[.]	Decimal Floating Encode BCD To DPD Quad
X	59	EC0006C4	DFP	diex[.]	Decimal Floating Insert Exponent
X	63	FC0006C4	DFP	diexq[.]	Decimal Floating Insert Exponent Quad
X	59	EC000044	DFP	dmul[.]	Decimal Floating Multiply
X	63	FC000044	DFP	dmulq[.]	Decimal Floating Multiply Quad
Z23	59	EC000006	DFP	dqua[.]	Decimal Quantize
Z23	59	EC000086	DFP	dquai[.]	Decimal Quantize Immediate
Z23	63	FC000086	DFP	dquaiq[.]	Decimal Quantize Immediate Quad
Z23	63	FC000006	DFP	dquaq[.]	Decimal Quantize Quad
X	63	FC000604	DFP	drdpq[.]	Decimal Floating Round To DFP Long
Z23	59	EC0001C6	DFP	drtn[.]	Decimal Floating Round To FP Integer Without Inexact
Z23	63	FC0001C6	DFP	drtnq[.]	Decimal Floating Round To FP Integer Without Inexact Quad
Z23	59	EC0000C6	DFP	drtx[.]	Decimal Floating Round To FP Integer With Inexact
Z23	63	FC0000C6	DFP	drtxq[.]	Decimal Floating Round To FP Integer With Inexact Quad
Z23	59	EC000046	DFP	drrnd[.]	Decimal Floating Reround
Z23	63	FC000046	DFP	drrndq[.]	Decimal Floating Reround Quad
X	59	EC000604	DFP	drsp[.]	Decimal Floating Round To DFP Short
Z22	59	EC000084	DFP	dscli[.]	Decimal Floating Shift Coefficient Left Immediate
Z22	63	FC000084	DFP	dscliq[.]	Decimal Floating Shift Coefficient Left Immediate Quad
Z22	59	EC0000C4	DFP	dscrl[.]	Decimal Floating Shift Coefficient Right Immediate
Z22	63	FC0000C4	DFP	dscrlq[.]	Decimal Floating Shift Coefficient Right Immediate Quad
X	59	EC000404	DFP	dsub[.]	Decimal Floating Subtract
X	63	FC000404	DFP	dsubq[.]	Decimal Floating Subtract Quad
Z22	59	EC000184	DFP	dtstdc	Decimal Floating Test Data Class
Z22	63	FC000184	DFP	dtstdcq	Decimal Floating Test Data Class Quad
Z22	59	EC0001C4	DFP	dtstdg	Decimal Floating Test Data Group
Z22	63	FC0001C4	DFP	dtstdgq	Decimal Floating Test Data Group Quad
X	59	EC000144	DFP	dtstex	Decimal Floating Test Exponent



Table A-2. POWER8 Instructions by Category (Sheet 8 of 24)

Instruction Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
X	63	FC000144	DFP	dtstexq	Decimal Floating Test Exponent Quad
X	59	EC000544	DFP	dtstsf	Decimal Floating Test Significance
X	63	FC000544	DFP	dtstsfq	Decimal Floating Test Significance Quad
X	59	EC0002C4	DFP	dxex[.]	Decimal Floating Extract Exponent
X	63	FC0002C4	DFP	dxexq[.]	Decimal Floating Extract Exponent Quad
X	63	FC000040	FP	fcmpo	Floating Compare Ordered
X	63	FC000000	FP	fcmpu	Floating Compare Unordered
X	63	FC000100	FP	ftdiv	Floating Test for Software Divide
X	63	FC000140	FP	ftsqrt	Floating Test for Software Square Root
D	50	C8000000	FP	lfd	Load Floating-Point Double
D	51	CC000000	FP	lfdw	Load Floating-Point Double with Update
X	31	7C0004EE	FP	lfdwx	Load Floating-Point Double with Update Indexed
X	31	7C0004AE	FP	lfdx	Load Floating-Point Double Indexed
X	31	7C0006AE	FP	lfiwax	Load Floating-Point as Integer Word Algebraic Indexed
X	31	7C0006EE	FP	lfiwzx	Load Floating-Point as Integer Word and Zero Indexed
D	48	C0000000	FP	lfs	Load Floating-Point Single
D	49	C4000000	FP	lfsu	Load Floating-Point Single with Update
X	31	7C00046E	FP	lfsux	Load Floating-Point Single with Update Indexed
X	31	7C00042E	FP	lfsx	Load Floating-Point Single Indexed
X	63	FC000080	FP	mcrfs	Move To Condition Register from FPSCR
D	54	D8000000	FP	stfd	Store Floating-Point Double
D	55	DC000000	FP	stfdw	Store Floating-Point Double with Update
X	31	7C0005EE	FP	stfdwx	Store Floating-Point Double with Update Indexed
X	31	7C0005AE	FP	stfdx	Store Floating-Point Double Indexed
X	31	7C0007AE	FP	stfiwx	Store Floating-Point as Integer Word Indexed
D	52	D0000000	FP	stfs	Store Floating-Point Single
D	53	D4000000	FP	stfsu	Store Floating-Point Single with Update
X	31	7C00056E	FP	stfsux	Store Floating-Point Single with Update Indexed
X	31	7C00052E	FP	stfsx	Store Floating-Point Single Indexed
DS	57	E4000000	FP.out	lfdp	Load Floating-Point Double Pair
X	31	7C00062E	FP.out	lfdpx	Load Floating-Point Double Pair Indexed
DS	61	F4000000	FP.out	stfdp	Store Floating-Point Double Pair
X	31	7C00072E	FP.out	stfdpx	Store Floating-Point Double Pair Indexed
X	63	FC000210	FP[R]	fabs[.]	Floating Absolute Value



Table A-2. POWER8 Instructions by Category (Sheet 9 of 24)

Instruction Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
A	63	FC00002A	FP[R]	fadd[.]	Floating Add
A	59	EC00002A	FP[R]	fadds[.]	Floating Add Single
X	63	FC00069C	FP[R]	fcfid[.]	Floating Convert From Integer Doubleword
X	59	EC00069C	FP[R]	fcfids[.]	Floating Convert From Integer Doubleword Single
X	63	FC00079C	FP[R]	fcfidu[.]	Floating Convert From Integer Doubleword Unsigned
X	59	EC00079C	FP[R]	fcfidus[.]	Floating Convert From Integer Doubleword Unsigned Single
X	63	FC000010	FP[R]	fcpsgn[.]	Floating Copy Sign
X	63	FC00065C	FP[R]	fctid[.]	Floating Convert To Integer Doubleword
X	63	FC00075C	FP[R]	fctidu[.]	Floating Convert To Integer Doubleword Unsigned
X	63	FC00075E	FP[R]	fctiduz[.]	Floating Convert To Integer Doubleword Unsigned with Round Toward Zero
X	63	FC00065E	FP[R]	fctidz[.]	Floating Convert To Integer Doubleword with Round Toward Zero
X	63	FC00001C	FP[R]	fctiw[.]	Floating Convert To Integer Word
X	63	FC00011C	FP[R]	fctiwu[.]	Floating Convert To Integer Word Unsigned
X	63	FC00011E	FP[R]	fctiwuz[.]	Floating Convert To Integer Word Unsigned with Round Toward Zero
X	63	FC00001E	FP[R]	fctiwz[.]	Floating Convert To Integer Word with round to Zero
A	63	FC000024	FP[R]	fdiv[.]	Floating Divide
A	59	EC000024	FP[R]	fdivs[.]	Floating Divide Single
A	63	FC00003A	FP[R]	fmadd[.]	Floating Multiply-Add
A	59	EC00003A	FP[R]	fmadds[.]	Floating Multiply-Add Single
X	63	FC000090	FP[R]	fmr[.]	Floating Move Register
A	63	FC000038	FP[R]	fmsub[.]	Floating Multiply-Subtract
A	59	EC000038	FP[R]	fmsubs[.]	Floating Multiply-Subtract Single
A	63	FC000032	FP[R]	fmul[.]	Floating Multiply
A	59	EC000032	FP[R]	fmuls[.]	Floating Multiply Single
X	63	FC000110	FP[R]	fnabs[.]	Floating Negative Absolute Value
X	63	FC000050	FP[R]	fneg[.]	Floating Negate
A	63	FC00003E	FP[R]	fnmadd[.]	Floating Negative Multiply-Add
A	59	EC00003E	FP[R]	fnmadds[.]	Floating Negative Multiply-Add Single
A	63	FC00003C	FP[R]	fnmsub[.]	Floating Negative Multiply-Subtract
A	59	EC00003C	FP[R]	fnmsubs[.]	Floating Negative Multiply-Subtract Single
A	59	EC000030	FP[R]	fres[.]	Floating Reciprocal Estimate Single
X	63	FC000018	FP[R]	frsp[.]	Floating Round to Single-Precision
A	63	FC000034	FP[R]	frsqte[.]	Floating Reciprocal Square Root Estimate



Table A-2. POWER8 Instructions by Category (Sheet 10 of 24)

Instruction Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
A	63	FC00002E	FP[R]	fsel[.]	Floating Select
A	63	FC00002C	FP[R]	fsqrt[.]	Floating Square Root
A	59	EC00002C	FP[R]	fsqrts[.]	Floating Square Root Single
A	63	FC000028	FP[R]	fsub[.]	Floating Subtract
A	59	EC000028	FP[R]	fsubs[.]	Floating Subtract Single
X	63	FC00048E	FP[R]	mffs[.]	Move From FPSCR
X	63	FC00008C	FP[R]	mtfsb0[.]	Move To FPSCR Bit 0
X	63	FC00004C	FP[R]	mtfsb1[.]	Move To FPSCR Bit 1
XFL	63	FC00058E	FP[R]	mtfsf[.]	Move To FPSCR Fields
X	63	FC00010C	FP[R]	mtfsfi[.]	Move To FPSCR Field Immediate
A	63	FC000030	FP[R].in	fre[.]	Floating Reciprocal Estimate
X	63	FC0003D0	FP[R].in	frim[.]	Floating Round To Integer Minus
X	63	FC000310	FP[R].in	frin[.]	Floating Round To Integer Nearest
X	63	FC000390	FP[R].in	frip[.]	Floating Round To Integer Plus
X	63	FC000350	FP[R].in	friz[.]	Floating Round To Integer toward Zero
A	59	EC000034	FP[R].in	frsqtes[.]	Floating Reciprocal Square Root Estimate Single
DQ	56	E0000000	LSQ	lq	Load Quadword
X	31	7C000228	LSQ	lqarx	Load Quadword And Reserve Indexed
DS	62	F8000002	LSQ	stq	Store Quadword
X	31	7C00016D	LSQ	stqcx.	Store Quadword Conditional Indexed and record CR0
X	31	7C0004AA	MA	lswi	Load String Word Immediate
X	31	7C00042A	MA	lswx	Load String Word Indexed
X	31	7C0005AA	MA	stswi	Store String Word Immediate
X	31	7C00052A	MA	stswx	Store String Word Indexed
X	31	7C00035C	S	clrbhrb	Clear BHRB
XL	19	4C000324	S	doze	Doze
X	31	7C0006AC	S	eieio	Enforce In-order Execution of I/O
XL	19	4C000224	S	hrfid	Return From Interrupt Doubleword Hypervisor
X	31	7C0006AA	S	lbzcix	Load Byte and Zero Caching Inhibited Indexed
X	31	7C0006EA	S	ldcix	Load Doubleword Caching Inhibited Indexed
X	31	7C00066A	S	lhzcix	Load Halfword and Zero Caching Inhibited Indexed
X	31	7C00062A	S	lwzcix	Load Word and Zero Caching Inhibited Indexed
XFX	31	7C00025C	S	mfbhrbe	Move From Branch History Rolling Buffer
X	31	7C0000A6	S	mfmsr	Move From Machine State Register



Table A-2. POWER8 Instructions by Category (Sheet 11 of 24)

Instruction Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
X	31	7C0004A6	S	mfsr	Move From Segment Register
X	31	7C000526	S	mfsrin	Move From Segment Register Indirect
X	31	7C0001DC	S	msgclr	Message Clear
X	31	7C00015C	S	msgclrp	Message Clear Privileged
X	31	7C00019C	S	msgsnd	Message Send
X	31	7C00011C	S	msgsndp	Message Send Privileged
X	31	7C000124	S	mtmsr	Move To Machine State Register
X	31	7C000164	S	mtmsrd	Move To Machine State Register Doubleword
X	31	7C0001A4	S	mtsrr	Move To Segment Register
X	31	7C0001E4	S	mtsrin	Move To Segment Register Indirect
XL	19	4C000364	S	nap	Nap
XL	19	4C000124	S	rfebb	Return from Event Based Branch
XL	19	4C000024	S	rfid	Return from Interrupt Doubleword
XL	19	4C0003E4	S	rvwinkle	Rip Van Winkle
X	31	7C0007A7	S	slbfee.	SLB Find Entry ESID
X	31	7C0003E4	S	slbia	SLB Invalidate All
X	31	7C000364	S	slbie	SLB Invalidate Entry
X	31	7C000726	S	slbmfee	SLB Move From Entry ESID
X	31	7C0006A6	S	slbmfev	SLB Move From Entry VSID
X	31	7C000324	S	slbmte	SLB Move To Entry
XL	19	4C0003A4	S	sleep	Sleep
X	31	7C0007AA	S	stbcix	Store Byte Caching Inhibited Indexed
X	31	7C0007EA	S	stdcix	Store Doubleword Caching Inhibited Indexed
X	31	7C00076A	S	sthcix	Store Halfword and Zero Caching Inhibited Indexed
X	31	7C00072A	S	stwcix	Store Word and Zero Caching Inhibited Indexed
X	31	7C000264	S	tlbie	TLB Invalidate Entry
X	31	7C000224	S	tlbiel	TLB Invalidate Entry Local
X	31	7C00046C	S	tlbsync	TLB Synchronize
XFX	31	7C0002E6	S.out	mftb	Move From Time Base
X	31	7C00071D	TM	tabort.	Transaction Abort
X	31	7C00065D	TM	tabortdc.	Transaction Abort Doubleword Conditional
X	31	7C0006DD	TM	tabortdci.	Transaction Abort Doubleword Conditional Immediate
X	31	7C00061D	TM	tabortwci.	Transaction Abort Word Conditional
X	31	7C00069D	TM	tabortwci.	Transaction Abort Word Conditional Immediate



Table A-2. POWER8 Instructions by Category (Sheet 12 of 24)

Instruction Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
X	31	7C00051D	TM	tbegin.	Transaction Begin
X	31	7C00059C	TM	tcheck	Transaction Check
X	31	7C00055C	TM	tend.	Transaction End
X	31	7C0007DD	TM	trechkpt.	Transaction Recheckpoint
X	31	7C00075D	TM	treclaim.	Transaction Reclaim
VX	4	10000401	V	bcdadd.	Decimal Add Modulo
VX	4	10000441	V	bcdsub.	Decimal Subtract Modulo
X	31	7C00000E	V	lvebx	Load Vector Element Byte Indexed
X	31	7C00004E	V	lvehx	Load Vector Element Halfword Indexed
X	31	7C00008E	V	lvewx	Load Vector Element Word Indexed
X	31	7C00000C	V	lvsl	Load Vector for Shift Left
X	31	7C00004C	V	lvsr	Load Vector for Shift Right
X	31	7C0000CE	V	lvx	Load Vector Indexed
X	31	7C0002CE	V	lvxl	Load Vector Indexed Last
VX	4	10000604	V	mfvscr	Move From Vector Status and Control Register
VX	4	10000644	V	mtvscr	Move To Vector Status and Control Register
X	31	7C00010E	V	stvebx	Store Vector Element Byte Indexed
X	31	7C00014E	V	stvehx	Store Vector Element Halfword Indexed
X	31	7C00018E	V	stvewx	Store Vector Element Word Indexed
X	31	7C0001CE	V	stvx	Store Vector Indexed
X	31	7C0003CE	V	stvxl	Store Vector Indexed Last
VX	4	10000140	V	vaddcuq	Vector Add and Write Carry Unsigned Quadword
VX	4	10000180	V	vaddcuw	Vector Add and Write Carry-Out Unsigned Word
VA	4	1000003D	V	vaddecuq	Vector Add Extended and Write Carry Unsigned Quadword
VA	4	1000003C	V	vaddeuqm	Vector Add Extended Unsigned Quadword Modulo
VX	4	1000000A	V	vaddfp	Vector Add Single-Precision
VX	4	10000300	V	vaddsbs	Vector Add Signed Byte Saturate
VX	4	10000340	V	vaddshs	Vector Add Signed Halfword Saturate
VX	4	10000380	V	vaddsws	Vector Add Signed Word Saturate
VX	4	10000000	V	vaddubm	Vector Add Unsigned Byte Modulo
VX	4	10000200	V	vaddubs	Vector Add Unsigned Byte Saturate
VX	4	100000C0	V	vaddudm	Vector Add Unsigned Doubleword Modulo
VX	4	10000040	V	vadduhm	Vector Add Unsigned Halfword Modulo
VX	4	10000240	V	vadduhs	Vector Add Unsigned Halfword Saturate



Table A-2. POWER8 Instructions by Category (Sheet 13 of 24)

Instruction Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
VX	4	10000100	V	vadduqm	Vector Add Unsigned Quadword Modulo
VX	4	10000080	V	vadduwrm	Vector Add Unsigned Word Modulo
VX	4	10000280	V	vadduws	Vector Add Unsigned Word Saturate
VX	4	10000404	V	vand	Vector Logical AND
VX	4	10000444	V	vandc	Vector Logical AND with Complement
VX	4	10000502	V	vavgsb	Vector Average Signed Byte
VX	4	10000542	V	vavgsh	Vector Average Signed Halfword
VX	4	10000582	V	vavgsw	Vector Average Signed Word
VX	4	10000402	V	vavgub	Vector Average Unsigned Byte
VX	4	10000442	V	vavguh	Vector Average Unsigned Halfword
VX	4	10000482	V	vavguw	Vector Average Unsigned Word
VX	4	1000054C	V	vbpermq	Vector Bit Permute Quadword
VX	4	1000054C	V	vbpermq	Vector Bit Permute Quadword
VX	4	1000034A	V	vcfsx	Vector Convert From Signed Fixed-Point Word To Single-Precision
VX	4	1000030A	V	vcfux	Vector Convert From Unsigned Fixed-Point Word
VX	4	10000702	V	vclzb	Vector Count Leading Zeros Byte
VX	4	100007C2	V	vclzd	Vector Count Leading Zeros Doubleword
VX	4	10000742	V	vclzh	Vector Count Leading Zeros Halfword
VX	4	10000782	V	vclzw	Vector Count Leading Zeros Word
VC	4	100003C6	V	vcmpbfp[.]	Vector Compare Bounds Single-Precision
VC	4	100000C6	V	vcmpeqfp[.]	Vector Compare Equal To Single-Precision
VC	4	10000006	V	vcmpequb[.]	Vector Compare Equal To Unsigned Byte
VC	4	100000C7	V	vcmpequd[.]	Vector Compare Equal To Unsigned Doubleword
VC	4	10000046	V	vcmpequh[.]	Vector Compare Equal To Unsigned Halfword
VC	4	10000086	V	vcmpequw[.]	Vector Compare Equal To Unsigned Word
VC	4	100001C6	V	vcmpgef[.]	Vector Compare Greater Than or Equal To Single-Precision
VC	4	100002C6	V	vcmpgtfp[.]	Vector Compare Greater Than Single-Precision
VC	4	10000306	V	vcmpgtsb[.]	Vector Compare Greater Than Signed Byte
VC	4	100003C7	V	vcmpgtsd[.]	Vector Compare Greater Than Signed Doubleword
VC	4	10000346	V	vcmpgtsh[.]	Vector Compare Greater Than Signed Halfword
VC	4	10000386	V	vcmpgtsw[.]	Vector Compare Greater Than Signed Word
VC	4	10000206	V	vcmpgtub[.]	Vector Compare Greater Than Unsigned Byte
VC	4	100002C7	V	vcmpgtud[.]	Vector Compare Greater Than Unsigned Doubleword
VC	4	10000246	V	vcmpgtuh[.]	Vector Compare Greater Than Unsigned Halfword



Table A-2. POWER8 Instructions by Category (Sheet 14 of 24)

Instruction Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
VC	4	10000286	V	vcmpgtuw[.]	Vector Compare Greater Than Unsigned Word
VX	4	100003CA	V	vctsxs	Vector Convert From Single-Precision To Signed Fixed-Point Word Saturate
VX	4	1000038A	V	vctuxs	Vector Convert From Single-Precision To Unsigned Fixed-Point Word Saturate
VX	4	10000684	V	veqv	Vector Equivalence
VX	4	1000018A	V	vexptefp	Vector 2 Raised to the Exponent Estimate Single-Precision
VX	4	1000050C	V	vgbbd	Vector Gather Bits by Byte by Doubleword
VX	4	100001CA	V	vlogefp	Vector Log Base 2 Estimate Single-Precision
VA	4	1000002E	V	vmaddfp	Vector Multiply-Add Single-Precision
VX	4	1000040A	V	vmaxfp	Vector Maximum Single-Precision
VX	4	10000102	V	vmaxsb	Vector Maximum Signed Byte
VX	4	100001C2	V	vmaxsd	Vector Maximum Signed Doubleword
VX	4	10000142	V	vmaxsh	Vector Maximum Signed Halfword
VX	4	10000182	V	vmaxsw	Vector Maximum Signed Word
VX	4	10000002	V	vmaxub	Vector Maximum Unsigned Byte
VX	4	100000C2	V	vmaxud	Vector Maximum Unsigned Doubleword
VX	4	10000042	V	vmaxuh	Vector Maximum Unsigned Halfword
VX	4	10000082	V	vmaxuw	Vector Maximum Unsigned Word
VA	4	10000020	V	vhaddshs	Vector Multiply-High-Add Signed Halfword Saturate
VA	4	10000021	V	vhhraddshs	Vector Multiply-High-Round-Add Signed Halfword Saturate
VX	4	1000044A	V	vminfp	Vector Minimum Single-Precision
VX	4	10000302	V	vminsb	Vector Minimum Signed Byte
X	4	100003C2	V	vminsd	Vector Minimum Signed Doubleword
VX	4	10000342	V	vminsh	Vector Minimum Signed Halfword
VX	4	10000382	V	vminsw	Vector Minimum Signed Word
VX	4	10000202	V	vminub	Vector Minimum Unsigned Byte
VX	4	100002C2	V	vminud	Vector Minimum Unsigned Doubleword
VX	4	10000242	V	vminuh	Vector Minimum Unsigned Halfword
VX	4	10000282	V	vminuw	Vector Minimum Unsigned Word
VA	4	10000022	V	vmladduhm	Vector Multiply-Low-Add Unsigned Halfword Modulo
VX	4	1000000C	V	vmrghb	Vector Merge High Byte
VX	4	1000004C	V	vmrghh	Vector Merge High Halfword
VX	4	1000008C	V	vmrghw	Vector Merge High Word
VX	4	1000010C	V	vmrglb	Vector Merge Low Byte



Table A-2. POWER8 Instructions by Category (Sheet 15 of 24)

Instruction Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
VX	4	1000014C	V	vmrglh	Vector Merge Low Halfword
VX	4	1000018C	V	vmrglw	Vector Merge Low Word
VA	4	10000025	V	vmsummbrm	Vector Multiply-Sum Mixed Byte Modulo
VA	4	10000028	V	vmsumshm	Vector Multiply-Sum Signed Halfword Modulo
VA	4	10000029	V	vmsumshs	Vector Multiply-Sum Signed Halfword Saturate
VA	4	10000024	V	vmsumubm	Vector Multiply-Sum Unsigned Byte Modulo
VA	4	10000026	V	vmsumuhm	Vector Multiply-Sum Unsigned Halfword Modulo
VA	4	10000027	V	vmsumuhs	Vector Multiply-Sum Unsigned Halfword Saturate
VX	4	10000308	V	vmulesb	Vector Multiply Even Signed Byte
VX	4	10000348	V	vmulesh	Vector Multiply Even Signed Halfword
VX	4	10000388	V	vmulesw	Vector Multiply Even Signed Word
VX	4	10000208	V	vmuleub	Vector Multiply Even Unsigned Byte
VX	4	10000248	V	vmuleuh	Vector Multiply Even Unsigned Halfword
VX	4	10000288	V	vmuleuw	Vector Multiply Even Unsigned Word
VX	4	10000108	V	vmulosb	Vector Multiply Odd Signed Byte
VX	4	10000148	V	vmulosh	Vector Multiply Odd Signed Halfword
VX	4	10000188	V	vmulosw	Vector Multiply Odd Signed Word
VX	4	10000008	V	vmuloub	Vector Multiply Odd Unsigned Byte
VX	4	10000048	V	vmulouh	Vector Multiply Odd Unsigned Halfword
VX	4	10000088	V	vmulouw	Vector Multiply Odd Unsigned Word
VX	4	10000089	V	vmuluwm	Vector Multiply Unsigned Word Modulo
VX	4	10000584	V	vnand	Vector NAND
VA	4	1000002F	V	vnmsubfp	Vector Negative Multiply-Subtract Single-Precision
VX	4	10000504	V	vnor	Vector Logical NOR
VX	4	10000484	V	vor	Vector Logical OR
VX	4	10000544	V	vorc	Vector OR with Complement
VA	4	1000002B	V	vperm	Vector Permute
VX	4	1000030E	V	vpkpx	Vector Pack Pixel
VX	4	100005CE	V	vpksdss	Vector Pack Signed Doubleword Signed Saturate
VX	4	1000054E	V	vpksdus	Vector Pack Signed Doubleword Unsigned Saturate
VX	4	1000018E	V	vpkshss	Vector Pack Signed Halfword Signed Saturate
VX	4	1000010E	V	vpkshus	Vector Pack Signed Halfword Unsigned Saturate
VX	4	100001CE	V	vpkswss	Vector Pack Signed Word Signed Saturate
VX	4	1000014E	V	vpkswus	Vector Pack Signed Word Unsigned Saturate



Table A-2. POWER8 Instructions by Category (Sheet 16 of 24)

Instruction Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
VX	4	1000044E	V	vpkudum	Vector Pack Unsigned Doubleword Unsigned Modulo
VX	4	100004CE	V	vpkudus	Vector Pack Unsigned Doubleword Unsigned Saturate
VX	4	1000000E	V	vpkuhum	Vector Pack Unsigned Halfword Unsigned Modulo
VX	4	1000008E	V	vpkuhus	Vector Pack Unsigned Halfword Unsigned Saturate
VX	4	1000004E	V	vpkuwum	Vector Pack Unsigned Word Unsigned Modulo
VX	4	100000CE	V	vpkuwus	Vector Pack Unsigned Word Unsigned Saturate
VX	4	10000408	V	vpmsumb	Vector Polynomial Multiply-Sum Byte
VX	4	100004C8	V	vpmsumd	Vector Polynomial Multiply-Sum Doubleword
VX	4	10000448	V	vpmsumh	Vector Polynomial Multiply-Sum Halfword
VX	4	10000488	V	vpmsumw	Vector Polynomial Multiply-Sum Word
VX	4	10000703	V	vpopcntb	Vector Population Count Byte
VX	4	100007C3	V	vpopcntd	Vector Population Count Doubleword
VX	4	10000743	V	vpopcnth	Vector Population Count Halfword
VX	4	10000783	V	vpopcntw	Vector Population Count Word
VX	4	1000010A	V	vrefp	Vector Reciprocal Estimate Single-Precision
VX	4	100002CA	V	vrfim	Vector Round to Single-Precision Integer toward -Infinity
VX	4	1000020A	V	vrfin	Vector Round to Single-Precision Integer Nearest
VX	4	1000028A	V	vrfip	Vector Round to Single-Precision Integer toward +Infinity
VX	4	1000024A	V	vrfiz	Vector Round to Single-Precision Integer toward Zero
VX	4	10000004	V	vrlb	Vector Rotate Left Byte
VX	4	100000C4	V	vrld	Vector Rotate Left Doubleword
VX	4	10000044	V	vrlh	Vector Rotate Left Halfword
VX	4	10000084	V	vrlw	Vector Rotate Left Word
VX	4	1000014A	V	vrsqrtefp	Vector Reciprocal Square Root Estimate Single-Precision
VA	4	1000002A	V	vsel	Vector Select
VX	4	100001C4	V	vsl	Vector Shift Left
VX	4	10000104	V	vslb	Vector Shift Left Byte
VX	4	100005C4	V	vslid	Vector Shift Left Doubleword
VA	4	1000002C	V	vsldoi	Vector Shift Left Double by Octet Immediate
VX	4	10000144	V	vslh	Vector Shift Left Halfword
VX	4	1000040C	V	vslo	Vector Shift Left by Octet
VX	4	10000184	V	vslw	Vector Shift Left Word
VX	4	1000020C	V	vspltb	Vector Splat Byte
VX	4	1000024C	V	vsplth	Vector Splat Halfword



Table A-2. POWER8 Instructions by Category (Sheet 17 of 24)

Instruction Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
VX	4	1000030C	V	vspltsb	Vector Splat Immediate Signed Byte
VX	4	1000034C	V	vspltsih	Vector Splat Immediate Signed Halfword
VX	4	1000038C	V	vspltsiw	Vector Splat Immediate Signed Word
VX	4	1000028C	V	vspltw	Vector Splat Word
VX	4	100002C4	V	vsr	Vector Shift Right
VX	4	10000304	V	vsrab	Vector Shift Right Algebraic Byte
VX	4	100003C4	V	vsrad	Vector Shift Right Algebraic Doubleword
VX	4	10000344	V	vsrah	Vector Shift Right Algebraic Halfword
VX	4	10000384	V	vsraw	Vector Shift Right Algebraic Word
VX	4	10000204	V	vsrb	Vector Shift Right Byte
VX	4	100006C4	V	vsrd	Vector Shift Right Doubleword
VX	4	10000244	V	vsrh	Vector Shift Right Halfword
VX	4	1000044C	V	vsro	Vector Shift Right by Octet
VX	4	10000284	V	vsrw	Vector Shift Right Word
VX	4	10000540	V	vsubcuq	Vector Subtract and Write Carry Unsigned Quadword
VX	4	10000580	V	vsubcuw	Vector Subtract and Write Carry-Out Unsigned Word
VA	4	1000003F	V	vsubecuq	Vector Subtract Extended and Write Carry Unsigned Quadword
VA	4	1000003E	V	vsubeuqm	Vector Subtract Extended Unsigned Quadword Modulo
VX	4	1000004A	V	vsubfp	Vector Subtract Single-Precision
VX	4	10000700	V	vsubbs	Vector Subtract Signed Byte Saturate
VX	4	10000740	V	vsubshs	Vector Subtract Signed Halfword Saturate
VX	4	10000780	V	vsubsws	Vector Subtract Signed Word Saturate
VX	4	10000400	V	vsububm	Vector Subtract Unsigned Byte Modulo
VX	4	10000600	V	vsububus	Vector Subtract Unsigned Byte Saturate
VX	4	100004C0	V	vsubudm	Vector Subtract Unsigned Doubleword Modulo
VX	4	10000440	V	vsubuhm	Vector Subtract Unsigned Halfword Modulo
VX	4	10000640	V	vsubuhs	Vector Subtract Unsigned Halfword Saturate
VX	4	10000500	V	vsubuqm	Vector Subtract Unsigned Quadword Modulo
VX	4	10000480	V	vsubuwm	Vector Subtract Unsigned Word Modulo
VX	4	10000680	V	vsubuws	Vector Subtract Unsigned Word Saturate
VX	4	10000688	V	vsum2sws	Vector Sum across Half Signed Word Saturate
VX	4	10000708	V	vsum4sbs	Vector Sum across Quarter Signed Byte Saturate
VX	4	10000648	V	vsum4shs	Vector Sum across Quarter Signed Halfword Saturate
VX	4	10000608	V	vsum4ubs	Vector Sum across Quarter Unsigned Byte Saturate



Table A-2. POWER8 Instructions by Category (Sheet 18 of 24)

Instruction Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
VX	4	10000788	V	vsumsws	Vector Sum across Signed Word Saturate
VX	4	1000034E	V	vupkhpx	Vector Unpack High Pixel
VX	4	1000020E	V	vupkhsb	Vector Unpack High Signed Byte
VX	4	1000024E	V	vupkhsh	Vector Unpack High Signed Halfword
VX	4	1000064E	V	vupkhsw	Vector Unpack High Signed Word
VX	4	100003CE	V	vupklpx	Vector Unpack Low Pixel
VX	4	1000028E	V	vupklsb	Vector Unpack Low Signed Byte
VX	4	100002CE	V	vupklsh	Vector Unpack Low Signed Halfword
VX	4	100006CE	V	vupklsw	Vector Unpack Low Signed Word
VX	4	100004C4	V	vxor	Vector Logical XOR
VX	4	10000508	V.AES	vcipher	Vector AES Cipher
VX	4	10000509	V.AES	vcipherlast	Vector AES Cipher Last
VX	4	10000548	V.AES	vncipher	Vector AES Inverse Cipher
VX	4	10000549	V.AES	vncipherlast	Vector AES Inverse Cipher Last
VX	4	100005C8	V.AES	vsbox	Vector AES S-Box
VA	4	1000002D	V.RAID	vpermxor	Vector Permute and Exclusive-OR
VX	4	100006C2	V.SHA2	vshasigmad	Vector SHA-512 Sigma Doubleword
VX	4	10000682	V.SHA2	vshasigmaw	Vector SHA-256 Sigma Word
X	63	FC00078C	VSX	fmrgew	Floating Merge Even Word
X	63	FC00068C	VSX	fmgow	Floating Merge Odd Word
XX1	31	7C000498	VSX	lxsdx	Load VSR Scalar Doubleword Indexed
XX1	31	7C000098	VSX	lxsiwax	Load VSX Scalar as Integer Word Algebraic Indexed
XX1	31	7C000018	VSX	lxsiwzx	Load VSX Scalar as Integer Word and Zero Indexed
XX1	31	7C000418	VSX	lxsspx	Load VSX Scalar Single-Precision Indexed
XX1	31	7C000698	VSX	lxvd2x	Load VSR Vector Doubleword*2 Indexed
XX1	31	7C000298	VSX	lxvdsx	Load VSR Vector Doubleword and Splat Indexed
XX1	31	7C000618	VSX	lxvw4x	Load VSR Vector Word*4 Indexed
XX1	31	7C000066	VSX	mfvsrd	Move From VSR Doubleword
XX1	31	7C0000E6	VSX	mfvsrwz	Move From VSR Word and Zero
XX1	31	7C000166	VSX	mtvsrd	Move To VSR Doubleword
XX1	31	7C0001A6	VSX	mtvsrwa	Move To VSR Word Algebraic
XX1	31	7C0001E6	VSX	mtvsrwz	Move To VSR Word and Zero
XX1	31	7C000598	VSX	stxsdx	Store VSR Scalar Doubleword Indexed
XX1	31	7C000118	VSX	stxsiwx	Store VSX Scalar as Integer Word Indexed



Table A-2. POWER8 Instructions by Category (Sheet 19 of 24)

Instruction Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
XX1	31	7C000518	VSX	stxsspx	Store VSR Scalar Word Indexed
XX1	31	7C000798	VSX	stxvd2x	Store VSR Vector Doubleword*2 Indexed
XX1	31	7C000718	VSX	stvw4x	Store VSR Vector Word*4 Indexed
VX	4	1000078C	VSX	vmrgew	Vector Merge Even Word
VX	4	1000068C	VSX	vmrgow	Vector Merge Odd Word
XX2	60	F0000564	VSX	xsabsdp	VSX Scalar Absolute Value Double-Precision
XX3	60	F0000100	VSX	xsadddp	VSX Scalar Add Double-Precision
XX3	60	F0000000	VSX	xsaddsp	VSX Scalar Add Single-Precision
XX3	60	F0000158	VSX	xscmpodp	VSX Scalar Compare Ordered Double-Precision
XX3	60	F0000118	VSX	xscmpudp	VSX Scalar Compare Unordered Double-Precision
XX3	60	F0000580	VSX	xscpsgrdp	VSX Scalar Copy Sign Double-Precision
XX2	60	F0000424	VSX	xscvdsp	VSX Scalar Convert Double-Precision to Single-Precision
XX2	60	F000042C	VSX	xscvdspn	VSX Scalar Convert Double-Precision to Single-Precision format Non-signalling
XX2	60	F0000560	VSX	xscvdpsxds	VSX Scalar Convert Double-Precision to Signed Fixed-Point Doubleword Saturate
XX2	60	F0000160	VSX	xscvdpsxws	VSX Scalar Convert Double-Precision to Signed Fixed-Point Word Saturate
XX2	60	F0000520	VSX	xscvdpxuds	VSX Scalar Convert Double-Precision to Unsigned Fixed-Point Doubleword Saturate
XX2	60	F0000120	VSX	xscvdpxuws	VSX Scalar Convert Double-Precision to Unsigned Fixed-Point Word Saturate
XX2	60	F0000524	VSX	xscvspdp	VSX Scalar Convert Single-Precision to Double-Precision (p=1)
XX2	60	F000052C	VSX	xscvspdpn	Scalar Convert Single-Precision to Double-Precision format Non-signalling
XX2	60	F00005E0	VSX	xscvsxddp	VSX Scalar Convert Signed Fixed-Point Doubleword to Double-Precision
XX2	60	F00004E0	VSX	xscvsxdsp	VSX Scalar Convert Signed Fixed-Point Doubleword to Single-Precision
XX2	60	F00005A0	VSX	xscvuxddp	VSX Scalar Convert Unsigned Fixed-Point Doubleword to Double-Precision
XX2	60	F00004A0	VSX	xscvuxdsp	VSX Scalar Convert Unsigned Fixed-Point Doubleword to Single-Precision
XX3	60	F00001C0	VSX	xsdvidp	VSX Scalar Divide Double-Precision
XX3	60	F00000C0	VSX	xsdivsp	VSX Scalar Divide Single-Precision
XX3	60	F0000108	VSX	xsmaddadp	VSX Scalar Multiply-Add Type-A Double-Precision
XX3	60	F0000008	VSX	xsmaddasp	VSX Scalar Multiply-Add Type-A Single-Precision
XX3	60	F0000148	VSX	xsmaddmdp	VSX Scalar Multiply-Add Type-M Double-Precision



Table A-2. POWER8 Instructions by Category (Sheet 20 of 24)

Instruction Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
XX3	60	F0000048	VSX	xsmaddmsp	VSX Scalar Multiply-Add Type-M Single-Precision
XX3	60	F0000500	VSX	xsmadp	VSX Scalar Maximum Double-Precision
XX3	60	F0000540	VSX	xsmindp	VSX Scalar Minimum Double-Precision
XX3	60	F0000188	VSX	xsmsubadp	VSX Scalar Multiply-Subtract Type-A Double-Precision
XX3	60	F0000088	VSX	xsmsubasp	VSX Scalar Multiply-Subtract Type-A Single-Precision
XX3	60	F00001C8	VSX	xsmsubmdp	VSX Scalar Multiply-Subtract Type-M Double-Precision
XX3	60	F00000C8	VSX	xsmsubmsp	VSX Scalar Multiply-Subtract Type-M Single-Precision
XX3	60	F0000180	VSX	xsmuldp	VSX Scalar Multiply Double-Precision
XX3	60	F0000080	VSX	xsmulsp	VSX Scalar Multiply Single-Precision
XX2	60	F00005A4	VSX	xsnabsdp	VSX Scalar Negative Absolute Value Double-Precision
XX2	60	F00005E4	VSX	xsnegdp	VSX Scalar Negate Double-Precision
XX3	60	F0000508	VSX	xsnmaddadp	VSX Scalar Negative Multiply-Add Type-A Double-Precision
XX3	60	F0000408	VSX	xsnmaddasp	VSX Scalar Negative Multiply-Add Type-A Single-Precision
XX3	60	F0000548	VSX	xsnmaddmdp	VSX Scalar Negative Multiply-Add Type-M Double-Precision
XX3	60	F0000448	VSX	xsnmaddmsp	VSX Scalar Negative Multiply-Add Type-M Single-Precision
XX3	60	F0000588	VSX	xsnmsubadp	VSX Scalar Negative Multiply-Subtract Type-A Double-Precision
XX3	60	F0000488	VSX	xsnmsubasp	VSX Scalar Negative Multiply-Subtract Type-A Single-Precision
XX3	60	F00005C8	VSX	xsnmsubmdp	VSX Scalar Negative Multiply-Subtract Type-M Double-Precision
XX3	60	F00004C8	VSX	xsnmsubmsp	VSX Scalar Negative Multiply-Subtract Type-M Single-Precision
XX2	60	F0000124	VSX	xsrdpi	VSX Scalar Round to Double-Precision Integer
XX2	60	F00001AC	VSX	xsrpic	VSX Scalar Round to Double-Precision Integer Using Current Rounding Mode
XX2	60	F00001E4	VSX	xsrpim	VSX Scalar Round to Double-Precision Integer toward -Infinity
XX2	60	F00001A4	VSX	xsrpip	VSX Scalar Round to Double-Precision Integer toward +Infinity
XX2	60	F0000164	VSX	xsrpiz	VSX Scalar Round to Double-Precision Integer toward Zero
XX1	60	F0000168	VSX	xsredp	VSX Scalar Reciprocal Estimate Double-Precision
XX2	60	F0000068	VSX	xsresp	VSX Scalar Reciprocal Estimate Single-Precision
XX2	60	F0000464	VSX	xrsp	VSX Scalar Round to Single-Precision
XX2	60	F0000128	VSX	xrsqrtepd	VSX Scalar Reciprocal Square Root Estimate Double-Precision
XX2	60	F0000028	VSX	xrsqrtesp	VSX Scalar Reciprocal Square Root Estimate Single-Precision



Table A-2. POWER8 Instructions by Category (Sheet 21 of 24)

Instruction Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
XX2	60	F000012C	VSX	xssqrtdp	VSX Scalar Square Root Double-Precision
XX2	60	F000002C	VSX	xssqrtsp	VSX Scalar Square Root Single-Precision
XX3	60	F0000140	VSX	xssubdp	VSX Scalar Subtract Double-Precision
XX3	60	F0000040	VSX	xssubsp	VSX Scalar Subtract Single-Precision
XX3	60	F00001E8	VSX	xstdivdp	VSX Scalar Test for Software Divide Double-Precision
XX2	60	F00001A8	VSX	xstsqrtdp	VSX Scalar Test for Software Square Root Double-Precision
XX2	60	F0000764	VSX	xvabsdp	VSX Vector Absolute Value Double-Precision
XX2	60	F0000664	VSX	xvabssp	VSX Vector Absolute Value Single-Precision
XX3	60	F0000300	VSX	xvadddp	VSX Vector Add Double-Precision
XX3	60	F0000200	VSX	xvaddsp	VSX Vector Add Single-Precision
XX3	60	F0000318	VSX	xvcmpqdp	VSX Vector Compare Equal To Double-Precision
XX3	60	F0000718	VSX	xvcmpqdp.	VSX Vector Compare Equal To Double-Precision and Record CR6
XX3	60	F0000218	VSX	xvcmpqsp	VSX Vector Compare Equal To Single-Precision
XX3	60	F0000618	VSX	xvcmpqsp.	VSX Vector Compare Equal To Single-Precision and Record CR6
XX3	60	F0000398	VSX	xvcmpgedp	VSX Vector Compare Greater Than or Equal To Double-Precision
XX3	60	F0000798	VSX	xvcmpgedp.	VSX Vector Compare Greater Than or Equal To Double-Precision and Record CR6
XX3	60	F0000298	VSX	xvcmpgesp	VSX Vector Compare Greater Than or Equal To Single-Precision
XX3	60	F0000698	VSX	xvcmpgesp.	VSX Vector Compare Greater Than or Equal To Single-Precision and Record CR6
XX3	60	F0000358	VSX	xvcmpgtdp	VSX Vector Compare Greater Than Double-Precision
XX3	60	F0000758	VSX	xvcmpgtdp.	VSX Vector Compare Greater Than Double-Precision and Record CR6
XX3	60	F0000258	VSX	xvcmpgtsp	VSX Vector Compare Greater Than Single-Precision
XX3	60	F0000658	VSX	xvcmpgtsp.	VSX Vector Compare Greater Than Single-Precision and Record CR6
XX3	60	F0000780	VSX	xvcpsgrndp	VSX Vector Copy Sign Double-Precision
XX3	60	F0000680	VSX	xvcpsgnsp	VSX Vector Copy Sign Single-Precision
XX2	60	F0000624	VSX	xvcvdpsp	VSX Vector Convert Double-Precision to Single-Precision
XX2	60	F0000760	VSX	xvcvdpsxds	VSX Vector Convert Double-Precision to Signed Fixed-Point Doubleword Saturate
XX2	60	F0000360	VSX	xvcvdpsxws	VSX Vector Convert Double-Precision to Signed Fixed-Point Word Saturate
XX2	60	F0000720	VSX	xvcvdpxuds	VSX Vector Convert Double-Precision to Unsigned Fixed-Point Doubleword Saturate



Table A-2. POWER8 Instructions by Category (Sheet 22 of 24)

Instruction Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
XX2	60	F0000320	VSX	xvcvdpuxws	VSX Vector Convert Double-Precision to Unsigned Fixed-Point Word Saturate
XX2	60	F0000724	VSX	xvcvspdp	VSX Vector Convert Single-Precision to Double-Precision
XX2	60	F0000660	VSX	xvcvpsxds	VSX Vector Convert Single-Precision to Signed Fixed-Point Doubleword Saturate
XX2	60	F0000260	VSX	xvcvspsxws	VSX Vector Convert Single-Precision to Signed Fixed-Point Word Saturate
XX2	60	F0000620	VSX	xvcvspuxds	VSX Vector Convert Single-Precision to Unsigned Fixed-Point Doubleword Saturate
XX2	60	F0000220	VSX	xvcvspuxws	VSX Vector Convert Single-Precision to Unsigned Fixed-Point Word Saturate
XX2	60	F00007E0	VSX	xvcvsxddp	VSX Vector Convert Signed Fixed-Point Doubleword to Double-Precision
XX2	60	F00006E0	VSX	xvcvsxdsp	VSX Vector Convert Signed Fixed-Point Doubleword to Single-Precision
XX2	60	F00003E0	VSX	xvcvsxwdp	VSX Vector Convert Signed Fixed-Point Word to Double-Precision
XX2	60	F00002E0	VSX	xvcvsxwsp	VSX Vector Convert Signed Fixed-Point Word to Single-Precision
XX2	60	F00007A0	VSX	xvcvuxddp	VSX Vector Convert Unsigned Fixed-Point Doubleword to Double-Precision
XX2	60	F00006A0	VSX	xvcvuxdsp	VSX Vector Convert Unsigned Fixed-Point Doubleword to Single-Precision
XX2	60	F00003A0	VSX	xvcvuxwdp	VSX Vector Convert Unsigned Fixed-Point Word to Double-Precision
XX2	60	F00002A0	VSX	xvcvuxwsp	VSX Vector Convert Unsigned Fixed-Point Word to Single-Precision
XX3	60	F00003C0	VSX	xvdvdp	VSX Vector Divide Double-Precision
XX3	60	F00002C0	VSX	xvdvsp	VSX Vector Divide Single-Precision
XX3	60	F0000308	VSX	xvmaddadp	VSX Vector Multiply-Add Type-A Double-Precision
XX3	60	F0000208	VSX	xvmaddasp	VSX Vector Multiply-Add Type-A Single-Precision
XX3	60	F0000348	VSX	xvmaddmdp	VSX Vector Multiply-Add Type-M Double-Precision
XX3	60	F0000248	VSX	xvmaddmsp	VSX Vector Multiply-Add Type-M Single-Precision
XX3	60	F0000700	VSX	xvmaxdp	VSX Vector Maximum Double-Precision
XX3	60	F0000600	VSX	xvmaxsp	VSX Vector Maximum Single-Precision
XX3	60	F0000740	VSX	xvmindp	VSX Vector Minimum Double-Precision
XX3	60	F0000640	VSX	xvminsp	VSX Vector Minimum Single-Precision
XX3	60	F0000388	VSX	xvmsubadp	VSX Vector Multiply-Subtract Type-A Double-Precision
XX3	60	F0000288	VSX	xvmsubasp	VSX Vector Multiply-Subtract Type-A Single-Precision
XX3	60	F00003C8	VSX	xvmsubmdp	VSX Vector Multiply-Subtract Type-M Double-Precision



Table A-2. POWER8 Instructions by Category (Sheet 23 of 24)

Instruction Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
XX3	60	F00002C8	VSX	xvmsubmsp	VSX Vector Multiply-Subtract Type-M Single-Precision
XX3	60	F0000380	VSX	xvmuldp	VSX Vector Multiply Double-Precision
XX3	60	F0000280	VSX	xvmulsp	VSX Vector Multiply Single-Precision
XX2	60	F00007A4	VSX	xvnabsdp	VSX Vector Negative Absolute Value Double-Precision
XX2	60	F00006A4	VSX	xvnabssp	VSX Vector Negative Absolute Value Single-Precision
XX2	60	F00007E4	VSX	xvnegdp	VSX Vector Negate Double-Precision
XX2	60	F00006E4	VSX	xvnegsp	VSX Vector Negate Single-Precision
XX3	60	F0000708	VSX	xvnmaddadp	VSX Vector Negative Multiply-Add Type-A Double-Precision
XX3	60	F0000608	VSX	xvnmaddasp	VSX Vector Negative Multiply-Add Type-A Single-Precision
XX3	60	F0000748	VSX	xvnmaddmdp	VSX Vector Negative Multiply-Add Type-M Double-Precision
XX3	60	F0000648	VSX	xvnmaddmsp	VSX Vector Negative Multiply-Add Type-M Single-Precision
XX3	60	F0000788	VSX	xvnmsubadp	VSX Vector Negative Multiply-Subtract Type-A Double-Precision
XX3	60	F0000688	VSX	xvnmsubasp	VSX Vector Negative Multiply-Subtract Type-A Single-Precision
XX3	60	F00007C8	VSX	xvnmsubmdp	VSX Vector Negative Multiply-Subtract Type-M Double-Precision
XX3	60	F00006C8	VSX	xvnmsubmsp	VSX Vector Negative Multiply-Subtract Type-M Single-Precision
XX2	60	F0000324	VSX	xvrdpi	VSX Vector Round to Double-Precision Integer
XX2	60	F00003AC	VSX	xvrdpic	VSX Vector Round to Double-Precision Integer using Current Rounding Mode
XX2	60	F00003E4	VSX	xvrdpim	VSX Vector Round to Double-Precision Integer toward -Infinity
XX2	60	F00003A4	VSX	xvrdpip	VSX Vector Round to Double-Precision Integer toward +Infinity
XX2	60	F0000364	VSX	xvrdpiz	VSX Vector Round to Double-Precision Integer toward Zero
XX2	60	F0000368	VSX	xvredp	VSX Vector Reciprocal Estimate Double-Precision
XX2	60	F0000268	VSX	xvresp	VSX Vector Reciprocal Estimate Single-Precision
XX2	60	F0000224	VSX	xvrspi	VSX Vector Round to Single-Precision Integer
XX2	60	F00002AC	VSX	xvrspic	VSX Vector Round to Single-Precision Integer using Current Rounding Mode
XX2	60	F00002E4	VSX	xvrspim	VSX Vector Round to Single-Precision Integer toward -Infinity
XX2	60	F00002A4	VSX	xvrspip	VSX Vector Round to Single-Precision Integer toward +Infinity
XX2	60	F0000264	VSX	xvrspiz	VSX Vector Round to Single-Precision Integer toward Zero
XX2	60	F0000328	VSX	xvrsqrtdp	VSX Vector Reciprocal Square Root Estimate Double-Precision



Table A-2. POWER8 Instructions by Category (Sheet 24 of 24)

Instruction Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
XX2	60	F0000228	VSX	xvrsqrtesp	VSX Vector Reciprocal Square Root Estimate Single-Precision
XX2	60	F000032C	VSX	xvsqrtdp	VSX Vector Square Root Double-Precision
XX2	60	F000022C	VSX	xvsqrtsp	VSX Vector Square Root Single-Precision
XX3	60	F0000340	VSX	xvsubdp	VSX Vector Subtract Double-Precision
XX3	60	F0000240	VSX	xvsubsp	VSX Vector Subtract Single-Precision
XX3	60	F00003E8	VSX	xvtdivdp	VSX Vector Test for Software Divide Double-Precision
XX3	60	F00002E8	VSX	xvtdivsp	VSX Vector Test for Software Divide Single-Precision
XX2	60	F00003A8	VSX	xvtsqrtdp	VSX Vector Test for Software Square Root Double-Precision
XX2	60	F00002A8	VSX	xvtsqrtsp	VSX Vector Test for Software Square Root Single-Precision
XX3	60	F0000410	VSX	xxland	VSX Logical AND
XX3	60	F0000450	VSX	xxlandc	VSX Logical AND with Complement
XX3	60	F00005D0	VSX	xxleqv	VSX Logical Equivalence
XX3	60	F0000590	VSX	xxlnand	VSX Logical NAND
XX3	60	F0000510	VSX	xxlnor	VSX Logical NOR
XX3	60	F0000490	VSX	xxlor	VSX Logical OR
XX3	60	F0000550	VSX	xxlorc	VSX Logical OR with Complement
XX3	60	F00004D0	VSX	xxlxor	VSX Logical XOR
XX3	60	F0000090	VSX	xxmrghw	VSX Merge High Word
XX3	60	F0000190	VSX	xxmrglw	VSX Merge Low Word
XX3	60	F0000050	VSX	xxpermdi	VSX Permute Doubleword Immediate
XX4	60	F0000030	VSX	xxsel	VSX Select
XX3	60	F0000010	VSX	xxsldwi	VSX Shift Left Double by Word Immediate
XX2	60	F0000290	VSX	xxspltw	VSX Splat Word



Appendix B. POWER8 Instruction Summary by Mnemonic

Table B-1. lists the instructions implemented in the POWER8 processor in alphabetical order by mnemonic. See *Table A-1. Category Listing* on page 313 for a description of the categories.

Table B-1. POWER8 Instructions by Mnemonic (Sheet 1 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
XO	31	7C000214	B	add[.]	Add
XO	31	7C000014	B	addc[.]	Add Carrying
XO	31	7C000414	B	addco[.]	Add Carrying and Record OV
XO	31	7C000114	B	adde[.]	Add Extended
XO	31	7C000514	B	addeo[.]	Add Extended and Record OV and Record OV
D	14	38000000	B	addi	Add Immediate
D	12	30000000	B	addic	Add Immediate Carrying
D	13	34000000	B	addic.	Add Immediate Carrying and Record CR0
D	15	3C000000	B	addis	Add Immediate Shifted
XO	31	7C0001D4	B	addme[.]	Add to Minus One Extended
XO	31	7C0005D4	B	addmeo[.]	Add to Minus One Extended and Record OV
XO	31	7C000614	B	addo[.]	Add and Record OV
XO	31	7C000194	B	addze[.]	Add to Zero Extended
XO	31	7C000594	B	addzeo[.]	Add to Zero Extended and Record OV
X	31	7C000038	B	and[.]	AND
X	31	7C000078	B	andc[.]	AND with Complement
D	28	70000000	B	andi.	AND Immediate and Record CR0
D	29	74000000	B	andis.	AND Immediate Shifted and Record CR0
I	18	48000000	B	b[I][a]	Branch
B	16	40000000	B	bc[l][a]	Branch Conditional
XL	19	4C000420	B	bcctr[l]	Branch Conditional to Count Register
VX	4	10000401	V	bcdadd.	Decimal Add Modulo
VX	4	10000441	V	bcdsub.	Decimal Subtract Modulo
XL	19	4C000020	B	bclr[l]	Branch Conditional to Link Register
X	19	4C000460	B	bctar[l]	Branch Conditional to Branch Target Address Register
X	31	7C0001F8	64	bpermd	Bit Permute Doubleword
X	31	7C00035C	S	clrbhrb	Clear BHRB
X	31	7C000000	B	cmp	Compare
X	31	7C0003F8	B	cmpb	Compare Byte
D	11	2C000000	B	cmpl	Compare Immediate
X	31	7C000040	B	cmpl	Compare Logical



Table B-1. POWER8 Instructions by Mnemonic (Sheet 2 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
D	10	28000000	B	cmpli	Compare Logical Immediate
X	31	7C000074	64	cntlzd[.]	Count Leading Zeros Doubleword
X	31	7C000034	B	cntlzw[.]	Count Leading Zeros Word
XL	19	4C000202	B	crand	Condition Register AND
XL	19	4C000102	B	crandc	Condition Register AND with Complement
XL	19	4C000242	B	creqv	Condition Register Equivalent
XL	19	4C0001C2	B	crnand	Condition Register NAND
XL	19	4C000042	B	crnor	Condition Register NOR
XL	19	4C000382	B	cror	Condition Register OR
XL	19	4C000342	B	crorc	Condition Register OR with Complement
XL	19	4C000182	B	crxor	Condition Register XOR
X	59	EC000004	DFP	dadd[.]	Decimal Floating Add
X	63	FC000004	DFP	daddq[.]	Decimal Floating Add Quad
X	31	7C0000AC	B	dcbf	Data Cache Block Flush
X	31	7C00006C	B	dcbst	Data Cache Block Store
X	31	7C00022C	B	dcbt	Data Cache Block Touch
X	31	7C0001EC	B	dcbtst	Data Cache Block Touch for Store
X	31	7C0007EC	B	dcbz	Data Cache Block Zero
X	59	EC000644	DFP	dcffix[.]	Decimal Floating Convert From Fixed
X	63	FC000644	DFP	dcffixq[.]	Decimal Floating Convert From Fixed Quad
X	59	EC000104	DFP	dcmpo	Decimal Floating Compare Ordered
X	63	FC000104	DFP	dcmpoq	Decimal Floating Compare Ordered Quad
X	59	EC000504	DFP	dcmpu	Decimal Floating Compare Unordered
X	63	FC000504	DFP	dcmpuq	Decimal Floating Compare Unordered Quad
X	59	EC000204	DFP	dctdp[.]	Decimal Floating Convert To DFP Long
X	59	EC000244	DFP	dctfix[.]	Decimal Floating Convert To Fixed
X	63	FC000244	DFP	dctfixq[.]	Decimal Floating Convert To Fixed Quad
X	63	FC000204	DFP	dctqpq[.]	Decimal Floating Convert To DFP Extended
X	59	EC000284	DFP	ddedpd[.]	Decimal Floating Decode DPD To BCD
X	63	FC000284	DFP	ddedpdq[.]	Decimal Floating Decode DPD To BCD Quad
X	59	EC000444	DFP	ddiv[.]	Decimal Floating Divide
X	63	FC000444	DFP	ddivq[.]	Decimal Floating Divide Quad
X	59	EC000684	DFP	denbcd[.]	Decimal Floating Encode BCD To DPD
X	63	FC000684	DFP	denbcdq[.]	Decimal Floating Encode BCD To DPD Quad



Table B-1. POWER8 Instructions by Mnemonic (Sheet 3 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
X	59	EC0006C4	DFP	diex[.]	Decimal Floating Insert Exponent
X	63	FC0006C4	DFP	diexq[.]	Decimal Floating Insert Exponent Quad
XO	31	7C0003D2	64	divd[.]	Divide Doubleword
XO	31	7C000352	64	divde[.]	Divide Doubleword Extended
XO	31	7C000752	64	divdeo[.]	Divide Doubleword Extended and Record OV
XO	31	7C000312	64	divdeu[.]	Divide Doubleword Extended Unsigned
XO	31	7C000712	64	divdeuo[.]	Divide Doubleword Extended Unsigned and Record OV
XO	31	7C0007D2	64	divdo[.]	Divide Doubleword and Record OV
XO	31	7C000392	64	divdu[.]	Divide Doubleword Unsigned
XO	31	7C000792	64	divduo[.]	Divide Doubleword Unsigned and Record OV
XO	31	7C0003D6	B	divw[.]	Divide Word
XO	31	7C000356	B	divwe[.]	Divide Word Extended
XO	31	7C000756	B	divweo[.]	Divide Word Extended and Record OV
XO	31	7C000316	B	divweu[.]	Divide Word Extended Unsigned
XO	31	7C000716	B	divweuo[.]	Divide Word Extended Unsigned and Record OV
XO	31	7C0007D6	B	divwo[.]	Divide Word and Record OV
XO	31	7C000396	B	divwu[.]	Divide Word Unsigned
XO	31	7C000796	B	divwuo[.]	Divide Word Unsigned and Record OV
X	59	EC000044	DFP	dmul[.]	Decimal Floating Multiply
X	63	FC000044	DFP	dmulq[.]	Decimal Floating Multiply Quad
XL	19	4C000324	S	doze	Doze
Z23	59	EC000006	DFP	dqua[.]	Decimal Quantize
Z23	59	EC000086	DFP	dquai[.]	Decimal Quantize Immediate
Z23	63	FC000086	DFP	dquaiq[.]	Decimal Quantize Immediate Quad
Z23	63	FC000006	DFP	dquaq[.]	Decimal Quantize Quad
X	63	FC000604	DFP	drdpq[.]	Decimal Floating Round To DFP Long
Z23	59	EC0001C6	DFP	drintn[.]	Decimal Floating Round To FP Integer Without Inexact
Z23	63	FC0001C6	DFP	drtnq[.]	Decimal Floating Round To FP Integer Without Inexact Quad
Z23	59	EC0000C6	DFP	drintx[.]	Decimal Floating Round To FP Integer With Inexact
Z23	63	FC0000C6	DFP	drtnql[.]	Decimal Floating Round To FP Integer With Inexact Quad
Z23	59	EC000046	DFP	drrnd[.]	Decimal Floating Reround
Z23	63	FC000046	DFP	drrndq[.]	Decimal Floating Reround Quad
X	59	EC000604	DFP	drsp[.]	Decimal Floating Round To DFP Short
Z22	59	EC000084	DFP	dscli[.]	Decimal Floating Shift Coefficient Left Immediate



Table B-1. POWER8 Instructions by Mnemonic (Sheet 4 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
Z22	63	FC000084	DFP	dscliq[.]	Decimal Floating Shift Coefficient Left Immediate Quad
Z22	59	EC0000C4	DFP	dscrif[.]	Decimal Floating Shift Coefficient Right Immediate
Z22	63	FC0000C4	DFP	dscrq[.]	Decimal Floating Shift Coefficient Right Immediate Quad
X	59	EC000404	DFP	dsub[.]	Decimal Floating Subtract
X	63	FC000404	DFP	dsubq[.]	Decimal Floating Subtract Quad
Z22	59	EC000184	DFP	dtstdc	Decimal Floating Test Data Class
Z22	63	FC000184	DFP	dtstdcq	Decimal Floating Test Data Class Quad
Z22	59	EC0001C4	DFP	dtstdg	Decimal Floating Test Data Group
Z22	63	FC0001C4	DFP	dtstdgq	Decimal Floating Test Data Group Quad
X	59	EC000144	DFP	dtstex	Decimal Floating Test Exponent
X	63	FC000144	DFP	dtstexq	Decimal Floating Test Exponent Quad
X	59	EC000544	DFP	dtstsf	Decimal Floating Test Significance
X	63	FC000544	DFP	dtstfq	Decimal Floating Test Significance Quad
X	59	EC0002C4	DFP	dxex[.]	Decimal Floating Extract Exponent
X	63	FC0002C4	DFP	dxeq[.]	Decimal Floating Extract Exponent Quad
X	31	7C0006AC	S	eieio	Enforce In-order Execution of I/O
X	31	7C000238	B	eqv[.]	Equivalent
X	31	7C000774	B	extsb[.]	Extend Sign Byte
X	31	7C000734	B	extsh[.]	Extend Sign Halfword
X	31	7C0007B4	64	extsw[.]	Extend Sign Word
X	63	FC000210	FP[R]	fabs[.]	Floating Absolute Value
A	63	FC00002A	FP[R]	fadd[.]	Floating Add
A	59	EC00002A	FP[R]	fadds[.]	Floating Add Single
X	63	FC00069C	FP[R]	fcfid[.]	Floating Convert From Integer Doubleword
X	59	EC00069C	FP[R]	fcfid[s.]	Floating Convert From Integer Doubleword Single
X	63	FC00079C	FP[R]	fcfidu[.]	Floating Convert From Integer Doubleword Unsigned
X	59	EC00079C	FP[R]	fcfidus[.]	Floating Convert From Integer Doubleword Unsigned Single
X	63	FC000040	FP	fcmwo	Floating Compare Ordered
X	63	FC000000	FP	fcmwu	Floating Compare Unordered
X	63	FC000010	FP[R]	fcpsgn[.]	Floating Copy Sign
X	63	FC00065C	FP[R]	fctid[.]	Floating Convert To Integer Doubleword
X	63	FC00075C	FP[R]	fctidu[.]	Floating Convert To Integer Doubleword Unsigned
X	63	FC00075E	FP[R]	fctiduz[.]	Floating Convert To Integer Doubleword Unsigned with Round Toward Zero



Table B-1. POWER8 Instructions by Mnemonic (Sheet 5 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
X	63	FC00065E	FP[R]	fctidz[.]	Floating Convert To Integer Doubleword with Round Toward Zero
X	63	FC00001C	FP[R]	fctiw[.]	Floating Convert To Integer Word
X	63	FC00011C	FP[R]	fctiwu[.]	Floating Convert To Integer Word Unsigned
X	63	FC00011E	FP[R]	fctiwuz[.]	Floating Convert To Integer Word Unsigned with Round Toward Zero
X	63	FC00001E	FP[R]	fctiwz[.]	Floating Convert To Integer Word with round to Zero
A	63	FC000024	FP[R]	fdiv[.]	Floating Divide
A	59	EC000024	FP[R]	fdivs[.]	Floating Divide Single
A	63	FC00003A	FP[R]	fmadd[.]	Floating Multiply-Add
A	59	EC00003A	FP[R]	fmadds[.]	Floating Multiply-Add Single
X	63	FC000090	FP[R]	fmr[.]	Floating Move Register
X	63	FC00078C	VSX	fmrgew	Floating Merge Even Word
X	63	FC00068C	VSX	fmrgow	Floating Merge Odd Word
A	63	FC000038	FP[R]	fmsub[.]	Floating Multiply-Subtract
A	59	EC000038	FP[R]	fmsubs[.]	Floating Multiply-Subtract Single
A	63	FC000032	FP[R]	fmul[.]	Floating Multiply
A	59	EC000032	FP[R]	fmuls[.]	Floating Multiply Single
X	63	FC000110	FP[R]	fnabs[.]	Floating Negative Absolute Value
X	63	FC000050	FP[R]	fneg[.]	Floating Negate
A	63	FC00003E	FP[R]	fnmadd[.]	Floating Negative Multiply-Add
A	59	EC00003E	FP[R]	fnmadds[.]	Floating Negative Multiply-Add Single
A	63	FC00003C	FP[R]	fnmsub[.]	Floating Negative Multiply-Subtract
A	59	EC00003C	FP[R]	fnmsubs[.]	Floating Negative Multiply-Subtract Single
A	63	FC000030	FP[R].in	fre[.]	Floating Reciprocal Estimate
A	59	EC000030	FP[R]	fres[.]	Floating Reciprocal Estimate Single
X	63	FC0003D0	FP[R].in	frim[.]	Floating Round To Integer Minus
X	63	FC000310	FP[R].in	frin[.]	Floating Round To Integer Nearest
X	63	FC000390	FP[R].in	frip[.]	Floating Round To Integer Plus
X	63	FC000350	FP[R].in	friz[.]	Floating Round To Integer toward Zero
X	63	FC000018	FP[R]	frsp[.]	Floating Round to Single-Precision
A	63	FC000034	FP[R]	frsqte[.]	Floating Reciprocal Square Root Estimate
A	59	EC000034	FP[R].in	frsqtes[.]	Floating Reciprocal Square Root Estimate Single
A	63	FC00002E	FP[R]	fsel[.]	Floating Select
A	63	FC00002C	FP[R]	fsqrt[.]	Floating Square Root



Table B-1. POWER8 Instructions by Mnemonic (Sheet 6 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
A	59	EC00002C	FP[R]	fqrts[.]	Floating Square Root Single
A	63	FC000028	FP[R]	fsub[.]	Floating Subtract
A	59	EC000028	FP[R]	fsubs[.]	Floating Subtract Single
X	63	FC000100	FP	ftdiv	Floating Test for Software Divide
X	63	FC000140	FP	ftsqrt	Floating Test for Software Square Root
XL	19	4C000224	S	hrfid	Return From Interrupt Doubleword Hypervisor
X	31	7C0007AC	B	icbi	Instruction Cache Block Invalidate
X	31	7C00002C	B	icbt	Instruction Cache Block Touch
A	31	7C00001E	B	isel	Integer Select
XL	19	4C00012C	B	isync	Instruction Synchronize
X	31	7C000068	B	lbarx	Load Byte And Reserve Indexed
D	34	88000000	B	lbz	Load Byte and Zero
X	31	7C0006AA	S	lbzcx	Load Byte and Zero Caching Inhibited Indexed
D	35	8C000000	B	lbzu	Load Byte and Zero with Update
X	31	7C0000EE	B	lbzux	Load Byte and Zero with Update Indexed
X	31	7C0000AE	B	lbzx	Load Byte and Zero Indexed
DS	58	E8000000	64	ld	Load Doubleword
X	31	7C0000A8	64	ldarx	Load Doubleword And Reserve Indexed
X	31	7C000428	64	ldbrx	Load Doubleword Byte-Reverse Indexed
X	31	7C0006EA	S	ldcix	Load Doubleword Caching Inhibited Indexed
X	31	7C00003A	64	ldepx	Load Doubleword by External PID Indexed
DS	58	E8000001	64	ldu	Load Doubleword with Update
X	31	7C00006A	64	ldux	Load Doubleword with Update Indexed
X	31	7C00002A	64	idx	Load Doubleword Indexed
D	50	C8000000	FP	lfd	Load Floating-Point Double
DS	57	E4000000	FP.out	lfdp	Load Floating-Point Double Pair
X	31	7C00062E	FP.out	lfdpx	Load Floating-Point Double Pair Indexed
D	51	CC000000	FP	lfdw	Load Floating-Point Double with Update
X	31	7C0004EE	FP	lfdwx	Load Floating-Point Double with Update Indexed
X	31	7C0004AE	FP	lfdx	Load Floating-Point Double Indexed
X	31	7C0006AE	FP	lfiwax	Load Floating-Point as Integer Word Algebraic Indexed
X	31	7C0006EE	FP	lfiwzx	Load Floating-Point as Integer Word and Zero Indexed
D	48	C0000000	FP	lfs	Load Floating-Point Single
D	49	C4000000	FP	lfsu	Load Floating-Point Single with Update



Table B-1. POWER8 Instructions by Mnemonic (Sheet 7 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
X	31	7C00046E	FP	lfsux	Load Floating-Point Single with Update Indexed
X	31	7C00042E	FP	lfsx	Load Floating-Point Single Indexed
D	42	A8000000	B	lha	Load Halfword Algebraic
X	31	7C0000E8	B	lhaxr	Load Halfword And Reserve Indexed Xform
D	43	AC000000	B	lhau	Load Halfword Algebraic with Update
X	31	7C0002EE	B	lhaux	Load Halfword Algebraic with Update Indexed
X	31	7C0002AE	B	lhax	Load Halfword Algebraic Indexed
X	31	7C00062C	B	lhbrx	Load Halfword Byte-Reverse Indexed
D	40	A0000000	B	lhz	Load Halfword and Zero
X	31	7C00066A	S	lhzcix	Load Halfword and Zero Caching Inhibited Indexed
D	41	A4000000	B	lhzu	Load Halfword and Zero with Update
X	31	7C00026E	B	lhzux	Load Halfword and Zero with Update Indexed
X	31	7C00022E	B	lhzx	Load Halfword and Zero Indexed
D	46	B8000000	B	lmw	Load Multiple Word
DQ	56	E0000000	LSQ	lq	Load Quadword
X	31	7C000228	LSQ	lqarx	Load Quadword And Reserve Indexed
X	31	7C0004AA	MA	lswi	Load String Word Immediate
X	31	7C00042A	MA	lswx	Load String Word Indexed
X	31	7C00000E	V	lvebx	Load Vector Element Byte Indexed
X	31	7C00004E	V	lvehx	Load Vector Element Halfword Indexed
X	31	7C00008E	V	lvewx	Load Vector Element Word Indexed
X	31	7C00000C	V	lvsl	Load Vector for Shift Left
X	31	7C00004C	V	lvsr	Load Vector for Shift Right
X	31	7C0000CE	V	lvx	Load Vector Indexed
X	31	7C0002CE	V	lvxl	Load Vector Indexed Last
DS	58	E8000002	64	lwa	Load Word Algebraic
X	31	7C000028	B	lwaxr	Load Word and Reserve Indexed
X	31	7C0002EA	64	lwaux	Load Word Algebraic with Update Indexed
X	31	7C0002AA	64	lwax	Load Word Algebraic Indexed
X	31	7C00042C	B	lwbrx	Load Word Byte-Reverse Indexed
D	32	80000000	B	lwz	Load Word and Zero
X	31	7C00062A	S	lwzcix	Load Word and Zero Caching Inhibited Indexed
D	33	84000000	B	lwzu	Load Word and Zero with Update
X	31	7C00006E	B	lwzux	Load Word and Zero with Update Indexed



Table B-1. POWER8 Instructions by Mnemonic (Sheet 8 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
X	31	7C00002E	B	lwzx	Load Word and Zero Indexed
XX1	31	7C000498	VSX	lxsdx	Load VSR Scalar Doubleword Indexed
XX1	31	7C000098	VSX	lxsiwax	Load VSX Scalar as Integer Word Algebraic Indexed
XX1	31	7C000018	VSX	lxsiwzx	Load VSX Scalar as Integer Word and Zero Indexed
XX1	31	7C000418	VSX	lxsspx	Load VSX Scalar Single-Precision Indexed
XX1	31	7C000698	VSX	lxvd2x	Load VSR Vector Doubleword*2 Indexed
XX1	31	7C000298	VSX	lxvdsx	Load VSR Vector Doubleword and Splat Indexed
XX1	31	7C000618	VSX	lxvw4x	Load VSR Vector Word*4 Indexed
XL	19	4C000000	B	mcrf	Move Condition Register Field
X	63	FC000080	FP	mcrfs	Move To Condition Register from FPSCR
XFX	31	7C00025C	S	mfblrbe	Move From Branch History Rolling Buffer
XFX	31	7C000026	B	mfcr	Move From Condition Register
X	63	FC00048E	FP[R]	mfefs[.]	Move From FPSCR
X	31	7C0000A6	S	mfmsr	Move From Machine State Register
XFX	31	7C100026	B	mfocrf	Move From One Condition Register Field
XFX	31	7C0002A6	B	mfsp	Move From Special Purpose Register
X	31	7C0004A6	S	mfsr	Move From Segment Register
X	31	7C000526	S	mfsrcin	Move From Segment Register Indirect
XFX	31	7C0002E6	S.out	mftb	Move From Time Base
VX	4	10000604	V	mfvscr	Move From Vector Status and Control Register
XX1	31	7C000066	VSX	mfvsrd	Move From VSR Doubleword
XX1	31	7C0000E6	VSX	mfvsrwz	Move From VSR Word and Zero
X	31	7C0001DC	S	msgclr	Message Clear
X	31	7C00015C	S	msgclrp	Message Clear Privileged
X	31	7C00019C	S	msgsnd	Message Send
X	31	7C00011C	S	msgsndp	Message Send Privileged
XFX	31	7C000120	B	mtcrcf	Move To Condition Register Fields
X	63	FC00008C	FP[R]	mtfsb0[.]	Move To FPSCR Bit 0
X	63	FC00004C	FP[R]	mtfsb1[.]	Move To FPSCR Bit 1
XFL	63	FC00058E	FP[R]	mtfsf[.]	Move To FPSCR Fields
X	63	FC00010C	FP[R]	mtfsfi[.]	Move To FPSCR Field Immediate
X	31	7C000124	S	mtmsr	Move To Machine State Register
X	31	7C000164	S	mtmsrd	Move To Machine State Register Doubleword
XFX	31	7C100120	B	mtocrf	Move To One Condition Register Field



Table B-1. POWER8 Instructions by Mnemonic (Sheet 9 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
XFX	31	7C0003A6	B	mtspr	Move To Special Purpose Register
X	31	7C0001A4	S	mtsr	Move To Segment Register
X	31	7C0001E4	S	mtsrin	Move To Segment Register Indirect
VX	4	10000644	V	mtvscr	Move To Vector Status and Control Register
XX1	31	7C000166	VSX	mtvsrd	Move To VSR Doubleword
XX1	31	7C0001A6	VSX	mtvsrwa	Move To VSR Word Algebraic
XX1	31	7C0001E6	VSX	mtvsrwz	Move To VSR Word and Zero
XO	31	7C000092	64	mulhd[.]	Multiply High Doubleword
XO	31	7C000012	64	mulhdu[.]	Multiply High Doubleword Unsigned
XO	31	7C000096	B	mulhw[.]	Multiply High Word
XO	31	7C000016	B	mulhwu[.]	Multiply High Word Unsigned
XO	31	7C0001D2	64	mulld[.]	Multiply Low Doubleword
XO	31	7C0005D2	64	mulldo[.]	Multiply Low Doubleword and Record OV
D	7	1C000000	B	mulli	Multiply Low Immediate
XO	31	7C0001D6	B	mullw[.]	Multiply Low Word
XO	31	7C0005D6	B	mullwo[.]	Multiply Low Word and Record OV
X	31	7C0003B8	B	nand[.]	NAND
XL	19	4C000364	S	nap	Nap
XO	31	7C0000D0	B	neg[.]	Negate
XO	31	7C0004D0	B	nego[.]	Negate and Record OV
X	31	7C0000F8	B	nor[.]	NOR
X	31	7C000378	B	or[.]	OR
X	31	7C000338	B	orc[.]	OR with Complement
D	24	60000000	B	ori	OR Immediate
D	25	64000000	B	oris	OR Immediate Shifted
X	31	7C0000F4	B	popcntb	Population Count Byte-wise
X	31	7C0003F4	64	popcntd	Population Count Doubleword
X	31	7C0002F4	B	popcntw	Population Count Words
X	31	7C000174	64	ptryd	Parity Doubleword
X	31	7C000134	B	ptryw	Parity Word
XL	19	4C000124	S	rfebb	Return from Event Based Branch
XL	19	4C000024	S	rfdi	Return from Interrupt Doubleword
MDS	30	78000010	64	rldcl[.]	Rotate Left Doubleword then Clear Left
MDS	30	78000012	64	rldcr[.]	Rotate Left Doubleword then Clear Right



Table B-1. POWER8 Instructions by Mnemonic (Sheet 10 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
MD	30	78000008	64	rldic[.]	Rotate Left Doubleword Immediate then Clear
MD	30	78000000	64	rldicl[.]	Rotate Left Doubleword Immediate then Clear Left
MD	30	78000004	64	rldicr[.]	Rotate Left Doubleword Immediate then Clear Right
MD	30	7800000C	64	rldimi[.]	Rotate Left Doubleword Immediate then Mask Insert
M	20	50000000	B	rlwimi[.]	Rotate Left Word Immediate then Mask Insert
M	21	54000000	B	rlwinm[.]	Rotate Left Word Immediate then AND with Mask
M	23	5C000000	B	rlwnm[.]	Rotate Left Word then AND with Mask
XL	19	4C0003E4	S	rvwinkle	Rip Van Winkle
SC	17	44000002	B	sc	System Call
X	31	7C0007A7	S	slbfee.	SLB Find Entry ESID
X	31	7C0003E4	S	slbia	SLB Invalidate All
X	31	7C000364	S	slbie	SLB Invalidate Entry
X	31	7C000726	S	slbmfee	SLB Move From Entry ESID
X	31	7C0006A6	S	slbmfev	SLB Move From Entry VSID
X	31	7C000324	S	slbmte	SLB Move To Entry
X	31	7C000036	64	sld[.]	Shift Left Doubleword
XL	19	4C0003A4	S	sleep	Sleep
X	31	7C000030	B	slw[.]	Shift Left Word
X	31	7C000634	64	srad[.]	Shift Right Algebraic Doubleword
XS	31	7C000674	64	sradi[.]	Shift Right Algebraic Doubleword Immediate
X	31	7C000630	B	sraw[.]	Shift Right Algebraic Word
X	31	7C000670	B	srawi[.]	Shift Right Algebraic Word Immediate
X	31	7C000436	64	srd[.]	Shift Right Doubleword
X	31	7C000430	B	srw[.]	Shift Right Word
D	38	98000000	B	stb	Store Byte
X	31	7C0007AA	S	stbcix	Store Byte Caching Inhibited Indexed
X	31	7C00056D	B	stbcx.	Store Byte Conditional Indexed
D	39	9C000000	B	stbu	Store Byte with Update
X	31	7C0001EE	B	stbux	Store Byte with Update Indexed
X	31	7C0001AE	B	stbx	Store Byte Indexed
DS	62	F8000000	64	std	Store Doubleword
X	31	7C000528	64	stdbrx	Store Doubleword Byte-Reverse Indexed
X	31	7C0007EA	S	stdcix	Store Doubleword Caching Inhibited Indexed
X	31	7C0001AD	64	stdcx.	Store Doubleword Conditional Indexed and Record CR0



Table B-1. POWER8 Instructions by Mnemonic (Sheet 11 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
X	31	7C00013A	64	stdpx	Store Doubleword by External PID Indexed
DS	62	F8000001	64	stdu	Store Doubleword with Update
X	31	7C00016A	64	stdux	Store Doubleword with Update Indexed
X	31	7C00012A	64	stdx	Store Doubleword Indexed
D	54	D8000000	FP	stfd	Store Floating-Point Double
DS	61	F4000000	FP.out	stfdp	Store Floating-Point Double Pair
X	31	7C00072E	FP.out	stfdpx	Store Floating-Point Double Pair Indexed
D	55	DC000000	FP	stfdw	Store Floating-Point Double with Update
X	31	7C0005EE	FP	stfdux	Store Floating-Point Double with Update Indexed
X	31	7C0005AE	FP	stfdx	Store Floating-Point Double Indexed
X	31	7C0007AE	FP	stfiwx	Store Floating-Point as Integer Word Indexed
D	52	D0000000	FP	stfs	Store Floating-Point Single
D	53	D4000000	FP	stfsu	Store Floating-Point Single with Update
X	31	7C00056E	FP	stfsux	Store Floating-Point Single with Update Indexed
X	31	7C00052E	FP	stfsx	Store Floating-Point Single Indexed
D	44	B0000000	B	sth	Store Halfword
X	31	7C00072C	B	sthbrx	Store Halfword Byte-Reverse Indexed
X	31	7C00076A	S	sthcix	Store Halfword and Zero Caching Inhibited Indexed
X	31	7C0005AD	B	sthcx.	Store Halfword Conditional Indexed Xform
D	45	B4000000	B	sthu	Store Halfword with Update
X	31	7C00036E	B	sthux	Store Halfword with Update Indexed
X	31	7C00032E	B	sthx	Store Halfword Indexed
D	47	BC000000	B	stmw	Store Multiple Word
DS	62	F8000002	LSQ	stq	Store Quadword
X	31	7C00016D	LSQ	stqcx.	Store Quadword Conditional Indexed and record CR0
X	31	7C0005AA	MA	stswi	Store String Word Immediate
X	31	7C00052A	MA	stswx	Store String Word Indexed
X	31	7C00010E	V	stvebx	Store Vector Element Byte Indexed
X	31	7C00014E	V	stvehx	Store Vector Element Halfword Indexed
X	31	7C00018E	V	stvewx	Store Vector Element Word Indexed
X	31	7C0001CE	V	stvx	Store Vector Indexed
X	31	7C0003CE	V	stvxl	Store Vector Indexed Last
D	36	90000000	B	stw	Store Word
X	31	7C00052C	B	stwbrx	Store Word Byte-Reverse Indexed



Table B-1. POWER8 Instructions by Mnemonic (Sheet 12 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
X	31	7C00072A	S	stwcix	Store Word and Zero Caching Inhibited Indexed
X	31	7C00012D	B	stwcx.	Store Word Conditional Indexed and Record CR0
D	37	94000000	B	stwu	Store Word with Update
X	31	7C00016E	B	stwux	Store Word with Update Indexed
X	31	7C00012E	B	stwx	Store Word Indexed
XX1	31	7C000598	VSX	stxsdx	Store VSR Scalar Doubleword Indexed
XX1	31	7C000118	VSX	stxsiwx	Store VSX Scalar as Integer Word Indexed
XX1	31	7C000518	VSX	stxsspx	Store VSR Scalar Word Indexed
XX1	31	7C000798	VSX	stxvd2x	Store VSR Vector Doubleword*2 Indexed
XX1	31	7C000718	VSX	stxvw4x	Store VSR Vector Word*4 Indexed
XO	31	7C000050	B	subf[.]	Subtract From
XO	31	7C000010	B	subfc[.]	Subtract From Carrying
XO	31	7C000410	B	subfc0[.]	Subtract From Carrying and Record OV
XO	31	7C000110	B	subfe[.]	Subtract From Extended
XO	31	7C000510	B	subfeo[.]	Subtract From Extended and Record OV
D	8	20000000	B	subfic	Subtract From Immediate Carrying
XO	31	7C0001D0	B	subfme[.]	Subtract From Minus One Extended
XO	31	7C0005D0	B	subfmeo[.]	Subtract From Minus One Extended and Record OV
XO	31	7C000450	B	subfo[.]	Subtract From and Record OV
XO	31	7C000190	B	subfze[.]	Subtract From Zero Extended
XO	31	7C000590	B	subfzeo[.]	Subtract From Zero Extended and Record OV
X	31	7C0004AC	B	sync	Synchronize
X	31	7C00071D	TM	tabort.	Transaction Abort
X	31	7C00065D	TM	tabortdc.	Transaction Abort Doubleword Conditional
X	31	7C0006DD	TM	tabortdci.	Transaction Abort Doubleword Conditional Immediate
X	31	7C00061D	TM	tabortwc.	Transaction Abort Word Conditional
X	31	7C00069D	TM	tabortwci.	Transaction Abort Word Conditional Immediate
X	31	7C00051D	TM	tbegin.	Transaction Begin
X	31	7C00059C	TM	tcheck	Transaction Check
X	31	7C000088	64	td	Trap Doubleword
D	2	08000000	64	tdi	Trap Doubleword Immediate
X	31	7C00055C	TM	tend.	Transaction End
X	31	7C000264	S	tlbie	TLB Invalidate Entry
X	31	7C000224	S	tlbiel	TLB Invalidate Entry Local



Table B-1. POWER8 Instructions by Mnemonic (Sheet 13 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
X	31	7C00046C	S	tlbsync	TLB Synchronize
X	31	7C0007DD	TM	trechkpt.	Transaction Recheckpoint
X	31	7C00075D	TM	treclaim.	Transaction Reclaim
X	31	7C000008	B	tw	Trap Word
D	3	0C000000	B	twi	Trap Word Immediate
VX	4	10000140	V	vaddcuq	Vector Add and Write Carry Unsigned Quadword
VX	4	10000180	V	vaddcuw	Vector Add and Write Carry-Out Unsigned Word
VA	4	1000003D	V	vaddecuq	Vector Add Extended and Write Carry Unsigned Quadword
VA	4	1000003C	V	vaddeuqm	Vector Add Extended Unsigned Quadword Modulo
VX	4	1000000A	V	vaddirfp	Vector Add Single-Precision
VX	4	10000300	V	vaddsbs	Vector Add Signed Byte Saturate
VX	4	10000340	V	vaddshs	Vector Add Signed Halfword Saturate
VX	4	10000380	V	vaddsws	Vector Add Signed Word Saturate
VX	4	10000000	V	vaddubm	Vector Add Unsigned Byte Modulo
VX	4	10000200	V	vaddubs	Vector Add Unsigned Byte Saturate
VX	4	100000C0	V	vaddudm	Vector Add Unsigned Doubleword Modulo
VX	4	10000040	V	vadduhm	Vector Add Unsigned Halfword Modulo
VX	4	10000240	V	vadduhs	Vector Add Unsigned Halfword Saturate
VX	4	10000100	V	vadduqm	Vector Add Unsigned Quadword Modulo
VX	4	10000080	V	vadduwrm	Vector Add Unsigned Word Modulo
VX	4	10000280	V	vadduws	Vector Add Unsigned Word Saturate
VX	4	10000404	V	vand	Vector Logical AND
VX	4	10000444	V	vandc	Vector Logical AND with Complement
VX	4	10000502	V	vavgsb	Vector Average Signed Byte
VX	4	10000542	V	vavgsh	Vector Average Signed Halfword
VX	4	10000582	V	vavgsw	Vector Average Signed Word
VX	4	10000402	V	vavgub	Vector Average Unsigned Byte
VX	4	10000442	V	vavguh	Vector Average Unsigned Halfword
VX	4	10000482	V	vavguw	Vector Average Unsigned Word
VX	4	1000054C	V	vbpermq	Vector Bit Permute Quadword
VX	4	1000054C	V	vbpermq	Vector Bit Permute Quadword
VX	4	1000034A	V	vcfsx	Vector Convert From Signed Fixed-Point Word To Single-Precision
VX	4	1000030A	V	vcfux	Vector Convert From Unsigned Fixed-Point Word
VX	4	10000508	V.AES	vcipher	Vector AES Cipher



Table B-1. POWER8 Instructions by Mnemonic (Sheet 14 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
VX	4	10000509	V.AES	vcipherlast	Vector AES Cipher Last
VX	4	10000702	V	vclzb	Vector Count Leading Zeros Byte
VX	4	100007C2	V	vclzd	Vector Count Leading Zeros Doubleword
VX	4	10000742	V	vclzh	Vector Count Leading Zeros Halfword
VX	4	10000782	V	vclzw	Vector Count Leading Zeros Word
VC	4	100003C6	V	vcmpbfp[.]	Vector Compare Bounds Single-Precision
VC	4	100000C6	V	vcmpeqfp[.]	Vector Compare Equal To Single-Precision
VC	4	10000006	V	vcmpequb[.]	Vector Compare Equal To Unsigned Byte
VC	4	100000C7	V	vcmpequd[.]	Vector Compare Equal To Unsigned Doubleword
VC	4	10000046	V	vcmpequh[.]	Vector Compare Equal To Unsigned Halfword
VC	4	10000086	V	vcmpequw[.]	Vector Compare Equal To Unsigned Word
VC	4	100001C6	V	vcmpgefp[.]	Vector Compare Greater Than or Equal To Single-Precision
VC	4	100002C6	V	vcmpgtfp[.]	Vector Compare Greater Than Single-Precision
VC	4	10000306	V	vcmpgtbsb[.]	Vector Compare Greater Than Signed Byte
VC	4	100003C7	V	vcmpgttsd[.]	Vector Compare Greater Than Signed Doubleword
VC	4	10000346	V	vcmpgtsh[.]	Vector Compare Greater Than Signed Halfword
VC	4	10000386	V	vcmpgtsw[.]	Vector Compare Greater Than Signed Word
VC	4	10000206	V	vcmpgtub[.]	Vector Compare Greater Than Unsigned Byte
VC	4	100002C7	V	vcmpgtud[.]	Vector Compare Greater Than Unsigned Doubleword
VC	4	10000246	V	vcmpgtuh[.]	Vector Compare Greater Than Unsigned Halfword
VC	4	10000286	V	vcmpgtuw[.]	Vector Compare Greater Than Unsigned Word
VX	4	100003CA	V	vctsxs	Vector Convert From Single-Precision To Signed Fixed-Point Word Saturate
VX	4	1000038A	V	vctuxs	Vector Convert From Single-Precision To Unsigned Fixed-Point Word Saturate
VX	4	10000684	V	veqv	Vector Equivalence
VX	4	1000018A	V	vexptefp	Vector 2 Raised to the Exponent Estimate Single-Precision
VX	4	1000050C	V	vgbbd	Vector Gather Bits by Byte by Doubleword
VX	4	100001CA	V	vlogefp	Vector Log Base 2 Estimate Single-Precision
VA	4	1000002E	V	vmaddfp	Vector Multiply-Add Single-Precision
VX	4	1000040A	V	vmaxfp	Vector Maximum Single-Precision
VX	4	10000102	V	vmaxsb	Vector Maximum Signed Byte
VX	4	100001C2	V	vmaxsd	Vector Maximum Signed Doubleword
VX	4	10000142	V	vmaxsh	Vector Maximum Signed Halfword
VX	4	10000182	V	vmaxsw	Vector Maximum Signed Word



Table B-1. POWER8 Instructions by Mnemonic (Sheet 15 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
VX	4	10000002	V	vmaxub	Vector Maximum Unsigned Byte
VX	4	100000C2	V	vmaxud	Vector Maximum Unsigned Doubleword
VX	4	10000042	V	vmaxuh	Vector Maximum Unsigned Halfword
VX	4	10000082	V	vmaxuw	Vector Maximum Unsigned Word
VA	4	10000020	V	vmhaddshs	Vector Multiply-High-Add Signed Halfword Saturate
VA	4	10000021	V	vmhraddshs	Vector Multiply-High-Round-Add Signed Halfword Saturate
VX	4	1000044A	V	vminfp	Vector Minimum Single-Precision
VX	4	10000302	V	vminsbt	Vector Minimum Signed Byte
X	4	100003C2	V	vminsd	Vector Minimum Signed Doubleword
VX	4	10000342	V	vminsh	Vector Minimum Signed Halfword
VX	4	10000382	V	vminsw	Vector Minimum Signed Word
VX	4	10000202	V	vminub	Vector Minimum Unsigned Byte
VX	4	100002C2	V	vminud	Vector Minimum Unsigned Doubleword
VX	4	10000242	V	vminuh	Vector Minimum Unsigned Halfword
VX	4	10000282	V	vminuw	Vector Minimum Unsigned Word
VA	4	10000022	V	vmladduhm	Vector Multiply-Low-Add Unsigned Halfword Modulo
VX	4	1000078C	VSX	vmrgew	Vector Merge Even Word
VX	4	1000000C	V	vmrghb	Vector Merge High Byte
VX	4	1000004C	V	vmrghh	Vector Merge High Halfword
VX	4	1000008C	V	vmrghw	Vector Merge High Word
VX	4	1000010C	V	vmrglb	Vector Merge Low Byte
VX	4	1000014C	V	vmrglh	Vector Merge Low Halfword
VX	4	1000018C	V	vmrglw	Vector Merge Low Word
VX	4	1000068C	VSX	vmrgow	Vector Merge Odd Word
VA	4	10000025	V	vmsummbm	Vector Multiply-Sum Mixed Byte Modulo
VA	4	10000028	V	vmsumshm	Vector Multiply-Sum Signed Halfword Modulo
VA	4	10000029	V	vmsumshs	Vector Multiply-Sum Signed Halfword Saturate
VA	4	10000024	V	vmsumubm	Vector Multiply-Sum Unsigned Byte Modulo
VA	4	10000026	V	vmsumuhm	Vector Multiply-Sum Unsigned Halfword Modulo
VA	4	10000027	V	vmsumuhs	Vector Multiply-Sum Unsigned Halfword Saturate
VX	4	10000308	V	vmulesb	Vector Multiply Even Signed Byte
VX	4	10000348	V	vmulesh	Vector Multiply Even Signed Halfword
VX	4	10000388	V	vmulesw	Vector Multiply Even Signed Word
VX	4	10000208	V	vmuleub	Vector Multiply Even Unsigned Byte



Table B-1. POWER8 Instructions by Mnemonic (Sheet 16 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
VX	4	10000248	V	vmuleuh	Vector Multiply Even Unsigned Halfword
VX	4	10000288	V	vmuleuw	Vector Multiply Even Unsigned Word
VX	4	10000108	V	vmulosb	Vector Multiply Odd Signed Byte
VX	4	10000148	V	vmulosh	Vector Multiply Odd Signed Halfword
VX	4	10000188	V	vmulosw	Vector Multiply Odd Signed Word
VX	4	10000008	V	vmuloub	Vector Multiply Odd Unsigned Byte
VX	4	10000048	V	vmulouh	Vector Multiply Odd Unsigned Halfword
VX	4	10000088	V	vmulouw	Vector Multiply Odd Unsigned Word
VX	4	10000089	V	vmuluwm	Vector Multiply Unsigned Word Modulo
VX	4	10000584	V	vnand	Vector NAND
VX	4	10000548	V.AES	vncipher	Vector AES Inverse Cipher
VX	4	10000549	V.AES	vncipherlast	Vector AES Inverse Cipher Last
VA	4	1000002F	V	vnmsubfp	Vector Negative Multiply-Subtract Single-Precision
VX	4	10000504	V	vnor	Vector Logical NOR
VX	4	10000484	V	vor	Vector Logical OR
VX	4	10000544	V	vorc	Vector OR with Complement
VA	4	1000002B	V	vperm	Vector Permute
VA	4	1000002D	V.RAID	vpermxor	Vector Permute and Exclusive-OR
VX	4	1000030E	V	vpkpx	Vector Pack Pixel
VX	4	100005CE	V	vpkdss	Vector Pack Signed Doubleword Signed Saturate
VX	4	1000054E	V	vpkdus	Vector Pack Signed Doubleword Unsigned Saturate
VX	4	1000018E	V	vpkshss	Vector Pack Signed Halfword Signed Saturate
VX	4	1000010E	V	vpkshus	Vector Pack Signed Halfword Unsigned Saturate
VX	4	100001CE	V	vpkswss	Vector Pack Signed Word Signed Saturate
VX	4	1000014E	V	vpkswus	Vector Pack Signed Word Unsigned Saturate
VX	4	1000044E	V	vpkudum	Vector Pack Unsigned Doubleword Unsigned Modulo
VX	4	100004CE	V	vpkudus	Vector Pack Unsigned Doubleword Unsigned Saturate
VX	4	1000000E	V	vpkuhum	Vector Pack Unsigned Halfword Unsigned Modulo
VX	4	1000008E	V	vpkuhus	Vector Pack Unsigned Halfword Unsigned Saturate
VX	4	1000004E	V	vpkuwum	Vector Pack Unsigned Word Unsigned Modulo
VX	4	100000CE	V	vpkuwus	Vector Pack Unsigned Word Unsigned Saturate
VX	4	10000408	V	vpmsumb	Vector Polynomial Multiply-Sum Byte
VX	4	100004C8	V	vpmsumd	Vector Polynomial Multiply-Sum Doubleword
VX	4	10000448	V	vpmsumh	Vector Polynomial Multiply-Sum Halfword



Table B-1. POWER8 Instructions by Mnemonic (Sheet 17 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
VX	4	10000488	V	vpmsumw	Vector Polynomial Multiply-Sum Word
VX	4	10000703	V	vpopcntb	Vector Population Count Byte
VX	4	100007C3	V	vpopcntd	Vector Population Count Doubleword
VX	4	10000743	V	vpopcnth	Vector Population Count Halfword
VX	4	10000783	V	vpopcntw	Vector Population Count Word
VX	4	1000010A	V	vrefp	Vector Reciprocal Estimate Single-Precision
VX	4	100002CA	V	vrfim	Vector Round to Single-Precision Integer toward -Infinity
VX	4	1000020A	V	vrfin	Vector Round to Single-Precision Integer Nearest
VX	4	1000028A	V	vrfip	Vector Round to Single-Precision Integer toward +Infinity
VX	4	1000024A	V	vrfiz	Vector Round to Single-Precision Integer toward Zero
VX	4	10000004	V	vrlb	Vector Rotate Left Byte
VX	4	100000C4	V	vrld	Vector Rotate Left Doubleword
VX	4	10000044	V	vrlh	Vector Rotate Left Halfword
VX	4	10000084	V	vrlw	Vector Rotate Left Word
VX	4	1000014A	V	vrsqrtefp	Vector Reciprocal Square Root Estimate Single-Precision
VX	4	100005C8	V.AES	vsbox	Vector AES S-Box
VA	4	1000002A	V	vsel	Vector Select
VX	4	100006C2	V.SHA2	vshasigmad	Vector SHA-512 Sigma Doubleword
VX	4	10000682	V.SHA2	vshasigmaw	Vector SHA-256 Sigma Word
VX	4	100001C4	V	vsl	Vector Shift Left
VX	4	10000104	V	vslb	Vector Shift Left Byte
VX	4	100005C4	V	vslld	Vector Shift Left Doubleword
VA	4	1000002C	V	vsldoi	Vector Shift Left Double by Octet Immediate
VX	4	10000144	V	vslh	Vector Shift Left Halfword
VX	4	1000040C	V	vslo	Vector Shift Left by Octet
VX	4	10000184	V	vslw	Vector Shift Left Word
VX	4	1000020C	V	vspltb	Vector Splat Byte
VX	4	1000024C	V	vsplth	Vector Splat Halfword
VX	4	1000030C	V	vspltsb	Vector Splat Immediate Signed Byte
VX	4	1000034C	V	vspltsih	Vector Splat Immediate Signed Halfword
VX	4	1000038C	V	vspltw	Vector Splat Immediate Signed Word
VX	4	1000028C	V	vspltw	Vector Splat Word
VX	4	100002C4	V	vsr	Vector Shift Right
VX	4	10000304	V	vsrab	Vector Shift Right Algebraic Byte



Table B-1. POWER8 Instructions by Mnemonic (Sheet 18 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
VX	4	100003C4	V	vsrad	Vector Shift Right Algebraic Doubleword
VX	4	10000344	V	vsrah	Vector Shift Right Algebraic Halfword
VX	4	10000384	V	vsraw	Vector Shift Right Algebraic Word
VX	4	10000204	V	vsrb	Vector Shift Right Byte
VX	4	100006C4	V	vsrd	Vector Shift Right Doubleword
VX	4	10000244	V	vsrh	Vector Shift Right Halfword
VX	4	1000044C	V	vsro	Vector Shift Right by Octet
VX	4	10000284	V	vsrw	Vector Shift Right Word
VX	4	10000540	V	vsubcuq	Vector Subtract and Write Carry Unsigned Quadword
VX	4	10000580	V	vsubcuw	Vector Subtract and Write Carry-Out Unsigned Word
VA	4	1000003F	V	vsubecuq	Vector Subtract Extended and Write Carry Unsigned Quadword
VA	4	1000003E	V	vsubeuqm	Vector Subtract Extended Unsigned Quadword Modulo
VX	4	1000004A	V	vsubfp	Vector Subtract Single-Precision
VX	4	10000700	V	vsubbsbs	Vector Subtract Signed Byte Saturate
VX	4	10000740	V	vsubshs	Vector Subtract Signed Halfword Saturate
VX	4	10000780	V	vsubsws	Vector Subtract Signed Word Saturate
VX	4	10000400	V	vsububm	Vector Subtract Unsigned Byte Modulo
VX	4	10000600	V	vsububs	Vector Subtract Unsigned Byte Saturate
VX	4	100004C0	V	vsubudm	Vector Subtract Unsigned Doubleword Modulo
VX	4	10000440	V	vsubuhm	Vector Subtract Unsigned Halfword Modulo
VX	4	10000640	V	vsubuhs	Vector Subtract Unsigned Halfword Saturate
VX	4	10000500	V	vsubuqm	Vector Subtract Unsigned Quadword Modulo
VX	4	10000480	V	vsubuw	Vector Subtract Unsigned Word Modulo
VX	4	10000680	V	vsubuws	Vector Subtract Unsigned Word Saturate
VX	4	10000688	V	vsum2sws	Vector Sum across Half Signed Word Saturate
VX	4	10000708	V	vsum4sbs	Vector Sum across Quarter Signed Byte Saturate
VX	4	10000648	V	vsum4shs	Vector Sum across Quarter Signed Halfword Saturate
VX	4	10000608	V	vsum4ubs	Vector Sum across Quarter Unsigned Byte Saturate
VX	4	10000788	V	vsumsws	Vector Sum across Signed Word Saturate
VX	4	1000034E	V	vupkhpx	Vector Unpack High Pixel
VX	4	1000020E	V	vupkhsb	Vector Unpack High Signed Byte
VX	4	1000024E	V	vupkhsh	Vector Unpack High Signed Halfword
VX	4	1000064E	V	vupkhw	Vector Unpack High Signed Word
VX	4	100003CE	V	vupklpx	Vector Unpack Low Pixel



Table B-1. POWER8 Instructions by Mnemonic (Sheet 19 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
VX	4	1000028E	V	vupklsb	Vector Unpack Low Signed Byte
VX	4	100002CE	V	vupklsh	Vector Unpack Low Signed Halfword
VX	4	100006CE	V	vupklsw	Vector Unpack Low Signed Word
VX	4	100004C4	V	vxor	Vector Logical XOR
X	26	68000000	B	xnop	Executed No Operation
X	31	7C000278	B	xor[.]	XOR
D	26	68000000	B	xori	XOR Immediate
D	27	6C000000	B	xoris	XOR Immediate Shifted
XX2	60	F0000564	VSX	xsabsdp	VSX Scalar Absolute Value Double-Precision
XX3	60	F0000100	VSX	xsadddp	VSX Scalar Add Double-Precision
XX3	60	F0000000	VSX	xsaddsp	VSX Scalar Add Single-Precision
XX3	60	F0000158	VSX	xscmpodp	VSX Scalar Compare Ordered Double-Precision
XX3	60	F0000118	VSX	xscmpudp	VSX Scalar Compare Unordered Double-Precision
XX3	60	F0000580	VSX	xscpsgndp	VSX Scalar Copy Sign Double-Precision
XX2	60	F0000424	VSX	xscvdpsp	VSX Scalar Convert Double-Precision to Single-Precision
XX2	60	F000042C	VSX	xscvdpspn	VSX Scalar Convert Double-Precision to Single-Precision format Non-signalling
XX2	60	F0000560	VSX	xscvdpsxds	VSX Scalar Convert Double-Precision to Signed Fixed-Point Doubleword Saturate
XX2	60	F0000160	VSX	xscvdpsxws	VSX Scalar Convert Double-Precision to Signed Fixed-Point Word Saturate
XX2	60	F0000520	VSX	xscvdpxuds	VSX Scalar Convert Double-Precision to Unsigned Fixed-Point Doubleword Saturate
XX2	60	F0000120	VSX	xscvdpxwds	VSX Scalar Convert Double-Precision to Unsigned Fixed-Point Word Saturate
XX2	60	F0000524	VSX	xscvspdp	VSX Scalar Convert Single-Precision to Double-Precision (p=1)
XX2	60	F000052C	VSX	xscvspdpn	Scalar Convert Single-Precision to Double-Precision format Non-signalling
XX2	60	F00005E0	VSX	xscvsxddp	VSX Scalar Convert Signed Fixed-Point Doubleword to Double-Precision
XX2	60	F00004E0	VSX	xscvsxdsp	VSX Scalar Convert Signed Fixed-Point Doubleword to Single-Precision
XX2	60	F00005A0	VSX	xscvuxddp	VSX Scalar Convert Unsigned Fixed-Point Doubleword to Double-Precision
XX2	60	F00004A0	VSX	xscvuxdsp	VSX Scalar Convert Unsigned Fixed-Point Doubleword to Single-Precision
XX3	60	F00001C0	VSX	xsdvidp	VSX Scalar Divide Double-Precision
XX3	60	F00000C0	VSX	xsdivsp	VSX Scalar Divide Single-Precision
XX3	60	F0000108	VSX	xsmaddadp	VSX Scalar Multiply-Add Type-A Double-Precision



Table B-1. POWER8 Instructions by Mnemonic (Sheet 20 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
XX3	60	F0000008	VSX	xsmaddasp	VSX Scalar Multiply-Add Type-A Single-Precision
XX3	60	F0000148	VSX	xsmaddmdp	VSX Scalar Multiply-Add Type-M Double-Precision
XX3	60	F0000048	VSX	xsmaddmsp	VSX Scalar Multiply-Add Type-M Single-Precision
XX3	60	F0000500	VSX	xsmaxdp	VSX Scalar Maximum Double-Precision
XX3	60	F0000540	VSX	xsmindp	VSX Scalar Minimum Double-Precision
XX3	60	F0000188	VSX	xsmsubadp	VSX Scalar Multiply-Subtract Type-A Double-Precision
XX3	60	F0000088	VSX	xsmsubasp	VSX Scalar Multiply-Subtract Type-A Single-Precision
XX3	60	F00001C8	VSX	xsmsubmdp	VSX Scalar Multiply-Subtract Type-M Double-Precision
XX3	60	F00000C8	VSX	xsmsubmsp	VSX Scalar Multiply-Subtract Type-M Single-Precision
XX3	60	F0000180	VSX	xsmuldp	VSX Scalar Multiply Double-Precision
XX3	60	F0000080	VSX	xsmulsp	VSX Scalar Multiply Single-Precision
XX2	60	F00005A4	VSX	xsnabsdp	VSX Scalar Negative Absolute Value Double-Precision
XX2	60	F00005E4	VSX	xsnegdp	VSX Scalar Negate Double-Precision
XX3	60	F0000508	VSX	xsnmaddadp	VSX Scalar Negative Multiply-Add Type-A Double-Precision
XX3	60	F0000408	VSX	xsnmaddasp	VSX Scalar Negative Multiply-Add Type-A Single-Precision
XX3	60	F0000548	VSX	xsnmaddmdp	VSX Scalar Negative Multiply-Add Type-M Double-Precision
XX3	60	F0000448	VSX	xsnmaddmsp	VSX Scalar Negative Multiply-Add Type-M Single-Precision
XX3	60	F0000588	VSX	xsnmsubadp	VSX Scalar Negative Multiply-Subtract Type-A Double-Precision
XX3	60	F0000488	VSX	xsnmsubasp	VSX Scalar Negative Multiply-Subtract Type-A Single-Precision
XX3	60	F00005C8	VSX	xsnmsubmdp	VSX Scalar Negative Multiply-Subtract Type-M Double-Precision
XX3	60	F00004C8	VSX	xsnmsubmsp	VSX Scalar Negative Multiply-Subtract Type-M Single-Precision
XX2	60	F0000124	VSX	xsrdpi	VSX Scalar Round to Double-Precision Integer
XX2	60	F00001AC	VSX	xsrpic	VSX Scalar Round to Double-Precision Integer using Current Rounding Mode
XX2	60	F00001E4	VSX	xsrpim	VSX Scalar Round to Double-Precision Integer toward -Infinity
XX2	60	F00001A4	VSX	xsrpip	VSX Scalar Round to Double-Precision Integer toward +Infinity
XX2	60	F0000164	VSX	xsrpiz	VSX Scalar Round to Double-Precision Integer toward Zero
XX1	60	F0000168	VSX	xsredp	VSX Scalar Reciprocal Estimate Double-Precision
XX2	60	F0000068	VSX	xsresp	VSX Scalar Reciprocal Estimate Single-Precision
XX2	60	F0000464	VSX	xrsp	VSX Scalar Round to Single-Precision
XX2	60	F0000128	VSX	xrsqrtdp	VSX Scalar Reciprocal Square Root Estimate Double-Precision
XX2	60	F0000028	VSX	xrsqrtesp	VSX Scalar Reciprocal Square Root Estimate Single-Precision
XX2	60	F000012C	VSX	xssqrtdp	VSX Scalar Square Root Double-Precision
XX2	60	F000002C	VSX	xssqrtsp	VSX Scalar Square Root Single-Precision



Table B-1. POWER8 Instructions by Mnemonic (Sheet 21 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
XX3	60	F0000140	VSX	xssubdp	VSX Scalar Subtract Double-Precision
XX3	60	F0000040	VSX	xssubsp	VSX Scalar Subtract Single-Precision
XX3	60	F00001E8	VSX	xstdivdp	VSX Scalar Test for Software Divide Double-Precision
XX2	60	F00001A8	VSX	xstsqrtdp	VSX Scalar Test for Software Square Root Double-Precision
XX2	60	F0000764	VSX	xvabsdp	VSX Vector Absolute Value Double-Precision
XX2	60	F0000664	VSX	xvabssp	VSX Vector Absolute Value Single-Precision
XX3	60	F0000300	VSX	xvaddirp	VSX Vector Add Double-Precision
XX3	60	F0000200	VSX	xvaddsp	VSX Vector Add Single-Precision
XX3	60	F0000318	VSX	xvcmeqdp	VSX Vector Compare Equal To Double-Precision
XX3	60	F0000718	VSX	xvcmeqdp.	VSX Vector Compare Equal To Double-Precision and Record CR6
XX3	60	F0000218	VSX	xvcmeqsp	VSX Vector Compare Equal To Single-Precision
XX3	60	F0000618	VSX	xvcmeqsp.	VSX Vector Compare Equal To Single-Precision and Record CR6
XX3	60	F0000398	VSX	xvcmpgedp	VSX Vector Compare Greater Than or Equal To Double-Precision
XX3	60	F0000798	VSX	xvcmpgedp.	VSX Vector Compare Greater Than or Equal To Double-Precision and Record CR6
XX3	60	F0000298	VSX	xvcmpgesp	VSX Vector Compare Greater Than or Equal To Single-Precision
XX3	60	F0000698	VSX	xvcmpgesp.	VSX Vector Compare Greater Than or Equal To Single-Precision and Record CR6
XX3	60	F0000358	VSX	xvcmpgtdp	VSX Vector Compare Greater Than Double-Precision
XX3	60	F0000758	VSX	xvcmpgtdp.	VSX Vector Compare Greater Than Double-Precision and Record CR6
XX3	60	F0000258	VSX	xvcmpgtsp	VSX Vector Compare Greater Than Single-Precision
XX3	60	F0000658	VSX	xvcmpgtsp.	VSX Vector Compare Greater Than Single-Precision and Record CR6
XX3	60	F0000780	VSX	xvcpsgndp	VSX Vector Copy Sign Double-Precision
XX3	60	F0000680	VSX	xvcpsgnsp	VSX Vector Copy Sign Single-Precision
XX2	60	F0000624	VSX	xvcvdpsp	VSX Vector Convert Double-Precision to Single-Precision
XX2	60	F0000760	VSX	xvcvdpsxds	VSX Vector Convert Double-Precision to Signed Fixed-Point Doubleword Saturate
XX2	60	F0000360	VSX	xvcvdpsxws	VSX Vector Convert Double-Precision to Signed Fixed-Point Word Saturate
XX2	60	F0000720	VSX	xvcvdpxuds	VSX Vector Convert Double-Precision to Unsigned Fixed-Point Doubleword Saturate
XX2	60	F0000320	VSX	xvcvdpxws	VSX Vector Convert Double-Precision to Unsigned Fixed-Point Word Saturate
XX2	60	F0000724	VSX	xvcvspdp	VSX Vector Convert Single-Precision to Double-Precision



Table B-1. POWER8 Instructions by Mnemonic (Sheet 22 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
XX2	60	F0000660	VSX	xvcvpsxds	VSX Vector Convert Single-Precision to Signed Fixed-Point Doubleword Saturate
XX2	60	F0000260	VSX	xvcvpsxws	VSX Vector Convert Single-Precision to Signed Fixed-Point Word Saturate
XX2	60	F0000620	VSX	xvcvspuxds	VSX Vector Convert Single-Precision to Unsigned Fixed-Point Doubleword Saturate
XX2	60	F0000220	VSX	xvcvspuxws	VSX Vector Convert Single-Precision to Unsigned Fixed-Point Word Saturate
XX2	60	F00007E0	VSX	xvcvsxddp	VSX Vector Convert Signed Fixed-Point Doubleword to Double-Precision
XX2	60	F00006E0	VSX	xvcvsxdsp	VSX Vector Convert Signed Fixed-Point Doubleword to Single-Precision
XX2	60	F00003E0	VSX	xvcvsxwdp	VSX Vector Convert Signed Fixed-Point Word to Double-Precision
XX2	60	F00002E0	VSX	xvcvsxwsp	VSX Vector Convert Signed Fixed-Point Word to Single-Precision
XX2	60	F00007A0	VSX	xvcvuxddp	VSX Vector Convert Unsigned Fixed-Point Doubleword to Double-Precision
XX2	60	F00006A0	VSX	xvcvuxdsp	VSX Vector Convert Unsigned Fixed-Point Doubleword to Single-Precision
XX2	60	F00003A0	VSX	xvcvuxwdp	VSX Vector Convert Unsigned Fixed-Point Word to Double-Precision
XX2	60	F00002A0	VSX	xvcvuxwsp	VSX Vector Convert Unsigned Fixed-Point Word to Single-Precision
XX3	60	F00003C0	VSX	xvdivdp	VSX Vector Divide Double-Precision
XX3	60	F00002C0	VSX	xvdivsp	VSX Vector Divide Single-Precision
XX3	60	F0000308	VSX	xvmaddadp	VSX Vector Multiply-Add Type-A Double-Precision
XX3	60	F0000208	VSX	xvmaddasp	VSX Vector Multiply-Add Type-A Single-Precision
XX3	60	F0000348	VSX	xvmaddmdp	VSX Vector Multiply-Add Type-M Double-Precision
XX3	60	F0000248	VSX	xvmaddmsp	VSX Vector Multiply-Add Type-M Single-Precision
XX3	60	F0000700	VSX	xvmaxdp	VSX Vector Maximum Double-Precision
XX3	60	F0000600	VSX	xvmaxsp	VSX Vector Maximum Single-Precision
XX3	60	F0000740	VSX	xvmindp	VSX Vector Minimum Double-Precision
XX3	60	F0000640	VSX	xvminsp	VSX Vector Minimum Single-Precision
XX3	60	F0000388	VSX	xvmsubadp	VSX Vector Multiply-Subtract Type-A Double-Precision
XX3	60	F0000288	VSX	xvmsubasp	VSX Vector Multiply-Subtract Type-A Single-Precision
XX3	60	F00003C8	VSX	xvmsubmdp	VSX Vector Multiply-Subtract Type-M Double-Precision
XX3	60	F00002C8	VSX	xvmsubmsp	VSX Vector Multiply-Subtract Type-M Single-Precision
XX3	60	F0000380	VSX	xvmuldp	VSX Vector Multiply Double-Precision
XX3	60	F0000280	VSX	xvmulsp	VSX Vector Multiply Single-Precision



Table B-1. POWER8 Instructions by Mnemonic (Sheet 23 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
XX2	60	F00007A4	VSX	xvnabsdp	VSX Vector Negative Absolute Value Double-Precision
XX2	60	F00006A4	VSX	xvnabssp	VSX Vector Negative Absolute Value Single-Precision
XX2	60	F00007E4	VSX	xvnegdp	VSX Vector Negate Double-Precision
XX2	60	F00006E4	VSX	xvnegsp	VSX Vector Negate Single-Precision
XX3	60	F0000708	VSX	xvnmaddadp	VSX Vector Negative Multiply-Add Type-A Double-Precision
XX3	60	F0000608	VSX	xvnmaddasp	VSX Vector Negative Multiply-Add Type-A Single-Precision
XX3	60	F0000748	VSX	xvnmaddmdp	VSX Vector Negative Multiply-Add Type-M Double-Precision
XX3	60	F0000648	VSX	xvnmaddmsp	VSX Vector Negative Multiply-Add Type-M Single-Precision
XX3	60	F0000788	VSX	xvnmsubadp	VSX Vector Negative Multiply-Subtract Type-A Double-Precision
XX3	60	F0000688	VSX	xvnmsubasp	VSX Vector Negative Multiply-Subtract Type-A Single-Precision
XX3	60	F00007C8	VSX	xvnmsubmdp	VSX Vector Negative Multiply-Subtract Type-M Double-Precision
XX3	60	F00006C8	VSX	xvnmsubmsp	VSX Vector Negative Multiply-Subtract Type-M Single-Precision
XX2	60	F0000324	VSX	xvrdpi	VSX Vector Round to Double-Precision Integer
XX2	60	F00003AC	VSX	xvrdpic	VSX Vector Round to Double-Precision Integer using Current Rounding Mode
XX2	60	F00003E4	VSX	xvrdpim	VSX Vector Round to Double-Precision Integer toward -Infinity
XX2	60	F00003A4	VSX	xvrdpip	VSX Vector Round to Double-Precision Integer toward +Infinity
XX2	60	F0000364	VSX	xvrdpiz	VSX Vector Round to Double-Precision Integer toward Zero
XX2	60	F0000368	VSX	xvredp	VSX Vector Reciprocal Estimate Double-Precision
XX2	60	F0000268	VSX	xvresp	VSX Vector Reciprocal Estimate Single-Precision
XX2	60	F0000224	VSX	xvrspi	VSX Vector Round to Single-Precision Integer
XX2	60	F00002AC	VSX	xvrspic	VSX Vector Round to Single-Precision Integer using Current Rounding Mode
XX2	60	F00002E4	VSX	xvrspim	VSX Vector Round to Single-Precision Integer toward -Infinity
XX2	60	F00002A4	VSX	xvrspip	VSX Vector Round to Single-Precision Integer toward +Infinity
XX2	60	F0000264	VSX	xvrspiz	VSX Vector Round to Single-Precision Integer toward Zero
XX2	60	F0000328	VSX	xvrsqrtdp	VSX Vector Reciprocal Square Root Estimate Double-Precision
XX2	60	F0000228	VSX	xvrsqrtesp	VSX Vector Reciprocal Square Root Estimate Single-Precision
XX2	60	F000032C	VSX	xvsqrtdp	VSX Vector Square Root Double-Precision
XX2	60	F000022C	VSX	xvsqrtsp	VSX Vector Square Root Single-Precision
XX3	60	F0000340	VSX	xvsubdp	VSX Vector Subtract Double-Precision
XX3	60	F0000240	VSX	xvsubsp	VSX Vector Subtract Single-Precision
XX3	60	F00003E8	VSX	xvtdivdp	VSX Vector Test for Software Divide Double-Precision
XX3	60	F00002E8	VSX	xvtdivsp	VSX Vector Test for Software Divide Single-Precision



Table B-1. POWER8 Instructions by Mnemonic (Sheet 24 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
XX2	60	F00003A8	VSX	xvtsqrtdp	VSX Vector Test for Software Square Root Double-Precision
XX2	60	F00002A8	VSX	xvtsqrtp	VSX Vector Test for Software Square Root Single-Precision
XX3	60	F0000410	VSX	xxland	VSX Logical AND
XX3	60	F0000450	VSX	xxlandc	VSX Logical AND with Complement
XX3	60	F00005D0	VSX	xxleqv	VSX Logical Equivalence
XX3	60	F0000590	VSX	xxlnand	VSX Logical NAND
XX3	60	F0000510	VSX	xxlnor	VSX Logical NOR
XX3	60	F0000490	VSX	xxlor	VSX Logical OR
XX3	60	F0000550	VSX	xxlorc	VSX Logical OR with Complement
XX3	60	F00004D0	VSX	xxlxor	VSX Logical XOR
XX3	60	F0000090	VSX	xxmrghw	VSX Merge High Word
XX3	60	F0000190	VSX	xxmrglw	VSX Merge Low Word
XX3	60	F0000050	VSX	xxpermdi	VSX Permute Doubleword Immediate
XX4	60	F0000030	VSX	xxsel	VSX Select
XX3	60	F0000010	VSX	xxsldwi	VSX Shift Left Double by Word Immediate
XX2	60	F0000290	VSX	xxspltw	VSX Splat Word

Appendix C. POWER8 Instruction Summary by Opcode

Table C-1. lists the instructions implemented in the POWER8 processor in order by opcode. See *Table A-1. Category Listing* on page 313 for a description of the categories.

Table C-1. POWER8 Instructions by Opcode (Sheet 1 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
D	2	08000000	64	tdi	Trap Doubleword Immediate
D	3	0C000000	B	twi	Trap Word Immediate
VX	4	10000000	V	vaddubm	Vector Add Unsigned Byte Modulo
VX	4	10000002	V	vmaxub	Vector Maximum Unsigned Byte
VX	4	10000004	V	vrlb	Vector Rotate Left Byte
VC	4	10000006	V	vcmpequb[.]	Vector Compare Equal To Unsigned Byte
VX	4	10000008	V	vmuloub	Vector Multiply Odd Unsigned Byte
VA	4	10000020	V	vhaddshs	Vector Multiply-High-Add Signed Halfword Saturate
VA	4	10000021	V	vhrraddshs	Vector Multiply-High-Round-Add Signed Halfword Saturate
VA	4	10000022	V	vmladduhm	Vector Multiply-Low-Add Unsigned Halfword Modulo
VA	4	10000024	V	vmsumubm	Vector Multiply-Sum Unsigned Byte Modulo
VA	4	10000025	V	vmsummbm	Vector Multiply-Sum Mixed Byte Modulo
VA	4	10000026	V	vmsumuhm	Vector Multiply-Sum Unsigned Halfword Modulo
VA	4	10000027	V	vmsumuhs	Vector Multiply-Sum Unsigned Halfword Saturate
VA	4	10000028	V	vmsumshm	Vector Multiply-Sum Signed Halfword Modulo
VA	4	10000029	V	vmsumshs	Vector Multiply-Sum Signed Halfword Saturate
VX	4	10000040	V	vadduhm	Vector Add Unsigned Halfword Modulo
VX	4	10000042	V	vmaxuh	Vector Maximum Unsigned Halfword
VX	4	10000044	V	vrlh	Vector Rotate Left Halfword
VC	4	10000046	V	vcmpequh[.]	Vector Compare Equal To Unsigned Halfword
VX	4	10000048	V	vmulouh	Vector Multiply Odd Unsigned Halfword
VX	4	10000080	V	vadduw	Vector Add Unsigned Word Modulo
VX	4	10000082	V	vmaxuw	Vector Maximum Unsigned Word
VX	4	10000084	V	vrlw	Vector Rotate Left Word
VC	4	10000086	V	vcmpequw[.]	Vector Compare Equal To Unsigned Word
VX	4	10000088	V	vmulouw	Vector Multiply Odd Unsigned Word
VX	4	10000089	V	vmuluwm	Vector Multiply Unsigned Word Modulo
VX	4	10000100	V	vadduqm	Vector Add Unsigned Quadword Modulo
VX	4	10000102	V	vmaxsb	Vector Maximum Signed Byte
VX	4	10000104	V	vslb	Vector Shift Left Byte
VX	4	10000108	V	vmulosb	Vector Multiply Odd Signed Byte



Table C-1. POWER8 Instructions by Opcode (Sheet 2 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
VX	4	10000140	V	vaddcuq	Vector Add and Write Carry Unsigned Quadword
VX	4	10000142	V	vmaxsh	Vector Maximum Signed Halfword
VX	4	10000144	V	vslh	Vector Shift Left Halfword
VX	4	10000148	V	vmulosh	Vector Multiply Odd Signed Halfword
VX	4	10000180	V	vaddcuw	Vector Add and Write Carry-Out Unsigned Word
VX	4	10000182	V	vmaxsw	Vector Maximum Signed Word
VX	4	10000184	V	vslw	Vector Shift Left Word
VX	4	10000188	V	vmulosw	Vector Multiply Odd Signed Word
VX	4	10000200	V	vaddubs	Vector Add Unsigned Byte Saturate
VX	4	10000202	V	vminub	Vector Minimum Unsigned Byte
VX	4	10000204	V	vsrb	Vector Shift Right Byte
VC	4	10000206	V	vcmpgtub[.]	Vector Compare Greater Than Unsigned Byte
VX	4	10000208	V	vmuleub	Vector Multiply Even Unsigned Byte
VX	4	10000240	V	vadduhs	Vector Add Unsigned Halfword Saturate
VX	4	10000242	V	vminuh	Vector Minimum Unsigned Halfword
VX	4	10000244	V	vsrh	Vector Shift Right Halfword
VC	4	10000246	V	vcmpgtuh[.]	Vector Compare Greater Than Unsigned Halfword
VX	4	10000248	V	vmuleuh	Vector Multiply Even Unsigned Halfword
VX	4	10000280	V	vadduws	Vector Add Unsigned Word Saturate
VX	4	10000282	V	vminuw	Vector Minimum Unsigned Word
VX	4	10000284	V	vsrw	Vector Shift Right Word
VC	4	10000286	V	vcmpgtuw[.]	Vector Compare Greater Than Unsigned Word
VX	4	10000288	V	vmuleuw	Vector Multiply Even Unsigned Word
VX	4	10000300	V	vaddsbs	Vector Add Signed Byte Saturate
VX	4	10000302	V	vminsbt	Vector Minimum Signed Byte
VX	4	10000304	V	vsrab	Vector Shift Right Algebraic Byte
VC	4	10000306	V	vcmpgtsb[.]	Vector Compare Greater Than Signed Byte
VX	4	10000308	V	vmulesb	Vector Multiply Even Signed Byte
VX	4	10000340	V	vaddshs	Vector Add Signed Halfword Saturate
VX	4	10000342	V	vminsh	Vector Minimum Signed Halfword
VX	4	10000344	V	vsrah	Vector Shift Right Algebraic Halfword
VC	4	10000346	V	vcmpgtsb[.]	Vector Compare Greater Than Signed Halfword
VX	4	10000348	V	vmulesh	Vector Multiply Even Signed Halfword
VX	4	10000380	V	vaddsws	Vector Add Signed Word Saturate

Table C-1. POWER8 Instructions by Opcode (Sheet 3 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
VX	4	10000382	V	vminsw	Vector Minimum Signed Word
VX	4	10000384	V	vsraw	Vector Shift Right Algebraic Word
VC	4	10000386	V	vcmpgtsw[.]	Vector Compare Greater Than Signed Word
VX	4	10000388	V	vmulesw	Vector Multiply Even Signed Word
VX	4	10000400	V	vsububm	Vector Subtract Unsigned Byte Modulo
VX	4	10000401	V	bcdadd.	Decimal Add Modulo
VX	4	10000402	V	vavgub	Vector Average Unsigned Byte
VX	4	10000404	V	vand	Vector Logical AND
VX	4	10000408	V	vpmsumb	Vector Polynomial Multiply-Sum Byte
VX	4	10000440	V	vsubuhm	Vector Subtract Unsigned Halfword Modulo
VX	4	10000441	V	bcdsub.	Decimal Subtract Modulo
VX	4	10000442	V	vavguh	Vector Average Unsigned Halfword
VX	4	10000444	V	vandc	Vector Logical AND with Complement
VX	4	10000448	V	vpmsumh	Vector Polynomial Multiply-Sum Halfword
VX	4	10000480	V	vsubuw	Vector Subtract Unsigned Word Modulo
VX	4	10000482	V	vavgw	Vector Average Unsigned Word
VX	4	10000484	V	vor	Vector Logical OR
VX	4	10000488	V	vpmsumw	Vector Polynomial Multiply-Sum Word
VX	4	10000500	V	vsubuqm	Vector Subtract Unsigned Quadword Modulo
VX	4	10000502	V	vavgsb	Vector Average Signed Byte
VX	4	10000504	V	vnor	Vector Logical NOR
VX	4	10000508	V.AES	vcipher	Vector AES Cipher
VX	4	10000509	V.AES	vcipherlast	Vector AES Cipher Last
VX	4	10000540	V	vsubcuq	Vector Subtract and Write Carry Unsigned Quadword
VX	4	10000542	V	vavgsh	Vector Average Signed Halfword
VX	4	10000544	V	vorc	Vector OR with Complement
VX	4	10000548	V.AES	vncipher	Vector AES Inverse Cipher
VX	4	10000549	V.AES	vncipherlast	Vector AES Inverse Cipher Last
VX	4	10000580	V	vsubcuw	Vector Subtract and Write Carry-Out Unsigned Word
VX	4	10000582	V	vavgsw	Vector Average Signed Word
VX	4	10000584	V	vnand	Vector NAND
VX	4	10000600	V	vsububs	Vector Subtract Unsigned Byte Saturate
VX	4	10000604	V	mfvscr	Move From Vector Status and Control Register
VX	4	10000608	V	vsum4ubs	Vector Sum across Quarter Unsigned Byte Saturate



Table C-1. POWER8 Instructions by Opcode (Sheet 4 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
VX	4	10000640	V	vsubuh	Vector Subtract Unsigned Halfword Saturate
VX	4	10000644	V	mtvscr	Move To Vector Status and Control Register
VX	4	10000648	V	vsum4shs	Vector Sum across Quarter Signed Halfword Saturate
VX	4	10000680	V	vsubuw	Vector Subtract Unsigned Word Saturate
VX	4	10000682	V.SHA2	vshasigmaw	Vector SHA-256 Sigma Word
VX	4	10000684	V	veqv	Vector Equivalence
VX	4	10000688	V	vsum2sws	Vector Sum across Half Signed Word Saturate
VX	4	10000700	V	vsubbs	Vector Subtract Signed Byte Saturate
VX	4	10000702	V	vclzb	Vector Count Leading Zeros Byte
VX	4	10000703	V	vpopcntb	Vector Population Count Byte
VX	4	10000708	V	vsum4sbs	Vector Sum across Quarter Signed Byte Saturate
VX	4	10000740	V	vsubshs	Vector Subtract Signed Halfword Saturate
VX	4	10000742	V	vclzh	Vector Count Leading Zeros Halfword
VX	4	10000743	V	vpopcnth	Vector Population Count Halfword
VX	4	10000780	V	vsubsws	Vector Subtract Signed Word Saturate
VX	4	10000782	V	vclzw	Vector Count Leading Zeros Word
VX	4	10000783	V	vpopcntw	Vector Population Count Word
VX	4	10000788	V	vsumsws	Vector Sum across Signed Word Saturate
VX	4	1000000A	V	vaddfp	Vector Add Single-Precision
VX	4	1000000C	V	vmrghb	Vector Merge High Byte
VX	4	1000000E	V	vpkuhum	Vector Pack Unsigned Halfword Unsigned Modulo
VA	4	1000002A	V	vsel	Vector Select
VA	4	1000002B	V	vperm	Vector Permute
VA	4	1000002C	V	vsldoi	Vector Shift Left Double by Octet Immediate
VA	4	1000002D	V.RAID	vpermxor	Vector Permute and Exclusive-OR
VA	4	1000002E	V	vmaddfp	Vector Multiply-Add Single-Precision
VA	4	1000002F	V	vnmsubfp	Vector Negative Multiply-Subtract Single-Precision
VA	4	1000003C	V	vaddeuqm	Vector Add Extended Unsigned Quadword Modulo
VA	4	1000003D	V	vaddecuq	Vector Add Extended and Write Carry Unsigned Quadword
VA	4	1000003E	V	vsubeuqm	Vector Subtract Extended Unsigned Quadword Modulo
VA	4	1000003F	V	vsubecuq	Vector Subtract Extended and Write Carry Unsigned Quadword
VX	4	1000004A	V	vsubfp	Vector Subtract Single-Precision
VX	4	1000004C	V	vmrghh	Vector Merge High Halfword
VX	4	1000004E	V	vpkuwum	Vector Pack Unsigned Word Unsigned Modulo



Table C-1. POWER8 Instructions by Opcode (Sheet 5 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
VX	4	1000008C	V	vmrghw	Vector Merge High Word
VX	4	1000008E	V	vpkuhus	Vector Pack Unsigned Halfword Unsigned Saturate
VX	4	100000C0	V	vaddudm	Vector Add Unsigned Doubleword Modulo
VX	4	100000C2	V	vmaxud	Vector Maximum Unsigned Doubleword
VX	4	100000C4	V	vrld	Vector Rotate Left Doubleword
VC	4	100000C6	V	vcmpeqfp[.]	Vector Compare Equal To Single-Precision
VC	4	100000C7	V	vcmpequd[.]	Vector Compare Equal To Unsigned Doubleword
VX	4	100000CE	V	vpkuwus	Vector Pack Unsigned Word Unsigned Saturate
VX	4	1000010A	V	vrefp	Vector Reciprocal Estimate Single-Precision
VX	4	1000010C	V	vmrglb	Vector Merge Low Byte
VX	4	1000010E	V	vpkshus	Vector Pack Signed Halfword Unsigned Saturate
VX	4	1000014A	V	vrsqrtefp	Vector Reciprocal Square Root Estimate Single-Precision
VX	4	1000014C	V	vmrglh	Vector Merge Low Halfword
VX	4	1000014E	V	vpkswus	Vector Pack Signed Word Unsigned Saturate
VX	4	1000018A	V	vexptefp	Vector 2 Raised to the Exponent Estimate Single-Precision
VX	4	1000018C	V	vmrglw	Vector Merge Low Word
VX	4	1000018E	V	vpkshss	Vector Pack Signed Halfword Signed Saturate
VX	4	100001C2	V	vmaxsd	Vector Maximum Signed Doubleword
VX	4	100001C4	V	vsl	Vector Shift Left
VC	4	100001C6	V	vcmpgefp[.]	Vector Compare Greater Than or Equal To Single-Precision
VX	4	100001CA	V	vlogefp	Vector Log Base 2 Estimate Single-Precision
VX	4	100001CE	V	vpkswss	Vector Pack Signed Word Signed Saturate
VX	4	1000020A	V	vrfin	Vector Round to Single-Precision Integer Nearest
VX	4	1000020C	V	vspltb	Vector Splat Byte
VX	4	1000020E	V	vupkhsb	Vector Unpack High Signed Byte
VX	4	1000024A	V	vrfiz	Vector Round to Single-Precision Integer toward Zero
VX	4	1000024C	V	vsplth	Vector Splat Halfword
VX	4	1000024E	V	vupkhsh	Vector Unpack High Signed Halfword
VX	4	1000028A	V	vrfip	Vector Round to Single-Precision Integer toward +Infinity
VX	4	1000028C	V	vspltw	Vector Splat Word
VX	4	1000028E	V	vupklsb	Vector Unpack Low Signed Byte
VX	4	100002C2	V	vminud	Vector Minimum Unsigned Doubleword
VX	4	100002C4	V	vsr	Vector Shift Right
VC	4	100002C6	V	vcmpgtfp[.]	Vector Compare Greater Than Single-Precision



Table C-1. POWER8 Instructions by Opcode (Sheet 6 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
VC	4	100002C7	V	vcmpgtud[.]	Vector Compare Greater Than Unsigned Doubleword
VX	4	100002CA	V	vrfim	Vector Round to Single-Precision Integer toward -Infinity
VX	4	100002CE	V	vupklsh	Vector Unpack Low Signed Halfword
VX	4	1000030A	V	vcfux	Vector Convert From Unsigned Fixed-Point Word
VX	4	1000030C	V	vspltsb	Vector Splat Immediate Signed Byte
VX	4	1000030E	V	vpkpx	Vector Pack Pixel
VX	4	1000034A	V	vcfsx	Vector Convert From Signed Fixed-Point Word To Single-Precision
VX	4	1000034C	V	vspltsb	Vector Splat Immediate Signed Halfword
VX	4	1000034E	V	vupkhpx	Vector Unpack High Pixel
VX	4	1000038A	V	vctuxs	Vector Convert From Single-Precision To Unsigned Fixed-Point Word Saturate
VX	4	1000038C	V	vspltsw	Vector Splat Immediate Signed Word
X	4	100003C2	V	vminsrd	Vector Minimum Signed Doubleword
VX	4	100003C4	V	vsrad	Vector Shift Right Algebraic Doubleword
VC	4	100003C6	V	vcmpbfp[.]	Vector Compare Bounds Single-Precision
VC	4	100003C7	V	vcmpgtsd[.]	Vector Compare Greater Than Signed Doubleword
VX	4	100003CA	V	vctsxs	Vector Convert From Single-Precision To Signed Fixed-Point Word Saturate
VX	4	100003CE	V	vupklpx	Vector Unpack Low Pixel
VX	4	1000040A	V	vmaxfp	Vector Maximum Single-Precision
VX	4	1000040C	V	vslo	Vector Shift Left by Octet
VX	4	1000044A	V	vminfp	Vector Minimum Single-Precision
VX	4	1000044C	V	vsro	Vector Shift Right by Octet
VX	4	1000044E	V	vpkudum	Vector Pack Unsigned Doubleword Unsigned Modulo
VX	4	100004C0	V	vsubudm	Vector Subtract Unsigned Doubleword Modulo
VX	4	100004C4	V	vxor	Vector Logical XOR
VX	4	100004C8	V	vpmsumd	Vector Polynomial Multiply-Sum Doubleword
VX	4	100004CE	V	vpkudus	Vector Pack Unsigned Doubleword Unsigned Saturate
VX	4	1000050C	V	vgbbd	Vector Gather Bits by Byte by Doubleword
VX	4	1000054C	V	vbpermq	Vector Bit Permute Quadword
VX	4	1000054C	V	vbpermq	Vector Bit Permute Quadword
VX	4	1000054E	V	vpkdsdus	Vector Pack Signed Doubleword Unsigned Saturate
VX	4	100005C4	V	vsld	Vector Shift Left Doubleword
VX	4	100005C8	V.AES	vsbox	Vector AES S-Box
VX	4	100005CE	V	vpkdsdss	Vector Pack Signed Doubleword Signed Saturate



Table C-1. POWER8 Instructions by Opcode (Sheet 7 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
VX	4	1000064E	V	vupkhsw	Vector Unpack High Signed Word
VX	4	1000068C	VSX	vmrgow	Vector Merge Odd Word
VX	4	100006C2	V.SHA2	vshasigmad	Vector SHA-512 Sigma Doubleword
VX	4	100006C4	V	vsrd	Vector Shift Right Doubleword
VX	4	100006CE	V	vupklsw	Vector Unpack Low Signed Word
VX	4	1000078C	VSX	vmrgew	Vector Merge Even Word
VX	4	100007C2	V	vclzd	Vector Count Leading Zeros Doubleword
VX	4	100007C3	V	vpopcntd	Vector Population Count Doubleword
D	7	1C000000	B	mulli	Multiply Low Immediate
D	8	20000000	B	subfic	Subtract From Immediate Carrying
D	10	28000000	B	cmpli	Compare Logical Immediate
D	11	2C000000	B	cmpi	Compare Immediate
D	12	30000000	B	addic	Add Immediate Carrying
D	13	34000000	B	addic.	Add Immediate Carrying and Record CR0
D	14	38000000	B	addi	Add Immediate
D	15	3C000000	B	addis	Add Immediate Shifted
B	16	40000000	B	bcl[l][a]	Branch Conditional
SC	17	44000002	B	sc	System Call
I	18	48000000	B	b[l][a]	Branch
XL	19	4C000000	B	mcrf	Move Condition Register Field
XL	19	4C000020	B	bclr[l]	Branch Conditional to Link Register
XL	19	4C000024	S	rfdi	Return from Interrupt Doubleword
XL	19	4C000042	B	crnor	Condition Register NOR
XL	19	4C000102	B	crandc	Condition Register AND with Complement
XL	19	4C000124	S	rfebb	Return from Event Based Branch
XL	19	4C00012C	B	isync	Instruction Synchronize
XL	19	4C000182	B	crxor	Condition Register XOR
XL	19	4C0001C2	B	crnand	Condition Register NAND
XL	19	4C000202	B	crand	Condition Register AND
XL	19	4C000224	S	hrfid	Return From Interrupt Doubleword Hypervisor
XL	19	4C000242	B	creqv	Condition Register Equivalent
XL	19	4C000324	S	doze	Doze
XL	19	4C000342	B	crorc	Condition Register OR with Complement
XL	19	4C000364	S	nap	Nap



Table C-1. POWER8 Instructions by Opcode (Sheet 8 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
XL	19	4C000382	B	cror	Condition Register OR
XL	19	4C0003A4	S	sleep	Sleep
XL	19	4C0003E4	S	rvwinkle	Rip Van Winkle
XL	19	4C000420	B	bcctr[.]	Branch Conditional to Count Register
X	19	4C000460	B	bctar[.]	Branch Conditional to Branch Target Address Register
M	20	50000000	B	rlwimi[.]	Rotate Left Word Immediate then Mask Insert
M	21	54000000	B	rlwinm[.]	Rotate Left Word Immediate then AND with Mask
M	23	5C000000	B	rlwnm[.]	Rotate Left Word then AND with Mask
D	24	60000000	B	ori	OR Immediate
D	25	64000000	B	oris	OR Immediate Shifted
X	26	68000000	B	xnop	Executed No Operation
D	26	68000000	B	xori	XOR Immediate
D	27	6C000000	B	xoris	XOR Immediate Shifted
D	28	70000000	B	andi.	AND Immediate and Record CR0
D	29	74000000	B	andis.	AND Immediate Shifted and Record CR0
MD	30	78000000	64	rldicl[.]	Rotate Left Doubleword Immediate then Clear Left
MD	30	78000004	64	rldicr[.]	Rotate Left Doubleword Immediate then Clear Right
MD	30	78000008	64	rldic[.]	Rotate Left Doubleword Immediate then Clear
MDS	30	78000010	64	rldcl[.]	Rotate Left Doubleword then Clear Left
MDS	30	78000012	64	rldcr[.]	Rotate Left Doubleword then Clear Right
MD	30	7800000C	64	rldimi[.]	Rotate Left Doubleword Immediate then Mask Insert
X	31	7C000000	B	cmp	Compare
X	31	7C000008	B	tw	Trap Word
X	31	7C00000C	V	lvsl	Load Vector for Shift Left
X	31	7C00000E	V	lvebx	Load Vector Element Byte Indexed
XO	31	7C000010	B	subfc[.]	Subtract From Carrying
XO	31	7C000012	64	mulhdu[.]	Multiply High Doubleword Unsigned
XO	31	7C000014	B	addc[.]	Add Carrying
XO	31	7C000016	B	mulhwu[.]	Multiply High Word Unsigned
XX1	31	7C000018	VSX	lxsiwzx	Load VSX Scalar as Integer Word and Zero Indexed
A	31	7C00001E	B	isel	Integer Select
XFX	31	7C000026	B	mfcr	Move From Condition Register
X	31	7C000028	B	lwaxr	Load Word and Reserve Indexed
X	31	7C00002A	64	Idx	Load Doubleword Indexed



Table C-1. POWER8 Instructions by Opcode (Sheet 9 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
X	31	7C00002C	B	icbt	Instruction Cache Block Touch
X	31	7C00002E	B	lwzx	Load Word and Zero Indexed
X	31	7C000030	B	slw[.]	Shift Left Word
X	31	7C000034	B	cntlzw[.]	Count Leading Zeros Word
X	31	7C000036	64	sld[.]	Shift Left Doubleword
X	31	7C000038	B	and[.]	AND
X	31	7C00003A	64	ldepsx	Load Doubleword by External PID Indexed
X	31	7C000040	B	cmpl	Compare Logical
X	31	7C00004C	V	lvsr	Load Vector for Shift Right
X	31	7C00004E	V	lvehx	Load Vector Element Halfword Indexed
XO	31	7C000050	B	subf[.]	Subtract From
XX1	31	7C000066	VSX	mfvsrd	Move From VSR Doubleword
X	31	7C000068	B	lbarx	Load Byte And Reserve Indexed
X	31	7C00006A	64	ldux	Load Doubleword with Update Indexed
X	31	7C00006C	B	dcbst	Data Cache Block Store
X	31	7C00006E	B	lwzux	Load Word and Zero with Update Indexed
X	31	7C000074	64	cntlzd[.]	Count Leading Zeros Doubleword
X	31	7C000078	B	andc[.]	AND with Complement
X	31	7C000088	64	td	Trap Doubleword
X	31	7C00008E	V	lvewx	Load Vector Element Word Indexed
XO	31	7C000092	64	mulhd[.]	Multiply High Doubleword
XO	31	7C000096	B	mulhw[.]	Multiply High Word
XX1	31	7C000098	VSX	lxsiwax	Load VSX Scalar as Integer Word Algebraic Indexed
X	31	7C0000A6	S	mfmsr	Move From Machine State Register
X	31	7C0000A8	64	ldarx	Load Doubleword And Reserve Indexed
X	31	7C0000AC	B	dcbf	Data Cache Block Flush
X	31	7C0000AE	B	lbzx	Load Byte and Zero Indexed
X	31	7C0000CE	V	lvx	Load Vector Indexed
XO	31	7C0000D0	B	neg[.]	Negate
XX1	31	7C0000E6	VSX	mfvsrwz	Move From VSR Word and Zero
X	31	7C0000E8	B	lharcx	Load Halfword And Reserve Indexed Xform
X	31	7C0000EE	B	lbzux	Load Byte and Zero with Update Indexed
X	31	7C0000F4	B	popcntb	Population Count Byte-wise
X	31	7C0000F8	B	nor[.]	NOR



Table C-1. POWER8 Instructions by Opcode (Sheet 10 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
X	31	7C00010E	V	stvebx	Store Vector Element Byte Indexed
XO	31	7C000110	B	subfe[.]	Subtract From Extended
XO	31	7C000114	B	adde[.]	Add Extended
XX1	31	7C000118	VSX	stxsiwx	Store VSX Scalar as Integer Word Indexed
X	31	7C00011C	S	msgsndp	Message Send Privileged
XFX	31	7C000120	B	mtcrf	Move To Condition Register Fields
X	31	7C000124	S	mtmsr	Move To Machine State Register
X	31	7C00012A	64	stdx	Store Doubleword Indexed
X	31	7C00012D	B	stwcx.	Store Word Conditional Indexed and Record CR0
X	31	7C00012E	B	stwx	Store Word Indexed
X	31	7C000134	B	ptryw	Parity Word
X	31	7C00013A	64	stdepsx	Store Doubleword by External PID Indexed
X	31	7C00014E	V	stvehx	Store Vector Element Halfword Indexed
X	31	7C00015C	S	msgclr	Message Clear Privileged
X	31	7C000164	S	mtmsrd	Move To Machine State Register Doubleword
XX1	31	7C000166	VSX	mtvsrd	Move To VSR Doubleword
X	31	7C00016A	64	stdux	Store Doubleword with Update Indexed
X	31	7C00016D	LSQ	stqcx.	Store Quadword Conditional Indexed and record CR0
X	31	7C00016E	B	stwux	Store Word with Update Indexed
X	31	7C000174	64	ptryd	Parity Doubleword
X	31	7C00018E	V	stvewx	Store Vector Element Word Indexed
XO	31	7C000190	B	subfze[.]	Subtract From Zero Extended
XO	31	7C000194	B	addze[.]	Add to Zero Extended
X	31	7C00019C	S	msgsnd	Message Send
X	31	7C0001A4	S	mtsr	Move To Segment Register
XX1	31	7C0001A6	VSX	mtvsrwa	Move To VSR Word Algebraic
X	31	7C0001AD	64	stdcx.	Store Doubleword Conditional Indexed and Record CR0
X	31	7C0001AE	B	stbx	Store Byte Indexed
X	31	7C0001CE	V	stvx	Store Vector Indexed
XO	31	7C0001D0	B	subfme[.]	Subtract From Minus One Extended
XO	31	7C0001D2	64	mulld[.]	Multiply Low Doubleword
XO	31	7C0001D4	B	addme[.]	Add to Minus One Extended
XO	31	7C0001D6	B	mullw[.]	Multiply Low Word
X	31	7C0001DC	S	msgclr	Message Clear



Table C-1. POWER8 Instructions by Opcode (Sheet 11 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
X	31	7C0001E4	S	mtrsri	Move To Segment Register Indirect
XX1	31	7C0001E6	VSX	mtvsrwz	Move To VSR Word and Zero
X	31	7C0001EC	B	dcbtst	Data Cache Block Touch for Store
X	31	7C0001EE	B	stbux	Store Byte with Update Indexed
X	31	7C0001F8	64	bpermd	Bit Permute Doubleword
XO	31	7C000214	B	add[.]	Add
X	31	7C000224	S	tlbiel	TLB Invalidate Entry Local
X	31	7C000228	LSQ	lqarx	Load Quadword And Reserve Indexed
X	31	7C00022C	B	dcbt	Data Cache Block Touch
X	31	7C00022E	B	lhzx	Load Halfword and Zero Indexed
X	31	7C000238	B	eqv[.]	Equivalent
XFX	31	7C00025C	S	mfbhrbe	Move From Branch History Rolling Buffer
X	31	7C000264	S	tlbie	TLB Invalidate Entry
X	31	7C00026E	B	lhzux	Load Halfword and Zero with Update Indexed
X	31	7C000278	B	xor[.]	XOR
XX1	31	7C000298	VSX	lxvdsx	Load VSR Vector Doubleword and Splat Indexed
XFX	31	7C0002A6	B	mfspr	Move From Special Purpose Register
X	31	7C0002AA	64	lwax	Load Word Algebraic Indexed
X	31	7C0002AE	B	lhax	Load Halfword Algebraic Indexed
X	31	7C0002CE	V	lvxl	Load Vector Indexed Last
XFX	31	7C0002E6	S.out	mftb	Move From Time Base
X	31	7C0002EA	64	lwaux	Load Word Algebraic with Update Indexed
X	31	7C0002EE	B	lhaux	Load Halfword Algebraic with Update Indexed
X	31	7C0002F4	B	popcntw	Population Count Words
XO	31	7C000312	64	divdeu[.]	Divide Doubleword Extended Unsigned
XO	31	7C000316	B	divweu[.]	Divide Word Extended Unsigned
X	31	7C000324	S	slbmte	SLB Move To Entry
X	31	7C00032E	B	sthx	Store Halfword Indexed
X	31	7C000338	B	orc[.]	OR with Complement
XO	31	7C000352	64	divde[.]	Divide Doubleword Extended
XO	31	7C000356	B	divwe[.]	Divide Word Extended
X	31	7C00035C	S	clrbhrb	Clear BHRB
X	31	7C000364	S	slbie	SLB Invalidate Entry
X	31	7C00036E	B	sthux	Store Halfword with Update Indexed



Table C-1. POWER8 Instructions by Opcode (Sheet 12 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
X	31	7C000378	B	or[.]	OR
XO	31	7C000392	64	divdu[.]	Divide Doubleword Unsigned
XO	31	7C000396	B	divwu[.]	Divide Word Unsigned
XFX	31	7C0003A6	B	mtspr	Move To Special Purpose Register
X	31	7C0003B8	B	nand[.]	NAND
X	31	7C0003CE	V	stvxl	Store Vector Indexed Last
XO	31	7C0003D2	64	divd[.]	Divide Doubleword
XO	31	7C0003D6	B	divw[.]	Divide Word
X	31	7C0003E4	S	slbia	SLB Invalidate All
X	31	7C0003F4	64	popcntd	Population Count Doubleword
X	31	7C0003F8	B	cmpb	Compare Byte
XO	31	7C000410	B	subfco[.]	Subtract From Carrying and Record OV
XO	31	7C000414	B	addco[.]	Add Carrying and Record OV
XX1	31	7C000418	VSX	lxsspx	Load VSX Scalar Single-Precision Indexed
X	31	7C000428	64	ldbrx	Load Doubleword Byte-Reverse Indexed
X	31	7C00042A	MA	lswx	Load String Word Indexed
X	31	7C00042C	B	lwbrx	Load Word Byte-Reverse Indexed
X	31	7C00042E	FP	lfsx	Load Floating-Point Single Indexed
X	31	7C000430	B	srw[.]	Shift Right Word
X	31	7C000436	64	srd[.]	Shift Right Doubleword
XO	31	7C000450	B	subfo[.]	Subtract From and Record OV
X	31	7C00046C	S	tlbsync	TLB Synchronize
X	31	7C00046E	FP	lfsux	Load Floating-Point Single with Update Indexed
XX1	31	7C000498	VSX	lxsdx	Load VSR Scalar Doubleword Indexed
X	31	7C0004A6	S	mfsr	Move From Segment Register
X	31	7C0004AA	MA	lswi	Load String Word Immediate
X	31	7C0004AC	B	sync	Synchronize
X	31	7C0004AE	FP	lfdx	Load Floating-Point Double Indexed
XO	31	7C0004D0	B	nego[.]	Negate and Record OV
X	31	7C0004EE	FP	lfdux	Load Floating-Point Double with Update Indexed
XO	31	7C000510	B	subfeo[.]	Subtract From Extended and Record OV
XO	31	7C000514	B	addeo[.]	Add Extended and Record OV and Record OV
XX1	31	7C000518	VSX	stxsspx	Store VSR Scalar Word Indexed
X	31	7C00051D	TM	tbegin.	Transaction Begin



Table C-1. POWER8 Instructions by Opcode (Sheet 13 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
X	31	7C000526	S	mfsrin	Move From Segment Register Indirect
X	31	7C000528	64	stdbrx	Store Doubleword Byte-Reverse Indexed
X	31	7C00052A	MA	stswx	Store String Word Indexed
X	31	7C00052C	B	stwbrx	Store Word Byte-Reverse Indexed
X	31	7C00052E	FP	stfsx	Store Floating-Point Single Indexed
X	31	7C00055C	TM	tend.	Transaction End
X	31	7C00056D	B	stbcx.	Store Byte Conditional Indexed
X	31	7C00056E	FP	stfsux	Store Floating-Point Single with Update Indexed
XO	31	7C000590	B	subfzeo[.]	Subtract From Zero Extended and Record OV
XO	31	7C000594	B	addzeo[.]	Add to Zero Extended and Record OV
XX1	31	7C000598	VSX	stxsdx	Store VSR Scalar Doubleword Indexed
X	31	7C00059C	TM	tcheck	Transaction Check
X	31	7C0005AA	MA	stswi	Store String Word Immediate
X	31	7C0005AD	B	sthcx.	Store Halfword Conditional Indexed Xform
X	31	7C0005AE	FP	stfdx	Store Floating-Point Double Indexed
XO	31	7C0005D0	B	subfmeo[.]	Subtract From Minus One Extended and Record OV
XO	31	7C0005D2	64	mulldo[.]	Multiply Low Doubleword and Record OV
XO	31	7C0005D4	B	addmeo[.]	Add to Minus One Extended and Record OV
XO	31	7C0005D6	B	mulwo[.]	Multiply Low Word and Record OV
X	31	7C0005EE	FP	stfdx	Store Floating-Point Double with Update Indexed
XO	31	7C000614	B	addo[.]	Add and Record OV
XX1	31	7C000618	VSX	lxvw4x	Load VSR Vector Word*4 Indexed
X	31	7C00061D	TM	tabortwc.	Transaction Abort Word Conditional
X	31	7C00062A	S	lwzcix	Load Word and Zero Caching Inhibited Indexed
X	31	7C00062C	B	lhbrx	Load Halfword Byte-Reverse Indexed
X	31	7C00062E	FP.out	lfdp	Load Floating-Point Double Pair Indexed
X	31	7C000630	B	sraw[.]	Shift Right Algebraic Word
X	31	7C000634	64	srad[.]	Shift Right Algebraic Doubleword
X	31	7C00065D	TM	tabortdc.	Transaction Abort Doubleword Conditional
X	31	7C00066A	S	lhzcix	Load Halfword and Zero Caching Inhibited Indexed
X	31	7C000670	B	srawi[.]	Shift Right Algebraic Word Immediate
XS	31	7C000674	64	sradi[.]	Shift Right Algebraic Doubleword Immediate
XX1	31	7C000698	VSX	lxvd2x	Load VSR Vector Doubleword*2 Indexed
X	31	7C00069D	TM	tabortwci.	Transaction Abort Word Conditional Immediate



Table C-1. POWER8 Instructions by Opcode (Sheet 14 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
X	31	7C0006A6	S	slbmfev	SLB Move From Entry VSID
X	31	7C0006AA	S	lbzcix	Load Byte and Zero Caching Inhibited Indexed
X	31	7C0006AC	S	eieio	Enforce In-order Execution of I/O
X	31	7C0006AE	FP	lfiwax	Load Floating-Point as Integer Word Algebraic Indexed
X	31	7C0006DD	TM	tabortdci.	Transaction Abort Doubleword Conditional Immediate
X	31	7C0006EA	S	ldcx	Load Doubleword Caching Inhibited Indexed
X	31	7C0006EE	FP	lfiwzx	Load Floating-Point as Integer Word and Zero Indexed
XO	31	7C000712	64	divdeuo[.]	Divide Doubleword Extended Unsigned and Record OV
XO	31	7C000716	B	divweuo[.]	Divide Word Extended Unsigned and Record OV
XX1	31	7C000718	VSX	stxvw4x	Store VSR Vector Word*4 Indexed
X	31	7C00071D	TM	tabort.	Transaction Abort
X	31	7C000726	S	slbmfee	SLB Move From Entry ESID
X	31	7C00072A	S	stwcix	Store Word and Zero Caching Inhibited Indexed
X	31	7C00072C	B	sthbrx	Store Halfword Byte-Reverse Indexed
X	31	7C00072E	FP.out	stfdpx	Store Floating-Point Double Pair Indexed
X	31	7C000734	B	extsh[.]	Extend Sign Halfword
XO	31	7C000752	64	divdeo[.]	Divide Doubleword Extended and Record OV
XO	31	7C000756	B	divweo[.]	Divide Word Extended and Record OV
X	31	7C00075D	TM	treclaim.	Transaction Reclaim
X	31	7C00076A	S	sthcix	Store Halfword and Zero Caching Inhibited Indexed
X	31	7C000774	B	extsb[.]	Extend Sign Byte
XO	31	7C000792	64	divduo[.]	Divide Doubleword Unsigned and Record OV
XO	31	7C000796	B	divwu[.]	Divide Word Unsigned and Record OV
XX1	31	7C000798	VSX	stxvd2x	Store VSR Vector Doubleword*2 Indexed
X	31	7C0007A7	S	slbfee.	SLB Find Entry ESID
X	31	7C0007AA	S	stbcix	Store Byte Caching Inhibited Indexed
X	31	7C0007AC	B	icbi	Instruction Cache Block Invalidate
X	31	7C0007AE	FP	stfiwx	Store Floating-Point as Integer Word Indexed
X	31	7C0007B4	64	extsw[.]	Extend Sign Word
XO	31	7C0007D2	64	divdo[.]	Divide Doubleword and Record OV
XO	31	7C0007D6	B	divwo[.]	Divide Word and Record OV
X	31	7C0007DD	TM	trechkpt.	Transaction Recheckpoint
X	31	7C0007EA	S	stdcix	Store Doubleword Caching Inhibited Indexed
X	31	7C0007EC	B	dcbz	Data Cache Block Zero



Table C-1. POWER8 Instructions by Opcode (Sheet 15 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
XFX	31	7C100026	B	mfocrf	Move From One Condition Register Field
XFX	31	7C100120	B	mtocrf	Move To One Condition Register Field
D	32	80000000	B	lwz	Load Word and Zero
D	33	84000000	B	lwzu	Load Word and Zero with Update
D	34	88000000	B	lbz	Load Byte and Zero
D	35	8C000000	B	lbzu	Load Byte and Zero with Update
D	36	90000000	B	stw	Store Word
D	37	94000000	B	stwu	Store Word with Update
D	38	98000000	B	stb	Store Byte
D	39	9C000000	B	stbu	Store Byte with Update
D	40	A0000000	B	lhz	Load Halfword and Zero
D	41	A4000000	B	lhzu	Load Halfword and Zero with Update
D	42	A8000000	B	lha	Load Halfword Algebraic
D	43	AC000000	B	lhau	Load Halfword Algebraic with Update
D	44	B0000000	B	sth	Store Halfword
D	45	B4000000	B	sthu	Store Halfword with Update
D	46	B8000000	B	lmw	Load Multiple Word
D	47	BC000000	B	stmw	Store Multiple Word
D	48	C0000000	FP	lfs	Load Floating-Point Single
D	49	C4000000	FP	lfsu	Load Floating-Point Single with Update
D	50	C8000000	FP	lfd	Load Floating-Point Double
D	51	CC000000	FP	lfdw	Load Floating-Point Double with Update
D	52	D0000000	FP	stfs	Store Floating-Point Single
D	53	D4000000	FP	stfsu	Store Floating-Point Single with Update
D	54	D8000000	FP	stfd	Store Floating-Point Double
D	55	DC000000	FP	stfdw	Store Floating-Point Double with Update
DQ	56	E0000000	LSQ	lq	Load Quadword
DS	57	E4000000	FP.out	lfdp	Load Floating-Point Double Pair
DS	58	E8000000	64	ld	Load Doubleword
DS	58	E8000001	64	ldu	Load Doubleword with Update
DS	58	E8000002	64	lwa	Load Word Algebraic
X	59	EC000004	DFP	dadd[.]	Decimal Floating Add
Z23	59	EC000006	DFP	dqua[.]	Decimal Quantize
A	59	EC000024	FP[R]	fdivs[.]	Floating Divide Single



Table C-1. POWER8 Instructions by Opcode (Sheet 16 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
A	59	EC000028	FP[R]	fsubs[.]	Floating Subtract Single
A	59	EC00002A	FP[R]	fadds[.]	Floating Add Single
A	59	EC00002C	FP[R]	fsqrts[.]	Floating Square Root Single
A	59	EC000030	FP[R]	fres[.]	Floating Reciprocal Estimate Single
A	59	EC000032	FP[R]	fmuls[.]	Floating Multiply Single
A	59	EC000034	FP[R].in	frsqrtes[.]	Floating Reciprocal Square Root Estimate Single
A	59	EC000038	FP[R]	fmsubs[.]	Floating Multiply-Subtract Single
A	59	EC00003A	FP[R]	fmadds[.]	Floating Multiply-Add Single
A	59	EC00003C	FP[R]	fnmsubs[.]	Floating Negative Multiply-Subtract Single
A	59	EC00003E	FP[R]	fnmadds[.]	Floating Negative Multiply-Add Single
X	59	EC000044	DFP	dmul[.]	Decimal Floating Multiply
Z23	59	EC000046	DFP	drrnd[.]	Decimal Floating Reround
Z22	59	EC000084	DFP	dscli[.]	Decimal Floating Shift Coefficient Left Immediate
Z23	59	EC000086	DFP	dquai[.]	Decimal Quantize Immediate
Z22	59	EC0000C4	DFP	dscri[.]	Decimal Floating Shift Coefficient Right Immediate
Z23	59	EC0000C6	DFP	drinx[.]	Decimal Floating Round To FP Integer With Inexact
X	59	EC000104	DFP	dcmpo	Decimal Floating Compare Ordered
X	59	EC000144	DFP	dtstex	Decimal Floating Test Exponent
Z22	59	EC000184	DFP	dtstdc	Decimal Floating Test Data Class
Z22	59	EC0001C4	DFP	dtstdg	Decimal Floating Test Data Group
Z23	59	EC0001C6	DFP	drintr[.]	Decimal Floating Round To FP Integer Without Inexact
X	59	EC000204	DFP	dctdp[.]	Decimal Floating Convert To DFP Long
X	59	EC000244	DFP	dctfix[.]	Decimal Floating Convert To Fixed
X	59	EC000284	DFP	ddedpd[.]	Decimal Floating Decode DPD To BCD
X	59	EC0002C4	DFP	dxex[.]	Decimal Floating Extract Exponent
X	59	EC000404	DFP	dsub[.]	Decimal Floating Subtract
X	59	EC000444	DFP	ddiv[.]	Decimal Floating Divide
X	59	EC000504	DFP	dcmplu	Decimal Floating Compare Unordered
X	59	EC000544	DFP	dtstsf	Decimal Floating Test Significance
X	59	EC000604	DFP	drsp[.]	Decimal Floating Round To DFP Short
X	59	EC000644	DFP	dcffix[.]	Decimal Floating Convert From Fixed
X	59	EC000684	DFP	denbcd[.]	Decimal Floating Encode BCD To DPD
X	59	EC00069C	FP[R]	fcfids[.]	Floating Convert From Integer Doubleword Single
X	59	EC0006C4	DFP	diex[.]	Decimal Floating Insert Exponent



Table C-1. POWER8 Instructions by Opcode (Sheet 17 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
X	59	EC00079C	FPR	fcfidus[.]	Floating Convert From Integer Doubleword Unsigned Single
XX3	60	F0000000	VSX	x saddsp	VSX Scalar Add Single-Precision
XX3	60	F0000008	VSX	x smaddasp	VSX Scalar Multiply-Add Type-A Single-Precision
XX3	60	F0000010	VSX	xxsldwi	VSX Shift Left Double by Word Immediate
XX2	60	F0000028	VSX	x srqtesp	VSX Scalar Reciprocal Square Root Estimate Single-Precision
XX2	60	F000002C	VSX	x ssqrtsp	VSX Scalar Square Root Single-Precision
XX4	60	F0000030	VSX	xxsel	VSX Select
XX3	60	F0000040	VSX	x ssubsp	VSX Scalar Subtract Single-Precision
XX3	60	F0000048	VSX	x smaddmsp	VSX Scalar Multiply-Add Type-M Single-Precision
XX3	60	F0000050	VSX	xxpermmdi	VSX Permute Doubleword Immediate
XX2	60	F0000068	VSX	xsresp	VSX Scalar Reciprocal Estimate Single-Precision
XX3	60	F0000080	VSX	x smulsp	VSX Scalar Multiply Single-Precision
XX3	60	F0000088	VSX	x smsubasp	VSX Scalar Multiply-Subtract Type-A Single-Precision
XX3	60	F0000090	VSX	xxmrghw	VSX Merge High Word
XX3	60	F00000C0	VSX	x sdivsp	VSX Scalar Divide Single-Precision
XX3	60	F00000C8	VSX	x smsubmsp	VSX Scalar Multiply-Subtract Type-M Single-Precision
XX3	60	F0000100	VSX	x saddrp	VSX Scalar Add Double-Precision
XX3	60	F0000108	VSX	x smaddadp	VSX Scalar Multiply-Add Type-A Double-Precision
XX3	60	F0000118	VSX	x scmpudp	VSX Scalar Compare Unordered Double-Precision
XX2	60	F0000120	VSX	x scvdpxws	VSX Scalar Convert Double-Precision to Unsigned Fixed-Point Word Saturate
XX2	60	F0000124	VSX	x srdpi	VSX Scalar Round to Double-Precision Integer
XX2	60	F0000128	VSX	x srqtdp	VSX Scalar Reciprocal Square Root Estimate Double-Precision
XX2	60	F000012C	VSX	x ssqrtdp	VSX Scalar Square Root Double-Precision
XX3	60	F0000140	VSX	x ssubdp	VSX Scalar Subtract Double-Precision
XX3	60	F0000148	VSX	x smaddmdp	VSX Scalar Multiply-Add Type-M Double-Precision
XX3	60	F0000158	VSX	x scmpodp	VSX Scalar Compare Ordered Double-Precision
XX2	60	F0000160	VSX	x scvdpsxws	VSX Scalar Convert Double-Precision to Signed Fixed-Point Word Saturate
XX2	60	F0000164	VSX	x srdpiz	VSX Scalar Round to Double-Precision Integer toward Zero
XX1	60	F0000168	VSX	x sredp	VSX Scalar Reciprocal Estimate Double-Precision
XX3	60	F0000180	VSX	x smuldp	VSX Scalar Multiply Double-Precision
XX3	60	F0000188	VSX	x smsubadp	VSX Scalar Multiply-Subtract Type-A Double-Precision
XX3	60	F0000190	VSX	xxmrglw	VSX Merge Low Word
XX2	60	F00001A4	VSX	x srdpip	VSX Scalar Round to Double-Precision Integer toward +Infinity



Table C-1. POWER8 Instructions by Opcode (Sheet 18 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
XX2	60	F00001A8	VSX	xstsqrtdp	VSX Scalar Test for Software Square Root Double-Precision
XX2	60	F00001AC	VSX	xsrdpic	VSX Scalar Round to Double-Precision Integer using Current Rounding Mode
XX3	60	F00001C0	VSX	xsdivdp	VSX Scalar Divide Double-Precision
XX3	60	F00001C8	VSX	xsmsubmdp	VSX Scalar Multiply-Subtract Type-M Double-Precision
XX2	60	F00001E4	VSX	xsrdpim	VSX Scalar Round to Double-Precision Integer toward -Infinity
XX3	60	F00001E8	VSX	xstdivdp	VSX Scalar Test for Software Divide Double-Precision
XX3	60	F0000200	VSX	xvaddsp	VSX Vector Add Single-Precision
XX3	60	F0000208	VSX	xvmaddasp	VSX Vector Multiply-Add Type-A Single-Precision
XX3	60	F0000218	VSX	xvccmpeqsp	VSX Vector Compare Equal To Single-Precision
XX2	60	F0000220	VSX	xvcvspuxws	VSX Vector Convert Single-Precision to Unsigned Fixed-Point Word Saturate
XX2	60	F0000224	VSX	xvrspi	VSX Vector Round to Single-Precision Integer
XX2	60	F0000228	VSX	xvrsqrtesp	VSX Vector Reciprocal Square Root Estimate Single-Precision
XX2	60	F000022C	VSX	xvsqrtsp	VSX Vector Square Root Single-Precision
XX3	60	F0000240	VSX	xvsubsp	VSX Vector Subtract Single-Precision
XX3	60	F0000248	VSX	xvmaddmsp	VSX Vector Multiply-Add Type-M Single-Precision
XX3	60	F0000258	VSX	xvcmpgtsp	VSX Vector Compare Greater Than Single-Precision
XX2	60	F0000260	VSX	xvcvspsxws	VSX Vector Convert Single-Precision to Signed Fixed-Point Word Saturate
XX2	60	F0000264	VSX	xvrspiz	VSX Vector Round to Single-Precision Integer toward Zero
XX2	60	F0000268	VSX	xvresp	VSX Vector Reciprocal Estimate Single-Precision
XX3	60	F0000280	VSX	xvmulsp	VSX Vector Multiply Single-Precision
XX3	60	F0000288	VSX	xvmsubasp	VSX Vector Multiply-Subtract Type-A Single-Precision
XX2	60	F0000290	VSX	xxspltw	VSX Splat Word
XX3	60	F0000298	VSX	xvcmpgesp	VSX Vector Compare Greater Than or Equal To Single-Precision
XX2	60	F00002A0	VSX	xvcvuxwsp	VSX Vector Convert Unsigned Fixed-Point Word to Single-Precision
XX2	60	F00002A4	VSX	xvrspip	VSX Vector Round to Single-Precision Integer toward +Infinity
XX2	60	F00002A8	VSX	xvtsqrtsp	VSX Vector Test for Software Square Root Single-Precision
XX2	60	F00002AC	VSX	xvrspic	VSX Vector Round to Single-Precision Integer using Current Rounding Mode
XX3	60	F00002C0	VSX	xdivsp	VSX Vector Divide Single-Precision
XX3	60	F00002C8	VSX	xvmsubmsp	VSX Vector Multiply-Subtract Type-M Single-Precision
XX2	60	F00002E0	VSX	xvcvsxwsp	VSX Vector Convert Signed Fixed-Point Word to Single-Precision



Table C-1. POWER8 Instructions by Opcode (Sheet 19 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
XX2	60	F00002E4	VSX	xvrspim	VSX Vector Round to Single-Precision Integer toward -Infinity
XX3	60	F00002E8	VSX	xvtdivsp	VSX Vector Test for Software Divide Single-Precision
XX3	60	F0000300	VSX	xvadddp	VSX Vector Add Double-Precision
XX3	60	F0000308	VSX	xvmaddadp	VSX Vector Multiply-Add Type-A Double-Precision
XX3	60	F0000318	VSX	xvcmpqdp	VSX Vector Compare Equal To Double-Precision
XX2	60	F0000320	VSX	xvcvdpxws	VSX Vector Convert Double-Precision to Unsigned Fixed-Point Word Saturate
XX2	60	F0000324	VSX	xvrdpi	VSX Vector Round to Double-Precision Integer
XX2	60	F0000328	VSX	xvrsqrtdp	VSX Vector Reciprocal Square Root Estimate Double-Precision
XX2	60	F000032C	VSX	xvsqrtdp	VSX Vector Square Root Double-Precision
XX3	60	F0000340	VSX	xvsubdp	VSX Vector Subtract Double-Precision
XX3	60	F0000348	VSX	xvmaddmdp	VSX Vector Multiply-Add Type-M Double-Precision
XX3	60	F0000358	VSX	xvcmpgtdp	VSX Vector Compare Greater Than Double-Precision
XX2	60	F0000360	VSX	xvcvdpsxws	VSX Vector Convert Double-Precision to Signed Fixed-Point Word Saturate
XX2	60	F0000364	VSX	xvrdpiz	VSX Vector Round to Double-Precision Integer toward Zero
XX2	60	F0000368	VSX	xvredp	VSX Vector Reciprocal Estimate Double-Precision
XX3	60	F0000380	VSX	xvmuldp	VSX Vector Multiply Double-Precision
XX3	60	F0000388	VSX	xvmsubadp	VSX Vector Multiply-Subtract Type-A Double-Precision
XX3	60	F0000398	VSX	xvcmpgedp	VSX Vector Compare Greater Than or Equal To Double-Precision
XX2	60	F00003A0	VSX	xvcvuxwdp	VSX Vector Convert Unsigned Fixed-Point Word to Double-Precision
XX2	60	F00003A4	VSX	xvrdpip	VSX Vector Round to Double-Precision Integer toward +Infinity
XX2	60	F00003A8	VSX	xvtsqrtdp	VSX Vector Test for Software Square Root Double-Precision
XX2	60	F00003AC	VSX	xvrdpic	VSX Vector Round to Double-Precision Integer using Current Rounding Mode
XX3	60	F00003C0	VSX	xvdivdp	VSX Vector Divide Double-Precision
XX3	60	F00003C8	VSX	xvmsubmdp	VSX Vector Multiply-Subtract Type-M Double-Precision
XX2	60	F00003E0	VSX	xvcvsxwdp	VSX Vector Convert Signed Fixed-Point Word to Double-Precision
XX2	60	F00003E4	VSX	xvrdpim	VSX Vector Round to Double-Precision Integer toward -Infinity
XX3	60	F00003E8	VSX	xvtdivdp	VSX Vector Test for Software Divide Double-Precision
XX3	60	F0000408	VSX	xsnmaddasp	VSX Scalar Negative Multiply-Add Type-A Single-Precision
XX3	60	F0000410	VSX	xxland	VSX Logical AND
XX2	60	F0000424	VSX	xscvdsp	VSX Scalar Convert Double-Precision to Single-Precision



Table C-1. POWER8 Instructions by Opcode (Sheet 20 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
XX2	60	F000042C	VSX	xscvdpspn	VSX Scalar Convert Double-Precision to Single-Precision format Non-signalling
XX3	60	F0000448	VSX	xsnmaddmsp	VSX Scalar Negative Multiply-Add Type-M Single-Precision
XX3	60	F0000450	VSX	xxlandc	VSX Logical AND with Complement
XX2	60	F0000464	VSX	xsrsp	VSX Scalar Round to Single-Precision
XX3	60	F0000488	VSX	xsnmsubasp	VSX Scalar Negative Multiply-Subtract Type-A Single-Precision
XX3	60	F0000490	VSX	xxlor	VSX Logical OR
XX2	60	F00004A0	VSX	xscvuxdsp	VSX Scalar Convert Unsigned Fixed-Point Doubleword to Single-Precision
XX3	60	F00004C8	VSX	xsnmsubmsp	VSX Scalar Negative Multiply-Subtract Type-M Single-Precision
XX3	60	F00004D0	VSX	xxlxor	VSX Logical XOR
XX2	60	F00004E0	VSX	xscvsxdsp	VSX Scalar Convert Signed Fixed-Point Doubleword to Single-Precision
XX3	60	F0000500	VSX	xsmaxdp	VSX Scalar Maximum Double-Precision
XX3	60	F0000508	VSX	xsnmaddadp	VSX Scalar Negative Multiply-Add Type-A Double-Precision
XX3	60	F0000510	VSX	xxlnor	VSX Logical NOR
XX2	60	F0000520	VSX	xscvdpxuds	VSX Scalar Convert Double-Precision to Unsigned Fixed-Point Doubleword Saturate
XX2	60	F0000524	VSX	xscvspdp	VSX Scalar Convert Single-Precision to Double-Precision (p=1)
XX2	60	F000052C	VSX	xscvspdpn	Scalar Convert Single-Precision to Double-Precision format Non-signalling
XX3	60	F0000540	VSX	xsmindp	VSX Scalar Minimum Double-Precision
XX3	60	F0000548	VSX	xsnmaddmdp	VSX Scalar Negative Multiply-Add Type-M Double-Precision
XX3	60	F0000550	VSX	xxlorc	VSX Logical OR with Complement
XX2	60	F0000560	VSX	xscvdpsxds	VSX Scalar Convert Double-Precision to Signed Fixed-Point Doubleword Saturate
XX2	60	F0000564	VSX	xsabsdp	VSX Scalar Absolute Value Double-Precision
XX3	60	F0000580	VSX	xscpsgndp	VSX Scalar Copy Sign Double-Precision
XX3	60	F0000588	VSX	xsnmsubadp	VSX Scalar Negative Multiply-Subtract Type-A Double-Precision
XX3	60	F0000590	VSX	xxlnand	VSX Logical NAND
XX2	60	F00005A0	VSX	xscvuxddp	VSX Scalar Convert Unsigned Fixed-Point Doubleword to Double-Precision
XX2	60	F00005A4	VSX	xsnabsdp	VSX Scalar Negative Absolute Value Double-Precision
XX3	60	F00005C8	VSX	xsnmsubmdp	VSX Scalar Negative Multiply-Subtract Type-M Double-Precision
XX3	60	F00005D0	VSX	xxleqv	VSX Logical Equivalence
XX2	60	F00005E0	VSX	xscvsxddp	VSX Scalar Convert Signed Fixed-Point Doubleword to Double-Precision



Table C-1. POWER8 Instructions by Opcode (Sheet 21 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
XX2	60	F00005E4	VSX	xsnegdp	VSX Scalar Negate Double-Precision
XX3	60	F0000600	VSX	xvmaxsp	VSX Vector Maximum Single-Precision
XX3	60	F0000608	VSX	xvnaddasp	VSX Vector Negative Multiply-Add Type-A Single-Precision
XX3	60	F0000618	VSX	xvccmpeqsp.	VSX Vector Compare Equal To Single-Precision and Record CR6
XX2	60	F0000620	VSX	xvcvspuds	VSX Vector Convert Single-Precision to Unsigned Fixed-Point Doubleword Saturate
XX2	60	F0000624	VSX	xvcvdpsp	VSX Vector Convert Double-Precision to Single-Precision
XX3	60	F0000640	VSX	xvminsp	VSX Vector Minimum Single-Precision
XX3	60	F0000648	VSX	xvnaddmsp	VSX Vector Negative Multiply-Add Type-M Single-Precision
XX3	60	F0000658	VSX	xvcmpgtsp.	VSX Vector Compare Greater Than Single-Precision and Record CR6
XX2	60	F0000660	VSX	xvcvspxsds	VSX Vector Convert Single-Precision to Signed Fixed-Point Doubleword Saturate
XX2	60	F0000664	VSX	xvabssp	VSX Vector Absolute Value Single-Precision
XX3	60	F0000680	VSX	xvcpsgnsp	VSX Vector Copy Sign Single-Precision
XX3	60	F0000688	VSX	xvnmsubasp	VSX Vector Negative Multiply-Subtract Type-A Single-Precision
XX3	60	F0000698	VSX	xvcmpgesp.	VSX Vector Compare Greater Than or Equal To Single-Precision and Record CR6
XX2	60	F00006A0	VSX	xvcvuxdsp	VSX Vector Convert Unsigned Fixed-Point Doubleword to Single-Precision
XX2	60	F00006A4	VSX	xvnabssp	VSX Vector Negative Absolute Value Single-Precision
XX3	60	F00006C8	VSX	xvnmsubmsp	VSX Vector Negative Multiply-Subtract Type-M Single-Precision
XX2	60	F00006E0	VSX	xvcvsxdsp	VSX Vector Convert Signed Fixed-Point Doubleword to Single-Precision
XX2	60	F00006E4	VSX	xvnegsp	VSX Vector Negate Single-Precision
XX3	60	F0000700	VSX	xvmaxdp	VSX Vector Maximum Double-Precision
XX3	60	F0000708	VSX	xvnaddadp	VSX Vector Negative Multiply-Add Type-A Double-Precision
XX3	60	F0000718	VSX	xvcmeqdp.	VSX Vector Compare Equal To Double-Precision and Record CR6
XX2	60	F0000720	VSX	xvcvdpuuds	VSX Vector Convert Double-Precision to Unsigned Fixed-Point Doubleword Saturate
XX2	60	F0000724	VSX	xvcvspdp	VSX Vector Convert Single-Precision to Double-Precision
XX3	60	F0000740	VSX	xvmindp	VSX Vector Minimum Double-Precision
XX3	60	F0000748	VSX	xvnaddmdp	VSX Vector Negative Multiply-Add Type-M Double-Precision
XX3	60	F0000758	VSX	xvcmpgtdp.	VSX Vector Compare Greater Than Double-Precision and Record CR6
XX2	60	F0000760	VSX	xvcvdpsxsds	VSX Vector Convert Double-Precision to Signed Fixed-Point Doubleword Saturate



Table C-1. POWER8 Instructions by Opcode (Sheet 22 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
XX2	60	F0000764	VSX	xvabsdp	VSX Vector Absolute Value Double-Precision
XX3	60	F0000780	VSX	xvcpsgndp	VSX Vector Copy Sign Double-Precision
XX3	60	F0000788	VSX	xvnmsubadp	VSX Vector Negative Multiply-Subtract Type-A Double-Precision
XX3	60	F0000798	VSX	xvcmpgedp.	VSX Vector Compare Greater Than or Equal To Double-Precision and Record CR6
XX2	60	F00007A0	VSX	xvcvuxddp	VSX Vector Convert Unsigned Fixed-Point Doubleword to Double-Precision
XX2	60	F00007A4	VSX	xvnabsdp	VSX Vector Negative Absolute Value Double-Precision
XX3	60	F00007C8	VSX	xvnmsubmdp	VSX Vector Negative Multiply-Subtract Type-M Double-Precision
XX2	60	F00007E0	VSX	xvcvsxddp	VSX Vector Convert Signed Fixed-Point Doubleword to Double-Precision
XX2	60	F00007E4	VSX	xvnegdp	VSX Vector Negate Double-Precision
DS	61	F4000000	FP.out	stfdp	Store Floating-Point Double Pair
DS	62	F8000002	LSQ	stq	Store Quadword
DS	62	F8000000	64	std	Store Doubleword
DS	62	F8000001	64	stdu	Store Doubleword with Update
X	63	FC000000	FP	fcmpu	Floating Compare Unordered
X	63	FC000004	DFP	daddq[.]	Decimal Floating Add Quad
Z23	63	FC000006	DFP	dquaq[.]	Decimal Quantize Quad
X	63	FC000010	FP[R]	fcpsgn[.]	Floating Copy Sign
X	63	FC000018	FP[R]	frsp[.]	Floating Round to Single-Precision
X	63	FC00001C	FP[R]	fctiw[.]	Floating Convert To Integer Word
X	63	FC00001E	FP[R]	fctiwz[.]	Floating Convert To Integer Word with round to Zero
A	63	FC000024	FP[R]	fdiv[.]	Floating Divide
A	63	FC000028	FP[R]	fsub[.]	Floating Subtract
A	63	FC00002A	FP[R]	fadd[.]	Floating Add
A	63	FC00002C	FP[R]	fsqrt[.]	Floating Square Root
A	63	FC00002E	FP[R]	fsel[.]	Floating Select
A	63	FC000030	FP[R].in	fre[.]	Floating Reciprocal Estimate
A	63	FC000032	FP[R]	fmul[.]	Floating Multiply
A	63	FC000034	FP[R]	frsqrte[.]	Floating Reciprocal Square Root Estimate
A	63	FC000038	FP[R]	fmsub[.]	Floating Multiply-Subtract
A	63	FC00003A	FP[R]	fmadd[.]	Floating Multiply-Add
A	63	FC00003C	FP[R]	fnmsub[.]	Floating Negative Multiply-Subtract
A	63	FC00003E	FP[R]	fnmadd[.]	Floating Negative Multiply-Add



Table C-1. POWER8 Instructions by Opcode (Sheet 23 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
X	63	FC000040	FP	fcmwo	Floating Compare Ordered
X	63	FC000044	DFP	dmulq[.]	Decimal Floating Multiply Quad
Z23	63	FC000046	DFP	drrndq[.]	Decimal Floating Reround Quad
X	63	FC00004C	FP[R]	mtfsb1[.]	Move To FPSCR Bit 1
X	63	FC000050	FP[R]	fneg[.]	Floating Negate
X	63	FC000080	FP	mcrfs	Move To Condition Register from FPSCR
Z22	63	FC000084	DFP	dscliq[.]	Decimal Floating Shift Coefficient Left Immediate Quad
Z23	63	FC000086	DFP	dquaiq[.]	Decimal Quantize Immediate Quad
X	63	FC00008C	FP[R]	mtfsb0[.]	Move To FPSCR Bit 0
X	63	FC000090	FP[R]	fmr[.]	Floating Move Register
Z22	63	FC0000C4	DFP	dscrq[.]	Decimal Floating Shift Coefficient Right Immediate Quad
Z23	63	FC0000C6	DFP	drintxq[.]	Decimal Floating Round To FP Integer With Inexact Quad
X	63	FC000100	FP	ftdiv	Floating Test for Software Divide
X	63	FC000104	DFP	dcmpoq	Decimal Floating Compare Ordered Quad
X	63	FC00010C	FP[R]	mtfsfi[.]	Move To FPSCR Field Immediate
X	63	FC000110	FP[R]	fnabs[.]	Floating Negative Absolute Value
X	63	FC00011C	FP[R]	fctiwu[.]	Floating Convert To Integer Word Unsigned
X	63	FC00011E	FP[R]	fctiwuz[.]	Floating Convert To Integer Word Unsigned with Round Toward Zero
X	63	FC000140	FP	ftsqrt	Floating Test for Software Square Root
X	63	FC000144	DFP	dtstexq	Decimal Floating Test Exponent Quad
Z22	63	FC000184	DFP	dtstdcq	Decimal Floating Test Data Class Quad
Z22	63	FC0001C4	DFP	dtstdgq	Decimal Floating Test Data Group Quad
Z23	63	FC0001C6	DFP	drintnq[.]	Decimal Floating Round To FP Integer Without Inexact Quad
X	63	FC000204	DFP	dctqpq[.]	Decimal Floating Convert To DFP Extended
X	63	FC000210	FP[R]	fabs[.]	Floating Absolute Value
X	63	FC000244	DFP	dctfixq[.]	Decimal Floating Convert To Fixed Quad
X	63	FC000284	DFP	ddedpdq[.]	Decimal Floating Decode DPD To BCD Quad
X	63	FC0002C4	DFP	dxexq[.]	Decimal Floating Extract Exponent Quad
X	63	FC000310	FP[R].in	frin[.]	Floating Round To Integer Nearest
X	63	FC000350	FP[R].in	friz[.]	Floating Round To Integer toward Zero
X	63	FC000390	FP[R].in	frip[.]	Floating Round To Integer Plus
X	63	FC0003D0	FP[R].in	frim[.]	Floating Round To Integer Minus
X	63	FC000404	DFP	dsubq[.]	Decimal Floating Subtract Quad
X	63	FC000444	DFP	ddivq[.]	Decimal Floating Divide Quad



Table C-1. POWER8 Instructions by Opcode (Sheet 24 of 24)

Format	Opcode		Category	Mnemonic	Instruction
	Primary (Decimal)	Instruction Image (operands set to 0's) (Hexadecimal)			
X	63	FC00048E	FP[R]	mffs[.]	Move From FPSCR
X	63	FC000504	DFP	dcmpuq	Decimal Floating Compare Unordered Quad
X	63	FC000544	DFP	dtstsfq	Decimal Floating Test Significance Quad
XFL	63	FC00058E	FP[R]	mttsf[.]	Move To FPSCR Fields
X	63	FC000604	DFP	drdpq[.]	Decimal Floating Round To DFP Long
X	63	FC000644	DFP	dcffixq[.]	Decimal Floating Convert From Fixed Quad
X	63	FC00065C	FP[R]	fctid[.]	Floating Convert To Integer Doubleword
X	63	FC00065E	FP[R]	fctidz[.]	Floating Convert To Integer Doubleword with Round Toward Zero
X	63	FC000684	DFP	denbcdq[.]	Decimal Floating Encode BCD To DPD Quad
X	63	FC00068C	VSX	fmrgow	Floating Merge Odd Word
X	63	FC00069C	FP[R]	fcfid[.]	Floating Convert From Integer Doubleword
X	63	FC0006C4	DFP	diexq[.]	Decimal Floating Insert Exponent Quad
X	63	FC00075C	FP[R]	fctidu[.]	Floating Convert To Integer Doubleword Unsigned
X	63	FC00075E	FP[R]	fctiduz[.]	Floating Convert To Integer Doubleword Unsigned with Round Toward Zero
X	63	FC00078C	VSX	fmrgew	Floating Merge Even Word
X	63	FC00079C	FP[R]	fcfidu[.]	Floating Convert From Integer Doubleword Unsigned



Appendix D. Performance Monitoring Events

The POWER8 processor has a built-in performance monitoring unit (PMU) for each hardware thread that provides instrumentation to aid in performance monitoring, workload characterization, system characterization and code analysis.

There are six thread-level performance monitor counters (PMC) in a PMU. PMC1 – PMC4 are programmable, PMC5 counts nonidle completed instructions, and PMC6 counts nonidle cycles.

The thread-level and core-level instrumentation have access to a set of performance events that cover essential statistics such as:

- Miss rates
- Unit utilization
- Thread balance
- Hazard conditions
- Translation-related misses
- Stall analysis
- Instruction mix
- L1 I-cache and D-cache reload source
- Effective cache counts
- Memory latency counts

This document covers all of the performance monitoring events supported by the POWER8 PMU. These events can be measured using hpmcount (AIX) and perf (Linux).

Table D-1. on page 388 lists all of the performance monitoring events for the POWER8 processor in alphabetical order by event name. Each event entry includes the event code used to program the PMU and a short description.

Table D-1. POWER8 Event List by Event Name (Sheet 1 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
000001F05E				PM_1LPAR_CYC	Number of cycles in single LPAR mode. All threads in the core are assigned to the same LPAR.
00000100F2				PM_1PLUS_PPC_CMPL	One or more PowerPC instructions finished.
		00000400F2		PM_1PLUS_PPC_DISP	Number of cycles that at least one instruction is dispatched. A group can be only microcode.
	000002006E			PM_2LPAR_CYC	Number of cycles in 2 LPAR mode. Threads 0 - 3 belong to LPAR0 and threads 4 - 7 belong to LPAR1.
		000004E05E		PM_4LPAR_CYC	Number of cycles in 4 LPAR mode. Threads 0 - 1 belong to LPAR0.
0000610050				PM_ALL_CHIP_PUMP_CPRED	Initial and final pump scope was chip pump (prediction is correct) for all data types (demand load).
	0000520050			PM_ALL_GRP_PUMP_CPRED	Initial and final pump scope and data sourced across this scope was group pump for all data types (demand load).
	0000620052			PM_ALL_GRP_PUMP_MPRED	Final pump scope (group) was either larger or smaller than initial pump scope for all data types (demand load).
0000610052				PM_ALL_GRP_PUMP_MPRED_RTY	Final pump scope (group) was larger than the initial pump scope (chip) for all data types (demand load).
0000610054				PM_ALL_PUMP_CPRED	Pump prediction correct. Counts are across all types of pumps for all data types (demand load).
		0000640052		PM_ALL_PUMP_MPRED	Pump misprediction. Counts are across all types of pumps for all data types (demand load).
		0000630050		PM_ALL_SYS_PUMP_CPRED	Initial and final pump scope was system pump for all data types (demand load).
		0000630052		PM_ALL_SYS_PUMP_MPRED	Final pump scope (system) mispredicted. Either the original scope was too small (chip/group) or the original scope was system and it should have been smaller. Counts are for all data types (demand load).
			0000640050	PM_ALL_SYS_PUMP_MPRED_RTY	Final pump scope (system) was larger than initial pump scope (chip/group) for all data types (demand load).
00000100FA				PM_ANY_THRD_RUN_CYC	One or more threads (or at least one thread) is in run_cycles.
	000002505E			PM_BACK_BR_CMPL	Branch instruction completed with a target address less than the current instruction address.
0000004082	0000004082	0000004082	0000004082	PM_BANK_CONFLICT	Read blocked due to interleave conflict. The ifar logic detects an interleave conflict and kills the data that was read in that cycle.
	0000020036		0000040036	PM_BR_2PATH	Two-path branch.
0000005086	0000005086	0000005086	0000005086	PM_BR_BC_8	Pairable BC+8 branch that has not been converted to a Resolve Finished in the BRU pipeline.
0000005084	0000005084	0000005084	0000005084	PM_BR_BC_8_CONV	Pairable BC+8 branch that was converted to a resolve finished in the BRU pipeline.
			0000040060	PM_BR_CMPL	Branch instruction completed.

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 2 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
00000040AC	00000040AC	00000040AC	00000040AC	PM_BR_MPRED_CCACHE	Conditional branch completed that was mispredicted due to the count cache target prediction.
			00000400F6	PM_BR_MPRED_CMPL	Number of branch mispredicts.
00000040B8	00000040B8	00000040B8	00000040B8	PM_BR_MPRED_CR	Conditional branch completed that was mispredicted due to the BHT direction prediction (taken/not taken).
00000040AE	00000040AE	00000040AE	00000040AE	PM_BR_MPRED_LSTACK	Conditional branch completed that was mispredicted due to the link stack target prediction.
00000040BA	00000040BA	00000040BA	00000040BA	PM_BR_MPRED_TA	Conditional branch completed that was mispredicted due to the target address prediction from the count cache or link stack. Only XL-form branches that resolved taken set this event.
0000010138			0000040138	PM_BR_MRK_2PATH	Marked 2-path branch.
000000489C	000000489C	000000489C	000000489C	PM_BR_PRED_BR_CMPL	Completion time event.
000000409C	000000409C	000000409C	000000409C	PM_BR_PRED_BR0	Conditional branch completed on BR0 (first branch in the group) in which the hardware predicted the direction or target.
000000409E	000000409E	000000409E	000000409E	PM_BR_PRED_BR1	Conditional branch completed on BR1 (second branch in the group) in which the hardware predicted the direction or target. Note: BR1 can only be used in single-thread mode in all of the <u>SMT</u> modes.
00000040A4	00000040A4	00000040A4	00000040A4	PM_BR_PRED_CCACHE_BR0	Conditional branch completed on BR0 that used the count cache for target prediction.
00000040A6	00000040A6	00000040A6	00000040A6	PM_BR_PRED_CCACHE_BR1	Conditional branch completed on BR1 that used the count cache for target prediction.
00000048A4	00000048A4	00000048A4	00000048A4	PM_BR_PRED_CCACHE_CMPL	Completion time event.
00000040B0	00000040B0	00000040B0	00000040B0	PM_BR_PRED_CR_BR0	Conditional branch completed on BR0 that had its direction predicted. I-form branches do not set this event.
00000040B2	00000040B2	00000040B2	00000040B2	PM_BR_PRED_CR_BR1	Conditional branch completed on BR1 that had its direction predicted. I-form branches do not set this event.
00000048B0	00000048B0	00000048B0	00000048B0	PM_BR_PRED_CR_CMPL	Completion time event.
00000040A8	00000040A8	00000040A8	00000040A8	PM_BR_PRED_LSTACK_BR0	Conditional branch completed on BR0 that used the link stack for target prediction.
00000040AA	00000040AA	00000040AA	00000040AA	PM_BR_PRED_LSTACK_BR1	Conditional branch completed on BR1 that used the link stack for target prediction.
00000048A8	00000048A8	00000048A8	00000048A8	PM_BR_PRED_LSTACK_CMPL	Completion time event.
00000040B4	00000040B4	00000040B4	00000040B4	PM_BR_PRED_TA_BR0	Conditional branch completed on BR0 that had its target address predicted. Only XL-form branches set this event.
00000040B6	00000040B6	00000040B6	00000040B6	PM_BR_PRED_TA_BR1	Conditional branch completed on BR1 that had its target address predicted. Only XL-form branches set this event.
00000048B4	00000048B4	00000048B4	00000048B4	PM_BR_PRED_TA_CMPL	Completion time event.
	00000200FA			PM_BR_TAKEN_CMPL	New event for branch taken.

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 3 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
00000040A0	00000040A0	00000040A0	00000040A0	PM_BR_UNCOND_BR0	Unconditional branch completed on BR0. Hardware branch prediction was not used for this branch. This can be an I-form branch.
00000040A2	00000040A2	00000040A2	00000040A2	PM_BR_UNCOND_BR1	Unconditional branch completed on BR1. Hardware branch prediction was not used for this branch. This can be an I-form branch.
00000048A0	00000048A0	00000048A0	00000048A0	PM_BR_UNCOND_CMPL	Completion time event. This event can also be calculated from the direct bus as follows: if_pc_br0_br_pred = '00' and if_pc_br0_completed.
0000010068				PM_BRU_FIN	Branch instruction finished.
0000003094	0000003094	0000003094	0000003094	PM_CASTOUT_ISSUED	Register file castout issued for both GPRs and VRs.
0000003096	0000003096	0000003096	0000003096	PM_CASTOUT_ISSUED_GPR	Register file castout issued for GPRs only.
0000010050				PM_CHIP_PUMP_CPRED	Initial and final pump scope was chip pump (prediction is correct) for all data types excluding data prefetch (demand load).
0000002090	0000002090	0000002090	0000002090	PM_CLB_HELD	CLB hold: any reason.
000001E054				PM_CMPLU_STALL ¹	Completion stall. No groups completed.
		000004000A			
		000004D018		PM_CMPLU_STALL_BRU	Completion stall due to a branch unit.
	000002D018			PM_CMPLU_STALL_BRU_CRU	Completion stall due to an IFU instruction.
		0000030026		PM_CMPLU_STALL_COQ_FULL	Completion stall due to CO queue full.
	000002C012			PM_CMPLU_STALL_DCACHE_MISS	Completion stall due to a D-cache miss.
	000002C018			PM_CMPLU_STALL_DMISS_L21_L31	Completion stall due to a D-cache miss that resolves on chip (excluding the local L2 or L3 cache).
	000002C016			PM_CMPLU_STALL_DMISS_L2L3	Completion stall due to a D-cache miss that resolves in the L2 or L3 cache.
		000004C016		PM_CMPLU_STALL_DMISS_L2L3_CONFLICT	Completion stall due to a cache miss that resolves in the L2 or L3 cache with a conflict.
		000004C01A		PM_CMPLU_STALL_DMISS_L3MISS	Completion stall due to a cache miss that resolves in the L3 cache.
		000004C018		PM_CMPLU_STALL_DMISS_LMEM	Completion stall due to a cache miss that resolves in the core's local memory.
	000002C01C			PM_CMPLU_STALL_DMISS_REMOTE	Completion stall due to a D-cache miss that resolves on chip (excluding the local L2 or L3 cache).
		000004C012		PM_CMPLU_STALL_ERAT_MISS	Completion stall due to an LSU reject ERAT miss.
		0000030038		PM_CMPLU_STALL_FLUSH	Completion stall due to flush by own thread.
		000004D016		PM_CMPLU_STALL_FXLONG	Completion stall due to a long latency fixed-point instruction.
	000002D016			PM_CMPLU_STALL_FXU	Completion stall due to an FXU instruction.

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 4 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
		0000030036		PM_CMPLU_STALL_HWSYNC	Completion stall due to an hwsync .
		000004D014		PM_CMPLU_STALL_LOAD_FINISH	Completion stall due to a load finish.
	000002C010			PM_CMPLU_STALL_LSU	Completion stall due to an LSU instruction.
0000010036				PM_CMPLU_STALL_LWSYNC	Completion stall due to an isync/lwsync .
		0000030028		PM_CMPLU_STALL_MEM_ECC_DELAY	Completion stall due to a memory ECC delay.
	000002E01C			PM_CMPLU_STALL_NO_NTF	Completion stall due to a NOP.
	000002E01E			PM_CMPLU_STALL_NTCG_FLUSH	Completion stall because the NTC instruction was flushed.
		0000030006		PM_CMPLU_STALL_OTHER_CMPL	Number of instructions completed by all other threads while this thread was stalled.
			000004C014	PM_CMPLU_STALL_REJ_LMQ_FULL	Completion stall due to an LSU reject because the LMQ is full.
			000004C010	PM_CMPLU_STALL_REJECT	Completion stall due to an LSU reject.
	000002C01A			PM_CMPLU_STALL_REJECT_LHS	Completion stall due to a reject (load-hit-store).
			000004D010	PM_CMPLU_STALL_SCALAR	Completion stall due to a VSU scalar instruction.
	000002D010			PM_CMPLU_STALL_SCALAR_LONG	Completion stall due to a VSU scalar long-latency instruction.
			000004C01C	PM_CMPLU_STALL_ST_FWD	Completion stall due to a store forward.
	000002C014			PM_CMPLU_STALL_STORE	Completion stall by store instructions, which include store agen finishes in pipe LS0 or LS1 and store data finishes in the LS2 or LS3.
000001001C				PM_CMPLU_STALL_THRD	Completion stalled because of thread conflict. Group is ready to complete but it was another thread's turn.
	000002D014			PM_CMPLU_STALL_VECTOR	Completion stall due to a VSU vector instruction.
			000004D012	PM_CMPLU_STALL_VECTOR_LONG	Completion stall due to a VSU vector long instruction.
	000002D012			PM_CMPLU_STALL_VSU	Completion stall due to a VSU instruction.
0000517082				PM_CO_DISP_FAIL	CO dispatch failed due to all CO machines being busy.
	0000527084			PM_CO_TM_SC_FOOTPRINT	L2 did a clean-if-dirty CO to the L3 cache (for example, created a store clean (SC) line in the L3 cache).
		000003608A		PM_CO_USAGE ¹	Continuous 16-cycle (2:1) window where this signal rotates through sampling each L2 CO machine busy. The PMU uses this wave to then do a 16-cycle count to sample the total number of machines running.
		000073608A			
0000016083				PM_CO0_ALLOC	CO mach 0 busy. Used by the PMU to sample average RC lifetime (mach0 used as sample point).

1. There are multiple methods to program this event.

**Table D-1. POWER8 Event List by Event Name (Sheet 5 of 46)**

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
0000016082				PM_CO0_BUSY ¹	CO machine 0 busy. Used by the PMU to sample average RC lifetime (mach0 used as a sample point).
0000716082				PM_CRU_FIN	IFU finished a nonbranch instruction.
000001001E	000002001E	000003001E	000004001E	PM_CYC ¹	Cycles.
00000100F0				PM_DATA_ALL_CHIP_PUMP_CPRED	Initial and final pump scope was chip pump (prediction is correct) for a demand load or a data prefetch.
000061C050				PM_DATA_ALL_FROM_DL2L3_MOD	The processor's data cache was reloaded with modified (M) data from another chip's L2 or L3 cache on a different Node or Group (distant).
			000064C048	PM_DATA_ALL_FROM_DL2L3_SHR	The processor's data cache was reloaded with shared (S) data from another chip's L2 or L3 cache on a different Node or Group (distant).
			000064C04C	PM_DATA_ALL_FROM_DMEM	The processor's data cache was reloaded from another chip's memory on the same Node or Group (distant) due to a demand load or a data prefetch.
000061C042				PM_DATA_ALL_FROM_L2	The processor's data cache was reloaded from the local core's L2 cache due to a demand load or a data prefetch.
		000063C040		PM_DATA_ALL_FROM_L2_DISP_CONFLICT_LDHTST	The processor's data cache was reloaded from the local core's L2 cache with a load-hit-store conflict due to either demand loads or a data prefetch.
			000064C040	PM_DATA_ALL_FROM_L2_DISP_CONFLICT_OTHER	The processor's data cache was reloaded from the local core's L2 cache with a dispatch conflict due to a demand load or a data prefetch.
000062C040				PM_DATA_ALL_FROM_L2_MEFP	The processor's data cache was reloaded from the local core's L2 hit without a dispatch conflict on an Mefp state due to a demand load or a data prefetch.
000061C040				PM_DATA_ALL_FROM_L2_NO_CONFLICT	The processor's data cache was reloaded from the local core's L2 cache without conflict due to a demand load or a data prefetch.
			000064C046	PM_DATA_ALL_FROM_L21_MOD	The processor's data cache was reloaded with modified (M) data from another core's L2 cache on the same chip due to a demand load or a data prefetch.
			000063C046	PM_DATA_ALL_FROM_L21_SHR	The processor's data cache was reloaded with shared (S) data from another core's L2 cache on the same chip due to a demand load or a data prefetch.
000061C04E				PM_DATA_ALL_FROM_L2MISS_MOD	The processor's data cache was reloaded from a location other than the local core's L2 cache due to a demand load or a data prefetch.
			000064C042	PM_DATA_ALL_FROM_L3	The processor's data cache was reloaded from the local core's L3 cache due to a demand load or a data prefetch.
			000063C042	PM_DATA_ALL_FROM_L3_DISP_CONFLICT	The processor's data cache was reloaded from the local core's L3 cache with a dispatch conflict due to a demand load or a data prefetch.
	000062C042			PM_DATA_ALL_FROM_L3_MEFP	The processor's data cache was reloaded from the local core's L3 cache without dispatch conflict hits on an Mefp state due to a demand load or a data prefetch.

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 6 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
000061C044				PM_DATA_ALL_FROM_L3_NO_CONFLICT	The processor's data cache was reloaded from the local core's L3 cache without conflict due to a demand load or a data prefetch.
			000064C044	PM_DATA_ALL_FROM_L31_ECO_MOD	The processor's data cache was reloaded with modified (M) data from another core's ECO L3 on the same chip due to a demand load or a data prefetch.
		000063C044		PM_DATA_ALL_FROM_L31_ECO_SHR	The processor's data cache was reloaded with shared (S) data from another core's ECO L3 on the same chip due to a demand load or a data prefetch.
	000062C044			PM_DATA_ALL_FROM_L31_MOD	The processor's data cache was reloaded with modified (M) data from another core's L3 cache on the same chip due to a demand load or a data prefetch.
000061C046				PM_DATA_ALL_FROM_L31_SHR	The processor's data cache was reloaded with shared (S) data from another core's L3 cache on the same chip due to a demand load or a data prefetch.
			000064C04E	PM_DATA_ALL_FROM_L3MISS_MOD	The processor's data cache was reloaded from a location other than the local core's L3 cache due to either a demand load s or a data prefetch.
	000062C048			PM_DATA_ALL_FROM_LMEM	The processor's data cache was reloaded from the local chip's memory due to a demand load or a data prefetch.
	000062C04C			PM_DATA_ALL_FROM_MEMORY	The processor's data cache was reloaded from a memory location from a local remote or distant due to a demand load or a data prefetch.
			000064C04A	PM_DATA_ALL_FROM_OFF_CHIP_CACHE	The processor's data cache was reloaded with either shared or modified data from another core's L2 or L3 cache on a different chip (remote or distant) due to a demand load or a data prefetch.
000061C048				PM_DATA_ALL_FROM_ON_CHIP_CACHE	The processor's data cache was reloaded with either shared or modified data from another core's L2 or L3 cache on the same chip due to a demand load or a data prefetch.
	000062C046			PM_DATA_ALL_FROM_RL2L3_MOD	The processor's data cache was reloaded with modified (M) data from another chip's L2 or L3 cache on the same Node or Group (remote).
000061C04A				PM_DATA_ALL_FROM_RL2L3_SHR	The processor's data cache was reloaded with shared (S) data from another chip's L2 or L3 cache on the same Node or Group (remote).
		000063C04A		PM_DATA_ALL_FROM_RMEM	The processor's data cache was reloaded from another chip's memory on the same Node or Group (remote) due to a demand load or a data prefetch.
	000062C050			PM_DATA_ALL_GRP_PUMP_CPRED	Initial and final pump scope was group pump (prediction is correct) for a demand load or a data prefetch.
	000062C052			PM_DATA_ALL_GRP_PUMP_MPRED	Final pump scope (group) was either larger or smaller than initial pump scope for a demand load or a data prefetch. Final pump scope (group) to get data sourced.
000061C052				PM_DATA_ALL_GRP_PUMP_MPRED_RTY	Final pump scope (group) was larger than the initial pump scope (chip) for a demand load or a data prefetch.
000061C054				PM_DATA_ALL_PUMP_CPRED	Pump prediction correct. Counts across all types of pumps for a demand load or a data prefetch.
			000064C052	PM_DATA_ALL_PUMP_MPRED	Pump misprediction. Counts are across all types of pumps for a demand load or a data prefetch.

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 7 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
		000063C050		PM_DATA_ALL_SYS_PUMP_CPRED	Initial and final pump scope was system pump (prediction = correct) for a demand load or a data prefetch.
		000063C052		PM_DATA_ALL_SYS_PUMP_MPRED	Final pump scope (system) mispredicted. Either the original scope was too small (chip/group) or the original scope was system and it should have been smaller. Counts are for a demand load or a data prefetch. Final pump scope (system) to get data sourced.
		000064C050	PM_DATA_ALL_SYS_PUMP_MPRED_RTY	Final pump scope (system) was larger than initial pump scope (chip/group) for a demand load or a data prefetch. Final pump scope (system) to get data sourced.	
000001C050				PM_DATA_CHIP_PUMP_CPRED	Initial and final pump scope was chip pump (prediction is correct) for a demand load.
		000004C048	PM_DATA_FROM_DL2L3_MOD	The processor's data cache was reloaded with modified (M) data from another chip's L2 or L3 cache on a different Node or Group (distant).	
		000003C048	PM_DATA_FROM_DL2L3_SHR	The processor's data cache was reloaded with shared (S) data from another chip's L2 or L3 cache on a different Node or Group (distant).	
		000004C04C	PM_DATA_FROM_DMEM	The processor's data cache was reloaded from another chip's memory on the same Node or Group (distant) due to either only demand loads or demand loads plus prefetches if MMCR1[16] = '1'.	
000001C042				PM_DATA_FROM_L2	The processor's data cache was reloaded from the local core's L2 cache due to either a demand load or demand loads plus prefetches if MMCR1[16] = '1'.
		000003C040	PM_DATA_FROM_L2_DISP_CONFLICT_LDHTST	The processor's data cache was reloaded from the local core's L2 cache with a load-hit-store conflict due to either only demand loads or demand loads plus prefetches if MMCR1[16] = '1'.	
		000004C040	PM_DATA_FROM_L2_DISP_CONFLICT_OTHER	The processor's data cache was reloaded from the local core's L2 cache with a dispatch conflict due to either only demand loads or demand loads plus prefetches if MMCR1[16] = '1'.	
000002C040				PM_DATA_FROM_L2_MEPF	The processor's data cache was reloaded from a local core's L2 hit without dispatch conflicts on an Mpf state due to a demand load or demand loads plus prefetches if MMCR1[16] = '1'.
000001C040				PM_DATA_FROM_L2_NO_CONFLICT	The processor's data cache was reloaded from the local core's L2 cache without conflict due to either demand loads or demand loads plus prefetches if MMCR1[16] = '1'.
		000004C046	PM_DATA_FROM_L21_MOD	The processor's data cache was reloaded with modified (M) data from another core's L2 on the same chip due to either only demand loads or demand loads plus prefetches if MMCR1[16] = '1'.	
		000003C046	PM_DATA_FROM_L21_SHR	The processor's data cache was reloaded with shared (S) data from another core's L2 cache on the same chip due to either only demand loads or demand loads plus prefetches if MMCR1[16] = '1'.	
	00000200FE			PM_DATA_FROM_L2MISS	Demand LD - L2 miss (not an L2 hit). The processor's data cache was reloaded but not from the local L2 cache.
000001C04E				PM_DATA_FROM_L2MISS_MOD	The processor's data cache was reloaded from a location other than the local core's L2 cache due to either only demand loads or demand loads plus prefetches if MMCR1[16] = '1'.
		000004C042	PM_DATA_FROM_L3	The processor's data cache was reloaded from the local core's L3 cache due to either only demand loads or demand loads plus prefetches if MMCR1[16] = '1'.	
		000003C042	PM_DATA_FROM_L3_DISP_CONFLICT	The processor's data cache was reloaded from the local core's L3 cache with dispatch conflict due to either only demand loads or demand loads plus prefetches if MMCR1[16] = '1'.	

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 8 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
	000002C042			PM_DATA_FROM_L3_MEFP	The processor's data cache was reloaded from a local core's L3 cache without dispatch conflicts hit on an Mfp state due to a demand load or demand loads plus prefetches if MMCR1[16] = '1'.
000001C044				PM_DATA_FROM_L3_NO_CONFLICT	The processor's data cache was reloaded from the local core's L3 cache without conflict due to a demand load or demand load plus prefetch if MMCR1[16] = '1'.
		000004C044	PM_DATA_FROM_L31_ECO_MOD	The processor's data cache was reloaded with modified (M) data from another core's ECO L3 on the same chip due to either only demand loads or demand loads plus prefetches if MMCR1[16] = '1'.	
	000003C044			PM_DATA_FROM_L31_ECO_SHR	The processor's data cache was reloaded with shared (S) data from another core's ECO L3 on the same chip due to either only a demand load or demand loads plus prefetches if MMCR1[16] = '1'.
	000002C044			PM_DATA_FROM_L31_MOD	The processor's data cache was reloaded with modified (M) data from another core's L3 cache on the same chip due to a demand load or demand loads plus prefetches if MMCR1[16] = '1'.
000001C046				PM_DATA_FROM_L31_SHR	The processor's data cache was reloaded with shared (S) data from another core's L3 cache on the same chip due to either only demand loads or demand loads plus prefetches if MMCR1[16] = '1'.
	00000300FE			PM_DATA_FROM_L3MISS	Demand load - L3 miss (not an L2 hit and not an L3 hit). The processor's data cache was reloaded but not from the local L3 cache.
		000004C04E	PM_DATA_FROM_L3MISS_MOD		The processor's data cache was reloaded from a location other than the local core's L3 cache due to either only demand loads or demand loads plus prefetches if MMCR1[16] = '1'.
	000002C048			PM_DATA_FROM_LMEM	The processor's data cache was reloaded from the local chip's memory due to a demand load or demand loads plus prefetches if MMCR1[16] = '1'.
		00000400FE	PM_DATA_FROM_MEM		The processor's data cache was reloaded from a memory location including from a local remote or distant due to a demand load.
	000002C04C			PM_DATA_FROM_MEMORY	The processor's data cache was reloaded from a memory location from local remote or distant due to only demand loads or demand loads plus prefetches if MMCR1[16] = '1'.
		000004C04A	PM_DATA_FROM_OFF_CHIP_CACHE		The processor's data cache was reloaded with either shared or modified data from another core's L2 or L3 cache on a different chip (remote or distant) due to either only demand loads or demand loads plus prefetches if MMCR1[16] = '1'.
000001C048				PM_DATA_FROM_ON_CHIP_CACHE	The processor's data cache was reloaded with either shared or modified data from another core's L2/L3 cache on the same chip due to either only demand loads or demand loads plus prefetches if MMCR1[16] = '1'.
	000002C046			PM_DATA_FROM_RL2L3_MOD	The processor's data cache was reloaded with modified (M) data from another chip's L2 or L3 cache on the same Node or Group (remote).
000001C04A				PM_DATA_FROM_RL2L3_SHR	The processor's data cache was reloaded with Shared (S) data from another chip's L2 or L3 cache on the same Node or Group (remote).
	000003C04A			PM_DATA_FROM_RMEM	The processor's data cache was reloaded from another chip's memory on the same Node or Group (remote) due to either only a demand load or demand loads plus prefetches if MMCR1[16] = '1'.

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 9 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
	000002C050			PM_DATA_GRP_PUMP_CPRED	Initial and final pump scope and data sourced across this scope was group pump (prediction is correct) for a demand load.
	000002C052			PM_DATA_GRP_PUMP_MPRED	Final pump scope (group) was either larger or smaller than initial pump scope for a demand load. Final pump scope (group) to get data sourced.
000001C052				PM_DATA_GRP_PUMP_MPRED_RTY	Final pump scope (Group) was larger than initial pump scope (chip) for a demand load.
000001C054				PM_DATA_PUMP_CPRED	Pump prediction correct. Counts across all types of pumps for a demand load.
		000004C052		PM_DATA_PUMP_MPRED	Pump misprediction. Counts across all types of pumps for a demand load.
		000003C050		PM_DATA_SYS_PUMP_CPRED	Initial and final pump scope and data sourced across this scope was system pump (prediction is correct) for a demand load.
		000003C052		PM_DATA_SYS_PUMP_MPRED	Final Pump Scope (system) mispredicted. Either the original scope was too small (Chip/Group) or the original scope was System and it should have been smaller. Counts for a demand load.
			000004C050	PM_DATA_SYS_PUMP_MPRED_RTY	Final pump scope (system) was larger than initial pump scope (chip/group) for a demand load.
		000003001A		PM_DATA_TABLEWALK_CYC	Tablewalk cycles (one or two cycles are active).
000000E0BC	000000E0BC	000000E0BC	000000E0BC	PM_DC_COLLISIONS	Data cache collisions.
000001E050				PM_DC_PREF_STREAM_ALLOC	Stream marked valid. The stream could have been allocated through the hardware prefetch mechanism or through software. This is combined LS0 and LS1.
	000002E050			PM_DC_PREF_STREAM_CONF	A demand load referenced a line in an active prefetch stream. The stream might have been allocated through the hardware prefetch mechanism or through software. Combine up and down.
		000004E050		PM_DC_PREF_STREAM_FUZZY_CONF	A demand load referenced a line in an active fuzzy prefetch stream. The stream might have been allocated through the hardware prefetch mechanism or through software.
		000003E050		PM_DC_PREF_STREAM_STRIDED_CONF	A demand load referenced a line in an active strided prefetch stream. The stream might have been allocated through the hardware prefetch mechanism or through software.
		000004C054		PM_DERAT_MISS_16G	Data ERAT miss (data TLB access) page size 16 GB.
		000003C054		PM_DERAT_MISS_16M	Data ERAT miss (data TLB access) page size 16 MB.
000001C056				PM_DERAT_MISS_4K	Data ERAT miss (data TLB access) page size 4 KB.
	000002C054			PM_DERAT_MISS_64K	Data ERAT miss (data TLB access) page size 64 KB.
000000B0BA	000000B0BA	000000B0BA	000000B0BA	PM_DFU	Finish DFU (all finish).
000000B0BE	000000B0BE	000000B0BE	000000B0BE	PM_DFU_DCFFIX	Convert from fixed-opcode finish (dcffix).
000000B0BC	000000B0BC	000000B0BC	000000B0BC	PM_DFU_DENBCD	<u>BCD-to-DPD</u> opcode finish (denbcd).
000000B0B8	000000B0B8	000000B0B8	000000B0B8	PM_DFU_MC	Finish DFU multicycle.
0000002092	0000002092	0000002092	0000002092	PM_DISP_CLB_HELD_BAL	Dispatch/CLB hold: balance.

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 10 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
0000002094	0000002094	0000002094	0000002094	PM_DISP_CLB_HELD_RES	Dispatch/CLB held: resource.
00000020A8	00000020A8	00000020A8	00000020A8	PM_DISP_CLB_HELD_SB	Dispatch/CLB held: scoreboard.
0000002098	0000002098	0000002098	0000002098	PM_DISP_CLB_HELD_SYNC	Dispatch/CLB held: sync type instruction. An instruction group is held at dispatch in the IBuffer due to a sync type instruction.
0000002096	0000002096	0000002096	0000002096	PM_DISP_CLB_HELD_TLBIE	An instruction group is held at dispatch in the IBuffer due to a tlbie instruction.
0000010006				PM_DISP_HELD	Dispatch held. The number of cycles that the dispatch to pipeline is held.
	0000020006			PM_DISP_HELD_IQ_FULL	Dispatch held because the issue queue is full.
000001002A				PM_DISP_HELD_MAP_FULL	Dispatch for this thread was held because the mappers are full.
	0000030018			PM_DISP_HELD_SRQ_FULL	Dispatch held because the SRQ has no space.
		000004003C		PM_DISP_HELD_SYNC_HOLD	Dispatch held due to a sync hold.
00000030A6	00000030A6	00000030A6	00000030A6	PM_DISP_HOLD_GCT_FULL	Dispatch held due to no space in the GCT.
	0000030008			PM_DISP_WT	Dispatched starved (not held).
		000004E048		PM_DPTEG_FROM_DL2L3_MOD	A PTE was loaded into the TLB with modified (M) data from another chip's L2 or L3 cache on a different Node or Group (distant).
		000003E048		PM_DPTEG_FROM_DL2L3_SHR	A PTE was loaded into the TLB with shared (S) data from another chip's L2 or L3 cache on a different Node or Group (distant).
		000004E04C		PM_DPTEG_FROM_DMEM	A PTE was loaded into the TLB from another chip's memory on the same Node or Group (distant) due to a data-side request.
000001E042				PM_DPTEG_FROM_L2	A PTE was loaded into the TLB from a local core's L2 cache due to a data-side request.
	000003E040			PM_DPTEG_FROM_L2_DISP_CONFLICT_LDHTST	A PTE was loaded into the TLB from the local core's L2 cache with a load-hit-store conflict due to a data-side request.
		000004E040		PM_DPTEG_FROM_L2_DISP_CONFLICT_OTHER	A PTE was loaded into the TLB from the local core's L2 cache with a dispatch conflict due to a data-side request.
000002E040				PM_DPTEG_FROM_L2_MEFP	A PTE was loaded into the TLB from the local core's L2 hit without a dispatch conflict on an Mefp state due to a data-side request.
000001E040				PM_DPTEG_FROM_L2_NO_CONFLICT	A PTE was loaded into the TLB from a local core's L2 cache without a conflict due to a data-side request.
		000004E046		PM_DPTEG_FROM_L21_MOD	A PTE was loaded into the TLB with modified (M) data from another core's L2 cache on the same chip due to a data-side request.
		000003E046		PM_DPTEG_FROM_L21_SHR	A PTE was loaded into the TLB with shared (S) data from another core's L2 cache on the same chip due to a data-side request.
000001E04E				PM_DPTEG_FROM_L2MISS	A PTE was loaded into the TLB from a location other than the local core's L2 due to a data-side request.
		000004E042		PM_DPTEG_FROM_L3	A PTE was loaded into the TLB from local core's L3 cache due to a data-side request.

1. There are multiple methods to program this event.

Table D-1. POWER8 Event List by Event Name (Sheet 11 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
		000003E042		PM_DPTEG_FROM_L3_DISP_CONFLICT	A PTE was loaded into the TLB from the local core's L3 cache with a dispatch conflict due to a data-side request.
	000002E042			PM_DPTEG_FROM_L3_MEPF	A PTE was loaded into the TLB from the local core's L3 cache without a dispatch conflict hit on an Mepf state due to a data-side request.
000001E044				PM_DPTEG_FROM_L3_NO_CONFLICT	A PTE was loaded into the TLB from local core's L3 cache without conflict due to a data-side request.
		000004E044		PM_DPTEG_FROM_L31_ECO_MOD	A PTE was loaded into the TLB with modified (M) data from another core's ECO L3 on the same chip due to a data-side request.
		000003E044		PM_DPTEG_FROM_L31_ECO_SHR	A PTE was loaded into the TLB with shared (S) data from another core's ECO L3 on the same chip due to a data-side request.
	000002E044			PM_DPTEG_FROM_L31_MOD	A PTE was loaded into the TLB with modified (M) data from another core's L3 cache on the same chip due to a data-side request.
000001E046				PM_DPTEG_FROM_L31_SHR	A PTE was loaded into the TLB with shared (S) data from another core's L3 cache on the same chip due to a data-side request.
		000004E04E		PM_DPTEG_FROM_L3MISS	A PTE was loaded into the TLB from a location other than the local core's L3 cache due to a data-side request.
	000002E048			PM_DPTEG_FROM_LMEM	A PTE was loaded into the TLB from the local chip's memory due to a data-side request.
	000002E04C			PM_DPTEG_FROM_MEMORY	A PTE was loaded into the TLB from a memory location from a local remote or distant due to a data-side request.
		000004E04A		PM_DPTEG_FROM_OFF_CHIP_CACHE	A PTE was loaded into the TLB either shared or modified data from another core's L2 or L3 cache on a different chip (remote or distant) due to a data-side request.
000001E048				PM_DPTEG_FROM_ON_CHIP_CACHE	A PTE was loaded into the TLB with either shared or modified data from another core's L2 or L3 cache on the same chip due to a data-side request.
	000002E046			PM_DPTEG_FROM_RL2L3_MOD	A PTE was loaded into the TLB with modified (M) data from another chip's L2 or L3 cache on the same Node or Group (remote).
000001E04A				PM_DPTEG_FROM_RL2L3_SHR	A PTE was loaded into the TLB with shared (S) data from another chip's L2 or L3 cache on the same Node or Group (remote).
		000003E04A		PM_DPTEG_FROM_RMEM	A PTE was loaded into the TLB from another chip's memory on the same Node or Group (remote) due to a data-side request.
0000010016				PM_DSLB_MISS ¹	Data SLB miss; total of all segment sizes. An SLB miss for a data request occurred. SLB misses trap to the operating system to resolve. This is a total count for all segment sizes.
000000D094	000000D094	000000D094	000000D094		
		00000300FC			
000001C058					
		000004C056		PM_DTLB_MISS_16G	Data TLB miss page size 16 GB.
		000002C056		PM_DTLB_MISS_16M	Data TLB miss page size 16 MB.
				PM_DTLB_MISS_4K	Data TLB miss page size 4 KB.

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 12 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
		000003C056		PM_DTLB_MISS_64K	Data TLB miss page size 64 KB.
00000050A8	00000050A8	00000050A8	00000050A8	PM_EAT_FORCE_MISPRED	The XL-form branch was mispredicted due to the predicted target address missing from the EAT. The EAT forces a mispredict in this case because there is no predicated target to validate. This rare case can occur when the EAT is full and a branch is issued.
0000004084	0000004084	0000004084	0000004084	PM_EAT_FULL_CYC	Cycles the thread was blocked because the EAT was full.
0000002080	0000002080	0000002080	0000002080	PM_EE_OFF_EXT_INT	EE off and external interrupt. Cycles when an interrupt due to an external exception is pending but external exceptions were masked.
	00000200F8			PM_EXT_INT	External interrupt.
00000020B4	00000020B4	00000020B4	00000020B4	PM_FAU_TBEGIN	Dispatch time favored tbegin .
00000100F4				PM_FLOP	Floating-point operation finished.
000000A0AE	000000A0AE	000000A0AE	000000A0AE	PM_FLOP_SUM_SCALAR	Flops summary scalar instructions.
000000A0AC	000000A0AC	000000A0AC	000000A0AC	PM_FLOP_SUM_VEC	Flops summary vector instructions.
		00000400F8		PM_FLUSH	Flush (any type). A flush has occurred. This includes all types of flushes such as branch mispredicts, load store unit flushes, partial flushes, and completion flushes.
0000002084	0000002084	0000002084	0000002084	PM_FLUSH_BR_MPRED	A flush was caused by a branch mispredict.
	0000030012			PM_FLUSH_COMPLETION	Completion flush.
0000002082	0000002082	0000002082	0000002082	PM_FLUSH_DISP	A dispatch flush occurred.
000000208C	000000208C	000000208C	000000208C	PM_FLUSH_DISP_SB	Dispatch flush: scoreboard. An instruction with the scoreboard bit set caused a dispatch flush to flush the decode pipe.
0000002088	0000002088	0000002088	0000002088	PM_FLUSH_DISP_SYNC	Dispatch flush: sync . A sync/lwsync/ptesync/tlbsync caused a dispatch flush to flush the decode pipe.
000000208A	000000208A	000000208A	000000208A	PM_FLUSH_DISP_TLBIE	Dispatch flush: tlbie . A tlbie caused a dispatch flush to flush the decode pipe.
000000208E	000000208E	000000208E	000000208E	PM_FLUSH_LSU	Flush initiated by the LSU.
0000002086	0000002086	0000002086	0000002086	PM_FLUSH_PARTIAL	Partial flush.
000000A0B0	000000A0B0	000000A0B0	000000A0B0	PM_FPU0_FCONV	Convert instruction executed.
000000A0B8	000000A0B8	000000A0B8	000000A0B8	PM_FPU0_FEST	Estimate instruction executed.
000000A0B4	000000A0B4	000000A0B4	000000A0B4	PM_FPU0_FRSP	Round to single-precision instruction executed.
000000A0B2	000000A0B2	000000A0B2	000000A0B2	PM_FPU1_FCONV	Convert instruction executed.
000000A0BA	000000A0BA	000000A0BA	000000A0BA	PM_FPU1_FEST	Estimate instruction executed.
000000A0B6	000000A0B6	000000A0B6	000000A0B6	PM_FPU1_FRSP	Round to single-precision instruction executed.

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 13 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
		000003000C		PM_FREQ_DOWN	Power management is below threshold B. Frequency is slewed down due to power management.
			000004000C	PM_FREQ_UP	Frequency is being slewed up due to power management.
00000050B0	00000050B0	00000050B0	00000050B0	PM_FUSION_TOC_GRP0_1	One pair of instructions fused with the TOC in Group0.
00000050AE	00000050AE	00000050AE	00000050AE	PM_FUSION_TOC_GRP0_2	Two pairs of instructions fused with the TOC in Group0.
00000050AC	00000050AC	00000050AC	00000050AC	PM_FUSION_TOC_GRP0_3	Three pairs of instructions fused with the TOC in Group0.
00000050B2	00000050B2	00000050B2	00000050B2	PM_FUSION_TOC_GRP1_1	One pair of instructions fused with the TOC in Group1.
00000050B8	00000050B8	00000050B8	00000050B8	PM_FUSION_VSX_GRP0_1	One pair of instructions fused with the VSX in Group0.
00000050B6	00000050B6	00000050B6	00000050B6	PM_FUSION_VSX_GRP0_2	Two pairs of instructions fused with the VSX in Group0.
00000050B4	00000050B4	00000050B4	00000050B4	PM_FUSION_VSX_GRP0_3	Three pairs of instructions fused with VSX in Group0.
00000050BA	00000050BA	00000050BA	00000050BA	PM_FUSION_VSX_GRP1_1	One pair of instructions fused with the VSX in Group1.
	000002000E			PM_FXU_BUSY	Cycles when both fixed point units (<u>FXU0</u> and FXU1) are busy.
000001000E				PM_FXU_IDLE	FXU0 is idle and FXU1 is idle.
		000003000E		PM_FXU0_BUSY_FXU1_IDLE	FXU0 is busy and FXU1 is idle.
0000010004				PM_FXU0_FIN	FXU0 finished an instruction. Instructions that finish might not complete.
			000004000E	PM_FXU1_BUSY_FXU0_IDLE	Number of cycles when FXU0 is idle and FXU1 is busy.
			0000040004	PM_FXU1_FIN	FXU1 finished an instruction. Instructions that finish might not complete.
	0000020008			PM_GCT_EMPTY_CYC	No itags assigned with either thread (GCT is empty).
00000030A4	00000030A4	00000030A4	00000030A4	PM_GCT_MERGE	Group dispatched on a merged GCT empty. GCT entries can be merged only within the same thread.
			000004D01E	PM_GCT_NOSLOT_BR_MPRED	GCT is empty for this thread due to branch mispredict.
			000004D01A	PM_GCT_NOSLOT_BR_MPRED_ICMISS	GCT is empty for this thread due to an l-cache miss and branch mispredicted.
00000100F8				PM_GCT_NOSLOT_CYC	Pipeline is empty. No itags assigned.
	000002D01E			PM_GCT_NOSLOT_DISP_HELD_ISSQ	GCT is empty for this thread because a dispatch hold on this thread is due to the issue queue being full.
			000004D01C	PM_GCT_NOSLOT_DISP_HELD_MAP	GCT is empty for this thread due to a dispatch hold on this thread due to the mapper being full.
	000002E010			PM_GCT_NOSLOT_DISP_HELD_OTHER	GCT is empty for this thread because a dispatch hold on this thread is due to a sync instruction.

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 14 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
	000002D01C			PM_GCT_NOSLOT_DISP_HELD_SRQ	GCT is empty for this thread because a dispatch hold on this thread is due to the SRQ being full.
			000004E010	PM_GCT_NOSLOT_IC_L3MISS	GCT is empty for this thread due to an I-cache L3 miss.
	000002D01A			PM_GCT_NOSLOT_IC_MISS	GCT is empty for this thread due to an I-cache miss.
000000209C	000000209C	000000209C	000000209C	PM_GCT_UTIL_1_2_ENTRIES	GCT utilization 1 - 2 entries.
00000020A2	00000020A2	00000020A2	00000020A2	PM_GCT_UTIL_11_14_ENTRIES	GCT utilization 11 - 14 entries.
00000020A4	00000020A4	00000020A4	00000020A4	PM_GCT_UTIL_15_17_ENTRIES	GCT utilization 15 - 17 entries.
00000020A6	00000020A6	00000020A6	00000020A6	PM_GCT_UTIL_18_ENTRIES	GCT utilization 18+ entries.
000000209E	000000209E	000000209E	000000209E	PM_GCT_UTIL_3_6_ENTRIES	GCT utilization 3 - 6 entries.
00000020A0	00000020A0	00000020A0	00000020A0	PM_GCT_UTIL_7_10_ENTRIES	GCT utilization 7 - 10 entries.
000001000A				PM_GRP_BR_MPRED_NONSPEC	Group experienced nonspeculative branch redirect (mispredict).
	0000030004			PM_GRP_CMPL	Group completed.
	000003000A			PM_GRP_DISP	Dispatch success; group dispatched.
000001000C				PM_GRP_IC_MISS_NONSPEC	Group experienced nonspeculative I-cache miss.
0000010130				PM_GRP_MRK	Instruction marked in the IDU.
000000509C	000000509C	000000509C	000000509C	PM_GRP_NON_FULL_GROUP	Number of groups that dispatched with less than six nonbranch instructions (ST mode).
	0000020050			PM_GRP_PUMP_CPRED	Initial and final pump scope and data sourced across this scope was group pump for all data types excluding data prefetch (demand load).
	0000020052			PM_GRP_PUMP_MPRED	Final pump scope (group) was either larger or smaller than initial pump scope for all data types excluding data prefetch (demand load).
0000010052				PM_GRP_PUMP_MPRED_RTY	Final pump scope (Group) was larger than initial pump scope (chip) for all data types excluding data prefetch (demand load).
00000050A4	00000050A4	00000050A4	00000050A4	PM_GRP_TERM_2ND_BRANCH	There were enough instructions in the IBuffer.
00000050A6	00000050A6	00000050A6	00000050A6	PM_GRP_TERM_FPU_AFTER_BR	There were enough instructions in the IBuffer.
000000509E	000000509E	000000509E	000000509E	PM_GRP_TERM_NOINST	Do not fill every slot in the group.
00000050A0	00000050A0	00000050A0	00000050A0	PM_GRP_TERM_OTHER	There were enough instructions in the IBuffer.
00000050A2	00000050A2	00000050A2	00000050A2	PM_GRP_TERM_SLOT_LIMIT	There were enough instructions in the IBuffer.
	000002000A			PM_HV_CYC	Cycles in which MSR[HV] is high (cycles in hypervisor mode). Note that this event does not take MSR[PR] into consideration.
0000004086	0000004086	0000004086	0000004086	PM_IBUF_FULL_CYC	Number of cycles that the thread was blocked because the instruction fetch buffer was full.

1. There are multiple methods to program this event.

**Table D-1. POWER8 Event List by Event Name (Sheet 15 of 46)**

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
0000010018				PM_IC_DEMAND_CYC	Number of cycles when a demand instruction fetch is pending.
0000004098	0000004098	0000004098	0000004098	PM_IC_DEMAND_L2_BHT_REDIRECT	An instruction cache demand request was sent to the L2 cache because of a <u>BHT</u> redirect.
000000409A	000000409A	000000409A	000000409A	PM_IC_DEMAND_L2_BR_REDIRECT	An instruction cache demand request was sent to the L2 cache because of a branch mispredict (15-cycle path).
0000004088	0000004088	0000004088	0000004088	PM_IC_DEMAND_REQ	Demand instruction fetch request.
000000508A	000000508A	000000508A	000000508A	PM_IC_INVALIDATE	Instruction cache line invalidated.
0000004092	0000004092	0000004092	0000004092	PM_IC_PREF_CANCEL_HIT	Prefetch canceled due to an I-cache hit.
0000004094	0000004094	0000004094	0000004094	PM_IC_PREF_CANCEL_L2	The L2 squashed either a demand or prefetch request that was no longer required by the instruction fetch unit.
0000004090	0000004090	0000004090	0000004090	PM_IC_PREF_CANCEL_PAGE	Prefetch canceled due to page boundary. A request to prefetch a line to the I-cache was canceled because of a page boundary crossing.
000000408A	000000408A	000000408A	000000408A	PM_IC_PREF_REQ	Instruction prefetch request.
000000408E	000000408E	000000408E	000000408E	PM_IC_PREF_WRITE	Instruction prefetch written into the instruction L1 cache.
0000004096	0000004096	0000004096	0000004096	PM_IC_RELOAD_PRIVATE	Reloading line was brought in private for a specific thread. Most lines are brought in shared for all eight threads. If the <u>RA</u> does not match the invalidates, bring it shared to the other thread. For the POWER8 processor, the line is brought in private.
00000100F6				PM_IERAT_RELOAD	Number of I-ERAT reloads.
		000004006A		PM_IERAT_RELOAD_16M	IERAT reloaded (miss) for a 16 MB page.
	0000020064			PM_IERAT_RELOAD_4K	IERAT reloaded (miss) for a 4 KB page.
		000003006A		PM_IERAT_RELOAD_64K	IERAT reloaded (miss) for a 64 KB page.
		000003405E		PM_IFETCH_THROTTLE	Number of cycles that the instruction fetch throttle was active in the IFU.
0000005088	0000005088	0000005088	0000005088	PM_IFU_L2_TOUCH	L2 touch to update MRU on a line.
0000514050				PM_INST_ALL_CHIP_PUMP_CPRED	Initial and final pump scope was chip pump (prediction is correct) for instruction fetches and prefetches.
		0000544048		PM_INST_ALL_FROM_DL2L3_MOD	The processor's instruction cache was reloaded with modified (M) data from another chip's L2 or L3 cache on a different Node or Group (distant).
		0000534048		PM_INST_ALL_FROM_DL2L3_SHR	The processor's instruction cache was reloaded with shared (S) data from another chip's L2 or L3 cache on a different Node or Group (distant).
		000054404C		PM_INST_ALL_FROM_DMEM	The processor's instruction cache was reloaded from another chip's memory on the same Node or Group (distant) due to instruction fetches and prefetches.
0000514042				PM_INST_ALL_FROM_L2	The processor's instruction cache was reloaded from the local core's L2 cache due to instruction fetches and prefetches.

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 16 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
		0000534040		PM_INST_ALL_FROM_L2_DISP_CONFLICT_LDHTST	The processor's instruction cache was reloaded from the local core's L2 cache with a load-hit-store conflict due to instruction fetches and prefetches.
			0000544040	PM_INST_ALL_FROM_L2_DISP_CONFLICT_OTHER	The processor's instruction cache was reloaded from the local core's L2 cache with a dispatch conflict due to instruction fetches and prefetches.
	0000524040			PM_INST_ALL_FROM_L2_MEFP	The processor's instruction cache was reloaded from the local core's L2 hit without a dispatch conflict on an Mfp state due to instruction fetches and prefetches.
0000514040				PM_INST_ALL_FROM_L2_NO_CONFLICT	The processor's instruction cache was reloaded from local core's L2 cache without a conflict due to instruction fetches and prefetches.
			0000544046	PM_INST_ALL_FROM_L21_MOD	The processor's instruction cache was reloaded with modified (M) data from another core's L2 cache on the same chip due to instruction fetches and prefetches.
		0000534046		PM_INST_ALL_FROM_L21_SHR	The processor's instruction cache was reloaded with shared (S) data from another core's L2 cache on the same chip due to instruction fetches and prefetches.
000051404E				PM_INST_ALL_FROM_L2MISS	The processor's instruction cache was reloaded from a location other than the local core's L2 cache due to instruction fetches and prefetches.
			0000544042	PM_INST_ALL_FROM_L3	The processor's instruction cache was reloaded from the local core's L3 cache due to instruction fetches and prefetches.
		0000534042		PM_INST_ALL_FROM_L3_DISP_CONFLICT	The processor's instruction cache was reloaded from the local core's L3 cache with a dispatch conflict due to instruction fetches and prefetches.
	0000524042			PM_INST_ALL_FROM_L3_MEFP	The processor's instruction cache was reloaded from the local core's L3 cache without a dispatch conflict hit on an Mfp state due to instruction fetches and prefetches.
0000514044				PM_INST_ALL_FROM_L3_NO_CONFLICT	The processor's instruction cache was reloaded from local core's L3 cache without a conflict due to instruction fetches and prefetches.
			0000544044	PM_INST_ALL_FROM_L31_ECO_MOD	The processor's instruction cache was reloaded with modified (M) data from another core's ECO L3 on the same chip due to instruction fetches and prefetches.
			0000534044	PM_INST_ALL_FROM_L31_ECO_SHR	The processor's instruction cache was reloaded with shared (S) data from another core's ECO L3 on the same chip due to instruction fetches and prefetches.
	0000524044			PM_INST_ALL_FROM_L31_MOD	The processor's instruction cache was reloaded with modified (M) data from another core's L3 cache on the same chip due to instruction fetches and prefetches.
0000514046				PM_INST_ALL_FROM_L31_SHR	The processor's instruction cache was reloaded with shared (S) data from another core's L3 cache on the same chip due to instruction fetches and prefetches.
			000054404E	PM_INST_ALL_FROM_L3MISS_MOD	The processor's instruction cache was reloaded from a location other than the local core's L3 cache due to a instruction fetch.
	0000524048			PM_INST_ALL_FROM_LMEM	The processor's instruction cache was reloaded from the local chip's memory due to instruction fetches and prefetches.
	000052404C			PM_INST_ALL_FROM_MEMORY	The processor's instruction cache was reloaded from a memory location from local remote or distant due to instruction fetches and prefetches.

1. There are multiple methods to program this event.

**Table D-1. POWER8 Event List by Event Name (Sheet 17 of 46)**

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
			000054404A	PM_INST_ALL_FROM_OFF_CHIP_CACHE	The processor's instruction cache was reloaded with either shared or modified data from another core's L2 or L3 cache on a different chip (remote or distant) due to instruction fetches and prefetches.
0000514048				PM_INST_ALL_FROM_ON_CHIP_CACHE	The processor's instruction cache was reloaded with either shared or modified data from another core's L2 or L3 cache on the same chip due to instruction fetches and prefetches.
	0000524046			PM_INST_ALL_FROM_RL2L3_MOD	The processor's instruction cache was reloaded with modified (M) data from another chip's L2 or L3 cache on the same Node or Group (remote).
000051404A				PM_INST_ALL_FROM_RL2L3_SHR	The processor's instruction cache was reloaded with shared (S) data from another chip's L2 or L3 cache on the same Node or Group (remote).
		000053404A		PM_INST_ALL_FROM_RMEM	The processor's instruction cache was reloaded from another chip's memory on the same Node or Group (remote) due to instruction fetches and prefetches.
	0000524050			PM_INST_ALL_GRP_PUMP_CPRED	Initial and final pump scope was group pump (prediction = correct) for instruction fetches and prefetches.
	0000524052			PM_INST_ALL_GRP_PUMP_MPRED	Final pump scope (group) was either larger or smaller than initial pump scope for instruction fetches and prefetches. Final pump scope (group) to get data sourced.
0000514052				PM_INST_ALL_GRP_PUMP_MPRED_RTY	Final pump scope (group) was larger than the initial pump scope (chip) for instruction fetches and prefetches.
0000514054				PM_INST_ALL_PUMP_CPRED	Pump prediction correct. Counts are across all types of pumps for instruction fetches and prefetches.
		0000544052		PM_INST_ALL_PUMP_MPRED	Pump misprediction. Counts are across all types of pumps for instruction fetches and prefetches.
		0000534050		PM_INST_ALL_SYS_PUMP_CPRED	Initial and final pump scope was system pump (prediction = correct) for instruction fetches and prefetches. Initial and final pump scope and data sourced across this scope was system pump for an instruction fetch.
		0000534052		PM_INST_ALL_SYS_PUMP_MPRED	Final pump scope (system) mispredicted. Either the original scope was too small (chip/group) or the original scope was system and it should have been smaller. Counts are for instruction fetches and prefetches. Final pump scope (system) to get data sourced.
			0000544050	PM_INST_ALL_SYS_PUMP_MPRED_RTY	Final pump scope (system) was larger than initial pump scope (chip/group) for instruction fetches and prefetches. Final pump scope (system) to get data sourced.
0000014050				PM_INST_CHIP_PUMP_CPRED	Initial and final pump scope was chip pump (prediction is correct) for an instruction fetch.
0000010002	0000020002	0000030002	0000040002	PM_INST_CMPL	Number of PowerPC instructions that completed.
	00000200F2	00000300F2		PM_INST_DISP	Number of PowerPC instructions that were successfully dispatched.
			0000044048	PM_INST_FROM_DL2L3_MOD	The processor's instruction cache was reloaded with modified (M) data from another chip's L2 or L3 cache on a different Node or Group (distant).
			0000034048	PM_INST_FROM_DL2L3_SHR	The processor's instruction cache was reloaded with shared (S) data from another chip's L2 or L3 cache on a different Node or Group (distant).

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 18 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
			000004404C	PM_INST_FROM_DMEM	The processor's instruction cache was reloaded from another chip's memory on the same Node or Group (distant) due to an instruction fetch (not prefetch) or an instruction fetch plus prefetch if MMCR1[17] = '1'.
0000004080	0000004080	0000004080	0000004080	PM_INST_FROM_L1	Instruction fetches from L1 cache.
0000014042				PM_INST_FROM_L2	The processor's instruction cache was reloaded from local core's L2 cache due to an instruction fetch (not prefetch) or an instruction fetch plus prefetch if MMCR1[17] = '1'.
		0000034040		PM_INST_FROM_L2_DISP_CONFLICT_LDHITST	The processor's instruction cache was reloaded from the local core's L2 cache with a load-hit-store conflict due to an instruction fetch (not prefetch) or an instruction fetch plus prefetch if MMCR1[17] = '1'.
			0000044040	PM_INST_FROM_L2_DISP_CONFLICT_OTHER	The processor's instruction cache was reloaded from the local core's L2 cache with dispatch conflict due to an instruction fetch (not prefetch) or an instruction fetch plus prefetch if MMCR1[17] = '1'.
	0000024040			PM_INST_FROM_L2_MEFP	The processor's instruction cache was reloaded from a local core's L2 hit without dispatch conflicts on an Mefp state, due to either an instruction fetch (not prefetch) or an instruction fetch plus prefetch if MMCR1[17] = '1'.
0000014040				PM_INST_FROM_L2_NO_CONFLICT	The processor's instruction cache was reloaded from local core's L2 cache without conflict due to an instruction fetch (not prefetch) or an instruction fetch plus prefetch if MMCR1[17] = '1'.
			0000044046	PM_INST_FROM_L21_MOD	The processor's instruction cache was reloaded with modified (M) data from another core's L2 cache on the same chip due to an instruction fetch (not prefetch) or an instruction fetch plus prefetch if MMCR1[17] = '1'.
		0000034046		PM_INST_FROM_L21_SHR	The processor's instruction cache was reloaded with shared (S) data from another core's L2 cache on the same chip due to an instruction fetch (not prefetch) or an instruction fetch plus prefetch if MMCR1[17] = '1'.
000001404E				PM_INST_FROM_L2MISS	The processor's instruction cache was reloaded from a location other than the local core's L2 cache due to an instruction fetch (not prefetch) or an instruction fetch plus prefetch if MMCR1[17] = '1'.
			0000044042	PM_INST_FROM_L3	The processor's instruction cache was reloaded from the local core's L3 cache due to an instruction fetch (not prefetch) or an instruction fetch plus prefetch if MMCR1[17] = '1'.
		0000034042		PM_INST_FROM_L3_DISP_CONFLICT	The processor's instruction cache was reloaded from the local core's L3 cache with dispatch conflict due to an instruction fetch (not prefetch) or an instruction fetch plus prefetch if MMCR1[17] = '1'.
	0000024042			PM_INST_FROM_L3_MEFP	The processor's Instruction cache was reloaded from the local core's L3 cache without dispatch conflicts hit on an Mefp state, due to either an instruction fetch (not prefetch) or an instruction fetch plus prefetch if MMCR1[17] = '1'.
0000014044				PM_INST_FROM_L3_NO_CONFLICT	The processor's instruction cache was reloaded from local core's L3 cache without conflict due to an instruction fetch (not prefetch) or an instruction fetch plus prefetch if MMCR1[17] = '1'.
			0000044044	PM_INST_FROM_L31_ECO_MOD	The processor's instruction cache was reloaded with modified (M) data from another core's ECO L3 on the same chip due to an instruction fetch (not prefetch) or an instruction fetch plus prefetch if MMCR1[17] = '1'.

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 19 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
		0000034044		PM_INST_FROM_L31_ECO_SHR	The processor's instruction cache was reloaded with shared (S) data from another core's ECO L3 on the same chip due to an instruction fetch (not prefetch) or an instruction fetch plus prefetch if MMCR1[17] = '1'.
	0000024044			PM_INST_FROM_L31_MOD	The processor's instruction cache was reloaded with modified (M) data from another core's L3 cache on the same chip due to an instruction fetch (not prefetch) or an instruction fetch plus prefetch if MMCR1[17] = '1'.
0000014046				PM_INST_FROM_L31_SHR	The processor's instruction cache was reloaded with shared (S) data from another core's L3 cache on the same chip due to an instruction fetch (not prefetch) or an instruction fetch plus prefetch if MMCR1[17] = '1'.
		00000300FA		PM_INST_FROM_L3MISS	Marked instruction was reloaded from a location beyond the local chiplet. Instruction from an L3 miss.
			000004404E	PM_INST_FROM_L3MISS_MOD	The processor's instruction cache was reloaded from a location other than the local core's L3 cache due to an instruction fetch or an instruction fetch plus prefetch if MMCR1[17] = '1'.
	0000024048			PM_INST_FROM_LMEM	The processor's instruction cache was reloaded from the local chip's memory due to either an instruction fetch (not prefetch) or an instruction fetch plus prefetch if MMCR1[17] = '1'.
	000002404C			PM_INST_FROM_MEMORY	The processor's instruction cache was reloaded from a memory location from local remote or distant due to either an instruction fetch (not prefetch) or an instruction fetch plus prefetch if MMCR1[17] = '1'.
			000004404A	PM_INST_FROM_OFF_CHIP_CACHE	The processor's instruction cache was reloaded with either shared or modified data from another core's L2 or L3 cache on a different chip (remote or distant) due to an instruction fetch (not prefetch) or an instruction fetch plus prefetch if MMCR1[17] = '1'.
0000014048				PM_INST_FROM_ON_CHIP_CACHE	The processor's instruction cache was reloaded with either shared or modified data from another core's L2 or L3 cache on the same chip due to an instruction fetch (not prefetch) or an instruction fetch plus prefetch if MMCR1[17] = '1'.
	0000024046			PM_INST_FROM_RL2L3_MOD	The processor's instruction cache was reloaded with modified (M) data from another chip's L2 or L3 cache on the same Node or Group (remote).
000001404A				PM_INST_FROM_RL2L3_SHR	The processor's instruction cache was reloaded with shared (S) data from another chip's L2 or L3 cache on the same Node or Group (remote).
		000003404A		PM_INST_FROM_RMEM	The processor's instruction cache was reloaded from another chip's memory on the same Node or Group (remote) due to an instruction fetch (not prefetch) or an instruction fetch plus prefetch if MMCR1[17] = '1'.
	0000024050			PM_INST_GRP_PUMP_CPRED	Initial and final pump scope and data sourced across this scope was group pump (prediction is correct) for an instruction fetch.
	0000024052			PM_INST_GRP_PUMP_MPRED	Final pump scope (group) was either larger or smaller than initial pump scope for an instruction fetch. Final pump scope (group) to get data sourced.
0000014052				PM_INST_GRP_PUMP_MPRED_RTY	Final pump scope (Group) was larger than initial pump scope (chip) for an instruction fetch.
		0000030016		PM_INST_IMC_MATCH_DISP	Matched instructions (IMC matches) dispatched.
0000014054				PM_INST_PUMP_CPRED	Pump prediction correct. Counts across all types of pumps for an instruction fetch.
			0000044052	PM_INST_PUMP_MPRED	Pump misprediction. Counts across all types of pumps for an instruction fetch.

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 20 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
		0000034050		PM_INST_SYS_PUMP_CPRED	Initial and final pump scope and data sourced across this scope was system pump (prediction is correct) for an instruction fetch.
		0000034052		PM_INST_SYS_PUMP_MPRED	Final pump scope (system) mispredicted (data sourced). Either the original scope was too small (Chip/Group) or the original scope was System and it should have been smaller. Counts for an instruction fetch.
		0000044050		PM_INST_SYS_PUMP_MPRED_RTY	Final pump scope (system) was larger than initial pump scope (chip/group) for an instruction fetch. Final pump scope (system) to get data sourced.
0000010014				PM_IOPS_CMPL	Number of internal operations (IOPs) that completed.
		0000030014		PM_IOPS_DISP	Number of internal operations (IOPs) dispatched.
		0000045048		PM_IPTEG_FROM_DL2L3_MOD	A PTE was loaded into the TLB with modified (M) data from another chip's L2 or L3 cache on a different Node or Group (distant).
		0000035048		PM_IPTEG_FROM_DL2L3_SHR	A PTE was loaded into the TLB with shared (S) data from another chip's L2 or L3 cache on a different Node or Group (distant).
		000004504C		PM_IPTEG_FROM_DMEM	A PTE was loaded into the TLB from another chip's memory on the same Node or Group (distant) due to an instruction-side request.
0000015042				PM_IPTEG_FROM_L2	A PTE was loaded into the TLB from the local core's L2 cache due to an instruction-side request.
		0000035040		PM_IPTEG_FROM_L2_DISP_CONFLICT_LDHTST	A PTE was loaded into the TLB from the local core's L2 cache with a load-hit-store conflict due to an instruction-side request.
		0000045040		PM_IPTEG_FROM_L2_DISP_CONFLICT_OTHER	A PTE was loaded into the TLB from the local core's L2 cache with dispatch conflict due to an instruction-side request.
	0000025040			PM_IPTEG_FROM_L2_MEFP	A PTE was loaded into the TLB from the local core's L2 hit without dispatch conflicts on an Mefp state. due to an instruction-side request.
0000015040				PM_IPTEG_FROM_L2_NO_CONFLICT	A page table entry (PTE) was loaded into the TLB from a local core's L2 cache without conflict due to an instruction-side request.
		0000045046		PM_IPTEG_FROM_L21_MOD	A PTE was loaded into the TLB with modified (M) data from another core's L2 cache on the same chip due to an instruction-side request.
		0000035046		PM_IPTEG_FROM_L21_SHR	A PTE was loaded into the TLB with shared (S) data from another core's L2 cache on the same chip due to an instruction-side request.
000001504E				PM_IPTEG_FROM_L2MISS	A PTE was loaded into the TLB from a location other than the local core's L2 cache due to an instruction-side request
		0000045042		PM_IPTEG_FROM_L3	A PTE was loaded into the TLB from the local core's L3 cache due to an instruction-side request.
		0000035042		PM_IPTEG_FROM_L3_DISP_CONFLICT	A PTE was loaded into the TLB from the local core's L3 cache with a dispatch conflict due to an instruction-side request.
	0000025042			PM_IPTEG_FROM_L3_MEFP	A PTE was loaded into the TLB from the local core's L3 cache without dispatch conflicts hit on an Mefp state due to an instruction-side request

1. There are multiple methods to program this event.

**Table D-1. POWER8 Event List by Event Name (Sheet 21 of 46)**

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
0000015044				PM_IPTEG_FROM_L3_NO_CONFLICT	A PTE was loaded into the TLB from a local core's L3 cache without conflict due to an instruction-side request.
			0000045044	PM_IPTEG_FROM_L31_ECO_MOD	A PTE was loaded into the TLB with modified (M) data from another core's ECO L3 on the same chip due to an instruction-side request.
		0000035044		PM_IPTEG_FROM_L31_ECO_SHR	A PTE was loaded into the TLB with shared (S) data from another core's ECO L3 on the same chip due to an instruction-side request.
	0000025044			PM_IPTEG_FROM_L31_MOD	A PTE was loaded into the TLB with modified (M) data from another core's L3 cache on the same chip due to an instruction-side request.
0000015046				PM_IPTEG_FROM_L31_SHR	A PTE was loaded into the TLB with shared (S) data from another core's L3 cache on the same chip due to an instruction-side request.
			000004504E	PM_IPTEG_FROM_L3MISS	A PTE was loaded into the TLB from a location other than the local core's L3 cache due to an instruction-side request.
	0000025048			PM_IPTEG_FROM_LMEM	A PTE was loaded into the TLB from the local chip's memory due to an instruction-side request.
	000002504C			PM_IPTEG_FROM_MEMORY	A PTE was loaded into the TLB from a memory location from a local remote or distant due to an instruction-side request.
			000004504A	PM_IPTEG_FROM_OFF_CHIP_CACHE	A PTE was loaded into the TLB either shared or modified data from another core's L2 or L3 cache on a different chip (remote or distant) due to an instruction-side request.
0000015048				PM_IPTEG_FROM_ON_CHIP_CACHE	A PTE was loaded into the TLB with either shared or modified data from another core's L2 or L3 cache on the same chip due to an instruction-side request.
	0000025046			PM_IPTEG_FROM_RL2L3_MOD	A PTE was loaded into the TLB with modified (M) data from another chip's L2 or L3 cache on the same Node or Group (remote).
000001504A				PM_IPTEG_FROM_RL2L3_SHR	A PTE was loaded into the TLB with shared (S) data from another chip's L2 or L3 cache on the same Node or Group (remote).
		000003504A		PM_IPTEG_FROM_RMEM	A PTE was loaded into the TLB from another chip's memory on the same Node or Group (remote) due to an instruction-side request.
0000617082				PM_ISIDE_DISP	All instruction-side dispatch attempts.
	0000627084			PM_ISIDE_DISP_FAIL	All instruction-side dispatch attempts that failed due to an address collision with another machine.
	0000627086			PM_ISIDE_DISP_FAIL_OTHER	All instruction-side dispatch attempts that failed due to a reason other than an address collision.
			000004608E	PM_ISIDE_L2MEMACC ¹	Valid when the first beat of data comes in for an instruction-side fetch where data came from memory.
			000074608E		
			000044608E	PM_ISIDE_MRU_TOUCH	Instruction-side L2 MRU touch.
			0000040006	PM_ISLB_MISS ¹	Instruction SLB miss. An SLB miss for an instruction fetch has occurred. SLB misses trap to the operating system to resolve. This is a total count for all segment sizes.
000000D096	000000D096	000000D096	000000D096		

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 22 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
00000038AC	00000038AC	00000038AC	00000038AC	PM_ISU_REF_FXU	FXU ISU reject from either pipe.
00000030AC	00000030AC	00000030AC	00000030AC	PM_ISU_REJ_FX0	FX0 ISU reject.
00000030AE	00000030AE	00000030AE	00000030AE	PM_ISU_REJ_FX1	FX1 ISU reject.
00000030B0	00000030B0	00000030B0	00000030B0	PM_ISU_REJ_LS0	LS0 ISU reject.
00000030B2	00000030B2	00000030B2	00000030B2	PM_ISU_REJ_LS1	LS1 ISU reject.
00000030B4	00000030B4	00000030B4	00000030B4	PM_ISU_REJ_LS2	LS2 ISU reject.
00000030B6	00000030B6	00000030B6	00000030B6	PM_ISU_REJ_LS3	LS3 ISU reject.
00000030A8	00000030A8	00000030A8	00000030A8	PM_ISU_REJ_VS0	VS0 ISU reject.
00000030AA	00000030AA	00000030AA	00000030AA	PM_ISU_REJ_VS1	VS1 ISU reject.
00000038A8	00000038A8	00000038A8	00000038A8	PM_ISU_REJ_VSU	VSU ISU reject from either pipe.
00000030A2	00000030A2	00000030A2	00000030A2	PM_ISU_REJECT_RES_NA	ISU reject due to resource not available.
000000309E	000000309E	000000309E	000000309E	PM_ISU_REJECT_SAR_BYPASS	Reject because of SAR bypass.
00000030A0	00000030A0	00000030A0	00000030A0	PM_ISU_REJECT_SRC_NA	ISU reject due to source not available.
000000309C	000000309C	000000309C	000000309C	PM_ISU_REJECTS_ALL	All ISU rejects could be more than one per cycle.
00000030B8	00000030B8	00000030B8	00000030B8	PM_ISYNC	Completion count per thread for isync .
		00000400FC		PM_ITLB_MISS	Instruction TLB is reloaded. A TLB miss for an instruction fetch has occurred.
		00000300F6		PM_L1_DCACHE_RELOAD_VALID	DL1 reloaded due to demand load.
000001002C				PM_L1_DCACHE_RELOADED_ALL	L1 data cache reloaded for demand or prefetch.
000000408C	000000408C	000000408C	000000408C	PM_L1_DEMAND_WRITE	The count of instruction demand sectors written in the instruction L1 cache. Writes are always written as <u>LRU</u> and a subsequent read makes it an <u>MRU</u> .
	00000200FD			PM_L1_ICACHE_MISS	Demand I-cache miss.
		0000040012		PM_L1_ICACHE_RELOADED_ALL	Counts all I-cache reloads (includes demands).
		0000030068		PM_L1_ICACHE_RELOADED_PREF	Counts all I-cache prefetch reloads (includes demand turned into prefetch).
		67200301EA		PM_L1MISS_LAT_EXC_1024	L1 misses that took longer than 1024 cycles to resolve (miss to reload). Reload latency exceeded 1024 cycles.
			67200401EC	PM_L1MISS_LAT_EXC_2048	L1 misses that took longer than 2048 cycles to resolve (miss to reload). Reload latency exceeded 2048 cycles.
67200101E8				PM_L1MISS_LAT_EXC_256	L1 misses that took longer than 256 cycles to resolve (miss to reload). Reload latency exceeded 256 cycles.

1. There are multiple methods to program this event.

Table D-1. POWER8 Event List by Event Name (Sheet 23 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
	67200201E6			PM_L1MISS_LAT_EXC_32	L1 misses that took longer than 32 cycles to resolve (miss to reload). Reload latency exceeded 32 cycles.
	0000026086			PM_L1PF_L2MEMACC ¹	Valid when the first beat of data comes in for an L1 prefetch where data came from memory.
	0000726086				
0000417080				PM_L2_CASTOUT_MOD	L2 castouts (modified).
0000417082				PM_L2_CASTOUT_SHR	L2 castouts (shared).
	0000027084			PM_L2_CHIP_PUMP ¹	RC requests that were local on chip pump attempts.
		000064608C			
	0000427086			PM_L2_DC_INV	D-cache invalidates from the L2 cache.
		000044608C		PM_L2_DISP_ALL_L2MISS	All successful load/store dispatches for this thread that were an L2 miss.
	0000027086			PM_L2_GROUP_PUMP ¹	RC requests that were on node pump attempts.
		000064608E			
	0000626084			PM_L2_GRP_GUESS_CORRECT	L2 guess group and guess is correct.
	0000626086			PM_L2_GRP_GUESS_WRONG	L2 guess group and guess is not correct.
	0000427084			PM_L2_IC_INV	I-cache invalidates from the L2 cache.
		0000436088		PM_L2_INST	All successful instruction-side dispatches for this thread (excludes i_l2mru_tch requests).
		000043608A		PM_L2_INST_MISS	All successful instruction-side dispatches that were an L2 miss for this thread.
0000416080				PM_L2_LD	All successful data-side load dispatches for this thread.
		0000437088		PM_L2_LD_DISP	All successful load dispatches.
		000043708A		PM_L2_LD_HIT	All successful load dispatches that were L2 hits.
	0000426084			PM_L2_LD_MISS	All successful data-side load dispatches that were an L2 miss for this thread.
0000616080				PM_L2_LOC_GUESS_CORRECT	L2 guess location (loc) and guess is correct.
0000616082				PM_L2_LOC_GUESS_WRONG	L2 guess location and guess is not correct.
		0000537088		PM_L2_RC_ST_DONE	RC stored to a line that was Tx or Sx.
0000516080				PM_L2_RCLD_DISP	L2 RC load dispatch attempt.
0000516082				PM_L2_RCLD_DISP_FAIL_ADDR	L2 RC load dispatch attempt failed due to an address collision with RC/CO/SN/SQ.
	0000526084			PM_L2_RCLD_DISP_FAIL_OTHER	L2 RC load dispatch attempt failed due to other reasons.
		0000536088		PM_L2_RCST_DISP	L2 RC store dispatch attempt.

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 24 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
		000053608A		PM_L2_RCST_DISP_FAIL_ADDR	L2 RC store dispatch attempt failed due to an address collision with RC/CO/SN/SQ.
		000054608C		PM_L2_RCST_DISP_FAIL_OTHER	L2 RC store dispatch attempt failed due to other reasons.
		000004708C		PM_L2_RTY_LD ¹	RC retries on the SMP interconnect for any load from the core.
	000063708A			PM_L2_RTY_ST	RC retries on SMP interconnect for any store from the core.
	000003708A			PM_L2_RTY_ST	RC retries on the SMP interconnect for any store from the core.
	0000637088			PM_L2_SN_M_RD_DONE	A snoop machine was dispatched for a read and was in a modified (M) state.
		000054708C		PM_L2_SN_M_WR_DONE	A snoop machine was dispatched for a write and was in a modified (M) state.
	000053708A			PM_L2_SN_SX_I_DONE	A snoop machine was dispatched and went from Sx or Tx to Ix.
0000017080				PM_L2_ST ¹	All successful data-side store dispatches for this thread.
0000416082				PM_L2_ST_DISP	All successful store dispatches.
		000044708C		PM_L2_ST_HIT	All successful store dispatches that were L2 hits.
0000017082				PM_L2_ST_MISS ¹	All successful data-side store dispatches for this thread that were an L2 miss.
	0000426086			PM_L2_SYS_GUESS_CORRECT	L2 guess system and guess is correct.
		0000636088		PM_L2_SYS_GUESS_WRONG	L2 guess system and guess was not correct.
		0000037088		PM_L2_SYS_PUMP ¹	RC requests that were system pump attempts.
0000617080				PM_L2_TM_REQ_ABORT	TM abort.
000001E15E				PM_L2_TM_ST_ABORT_SISTER	TM marked store abort.
	0000128084			PM_L3_CI_HIT	Total count of L3 castin hits.
	0000128086			PM_L3_CI_MISS	Total count of L3 castin misses.
0000819082				PM_L3_CI_USAGE	Rotating sample of 16 CI or CO actives.
		000023808A		PM_L3_CINJ	L3 castin of cache inject.
		0000438088		PM_L3_CO	L3 castout occurring (does not include castthrough or log writes).
	0000028086			PM_L3_CO_L31	L3 CO to L3.1 or of port 0 and port 1 (lossy).
		0000238088		PM_L3_CO_LCO	Total L3 castouts occurred on LCO.

1. There are multiple methods to program this event.

**Table D-1. POWER8 Event List by Event Name (Sheet 25 of 46)**

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
	0000028084			PM_L3_CO_MEM	L3 CO to memory or of port 0 and port 1 (lossy).
		000003E05E		PM_L3_CO_MEPE ¹	L3 CO of a line in the Mepf state (includes castthrough).
0000018082				PM_L3_CO0_ALLOC	Lifetime.
		000083908B		PM_L3_CO0_BUSY	Lifetime.
0000B19082				PM_L3_GRP_GUESS_CORRECT	Initial scope is the group and data from same group (near). Prediction is successful.
		0000B3908A		PM_L3_GRP_GUESS_WRONG_HIGH	Initial scope is the group but data from the local node. Prediction is too high.
		0000B39088		PM_L3_GRP_GUESS_WRONG_LOW	Initial scope is the group but data from outside the group (far or remote). Prediction is too low.
0000218080				PM_L3_HIT	L3 hits.
		0000138088		PM_L3_L2_CO_HIT	L2 castout hits.
		000013808A		PM_L3_L2_CO_MISS	L2 castout misses.
		000014808C		PM_L3_LAT_CI_HIT	L3 lateral castin hits.
		000014808E		PM_L3_LAT_CI_MISS	L3 lateral castin misses.
0000228084				PM_L3_LD_HIT	L3 demand load hits.
0000228086				PM_L3_LD_MISS	L3 demand load miss.
000001E052				PM_L3_LD_PREF	L3 load prefetches.
0000B19080				PM_L3_LOC_GUESS_CORRECT	Initial scope is the node/chip and data from the local node (local). Prediction is successful.
	0000B29086			PM_L3_LOC_GUESS_WRONG	Initial scope is the node but data from outside the local node (near or far or remote). Prediction is too low.
0000218082				PM_L3_MISS	L3 misses.
		000054808C		PM_L3_P0_CO_L31	L3 CO to L3.1 port 0.
		0000538088		PM_L3_P0_CO_MEM	L3 CO to memory port 0.
	0000929084			PM_L3_P0_CO_RTY	L3 CO received retry port 0.
0000A29084				PM_L3_P0_GRP_PUMP	L3 prefetch sent with group scope port 0.
0000528084				PM_L3_P0_LCO_DATA	LCO sent with data port 0.
0000518080				PM_L3_P0_LCO_NO_DATA	Dataless L3 LCO sent to port 0.
		0000A4908C		PM_L3_P0_LCO_RTY	L3 LCO received retry port 0.
0000A19080				PM_L3_P0_NODE_PUMP	L3 prefetch sent with nodal scope port 0.

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 26 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
0000919080				PM_L3_P0_PF_RTY	L3 PF received retry port 0.
	0000939088			PM_L3_P0_SN_HIT	L3 snoop hit port 0.
0000118080				PM_L3_P0_SN_INV	Port 0 snooper detects a store to a line that is in any shared state (Sx).
		000094908C		PM_L3_P0_SN_MISS	L3 snoop miss port 0.
		0000A39088		PM_L3_P0_SYS_PUMP	L3 prefetch sent with system scope port 0.
		000054808E		PM_L3_P1_CO_L31	L3 CO to L3.1 port 1.
		000053808A		PM_L3_P1_CO_MEM	L3 CO to memory port 1.
0000929086				PM_L3_P1_CO_RTY	L3 CO received retry port 1.
0000A29086				PM_L3_P1_GRP_PUMP	L3 prefetch sent with group scope port 1.
0000528086				PM_L3_P1_LCO_DATA	LCO sent with data port 1.
0000518082				PM_L3_P1_LCO_NO_DATA	Dataless L3 LCO sent to port 1.
		0000A4908E		PM_L3_P1_LCO_RTY	L3 LCO received retry port 1.
0000A19082				PM_L3_P1_NODE_PUMP	L3 prefetch sent with nodal scope port 1.
0000919082				PM_L3_P1_PF_RTY	L3 PF received retry port 1.
	000093908A			PM_L3_P1_SN_HIT	L3 snoop hit port 1.
0000118082				PM_L3_P1_SN_INV	Port1 snooper detects a store to a line that is in any shared state.
		000094908E		PM_L3_P1_SN_MISS	L3 snoop miss port 1.
		0000A3908A		PM_L3_P1_SYS_PUMP	L3 prefetch sent with system scope port 1.
	0000428084			PM_L3_PF_HIT_L3	L3 prefetch hit in L3 cache.
0000018080				PM_L3_PF_MISS_L3	L3 prefetch missed in L3 cache.
		000003808A		PM_L3_PF_OFF_CHIP_CACHE	L3 prefetch from off-chip cache.
			000004808E	PM_L3_PF_OFF_CHIP_MEM	L3 prefetch from off-chip memory.
			0000038088	PM_L3_PF_ON_CHIP_CACHE	L3 prefetch from on-chip cache.
			000004808C	PM_L3_PF_ON_CHIP_MEM	L3 prefetch from on-chip memory.
	0000829084			PM_L3_PF_USAGE	Rotating sample of 32 prefetch actives.
			000084908D	PM_L3_PF0_ALLOC	Lifetime.
			000084908C	PM_L3_PF0_BUSY	Lifetime.
			000004E052	PM_L3_PREF_ALL	Total hardware L3 prefetches (load and store).

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 27 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
	0000829086			PM_L3_RD_USAGE	Rotating sample of 16 read actives.
			000084908F	PM_L3_RD0_ALLOC	Lifetime.
			000084908E	PM_L3_RD0_BUSY	Lifetime.
0000819080				PM_L3_SN_USAGE	Rotating sample of eight snoop valids.
		0000839089		PM_L3_SN0_ALLOC	Lifetime.
		0000839088		PM_L3_SN0_BUSY	Lifetime.
	000002E052			PM_L3_ST_PREF	L3 store prefetches.
		000003E052		PM_L3_SW_PREF	Data stream touch to L3 cache.
	0000B29084			PM_L3_SYS_GUESS_CORRECT	Initial scope is the system and data from outside the group (far or remote). Prediction is successful.
			0000B4908C	PM_L3_SYS_GUESS_WRONG	Initial scope is the system but data from local or near. Prediction is too high.
			000024808E	PM_L3_TRANS_PF	L3 transient prefetch.
0000418082				PM_L3_WI_USAGE	Rotating sample of eight write inject (WI) actives.
0000018081				PM_L3_WI0_ALLOC	Lifetime.
0000418080				PM_L3_WI0_BUSY	Lifetime.
		000003C058		PM_LARX_FIN	Larx finished.
000001002E				PM_LD_CMPL	Count of loads completed.
0000010062				PM_LD_L3MISS_PEND_CYC	Number of cycles an L3 miss was pending for this thread.
		000003E054		PM_LD_MISS_L1 ¹	Load missed L1 cache.
00000100EE			00000400F0		
000000C080	000000C080	000000C080	000000C080	PM_LD_REF_L1_LSU0	All L1 D-cache load references counted at finish.
000000C082	000000C082	000000C082	000000C082	PM_LD_REF_L1_LSU1	LS0 L1 D-cache load references counted at finish.
000000C094	000000C094	000000C094	000000C094	PM_LD_REF_L1_LSU2	LS1 L1 D-cache load references counted at finish.
000000C096	000000C096	000000C096	000000C096	PM_LD_REF_L1_LSU3	LS2 L1 D-cache load references counted at finish.
000000509A	000000509A	000000509A	000000509A	PM_LINK_STACK_INVALID_PTR	A flush where the LS pointer is invalid.
0000005098	0000005098	0000005098	0000005098	PM_LINK_STACK_WRONG_ADD_PRED	Link stack predicts a wrong address, because of link stack design limitations.
000000E080	000000E080	000000E080	000000E080	PM_LS0_ERAT_MISS_PREF	LS0 ERAT miss due to prefetch.

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 28 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
000000D0B8	000000D0B8	000000D0B8	000000D0B8	PM_LSO_L1_PREF	LS0 L1 cache data prefetches.
000000C098	000000C098	000000C098	000000C098	PM_LSO_L1_SW_PREF	Software L1 prefetches.
000000E082	000000E082	000000E082	000000E082	PM_LS1_ERAT_MISS_PREF	LS1 ERAT miss due to prefetch.
000000D0BA	000000D0BA	000000D0BA	000000D0BA	PM_LS1_L1_PREF	LS1 L1 cache data prefetches.
000000C09A	000000C09A	000000C09A	000000C09A	PM_LS1_L1_SW_PREF	Software L1 prefetches.
	00000200F6			PM_LSU_DERAT_MISS	DERAT reloaded due to a DERAT miss.
000000E880	000000E880	000000E880	000000E880	PM_LSU_ERAT_MISS_PREF	ERAT miss due to prefetch.
	0000030066			PM_LSU_FIN	LSU finished an instruction (up to two per cycle).
000000C8AC	000000C8AC	000000C8AC	000000C8AC	PM_LSU_FLUSH_UST	Unaligned store flush on either pipe.
000000D0A4	000000D0A4	000000D0A4	000000D0A4	PM_LSU_FOUR_TABLEWALK_CYC	Cycles when four tablewalks are pending on this thread.
0000010066				PM_LSU_FX_FIN	LSU finished an FX operation (up to two per cycle).
000000D8B8	000000D8B8	000000D8B8	000000D8B8	PM_LSU_L1_PREF	Hardware initiated.
000000C898	000000C898	000000C898	000000C898	PM_LSU_L1_SW_PREF	Software L1 prefetches.
000000C884	000000C884	000000C884	000000C884	PM_LSU_LDF	FPU loads only on LS2/LS3 (that is, LU0/LU1).
000000C888	000000C888	000000C888	000000C888	PM_LSU_LDX	Vector loads can issue only on LS2/LS3.
000000D0A2	000000D0A2	000000D0A2	000000D0A2	PM_LSU_LMQ_FULL_CYC	LMQ is full.
000000D0A1	000000D0A1	000000D0A1	000000D0A1	PM_LSU_LMQ_S0_ALLOC	Per thread: use edge detect to count allocates on a per-thread basis.
000000D0A0	000000D0A0	000000D0A0	000000D0A0	PM_LSU_LMQ_S0_VALID	Slot 0 of LMQ is valid.
	000003001C			PM_LSU_LMQ_SRQ_EMPTY_ALL_CYC	All threads LSU are empty (LMQ and SRQ empty).
	000002003E			PM_LSU_LMQ_SRQ_EMPTY_CYC	LSU empty (LMQ and SRQ empty).
000000D09F	000000D09F	000000D09F	000000D09F	PM_LSU_LRQ_S0_ALLOC	Use edge detect to count allocates on a per-thread basis.
000000D09E	000000D09E	000000D09E	000000D09E	PM_LSU_LRQ_S0_VALID	Slot 0 of LRQ is valid.
000000F091	000000F091	000000F091	000000F091	PM_LSU_LRQ_S43_ALLOC	LRQ slot 43 was released.
000000F090	000000F090	000000F090	000000F090	PM_LSU_LRQ_S43_VALID	LRQ slot 43 is busy.
	0004030162			PM_LSU_MRK_DERAT_MISS	D-ERAT reloaded (miss).
000000C88C	000000C88C	000000C88C	000000C88C	PM_LSU_NCLD	Count at finish so it can return only on LS0 or LS1.
000000C092	000000C092	000000C092	000000C092	PM_LSU_NCST	Noncacheable stores sent to nest.

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 29 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
0000010064				PM_LSU_REJECT	LSU reject (up to four per cycle).
	000002E05C			PM_LSU_REJECT_ERAT_MISS	LSU reject due to an ERAT (up to four per cycles).
		000004E05C		PM_LSU_REJECT_LHS	LSU reject due to LHS (up to 4 per cycle).
000001E05C				PM_LSU_REJECT_LMQ_FULL	LSU reject due to LMQ full (4 per cycle).
000000D082	000000D082	000000D082	000000D082	PM_LSU_SET_MPRED	Line already in cache at reload time.
			0000040008	PM_LSU_SRQ_EMPTY_CYC	Number of cycles the store reorder queue (SRQ) was empty for all threads.
000001001A				PM_LSU_SRQ_FULL_CYC	Storage reorder queue (SRQ) is full and is blocking dispatch.
000000D09D	000000D09D	000000D09D	000000D09D	PM_LSU_SRQ_S0_ALLOC	Per thread: use edge detect to count allocates on a per-thread basis.
000000D09C	000000D09C	000000D09C	000000D09C	PM_LSU_SRQ_S0_VALID	Slot 0 of SRQ is valid.
000000F093	000000F093	000000F093	000000F093	PM_LSU_SRQ_S39_ALLOC	SRQ slot 39 was released.
000000F092	000000F092	000000F092	000000F092	PM_LSU_SRQ_S39_VALID	SRQ slot 39 is busy.
000000D09B	000000D09B	000000D09B	000000D09B	PM_LSU_SRQ_SYNC	A sync in the SRQ ended.
000000D09A	000000D09A	000000D09A	000000D09A	PM_LSU_SRQ_SYNC_CYC	Number of cycles a sync is in the SRQ (edge detect to count).
000000F084	000000F084	000000F084	000000F084	PM_LSU_STORE_REJECT	Store reject on either pipe.
000000D0A6	000000D0A6	000000D0A6	000000D0A6	PM_LSU_TWO_TABLEWALK_CYC	Cycles when two tablewalks are pending on this thread.
000000C0B0	000000C0B0	000000C0B0	000000C0B0	PM_LSU0_FLUSH_LRQ	LSU0 flush: LRQ.
000000C0B8	000000C0B8	000000C0B8	000000C0B8	PM_LSU0_FLUSH_SRQ	LSU0 flush: SRQ load-hit-store flushes.
000000C0A4	000000C0A4	000000C0A4	000000C0A4	PM_LSU0_FLUSH_ULD	LSU0 flush: unaligned load.
000000C0AC	000000C0AC	000000C0AC	000000C0AC	PM_LSU0_FLUSH_UST	LS0 flush: unaligned store.
000000F088	000000F088	000000F088	000000F088	PM_LSU0_L1_CAM_CANCEL	LSU0 L1 TM CAM cancel.
000001E056				PM_LSU0_LARX_FIN	Larx finished in LSU Pipe0.
000000D08C	000000D08C	000000D08C	000000D08C	PM_LSU0_LMQ_LHR_MERGE	LSU0 load merged with another cache-line request.
000000C08C	000000C08C	000000C08C	000000C08C	PM_LSU0_NCLD	LS0 noncacheable loads counted at finish.
000000E090	000000E090	000000E090	000000E090	PM_LSU0_PRIMARY_ERAT_HIT	Primary ERAT hit.
000001E05A				PM_LSU0_REJECT	LSU0 reject.
000000C09C	000000C09C	000000C09C	000000C09C	PM_LSU0_SRQ_STFWD	LSU0 SRQ forwarded data to a load.
000000F084	000000F084	000000F084	000000F084	PM_LSU0_STORE_REJECT	LSU0 store reject.
000000E098	000000E098	000000E098	000000E098	PM_LSU0_TM_L1_HIT	Load TM hit in L1 cache.

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 30 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
000000E0A0	000000E0A0	000000E0A0	000000E0A0	PM_LSU0_TM_L1_MISS	Load TM L1 miss.
000000E0A8	000000E0A8	000000E0A8	000000E0A8	PM_LSU0_TMA_REQ_L2	Address-only requests to the L2 cache. Multiple requests to the same line are only reported once.
000000C0B2	000000C0B2	000000C0B2	000000C0B2	PM_LSU1_FLUSH_LRQ	LSU1 flush: LRQ.
000000C0BA	000000C0BA	000000C0BA	000000C0BA	PM_LSU1_FLUSH_SRQ	LSU1 flush: SRQ load-hit-store flushes.
000000C0A6	000000C0A6	000000C0A6	000000C0A6	PM_LSU1_FLUSH_ULD	LSU1 flush: unaligned load.
000000C0AE	000000C0AE	000000C0AE	000000C0AE	PM_LSU1_FLUSH_UST	LSU1 flush: unaligned store.
000000F08A	000000F08A	000000F08A	000000F08A	PM_LSU1_L1_CAM_CANCEL	LSU1 L1 TM CAM cancel.
	000002E056			PM_LSU1_LARX_FIN	Larx finished in LSU Pipe1.
000000D08E	000000D08E	000000D08E	000000D08E	PM_LSU1_LMQ_LHR_MERGE	LSU1 load merge with another cache-line request.
000000C08E	000000C08E	000000C08E	000000C08E	PM_LSU1_NCLD	LS1 noncacheable loads counted at finish.
000000E092	000000E092	000000E092	000000E092	PM_LSU1_PRIMARY_ERAT_HIT	Primary ERAT hit.
	000002E05A			PM_LSU1_REJECT	LSU1 reject.
000000C09E	000000C09E	000000C09E	000000C09E	PM_LSU1_SRQ_STFWD	LSU1 SRQ forwarded data to a load.
000000F086	000000F086	000000F086	000000F086	PM_LSU1_STORE_REJECT	LSU1 store reject.
000000E09A	000000E09A	000000E09A	000000E09A	PM_LSU1_TM_L1_HIT	Load TM hit in L1 cache.
000000E0A2	000000E0A2	000000E0A2	000000E0A2	PM_LSU1_TM_L1_MISS	Load TM L1 miss.
000000E0AA	000000E0AA	000000E0AA	000000E0AA	PM_LSU1_TMA_REQ_L2	Address-only requests to the L2 cache. Multiple requests to the same line are only reported once.
000000C0B4	000000C0B4	000000C0B4	000000C0B4	PM_LSU2_FLUSH_LRQ	LSU2 flush: LRQ.
000000C0BC	000000C0BC	000000C0BC	000000C0BC	PM_LSU2_FLUSH_SRQ	LSU2 flush: SRQ.
000000C0A8	000000C0A8	000000C0A8	000000C0A8	PM_LSU2_FLUSH_ULD	LSU2 flush: unaligned load.
000000F08C	000000F08C	000000F08C	000000F08C	PM_LSU2_L1_CAM_CANCEL	LSU2 L1 TM CAM cancel.
	000003E056			PM_LSU2_LARX_FIN	Larx finished in LSU Pipe2.
000000C084	000000C084	000000C084	000000C084	PM_LSU2_LDF	LS2 scalar loads.
000000C088	000000C088	000000C088	000000C088	PM_LSU2_LDX	LS0 vector loads.
000000D090	000000D090	000000D090	000000D090	PM_LSU2_LMQ_LHR_MERGE	LSU2 load merged with another cache-line request.
000000E094	000000E094	000000E094	000000E094	PM_LSU2_PRIMARY_ERAT_HIT	Primary ERAT hit.
	000003E05A			PM_LSU2_REJECT	LSU2 reject.

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 31 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
000000C0A0	000000C0A0	000000C0A0	000000C0A0	PM_LSU2_SRQ_STFWD	LSU2 SRQ forwarded data to a load.
000000E09C	000000E09C	000000E09C	000000E09C	PM_LSU2_TM_L1_HIT	Load TM hit in L1 cache.
000000E0A4	000000E0A4	000000E0A4	000000E0A4	PM_LSU2_TM_L1_MISS	Load TM L1 miss.
000000E0AC	000000E0AC	000000E0AC	000000E0AC	PM_LSU2_TMA_REQ_L2	Address-only requests to the L2 cache. Multiple requests to the same line are only reported once.
000000C0B6	000000C0B6	000000C0B6	000000C0B6	PM_LSU3_FLUSH_LRQ	LSU3 flush: LRQ.
000000C0BE	000000C0BE	000000C0BE	000000C0BE	PM_LSU3_FLUSH_SRQ	LSU3 flush: SRQ.
000000C0AA	000000C0AA	000000C0AA	000000C0AA	PM_LSU3_FLUSH_ULD	LSU3 flush: unaligned load.
000000F08E	000000F08E	000000F08E	000000F08E	PM_LSU3_L1_CAM_CANCEL	LSU3 L1 TM CAM cancel.
			000004E056	PM_LSU3_LARX_FIN	Larx finished in LSU Pipe3.
000000C086	000000C086	000000C086	000000C086	PM_LSU3_LDF	LS3 scalar loads
000000C08A	000000C08A	000000C08A	000000C08A	PM_LSU3_LDX	LS1 vector loads
000000D092	000000D092	000000D092	000000D092	PM_LSU3_LMQ_LHR_MERGE	LSU3 load merged with another cache-line request.
000000E096	000000E096	000000E096	000000E096	PM_LSU3_PRIMARY_ERAT_HIT	Primary ERAT hit.
			000004E05A	PM_LSU3_REJECT	LSU3 reject.
000000C0A2	000000C0A2	000000C0A2	000000C0A2	PM_LSU3_SRQ_STFWD	LSU3 SRQ forwarded data to a load.
000000E09E	000000E09E	000000E09E	000000E09E	PM_LSU3_TM_L1_HIT	Load TM hit in L1 cache.
000000E0A6	000000E0A6	000000E0A6	000000E0A6	PM_LSU3_TM_L1_MISS	Load TM L1 miss.
000000E0AE	000000E0AE	000000E0AE	000000E0AE	PM_LSU3_TMA_REQ_L2	Address-only requests to the L2 cache. Multiple requests to the same line are only reported once.
0000005094	0000005094	0000005094	0000005094	PM_LWSYNC ¹	Threaded version. An IWSYNC count (which is easier to use than IMC). An IWSYNC issued to the LSU was sent to the L2 cache. An IWSYNC does not require a sync_ack back from the L2 cache. Therefore, the store can be completed in the core and the SRQ entry can be released like other stores.
000000D098	000000D098	000000D098	000000D098		
000000209A	000000209A	000000209A	000000209A	PM_LWSYNC_HELD	Cycles an IWSYNC instruction was held at dispatch.
			000004C058	PM_MEM_CO	Memory castouts from this LPAR.
0000010058				PM_MEM_LOC_THRESH_IFU	Local memory above threshold for IFU speculation control.
			0000040056	PM_MEM_LOC_THRESH_LSU_HIGH	Local memory above threshold for LSU medium.
000001C05E				PM_MEM_LOC_THRESH_LSU_MED	Local memory above threshold for data prefetch.
	000002C058			PM_MEM_PREF	Memory prefetch for this LPAR.

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 32 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
0000010056				PM_MEM_READ	Reads from memory from this LPAR (includes demand data/instruction/translate /L1 prefetch/instruction prefetch).
	000003C05E			PM_MEM_RWITM	Memory <u>rwitm</u> for this LPAR.
	000003515E			PM_MRK_BACK_BR_CMPL	Marked branch instruction completed with a target address less than the current instruction address.
000001016E				PM_MRK_BR_CMPL	Branch instruction completed.
	00020301E4			PM_MRK_BR_MPRED_CMPL	Marked branch mispredicted.
000E0101E2				PM_MRK_BR_TAKEN_CMPL	Marked branch taken completed.
	000002013A			PM_MRK_BRU_FIN	BRU marked instruction finish.
	000003013A			PM_MRK_CRU_FIN	IFU non-branch marked instruction finished.
		000104D148	PM_MRK_DATA_FROM_DL2L3_MOD	The processor's data cache was reloaded with modified (M) data from another chip's L2 or L3 cache on a different Node or Group (distant).	
	000102D128		PM_MRK_DATA_FROM_DL2L3_MOD_CYC	Duration in cycles to reload with modified (M) data from another chip's L2 or L3 cache on a different Node or Group (distant).	
	000103D148		PM_MRK_DATA_FROM_DL2L3_SHR	The processor's data cache was reloaded with shared (S) data from another chip's L2 or L3 cache on a different Node or Group (distant).	
	000102C128		PM_MRK_DATA_FROM_DL2L3_SHR_CYC	Duration in cycles to reload with shared (S) data from another chip's L2 or L3 cache on a different Node or Group (distant).	
		000104D14C	PM_MRK_DATA_FROM_DMEM	The processor's data cache was reloaded from another chip's memory on the same Node or Group (distant) due to a marked load.	
	000102D12C		PM_MRK_DATA_FROM_DMEM_CYC	Duration in cycles to reload from another chip's memory on the same Node or Group (distant) due to a marked load.	
000101D142			PM_MRK_DATA_FROM_L2	The processor's data cache was reloaded from the local core's L2 cache due to a marked load.	
		000104C122	PM_MRK_DATA_FROM_L2_CYC	Duration in cycles to reload from the local core's L2 cache due to a marked load.	
	000103D140		PM_MRK_DATA_FROM_L2_DISP_CONFLICT_LDHTST	The processor's data cache was reloaded from the local core's L2 cache with a load-hit-store conflict due to a marked load.	
	000102C120		PM_MRK_DATA_FROM_L2_DISP_CONFLICT_LDHTST_CYC	Duration in cycles to reload from the local core's L2 cache with a load-hit-store conflict due to a marked load.	
		000104D140	PM_MRK_DATA_FROM_L2_DISP_CONFLICT_OTHER	The processor's data cache was reloaded from the local core's L2 cache with a dispatch conflict due to a marked load.	
	000102D120		PM_MRK_DATA_FROM_L2_DISP_CONFLICT_OTHER_CYC	Duration in cycles to reload from the local core's L2 cache with a dispatch conflict due to a marked load.	
	000102D140		PM_MRK_DATA_FROM_L2_MEFP	The processor's data cache was reloaded from the local core's L2 hit without a dispatch conflict on an Mepf state due to a marked load.	

1. There are multiple methods to program this event.

**Table D-1. POWER8 Event List by Event Name (Sheet 33 of 46)**

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
			000104D120	PM_MRK_DATA_FROM_L2_MEPEF_CYC	Duration in cycles to reload from the local core's L2 hit without a dispatch conflict on an Mepf state due to a marked load.
000101D140				PM_MRK_DATA_FROM_L2_NO_CONFLICT	The processor's data cache was reloaded from local core's L2 cache without conflict due to a marked load.
			000104C120	PM_MRK_DATA_FROM_L2_NO_CONFLICT_CYC	Duration in cycles to reload from the local core's L2 cache without conflict due to a marked load.
			000104D146	PM_MRK_DATA_FROM_L21_MOD	The processor's data cache was reloaded with modified (M) data from another core's L2 cache on the same chip due to a marked load.
	000102D126			PM_MRK_DATA_FROM_L21_MOD_CYC	Duration in cycles to reload with modified (M) data from another core's L2 cache on the same chip due to a marked load.
		000103D146		PM_MRK_DATA_FROM_L21_SHR	The processor's data cache was reloaded with shared (S) data from another core's L2 cache on the same chip due to a marked load.
	000102C126			PM_MRK_DATA_FROM_L21_SHR_CYC	Duration in cycles to reload with shared (S) data from another core's L2 cache on the same chip due to a marked load.
000101D14E			00010401E8	PM_MRK_DATA_FROM_L2MISS ¹	The processor's data cache was reloaded from a location other than the local core's L2 cache due to a marked load.
			000104C12E	PM_MRK_DATA_FROM_L2MISS_CYC	Duration in cycles to reload from a location other than the local core's L2 cache due to a marked load.
			000104D142	PM_MRK_DATA_FROM_L3	The processor's data cache was reloaded from the local core's L3 cache due to a marked load.
	000102D122			PM_MRK_DATA_FROM_L3_CYC	Duration in cycles to reload from the local core's L3 cache due to a marked load.
		000103D142		PM_MRK_DATA_FROM_L3_DISP_CONFLICT	The processor's data cache was reloaded from the local core's L3 cache with a dispatch conflict due to a marked load.
	000102C122			PM_MRK_DATA_FROM_L3_DISP_CONFLICT_CYC	Duration in cycles to reload from the local core's L3 cache with a dispatch conflict due to a marked load.
	000102D142			PM_MRK_DATA_FROM_L3_MEPEF	The processor's data cache was reloaded from the local core's L3 cache without dispatch conflicts hit on an Mepf state due to a marked load.
			000104D122	PM_MRK_DATA_FROM_L3_MEPEF_CYC	Duration in cycles to reload from the local core's L3 cache without a dispatch conflict hit on an Mepf state due to a marked load.
000101D144				PM_MRK_DATA_FROM_L3_NO_CONFLICT	The processor's data cache was reloaded from the local core's L3 cache without conflict due to a marked load.
			000104C124	PM_MRK_DATA_FROM_L3_NO_CONFLICT_CYC	Duration in cycles to reload from the local core's L3 cache without conflict due to a marked load.
			000104D144	PM_MRK_DATA_FROM_L31_ECO_MOD	The processor's data cache was reloaded with modified (M) data from another core's ECO L3 on the same chip due to a marked load.
	000102D124			PM_MRK_DATA_FROM_L31_ECO_MOD_CYC	Duration in cycles to reload with modified (M) data from another core's ECO L3 on the same chip due to a marked load.

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 34 of 46)

Version 1.3
16 March 2016Performance Monitoring Events
Page 421 of 450

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
		000103D144		PM_MRK_DATA_FROM_L31_ECO_SHR	The processor's data cache was reloaded with shared (S) data from another core's ECO L3 on the same chip due to a marked load.
	000102C124			PM_MRK_DATA_FROM_L31_ECO_SHR_CYC	Duration in cycles to reload with shared (S) data from another core's ECO L3 on the same chip due to a marked load.
	000102D144			PM_MRK_DATA_FROM_L31_MOD	The processor's data cache was reloaded with modified (M) data from another core's L3 cache on the same chip due to a marked load.
		000104D124		PM_MRK_DATA_FROM_L31_MOD_CYC	Duration in cycles to reload with modified (M) data from another core's L3 cache on the same chip due to a marked load.
000101D146				PM_MRK_DATA_FROM_L31_SHR	The processor's data cache was reloaded with shared (S) data from another core's L3 cache on the same chip due to a marked load.
		000104C126		PM_MRK_DATA_FROM_L31_SHR_CYC	Duration in cycles to reload with shared (S) data from another core's L3 cache on the same chip due to a marked load.
00010201E4				PM_MRK_DATA_FROM_L3MISS ¹	The processor's data cache was reloaded from a location other than the local core's L3 cache due to a marked load.
		000104D14E			
000102D12E				PM_MRK_DATA_FROM_L3MISS_CYC	Duration in cycles to reload from a location other than the local core's L3 cache due to a marked load.
000102D148				PM_MRK_DATA_FROM_LMEM	The processor's data cache was reloaded from the local chip's memory due to a marked load.
		000104D128		PM_MRK_DATA_FROM_LMEM_CYC	Duration in cycles to reload from the local chip's memory due to a marked load.
00010201E0				PM_MRK_DATA_FROM_MEM	The processor's data cache was reloaded from a memory location from a local remote or distant due to a marked load.
000102D14C				PM_MRK_DATA_FROM_MEMORY	The processor's data cache was reloaded from a memory location from local remote or distant due to a marked load.
		000104D12C		PM_MRK_DATA_FROM_MEMORY_CYC	Duration in cycles to reload from a memory location from the local remote or distant due to a marked load.
		000104D14A		PM_MRK_DATA_FROM_OFF_CHIP_CACHE	The processor's data cache was reloaded with either shared or modified data from another core's L2 or L3 cache on a different chip (remote or distant) due to a marked load.
000102D12A				PM_MRK_DATA_FROM_OFF_CHIP_CACHE_CYC	Duration in cycles to reload either shared or modified data from another core's L2 or L3 cache on a different chip (remote or distant) due to a marked load.
000101D148				PM_MRK_DATA_FROM_ON_CHIP_CACHE	The processor's data cache was reloaded with either shared or modified data from another core's L2 or L3 cache on the same chip due to a marked load.
		000104C128		PM_MRK_DATA_FROM_ON_CHIP_CACHE_CYC	Duration in cycles to reload either shared or modified data from another core's L2 or L3 cache on the same chip due to a marked load.
000102D146				PM_MRK_DATA_FROM_RL2L3_MOD	The processor's data cache was reloaded with modified (M) data from another chip's L2 or L3 cache on the same Node or Group (remote).
		000104D126		PM_MRK_DATA_FROM_RL2L3_MOD_CYC	Duration in cycles to reload with modified (M) data from another chip's L2 or L3 cache on the same Node or Group (remote).

1. There are multiple methods to program this event.

Table D-1. POWER8 Event List by Event Name (Sheet 35 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
000101D14A				PM_MRK_DATA_FROM_RL2L3_SHR	The processor's data cache was reloaded with shared (S) data from another chip's L2 or L3 cache on the same Node or Group (remote).
			000104C12A	PM_MRK_DATA_FROM_RL2L3_SHR_CYC	Duration in cycles to reload with shared (S) data from another chip's L2 or L3 cache on the same Node or Group (remote).
		000103D14A		PM_MRK_DATA_FROM_RMEM	The processor's data cache was reloaded from another chip's memory on the same Node or Group (remote) due to a marked load.
	000102C12A			PM_MRK_DATA_FROM_RMEM_CYC	Duration in cycles to reload from another chip's memory on the same Node or Group (remote) due to a marked load.
			0001040118	PM_MRK_DCACHE_RELOAD_INTV	Combined intervention event.
		00040301E6		PM_MRK_DERAT_MISS	ERAT miss (TLB access) all page sizes.
			000404D154	PM_MRK_DERAT_MISS_16G	Marked data ERAT miss (data TLB access) page size 16 GB.
		000403D154		PM_MRK_DERAT_MISS_16M	Marked data ERAT miss (data TLB access) page size 16 MB.
000401D156				PM_MRK_DERAT_MISS_4K	Marked data ERAT miss (data TLB access) page size is 4 KB.
	000402D154			PM_MRK_DERAT_MISS_64K	Marked data ERAT miss (data TLB access) page size 64 KB.
	0000020132			PM_MRK_DFU_FIN	Decimal unit marked instruction finish.
			000404F148	PM_MRK_DPTEG_FROM_DL2L3_MOD	A PTE was loaded into the TLB with modified (M) data from another chip's L2 or L3 cache on a different Node or Group (distant).
			000403F148	PM_MRK_DPTEG_FROM_DL2L3_SHR	A PTE was loaded into the TLB with shared (S) data from another chip's L2 or L3 cache on a different Node or Group (distant).
			000404F14C	PM_MRK_DPTEG_FROM_DMEM	A PTE was loaded into the TLB from another chip's memory on the same Node or Group (distant) due to a marked data-side request.
000401F142				PM_MRK_DPTEG_FROM_L2	A PTE was loaded into the TLB from local core's L2 cache due to a marked data-side request.
		000403F140		PM_MRK_DPTEG_FROM_L2_DISP_CONFLICT_LDHTST	A PTE was loaded into the TLB from the local core's L2 cache with a load-hit-store conflict due to a marked data-side request.
			000404F140	PM_MRK_DPTEG_FROM_L2_DISP_CONFLICT_OTHER	A PTE was loaded into the TLB from local core's L2 cache with dispatch conflict due to a marked data-side request.
	000402F140			PM_MRK_DPTEG_FROM_L2_MEFP	A PTE was loaded into the TLB from the local core's L2 hit without dispatch conflicts on an Mefp state due to a marked data-side request.
000401F140				PM_MRK_DPTEG_FROM_L2_NO_CONFLICT	A PTE was loaded into the TLB from local core's L2 cache without conflict due to a marked data-side request.
			000404F146	PM_MRK_DPTEG_FROM_L21_MOD	A PTE was loaded into the TLB with modified (M) data from another core's L2 cache on the same chip due to a marked data-side request.
		000403F146		PM_MRK_DPTEG_FROM_L21_SHR	A PTE was loaded into the TLB with shared (S) data from another core's L2 cache on the same chip due to a marked data-side request.

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 36 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
000401F14E				PM_MRK_DPTEG_FROM_L2MISS	A PTE was loaded into the TLB from a location other than the local core's L2 cache due to a marked data-side request.
			000404F142	PM_MRK_DPTEG_FROM_L3	A PTE was loaded into the TLB from local core's L3 cache due to a marked data-side request.
		000403F142		PM_MRK_DPTEG_FROM_L3_DISP_CONFLICT	A PTE was loaded into the TLB from the local core's L3 cache with a dispatch conflict due to a marked data-side request.
	000402F142			PM_MRK_DPTEG_FROM_L3_MEFP	A PTE was loaded into the TLB from the local core's L3 cache without dispatch conflicts hit on an Mfpf state due to a marked data-side request.
000401F144				PM_MRK_DPTEG_FROM_L3_NO_CONFLICT	A PTE was loaded into the TLB from local core's L3 cache without conflict due to a marked data-side request.
			000404F144	PM_MRK_DPTEG_FROM_L31_ECO_MOD	A PTE was loaded into the TLB with modified (M) data from another core's ECO L3 on the same chip due to a marked data-side request.
		000403F144		PM_MRK_DPTEG_FROM_L31_ECO_SR	A PTE was loaded into the TLB with shared (S) data from another core's ECO L3 on the same chip due to a marked data-side request.
	000402F144			PM_MRK_DPTEG_FROM_L31_MOD	A PTE was loaded into the TLB with modified (M) data from another core's L3 cache on the same chip due to a marked data-side request.
000401F146				PM_MRK_DPTEG_FROM_L31_SHR	A PTE was loaded into the TLB with shared (S) data from another core's L3 cache on the same chip due to a marked data-side request.
			000404F14E	PM_MRK_DPTEG_FROM_L3MISS	A PTE was loaded into the TLB from a location other than the local core's L3 cache due to a marked data-side request.
	000402F148			PM_MRK_DPTEG_FROM_LMEM	A PTE was loaded into the TLB from the local chip's memory due to a marked data-side request.
	000402F14C			PM_MRK_DPTEG_FROM_MEMORY	A PTE was loaded into the TLB from a memory location from a local remote or distant due to a marked data-side request.
			000404F14A	PM_MRK_DPTEG_FROM_OFF_CHIP_CACHE	A PTE was loaded into the TLB either shared or modified data from another core's L2 or L3 cache on a different chip (remote or distant) due to a marked data-side request.
000401F148				PM_MRK_DPTEG_FROM_ON_CHIP_CACHE	A PTE was loaded into the TLB with either shared or modified data from another core's L2 or L3 cache on the same chip due to a marked data-side request.
	000402F146			PM_MRK_DPTEG_FROM_RL2L3_MOD	A PTE was loaded into the TLB with modified (M) data from another chip's L2 or L3 cache on the same Node or Group (remote).
000401F14A				PM_MRK_DPTEG_FROM_RL2L3_SHR	A PTE was loaded into the TLB with shared (S) data from another chip's L2 or L3 cache on the same Node or Group (remote).
		000403F14A		PM_MRK_DPTEG_FROM_RMEM	A PTE was loaded into the TLB from another chip's memory on the same Node or Group (remote) due to a marked data-side request.
			00040401E4	PM_MRK_DTLB_MISS	Marked DTLB miss.
000401D158				PM_MRK_DTLB_MISS_16G	Marked data TLB miss page size 16 GB.
			000404D156	PM_MRK_DTLB_MISS_16M	Marked data TLB miss page size 16 MB.

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 37 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
	000402D156			PM_MRK_DTLB_MISS_4K	Marked data TLB miss page size 4 KB.
		000403D156		PM_MRK_DTLB_MISS_64K	Marked data TLB miss page size 64 KB.
		0004040154		PM_MRK_FAB_RSP_BKILL	Marked store had to do a bkill.
	000402F150			PM_MRK_FAB_RSP_BKILL_CYC	Number of cycles L2 RC took for a bkill.
		000403015E		PM_MRK_FAB_RSP_CLAIM_RTY	Sampled store did a rwitm and got a retry.
		0004030154		PM_MRK_FAB_RSP_DCLAIM	Marked store had to do a dclaim.
	000402F152			PM_MRK_FAB_RSP_DCLAIM_CYC	Number of cycles L2 RC took for a dclaim.
		0004030156		PM_MRK_FAB_RSP_MATCH	Ttype and cresp matched as specified in MMCR1.
		000404F152		PM_MRK_FAB_RSP_MATCH_CYC	Cresp/ttype match cycles.
		000404015E		PM_MRK_FAB_RSP_RD_RTY	Sampled L2 cache reads retry count.
000101015E				PM_MRK_FAB_RSP_RD_T_INTV	Sampled read got a T intervention.
			000404F150	PM_MRK_FAB_RSP_RWITM_CYC	Number of cycles an L2 RC took for a rwitm.
	000402015E			PM_MRK_FAB_RSP_RWITM_RTY	A sampled store did a rwitm resulting in a retry.
	000002013C	000003012E		PM_MRK_filt_MATCH	Marked filter match.
000001013C				PM_MRK_FIN_STALL_CYC	Number of marked instruction finish stall cycles (marked finish after NTC). Use edge detect to count the number.
	0000020134			PM_MRK_FXU_FIN	FXU marked instruction finish.
		0000040130		PM_MRK_GRP_CMPL	A marked instruction finished (completed).
		000004013A		PM_MRK_GRP_IC_MISS	Marked group experienced an I-cache miss.
		000003013C		PM_MRK_GRP_NTC	Marked group NTC cycles.
		00000401E0		PM_MRK_INST_CMPL	Marked instruction completed.
	0000020130			PM_MRK_INST_DECODED	Marked instruction decoded.
00000101E0				PM_MRK_INST_DISP	The thread has dispatched a randomly sampled marked instruction.
		0000030130		PM_MRK_INST_FIN	Marked instruction finished (any unit).
			00000401E6	PM_MRK_INST_FROM_L3MISS	Marked instruction was reloaded from a location beyond the local chiplet.
0000010132				PM_MRK_INST_ISSUED	Marked instruction issued.
			0000040134	PM_MRK_INST_TIMEOUT	Marked instruction finished, timeout (instruction lost).
00000101E4				PM_MRK_L1_ICACHE_MISS	Sampled instruction had an I-cache miss.

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 38 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
00010101EA				PM_MRK_L1_RELOAD_VALID	Marked demand reload.
	0004020114			PM_MRK_L2_RC_DISP	Marked instruction RC dispatched in L2 cache.
		000403012A		PM_MRK_L2_RC_DONE	Marked RC done.
			0011040116	PM_MRK_LARX_FIN	Larx finished.
000101013E				PM_MRK_LD_MISS_EXPOSED_CYC	Marked load exposed miss cycles. Use edge detect to count number.
000001013F				PM_MRK_LD_MISS_EXPOSED	Marked load exposed miss (exposed period ended, use edge detect to count number).
	00010201E2			PM_MRK_LD_MISS_L1	Marked data L1 cache demand miss counted at execution time.
			000104013E	PM_MRK_LD_MISS_L1_CYC	Marked load latency.
			0004040132	PM_MRK_LSU_FIN	An LSU marked instruction finished.
000400D180	000400D180	000400D180	000400D180	PM_MRK_LSU_FLUSH	Flush: (marked) all cases.
000400D188	000400D188	000400D188	000400D188	PM_MRK_LSU_FLUSH_LRQ	Flush: (marked) LRQ.
000400D18A	000400D18A	000400D18A	000400D18A	PM_MRK_LSU_FLUSH_SRQ	Flush: (marked) SRQ.
000400D184	000400D184	000400D184	000400D184	PM_MRK_LSU_FLUSH_ULD	Flush: (marked) unaligned load.
000400D186	000400D186	000400D186	000400D186	PM_MRK_LSU_FLUSH_UST	Flush: (marked) unaligned store.
			0004040164	PM_MRK_LSU_REJECT	LSU marked reject (up to two per cycle).
			0004030164	PM_MRK_LSU_REJECT_ERAT_MISS	LSU marked reject due to ERAT (up to two per cycle).
	0000020112			PM_MRK_NTF_FIN	Marked next-to-finish instruction as finished.
000001D15E				PM_MRK_RUN_CYC	Marked run cycles.
001501D15A				PM_MRK_SRC_PREF_TRACK_EFF	Marked source prefetch tracked was effective.
		001503D15A		PM_MRK_SRC_PREF_TRACK_INEFF	Prefetch tracked was ineffective for marked source.
			001504D15C	PM_MRK_SRC_PREF_TRACK_MOD	Prefetch tracked was moderate for marked source.
001501D15C				PM_MRK_SRC_PREF_TRACK_MOD_L2	Prefetch tracked was moderate for marked source (source L2 cache).
		001503D15C		PM_MRK_SRC_PREF_TRACK_MOD_L3	Prefetch tracked was moderate (L3 hit) for marked source.
0004010134				PM_MRK_ST_CMPL ¹	A marked store completed in the L2 cache (RC machine done).
			00040301E2		
			0004030134		Marked store complete (data home) with intervention.
			000403F150	PM_MRK_ST_DRAIN_TO_L2DISP_CYC	Number of cycles it takes to drain a store from the SRQ to the L2 cache.

1. There are multiple methods to program this event.

Table D-1. POWER8 Event List by Event Name (Sheet 39 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
		000403012C		PM_MRK_ST_FWD	Marked store forwards.
000401F150				PM_MRK_ST_L2DISP_TO_CMPL_CYC	Cycles from L2 read/claim dispatch to L2 read/claim completion.
	0004020138			PM_MRK_ST_NEST	Marked store sent to nest.
		000003013E		PM_MRK_STALL_CMPLU_CYC	Marked group completion stall cycles (use edge detect to count number of cycles).
		001103E158		PM_MRK_STCX_FAIL	Marked stcx failed.
001501C15A				PM_MRK_TGT_PREF_TRACK_EFF	Marked target prefetch track was effective.
		001503C15A		PM_MRK_TGT_PREF_TRACK_INEFF	Prefetch tracked was ineffective for marked target.
			001504C15C	PM_MRK_TGT_PREF_TRACK_MOD	Prefetch tracked was moderate for marked target.
001501C15C				PM_MRK_TGT_PREF_TRACK_MOD_L2	Marked target prefetch tracked was moderate (source L2 cache).
		001503C15C		PM_MRK_TGT_PREF_TRACK_MOD_L3	Prefetch tracked was moderate (L3 hit) for marked target.
		0000030132		PM_MRK_VSU_FIN	VSU (FPU) marked instruction finished.
		000003D15E		PM_MULT_MRK	Multiple marked instructions.
		000003006E		PM_NEST_REF_CLK	This event increments at the frequency of the SMP interconnect. Multiply by 4 to obtain the number of SMP interconnect cycles.
00000020B0	00000020B0	00000020B0	00000020B0	PM_NESTED_TEND	Completion time nested tend .
00000020B6	00000020B6	00000020B6	00000020B6	PM_NON_FAV_TBEGIN	Dispatch time nonfavored tbegin .
	0000328084			PM_NON_TM_RST_SC	A snoop from another core that was not working on a transaction hit on an L3 cache line that was in the SC state for a transaction that was no longer in flight. In this case, the L3 cache honors the snoop request and gives away the line after resetting the SC state.
	000002001A			PM_NTCCG_ALL_FIN	Cycles after all instructions have finished to group completed.
00000020AC	00000020AC	00000020AC	00000020AC	PM_OUTER_TBEGIN	Completion time outer tbegin .
00000020AE	00000020AE	00000020AE	00000020AE	PM_OUTER_TEND	Completion time outer tend .
	0000020010			PM_PMC1_OVERFLOW	Overflow from counter 1.
		0000030010		PM_PMC2_OVERFLOW	Overflow from counter 2.
		0000030020		PM_PMC2_REWIND	PMC2 rewind event (does not match condition).
0000010022			0000040010	PM_PMC2_SAVED	PMC2 rewind value saved (matched condition).
			PM_PMC3_OVERFLOW	Overflow from counter 3.	

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 40 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
0000010010				PM_PMC4_OVERFLOW	Overflows from PMC4 are counted. This effectively widens the PMC. The overflow from the original PMC does not trigger an exception even if the PMU is configured to generate exceptions on overflow.
0000010020				PM_PMC4_REWIND	PMC4 rewind event (did not match condition).
	0000030022			PM_PMC4_SAVED	PMC4 rewind value saved (matched condition).
0000010024				PM_PMC5_OVERFLOW	Overflow from counter 5. Overflows from PMC5 are counted. This effectively widens the PMC. The overflow from the original PMC does not trigger an exception even if the PMU is configured to generate exceptions on overflow.
	0000030024			PM_PMC6_OVERFLOW	Overflow from counter 6. Overflows from PMC6 are counted. This effectively widens the PMC. The overflow from the original PMC does not trigger an exception even if the PMU is configured to generate exceptions on overflow.
000001005A				PM_PREF_TRACK_EFF	Prefetch tracked was effective.
	000003005A			PM_PREF_TRACK_INEFF	Prefetch tracked was ineffective.
		000004005A		PM_PREF_TRACK_MOD	Prefetch tracked was moderate.
000001005C				PM_PREF_TRACK_MOD_L2	Prefetch tracked was moderate (source L2 cache).
	000003005C			PM_PREF_TRACK_MOD_L3	Prefetch tracked was moderate (L3 cache).
	000002005A			PM_PREF_TRACKED	Total number of prefetch operations that were tracked.
		0000040014		PM_PROBE_NOP_DISP	Probe NOPs dispatched.
000000E084	000000E084	000000E084	000000E084	PM_PTE_PREFETCH	PTE prefetches.
0000010054				PM_PUMP_CPRED	Pump prediction correct. Counts across all types of pumps for all data types excluding data prefetch (demand load).
		0000040052		PM_PUMP_MPRED	Pump misprediction. Counts across all types of pumps for all data types excluding data prefetch (demand load).
	DE200301EA			PM_RC_LIFETIME_EXC_1024	Number of times the RC machine for a sampled instruction was active for more than 1024 cycles. Reload latency exceeded 1024 cycles.
		DE200401EC		PM_RC_LIFETIME_EXC_2048	Number of times the RC machine for a sampled instruction was active for more than 2048 cycles.
DE200101E8				PM_RC_LIFETIME_EXC_256	Number of times the RC machine for a sampled instruction was active for more than 256 cycles.
	DE200201E6			PM_RC_LIFETIME_EXC_32	Number of times the RC machine for a sampled instruction was active for more than 32 cycles. Reload latency exceeded 32 cycles.
	0000036088			PM_RC_USAGE ¹	Continuous 16-cycle (2:1) window where this signal rotates through sampling each L2 RC machine busy. The PMU uses this wave to then do a 16-cycle count to sample the total number of machines running.
	0000736088				
0000016081				PM_RC0_ALLOC	RC machine 0 busy. Used by the PMU to sample average RC lifetime (mach0 used as sample point).

1. There are multiple methods to program this event.

**Table D-1. POWER8 Event List by Event Name (Sheet 41 of 46)**

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
0000016080				PM_RC0_BUSY ¹	RC machine 0 (mach0) busy. Used by the PMU to sample average RC lifetime (mach0 is used as a sample point).
0000716080				PM_RD_CLEARING_SC	Read clearing SC.
		000034808E		PM_RD_FORMING_SC	Read forming SC.
		000034808C		PM_RD_HIT_PF	Read machine hit in L3 prefetch machine.
	0000020004			PM_REAL_SRQ_FULL	Out of real SRQ entries.
	00000200F4			PM_RUN_CYC	Run cycles. Processor cycles gated by the run latch. Operating systems use the run latch to indicate when they are doing useful work. The run latch is typically cleared in the operating system idle loop. Gating by the run latch filters out the idle loop.
		000003006C		PM_RUN_CYC_SMT2_MODE	Number of cycles run latch is set and core is in SMT2 mode.
	000002006A			PM_RUN_CYC_SMT2_SHRD_MODE	Number of cycles that this thread's run latch is set and the core is in SMT2-shared mode.
000001006A				PM_RUN_CYC_SMT2_SPLIT_MODE	Number of cycles that the run latch is set and the core is in SMT2-split mode.
	000002006C			PM_RUN_CYC_SMT4_MODE	Number of cycles that this thread's run latch is set and the core is in SMT4 mode.
		000004006C		PM_RUN_CYC_SMT8_MODE	Number of cycles that the run latch is set and the core is in SMT8 mode.
000001006C				PM_RUN_CYC_ST_MODE	Number of cycles that the run latch is set and the core is in ST mode.
		00000400FA		PM_RUN_INST_CMPL	Number of PowerPC instructions that completed; gated by the run latch.
		00000400F4		PM_RUN_PURR	Run PURR.
0000010008				PM_RUN_SPURR	Run SPURR. The Scaled Processor Utilization of Resources Register (SPURR) was incremented while the run latch was set.
000000F082	000000F082	000000F082	000000F082	PM_SEC_ERAT_HIT	Secondary ERAT hit.
000000508C	000000508C	000000508C	000000508C	PM_SHL_CREATED	Store-hit-load table entry created.
000000508E	000000508E	000000508E	000000508E	PM_SHL_ST_CONVERT	Store-hit-load table read hit with entry enabled.
0000005090	0000005090	0000005090	0000005090	PM_SHL_ST_DISABLE	Store-hit-load table read hit with entry disabled.
		000004608C		PM_SN_USAGE ¹	Continuous 16-cycle (2:1) window where this signal rotates through sampling each L2 SN machine busy. The PMU uses this wave to then do a 16-cycle count to sample the total number of machines running.
		000074608C			
0000026085				PM_SNO_ALLOC	SN mach 0 busy. Used by the PMU to sample average RC lifetime (mach0 is used as a sample point).
0000026084				PM_SNO_BUSY ¹	SN mach0 busy. Used by the PMU to sample average RC lifetime (mach0 is used as a sample point).
0000726084					
000000D0B2	000000D0B2	000000D0B2	000000D0B2	PM_SNOOP_TLBIE	TLBIE snoop.

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 42 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
		0000338088		PM_SNP_TM_HIT_M	A snoop from another core for a TM store hit a line in the M or Mu state.
		000033808A		PM_SNP_TM_HIT_T	A snoop from another core for a TM store hit a line in the T, Tn, or Te state.
0000717080				PM_ST CAUSED FAIL	A non-TM store caused any thread to fail.
	0000020016			PM_ST_CMPL	Store completion count.
	00000200F0			PM_ST_FIN	Store instructions finished (store sent to nest).
	0000020018			PM_ST_FWD	Store forwards that finished.
		00000300F0		PM_ST_MISS_L1	Store missed L1.
0000010028				PM_STALL_END_GCT_EMPTY	Count ended because GCT went empty.
000001E058				PM_STCX_FAIL	Stcx failed.
000000C090	000000C090	000000C090	000000C090	PM_STCX_LSU	Stcx executed reported at sent to nest.
0000010000	0000020000	0000030000	0000040000	PM_SUSPENDED	Counter off. The counter is suspended (does not count).
0000003090	0000003090	0000003090	0000003090	PM_SWAP_CANCEL	Swap cancel.
0000003092	0000003092	0000003092	0000003092	PM_SWAP_CANCEL_GPR	Swap cancel.
000000308C	000000308C	000000308C	000000308C	PM_SWAP_COMPLETE	Swap cast-in completed.
000000308E	000000308E	000000308E	000000308E	PM_SWAP_COMPLETE_GPR	Swap cast-in completed for <u>GPR</u> .
0000015152				PM_SYNC_MRK_BR_LINK	Marked branch and link branch that can cause a synchronous interrupt.
000201515C				PM_SYNC_MRK_BR_MPRED	Marked branch mispredict that can cause a synchronous interrupt.
0018015156				PM_SYNC_MRK_FX_DIVIDE	Marked fixed-point divide that can cause a synchronous interrupt.
0001015158				PM_SYNC_MRK_L2HIT	Marked L2 hits that can throw a synchronous interrupt.
000101515A				PM_SYNC_MRK_L2MISS	Marked L2 miss that can throw a synchronous interrupt.
0001015154				PM_SYNC_MRK_L3MISS	Marked L3 misses that can throw a synchronous interrupt.
0008015150				PM_SYNC_MRK_PROBE_NOP	Marked probe NOPs that can cause synchronous interrupts.
		0000030050		PM_SYS_PUMP_CPRED	Initial and final pump scope was system pump for all data types excluding data prefetch (demand load).
		0000030052		PM_SYS_PUMP_MPRED	Final pump scope (system) mispredicted. Either the original scope was too small (chip/group) or the original scope was system and it should have been smaller. Number of counts for all data types excluding data prefetch (demand load).
			0000040050	PM_SYS_PUMP_MPRED_RTY	Final pump scope (system) was larger than initial pump scope (chip/group) for all data types excluding data prefetch (demand load).
0000010026				PM_TABLEWALK_CYC	Number of cycles when a tablewalk (instruction or data) is active.

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 43 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
000000E086	000000E086	000000E086	000000E086	PM_TABLEWALK_CYC_PREF	Tablewalk qualified for PTE prefetches.
00000020B2	00000020B2	00000020B2	00000020B2	PM_TABORT_TRECLAIM	Completion time treclaim .
		00000300F8		PM_TB_BIT_TRANS	Timebase event.
000000E0BA	000000E0BA	000000E0BA	000000E0BA	PM_TEND_PEND_CYC	TEND latency per thread.
	000002000C			PM_THRD_ALL_RUN_CYC	All threads in run_cycles.
		00000300F4		PM_THRD_CONC_RUN_INST	PowerPC instructions finished when both threads are in run_cycles; concurrent run instructions.
0000010012				PM_THRD_GRP_CMPL_BOTH_CYC	Cycles group completed on both completion slots by any thread. Two threads finished in the same cycle (gated by a run latch).
00000040BC	00000040BC	00000040BC	00000040BC	PM_THRD_PRIO_0_1_CYC	Number of cycles the thread is running at priority level 0 or 1.
00000040BE	00000040BE	00000040BE	00000040BE	PM_THRD_PRIO_2_3_CYC	Number of cycles the thread is running at priority level 2 or 3.
0000005080	0000005080	0000005080	0000005080	PM_THRD_PRIO_4_5_CYC	Number of cycles the thread is running at priority level 4 or 5.
0000005082	0000005082	0000005082	0000005082	PM_THRD_PRIO_6_7_CYC	Number of cycles the thread is running at priority level 6 or 7.
0000003098	0000003098	0000003098	0000003098	PM_THRD_REBAL_CYC	Number of cycles rebalance is active.
	00000301EA			PM_THRESH_EXC_1024	Threshold counter exceeded a value of 1024.
		00000401EA		PM_THRESH_EXC_128	Threshold counter exceeded a value of 128.
		00000401EC		PM_THRESH_EXC_2048	Threshold counter exceeded a value of 2048.
00000101E8				PM_THRESH_EXC_256	Threshold counter exceeded a count of 256.
	00000201E6			PM_THRESH_EXC_32	Threshold counter exceeded a value of 32.
00000101E6				PM_THRESH_EXC_4096	Threshold counter exceeded a count of 4096.
	00000201E8			PM_THRESH_EXC_512	Threshold counter exceeded a value of 512.
	00000301E8			PM_THRESH_EXC_64	Threshold counter exceeded a value of 64.
00000101EC				PM_THRESH_MET	Threshold exceeded.
		000004016E		PM_THRESH_NOT_MET	Threshold counter did not meet threshold.
	0000020066			PM_TLB_MISS	TLB miss (instruction and data).
		0000030058		PM_TLBIE_FIN	Tlbie finished.
00000020B8	00000020B8	00000020B8	00000020B8	PM_TM_BEGIN_ALL	TM any tbegin .
0000318082				PM_TM_CAM_OVERFLOW	L3 transactional memory CAM overflow during L2 castout of SC.
		000074708C		PM_TM_CAP_OVERFLOW	TM footprint capacity overflow.

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 44 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
00000020BA	00000020BA	00000020BA	00000020BA	PM_TM_END_ALL	TM any tend .
0000003088	0000003088	0000003088	0000003088	PM_TM_FAIL_CON_TM	TEXAS fail reason at completion.
0000003086	0000003086	0000003086	0000003086	PM_TM_FAIL_CONF_NON_TM	TEXAS fail reason at completion.
000000E0B2	000000E0B2	000000E0B2	000000E0B2	PM_TM_FAIL_DISALLOW	TM fail disallowed.
0000003084	0000003084	0000003084	0000003084	PM_TM_FAIL_FOOTPRINT_OVERFLOW	TEXAS fail reason at completion.
000000E0B8	000000E0B8	000000E0B8	000000E0B8	PM_TM_FAIL_NON_TX_CONFLICT	Nontransactional conflict from the LSU as reported to the TEXAS Register.
000000308A	000000308A	000000308A	000000308A	PM_TM_FAIL_SELF	TEXAS fail reason at completion.
000000E0B4	000000E0B4	000000E0B4	000000E0B4	PM_TM_FAIL_TLBIE	TLBIE hit bloom filter.
000000E0B6	000000E0B6	000000E0B6	000000E0B6	PM_TM_FAIL_TX_CONFLICT	Transactional conflict from the LSU.
	0000727086			PM_TM_FAV CAUSED FAIL	TM load (favored) caused another thread to fail.
0000717082				PM_TM_LD CAUSED FAIL	A non-TM load caused any thread to fail.
	0000727084			PM_TM_LD_CONF	A TM load (favored or nonfavored) ran into conflict (failed).
	0000328086			PM_TM_RST_SC	A snoop from another core that is working on a transaction hit on an L3 cache line that was in the SC state for a transaction that was no longer in flight. In this case, the L3 cache honors the snoop request and gives away the line after resetting the SC state.
0000318080				PM_TM_SC_CO	L3 cast out of a backup store image. During a transaction, the L2 cache sends a back-up image to the L3 cache, so that if the transaction fails, the state of the line can be recovered. This event counts the number of times one of these back-up image lines is evicted from the L3 cache.
	000073708A			PM_TM_ST CAUSED FAIL	TM store (favored or nonfavored) caused another thread to fail.
	0000737088			PM_TM_ST_CONF	TM store (favored or nonfavored) ran into conflict (failed).
00000020BC	00000020BC	00000020BC	00000020BC	PM_TM_TBEGIN	TM nested tbegin .
0000010060				PM_TM_TRANS_RUN_CYC	Run cycles in transactional state.
	0000030060			PM_TM_TRANS_RUN_INST	Instructions completed in transactional state.
0000003080	0000003080	0000003080	0000003080	PM_TM_TRESUME	TM resume.
00000020BE	00000020BE	00000020BE	00000020BE	PM_TM_TSUSPEND	TM suspend.
	000002E012			PM_TM_TX_PASS_RUN_CYC	Run cycles spent in successful transactions.
		000004E014		PM_TM_TX_PASS_RUN_INST	Run instructions spent in successful transactions.
000000E08C	000000E08C	000000E08C	000000E08C	PM_UP_PREF_L3	Micropartition prefetch.
000000E08E	000000E08E	000000E08E	000000E08E	PM_UP_PREF_POINTER	Micropartition pointer prefetches.

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 45 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
	000002405E			PM_UTHROTTLE	Cycles in which an instruction issue throttle was active in the ISU.
000000A0A4	000000A0A4	000000A0A4	000000A0A4	PM_VSU0_16FLOP	Sixteen flops operation (SP vector versions of fdiv).
000000A080	000000A080	000000A080	000000A080	PM_VSU0_1FLOP	One flop (fadd).
000000A098	000000A098	000000A098	000000A098	PM_VSU0_2FLOP	Two flops operation (scalar fmadd).
000000A09C	000000A09C	000000A09C	000000A09C	PM_VSU0_4FLOP	Four flops operation (scalar fdiv).
000000A0A0	000000A0A0	000000A0A0	000000A0A0	PM_VSU0_8FLOP	Eight flops operation (DP vector versions of fdiv).
000000B0A4	000000B0A4	000000B0A4	000000B0A4	PM_VSU0_COMPLEX_ISSUED	Complex VMX instruction issued.
000000B0B4	000000B0B4	000000B0B4	000000B0B4	PM_VSU0_CY_ISSUED	Cryptographic instruction.
000000B0A8	000000B0A8	000000B0A8	000000B0A8	PM_VSU0_DD_ISSUED	64-bit decimal issued.
000000A08C	000000A08C	000000A08C	000000A08C	PM_VSU0_DP_2FLOP	DP vector version of fmul .
000000A090	000000A090	000000A090	000000A090	PM_VSU0_DP_FMA	DP vector version of fmadd .
000000A094	000000A094	000000A094	000000A094	PM_VSU0_DP_FSQRT_FDIV	DP vector versions of fdiv .
000000B0AC	000000B0AC	000000B0AC	000000B0AC	PM_VSU0_DQ_ISSUED	128-bit decimal issued.
000000B0B0	000000B0B0	000000B0B0	000000B0B0	PM_VSU0_EX_ISSUED	Direct move 32- or 64-bit VRF-to-GPR.
000000A0BC	000000A0BC	000000A0BC	000000A0BC	PM_VSU0_FIN	VSU0 finished an instruction.
000000A084	000000A084	000000A084	000000A084	PM_VSU0_FMA	Two flops operation (fmadd).
000000B098	000000B098	000000B098	000000B098	PM_VSU0_FPSCR	Move to/from an FPSCR type instruction issued on Pipe0.
000000A088	000000A088	000000A088	000000A088	PM_VSU0_FSQRT_FDIV	Four flops operation (fdiv).
000000B090	000000B090	000000B090	000000B090	PM_VSU0_PERMUTE_ISSUED	Permute VMX instruction Issued.
000000B088	000000B088	000000B088	000000B088	PM_VSU0_SCALAR_DP_ISSUED	Double-precision scalar instruction issued on Pipe0.
000000B094	000000B094	000000B094	000000B094	PM_VSU0_SIMPLE_ISSUED	Simple VMX instruction issued.
000000A0A8	000000A0A8	000000A0A8	000000A0A8	PM_VSU0_SINGLE	FPU single precision.
000000B09C	000000B09C	000000B09C	000000B09C	PM_VSU0_SQ	Store vector issued.
000000B08C	000000B08C	000000B08C	000000B08C	PM_VSU0_STF	FPU store (SP or DP) issued on Pipe0.
000000B080	000000B080	000000B080	000000B080	PM_VSU0_VECTOR_DP_ISSUED	Double-precision vector instruction issued on Pipe0.
000000B084	000000B084	000000B084	000000B084	PM_VSU0_VECTOR_SP_ISSUED	Single-precision vector instruction issued (executed).
000000A0A6	000000A0A6	000000A0A6	000000A0A6	PM_VSU1_16FLOP	Sixteen flops operation (SP vector versions of fdiv).
000000A082	000000A082	000000A082	000000A082	PM_VSU1_1FLOP	One flop (fadd).

1. There are multiple methods to program this event.



Table D-1. POWER8 Event List by Event Name (Sheet 46 of 46)

PMC1 (Hexadecimal)	PMC2 (Hexadecimal)	PMC3 (Hexadecimal)	PMC4 (Hexadecimal)	Event Name	Event Description
000000A09A	000000A09A	000000A09A	000000A09A	PM_VSU1_2FLOP	Two flops operation (scalar fmad).
000000A09E	000000A09E	000000A09E	000000A09E	PM_VSU1_4FLOP	Four flops operation (scalar fdiv).
000000A0A2	000000A0A2	000000A0A2	000000A0A2	PM_VSU1_8FLOP	Eight flops operation (DP vector versions of fdiv).
000000B0A6	000000B0A6	000000B0A6	000000B0A6	PM_VSU1_COMPLEX_ISSUED	Complex VMX instruction issued.
000000B0B6	000000B0B6	000000B0B6	000000B0B6	PM_VSU1_CY_ISSUED	Cryptographic instruction.
000000B0AA	000000B0AA	000000B0AA	000000B0AA	PM_VSU1_DD_ISSUED	64-bit decimal issued.
000000A08E	000000A08E	000000A08E	000000A08E	PM_VSU1_DP_2FLOP	DP vector version of fmul .
000000A092	000000A092	000000A092	000000A092	PM_VSU1_DP_FMA	DP vector version of fmad .
000000A096	000000A096	000000A096	000000A096	PM_VSU1_DP_FSQRT_FDIV	DP vector versions of fdiv .
000000B0AE	000000B0AE	000000B0AE	000000B0AE	PM_VSU1_DQ_ISSUED	128-bit decimal issued.
000000B0B2	000000B0B2	000000B0B2	000000B0B2	PM_VSU1_EX_ISSUED	Direct move 32- or 64-bit VRF-to-GPR.
000000A0BE	000000A0BE	000000A0BE	000000A0BE	PM_VSU1_FIN	VSU1 finished an instruction.
000000A086	000000A086	000000A086	000000A086	PM_VSU1_FMA	Two flops operation (fmad).
000000B09A	000000B09A	000000B09A	000000B09A	PM_VSU1_FPSCR	Move to/from FPSCR type instruction issued on Pipe0.
000000A08A	000000A08A	000000A08A	000000A08A	PM_VSU1_FSQRT_FDIV	Four flops operation (fdiv).
000000B092	000000B092	000000B092	000000B092	PM_VSU1_PERMUTE_ISSUED	Permute VMX instruction Issued.
000000B08A	000000B08A	000000B08A	000000B08A	PM_VSU1_SCALAR_DP_ISSUED	Double-precision scalar instruction issued on Pipe1.
000000B096	000000B096	000000B096	000000B096	PM_VSU1_SIMPLE_ISSUED	Simple VMX instruction issued.
000000A0AA	000000A0AA	000000A0AA	000000A0AA	PM_VSU1_SINGLE	FPU single precision.
000000B09E	000000B09E	000000B09E	000000B09E	PM_VSU1_SQ	Store vector issued.
000000B08E	000000B08E	000000B08E	000000B08E	PM_VSU1_STF	FPU store (SP or DP) issued on Pipe1.
000000B082	000000B082	000000B082	000000B082	PM_VSU1_VECTOR_DP_ISSUED	Double-precision vector instruction issued on Pipe1.
000000B086	000000B086	000000B086	000000B086	PM_VSU1_VECTOR_SP_ISSUED	Single-precision vector instruction issued (executed).

1. There are multiple methods to program this event.





Appendix E. SPMC Performance Monitoring Events

Appendix E lists all of the SPMC performance monitoring events for the POWER8 processor by event name.

There are two programmable SPMCs. Each SPMC has access to a unique set of events. The desired events must be programmed and read from the specified SPMC.

The SPMC_xSEL bits correspond to MMCRS[SPMC_xSEL], where x is the SPMC number specified in the SPMC column.

The Unit column provides information about which core unit an event comes from.

The Select Event is a unique identifier for an event, consumable by certain tools.

Table E-1. SPMC Performance Monitoring Events (Sheet 1 of 7)

SPMC	SPMC _x SEL(0:3) (Binary)	SPMC _x SEL(4:6) (Binary)	Event Name	Description	Unit	Select Event (Hexadecimal)
SPMC1	1110	001	SPM_1PLUS_PPC_CMPL	One or more PowerPC instructions finished (completed).	ISU	00E2
SPMC2	1111	001	SPM_1PLUS_PPC_DISP	Cycles at least one instruction dispatched.	ISU	00F2
SPMC1	0100	100	SPM_BR_CMPL	Branch instruction finished.	IFU	0048
SPMC2	1111	011	SPM_BR_MPRED_CMPL	Number of branch mispredicts.	IFU	00F6
SPMC1	0100	101	SPM_BR_TAKEN_CMPL	Branch taken.	IFU	004A
SPMC1	0011	000	SPM_CHIP_PUMP_CPRED	Initial and final pump scope and data sourced across this scope was chip pump (prediction = correct) for all data types (demand load, instruction fetch, instruction or data translation).	LSU	0030
SPMC2	0011	001	SPM_CMPLU_STALL	Completion stall (any reason).	ISU	0032
SPMC1	1110	000	SPM_CYC	Cycles.	PMU	00E0
SPMC2	0100	000	SPM_DATA_FROM_DL2L3_MOD	The processor's data cache was reloaded with modified (M) data from another chip's L2 or L3 cache on a different node or group (distant) from this chip due to a demand load.	LSU	0040
SPMC2	0001	001	SPM_DATA_FROM_DL2L3_SHR	The processor's data cache was reloaded with shared (S) data from another chip's L2 or L3 cache on a different node or group (distant) from this chip due to a demand load.	LSU	0012
SPMC2	0100	010	SPM_DATA_FROM_DMEM	The processor's data cache was reloaded from another chip's memory on the same node or group (distant) due to a demand load.	LSU	0044
SPMC1	0001	000	SPM_DATA_FROM_L2	The processor's data cache was reloaded from the local core's L2 cache due to a demand load.	LSU	0010
SPMC2	0000	101	SPM_DATA_FROM_L2_DISP_CONFLICT_LDHTST	The processor's data cache was reloaded from the local core's L2 cache with a load-hit-store conflict due to a demand load.	LSU	000A
SPMC2	0011	100	SPM_DATA_FROM_L2_DISP_CONFLICT_OTHER	The processor's data cache was reloaded from the local core's L2 cache with a dispatch conflict due to a demand load.	LSU	0038

**Table E-1. SPMC Performance Monitoring Events (Sheet 2 of 7)**

SPMC	SPMCxSEL(0:3) (Binary)	SPMCxSEL(4:6) (Binary)	Event Name	Description	Unit	Select Event (Hexadecimal)
SPMC1	0101	010	SPM_DATA_FROM_L2_MEPF	The processor's data cache was reloaded from the local core's L2 hit without a dispatch conflicts on an Mepf state due to a demand load.	LSU	0054
SPMC1	0000	111	SPM_DATA_FROM_L2_NO_CONFLICT	The processor's data cache was reloaded from the local core's L2 cache without a conflict due to a demand load.	LSU	000E
SPMC2	0011	111	SPM_DATA_FROM_L2.1_MOD	The processor's data cache was reloaded with modified (M) data from another core's L2 cache on the same chip due to a demand load.	LSU	003E
SPMC2	0001	000	SPM_DATA_FROM_L2.1_SHR	The processor's data cache was reloaded with shared (S) data from another core's L2 cache on the same chip due to a demand load.	LSU	0010
SPMC1	1111	111	SPM_DATA_FROM_L2MISS	Demand load from an L2 miss (not an L2 hit).	LSU/PMU	00FE
SPMC2	0011	101	SPM_DATA_FROM_L3	The processor's data cache was reloaded from the local core's L3 cache due to a demand load.	LSU	003A
SPMC2	0000	110	SPM_DATA_FROM_L3_DISP_CONFLICT	The processor's data cache was reloaded from the local core's L3 cache with a dispatch conflict due to a demand load.	LSU	000C
SPMC1	0101	011	SPM_DATA_FROM_L3_MEPF	The processor's data cache was reloaded from the local core's L3 cache without a dispatch conflict hit on an Mepf state due to a demand load.	LSU	0056
SPMC1	0001	001	SPM_DATA_FROM_L3_NO_CONFLICT	The processor's data cache was reloaded from the local core's L3 cache without a conflict due to a demand load.	LSU	0012
SPMC2	0011	110	SPM_DATA_FROM_L3.1_ECO_MOD	The processor's data cache was reloaded with modified (M) data from another core's ECO L3 on the same chip due to a demand load.	LSU	003C
SPMC2	0000	111	SPM_DATA_FROM_L3.1_ECO_SHR	The processor's data cache was reloaded with shared (S) data from another core's ECO L3 on the same chip due to a demand load.	LSU	000E
SPMC1	0101	100	SPM_DATA_FROM_L3.1_MOD	The processor's data cache was reloaded with modified (M) data from another core's L3 cache on the same chip due to a demand load.	LSU	0058
SPMC1	0001	010	SPM_DATA_FROM_L3.1_SHR	The processor's data cache was reloaded with shared (S) data from another core's L3 cache on the same chip due to a demand load.	LSU	0014
SPMC2	1110	111	SPM_DATA_FROM_L3MISS	Demand load from an L3 miss (not an L2 hit and not an L3 hit).	LSU	00EE
SPMC1	0101	110	SPM_DATA_FROM_LMEM	The processor's data cache was reloaded from the local chip's memory due to a demand load.	LSU	005C
SPMC1	0110	000	SPM_DATA_FROM_MEMORY	The processor's data cache was reloaded from a memory location due to a demand load.	LSU	0060
SPMC2	1111	111	SPM_DATA_FROM_MEMORY	Data cache reload from memory.	LSU	00FE
SPMC2	0100	001	SPM_DATA_FROM_OFF_CHIP_CACHE	The processor's data cache was reloaded with either shared or modified data from another core's L2/L3 cache on a different chip (remote or distant) due to a demand load.	LSU	0042
SPMC1	0001	011	SPM_DATA_FROM_ON_CHIP_CACHE	The processor's data cache was reloaded with either shared or modified data from another core's L2 or L3 cache on the same chip due to a demand load.	LSU	0016
SPMC1	0101	101	SPM_DATA_FROM_RL2L3_MOD	The processor's data cache was reloaded with modified (M) data from another chip's L2 or L3 cache on the same node or group (remote) as this chip due to a demand load.	LSU	005A



Table E-1. SPMC Performance Monitoring Events (Sheet 3 of 7)

SPMC	SPMCxSEL(0:3) (Binary)	SPMCxSEL(4:6) (Binary)	Event Name	Description	Unit	Select Event (Hexadecimal)
SPMC1	0001	100	SPM_DATA_FROM_RL2L3_SHR	The processor's data cache was reloaded with shared (S) data from another chip's L2 or L3 cache on the same node or group (remote) as this chip due to a demand load.	LSU	0018
SPMC2	0001	010	SPM_DATA_FROM_RMEM	The processor's data cache was reloaded from another chip's memory on the same node or group (remote) due to a demand load.	LSU	0014
SPMC2	0000	011	SPM_DATA_TABLEWALK_CYC	Data tablewalk active.	LSU	0006
SPMC2	0110	000	SPM_DERAT_MISS_16G	Data ERAT miss (data TLB Access) page size 16 GB.	LSU	0060
SPMC2	0010	000	SPM_DERAT_MISS_16M	Data ERAT miss (data TLB access) page size 16 MB.	LSU	0020
SPMC1	0011	101	SPM_DERAT_MISS_4K	Data ERAT miss (data TLB access) page size 4 KB.	LSU	003A
SPMC1	0111	000	SPM_DERAT_MISS_64K	Data ERAT miss (data TLB access) page size 64 KB.	LSU	0070
SPMC2	0000	010	SPM_DISP_HELD_SRQ_FULL	Dispatch held because the SRQ is full.	ISU	0004
SPMC1	0010	101	SPM_DPTEG_FROM_L2	A page table entry was loaded into the TLB from the local core's L2 cache due to a data-side request.	LSU	002A
SPMC1	0010	111	SPM_DPTEG_FROM_L2MISS	A page table entry was loaded into the TLB from a location other than the local core's L2 cache due to a data-side request.	LSU	002E
SPMC2	0101	010	SPM_DPTEG_FROM_L3	A page table entry was loaded into the TLB from the local core's L3 cache due to a data-side request.	LSU	0054
SPMC2	0101	100	SPM_DPTEG_FROM_L3MISS	A page table entry was loaded into the TLB from a location other than the local core's L3 cache due to a data-side request.	LSU	0058
SPMC1	0110	101	SPM_DPTEG_FROM_LMEM	A page table entry was loaded into the TLB from the local chip's memory due to a data-side request.	LSU	006A
SPMC1	0110	110	SPM_DPTEG_FROM_MEMORY	A page table entry was loaded into the TLB from a memory location due to a data-side request.	LSU	006C
SPMC2	0101	011	SPM_DPTEG_FROM_OFF_CHIP_CACHE	A page table entry was loaded into the TLB with either shared or modified data from another core's L2 or L3 cache on a different chip (remote or distant) due to a data-side request.	LSU	0056
SPMC1	0010	110	SPM_DPTEG_FROM_ON_CHIP_CACHE	A page table entry was loaded into the TLB with either shared or modified data from another core's L2 or L3 cache on the same chip due to a data-side request.	LSU	002C
SPMC1	1110	111	SPM_DSLB_MISS	Data SLB miss.	ISU	00EE
SPMC2	1110	110	SPM_DTLB_MISS	Data PTEG reloaded (DTLB miss).	LSU	00EC
SPMC1	0011	110	SPM_DTLB_MISS_16G	Data TLB miss page size 16 GB.	LSU	003C
SPMC2	0110	001	SPM_DTLB_MISS_16M	Data TLB miss page size 16 MB.	LSU	0062
SPMC1	0111	001	SPM_DTLB_MISS_4K	Data TLB miss page size 4 KB.	LSU	0072
SPMC2	0010	001	SPM_DTLB_MISS_64K	Data TLB miss page size 64 KB.	LSU	0022
SPMC1	1111	100	SPM_EXT_INT	External interrupt.	ISU	00F8



Table E-1. SPMC Performance Monitoring Events (Sheet 4 of 7)

SPMC	SPMCxSEL(0:3) (Binary)	SPMCxSEL(4:6) (Binary)	Event Name	Description	Unit	Select Event (Hexadecimal)
SPMC1	1110	010	SPM_FLOP	Floating-point operations finished.	VSU	00E4
SPMC2	1111	100	SPM_FLUSH	Flush (any type).	ISU	00F8
SPMC2	0011	010	SPM_FREQ_UP	Frequency is slewed up due to power management.	PC	0034
SPMC1	0101	000	SPM_RXU_BUSY	RXU0 busy and RXU1 busy.	FXU	0050
SPMC1	0000	010	SPM_RXU_IDLE	RXU0 is idle and RXU1 is idle.	FXU	0004
SPMC1	0000	001	SPM_RXU0_FIN	RXU0 finished.	FXU	0002
SPMC2	0011	011	SPM_RXU1_BUSY_RXU0_IDLE	RXU0 is idle and RXU1 is busy.	FXU	0036
SPMC2	0011	000	SPM_RXU1_FIN	RXU1 finished.	FXU	0030
SPMC1	0100	110	SPM_GCT_EMPTY_CYC	No itags are assigned to either thread (GCT empty).	ISU	004C
SPMC1	1110	100	SPM_GCT_NOSLOT_CYC	Pipeline empty (no itags assigned, no GCT slots used).	ISU	00E8
SPMC2	0000	000	SPM_GRP_CMPL	Group completed.	ISU	0000
SPMC2	0000	001	SPM_GRP_DISP	Dispatch success (group dispatched).	ISU	0002
SPMC1	0110	111	SPM_GRP_PUMP_CPRED	Initial and final pump scope and data sourced across this scope was group pump for all data types (demand load, instruction fetch, instruction or data translation).	LSU	006E
SPMC1	0111	101	SPM_GRP_PUMP_MPRED	The final pump scope (group) that is to be data sourced ended up larger than initial pump scope or final pump scope (group) received data from a source that was at a smaller scope (chip). Final pump was group pump and initial pump was chip or final and initial pump was group but data was sourced at chip scope level for all data types (demand load, instruction fetch, instruction or data translation).	LSU	007A
SPMC1	0011	001	SPM_GRP_PUMP_MPRED_RTY	The final pump scope (group) to be data sourced ended up larger than initial pump scope (chip). Final pump was group pump and initial pump was chip pump for all data types (demand load, instruction fetch, instruction or data translation).	LSU	0032
SPMC1	0000	011	SPM_IC_DEMAND_CYC	Demand ifetch pending.	IFU	0006
SPMC1	1110	011	SPM_IERAT_RELOAD	I-ERAT reloaded (miss).	IFU	00E6
SPMC1	1110	110	SPM_IERAT_RELOAD_16M	I-ERAT reload for 16 MB page.	IFU	00EC
SPMC1	1110	101	SPM_IERAT_RELOAD_4K	I-ERAT reload for 4 KB page.	IFU	00EA
SPMC2	0001	110	SPM_IERAT_RELOAD_64K	Instruction ERAT reload for 64 KB page.	PMU	001C
SPMC2	0110	101	SPM_IFU_FIN	IFU finished a (nonbranch) instruction.	IFU	006A
SPMC1	1111	001	SPM_INST_DISP	Number of PowerPC instructions dispatched.	ISU	00F2
SPMC2	1110	001	SPM_INST_DISP	Number of PowerPC instructions dispatched.	ISU	00E2
SPMC1	0001	111	SPM_INST_FROM_L2	The processor's instruction cache was reloaded from the local core's L2 cache due to an instruction fetch.	IFU	001E

Table E-1. SPMC Performance Monitoring Events (Sheet 5 of 7)

SPMC	SPMCxSEL(0:3) (Binary)	SPMCxSEL(4:6) (Binary)	Event Name	Description	Unit	Select Event (Hexadecimal)
SPMC1	0010	001	SPM_INST_FROM_L2MISS	The processor's instruction cache was reloaded from a location other than the local core's L2 cache due to an instruction fetch.	IFU	0022
SPMC2	0100	100	SPM_INST_FROM_L3	The processor's instruction cache was reloaded from the local core's L3 cache due to an instruction fetch.	IFU	0048
SPMC2	1110	101	SPM_INST_FROM_L3MISS	Instruction from an L3 miss.	IFU	00EA
SPMC1	0110	001	SPM_INST_FROM_LMEM	The processor's instruction cache was reloaded from the local chip's memory due to an instruction fetch.	IFU	0062
SPMC1	0110	010	SPM_INST_FROM_MEMORY	The processor's instruction cache was reloaded from a memory location due to an instruction fetch.	IFU	0064
SPMC2	0100	101	SPM_INST_FROM_OFF_CHIP_CACHE	The processor's instruction cache was reloaded with either shared or modified data from another core's L2/L3 cache on a different chip (remote or distant) due to an instruction fetch.	IFU	004A
SPMC1	0010	000	SPM_INST_FROM_ON_CHIP_CACHE	The processor's instruction cache was reloaded with either shared or modified data from another core's L2 or L3 cache on the same chip due to an instruction fetch.	IFU	0020
SPMC1	0010	010	SPM_IPTEG_FROM_L2	A page table entry was loaded into the TLB from the local core's L2 cache due to an instruction-side request.	IFU	0024
SPMC1	0010	100	SPM_IPTEG_FROM_L2MISS	A page table entry was loaded into the TLB from a location other than the local core's L2 cache due to an instruction-side request.	IFU	0028
SPMC2	0100	111	SPM_IPTEG_FROM_L3	A page table entry was loaded into the TLB from the local core's L3 cache due to an instruction-side request.	IFU	004E
SPMC2	0101	001	SPM_IPTEG_FROM_L3MISS	A page table entry was loaded into the TLB from a location other than the local core's L3 cache due to an instruction-side request.	IFU	0052
SPMC1	0110	011	SPM_IPTEG_FROM_LMEM	A page table entry was loaded into the TLB from the local chip's memory due to an instruction-side request.	IFU	0066
SPMC1	0110	100	SPM_IPTEG_FROM_MEMORY	A page table entry was loaded into the TLB from a memory location due to an instruction-side request.	IFU	0068
SPMC2	0101	000	SPM_IPTEG_FROM_OFF_CHIP_CACHE	A page table entry was loaded into the TLB with either shared or modified data from another core's L2/L3 cache on a different chip (remote or distant) due to an instruction-side request.	IFU	0050
SPMC1	0010	011	SPM_IPTEG_FROM_ON_CHIP_CACHE	A page table entry was loaded into the TLB with either shared or modified data from another core's L2 or L3 cache on the same chip due to an instruction-side request.	IFU	0026
SPMC2	0110	110	SPM_ISLB_MISS	Instruction SLB (ISLB) miss.	ISU	006C
SPMC2	1111	110	SPM_ITLB_MISS	ITLB reloaded.	LSU	00FC
SPMC2	1110	011	SPM_L1_DCACHE_RELOAD_VALID	DL1 reloaded due to a demand load.	LSU	00E6
SPMC1	0000	110	SPM_L1_DCACHE_RELOADED_ALL	L1 data cache reloaded for demand or prefetch	LSU	000C
SPMC1	1111	110	SPM_L1_ICACHE_MISS	Demand I-cache miss.	IFU	00FD
SPMC1	0111	110	SPM_L3_CO_MEPF	Total number of prefetch operations that were tracked.	LSU	007C





Table E-1. SPMC Performance Monitoring Events (Sheet 6 of 7)

SPMC	SPMCxSEL(0:3) (Binary)	SPMCxSEL(4:6) (Binary)	Event Name	Description	Unit	Select Event (Hexadecimal)
SPMC2	0010	010	SPM_LARX_FIN	Larx finished.	LSU	0024
SPMC2	0100	011	SPM_LD_CMPL	Loads completed.	ISU	0046
SPMC1	0001	110	SPM_LD_L3MISS_PEND_CYC	L3 miss pending.	LSU	001C
SPMC2	1111	000	SPM_LD_MISS_L1	Load missed L1 cache.	LSU	00F0
SPMC1	1111	011	SPM LSU_DERAT_MISS	D-ERAT reloaded (miss).	LSU	00F6
SPMC2	0010	110	SPM LSU_FIN	LSU finished an instruction (up to two per cycle).	LSU	002C
SPMC1	0100	011	SPM LSU_FX_FIN	LSU finished an FX operation (up to two per cycle).	LSU	0046
SPMC2	0000	100	SPM LSU_LMQ_SRQ_EMPTY_ALL_CYC	All threads in the LSU are empty (LMQ and SRQ are empty).	ISU	0008
SPMC1	0101	001	SPM LSU_LMQ_SRQ_EMPTY_CYC	LSU empty (LMQ and SRQ empty).	ISU	0052
SPMC1	0100	010	SPM LSU_REJECT	LSU reject (up to four per cycle).	LSU	0044
SPMC2	0110	010	SPM LSU_REJECT_LHS	LSU reject due to LHS (up to four per cycle).	LSU	0064
SPMC1	0000	100	SPM LSU_SRQ_FULL_CYC	SRQ is full.	ISU	0008
SPMC2	0110	011	SPM_MEM_CO	Castouts to memory from this LPAR.	L2	0066
SPMC1	0011	100	SPM_MEM_PREF	Prefetch tracked was moderate (source L2 cache).	PMU	0038
SPMC1	0011	011	SPM_MEM_READ	Prefetch tracked was effective.	PMU	0036
SPMC2	0110	100	SPM_MEM_RWITM	RWITM to memory from this LPAR.	L2	0068
SPMC1	0011	010	SPM_PUMP_CPRED	Pump prediction correct. Counts across all types of pumps for all data types (demand load, instruction fetch, instruction or data translation).	LSU	0034
SPMC2	0101	110	SPM_PUMP_MPRED	Pump misprediction. Counts across all types of pumps for all data types (demand load, instruction fetch, instruction or data translation).	PMU	005C
SPMC1	1111	010	SPM_RUN_CYC	Run cycles.	ISU	00F4
SPMC1	0100	001	SPM_RUN_CYC_SMT2_MODE	Number of cycles that this thread was in SMT mode.	PMU	0042
SPMC1	0111	011	SPM_RUN_CYC_SMT2_SHRD_MODE	Number of cycles that this thread was in SMT-shared mode.	PMU	0076
SPMC1	0111	010	SPM_RUN_CYC_SMT2_SPLIT_MODE	Number of cycles that this thread was in SMT-split mode.	PMU	0074
SPMC2	0010	100	SPM_RUN_CYC_SMT4_MODE	Number of cycles that this thread is in SMT4 shared mode.	PMU	0028
SPMC2	0010	101	SPM_RUN_CYC_SMT8_MODE	Number of cycles that this thread is in SMT8 shared mode.	PMU	002A
SPMC1	0100	000	SPM_RUN_CYC_ST_MODE	Number of cycles that this thread was in ST mode.	PMU	0040
SPMC2	1111	101	SPM_RUN_INST_CMPL	Run instructions.	ISU	00FA
SPMC2	1111	010	SPM_RUN_PURR	Run PURR.	PC	00F4
SPMC2	0010	111	SPM_ST_CMPL	Store completion count.	LSU	002E



Table E-1. SPMC Performance Monitoring Events (Sheet 7 of 7)

SPMC	SPMCxSEL(0:3) (Binary)	SPMCxSEL(4:6) (Binary)	Event Name	Description	Unit	Select Event (Hexadecimal)
SPMC1	1111	000	SPM_ST_FIN	Store instructions finished (store sent to nest).	LSU	00F0
SPMC1	0111	111	SPM_ST_FWD	Store forwarding occurred.	LSU	007E
SPMC2	1110	000	SPM_ST_MISS_L1	Store missed the L1 cache.	LSU	00E0
SPMC1	0011	111	SPM_STCX_FAIL	Stcx failed.	LSU	003E
SPMC1	0000	000	SPM_SUSPENDED	Counter off.	PMU	0000
SPMC2	0001	100	SPM_SYS_PUMP_CPRED	Initial and final pump scope and data sourced across this scope was system pump for all data types (demand load, instruction fetch, instruction or data translation).	PMU	0018
SPMC2	0001	101	SPM_SYS_PUMP_MPRED	The final pump scope (system) that is to get data sourced, ended up larger than the initial pump scope (chip/group) or the final pump scope (system) received data from a source that was at a smaller scope (chip/group). Final pump was system pump and initial pump was chip or group or final and initial pump were system but data was sourced at chip/group scope level for all data types (demand load, instruction fetch, instruction or data translation).	PMU	001A
SPMC2	0101	101	SPM_SYS_PUMP_MPRED_RTY	The final pump scope (system) that is to get data sourced, ended up larger than initial pump scope (chip or group) for all data types (demand load, instruction fetch, instruction or data translation).	PMU	005A
SPMC1	0000	101	SPM_TABLEWALK_CYC	Tablewalk active.	LSU	000A
SPMC2	1110	100	SPM_TB_BIT_TRANS	Timebase event.	PC	00E8
SPMC1	0100	111	SPM_THRD_ALL_RUN_CYC	All threads in run_cycles.	ISU	004E
SPMC2	1110	010	SPM_THRD_CONC_RUN_INST	Concurrent run instructions.	ISU	00E4
SPMC1	0111	100	SPM_TLB_MISS	TLB miss (instruction and data).	LSU	0078



Glossary

ABIST	Array built-in self test
AES	Advanced Encryption Standard
AMC	Architected mapper cache
ARF	Architected register file
ASIC	Application-specific integrated circuit
BA	Base address
BCD	Binary coded decimal
BFP	Binary floating-point
BFU	Binary floating-point unit
BHT	Branch history table
BIST	Built-in self-test
BMC	Baseboard management control
BR	Branch register unit
BTAC	Branch target address cache
CAI	Coherent accelerator
CAM	Content-addressable memory
CEC	Central electronics complex
CI	Cast-in
CIABR	Current Instruction Address Breakpoint Register
CIR	Chip information register
CIU	Core interface unit
CLB	Cache load buffer (IBuffer)
CMOS	Complementary metal–oxide–semiconductor
CO	Cast-out
CPI	Cycles per instruction
CPU	Central processing unit
CR	Condition Register
CRC	Cyclic redundancy check

Cresp	Combined response
CS	Continuous sampling
DAWR	Data Address Watch Register
DDR	Double data rate
DECFP	Decimal floating-point unit
DFP	Decimal floating-point
DFU	Decimal floating-point unit
DIMM	Dual in-line memory module
DMA	Direct memory attach
DMI	Differential memory interface
DP	Double precision
DPD	Densely packed decimal
DPLL	Digital phase-locked loop
DTS	Digital thermal sensor
EADIR	Effective address directory
EAT	Effective address translation
ECC	Error correcting code
ECID	Electronic chip identification
ECO	Extended cache option
ECRC	End-to-end cyclic redundancy check
EDI	Elastic differential I/O
EEH	Enhance error handling
EEPROM	Electrically erasable programmable read-only memory
EMQ	ERAT miss queue
ERAT	Effective-to-real address translation
ESID	Effective segment identifier
FBC	SMP interconnect controller
FCPBGA	Flip-chip plastic ball grid array
FIFO	First-in, first-out

FIR	Fault Isolation Register
FLOPs	Floating-point operations per second
FPR	Floating-point register
FPSCR	Floating-Point Status and Control Register
FPU	Floating-point unit
FXU	Fixed-point units
GCM	Galios counter mode
GCT	Global completion table
GFW	Global firmware
GHV	Global history vector
GPE	General purpose engine
GPU	Graphics processor unit
GPR	General purpose register
GPS	Global Pstate
GPST	Global Pstates table
HCSL	Host clock signal level
HDEC	Hypervisor decrementer
HMER	Hypervisor Maintenance Exception Register
HMEER	Hypervisor Maintenance Exception Enable Register
HMI	Hypervisor maintenance interrupt
HRMOR	Hypervisor Real Mode Offset Register
HSS	High Speed Serial
HTM	Hardware trace monitor
ICA	Instruction cache access
ICP	Interrupt control presenter
ICS	Interrupt controller source
IEEE	Institute of Electrical and Electronics Engineers
IFAR	Instruction fetch address register
IFM	Instruction filtering mode

IFU	Instruction fetch and decode unit
IMA	In memory accumulate
IOP	Internal operation
IPC	Instruction per cycle
IPL	Interrupt presenter layer; or initial program load
IRL	Interrupt routing layer
ISU	Instruction sequencing unit
JIT	Just in time
LBIST	Logic built-in self test
LHR	Load-hit-reload
LMQ	Load miss queue
LPAR	Logical partition
LPID	Logic partition ID
LPST	Local Pstate table
LRDIMM	Load-reduced dual in-line memory module
LRQ	Load reorder queue
LRU	Least-recently used
LS	Link stack
LSI	Level signalled interrupt
LSSD	Level-sensitive scan design
LSU	Load store units
LU	Load-only unit
MBA	Memory buffer asynchronous
MCI	Memory channel interface
MCS	Memory controller synchronous
MDS	Memory domain status
MHCRO	Model hardware correlation ring oscillator
MPG	Multi-protocol gateway
MMCRA	Monitor Mode Control Register A

MMCR0	Monitor Mode Control Register 0
MMCR1	Monitor Mode Control Register 1
MMCR2	Monitor Mode Control Register 2
MMCRH	Monitor Mode Control Register H
MRU	Most-recently used
MSI	Message signalled interrupt
NaN	Not a number
NCU	Noncacheable unit
NPS	Nap Pstate
NTC	Next-to-complete
NTF	Next-to-finish
OCC	On-chip controller
OCTS	On-chip thermal sensor
OEM	Original equipment manufacturer
OHA	On-chiplet hardware assist
PAPR	Power Architecture Platform Reference
PB	Processor bus
PC	Pervasive core unit
PCR	Processor Compatibility Register
PE	Partitionable endpoints
PEC	PCI Express controller
PID	Process ID
PLL	Phase-locked loop
PMC	Performance monitor counter or Power Management Control Register
POR	Power-on reset
PRQ	Prefetch request queue
PSN	Power state number
PSPB	Problem-state priority boost
PSRO	Performance sort-ring oscillator

Pstate	Performance state
PTEG	Page table entry group
PURR	Processor Utilization Resource Register
PVR	Processor Version Register
QNaN	Quiet Not a number
qpos	Queue position
RAIM	Redundant array of independent memory
RAS	Reliability, availability, and serviceability
RAW	Read after write
RC	Root complex or read/claim
RCD	Register clock driver
RDIMM	Registered dual in-line memory module
RES	Random event sampling
RIS	Random instruction sampling
RMSC	Real mode storage control
RNG	Random number generator
ROB	Re-order buffer
SAO	Strong access ordering
SAR	Second-level Architected Register
SBE	Self-boot engine
SC	Store clean
SCM	Single-chip module
SCOM	Scan communications
SDAR	Sampled Data Address Register
SDQ	Store data queue
SECDED	Single-error correction, double-error detection
SEEPROM	Serial electrically erasable programmable read-only memory
SHA	Secure hash algorithm
SHR	Store-hit-reload

SIAR	Sampled Instruction Address Register
SIER	Sampled Instruction Event Register
SIMD	Single-instruction, multiple-data
SLB	Segment lookaside buffer
SMP	Symmetric multiprocessing
SMT	Simultaneous multithreading
SOI	Silicon-on-insulator
SP	Single precision
SPI	Serial peripheral interconnect
SPIVID	Serial Peripheral Interface - Voltage ID
SPMC	Supervisor-level performance monitor counter
SPR	Special purpose register
SPS	Sleep Pstate
SPURR	Scaled Processor Utilization Resource Register
SRAM	Static random access memory
SRQ	Store reorder queue
ST	Single thread
STAG	Storage tag
SVIC	Slave VME interface controller
TCE	Translation control entry
TDP	Thermal design point
TEXASR	Transaction Exception And Summary Register
TEXASRU	Transaction Exception And Summary Register Upper
TID	Thread ID
TLB	Translation lookaside buffer
TLP	Translation layer packet
TM	Transactional memory
TOD	Time of day
UE	Uncorrectable error



UniQ	Unified issue queues
VLE	Variable length encoding
VMX	Virtual machine extensions
VPD	Vital product data
VPN	Virtual page number
VRF	Vector scalar register file
VRM	Voltage regulator module
VRMA	Virtualized real mode area
VS	Vector scalar
VSCR	Vector Status and Control Register
VSU	Vector and scalar unit
VSX	Vector-scalar extension
WAW	Write after write
WI	Write inject
WPS	Winkle Pstate
XER	Fixed-Point Exception Register